


Reliability-Aware Scheduling on Heterogeneous Multicore Processors

Ajeya Naithani, Stijn Eyerman, and Lieven Eeckhout

Benefits of Heterogeneous Multicores


Benefits of Heterogeneous Multicores

1. Multiple core types



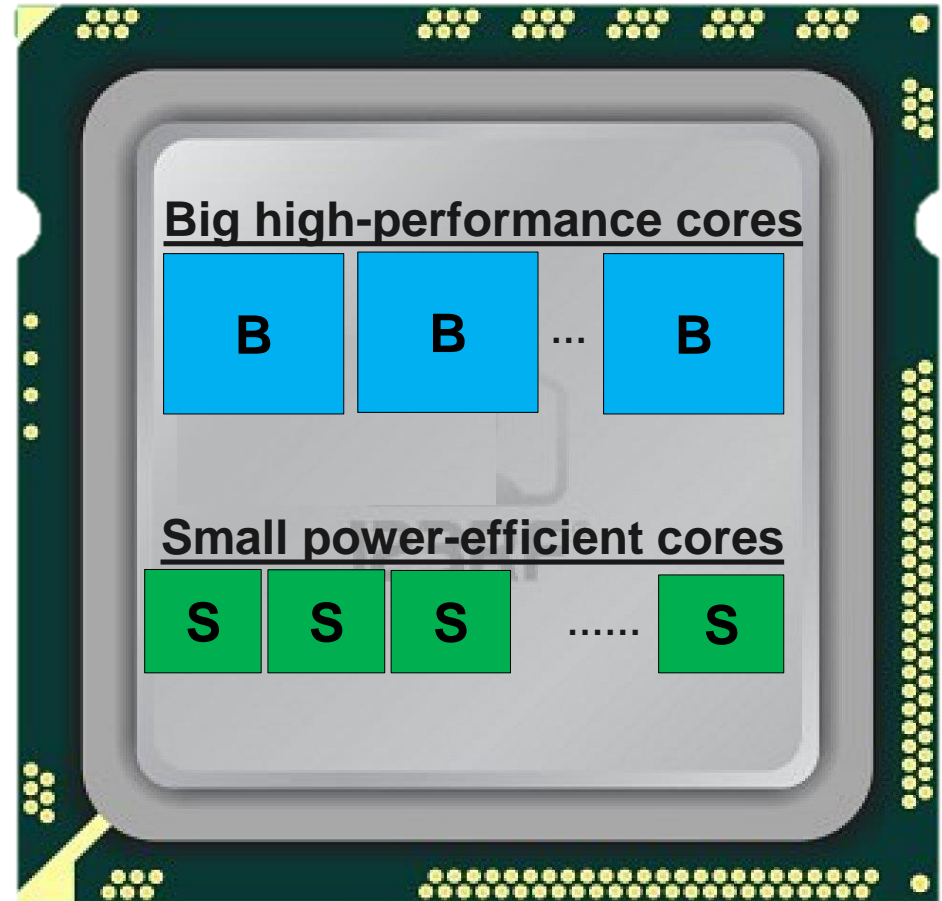
Benefits of Heterogeneous Multicores

1. Multiple core types




Benefits of Heterogeneous Multicores

1. Multiple core types



Benefits of Heterogeneous Multicores

1. Multiple core types
2. Well established **power benefits***




*[Kumar et al. MICRO'03, ISCA'04]

Benefits of Heterogeneous Multicores

1. Multiple core types
2. Well established **power benefits***
3. Well established **scheduling techniques****

e.g. Big.LITTLE and Kal-EI




*[Kumar et al. MICRO'03, ISCA'04]

**[Van Craeynest et al. ISCA'12] 2

Benefits of Heterogeneous Multicores

1. Multiple core types
2. Well established power benefits*

3. e.g. Big.LITTLE and Kal-EI

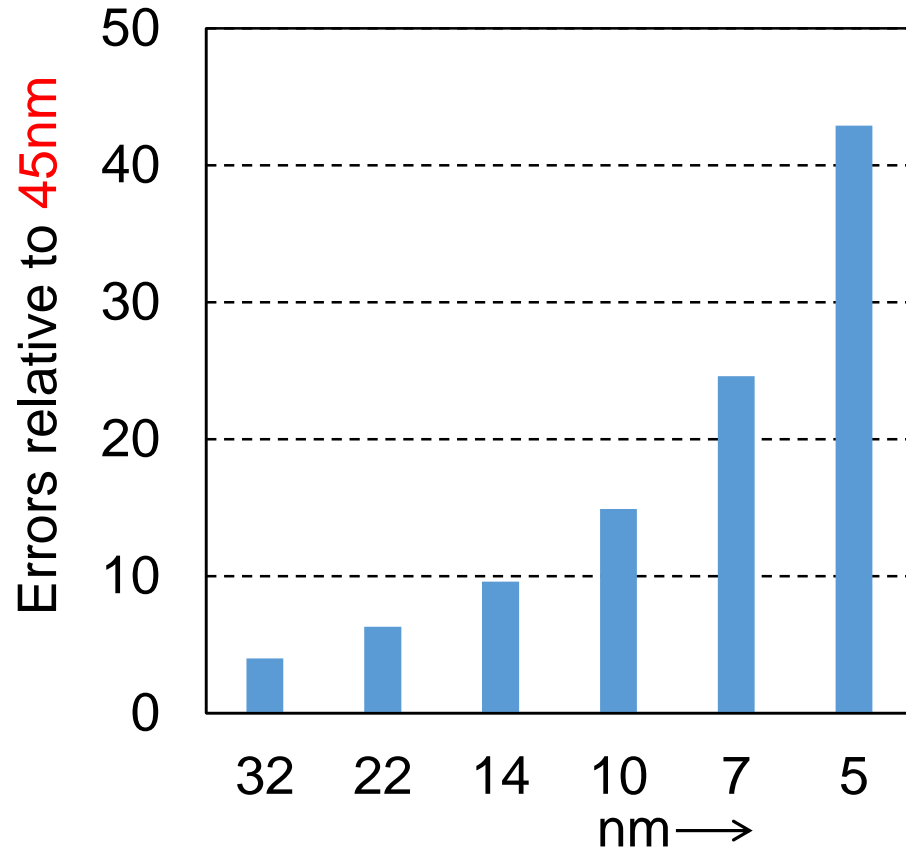


No prior work → How heterogeneous chip-multiprocessors affect reliability?

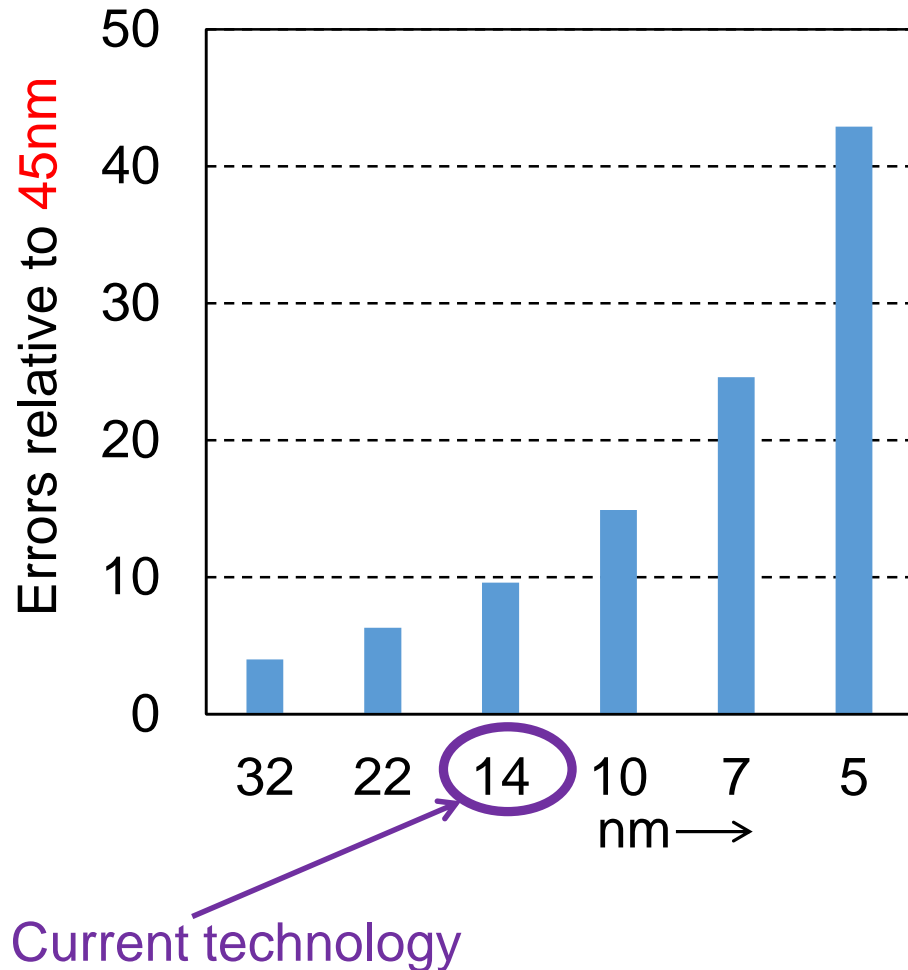
*[Kumar et al. MICRO'03, ISCA'04]

**[Van Craeynest et al. ISCA'12] 2

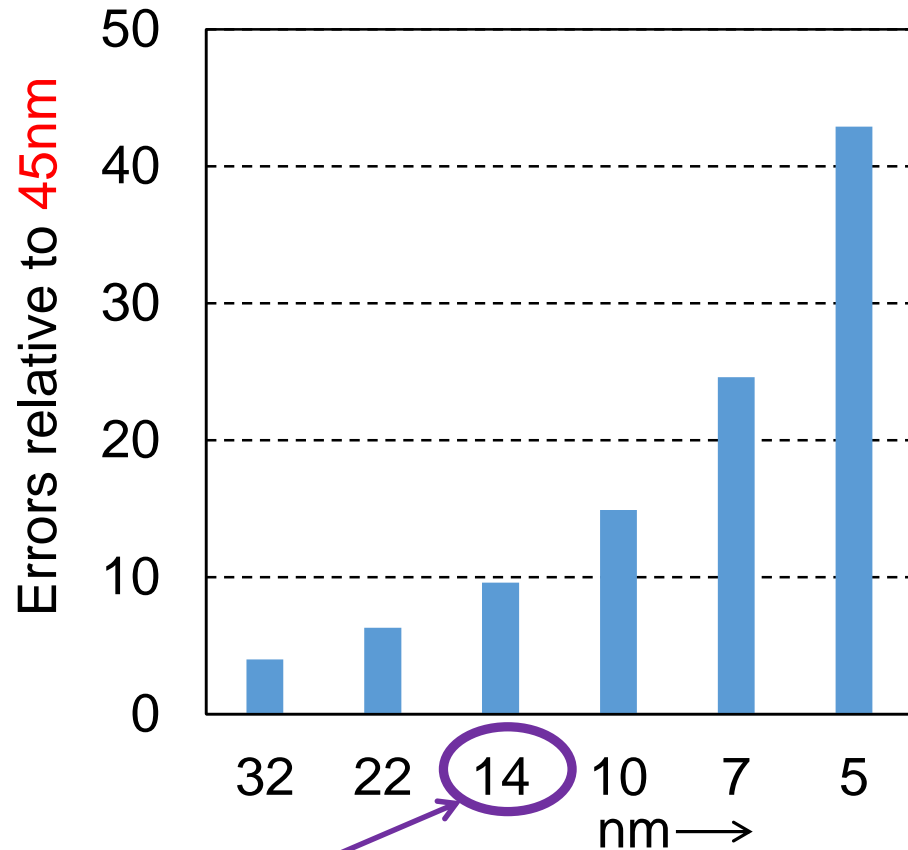
Errors vs. Technology




Errors vs. Technology



Errors vs. Technology




Current technology




SER memory SER logic Variability Aging

Errors vs. Technology

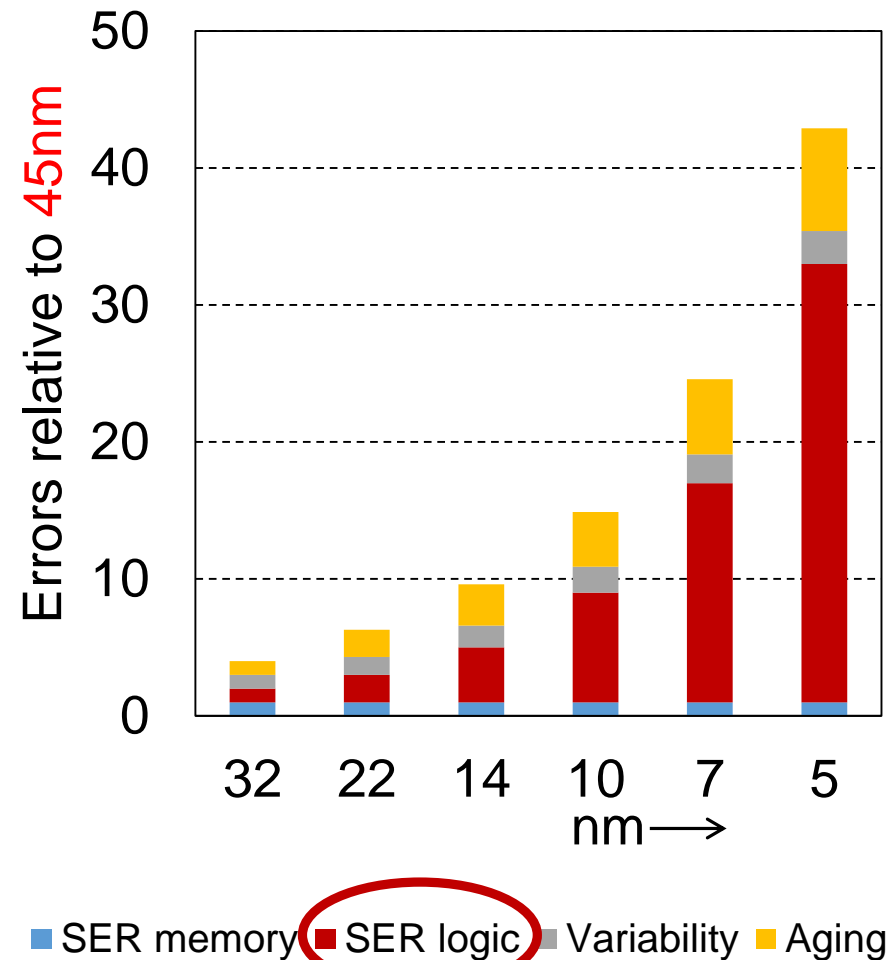



Current technology




SER memory SER logic Variability Aging

Errors vs. Technology




Soft errors are the major source of errors!

Soft Errors vs. Core Type



Soft Errors vs. Core Type



Key Contributions

Key Contributions

1. Analyze the difference in reliability characteristics between **big** and **small** cores

Key Contributions

1. Analyze the difference in reliability characteristics between **big** and **small** cores
2. Show the potential for improving reliability through scheduling

Key Contributions

1. Analyze the difference in reliability characteristics between **big** and **small** cores
2. Show the potential for improving reliability through scheduling
3. Introduce a **novel metric** called **System-level Soft Error Rate (SSER)** for multiprogram workloads on (heterogeneous) multicore processors

Key Contributions

1. Analyze the difference in reliability characteristics between **big** and **small** cores
2. Show the potential for improving reliability through scheduling
3. Introduce a **novel metric** called **System-level Soft Error Rate (SSER)** for multiprogram workloads on (heterogeneous) multicore processors
4. Propose a **dynamic scheduler** to optimize reliability
 - Scheduler improves reliability by **25.4%**

Talk Outline

Talk Outline

1. Background and Terminology

Talk Outline

1. Background and Terminology
2. Analyzing Reliability on Heterogeneous Multicores

Talk Outline

1. Background and Terminology
2. Analyzing Reliability on Heterogeneous Multicores
3. Reliability Metric for Multiprogram Workloads

Talk Outline

1. Background and Terminology
2. Analyzing Reliability on Heterogeneous Multicores
3. Reliability Metric for Multiprogram Workloads
4. Reliability-Aware Scheduler

Talk Outline

1. Background and Terminology
2. Analyzing Reliability on Heterogeneous Multicores
3. Reliability Metric for Multiprogram Workloads
4. Reliability-Aware Scheduler
5. Evaluation

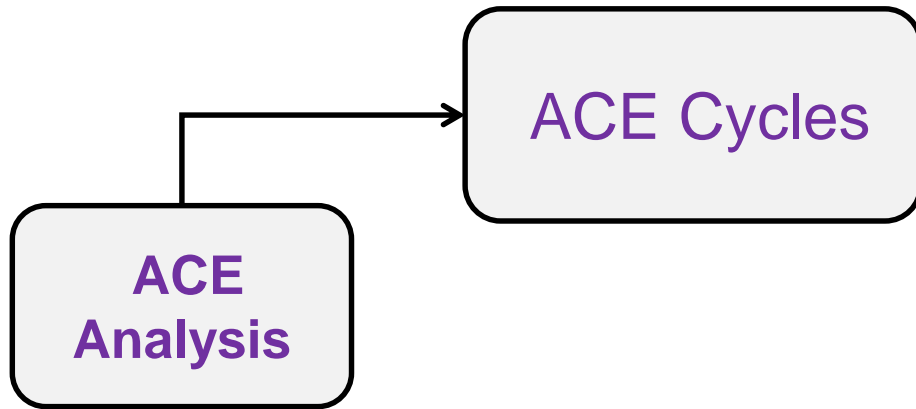
Talk Outline

1. Background and Terminology
2. Analyzing Reliability on Heterogeneous Multicores
3. Reliability Metric for Multiprogram Workloads
4. Reliability-Aware Scheduler
5. Evaluation
6. Conclusions


Estimating Soft Errors

**ACE
Analysis**

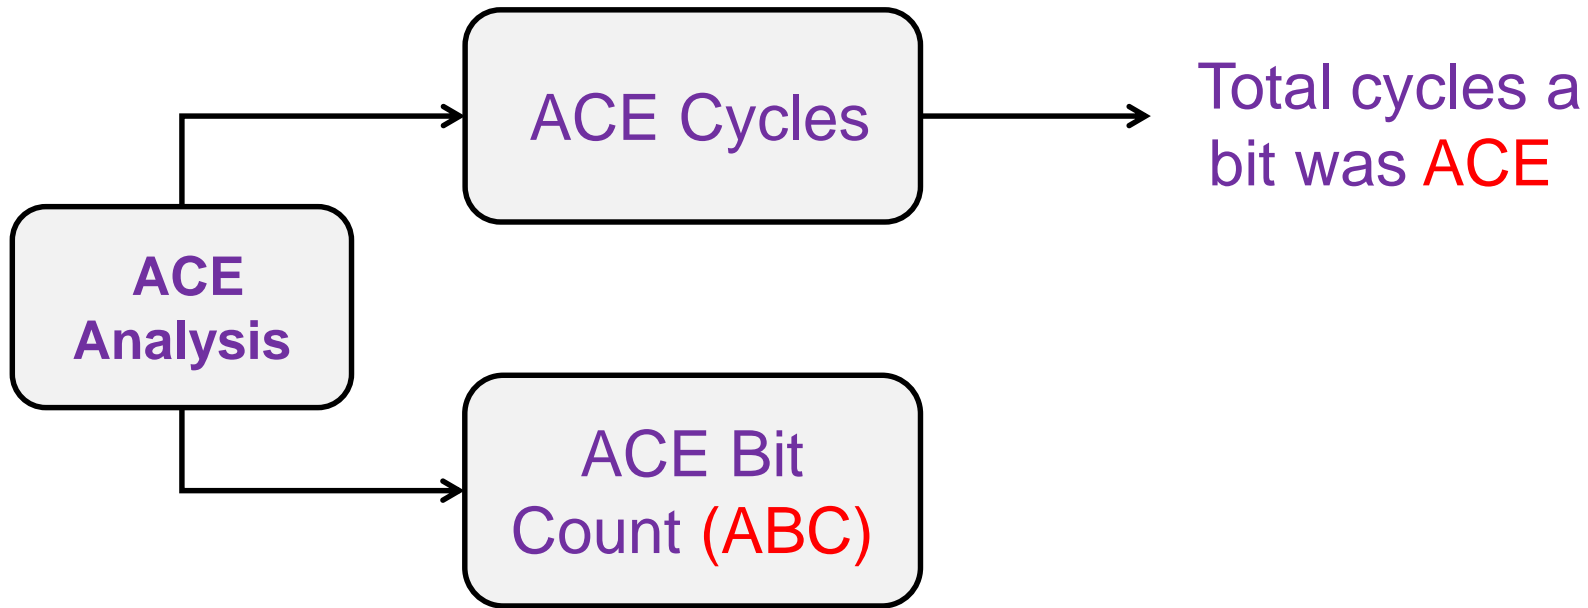
Estimating Soft Errors



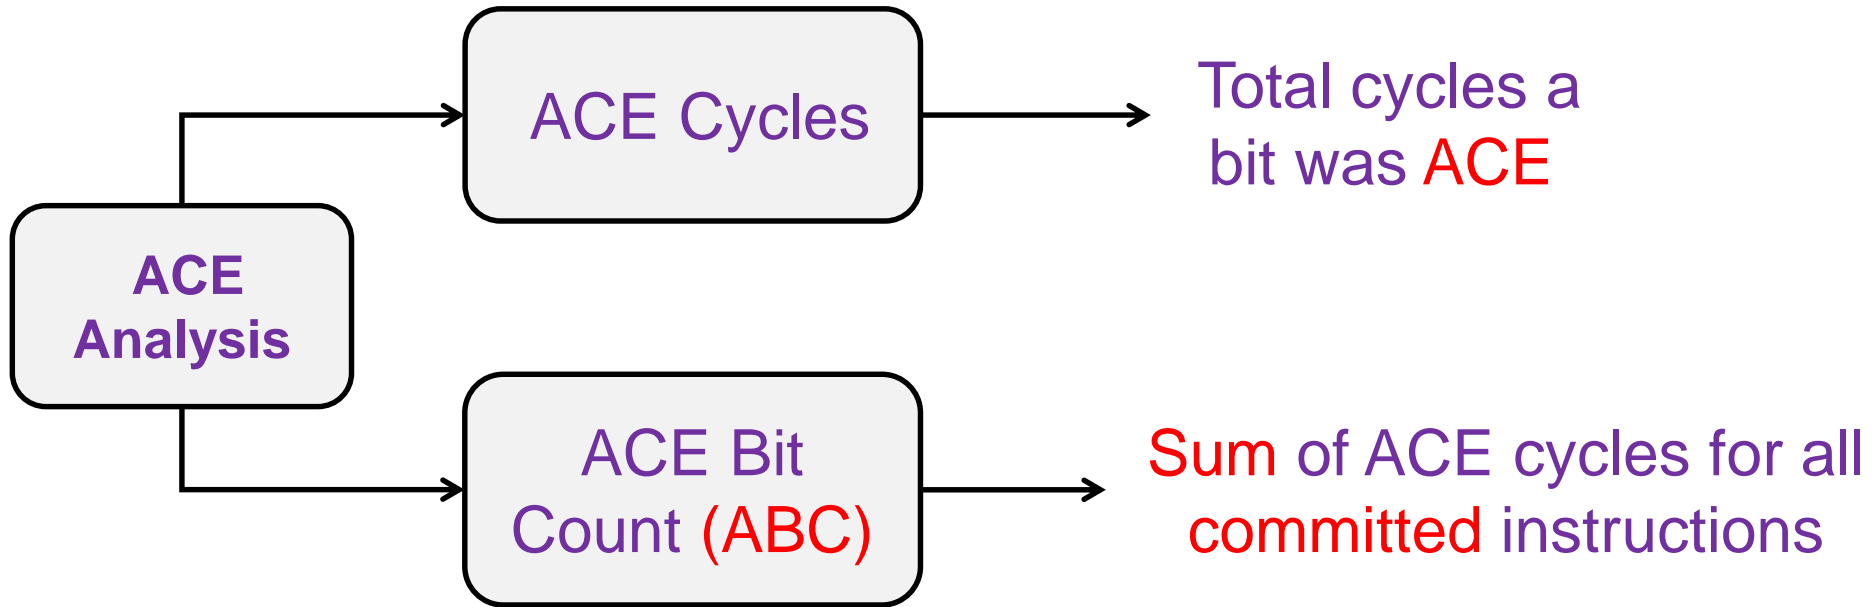
Estimating Soft Errors



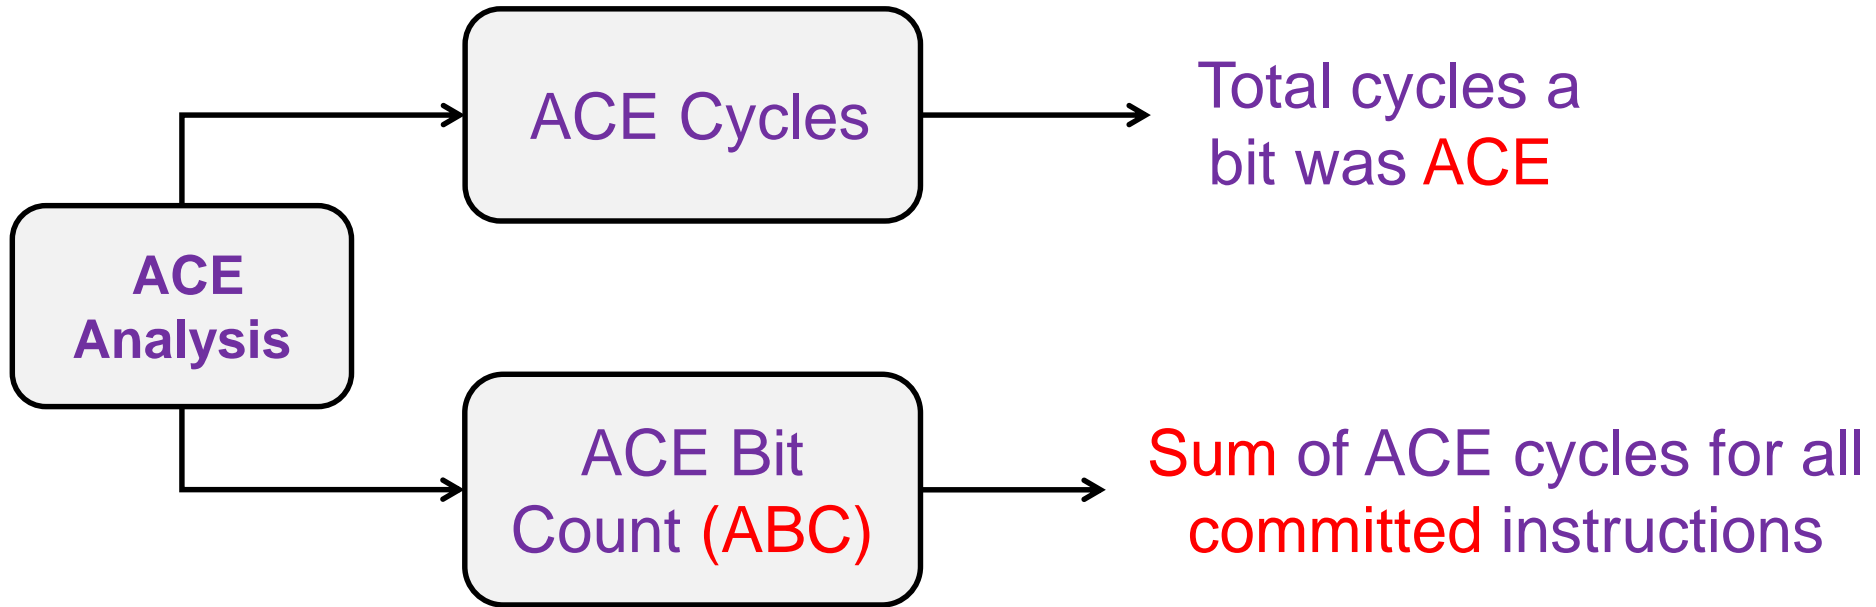
Estimating Soft Errors



Estimating Soft Errors



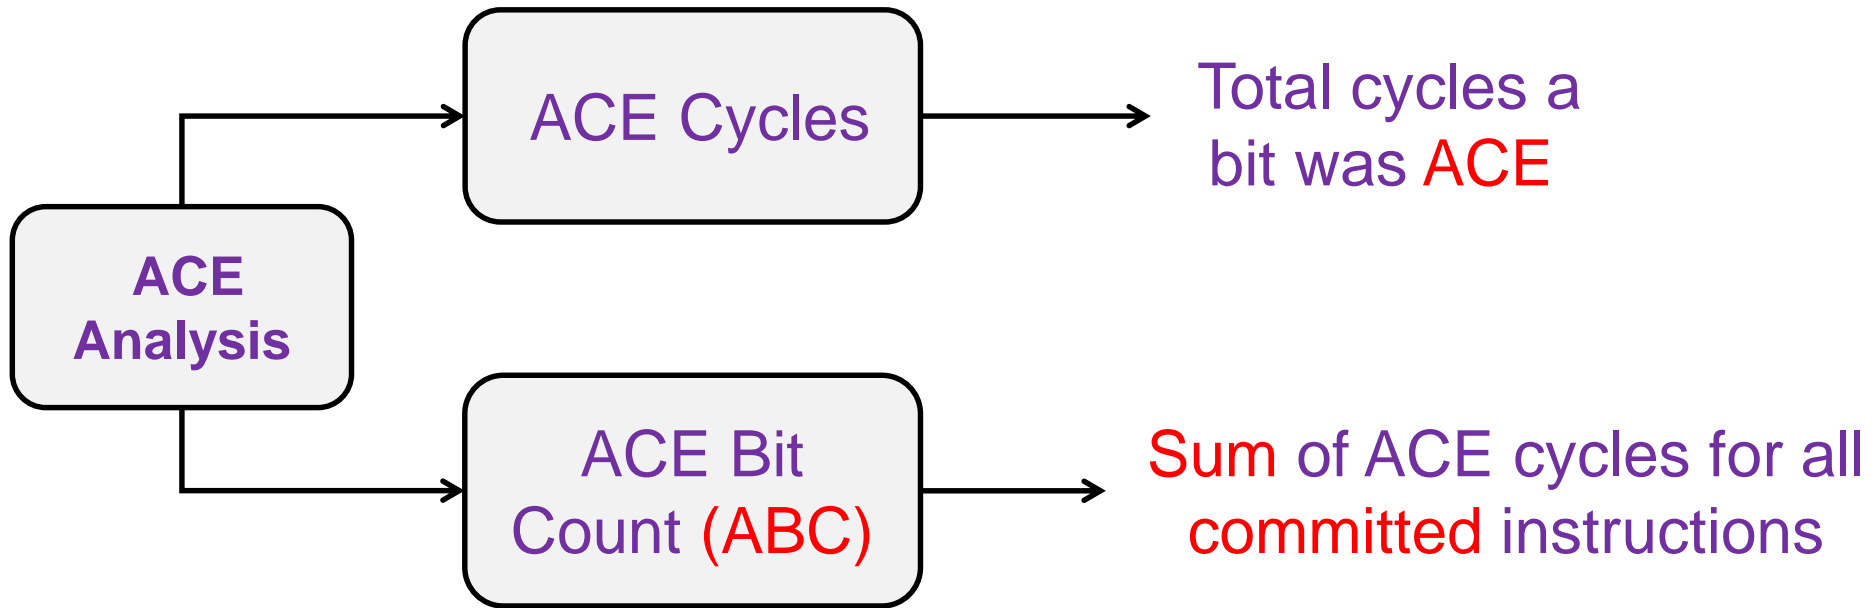
Estimating Soft Errors



$$\text{Soft Error Rate (SER)} = \frac{\text{ABC}}{\text{Total cycles}} \times \text{IFR}$$

IFR → Probability
of a soft error per
second

Estimating Soft Errors




$$\text{Soft Error Rate (SER)} = \frac{ABC}{\text{Total cycles}} \times \text{IFR}$$


IFR → Probability
of a soft error per
second

$$\text{Architectural Vulnerability Factor (AVF)} = \frac{ABC}{\text{Total cycles} \times \text{Core size}}$$

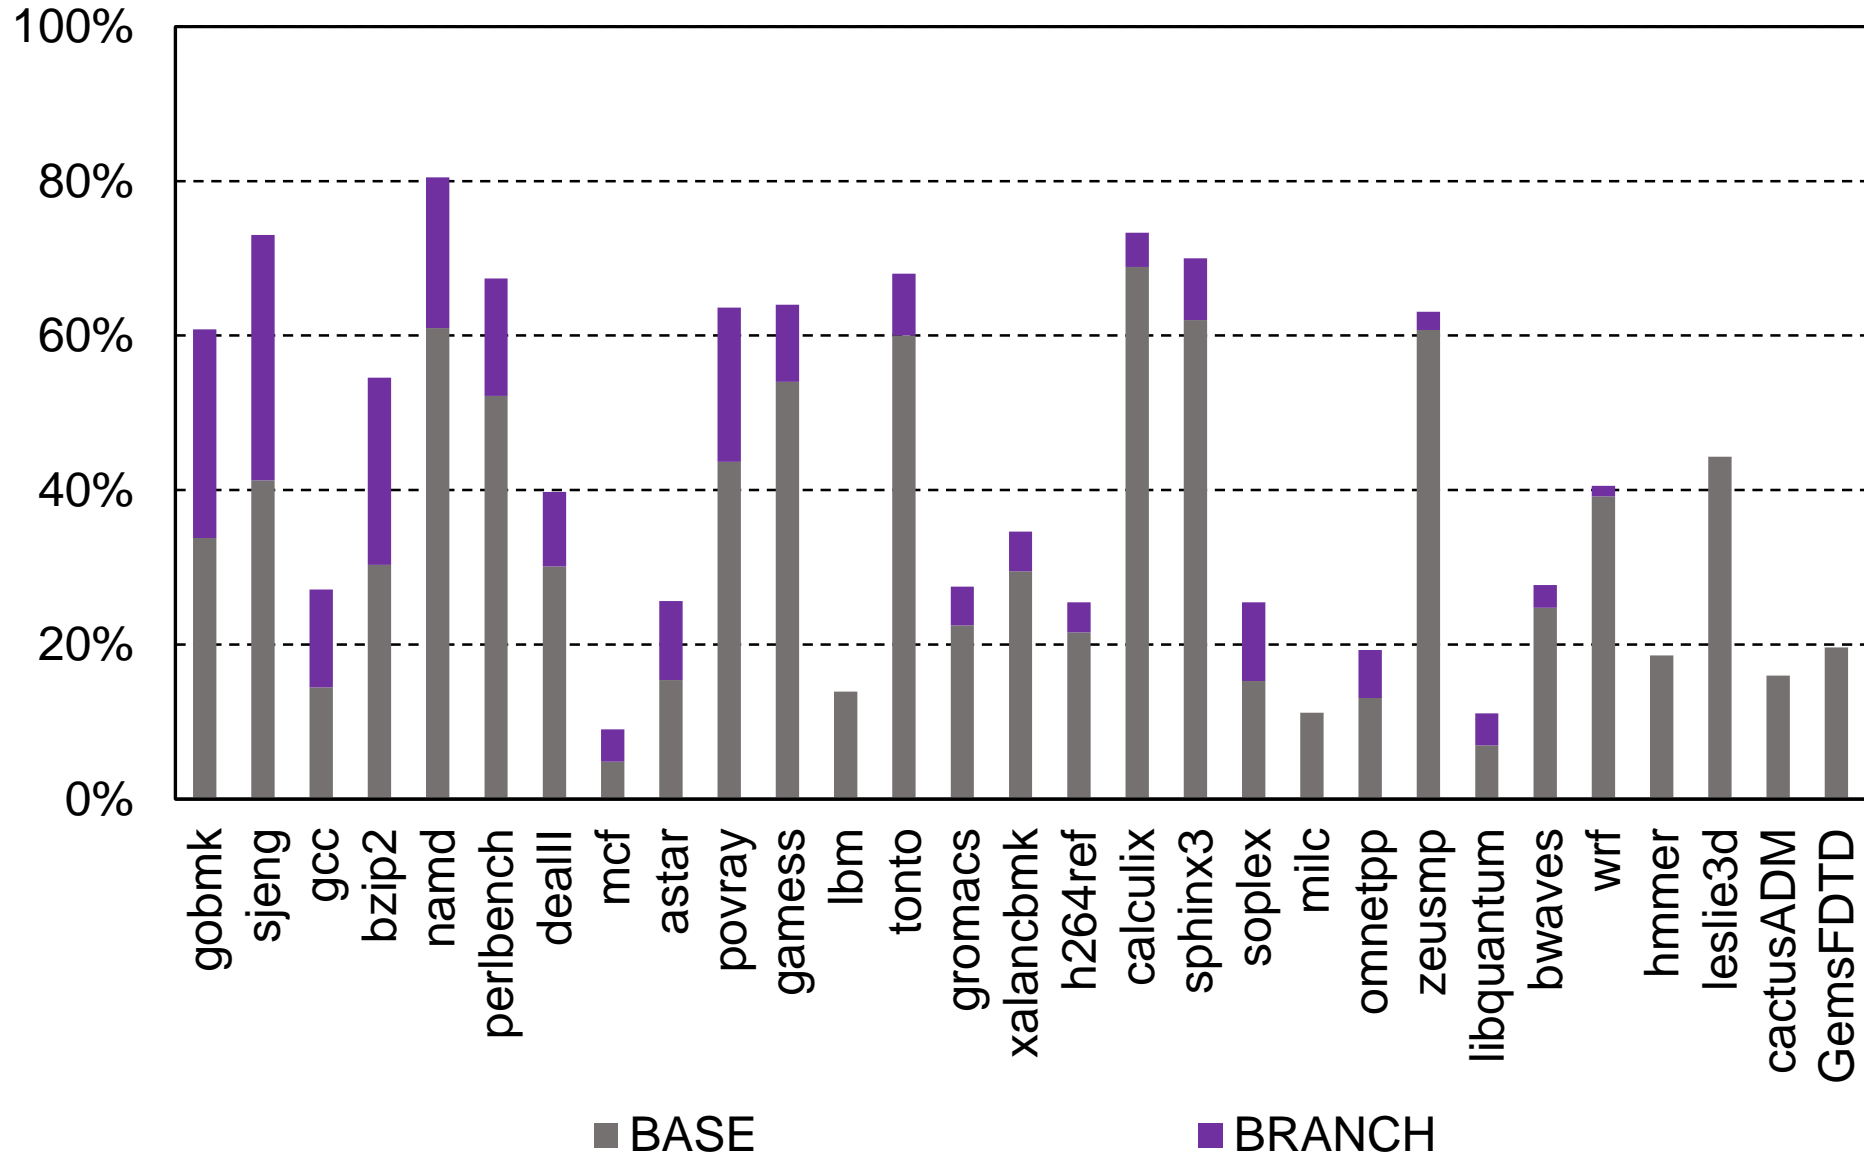
Vulnerability vs. Core Type



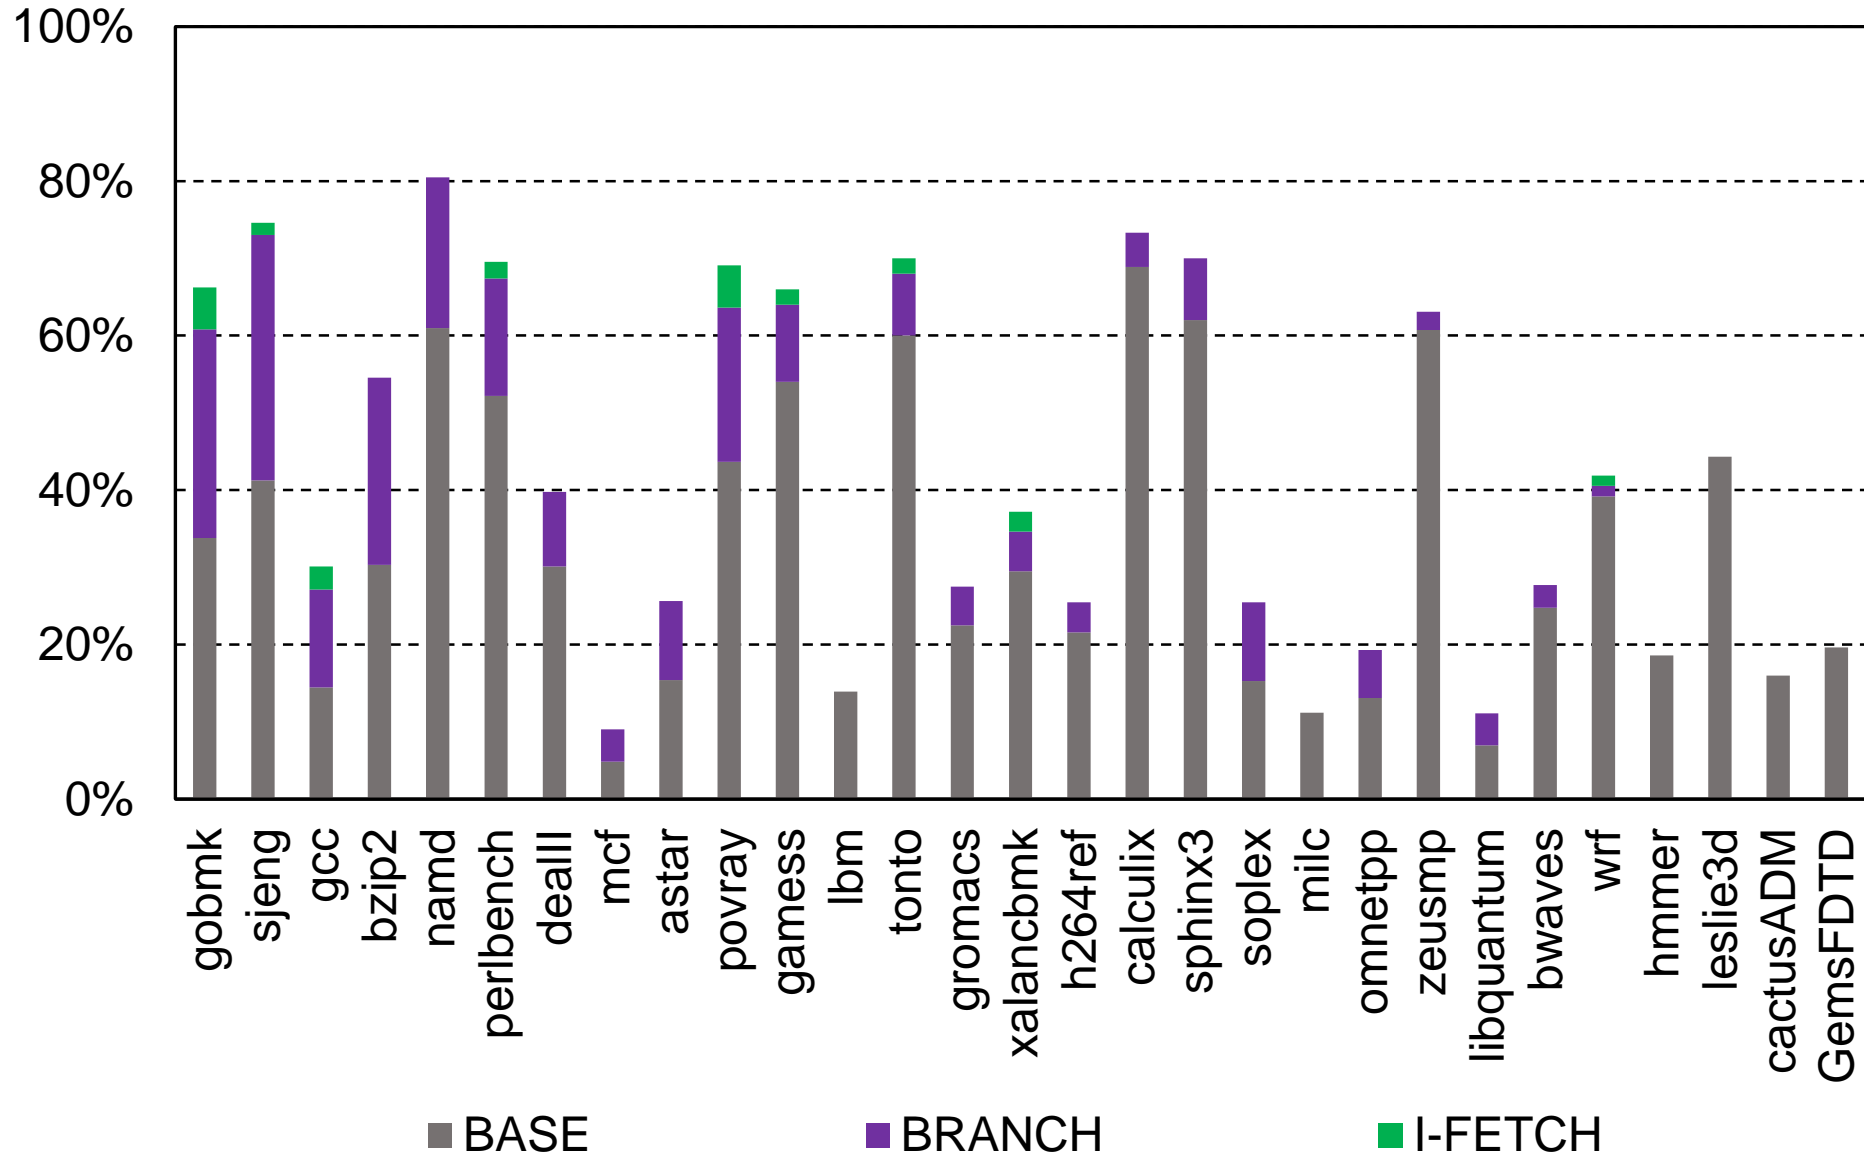
Vulnerability vs. CPI Stacks




Vulnerability vs. CPI Stacks



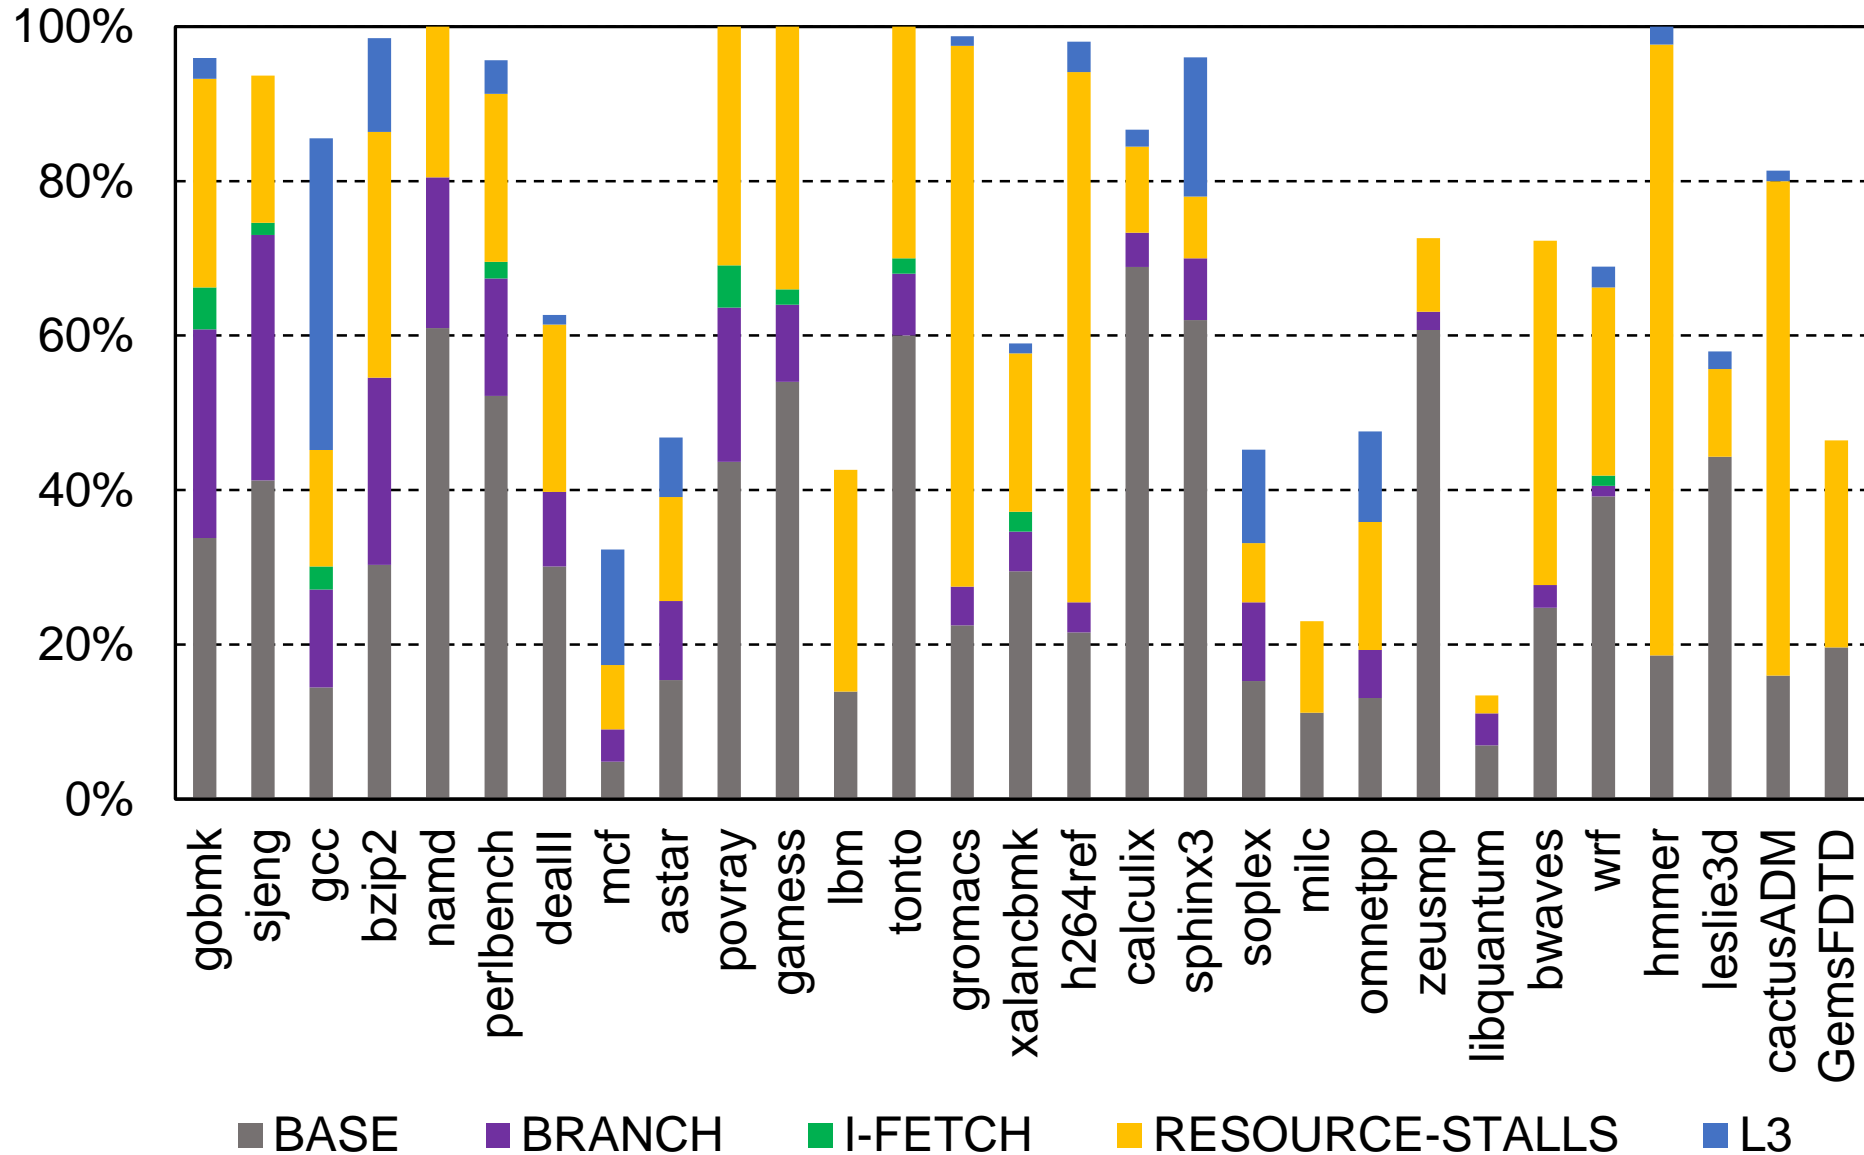
Vulnerability vs. CPI Stacks



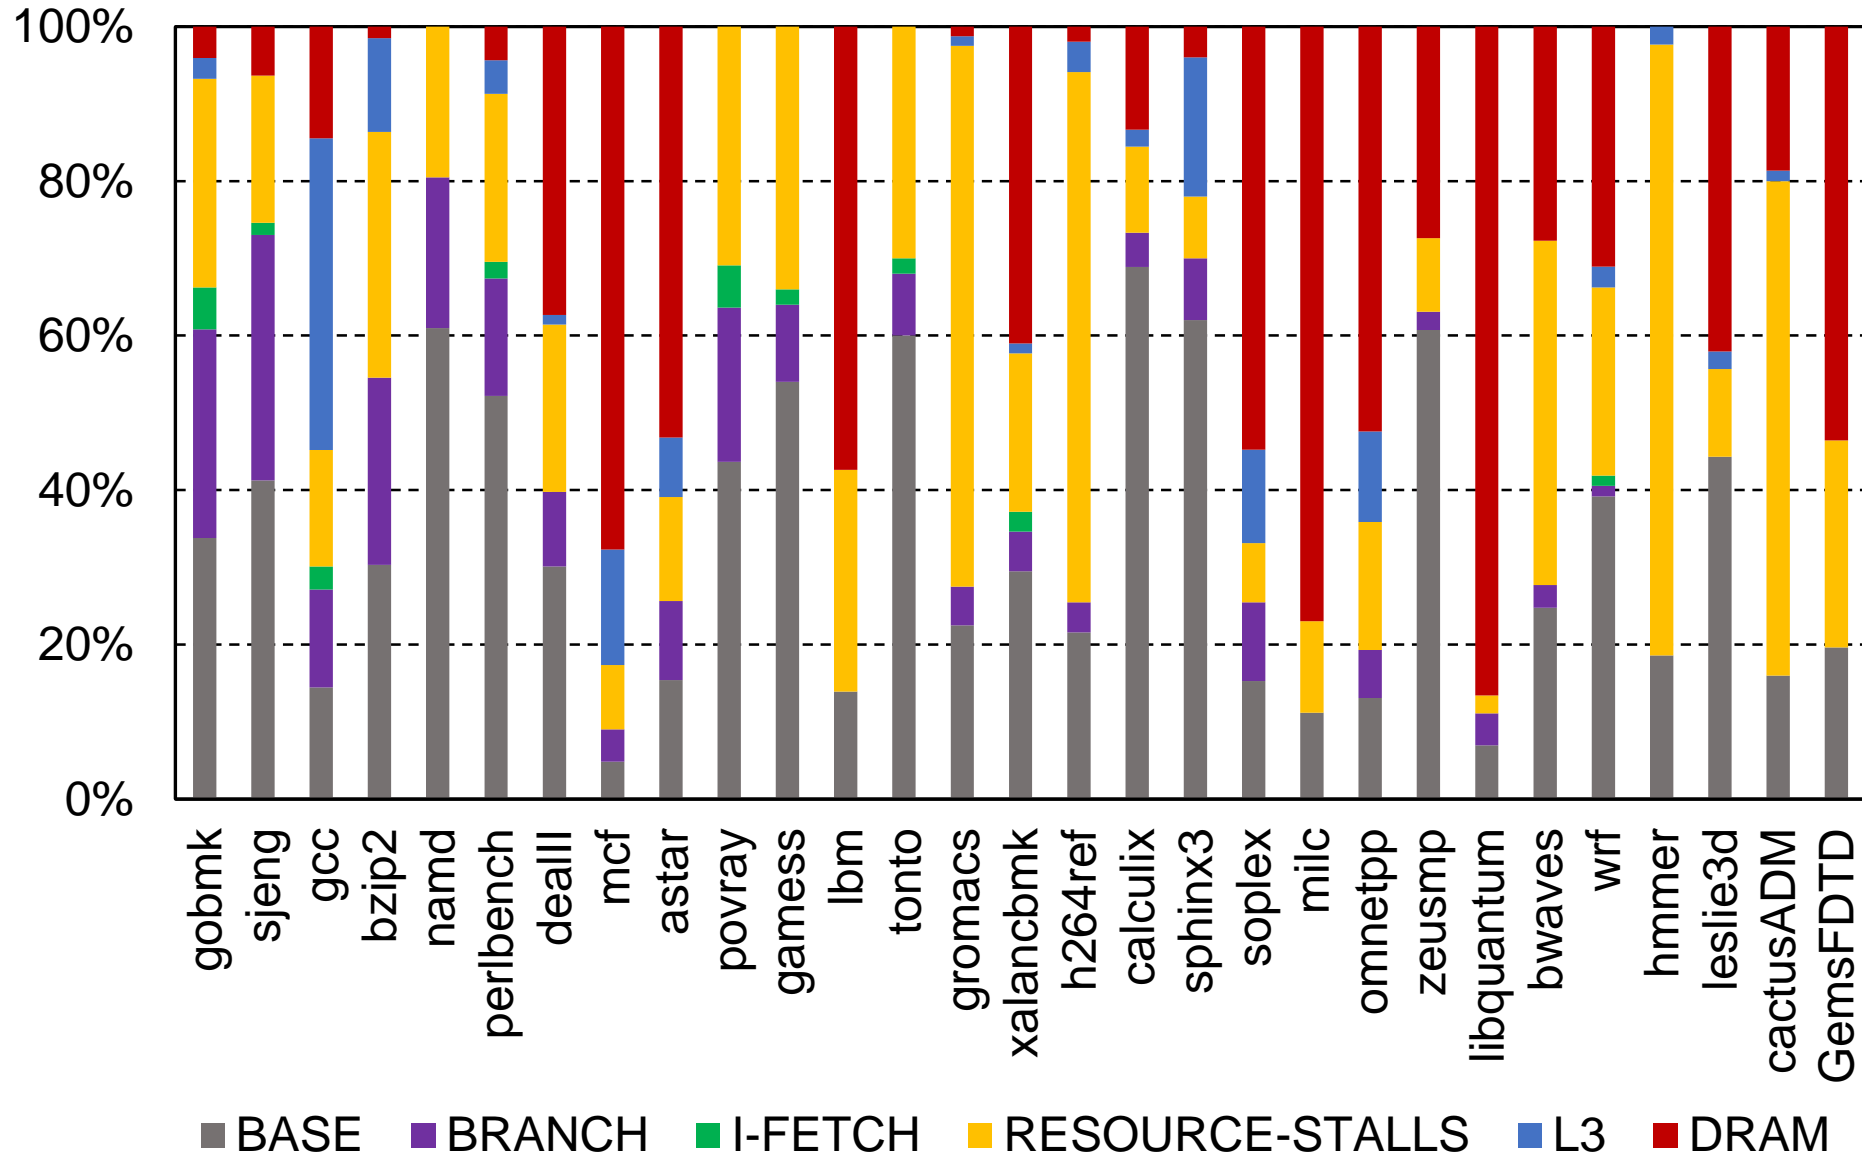
Vulnerability vs. CPI Stacks



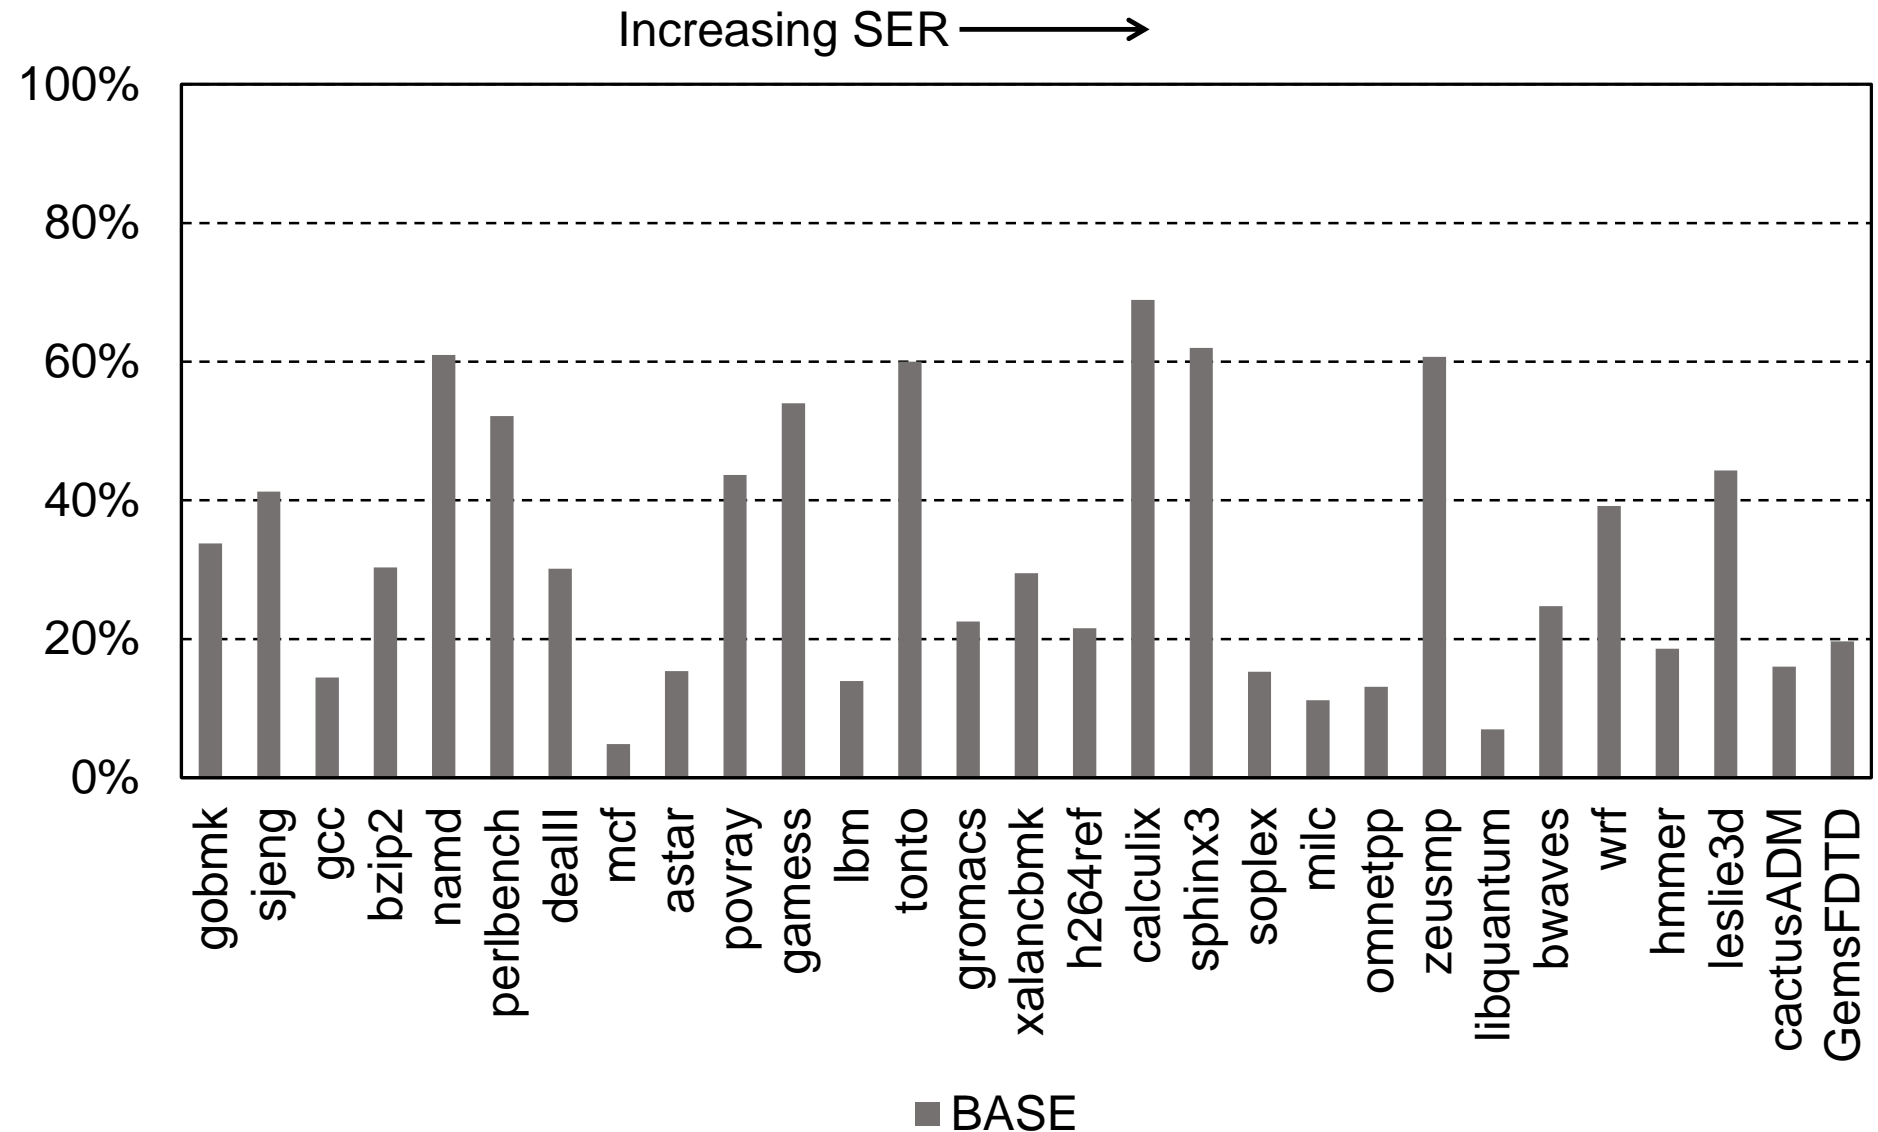
Vulnerability vs. CPI Stacks



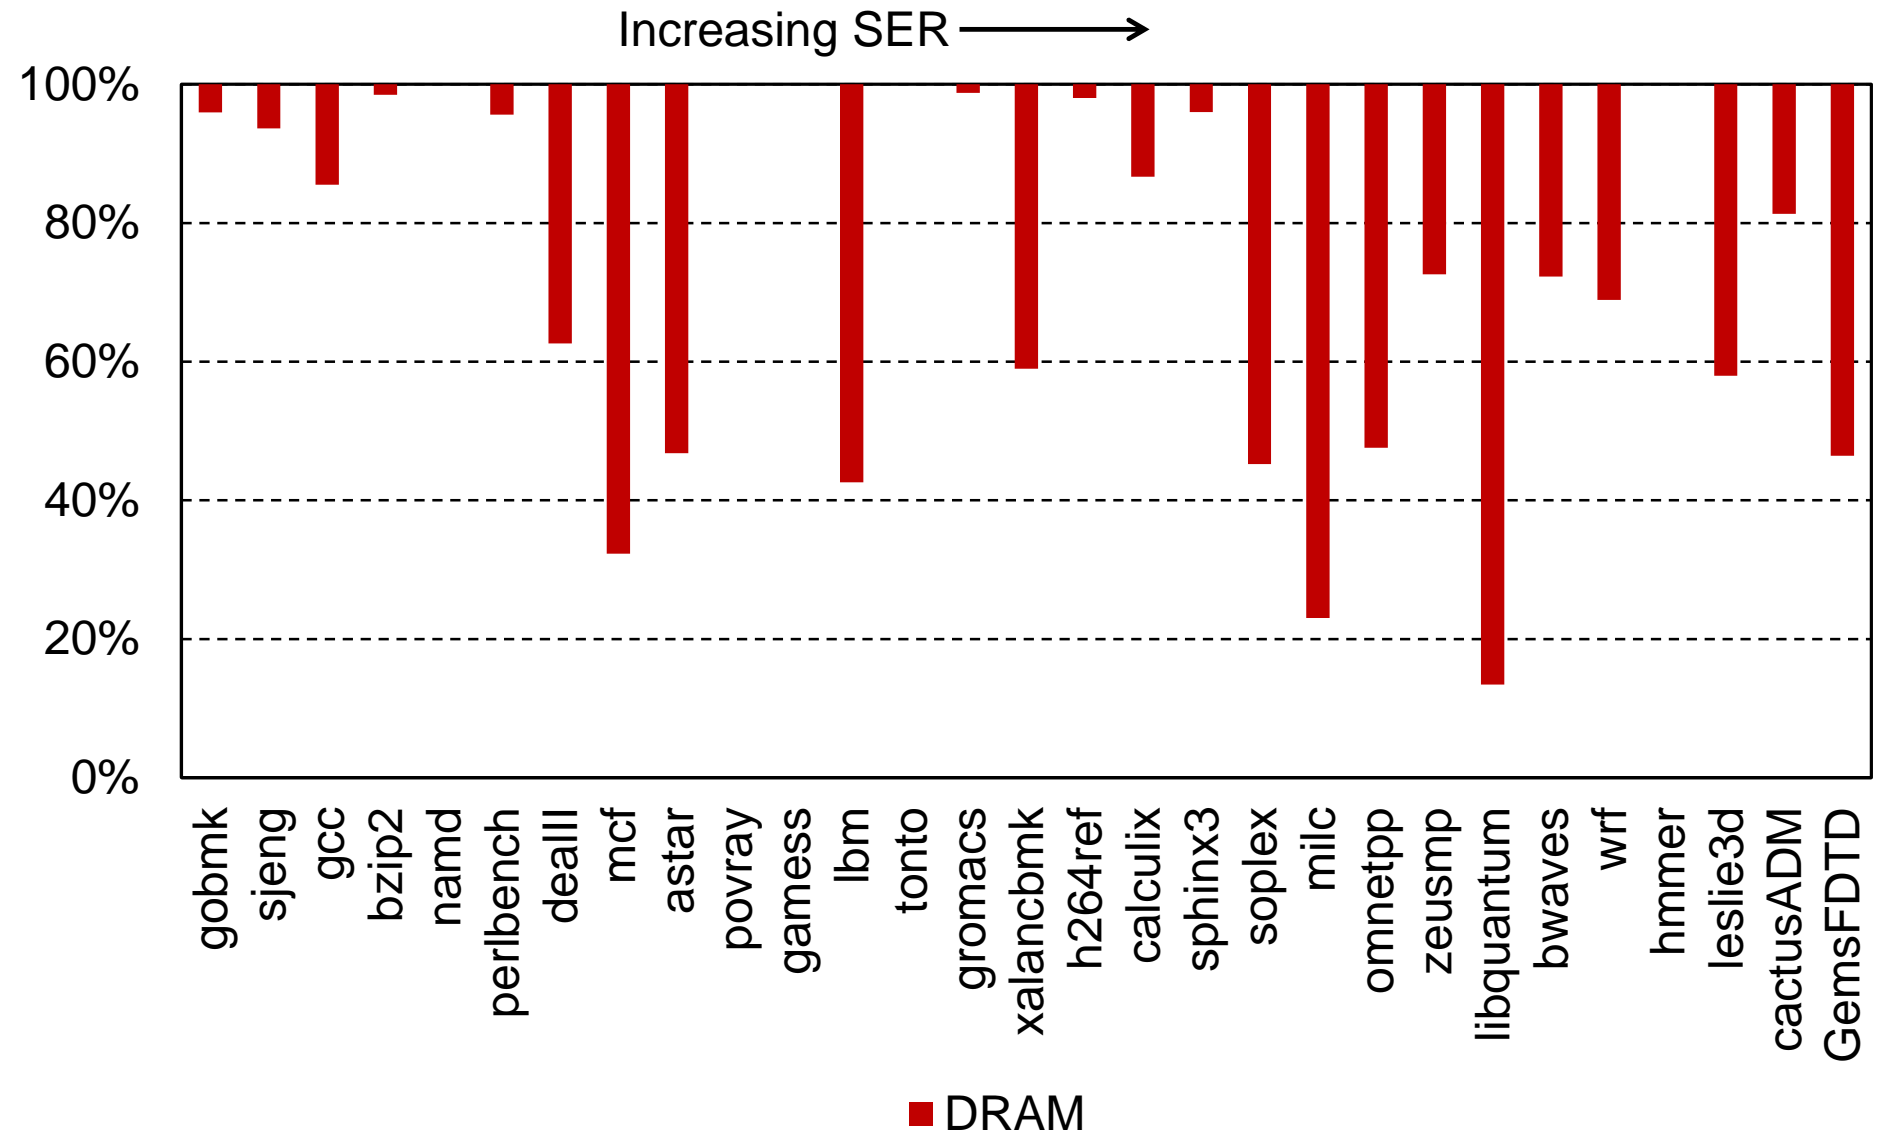
Vulnerability vs. CPI Stacks



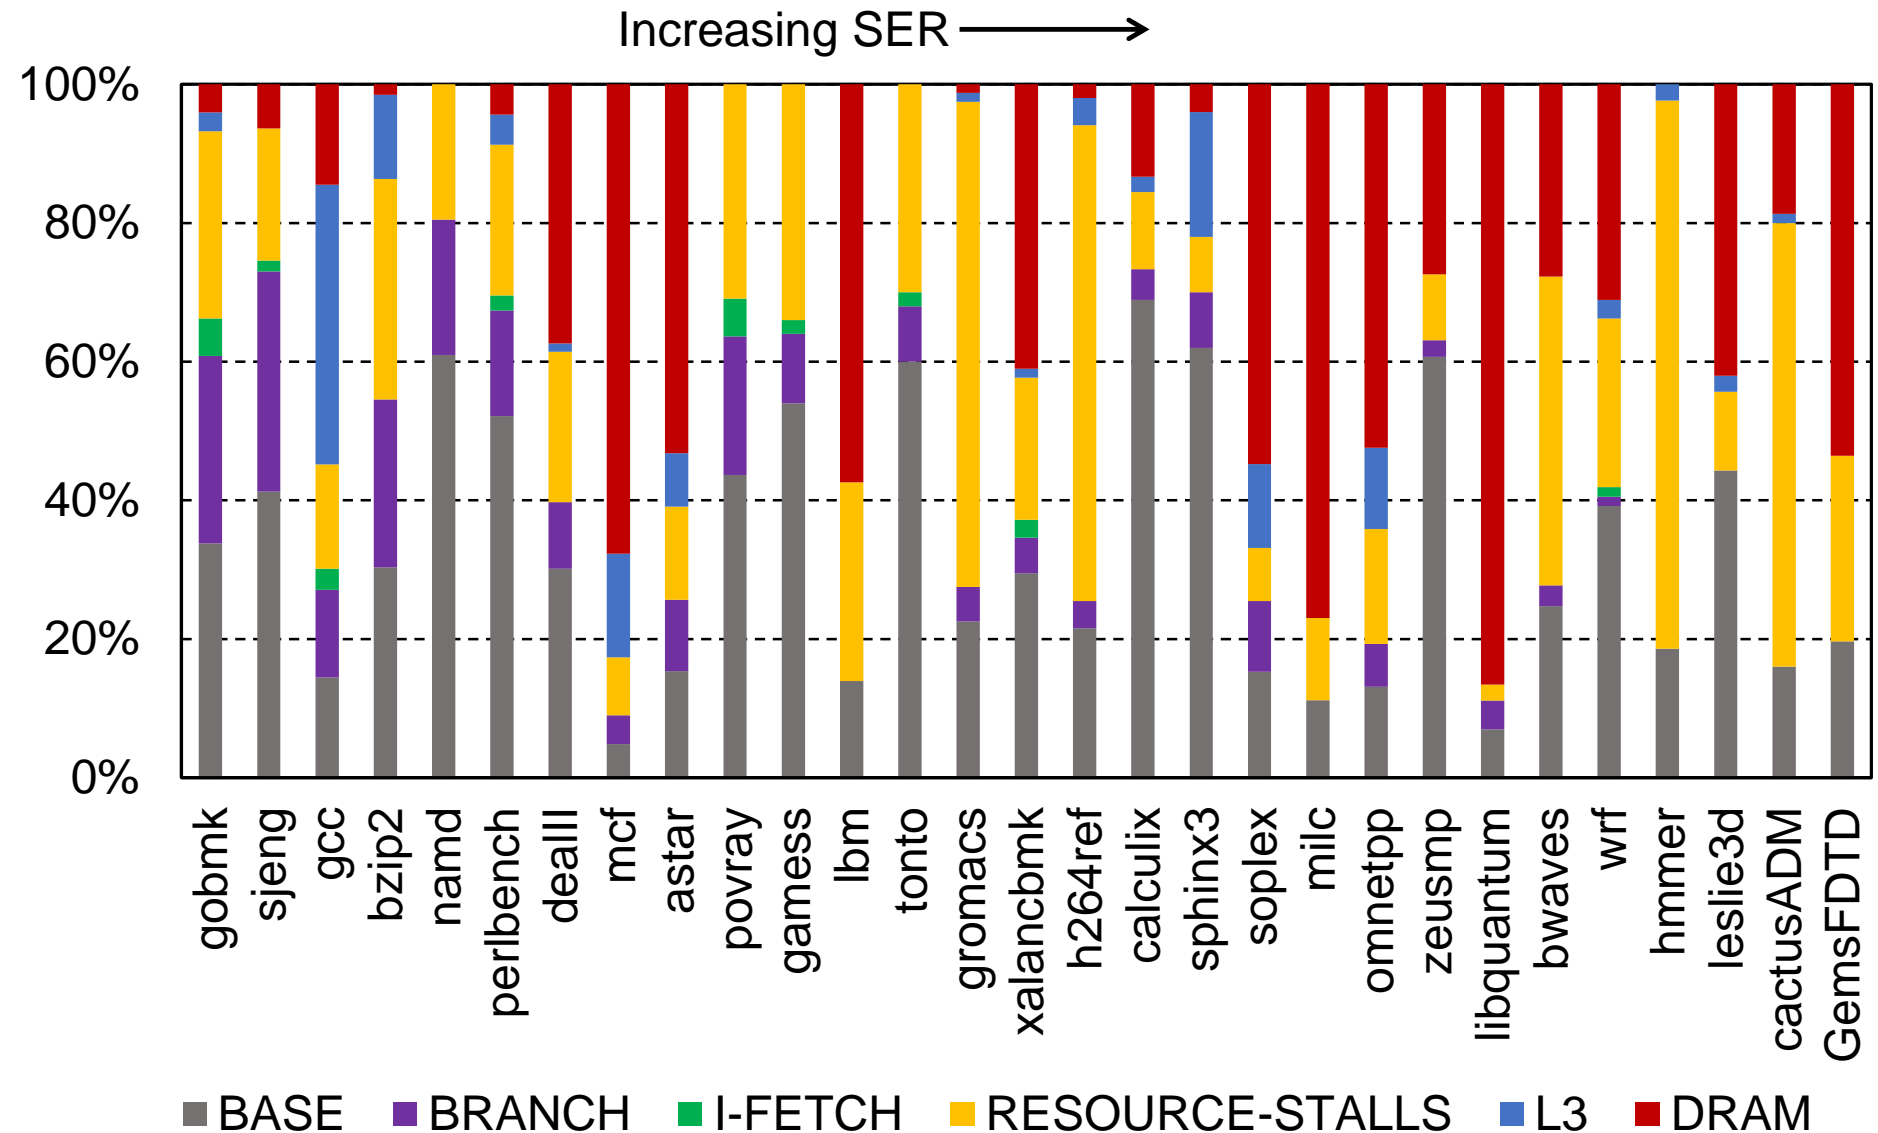
Vulnerability vs. Compute Intensity




Vulnerability vs. Memory Intensity




Vulnerability vs. CPI Stacks




Vulnerability vs. CPI Stacks




Vulnerability vs. CPI Stacks



Vulnerability vs. CPI Stacks



Vulnerability vs. CPI Stacks



Reliability vs. Application Type

Reliability vs. Application Type

1. No simple workload characteristic
 - For example, memory-intensity or compute-intensity

Reliability vs. Application Type


1. No simple workload characteristic
 - For example, **memory-intensity** or **compute-intensity**
2. Predicting SER is not straightforward*
 - **Complex interaction** of various workload characteristics

Reliability vs. Application Type

1. No simple workload characteristic
 - For example, **memory-intensity** or **compute-intensity**
2. Predicting SER is not straightforward*
 - **Complex interaction** of various workload characteristics
3. A **dynamic mechanism** required
 - To monitor reliability on either core type

Reliability vs. Application Type

1. No simple workload characteristic
 - For example, **memory-intensity** or **compute-intensity**
2. Predicting SER is not straightforward*
 - **Complex interaction** of various workload characteristics
3. A **dynamic mechanism** required
 - To monitor reliability on either core type



Reliability-Aware Scheduler

Reliability Metric for Multiprogram Workloads

Reliability Metric for Multiprogram Workloads

In isolation on one big core:

Reliability Metric for Multiprogram Workloads

In isolation on one big core:

A • 

time = 1

SER = 4

Reliability Metric for Multiprogram Workloads

In isolation on one big core:

A •————→

time = 1

SER = 4

B •————→

time = 1

SER = 1

Reliability Metric for Multiprogram Workloads

In isolation on one big core:

A •————→

time = 1

SER = 4

B •————→

time = 1

SER = 1

As a 2-program workload on an CMP with two big cores:

Reliability Metric for Multiprogram Workloads

In isolation on one big core:

A •————→

time = 1

SER = 4

B •————→

time = 1

SER = 1

As a 2-program workload on an CMP with two big cores:

A •————→

time = 1

SER = 4

Reliability Metric for Multiprogram Workloads

In isolation on one big core:

A •————→

time = 1

SER = 4

B •————→

time = 1

SER = 1

As a 2-program workload on an CMP with two big cores:

A •————→

time = 1

SER = 4

B •————→

time = 2

SER = 1

Reliability Metric for Multiprogram Workloads

In isolation on one big core:

A •————→

time = 1

SER = 4

B •————→

time = 1

SER = 1

As a 2-program workload on an CMP with two big cores:

A •————→

time = 1

SER = 4

B •————→

time = 2

SER = 1

+

total SER

=

5

Reliability Metric for Multiprogram Workloads

Adding SER

Reliability Metric for Multiprogram Workloads

Adding SER

```
graph LR; A[Adding SER] --> B[More weight to high-SER applications]
```

A diagram consisting of two light gray rounded rectangular boxes with black borders. The first box on the left contains the text 'Adding SER' in red and purple. A black line extends from the right side of this box, turns upwards, and then turns right to point at the first box of the second box on the right. The second box contains the text 'More weight to high-SER applications' in red and purple.

More weight to high-SER applications

Reliability Metric for Multiprogram Workloads

Adding SER

```
graph LR; A[Adding SER] --> B[More weight to high-SER applications]; A --> C[Ignores slowdown];
```

More weight to high-SER applications

Ignores slowdown

Reliability Metric for Multiprogram Workloads

Adding SER

More weight to high-SER applications

Ignores slowdown

Problem

Ignores the impact of performance on SER

Reliability Metric for Multiprogram Workloads

Adding SER

More weight to high-SER applications

Ignores slowdown

Problem

Ignores the impact of performance on SER

Solution

New metric: System-level Soft Error Rate (SSER)


System-level Soft Error Rate (SSER)

System-level Soft Error Rate (SSER)

Weight per-application SER by its slowdown

System-level Soft Error Rate (SSER)

Weight per-application SER by its slowdown




Slowdown \rightarrow More errors

$$wSER_i = SER_i \times slowdown_i$$


System-level Soft Error Rate (SSER)

Weight per-application SER by its slowdown



Slowdown \rightarrow More errors


$$wSER_i = SER_i \times slowdown_i$$



Add weighted SER for all applications


System-level Soft Error Rate (SSER)

Weight per-application SER by its slowdown



Slowdown \rightarrow More errors

$$wSER_i = SER_i \times slowdown_i$$




Add weighted SER for all applications

$$SSER = \sum_{i=1}^n wSER_i$$


System-level Soft Error Rate (SSER)

Weight per-application SER by its slowdown



Slowdown \rightarrow More errors

$$wSER_i = SER_i \times slowdown_i$$



Add weighted SER for all applications

$$SSER = \sum_{i=1}^n wSER_i$$


A novel metric to assess soft error vulnerability

Scheduling Algorithm

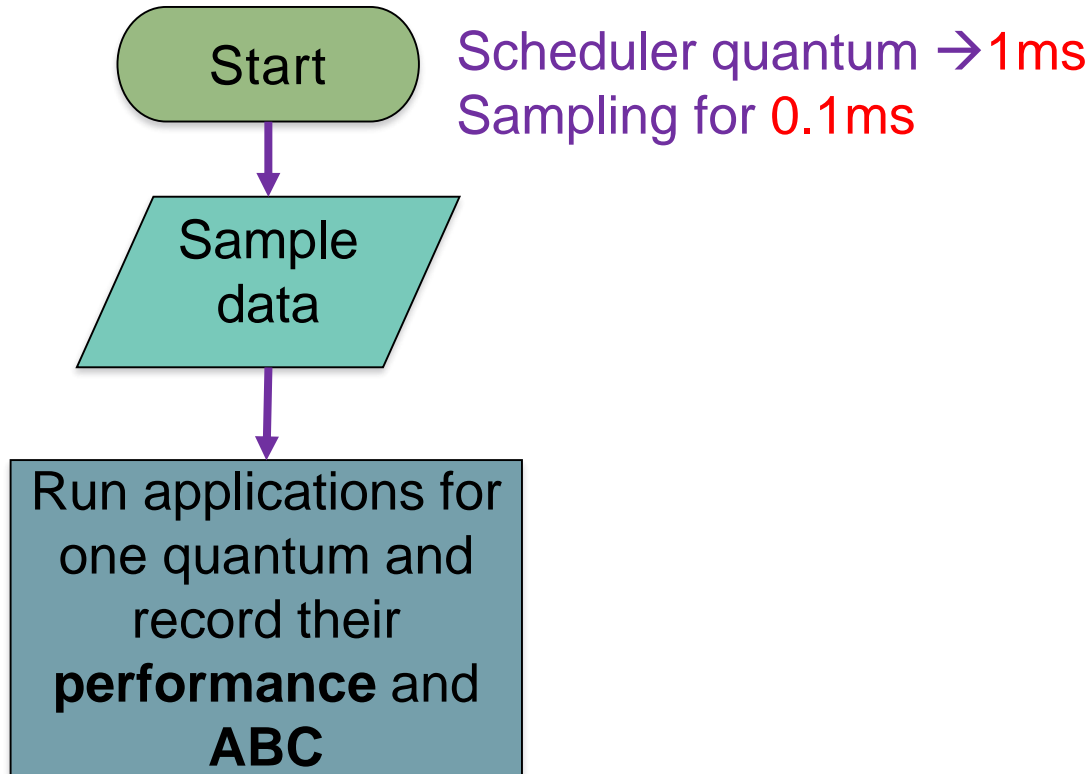
Scheduling Algorithm

Start

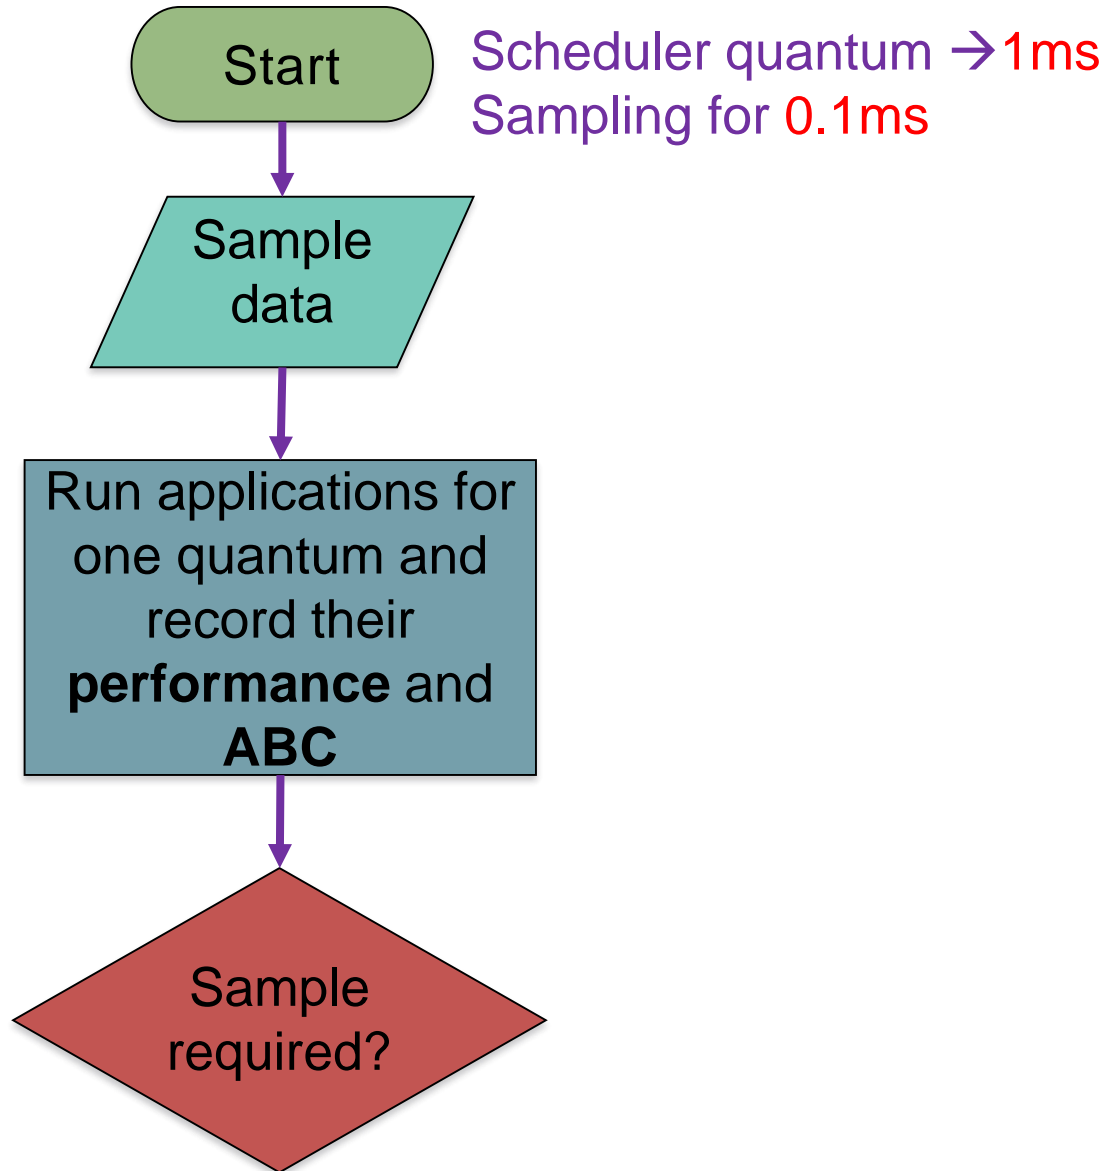
Scheduling Algorithm



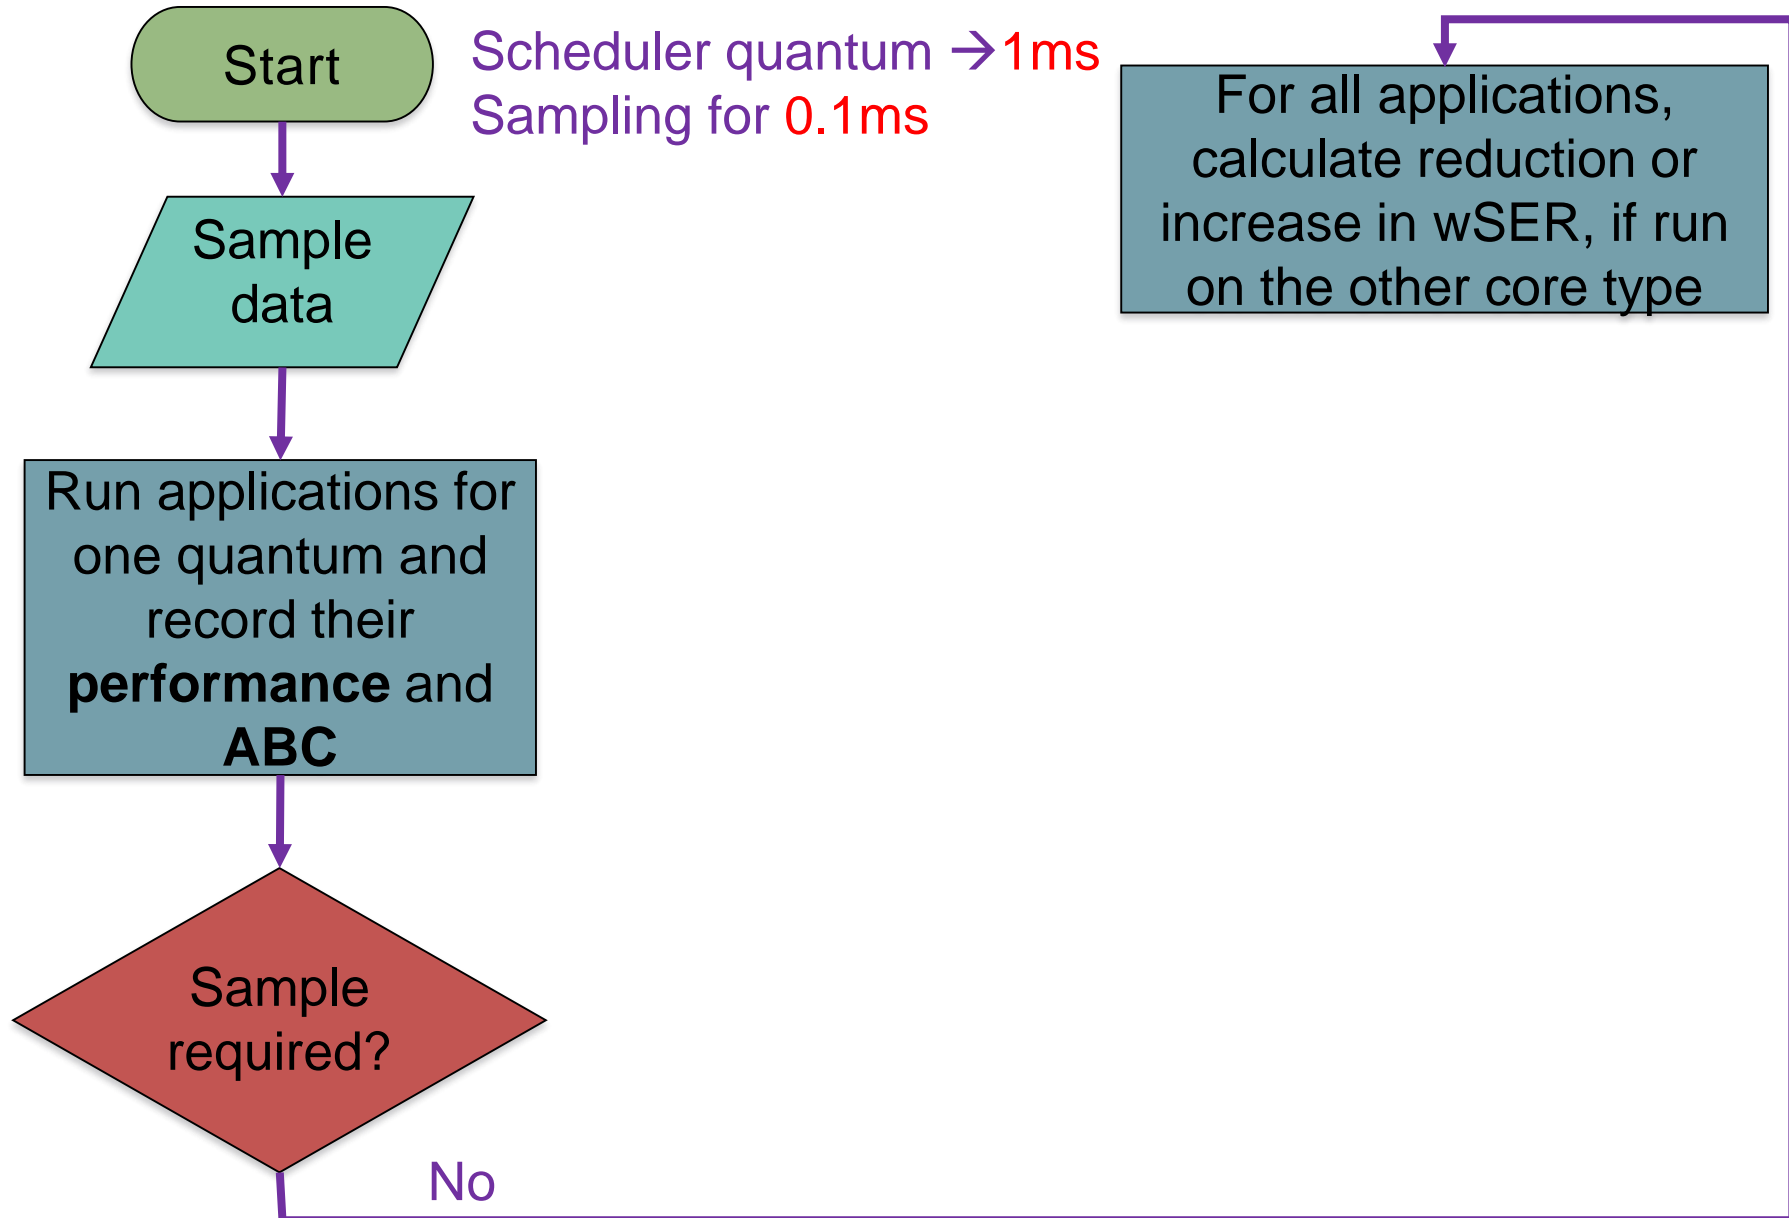
Scheduling Algorithm



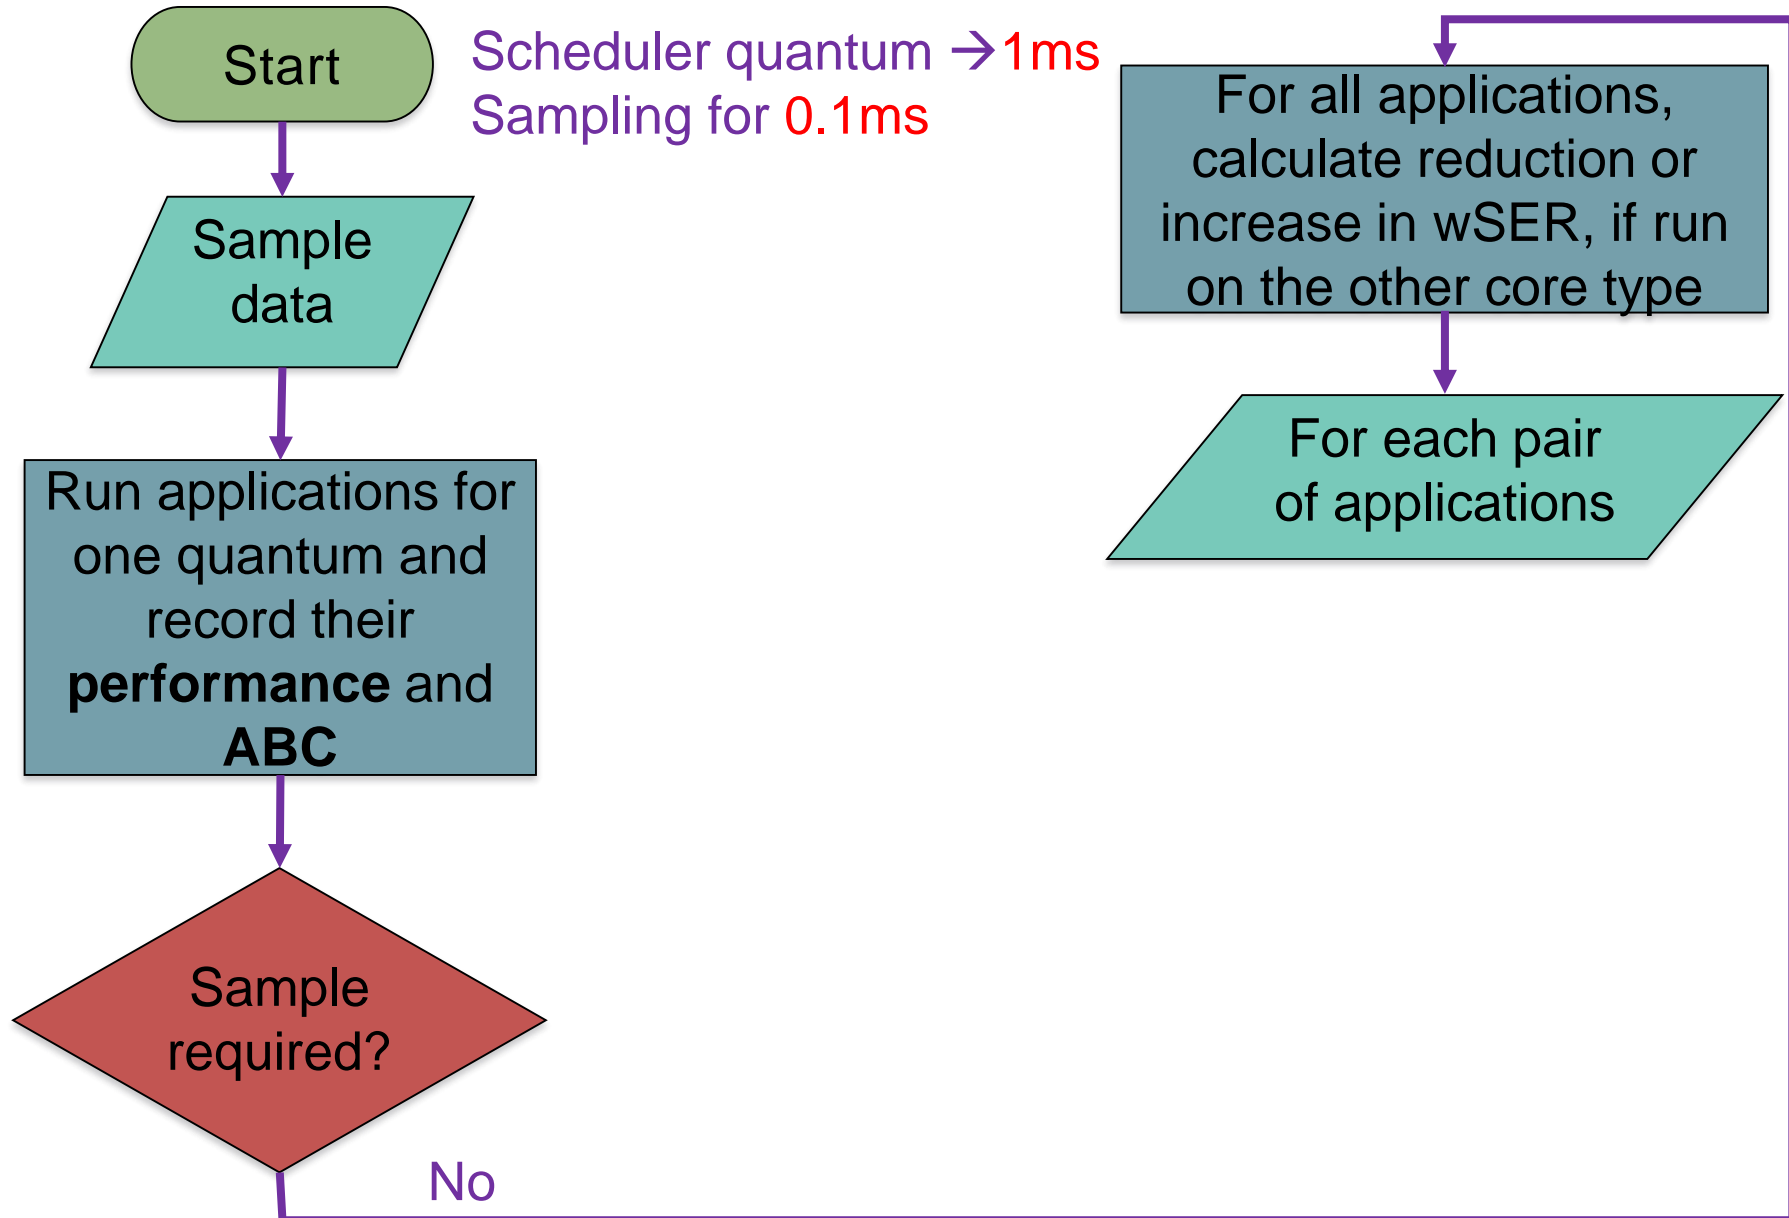
Scheduling Algorithm



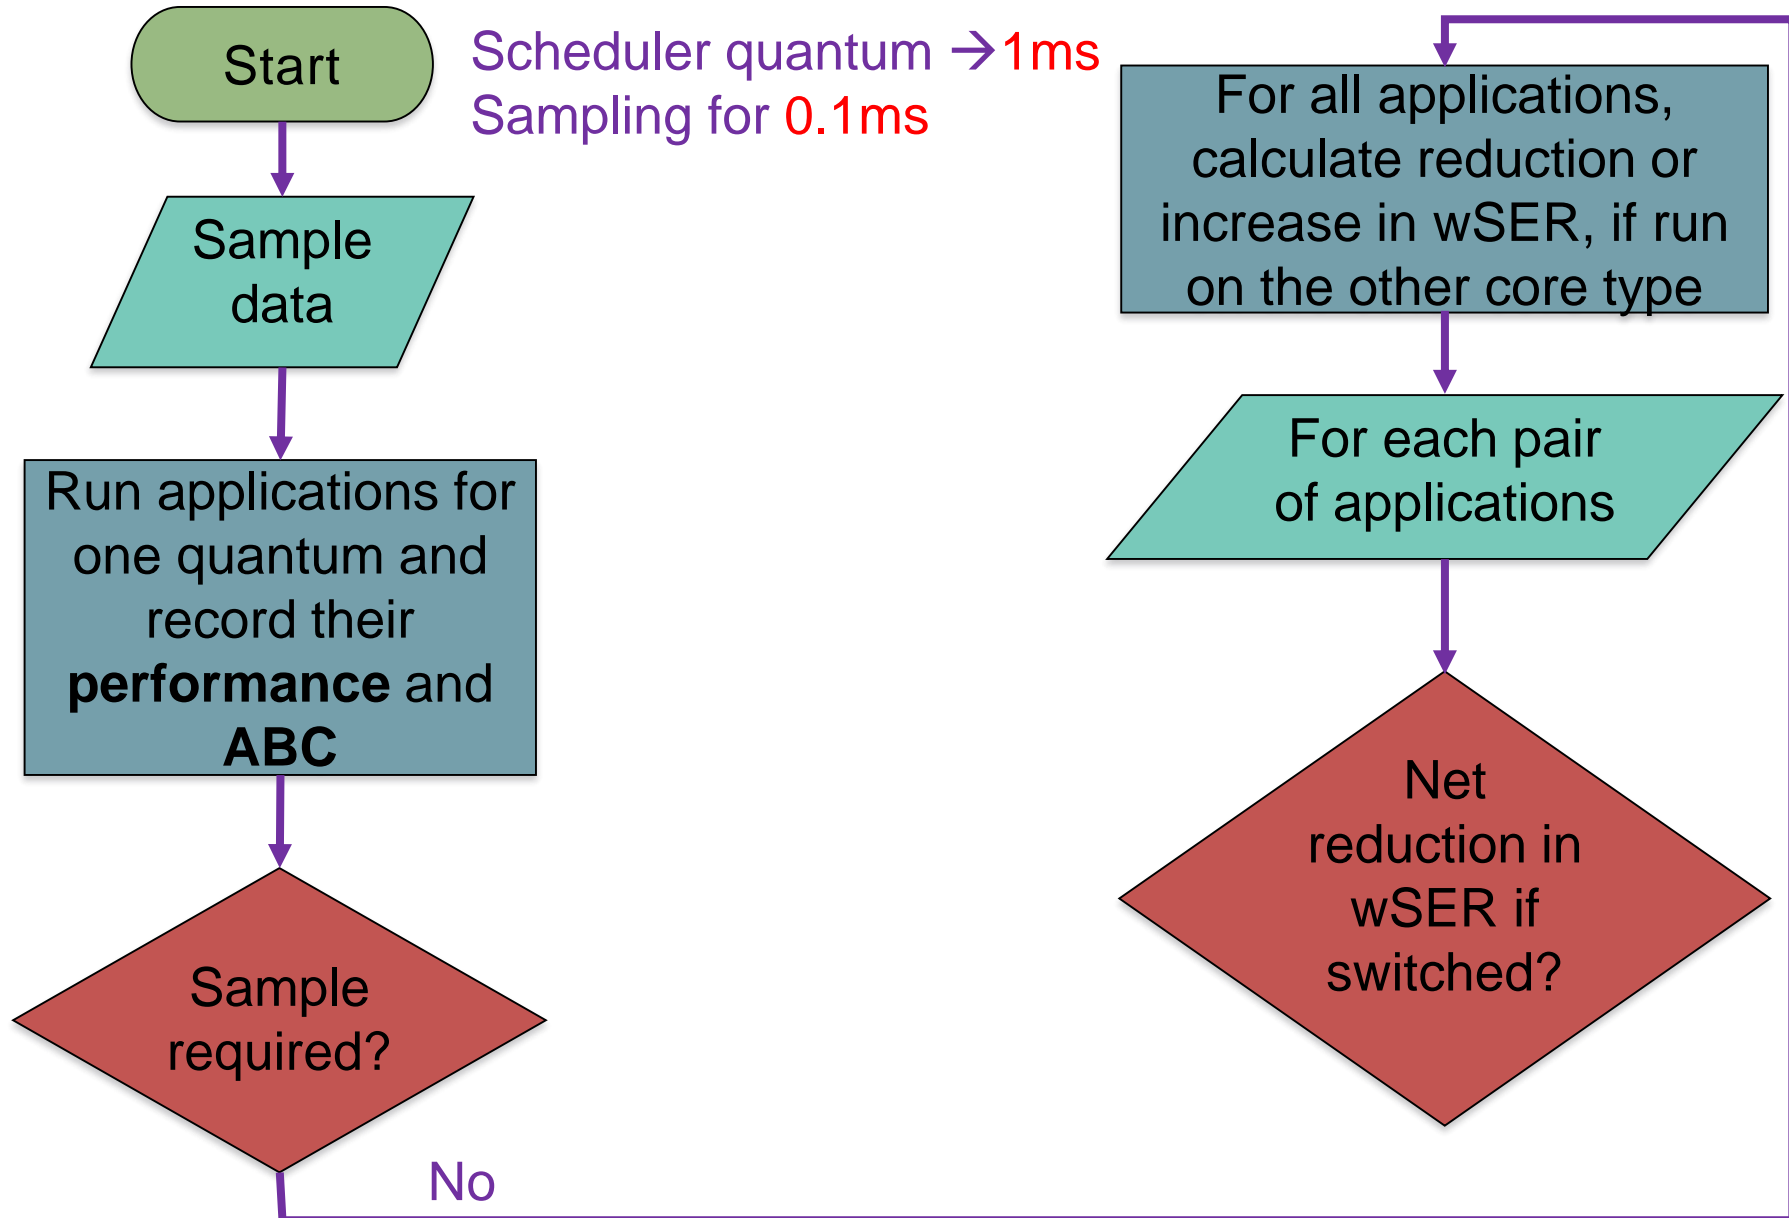
Scheduling Algorithm



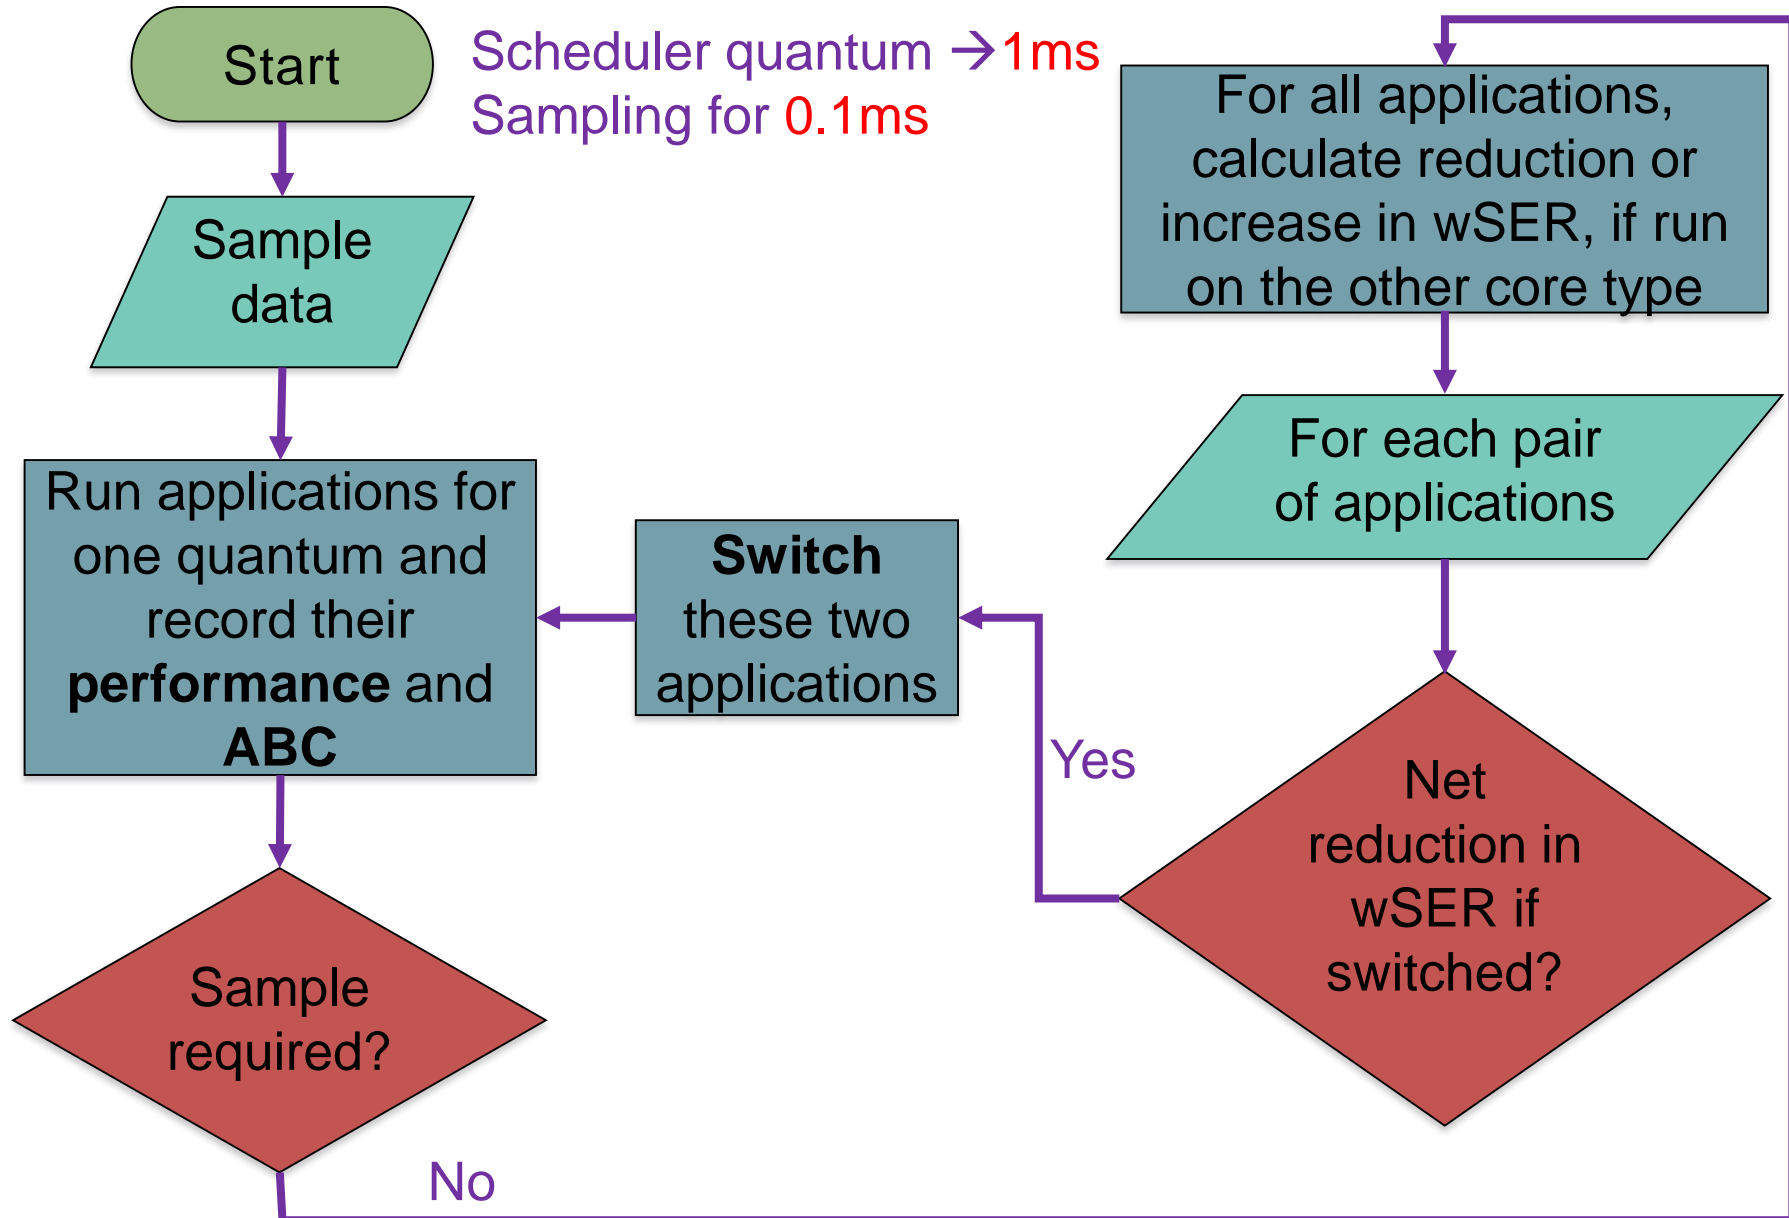
Scheduling Algorithm



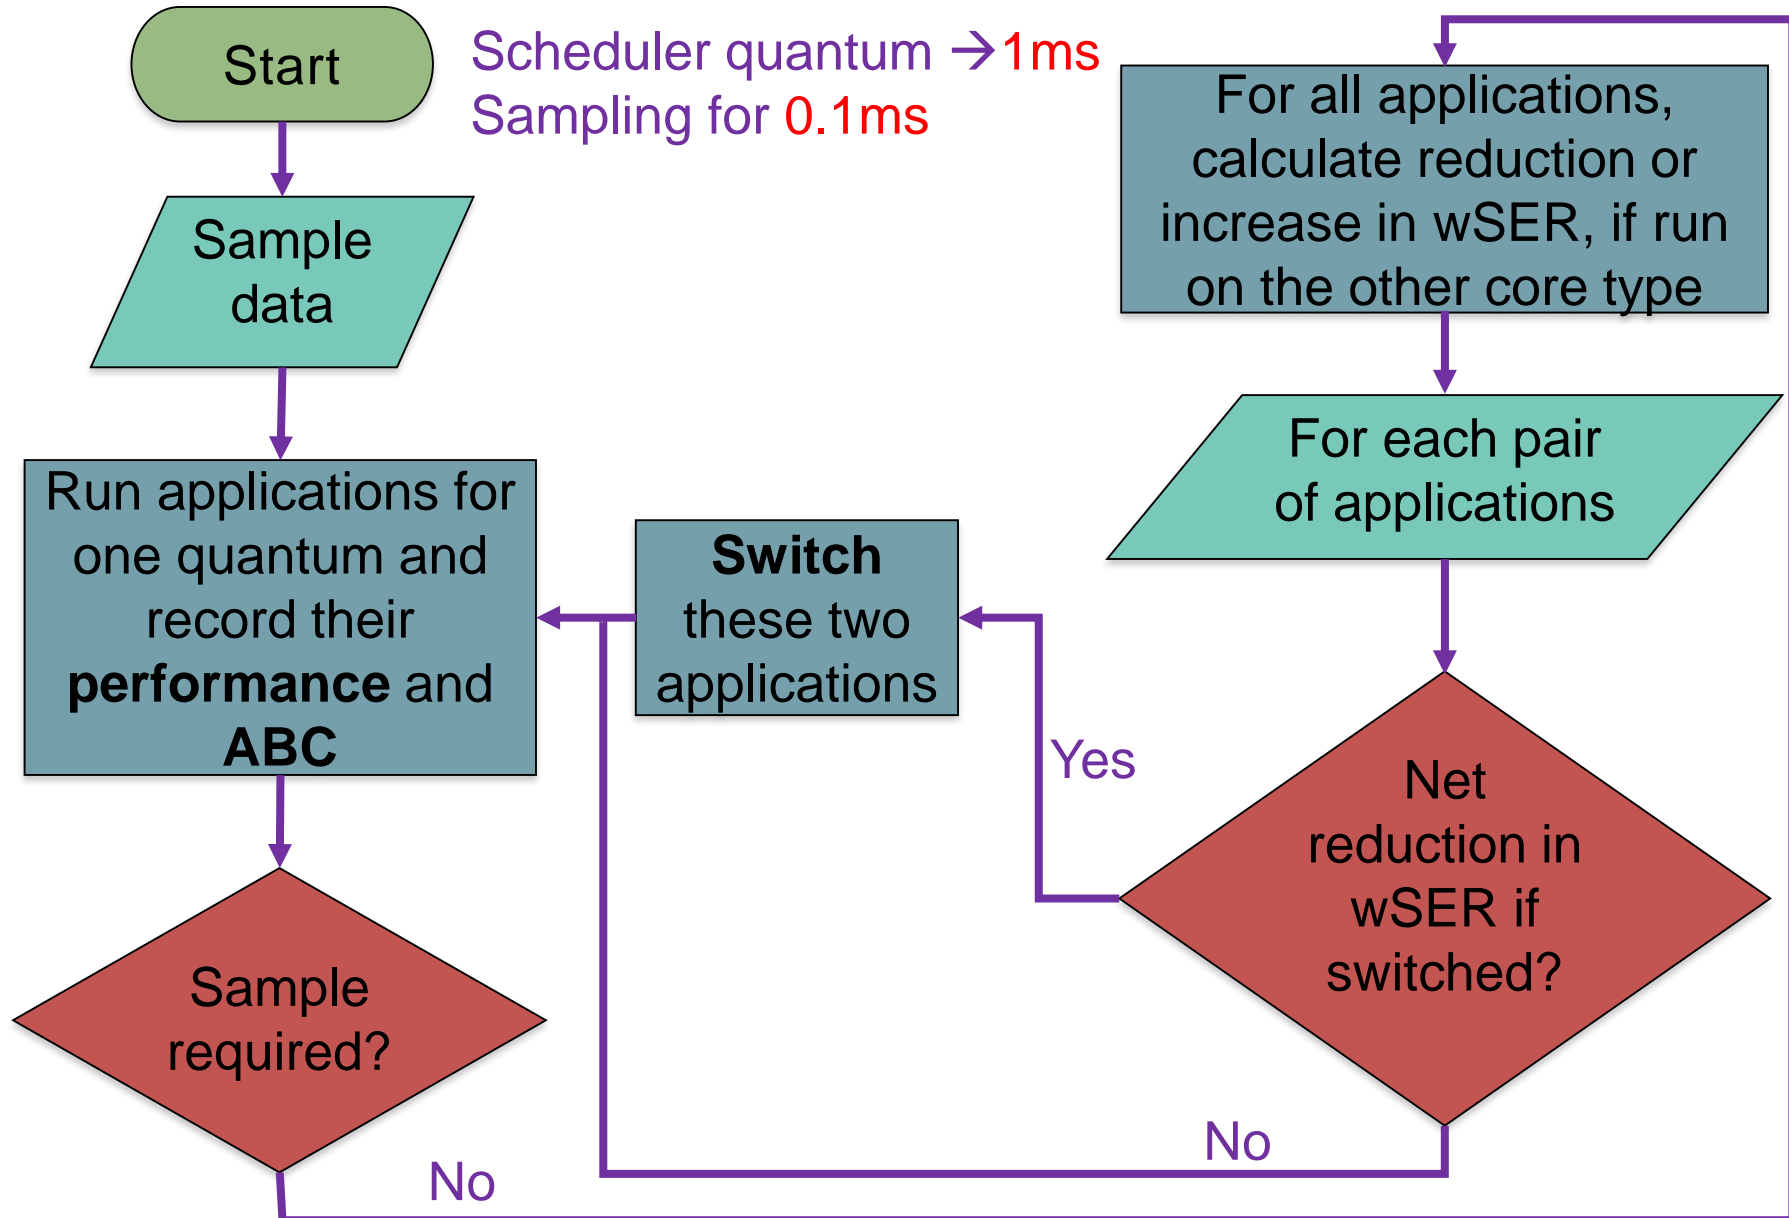
Scheduling Algorithm



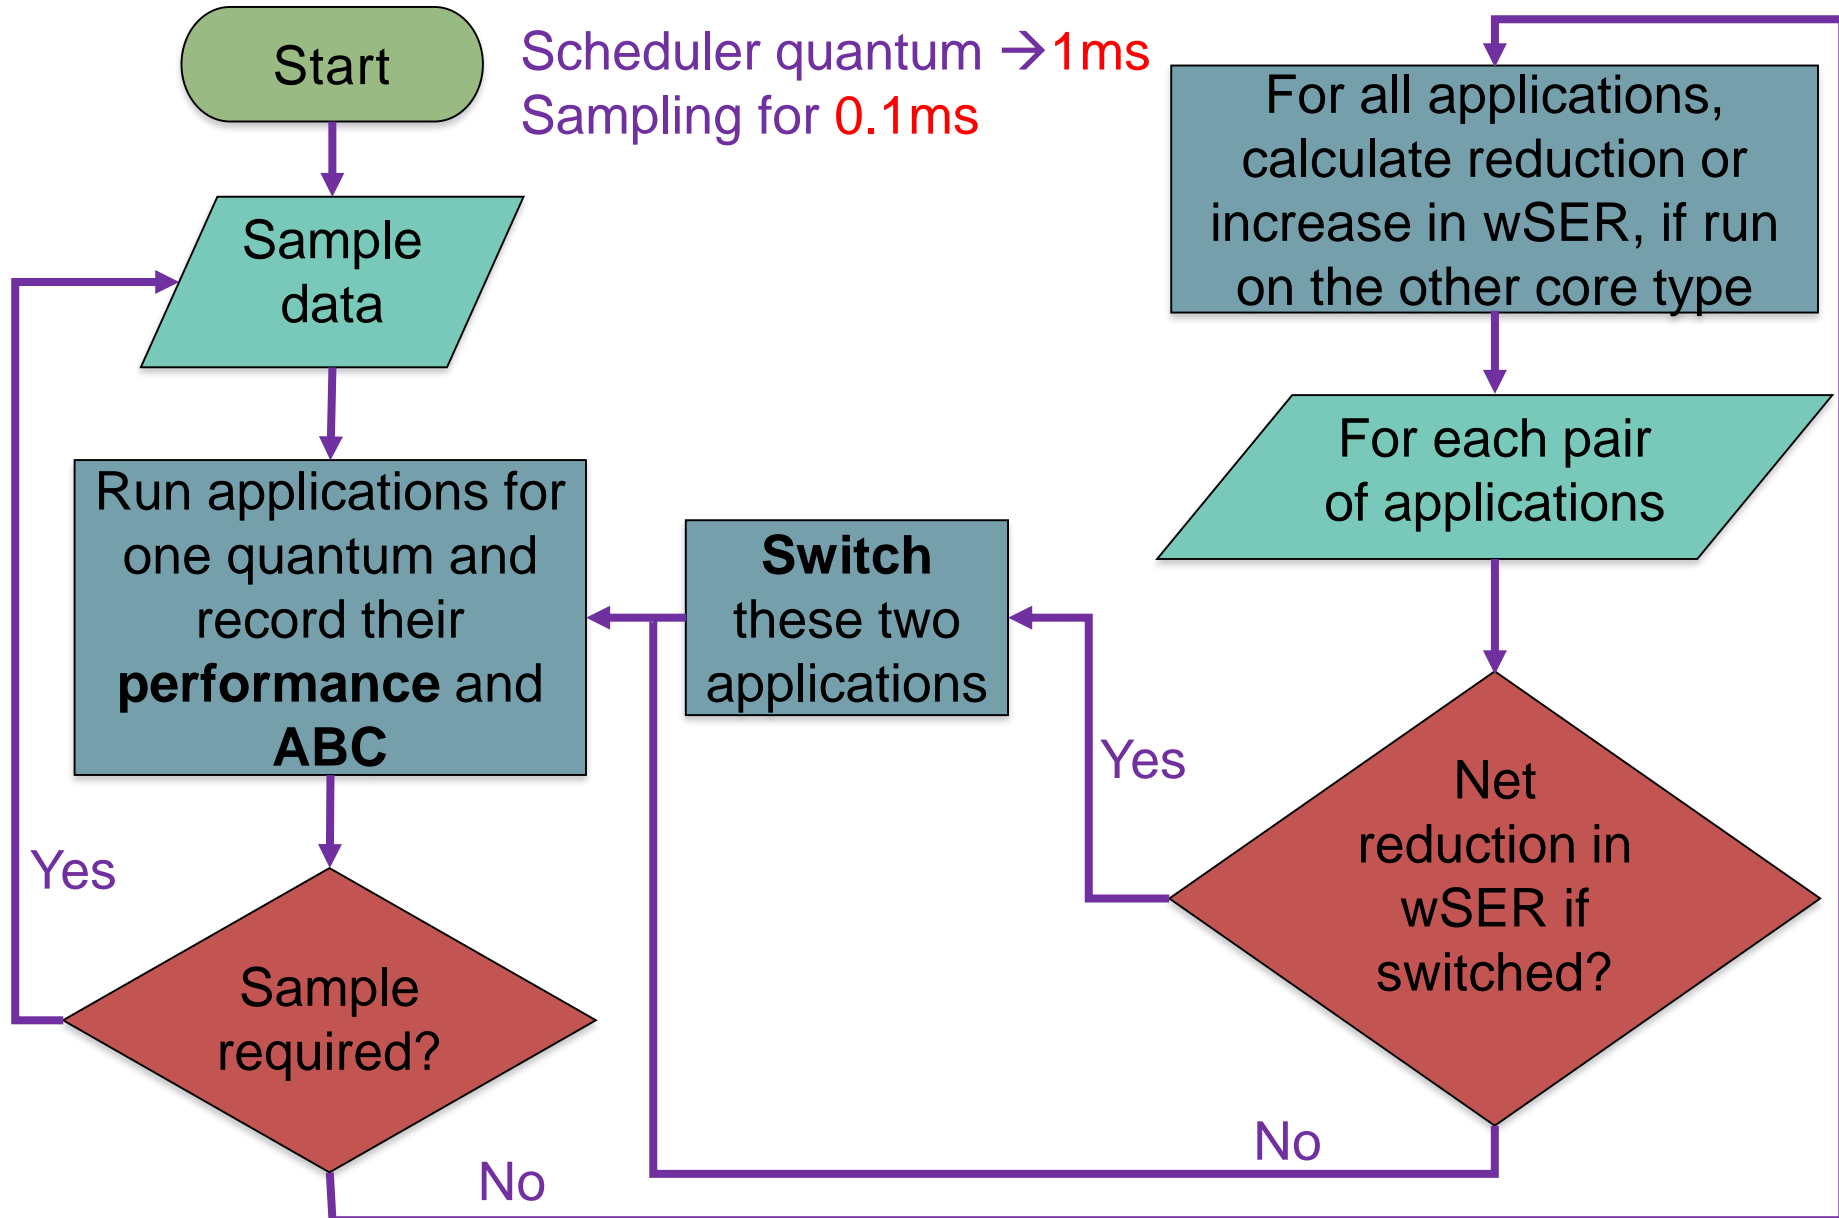
Scheduling Algorithm



Scheduling Algorithm



Scheduling Algorithm



Hardware Overhead

Hardware Overhead

Two 12-bit counters per ROB entry to record dispatch and issue times

Hardware Overhead

Two 12-bit counters per ROB entry to record dispatch and issue times



Hardware Overhead

Two 12-bit counters per ROB entry to record dispatch and issue times




Hardware Overhead

Two 12-bit counters per ROB entry to record dispatch and issue times




Hardware Overhead

Two 12-bit counters per ROB entry to record dispatch and issue times



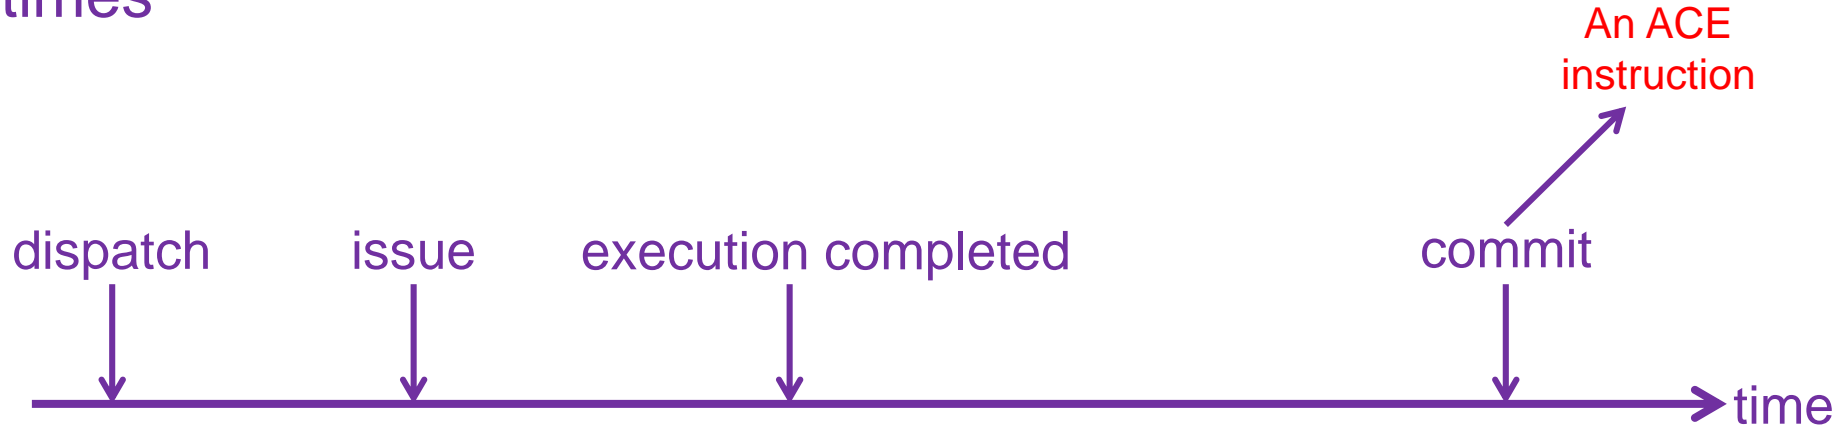
Hardware Overhead

Two 12-bit counters per ROB entry to record dispatch and issue times



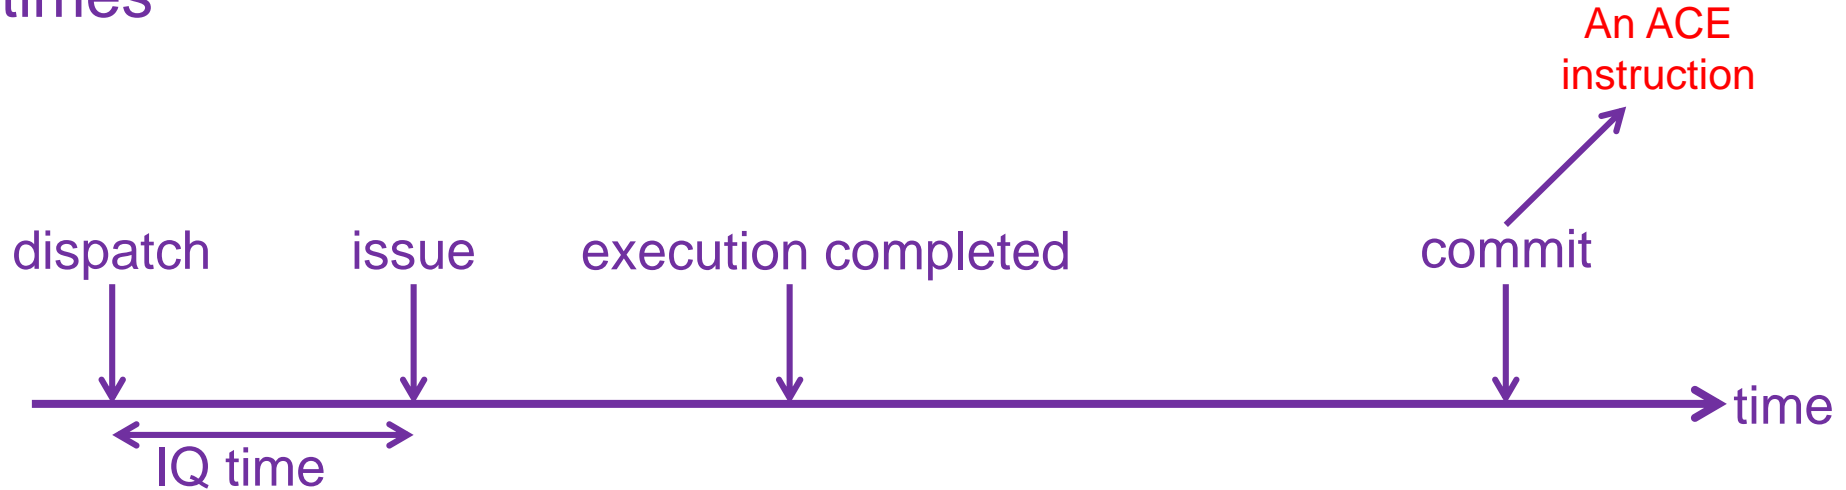
Hardware Overhead

Two 12-bit counters per ROB entry to record **dispatch** and **issue** times



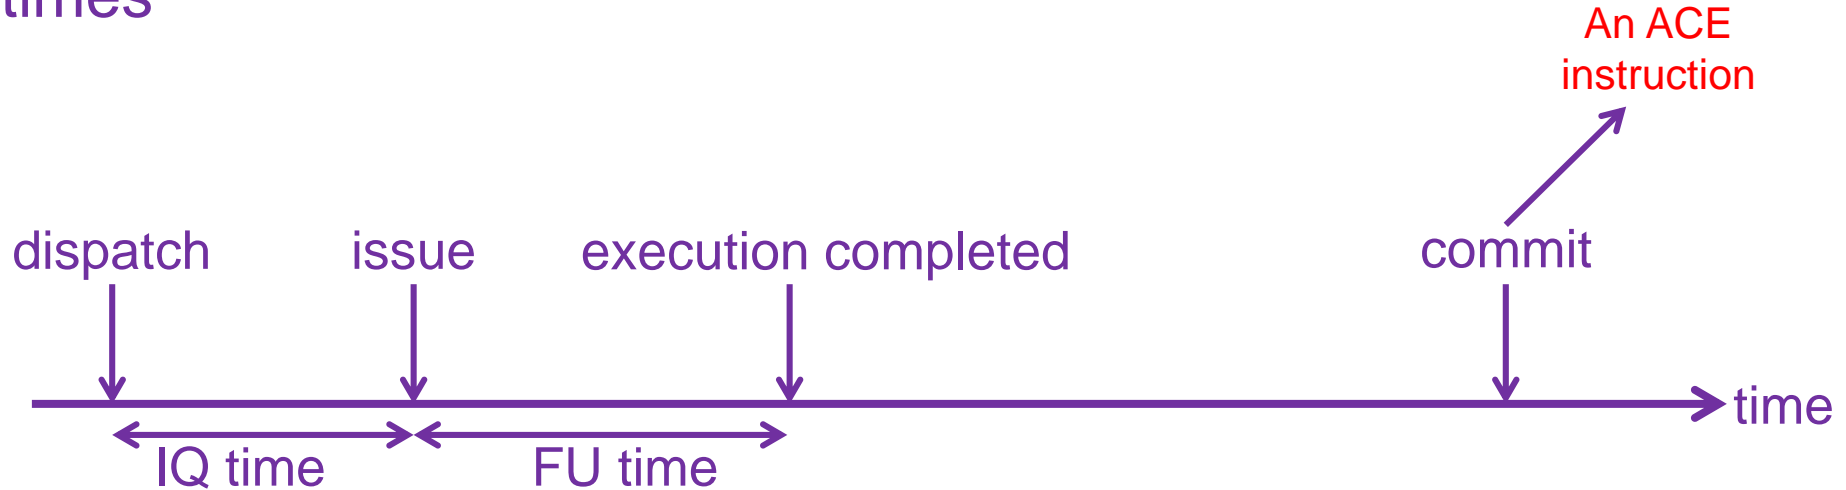
Hardware Overhead

Two 12-bit counters per ROB entry to record **dispatch** and **issue** times



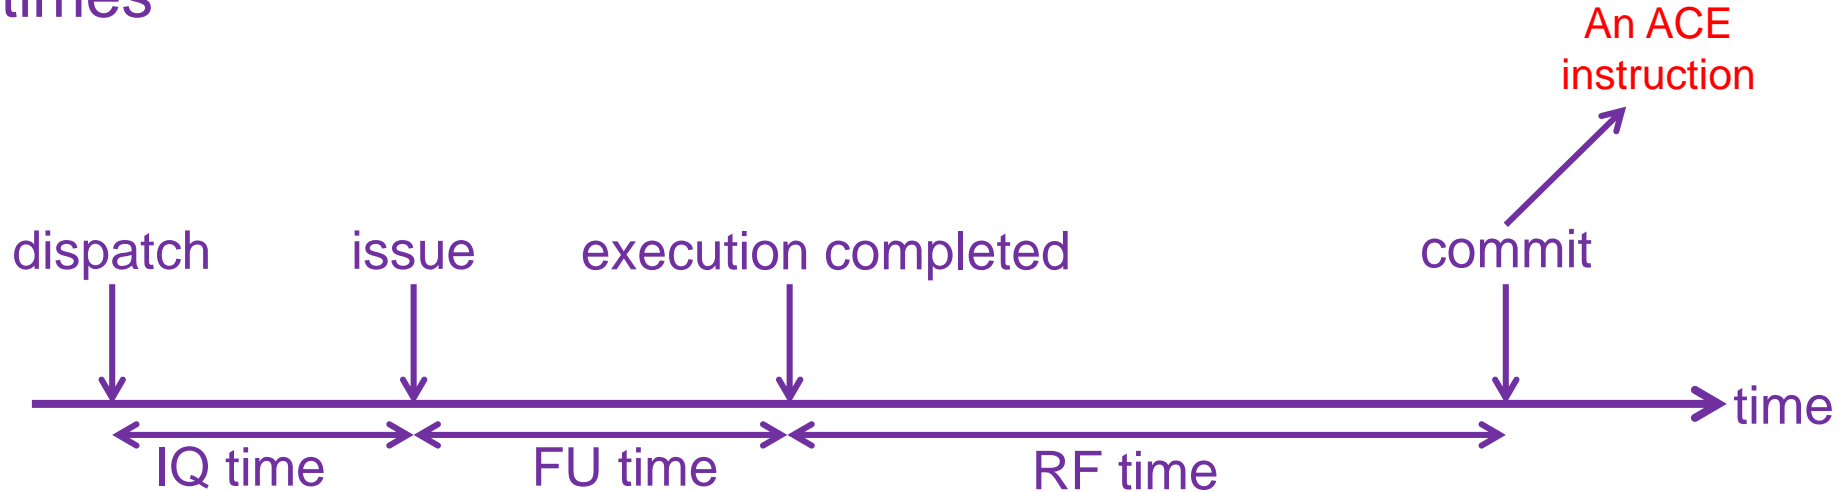
Hardware Overhead

Two 12-bit counters per ROB entry to record **dispatch** and **issue** times



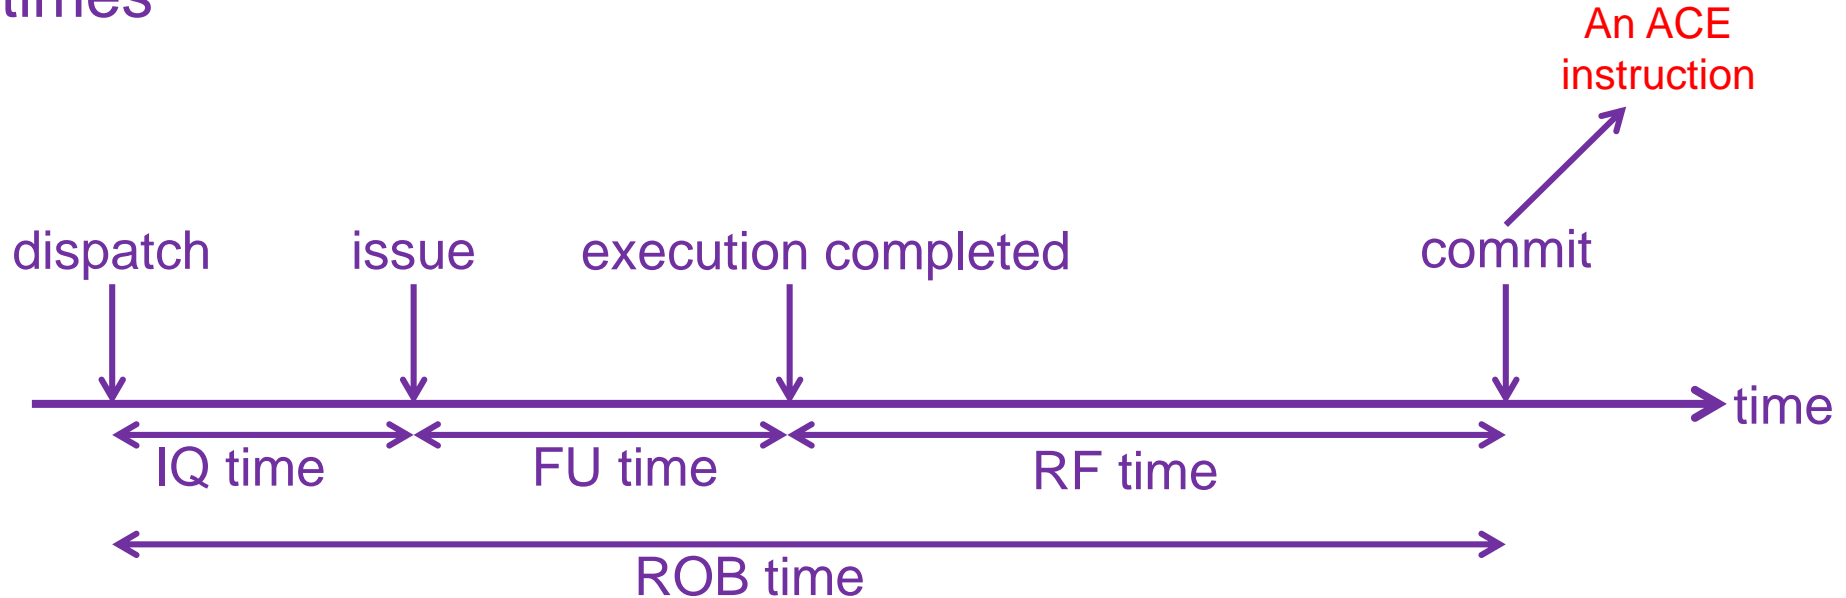
Hardware Overhead

Two 12-bit counters per ROB entry to record **dispatch** and **issue** times




Hardware Overhead

Two 12-bit counters per ROB entry to record **dispatch** and **issue** times



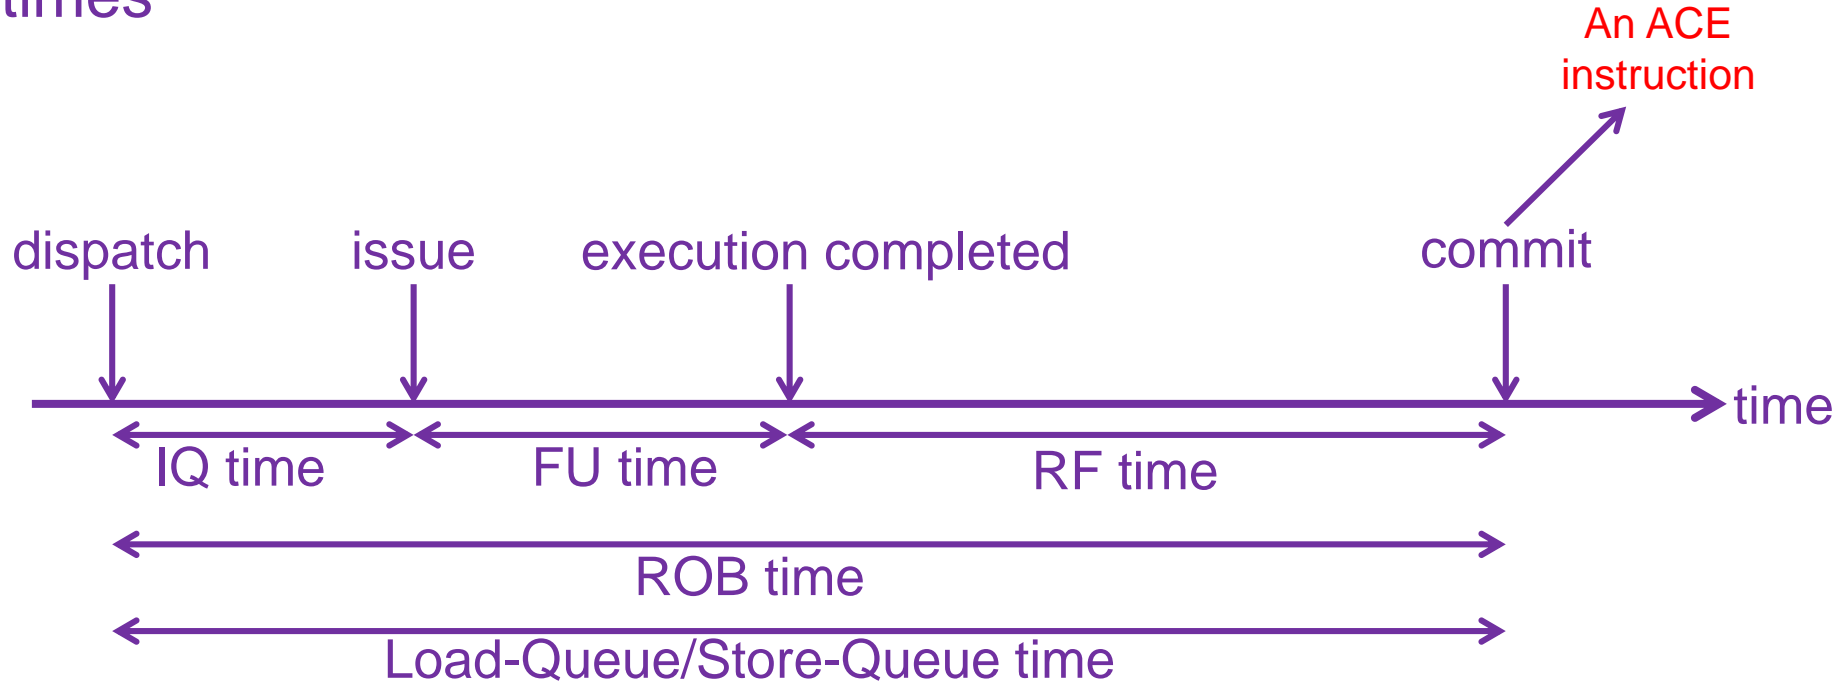
Hardware Overhead

Two 12-bit counters per ROB entry to record **dispatch** and **issue** times



Hardware Overhead

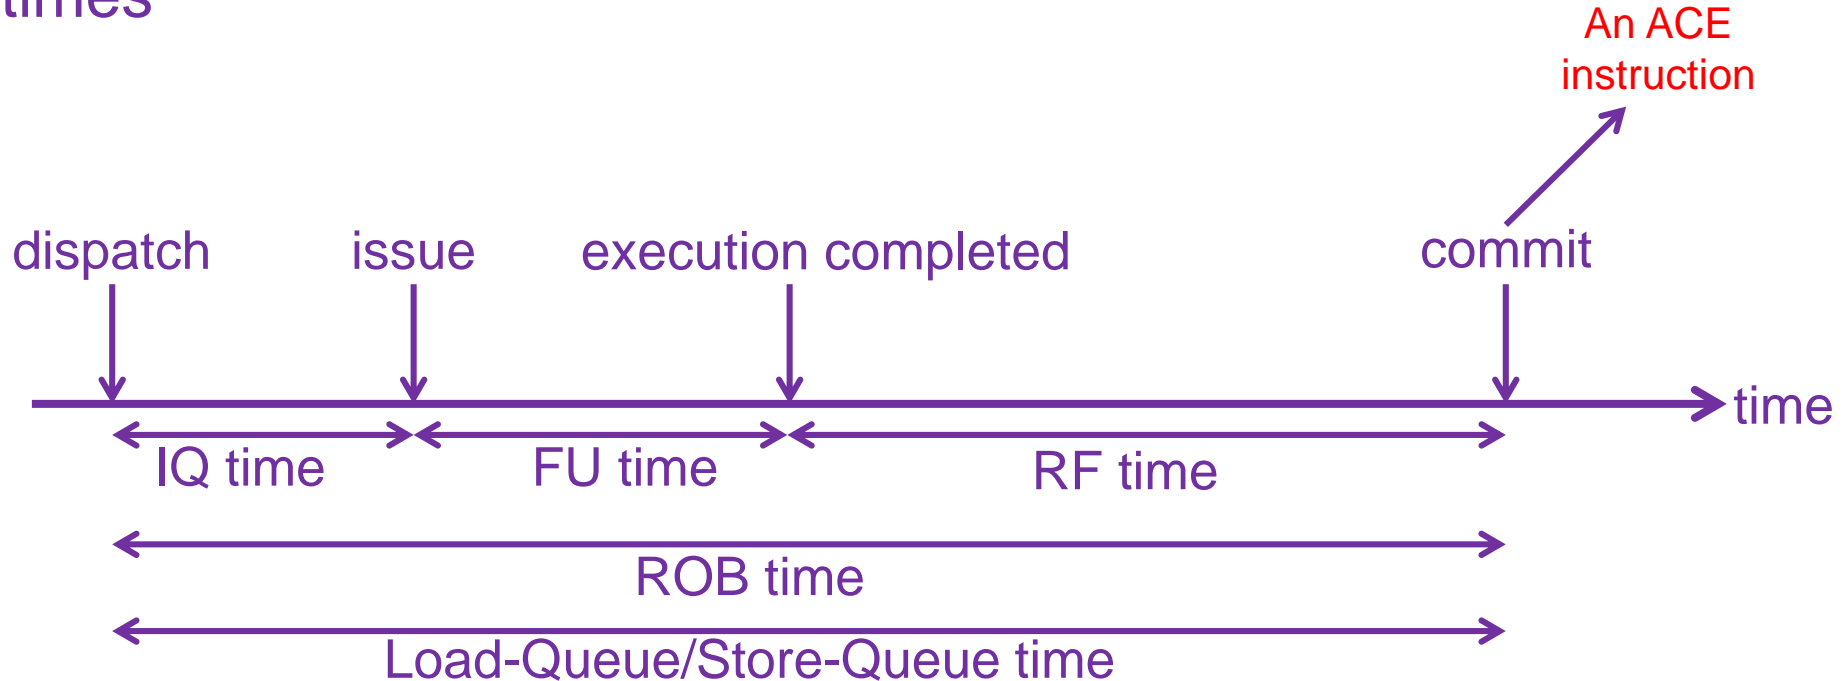
Two 12-bit counters per ROB entry to record **dispatch** and **issue** times



Total hardware overhead → **904 bytes** per core

Hardware Overhead

Two 12-bit counters per ROB entry to record **dispatch** and **issue** times



Total hardware overhead → **904 bytes** per core

Hardware-optimized version → **296 bytes** per core

Experimental Setup

Experimental Setup

Sniper 6.0:

- Augmented with ACE Bit Counters

Experimental Setup

Sniper 6.0:

- Augmented with ACE Bit Counters
- Migration overhead between cores → 20 microseconds
- Scheduler quantum → 1ms, sampling for 0.1ms

Experimental Setup

Sniper 6.0:

- Augmented with ACE Bit Counters
- Migration overhead between cores → 20 microseconds
- Scheduler quantum → 1ms, sampling for 0.1ms

SPEC CPU2006:

- 1B instruction SimPoints

Experimental Setup

Sniper 6.0:

- Augmented with ACE Bit Counters
- Migration overhead between cores → 20 microseconds
- Scheduler quantum → 1ms, sampling for 0.1ms

SPEC CPU2006:

- 1B instruction SimPoints
- Based on AVF-intensity

Experimental Setup

Sniper 6.0:

- Augmented with ACE Bit Counters
- Migration overhead between cores → 20 microseconds
- Scheduler quantum → 1ms, sampling for 0.1ms

✓ L → Low AVF

SPEC CPU2006:

- 1B instruction SimPoints
- Based on AVF-intensity

Experimental Setup

Sniper 6.0:

- Augmented with ACE Bit Counters
- Migration overhead between cores → 20 microseconds
- Scheduler quantum → 1ms, sampling for 0.1ms

SPEC CPU2006:

- 1B instruction SimPoints
- Based on AVF-intensity

✓ L → Low AVF

✓ M → Medium AVF

Experimental Setup

Sniper 6.0:

- Augmented with ACE Bit Counters
- Migration overhead between cores → 20 microseconds
- Scheduler quantum → 1ms, sampling for 0.1ms

SPEC CPU2006:

- 1B instruction SimPoints
- Based on AVF-intensity

✓ L → Low AVF

✓ M → Medium AVF

✓ H → High AVF

Experimental Setup

Sniper 6.0:

- Augmented with ACE Bit Counters
- Migration overhead between cores → 20 microseconds
- Scheduler quantum → 1ms, sampling for 0.1ms

SPEC CPU2006:

- 1B instruction SimPoints
- Based on AVF-intensity

✓ L → Low AVF

✓ M → Medium AVF

✓ H → High AVF

Experimental Setup

Sniper 6.0:

- Augmented with ACE Bit Counters
- Migration overhead between cores → 20 microseconds
- Scheduler quantum → 1ms, sampling for 0.1ms

SPEC CPU2006:

- 1B instruction SimPoints
- Based on AVF-intensity

- ✓ L → Low AVF
- ✓ M → Medium AVF
- ✓ H → High AVF

4-program workloads

- 6 categories → HHHH, HHMM, HHLL, MMMM, MMLL, LLLL

Experimental Setup

Sniper 6.0:

- Augmented with ACE Bit Counters
- Migration overhead between cores → 20 microseconds
- Scheduler quantum → 1ms, sampling for 0.1ms

SPEC CPU2006:

- 1B instruction SimPoints
- Based on AVF-intensity

- ✓ L → Low AVF
- ✓ M → Medium AVF
- ✓ H → High AVF

4-program workloads

- 6 categories → HHHH, HHMM, HHLL, MMMM, MMLL, LLLL
- No duplication

Experimental Setup

Sniper 6.0:

- Augmented with ACE Bit Counters
- Migration overhead between cores → 20 microseconds
- Scheduler quantum → 1ms, sampling for 0.1ms

SPEC CPU2006:

- 1B instruction SimPoints
- Based on AVF-intensity

- ✓ L → Low AVF
- ✓ M → Medium AVF
- ✓ H → High AVF

4-program workloads

- 6 categories → HHHH, HHMM, HHLL, MMMM, MMLL, LLLL
- No duplication
- 6 mixes in each category → 36 workloads

Experimental Setup

Sniper 6.0:

- Augmented with ACE Bit Counters
- Migration overhead between cores → 20 microseconds
- Scheduler quantum → 1ms, sampling for 0.1ms

SPEC CPU2006:

- 1B instruction SimPoints
- Based on AVF-intensity

- ✓ L → Low AVF
- ✓ M → Medium AVF
- ✓ H → High AVF

4-program workloads

- 6 categories → HHHH, HHMM, HHLL, MMMM, MMLL, LLLL
- No duplication
- 6 mixes in each category → 36 workloads

HCMP → two big and two small cores (2B2S)

Evaluation

Reliability-Optimized Scheduler

- Optimizes **SSER** using our scheduling algorithm

Evaluation

Reliability-Optimized Scheduler

- Optimizes **SSER** using our scheduling algorithm

Performance-Optimized Scheduler

- Optimizes System Throughput (**STP**)

Evaluation

Reliability-Optimized Scheduler

- Optimizes **SSER** using our scheduling algorithm


Performance-Optimized Scheduler

- Optimizes System Throughput (**STP**)


Random Scheduler

- Randomly selects applications to run on the **big core(s)**

2B2S Results – Reliability




2B2S Results – Reliability



Compared to the **random** scheduler:


2B2S Results – Reliability



Compared to the **random** scheduler:

Reliability-optimized scheduler improves SSER by **32%** on average and up to **55.6%**


2B2S Results – Reliability



Compared to the **random** scheduler:

Reliability-optimized scheduler improves SSER by **32%** on average and up to **55.6%**

2B2S Results – Reliability




Compared to the **random** scheduler:

Reliability-optimized scheduler improves SSER by **32%** on average and up to **55.6%**


Performance-optimized scheduler also improves SSER by **7.3%** on average

2B2S Results – Reliability



Compared to the **performance-optimized** scheduler:


2B2S Results – Reliability




Compared to the **performance-optimized** scheduler:

Reliability-optimized scheduler improves SSER by **25.4%** on average and up to **60.2%**

2B2S Results – Performance




2B2S Results – Performance



Compared to the random scheduler:


2B2S Results – Performance



Compared to the **random** scheduler:

Reliability-optimized scheduler **does not degrade** STP on average


2B2S Results – Performance



Compared to the **random** scheduler:

Reliability-optimized scheduler **does not degrade** STP on average

2B2S Results – Performance




Compared to the **random** scheduler:

Reliability-optimized scheduler **does not degrade** STP on average


Performance-optimized scheduler improves STP by **6.3%** on average and up to **18.7%**

2B2S Results – Performance



Compared to the **performance-optimized** scheduler:


2B2S Results – Performance




Compared to the **performance-optimized** scheduler:

Reliability-optimized scheduler **degrades** STP by **6.3%** on average and up to **18.7%**

2B2S Results – by Workload Category




2B2S Results – by Workload Category



Highest improvement in reliability for HHLL (53.2%)


2B2S Results – by Workload Category



Highest improvement in reliability for **HHLL (53.2%)**

- Next for HHMM (15.8%) and MMLL (22.7%)

2B2S Results – by Workload Category



Highest improvement in reliability for **HHLL (53.2%)**

- Next for HHMM (15.8%) and MMLL (22.7%)

- Significant improvement for HHHH (10.2%), MMMM (7.3%) and LLLL (10.6%)

Also in the Paper

1. Detailed hardware overhead analysis
2. Asymmetric HCMPs → 1B3S and 3B1S
3. Changing core count → 1B1S, 2B2S and 4B4S
4. Lowering small core frequency from 2.66 GHz → 1.33 GHz
5. Impact on power consumption → Improves by 6.2%
6. Impact of changing sampling and scheduler quanta

Conclusions

Conclusions

1. Reliability characteristics of applications vary on big and small cores

Conclusions

1. Reliability characteristics of applications vary on big and small cores

- Provides an opportunity for improving reliability on heterogeneous (multicore) processors

Conclusions

1. Reliability characteristics of applications vary on big and small cores
 - Provides an opportunity for improving reliability on heterogeneous (multicore) processors
2. Proposed a dynamic reliability-aware scheduler
 - Improved reliability by 25.4% with 6.3% performance degradation

Conclusions

1. Reliability characteristics of applications vary on big and small cores
 - Provides an opportunity for improving reliability on heterogeneous (multicore) processors
2. Proposed a dynamic reliability-aware scheduler
 - Improved reliability by 25.4% with 6.3% performance degradation
3. Introduced System-level Soft Error Rate (SSER), a novel soft error reliability metric for multiprogrammed workloads

Reliability-Aware Scheduling on Heterogeneous Multicore Processors

Thank you!

Questions?

Ajeya Naithani, Stijn Eyerman, and Lieven Eeckhout