

Precise Runahead Execution

Ajeya Naithani, Josue Feliu,
Almutaz Adileh, Lieven Eeckhout




UNIVERSITAT
POLITÈCNICA
DE VALÈNCIA

Full-Window Stalls Degrade Performance

Full-Window Stalls Degrade Performance


_____→ time

Full-Window Stalls Degrad Performance




Full-Window Stalls Degrade Performance

L → Loads



Full-Window Stalls Degrade Performance

L → Loads




Full-Window Stalls Degrade Performance

L → Loads




Full-Window Stalls Degrade Performance


L → Loads



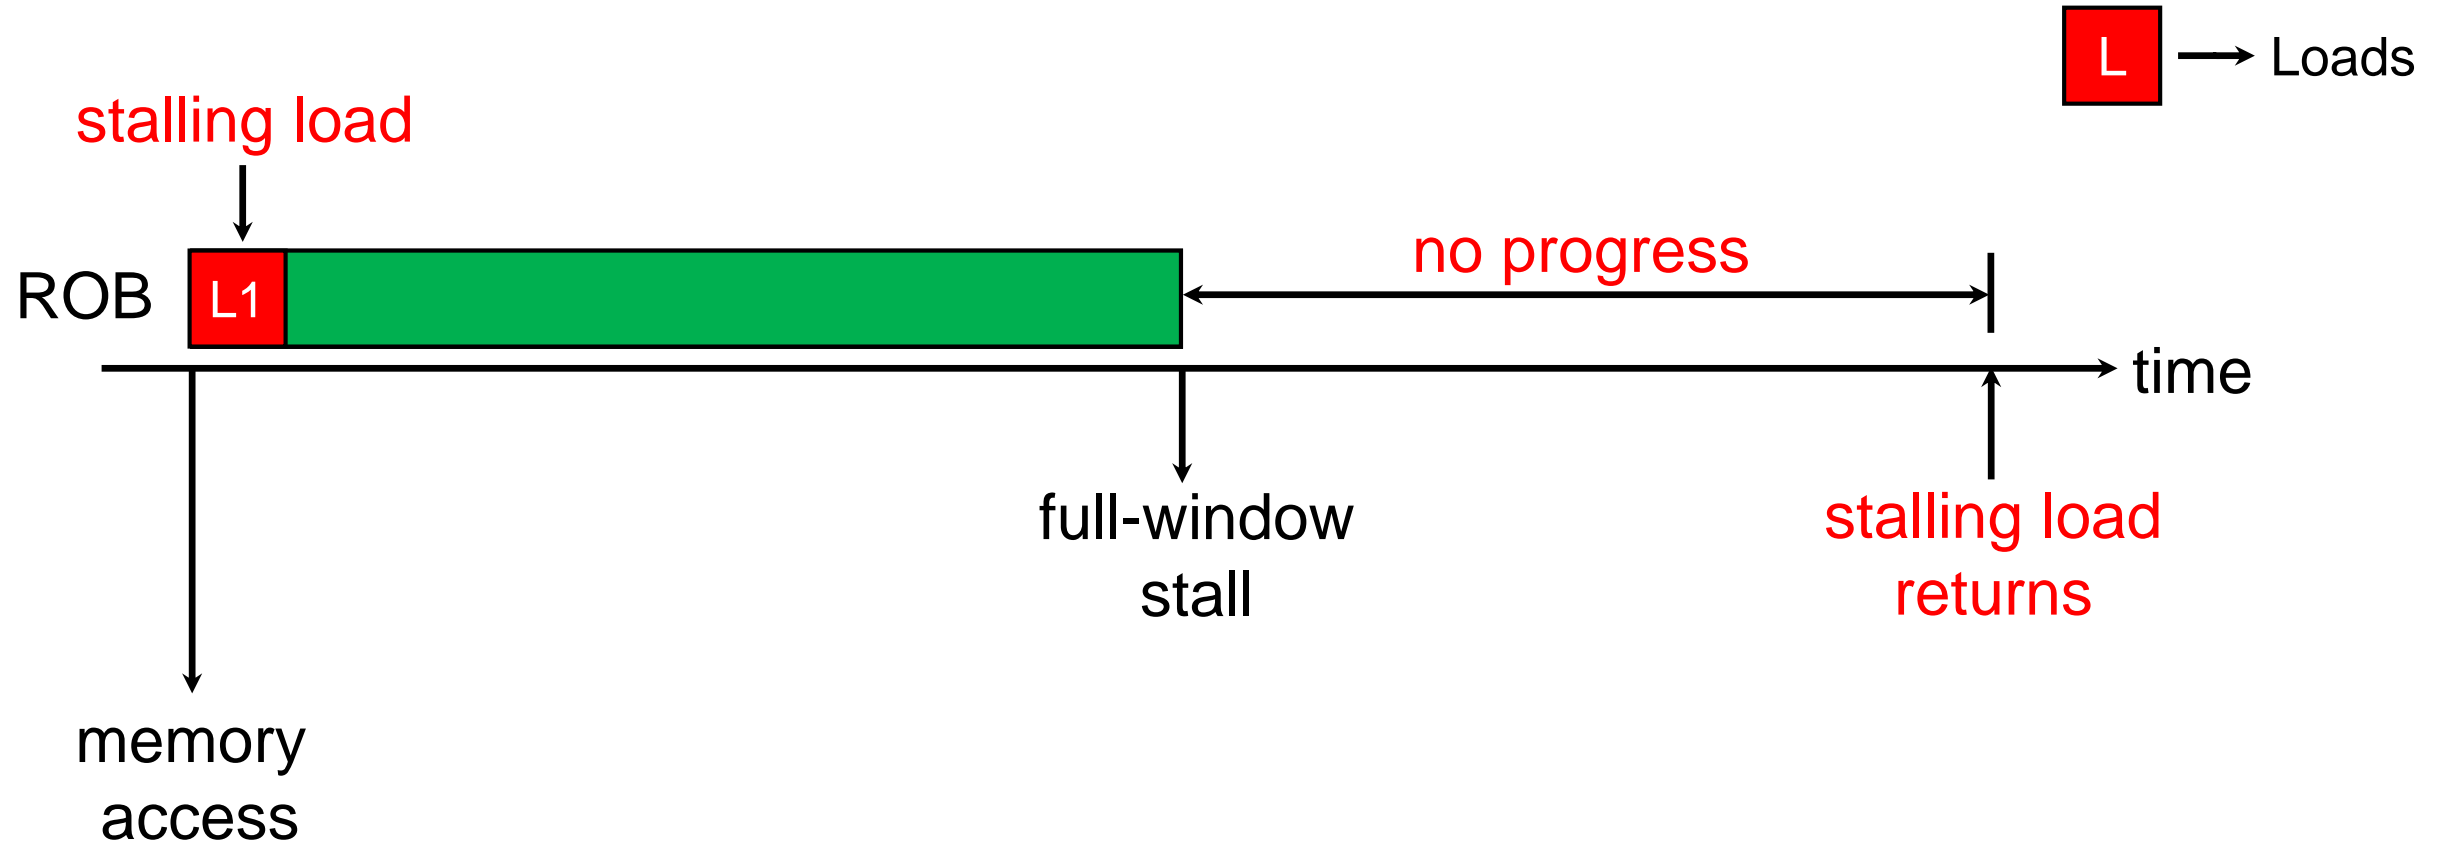
Full-Window Stalls Degrade Performance



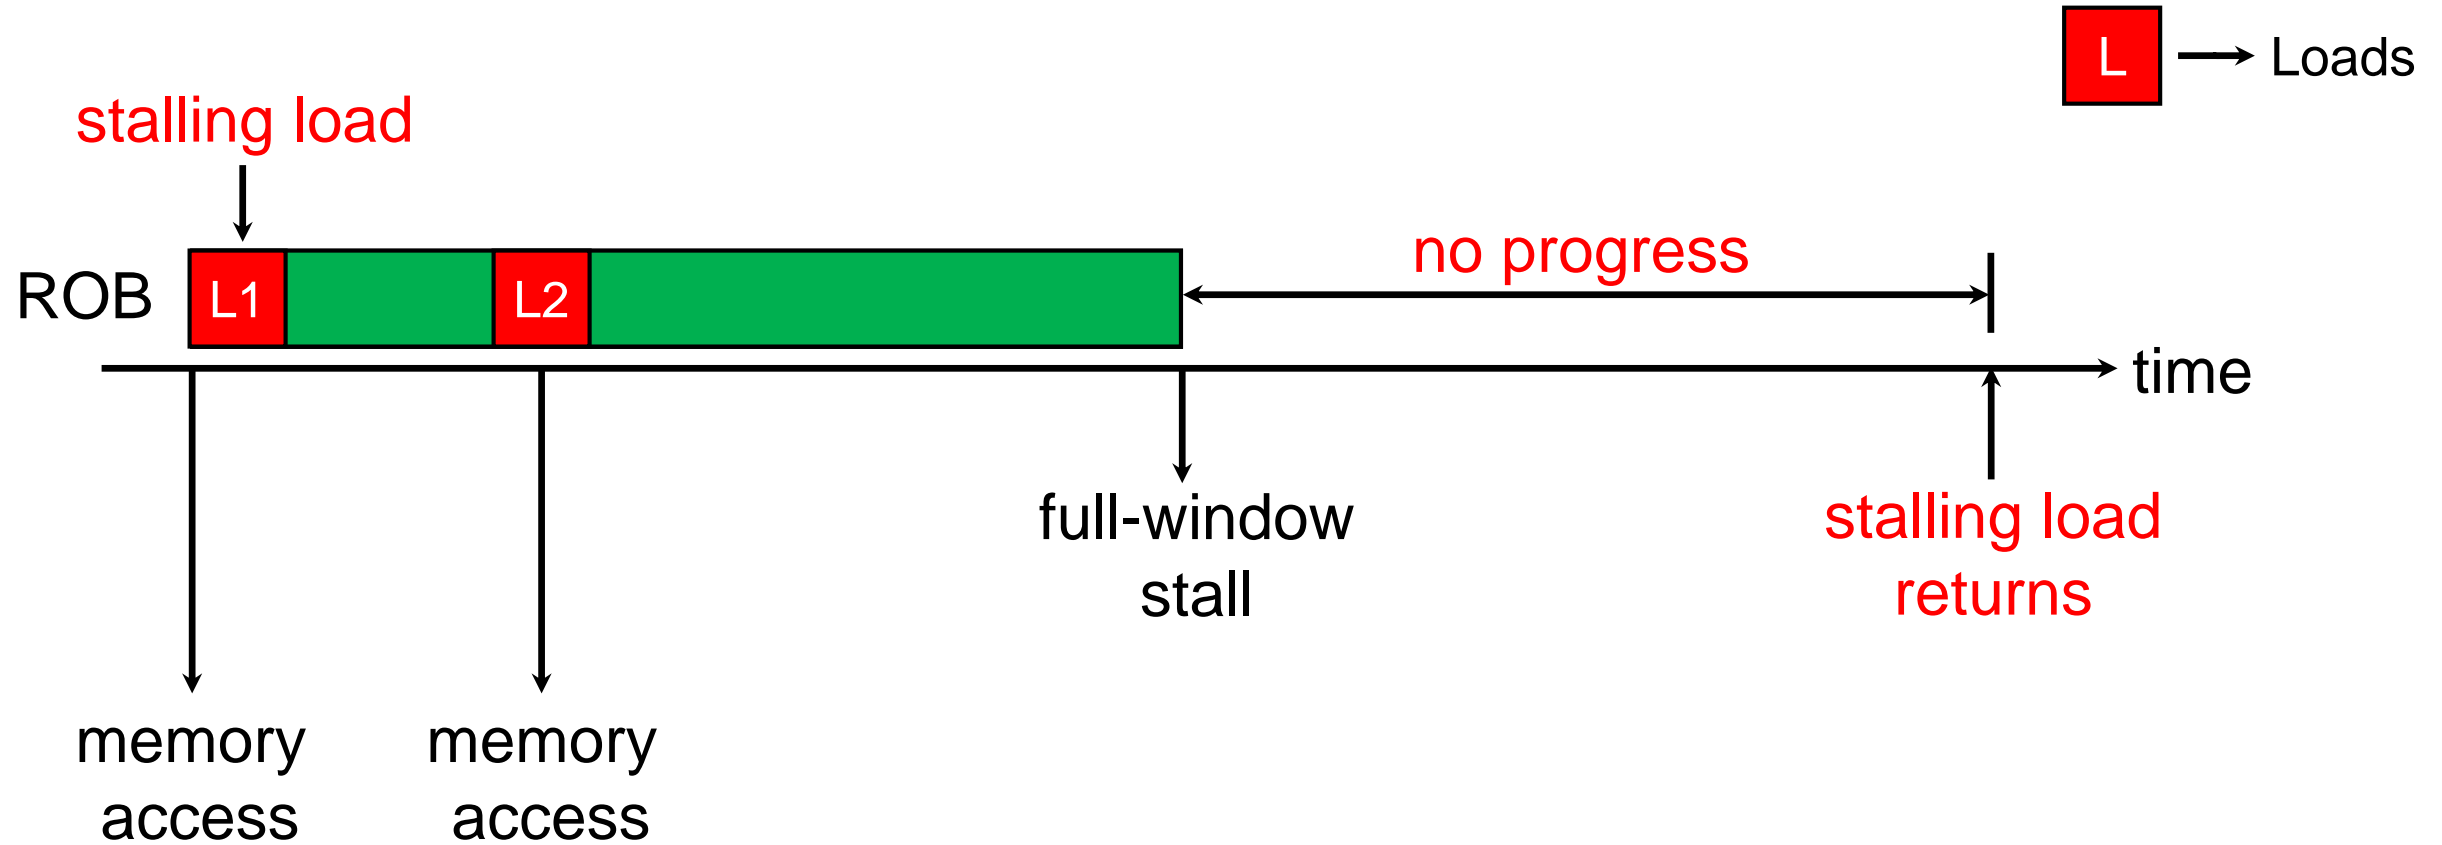
Full-Window Stalls Degrade Performance




Full-Window Stalls Degrade Performance




Full-Window Stalls Degrade Performance




Full-Window Stalls Degrade Performance




Full-Window Stalls Degrade Performance




Runahead Execution Prefetches under a Full-Window Stall




Runahead Execution Prefetches under a Full-Window Stall




Runahead Execution Prefetches under a Full-Window Stall




Runahead Execution Prefetches under a Full-Window Stall



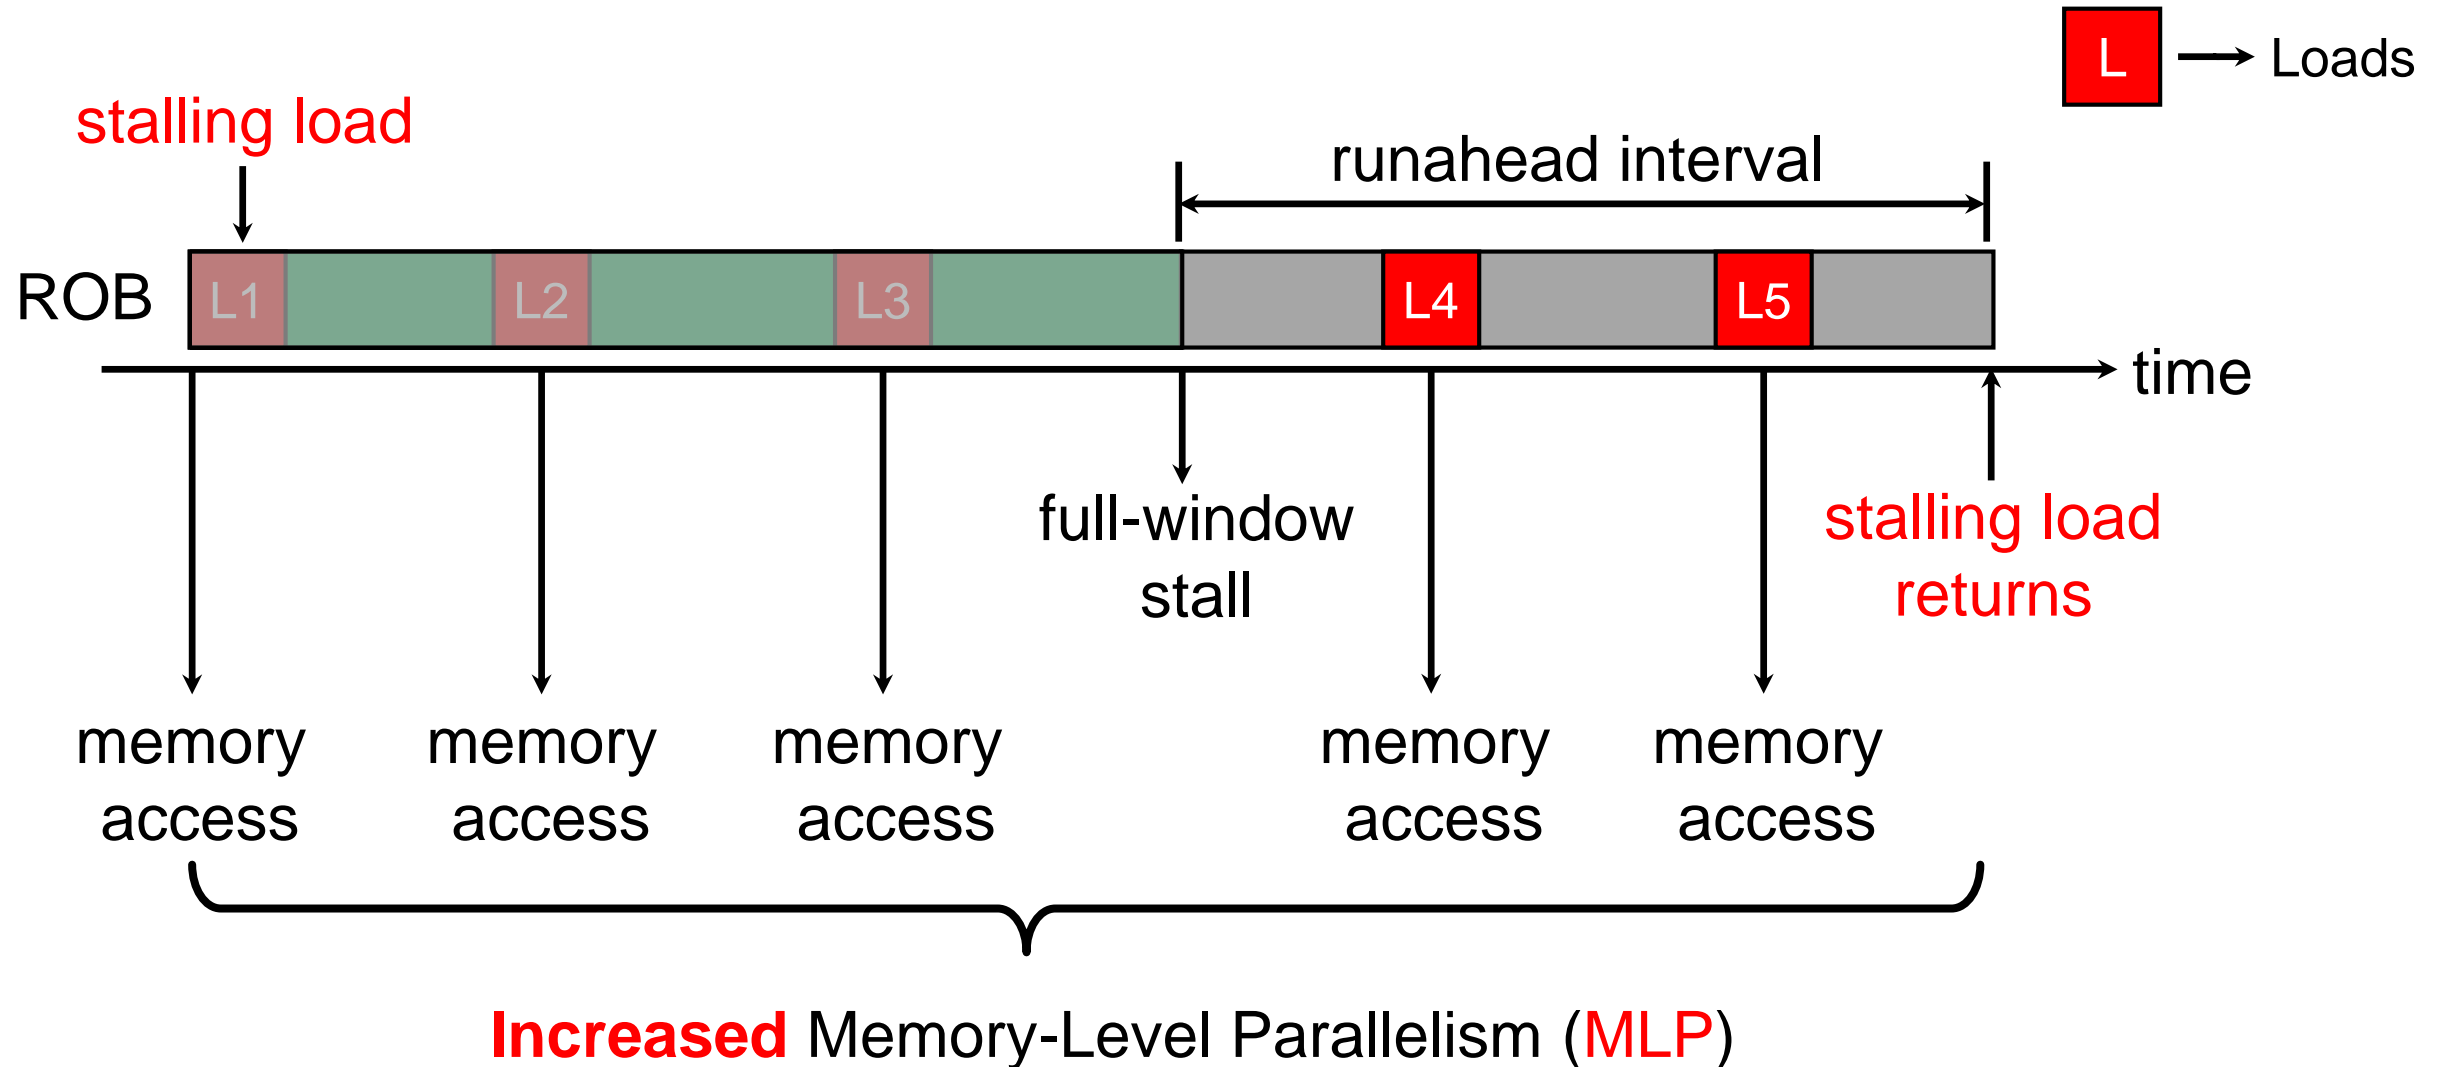
Runahead Execution Prefetches under a Full-Window Stall



Runahead Execution Prefetches under a Full-Window Stall



Runahead Execution Prefetches under a Full-Window Stall

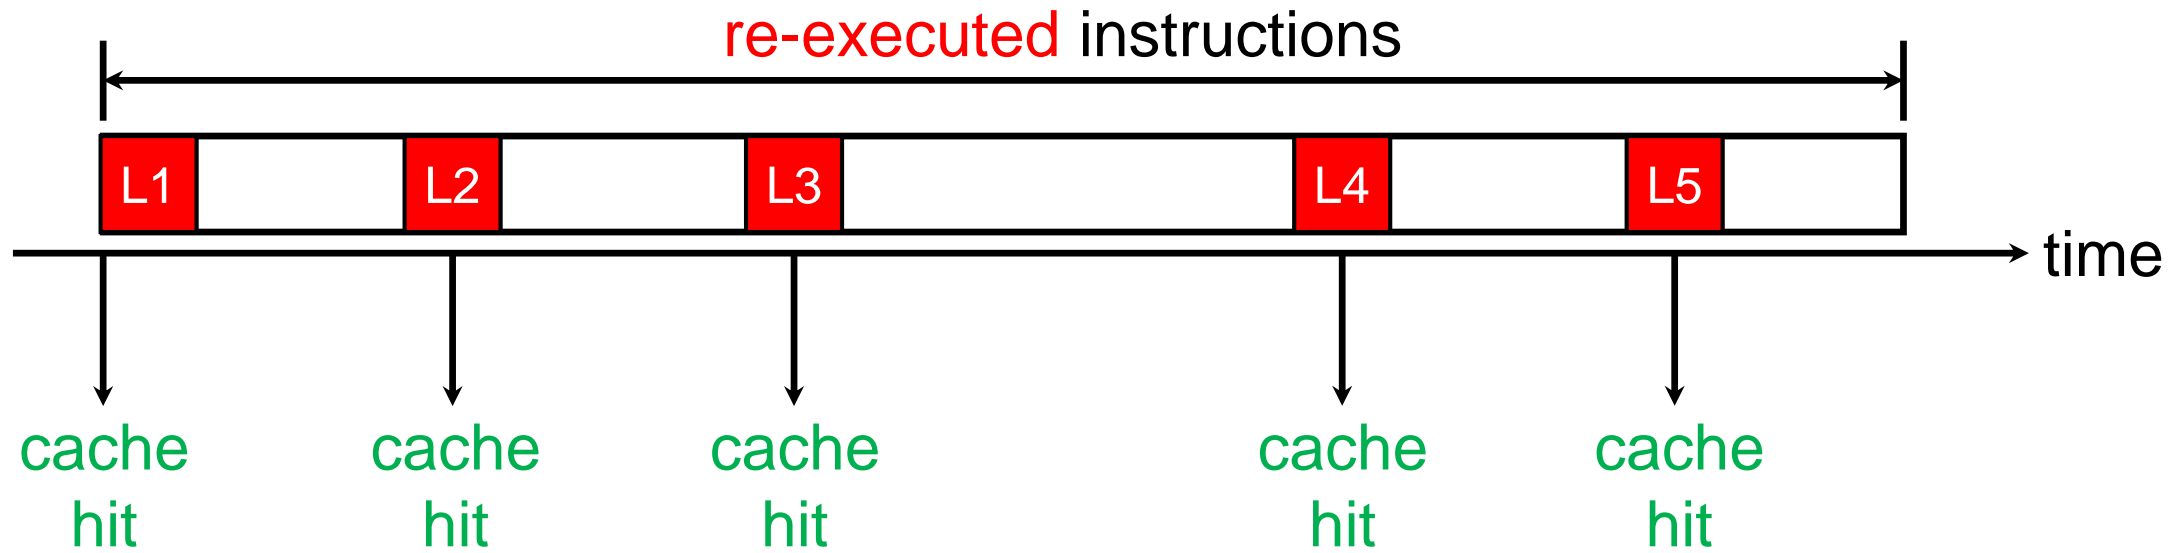


Runahead Execution Re-Executes All Instructions

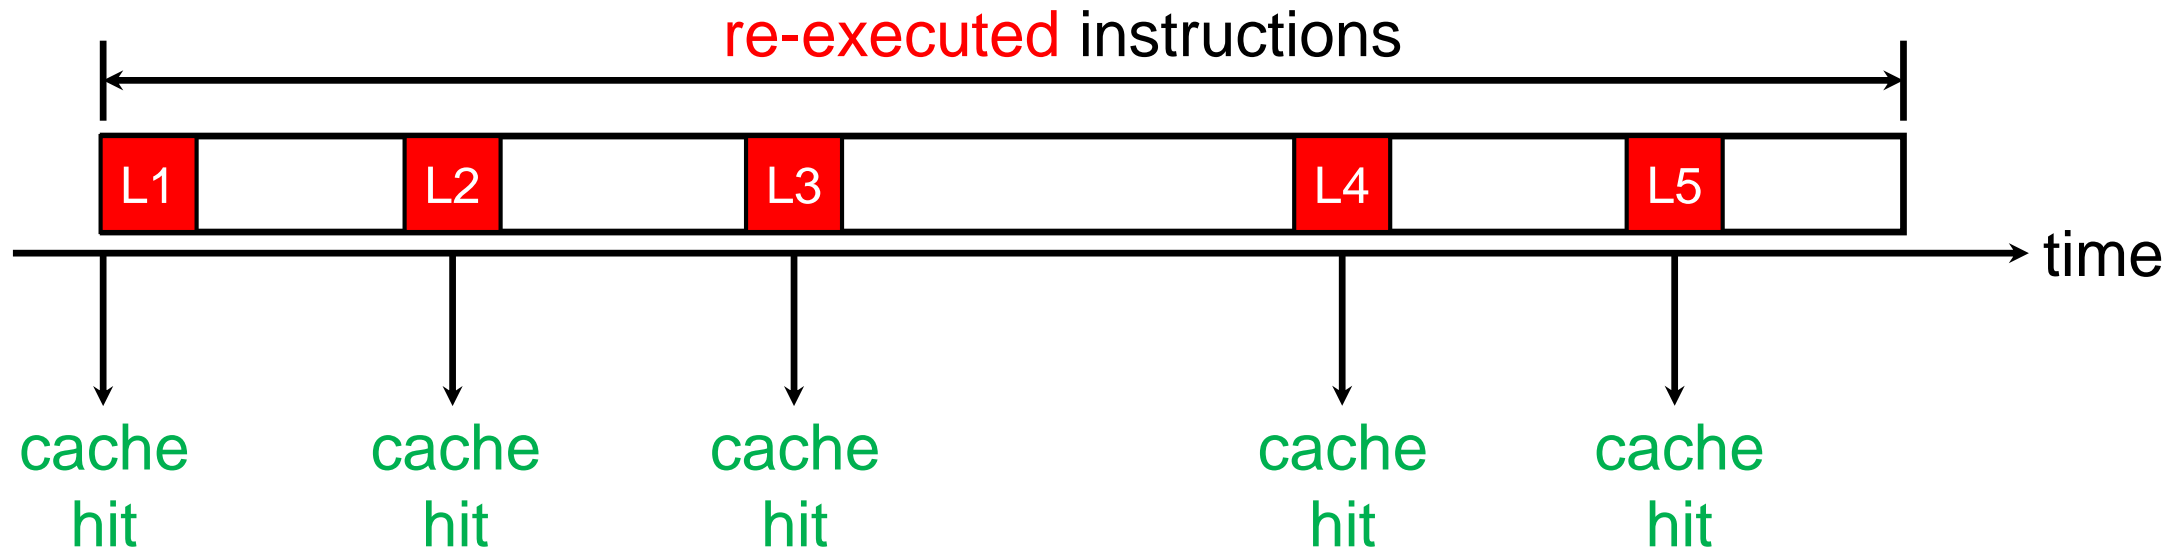
Runahead Execution Re-Executes All Instructions

—————→ time

Runahead Execution Re-Executes All Instructions

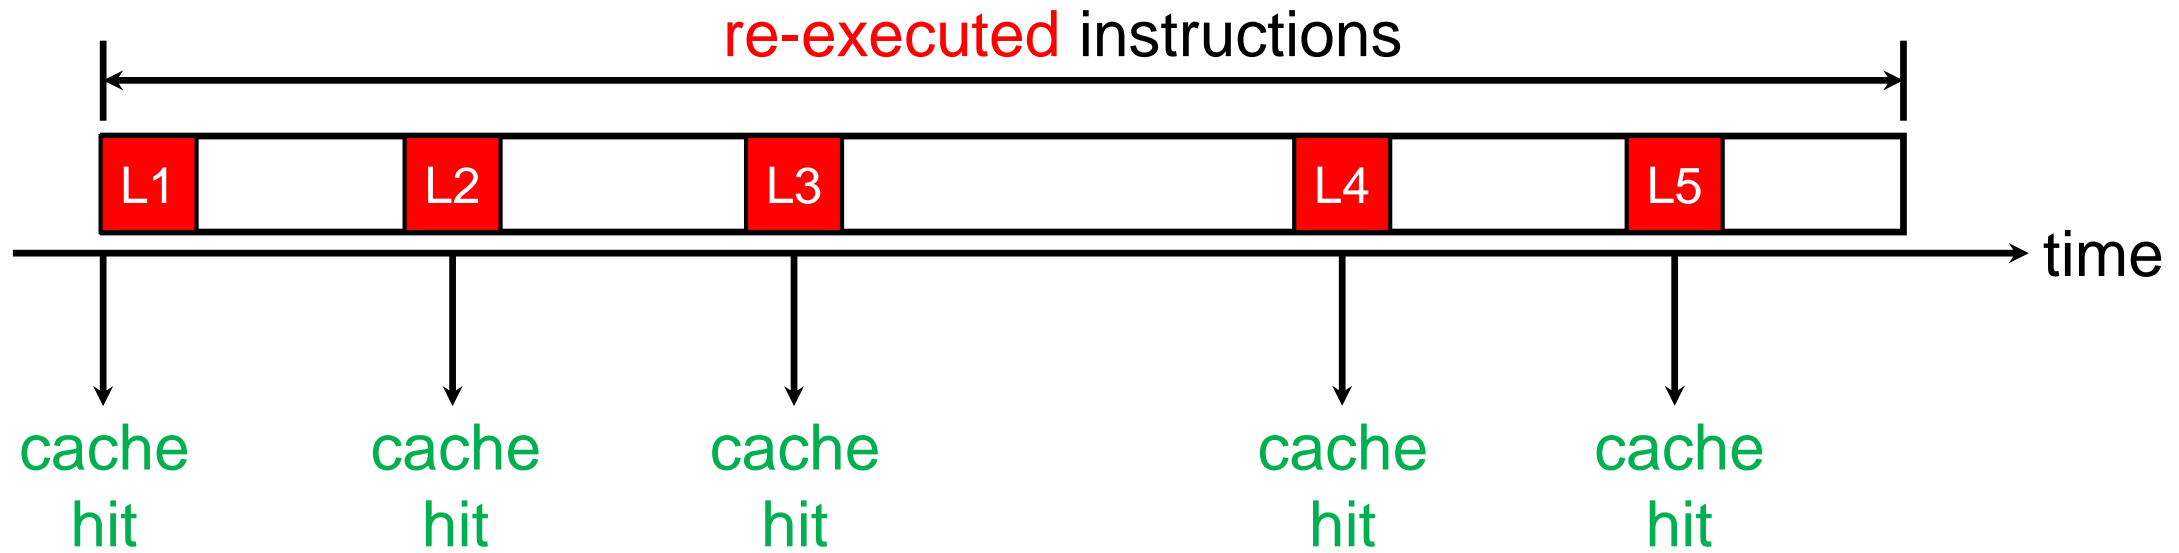


Runahead Execution Re-Executes All Instructions



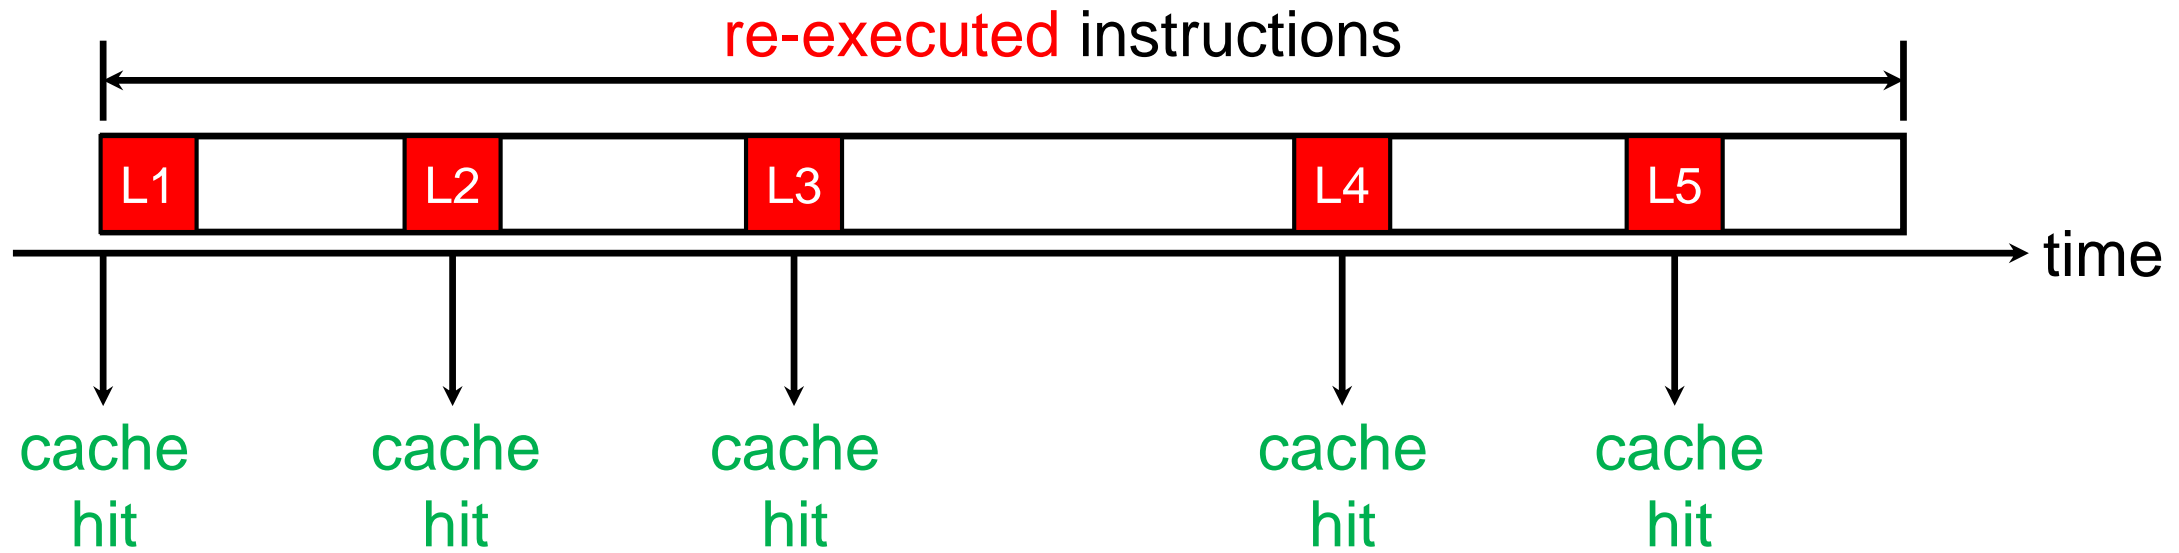
fetch

Runahead Execution Re-Executes All Instructions



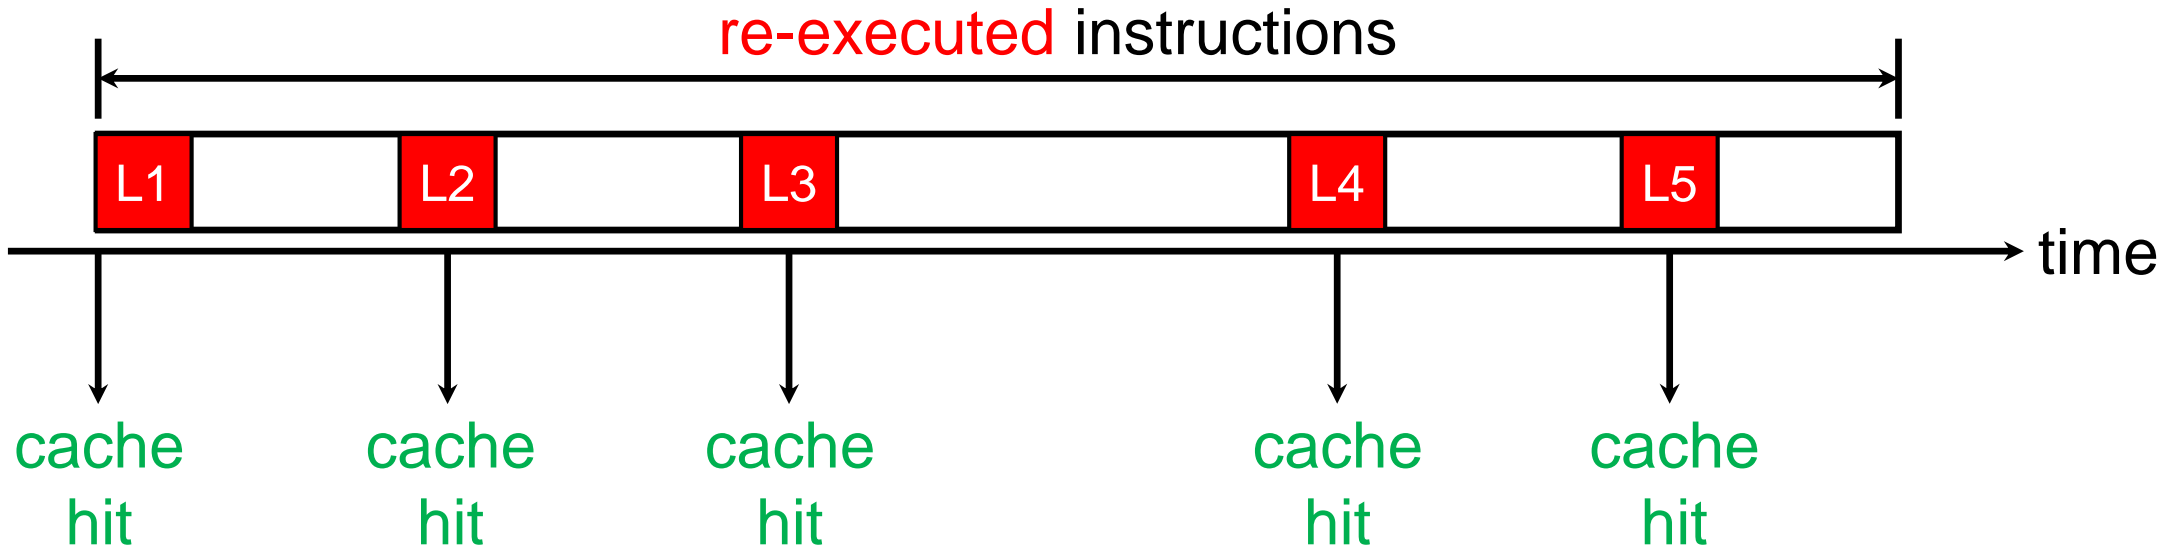
fetch → decode

Runahead Execution Re-Executes All Instructions



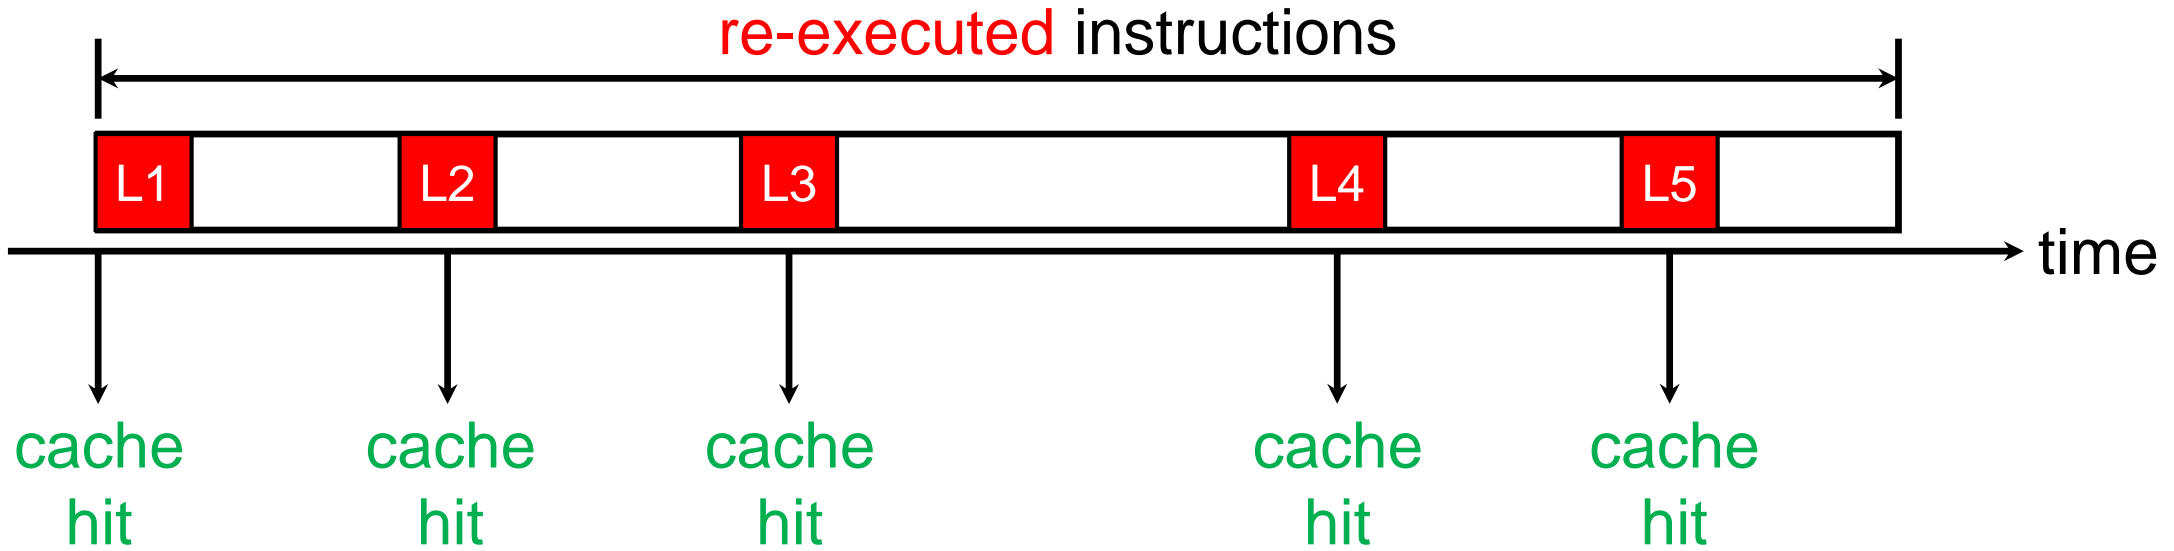
fetch → decode → rename

Runahead Execution Re-Executes All Instructions



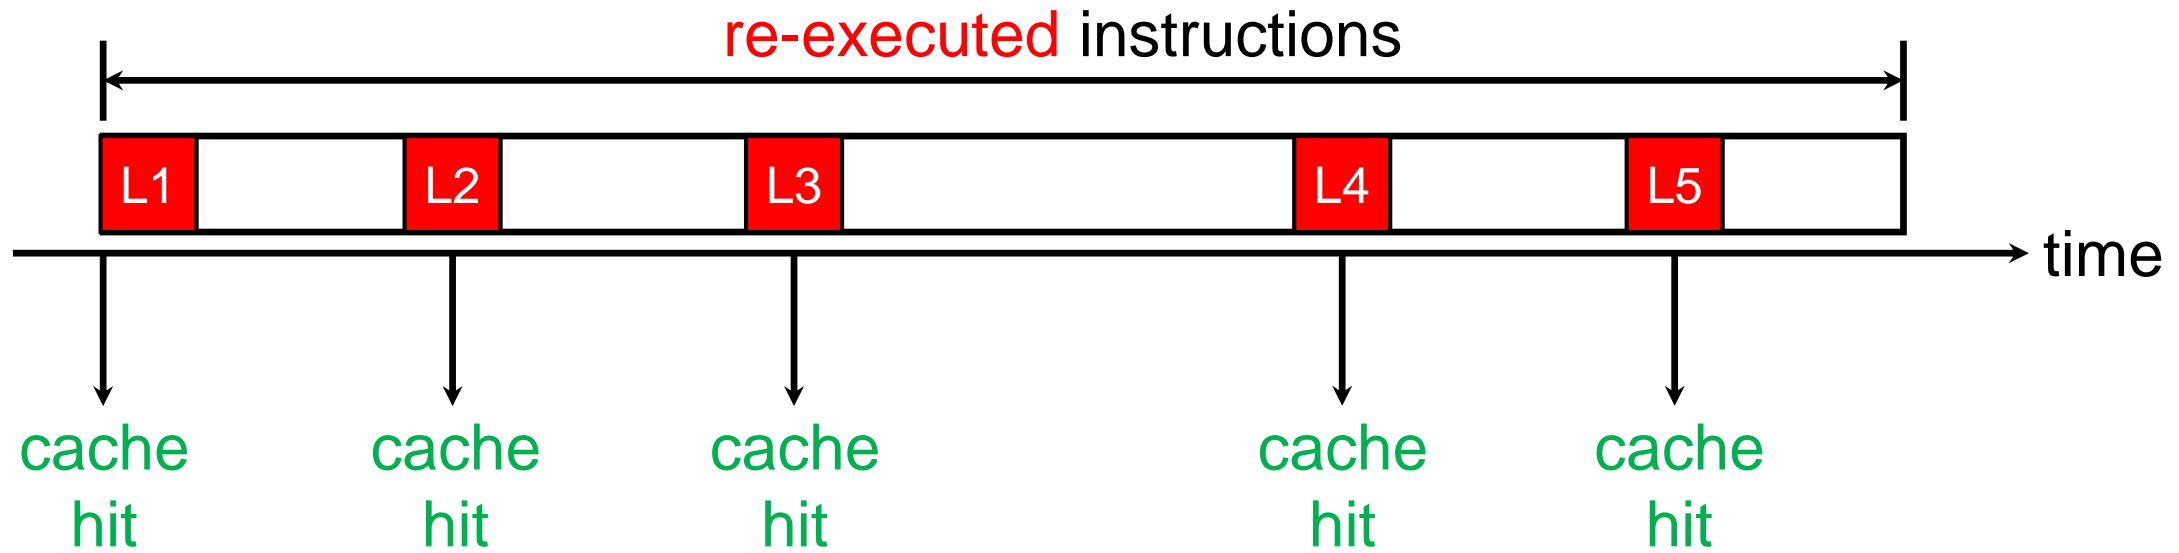
fetch → decode → rename → dispatch

Runahead Execution Re-Executes All Instructions



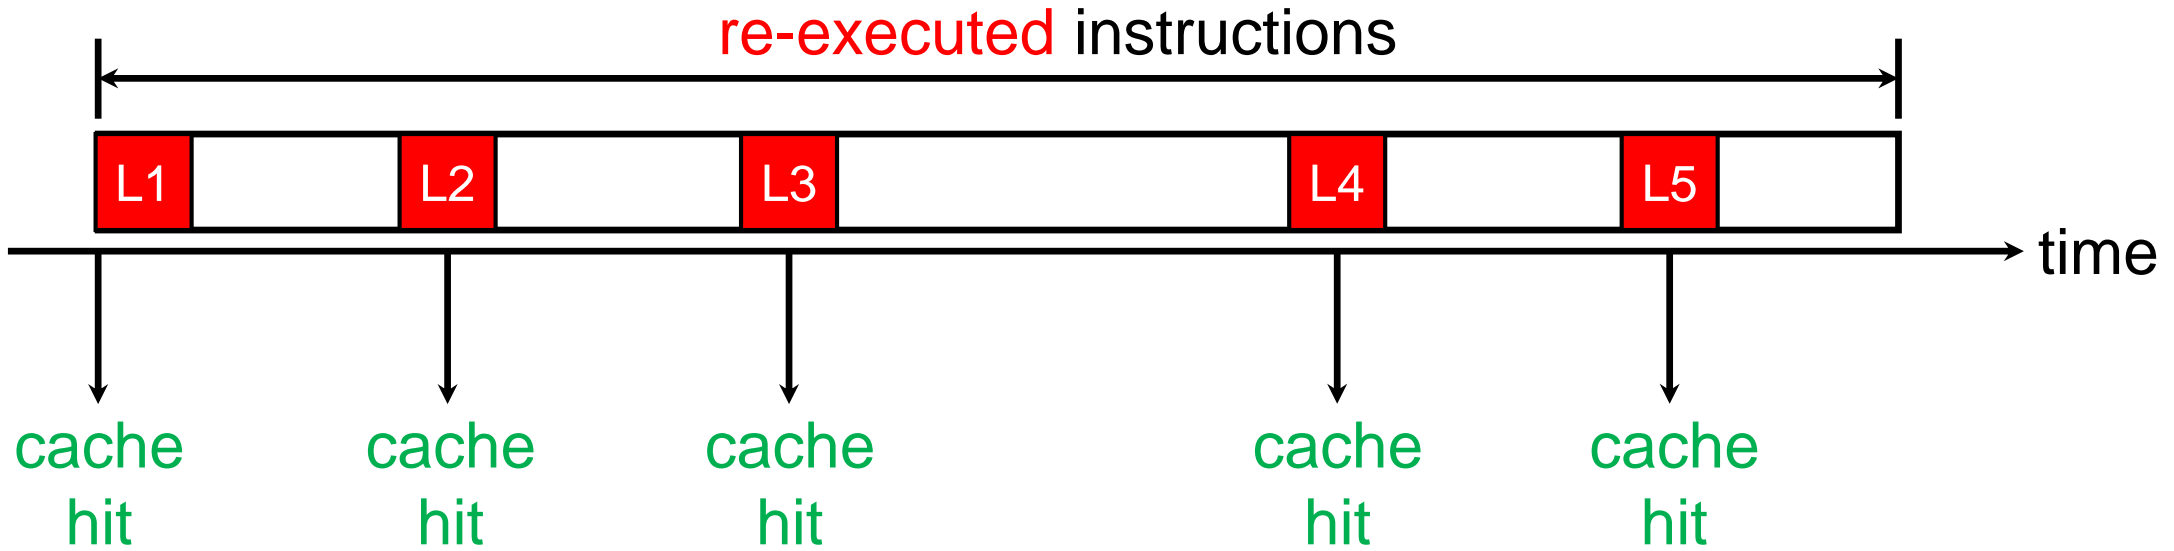
fetch → decode → rename → dispatch → issue

Runahead Execution Re-Executes All Instructions




fetch → decode → rename → dispatch → issue → execute

Runahead Execution Re-Executes All Instructions




fetch → decode → rename → dispatch → issue → execute → commit


Runahead Buffer Finds Blocking Chain in the ROB




Runahead Buffer Finds Blocking Chain in the ROB




Runahead Buffer Finds Blocking Chain in the ROB




Runahead Buffer Finds Blocking Chain in the ROB




Runahead Buffer Executes Blocking Chain Speculatively




Runahead Buffer Executes Blocking Chain Speculatively




Runahead Buffer Executes Blocking Chain Speculatively



Runahead Buffer Executes Blocking Chain Speculatively



Runahead Buffer Executes Blocking Chain Speculatively




Runahead Buffer Re-Executes the Window


Runahead Buffer Re-Executes the Window

_____→ time

Runahead Buffer Re-Executes the Window



Runahead Buffer Re-Executes the Window



fetch → decode → rename → dispatch → issue → execute → commit

Runahead Techniques Relative to OoO Core

Runahead Techniques Relative to OoO Core

Runahead
execution*

Runahead
buffer**

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]


Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB		

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB		

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]

Flushing and Re-Filling Incur High Overhead

Flushing and Re-Filling Incur High Overhead

- Front-end refill = 8 cycles

Flushing and Re-Filling Incur High Overhead

- Front-end refill = 8 cycles
- ROB = 192, width = 4
ROB fill time = 48 cycles

Flushing and Re-Filling Incur High Overhead


- Front-end refill = 8 cycles
- ROB = 192, width = 4
ROB fill time = 48 cycles
- Total overhead = 56 cycles

Flushing and Re-Filling Incur High Overhead


- Front-end refill = 8 cycles
- ROB = 192, width = 4
ROB fill time = 48 cycles
- Total overhead = 56 cycles

Runahead causes a pipeline bubble of 56 cycles **per invocation**

Flushing and Re-Filling Incur High Overhead




Flushing and Re-Filling Incur High Overhead




runahead: 15.9%

Flushing and Re-Filling Incur High Overhead



runahead: 15.9%


Flushing and Re-Filling Incur High Overhead



runahead: 15.9%

runahead without flushing: 22.7%

Flushing and Re-Filling Incur High Overhead



runahead: 15.9%

runahead without flushing: 22.7%

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals		

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals	✗	

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals	✗	✗

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals	✗	✗
Instructions executed		

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals	✗	✗
Instructions executed	All	

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals	✗	✗
Instructions executed	All	Only one slice

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]


Runahead Techniques Provide Limited Prefetch Coverage

- Runahead execution: Executes **useless** instructions


Runahead Techniques Provide Limited Prefetch Coverage

- Runahead execution: Executes **useless** instructions
- Runahead buffer: High coverage for **only one slice**

Only One Load does not Lead to Majority of Memory Accesses




Only One Load does not Lead to Majority of Memory Accesses




Most of the long-latency loads during runahead **differ from the stalling load**

Applications Access Memory through Multiple Slices



Applications Access Memory through Multiple Slices



There are more than eight unique load instructions accessing memory during each runahead interval

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals	✗	✗
Instructions executed	All	Only one slice

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals	✗	✗
Instructions executed	All	Only one slice
Performance		

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals	✗	✗
Instructions executed	All	Only one slice
Performance	High ↑	

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals	✗	✗
Instructions executed	All	Only one slice
Performance	High ↑	High ↑

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals	✗	✗
Instructions executed	All	Only one slice
Performance	High ↑	High ↑
Energy-Efficiency		

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals	✗	✗
Instructions executed	All	Only one slice
Performance	High ↑	High ↑
Energy-Efficiency	Low ↓	

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**
Flush ROB	✓	✓
Short intervals	✗	✗
Instructions executed	All	Only one slice
Performance	High ↑	High ↑
Energy-Efficiency	Low ↓	Same

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**	
Flush ROB	✓	✓	
Short intervals	✗	✗	
Instructions executed	All	Only one slice	
Performance	High ↑	High ↑	
Energy-Efficiency	Low ↓	Same	

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**	
Flush ROB	✓	✓	✗
Short intervals	✗	✗	
Instructions executed	All	Only one slice	
Performance	High ↑	High ↑	
Energy-Efficiency	Low ↓	Same	

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**	
Flush ROB	✓	✓	✗
Short intervals	✗	✗	✓
Instructions executed	All	Only one slice	
Performance	High ↑	High ↑	
Energy-Efficiency	Low ↓	Same	

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**	
Flush ROB	✓	✓	✗
Short intervals	✗	✗	✓
Instructions executed	All	Only one slice	All slices
Performance	High ↑	High ↑	
Energy-Efficiency	Low ↓	Same	

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**	
Flush ROB	✓	✓	✗
Short intervals	✗	✗	✓
Instructions executed	All	Only one slice	All slices
Performance	High ↑	High ↑	Very high ↑↑
Energy-Efficiency	Low ↓	Same	

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**	
Flush ROB	✓	✓	✗
Short intervals	✗	✗	✓
Instructions executed	All	Only one slice	All slices
Performance	High ↑	High ↑	Very high ↑↑
Energy-Efficiency	Low ↓	Same	High ↑

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

Runahead Techniques Relative to OoO Core

	Runahead execution*	Runahead buffer**	Precise runahead***
Flush ROB	✓	✓	✗
Short intervals	✗	✗	✓
Instructions executed	All	Only one slice	All slices
Performance	High ↑	High ↑	Very high ↑↑
Energy-Efficiency	Low ↓	Same	High ↑

*[Mutlu et al. ISCA' 05]

**[Hashemi et al. MICRO' 15]

***[Naithani et al. HPCA' 20]

Precise Runahead Execution (PRE)

Precise Runahead Execution (PRE)

Key insight: There are sufficient resources to (start) run ahead without flushing the ROB

Precise Runahead Execution (PRE)

Key insight: There are sufficient resources to (start) run ahead without flushing the ROB

When running ahead:

Precise Runahead Execution (PRE)

Key insight: There are sufficient resources to (start) run ahead without flushing the ROB

When running ahead:

1. Executes only useful instructions in runahead mode


Precise Runahead Execution (PRE)

Key insight: There are sufficient resources to (start) run ahead without flushing the ROB


When running ahead:

1. Executes only useful instructions in runahead mode
2. Efficiently manages microarchitectural resources


Processor Resources at Full-Window Stall




Processor Resources at Full-Window Stall




Processor Resources at Full-Window Stall




Processor Resources at Full-Window Stall




Processor Resources at Full-Window Stall



Processor Resources at Full-Window Stall




Processor Resources at Full-Window Stall




**There are sufficient resources to start runahead
without flushing the ROB**

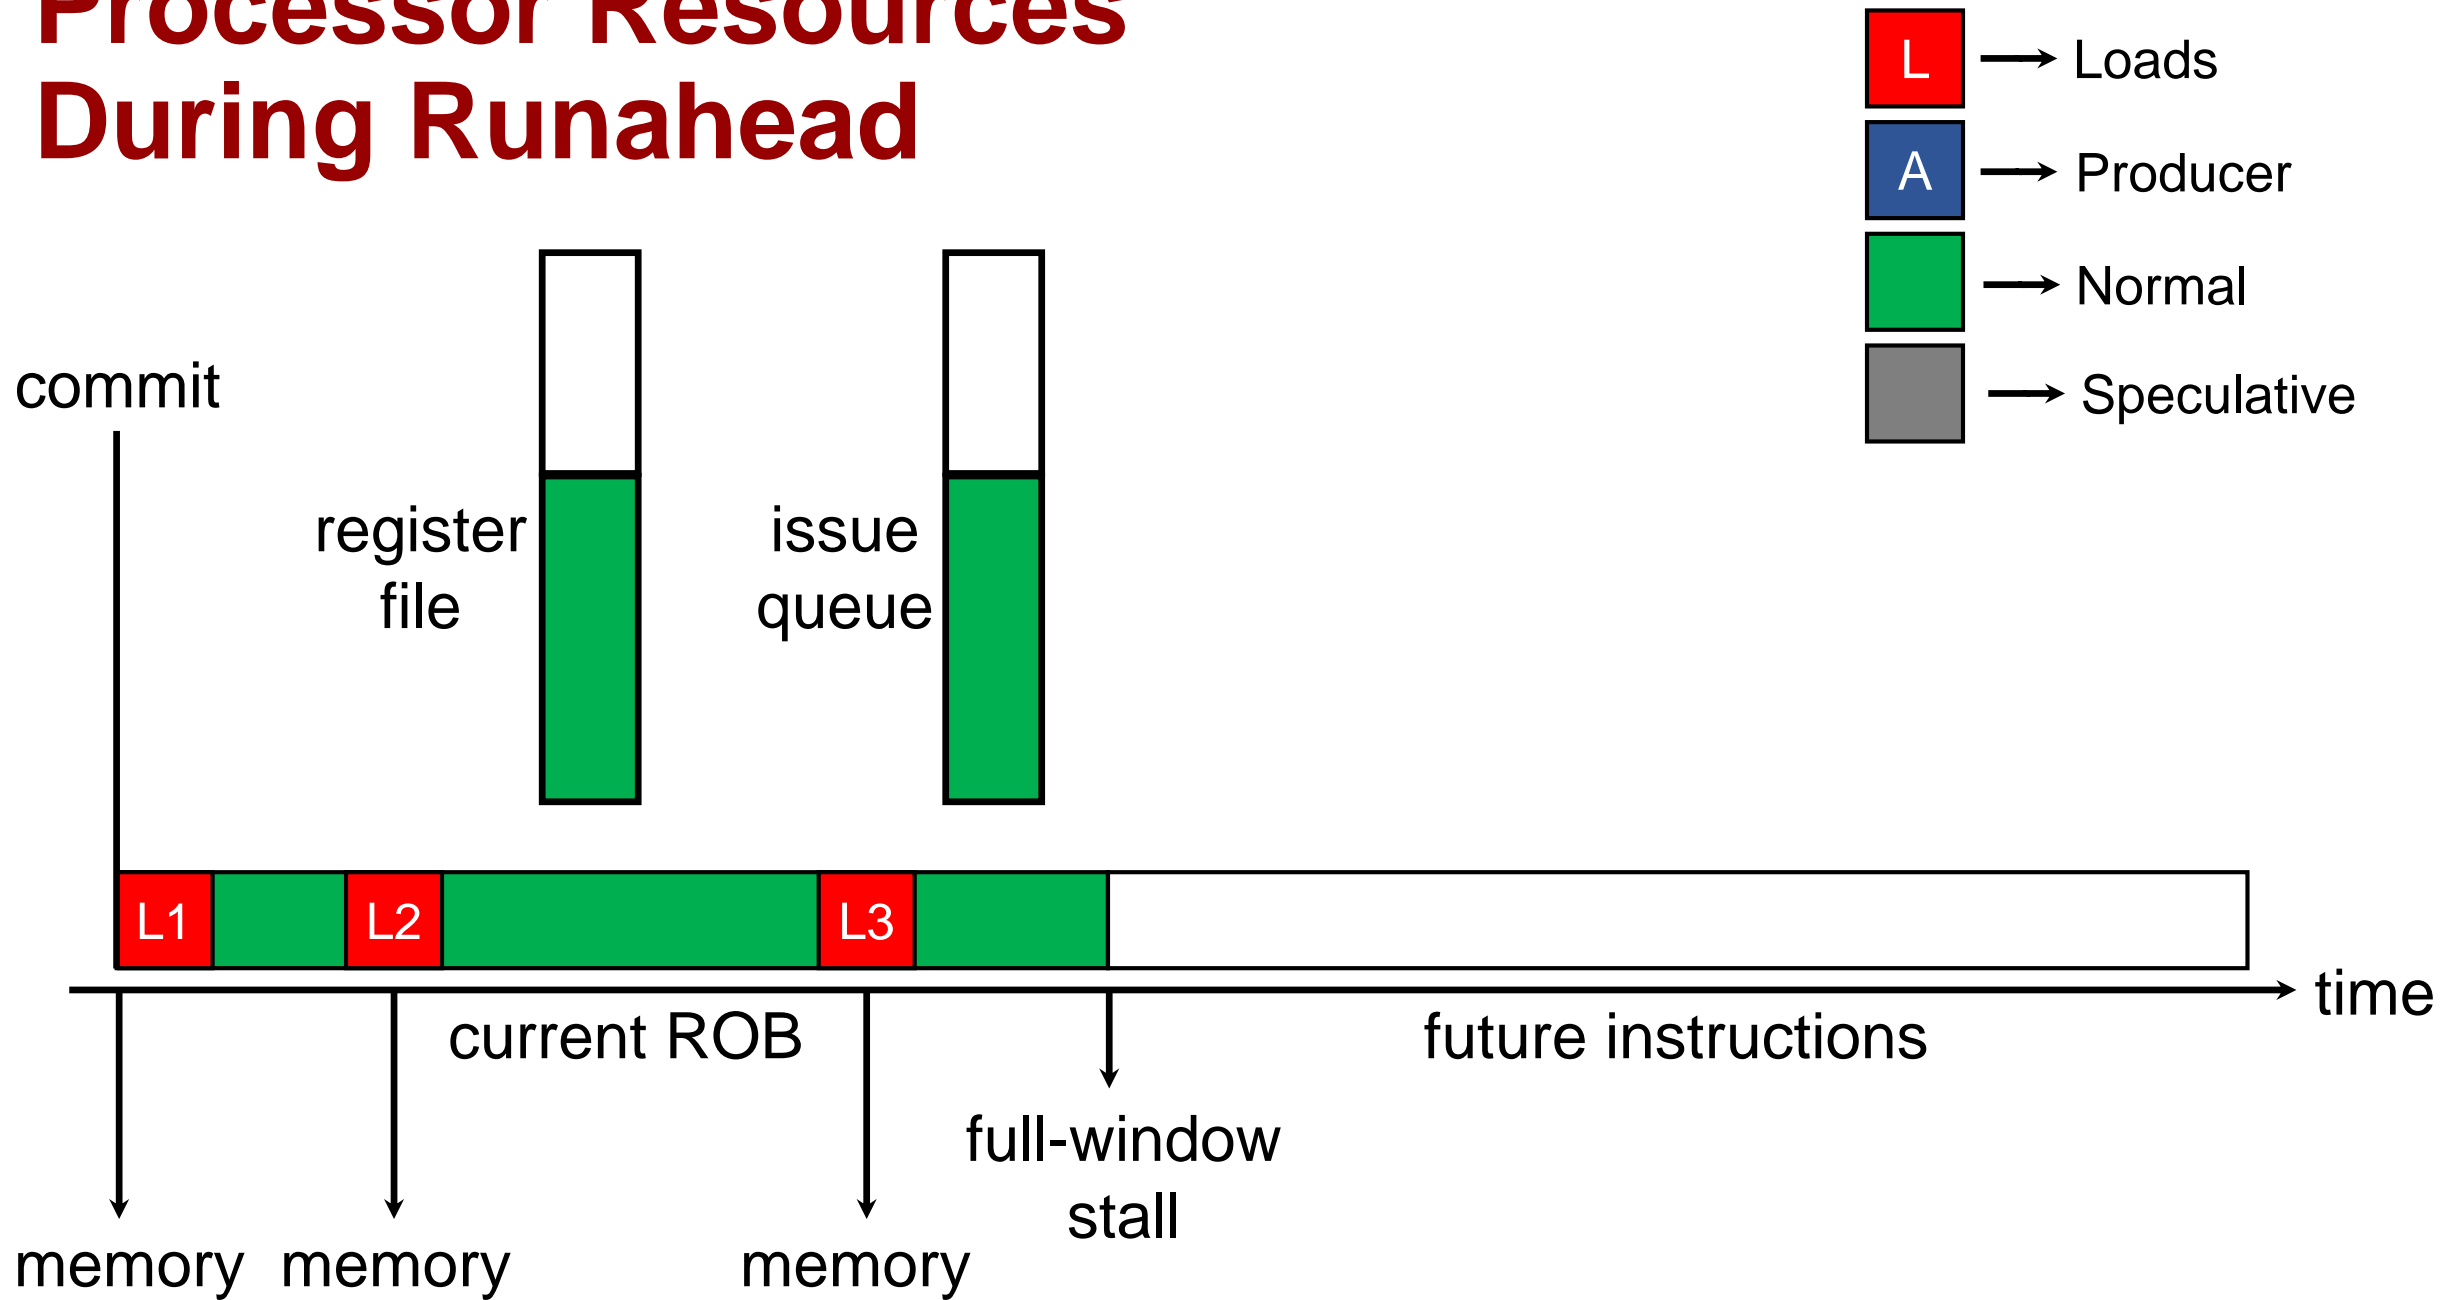
Processor Resources During Runahead



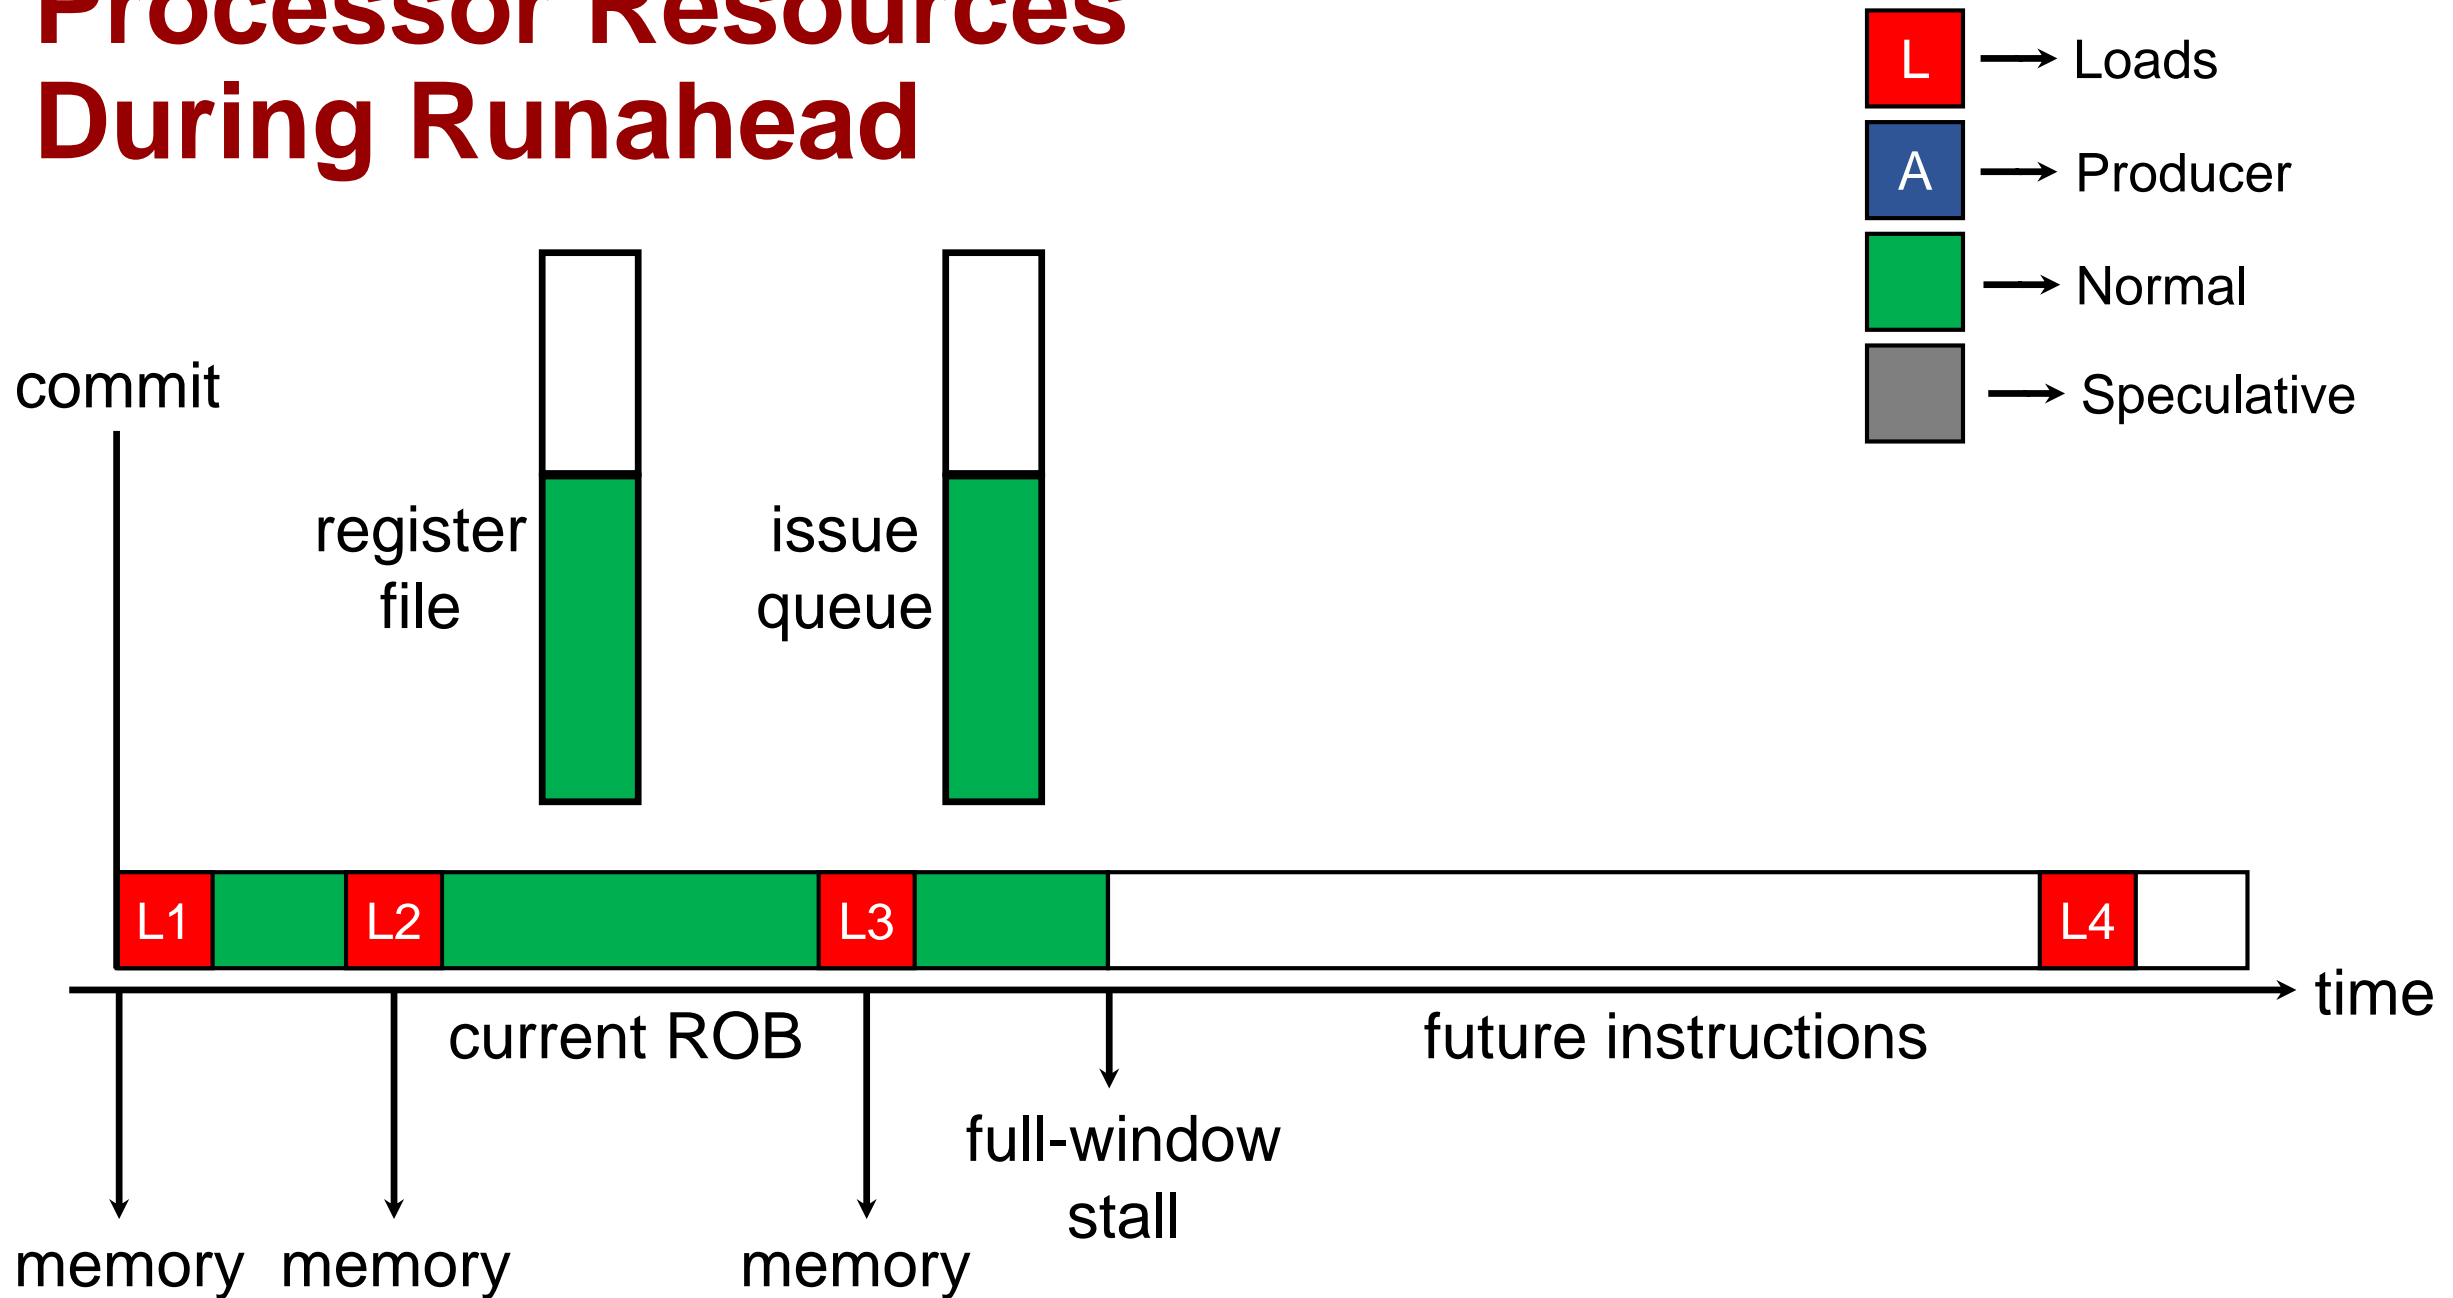
Processor Resources During Runahead



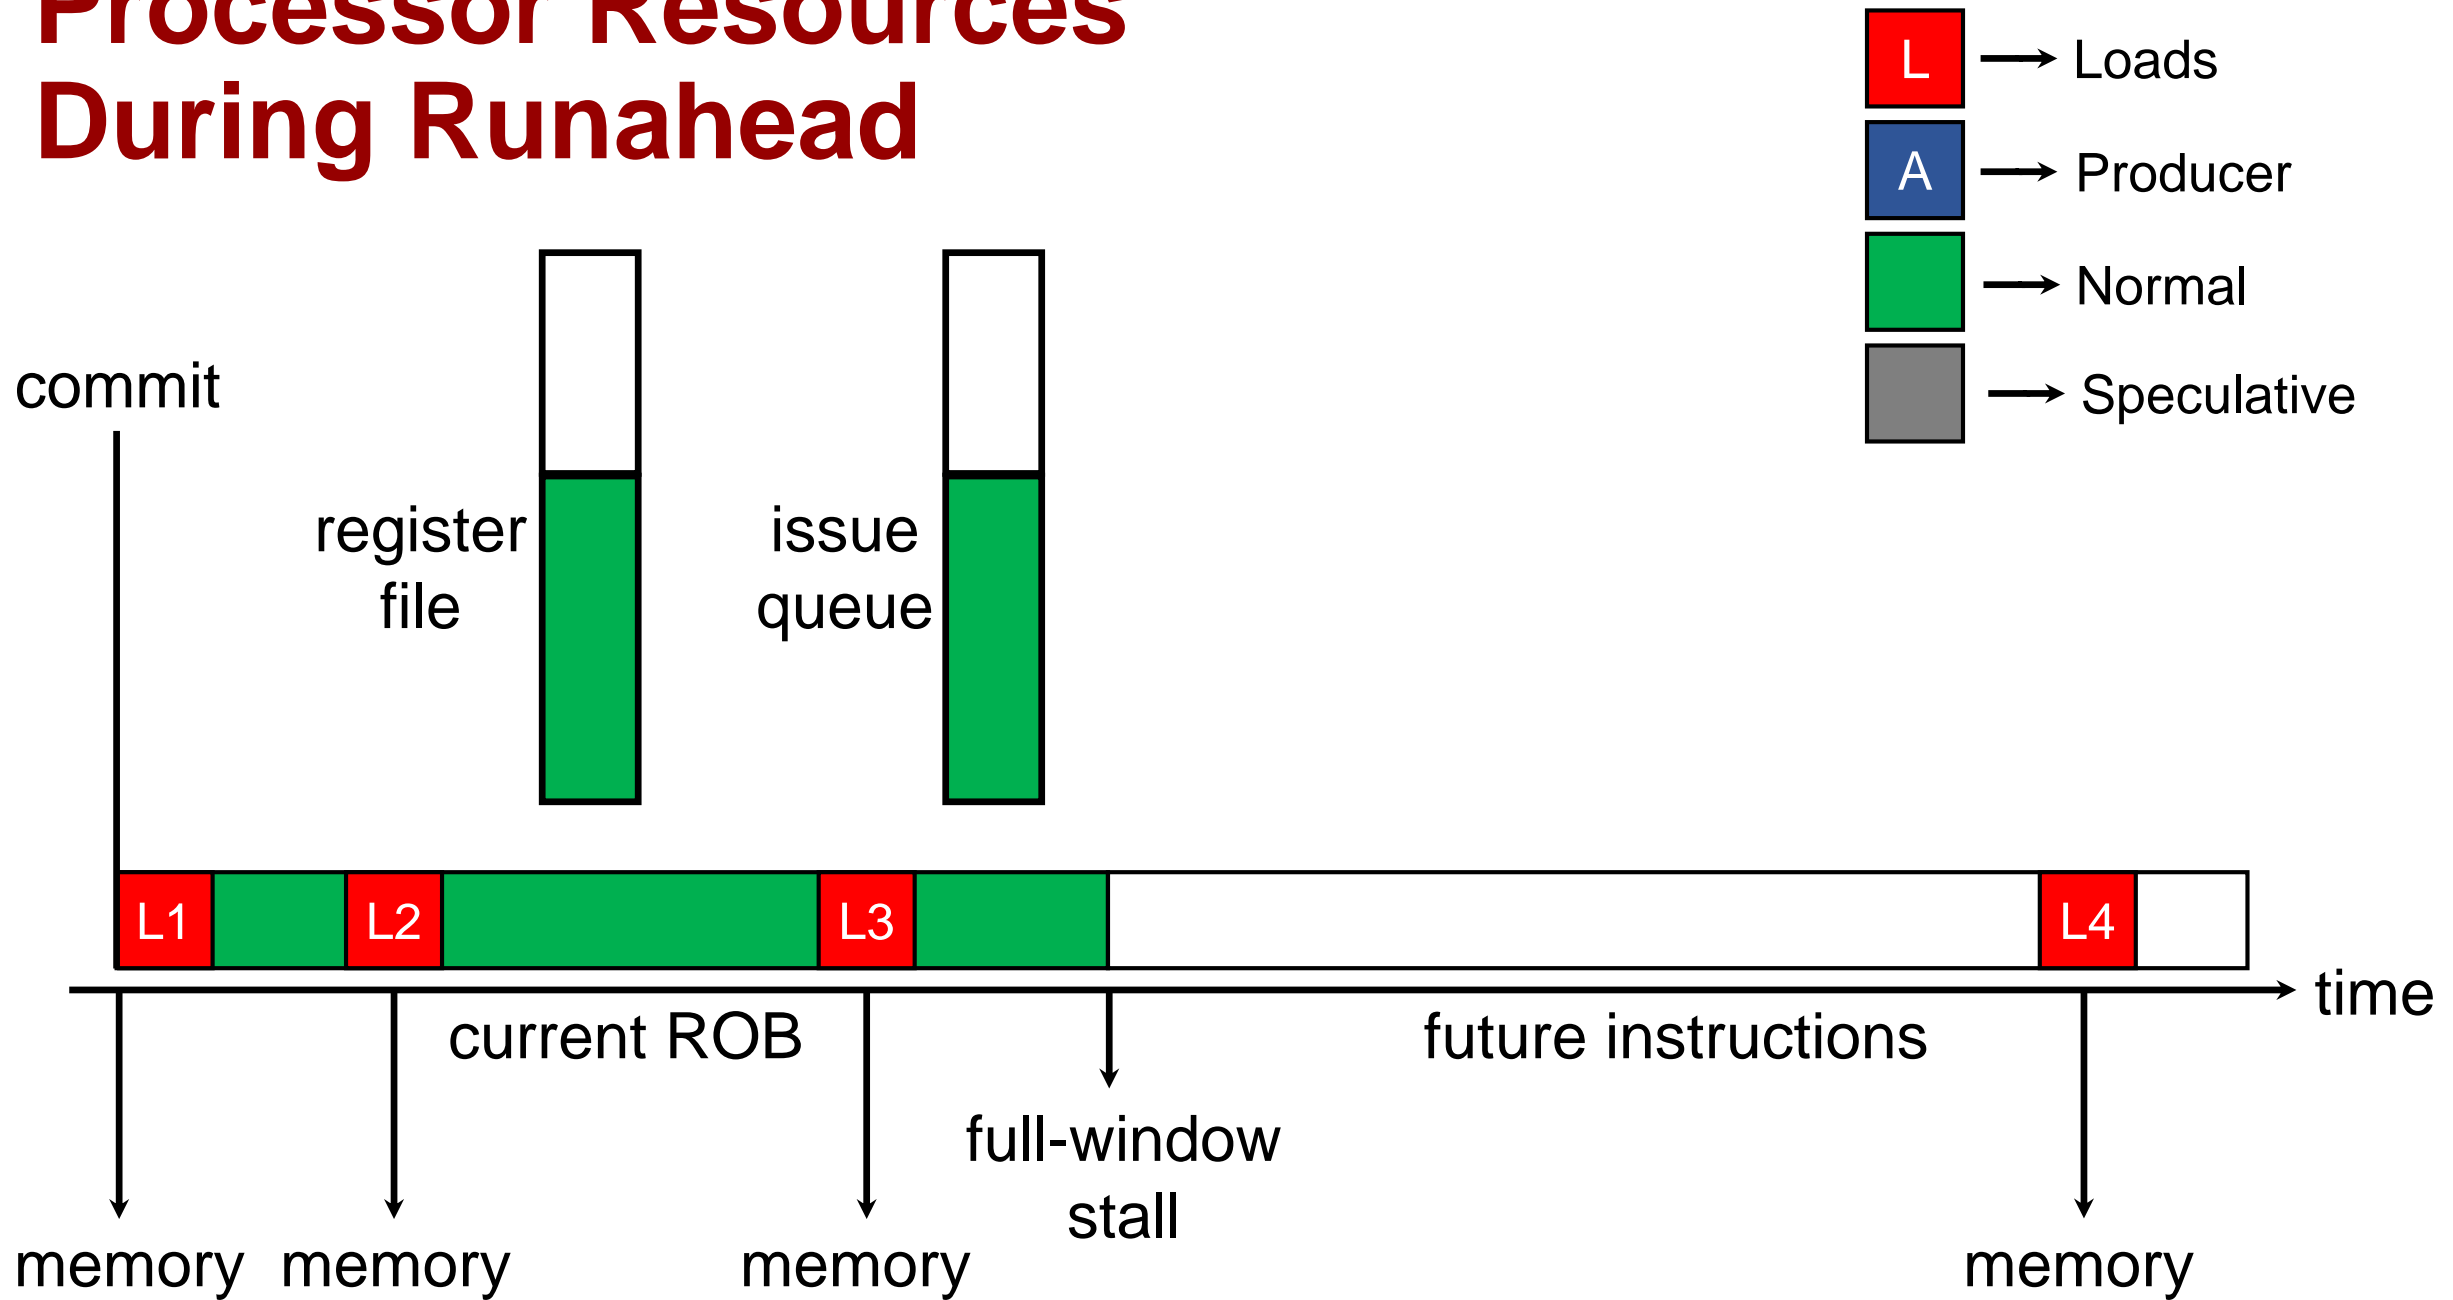
Processor Resources During Runahead



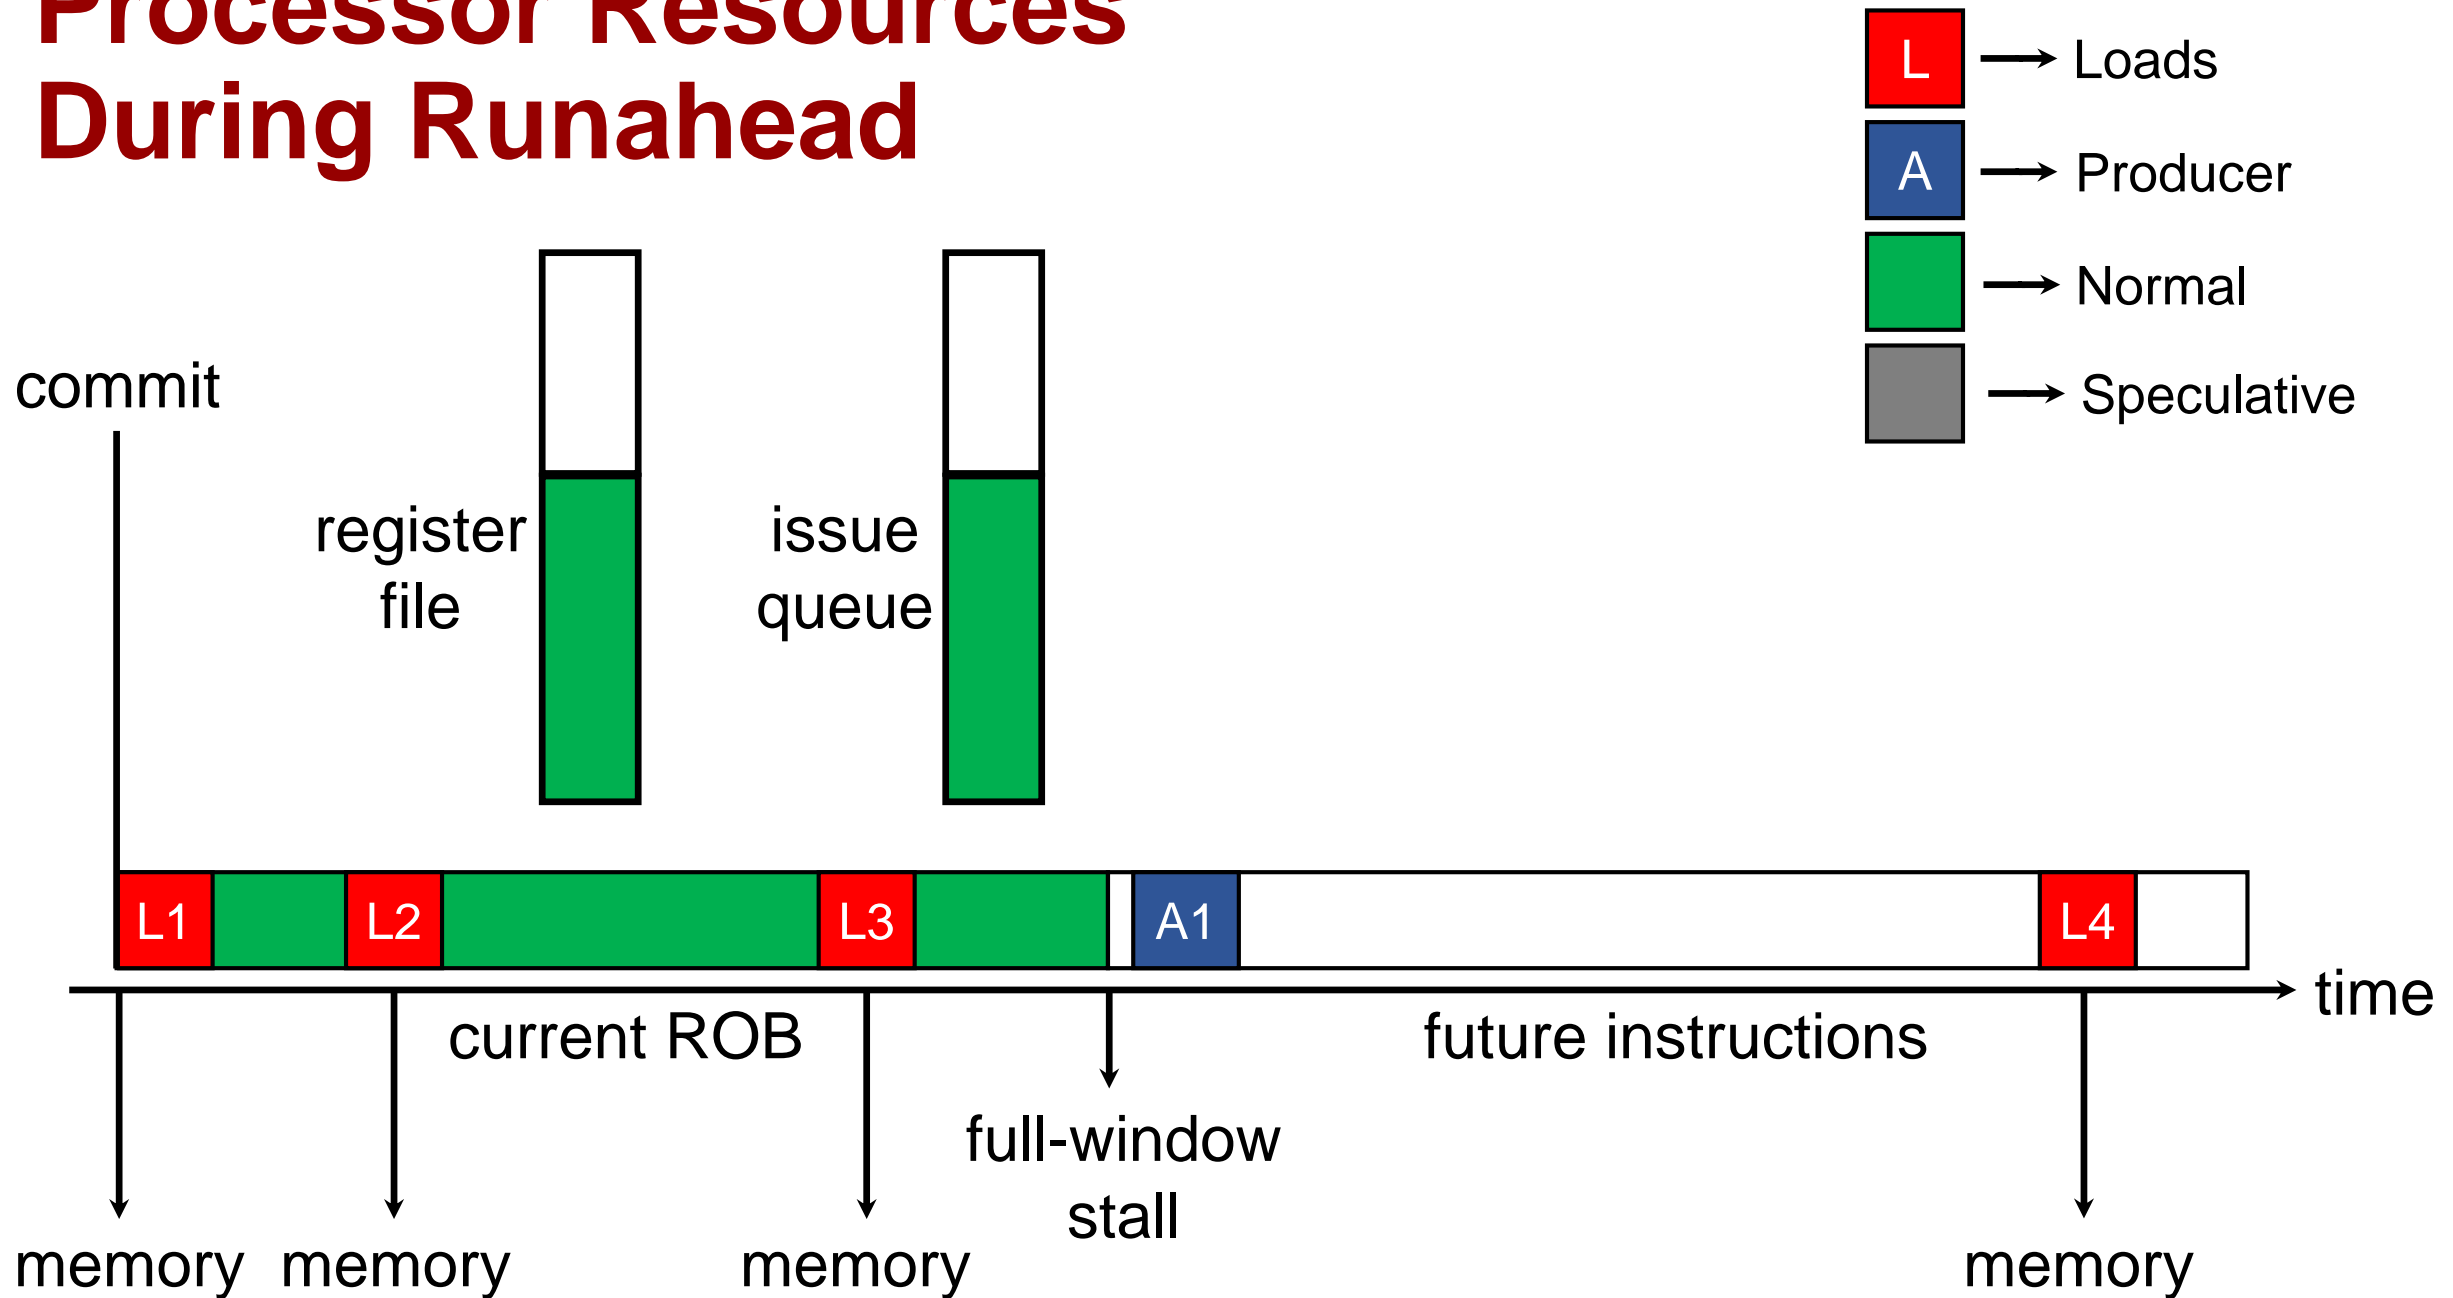
Processor Resources During Runahead



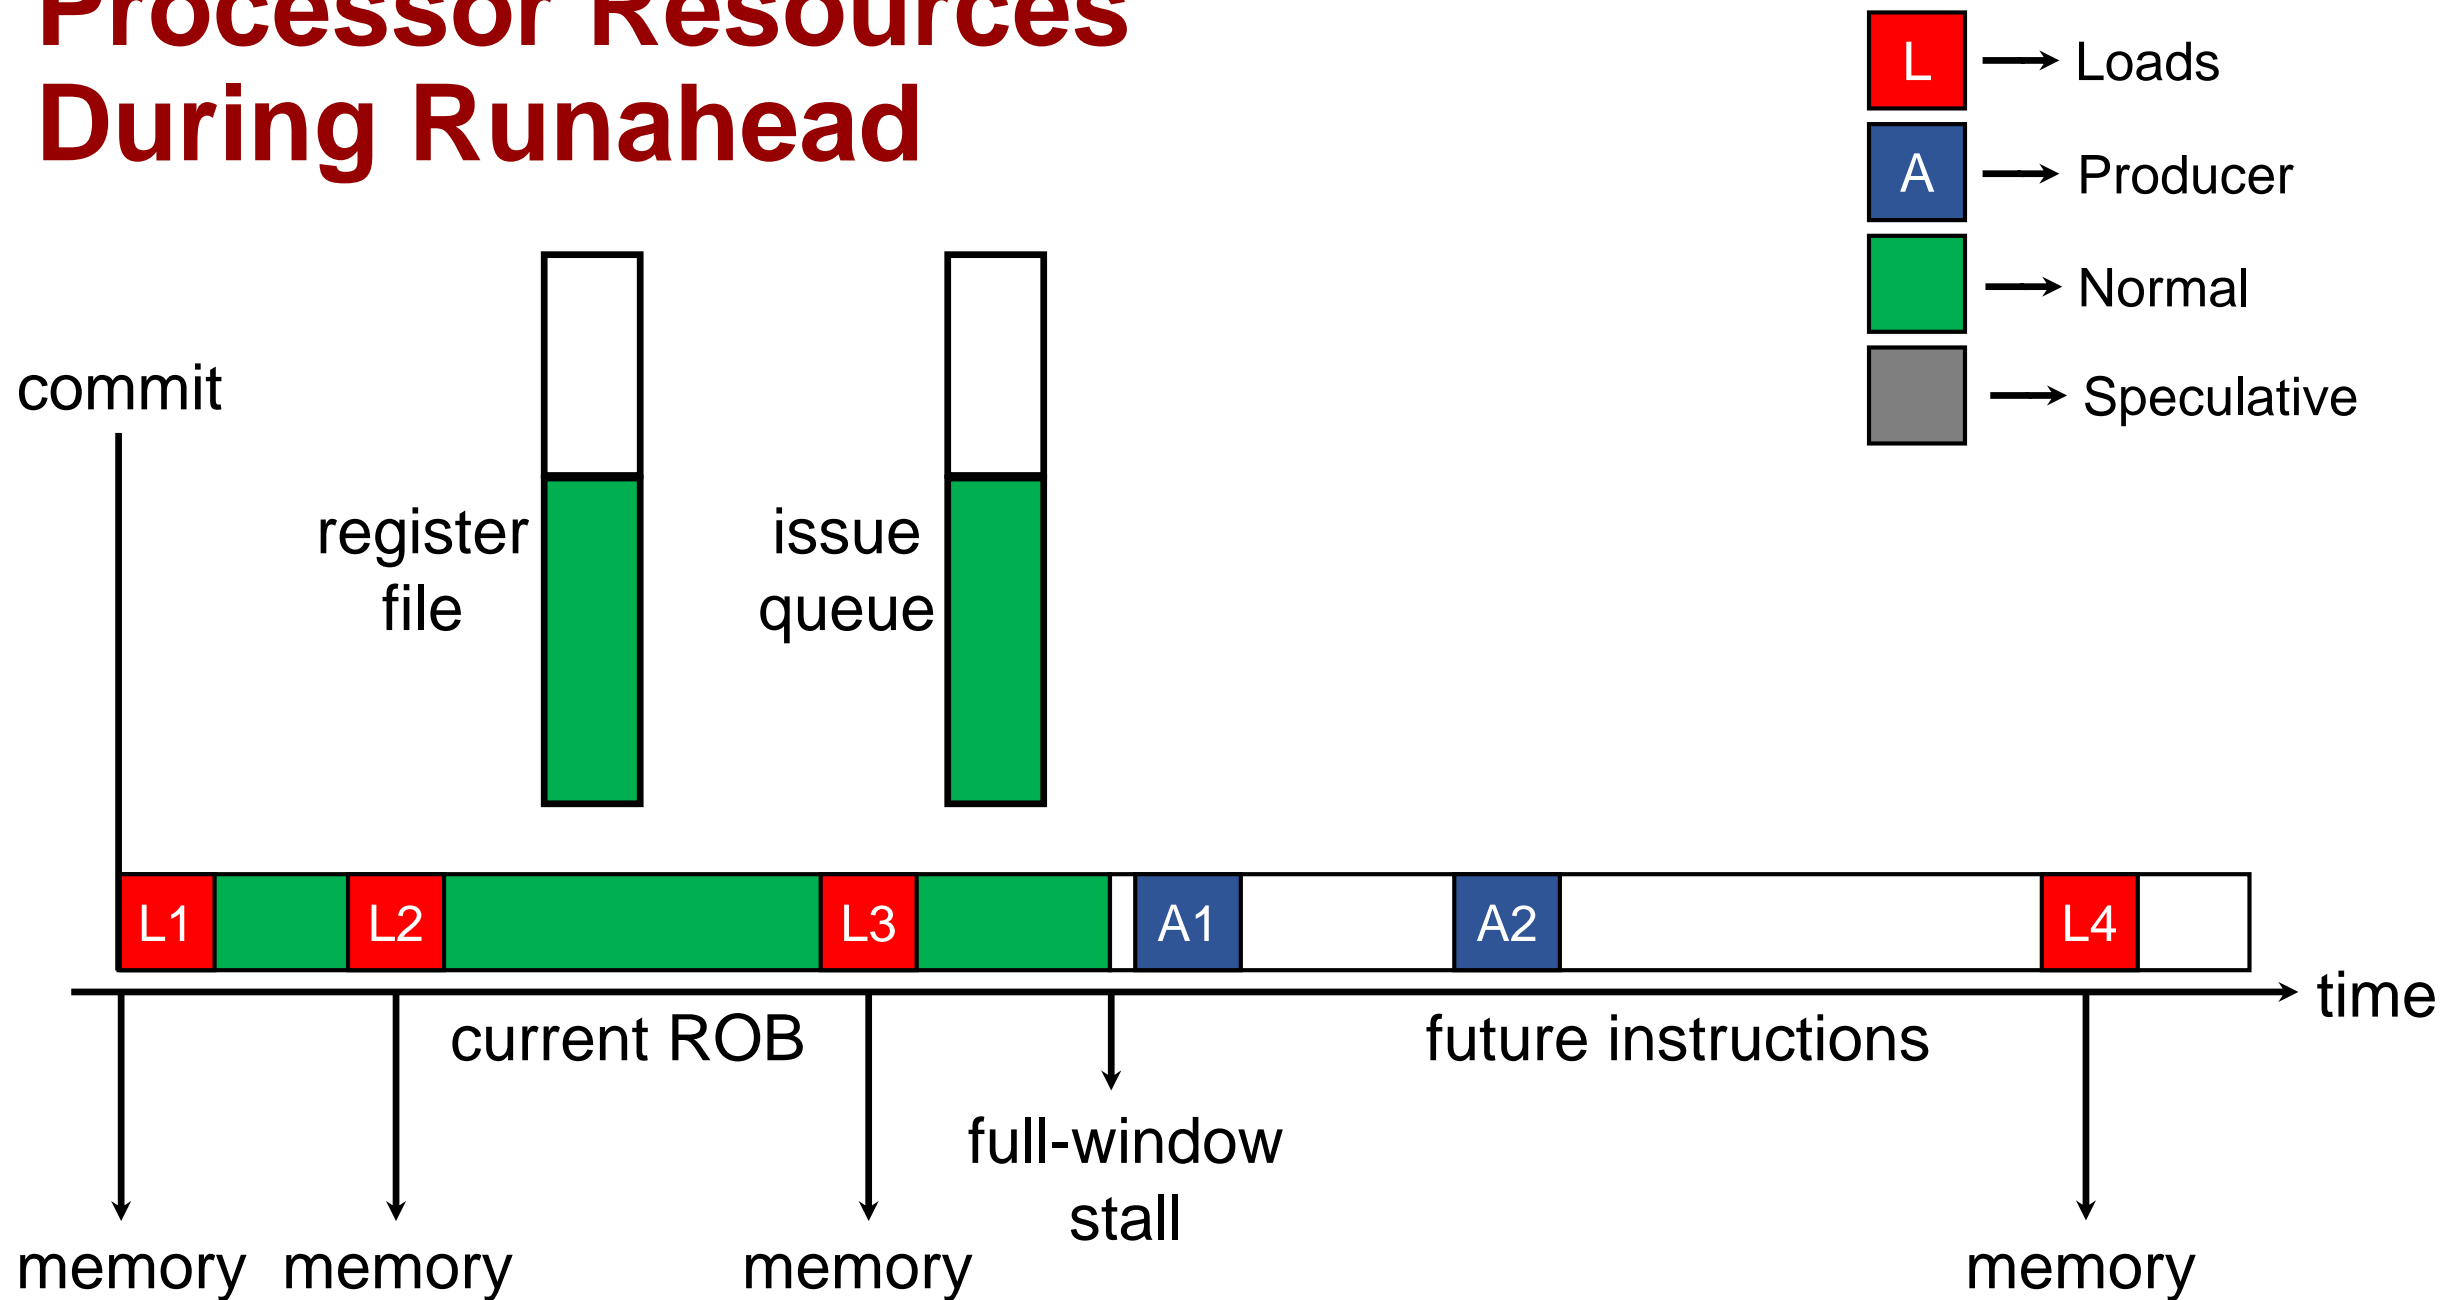
Processor Resources During Runahead



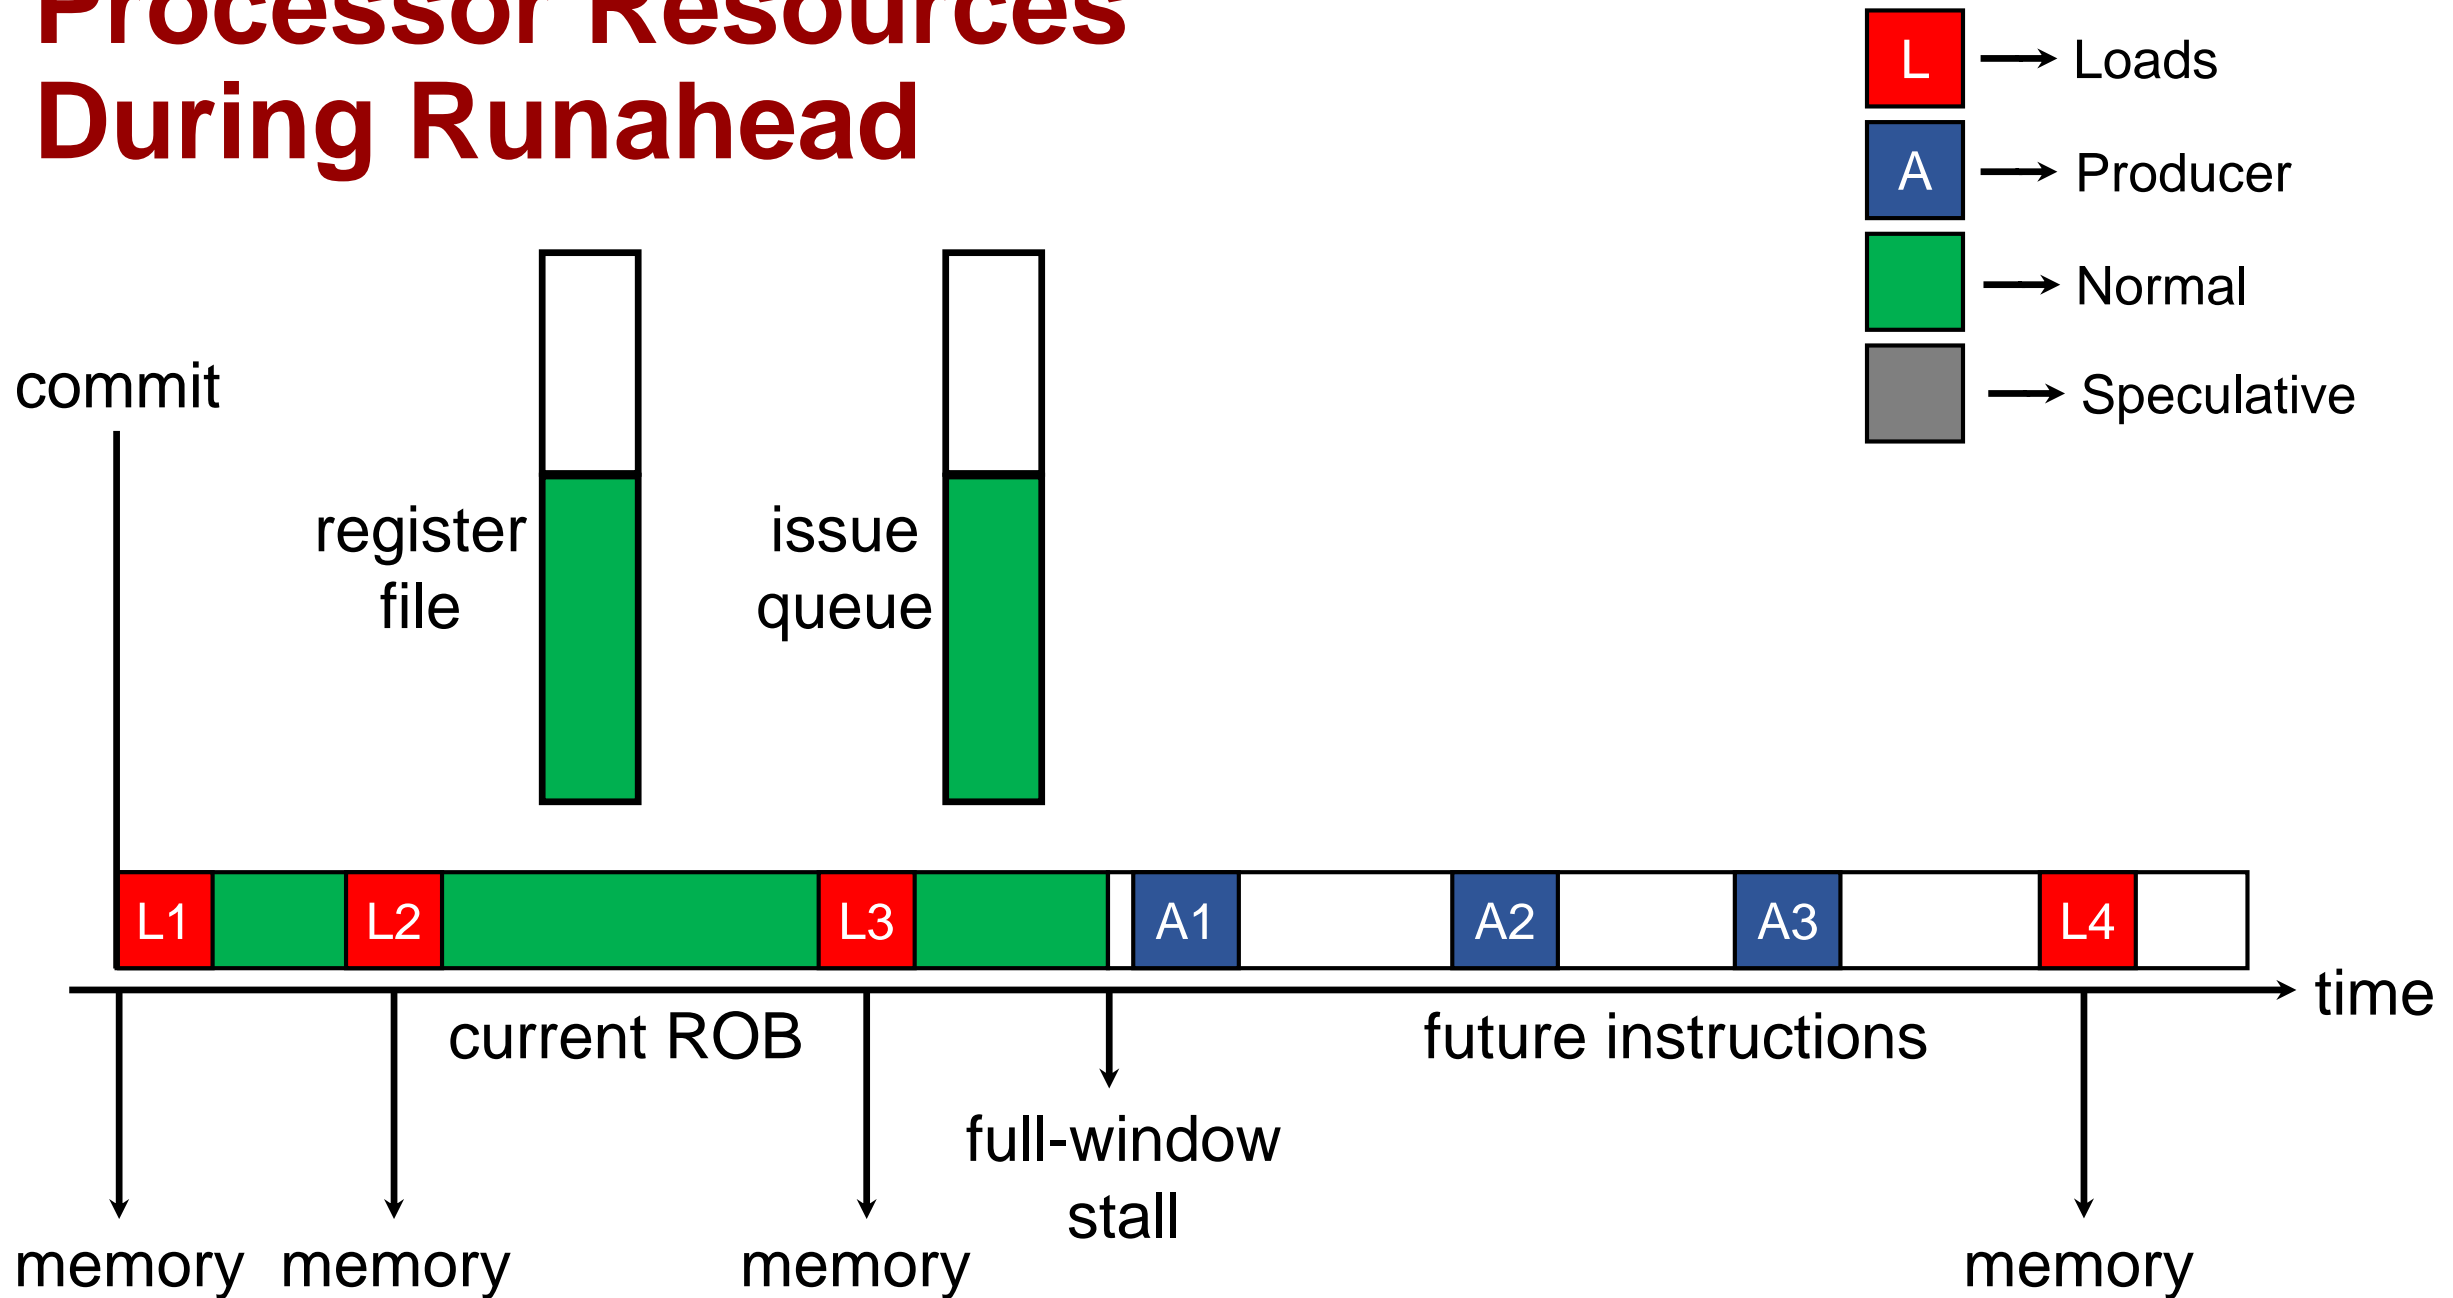
Processor Resources During Runahead




Processor Resources During Runahead




Processor Resources During Runahead




Processor Resources During Runahead




Processor Resources During Runahead




Processor Resources During Runahead




Processor Resources During Runahead




Processor Resources During Runahead




Processor Resources During Runahead




Processor Resources During Runahead




Processor Resources During Runahead




Processor Resources During Runahead



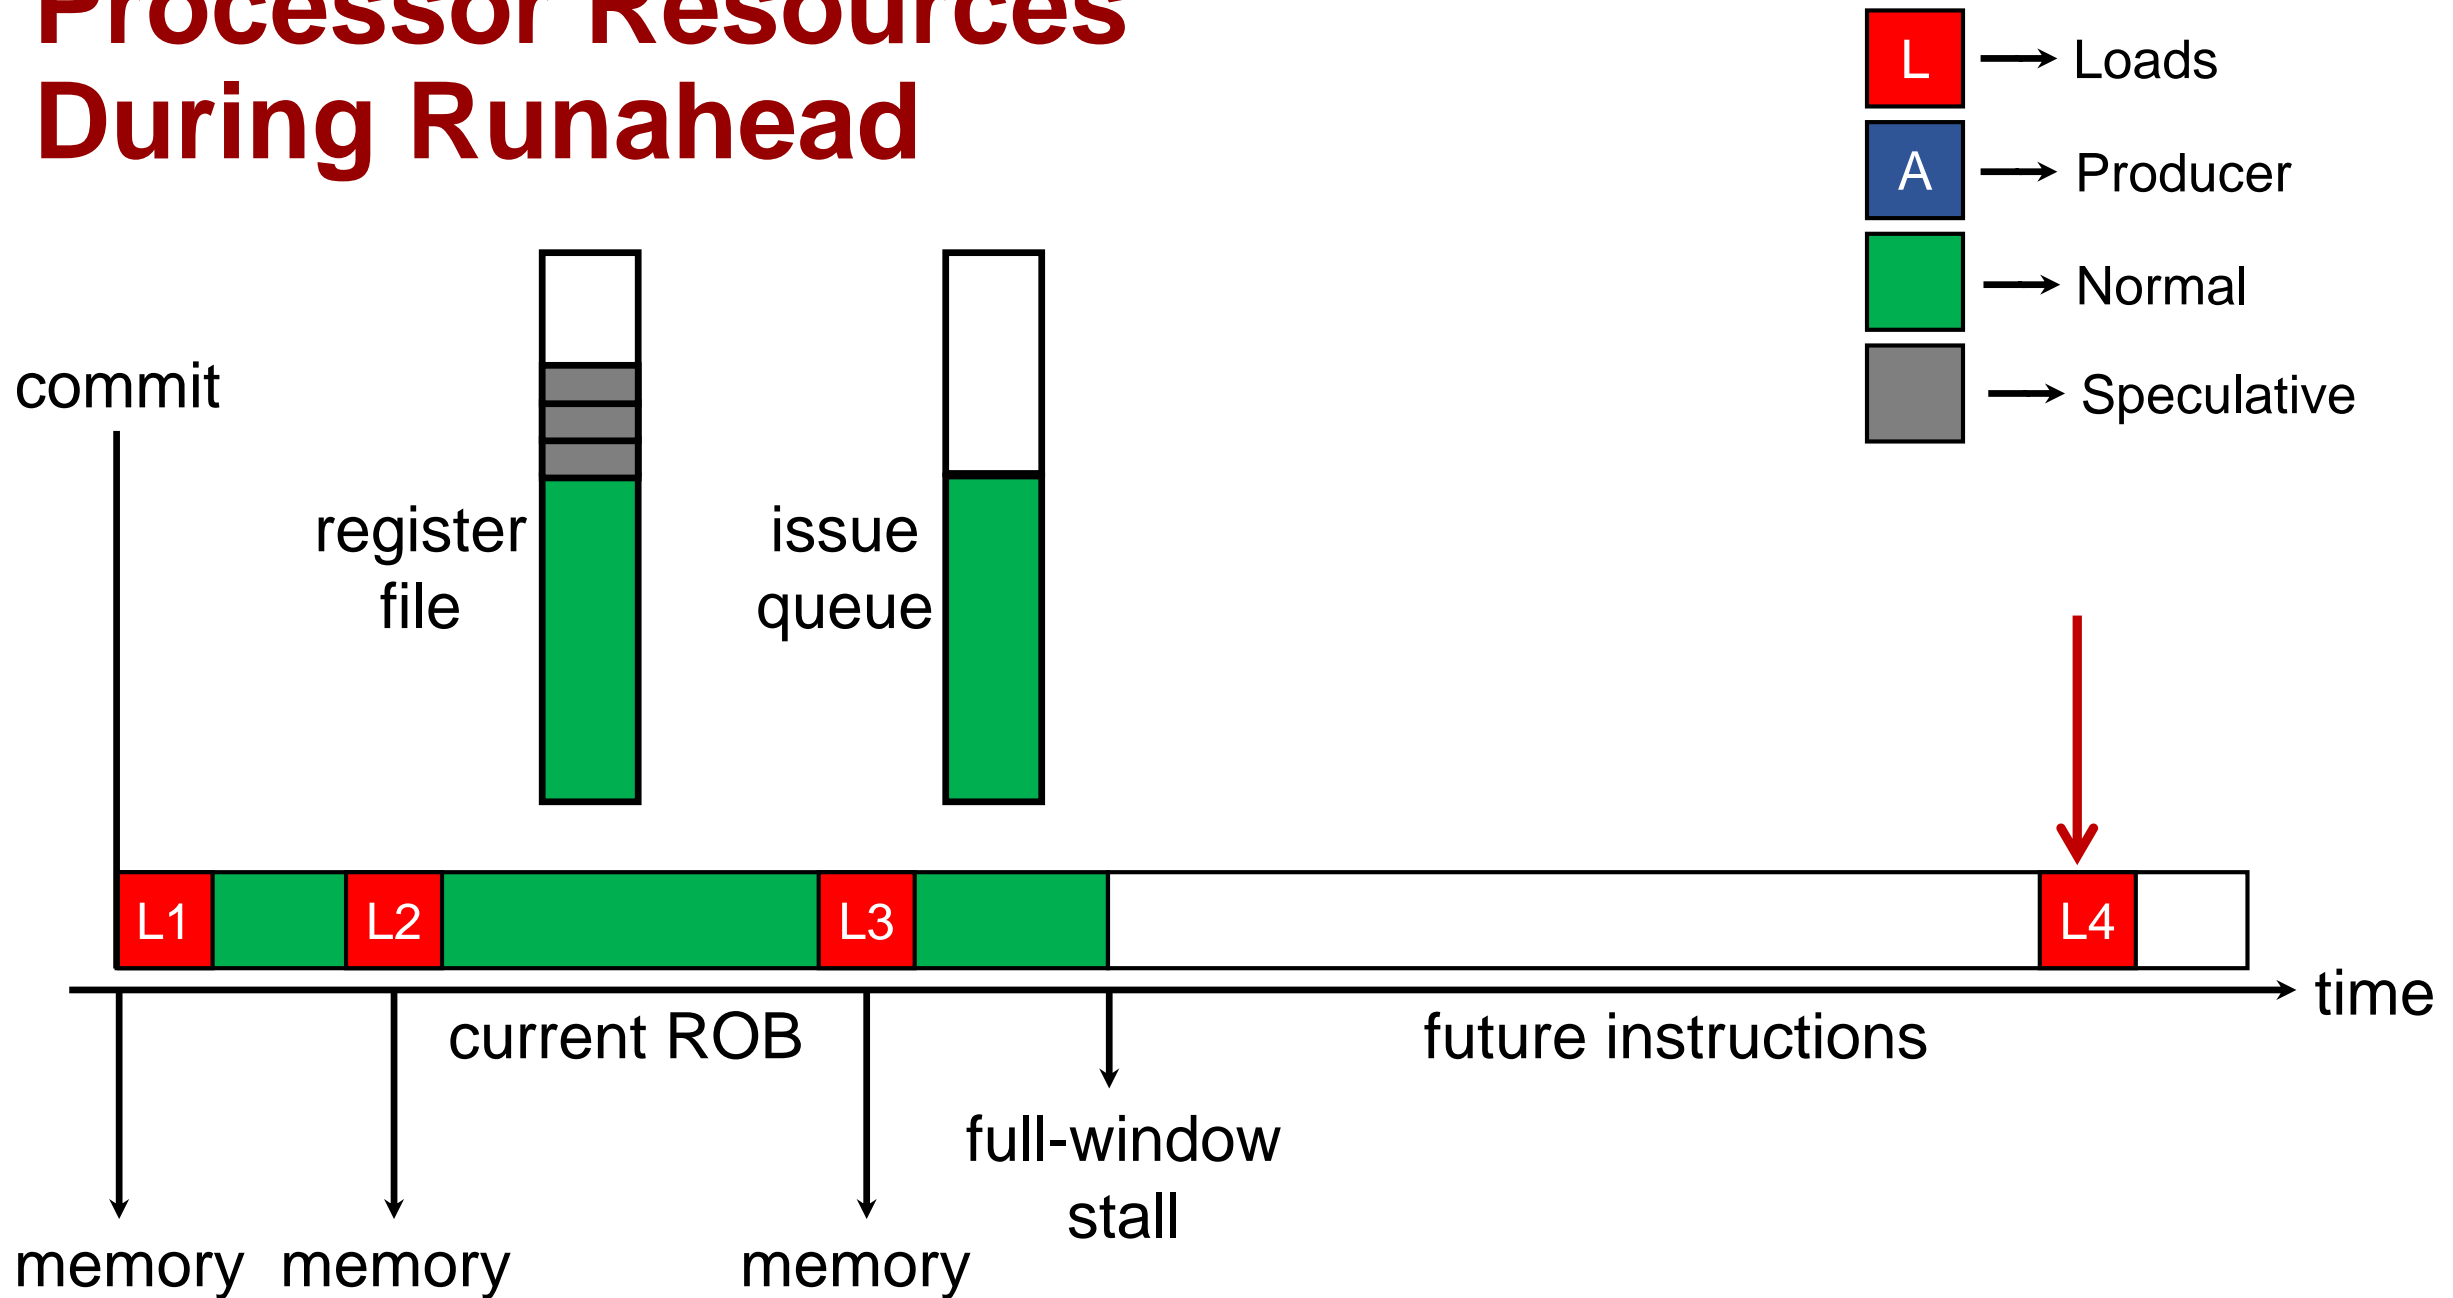
Processor Resources During Runahead




Processor Resources During Runahead



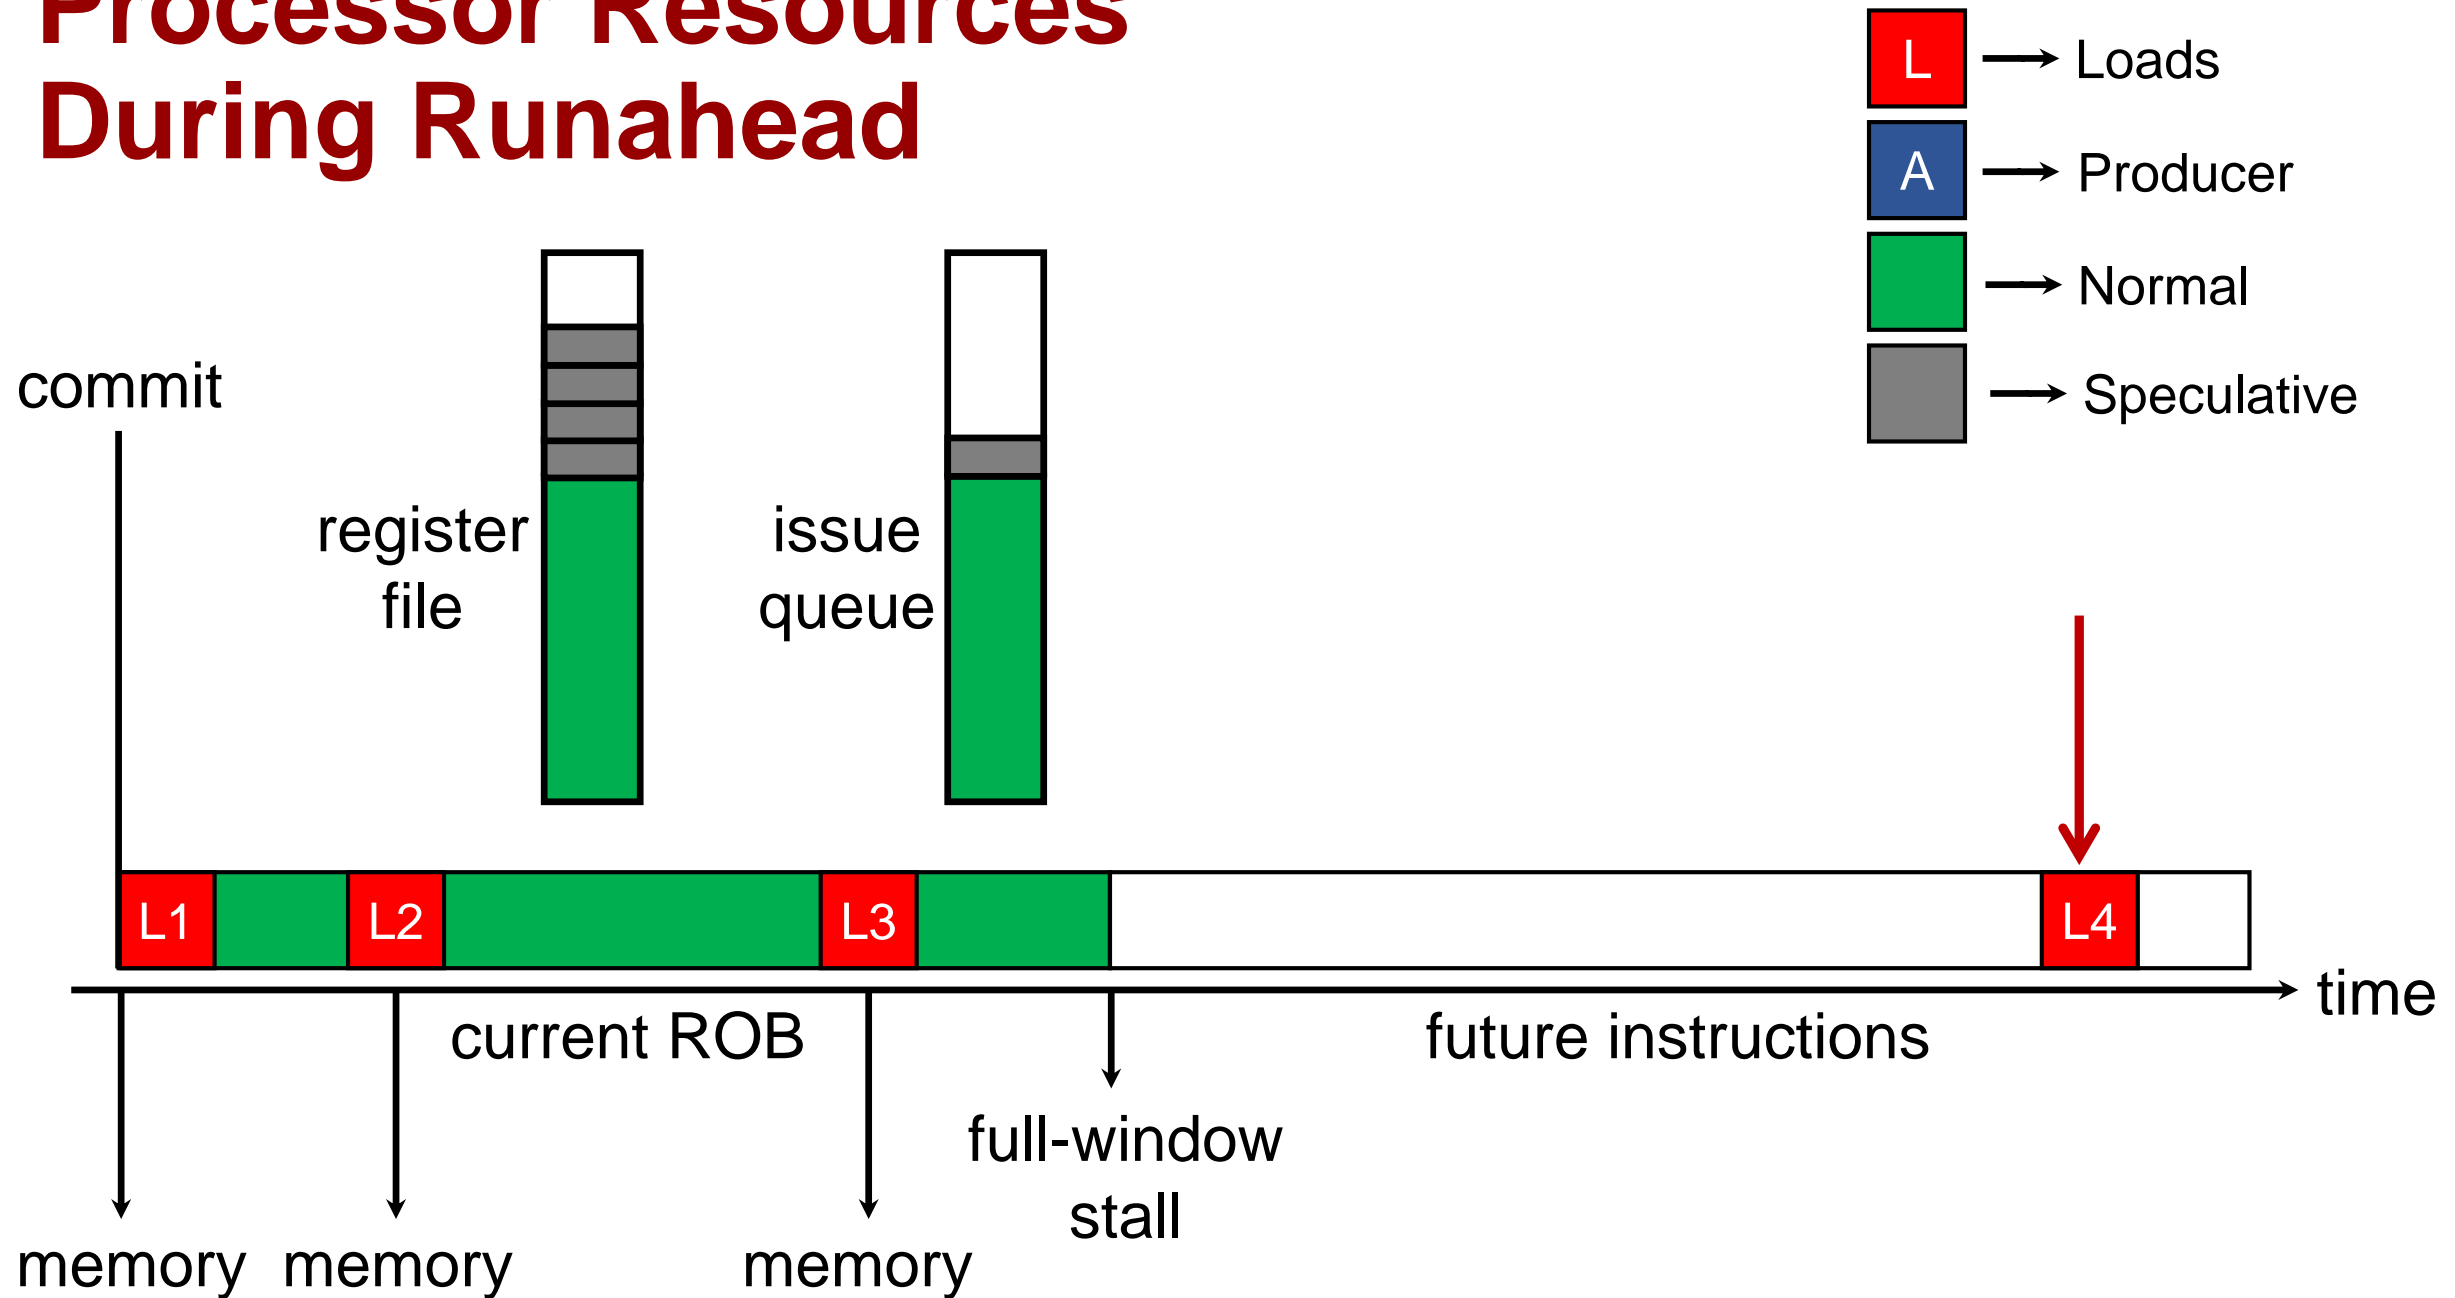
Processor Resources During Runahead



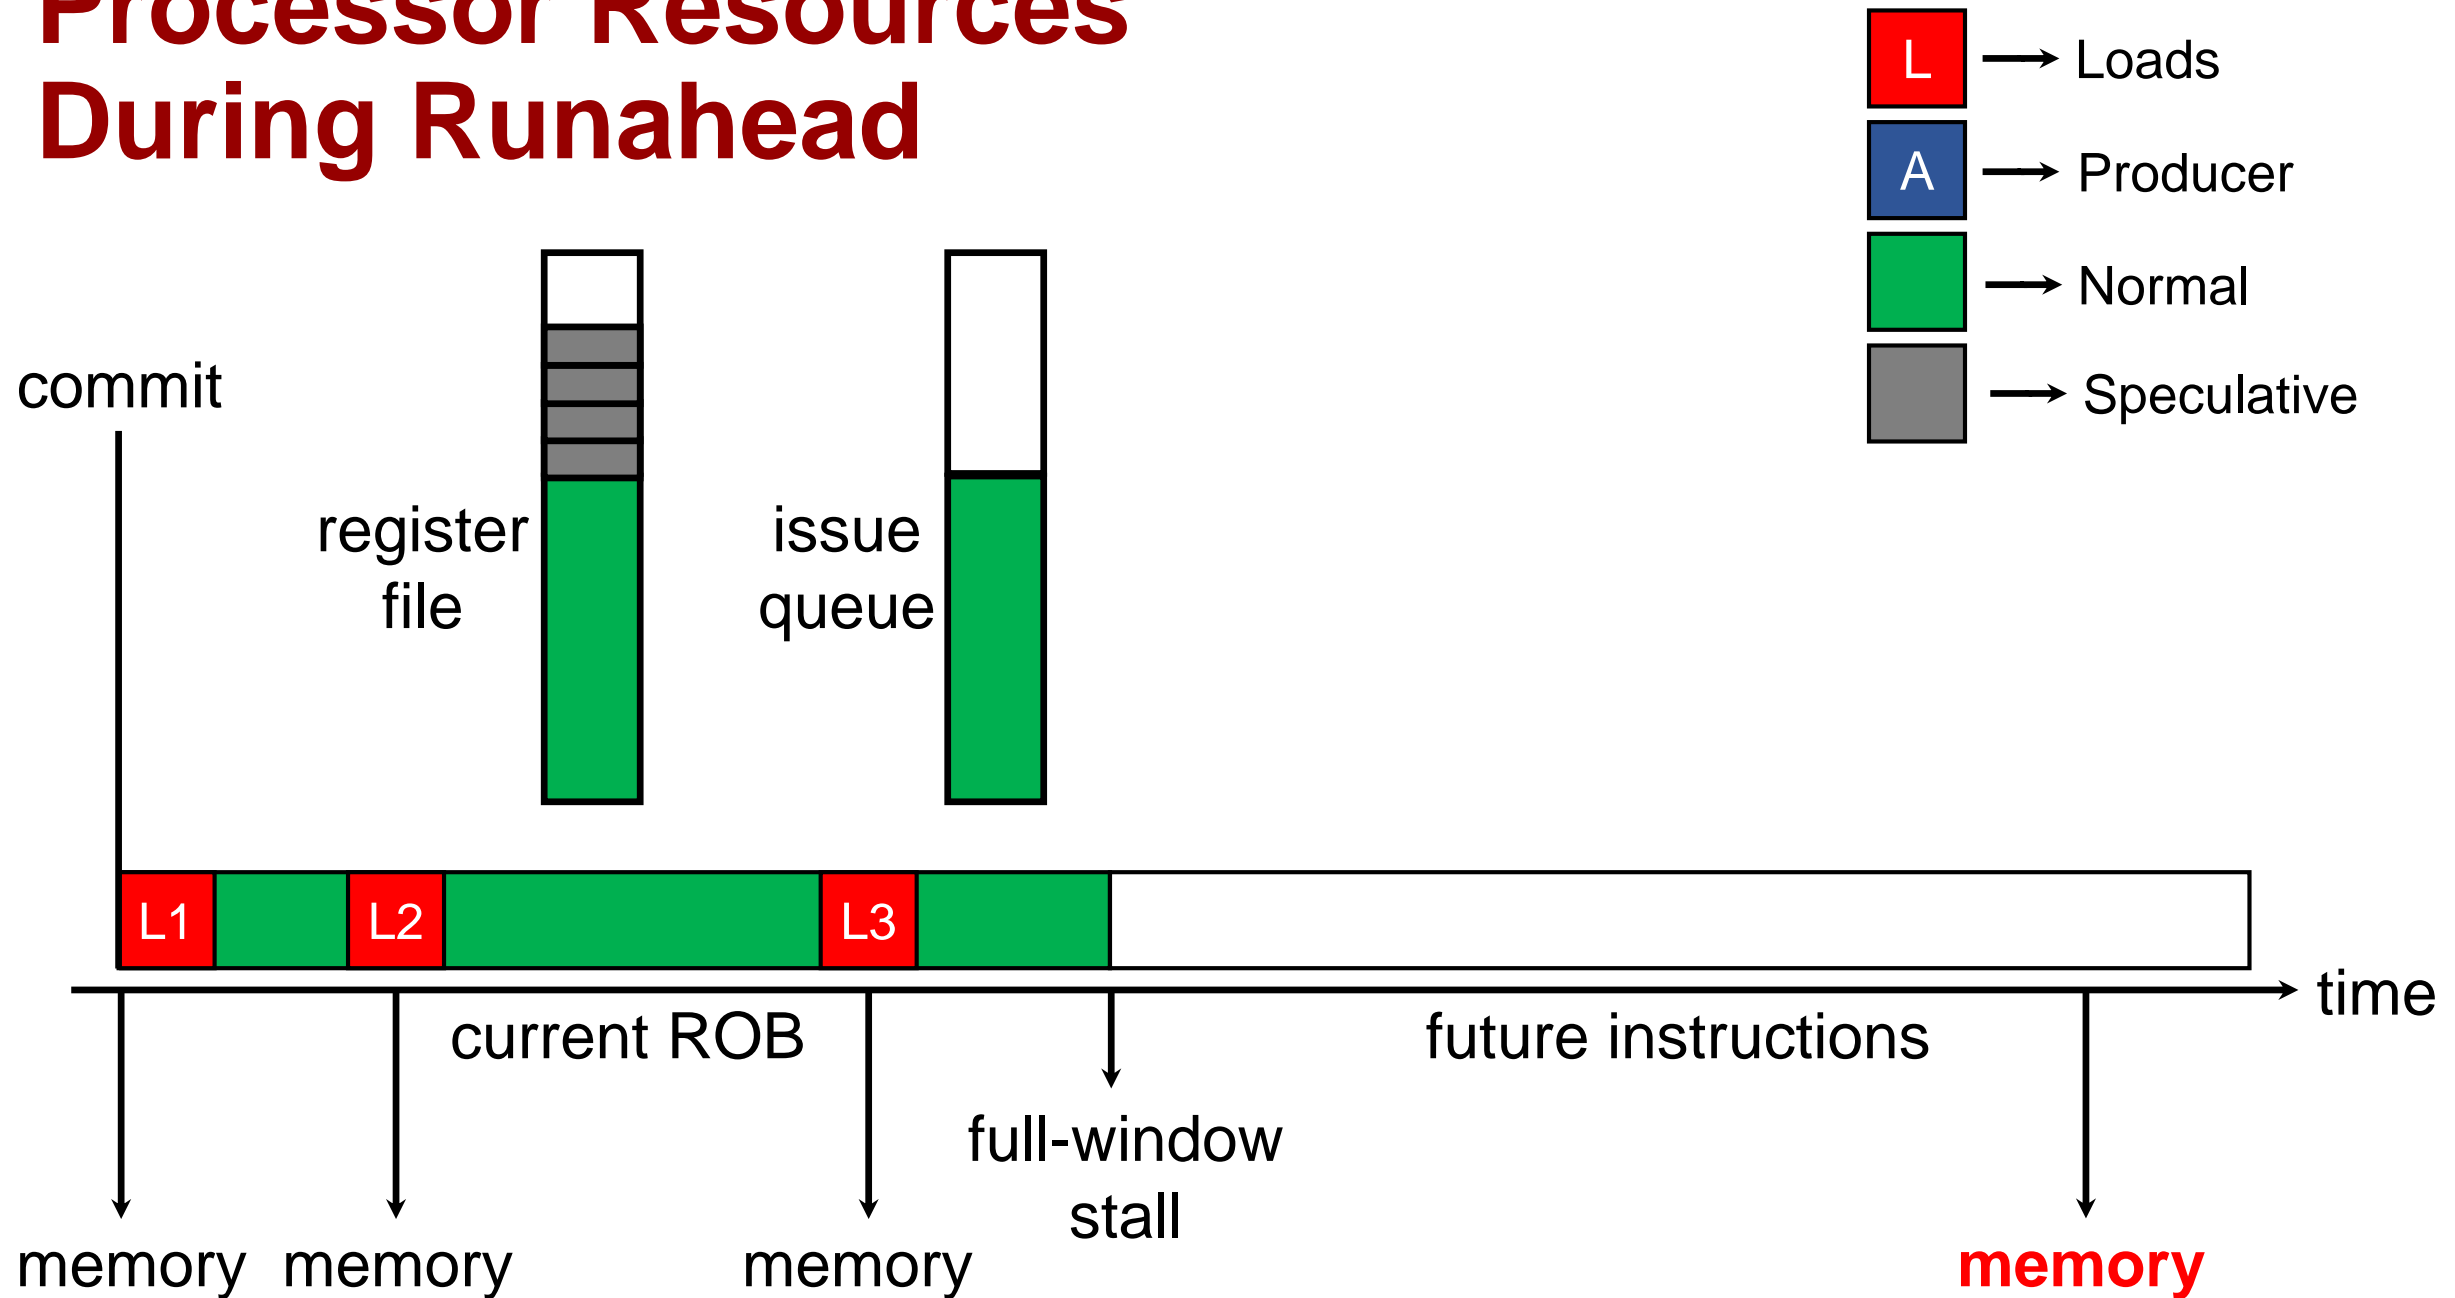
Processor Resources During Runahead




Processor Resources During Runahead



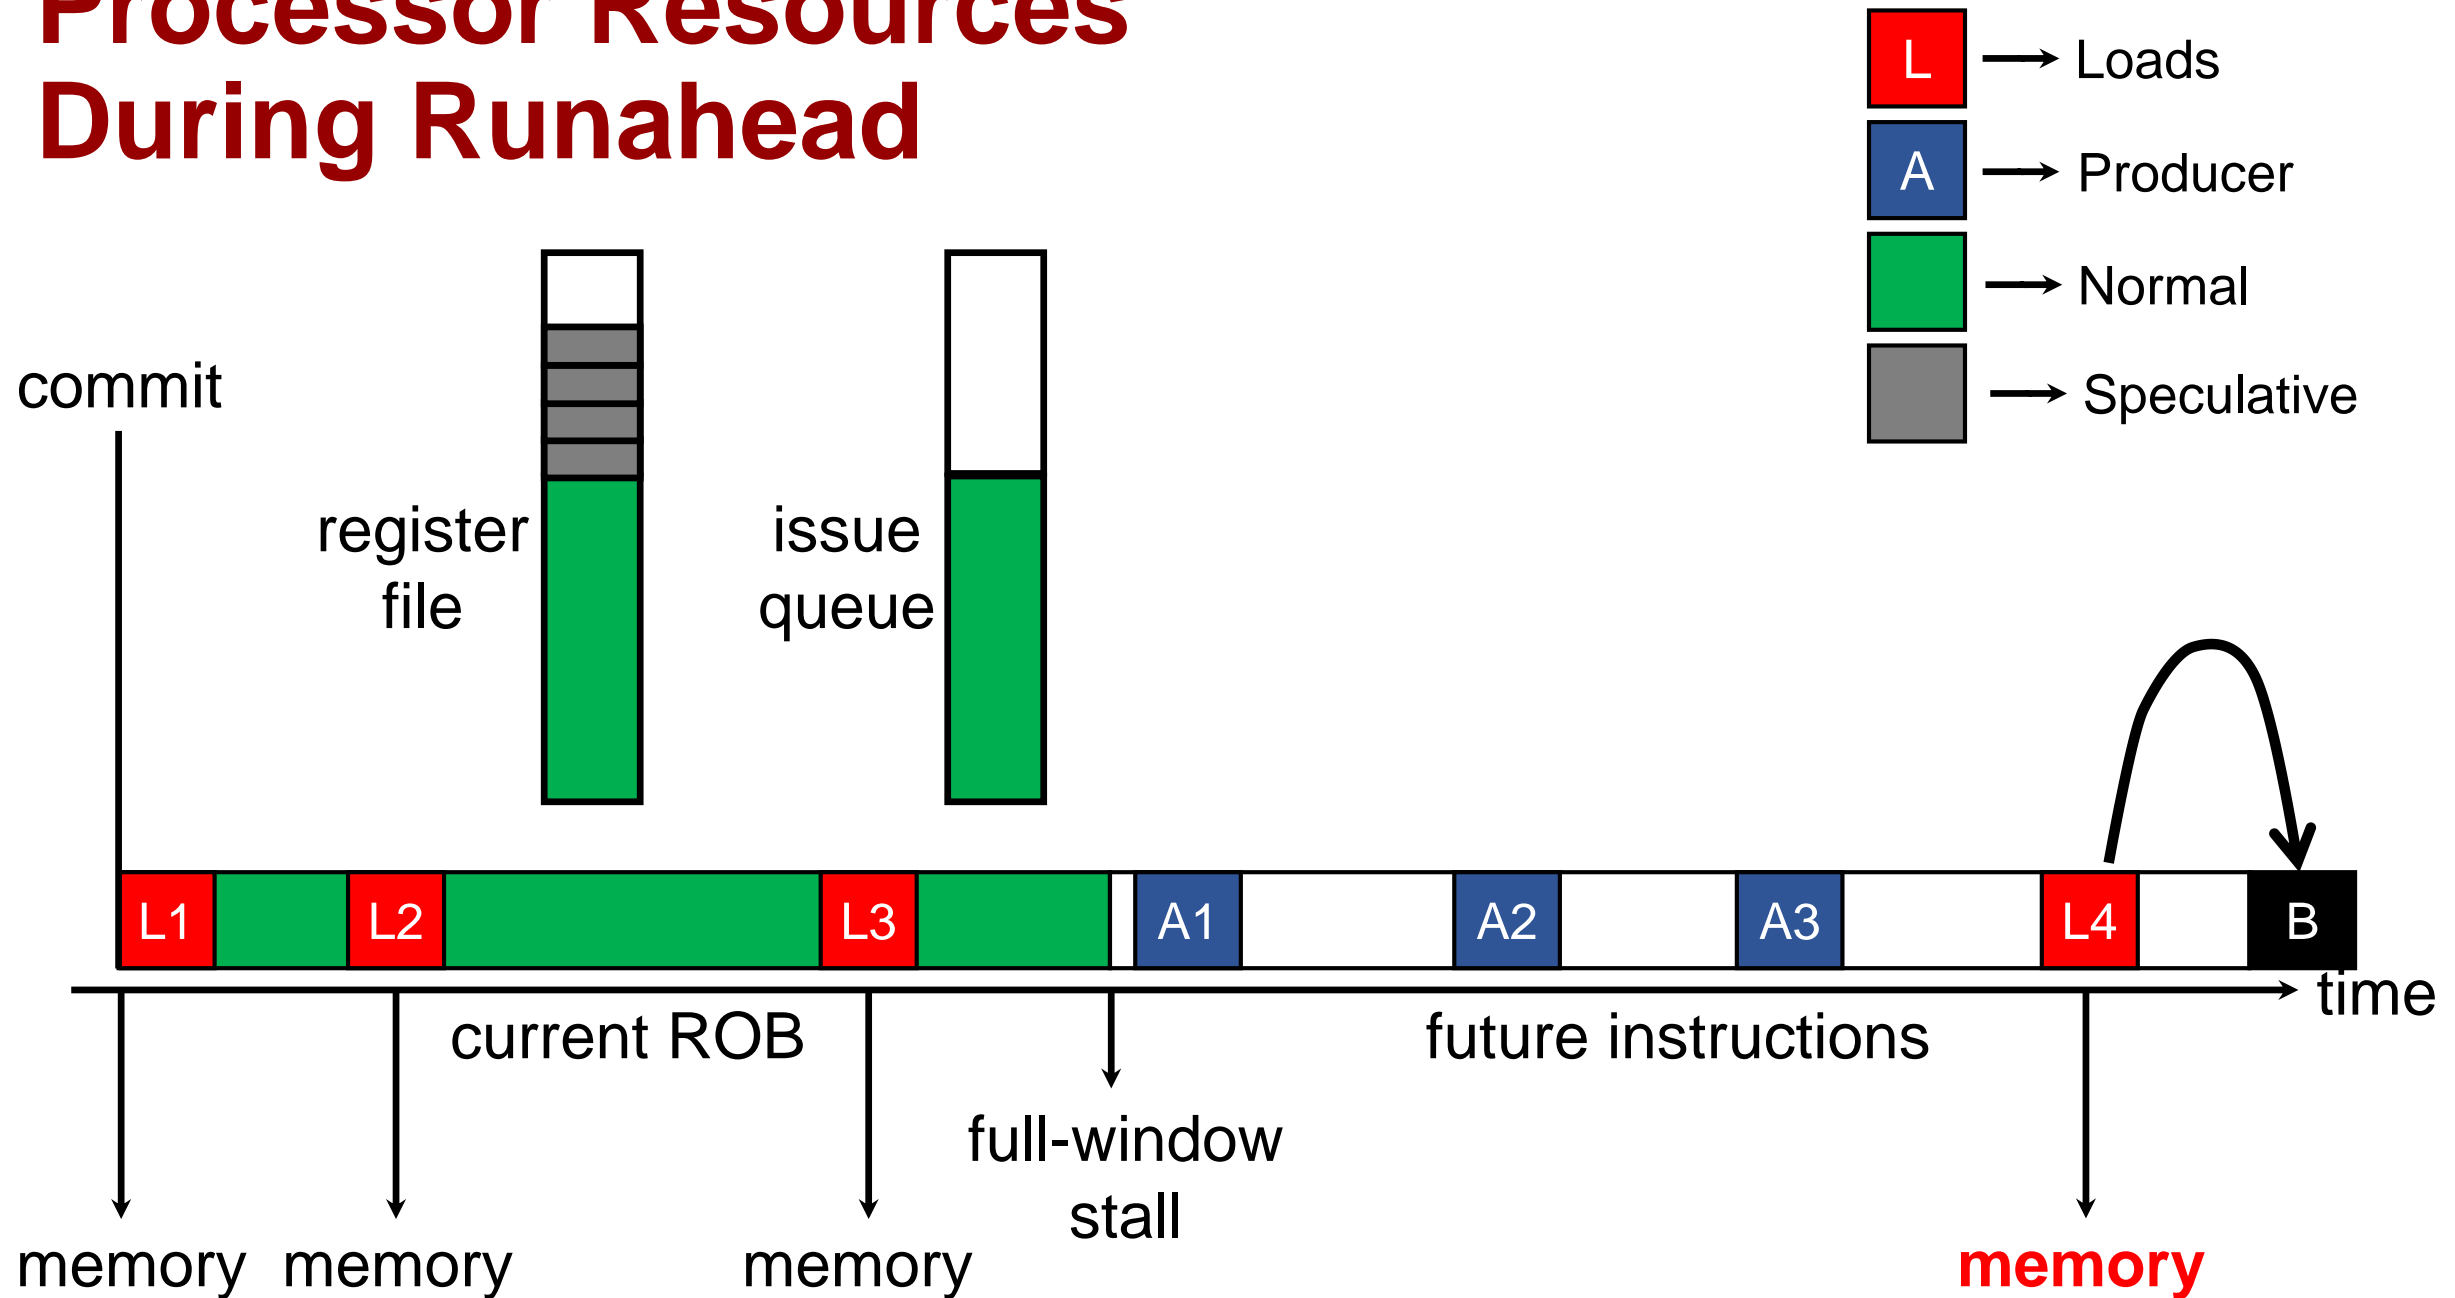
Processor Resources During Runahead



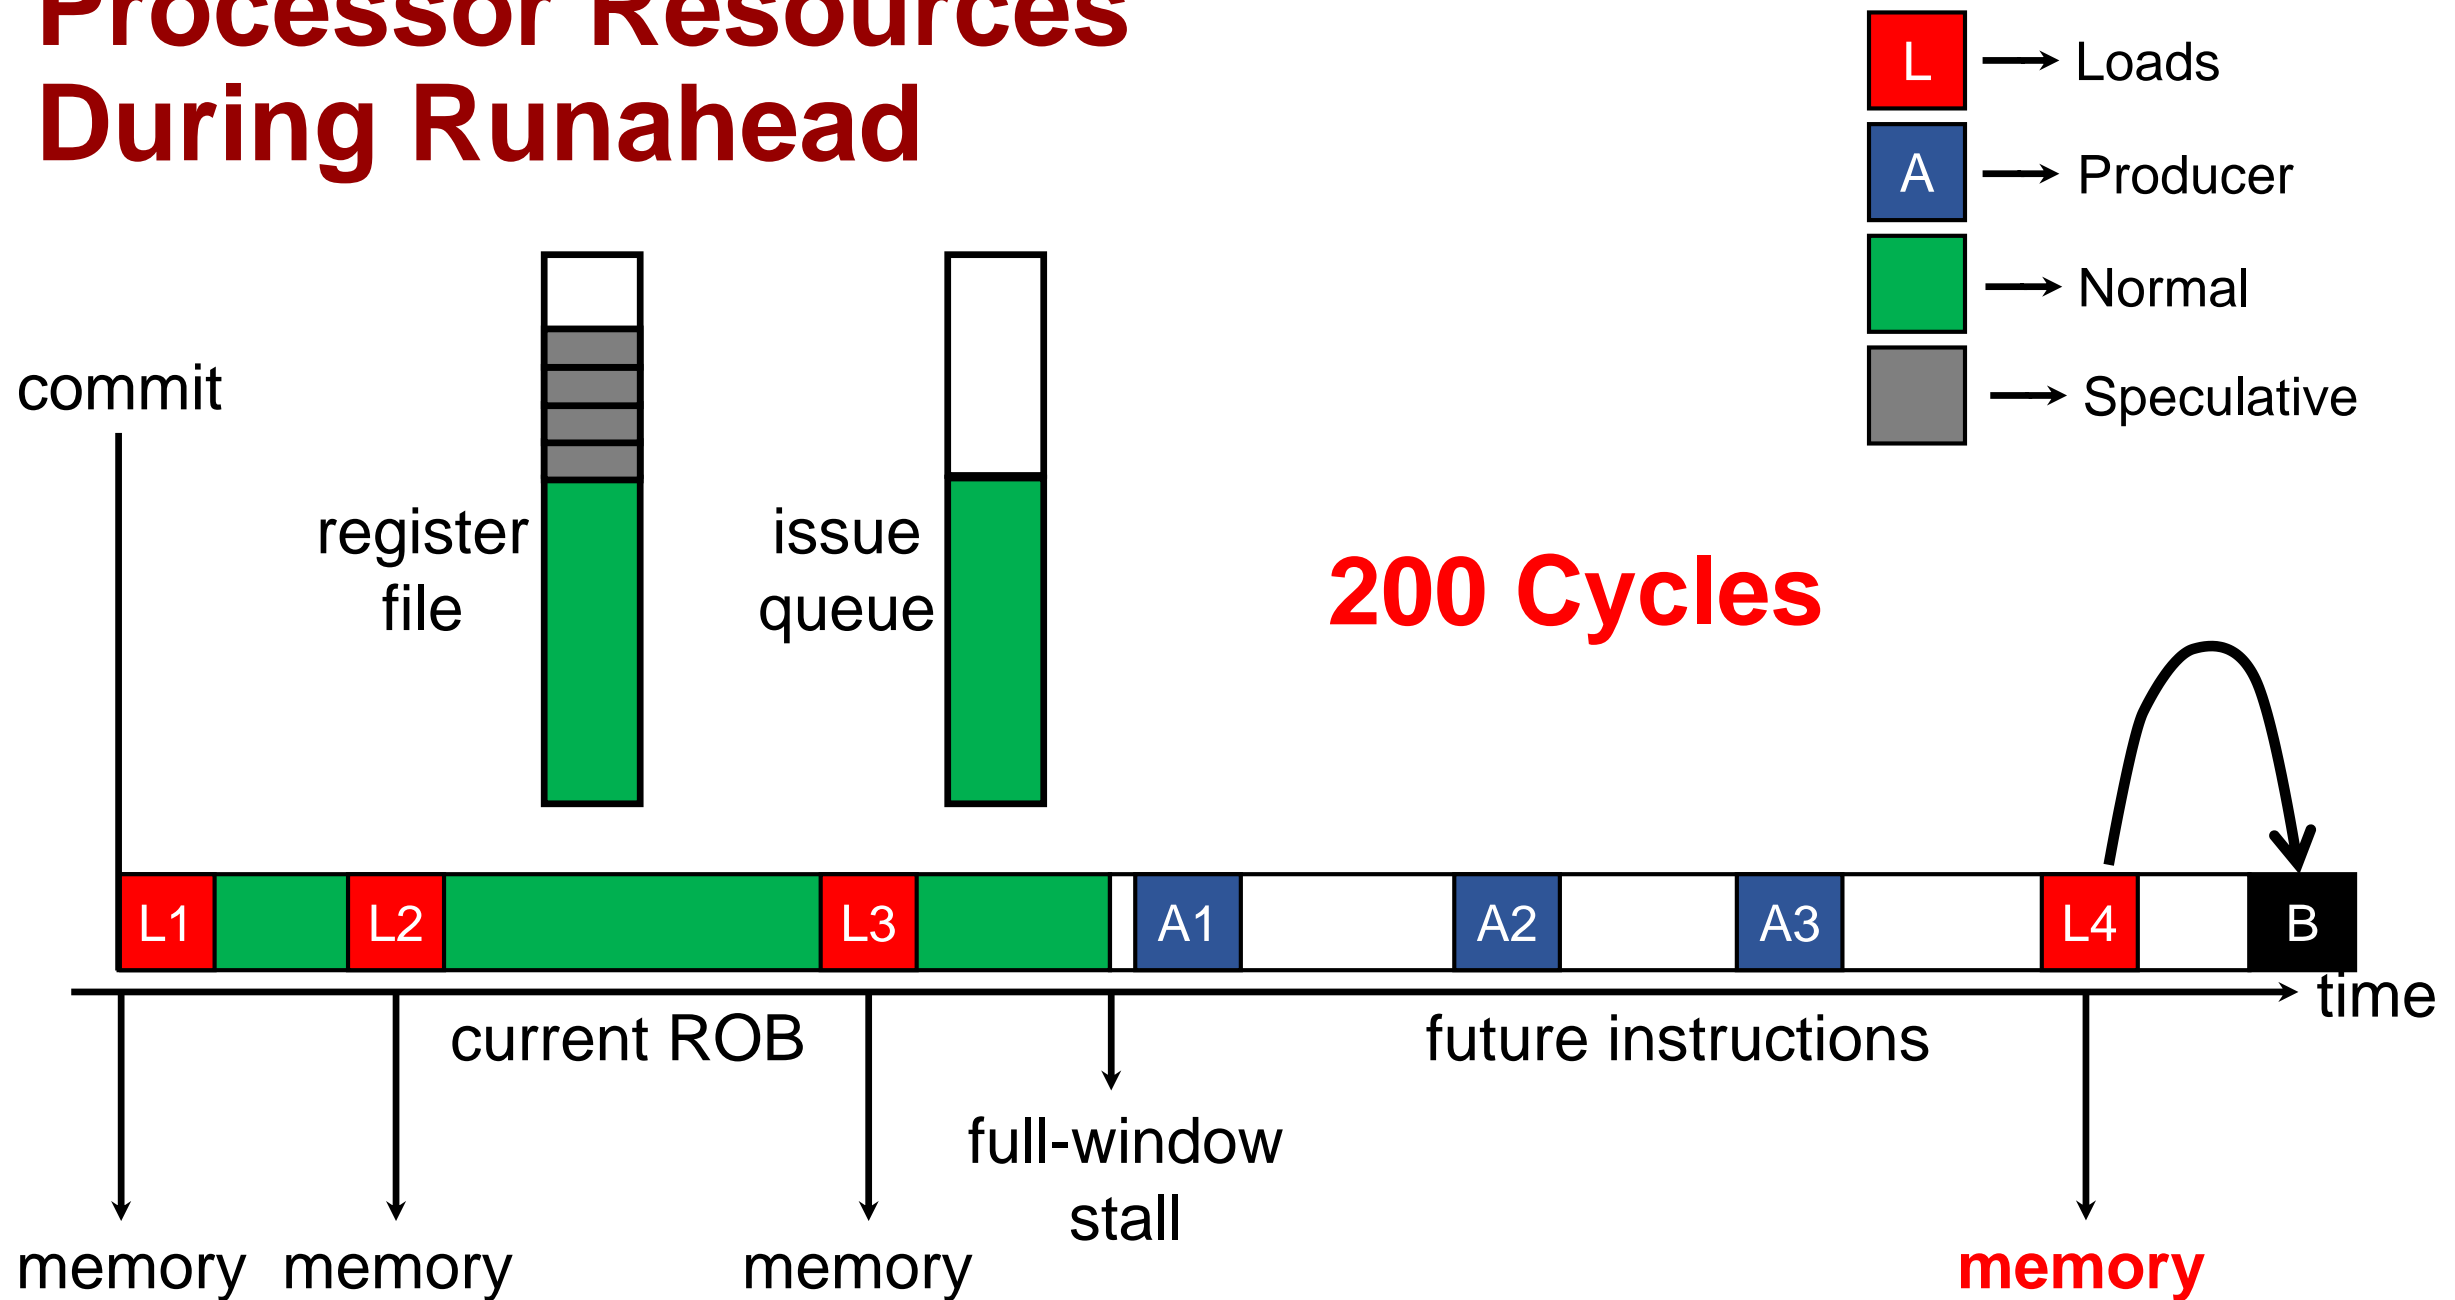
Processor Resources During Runahead



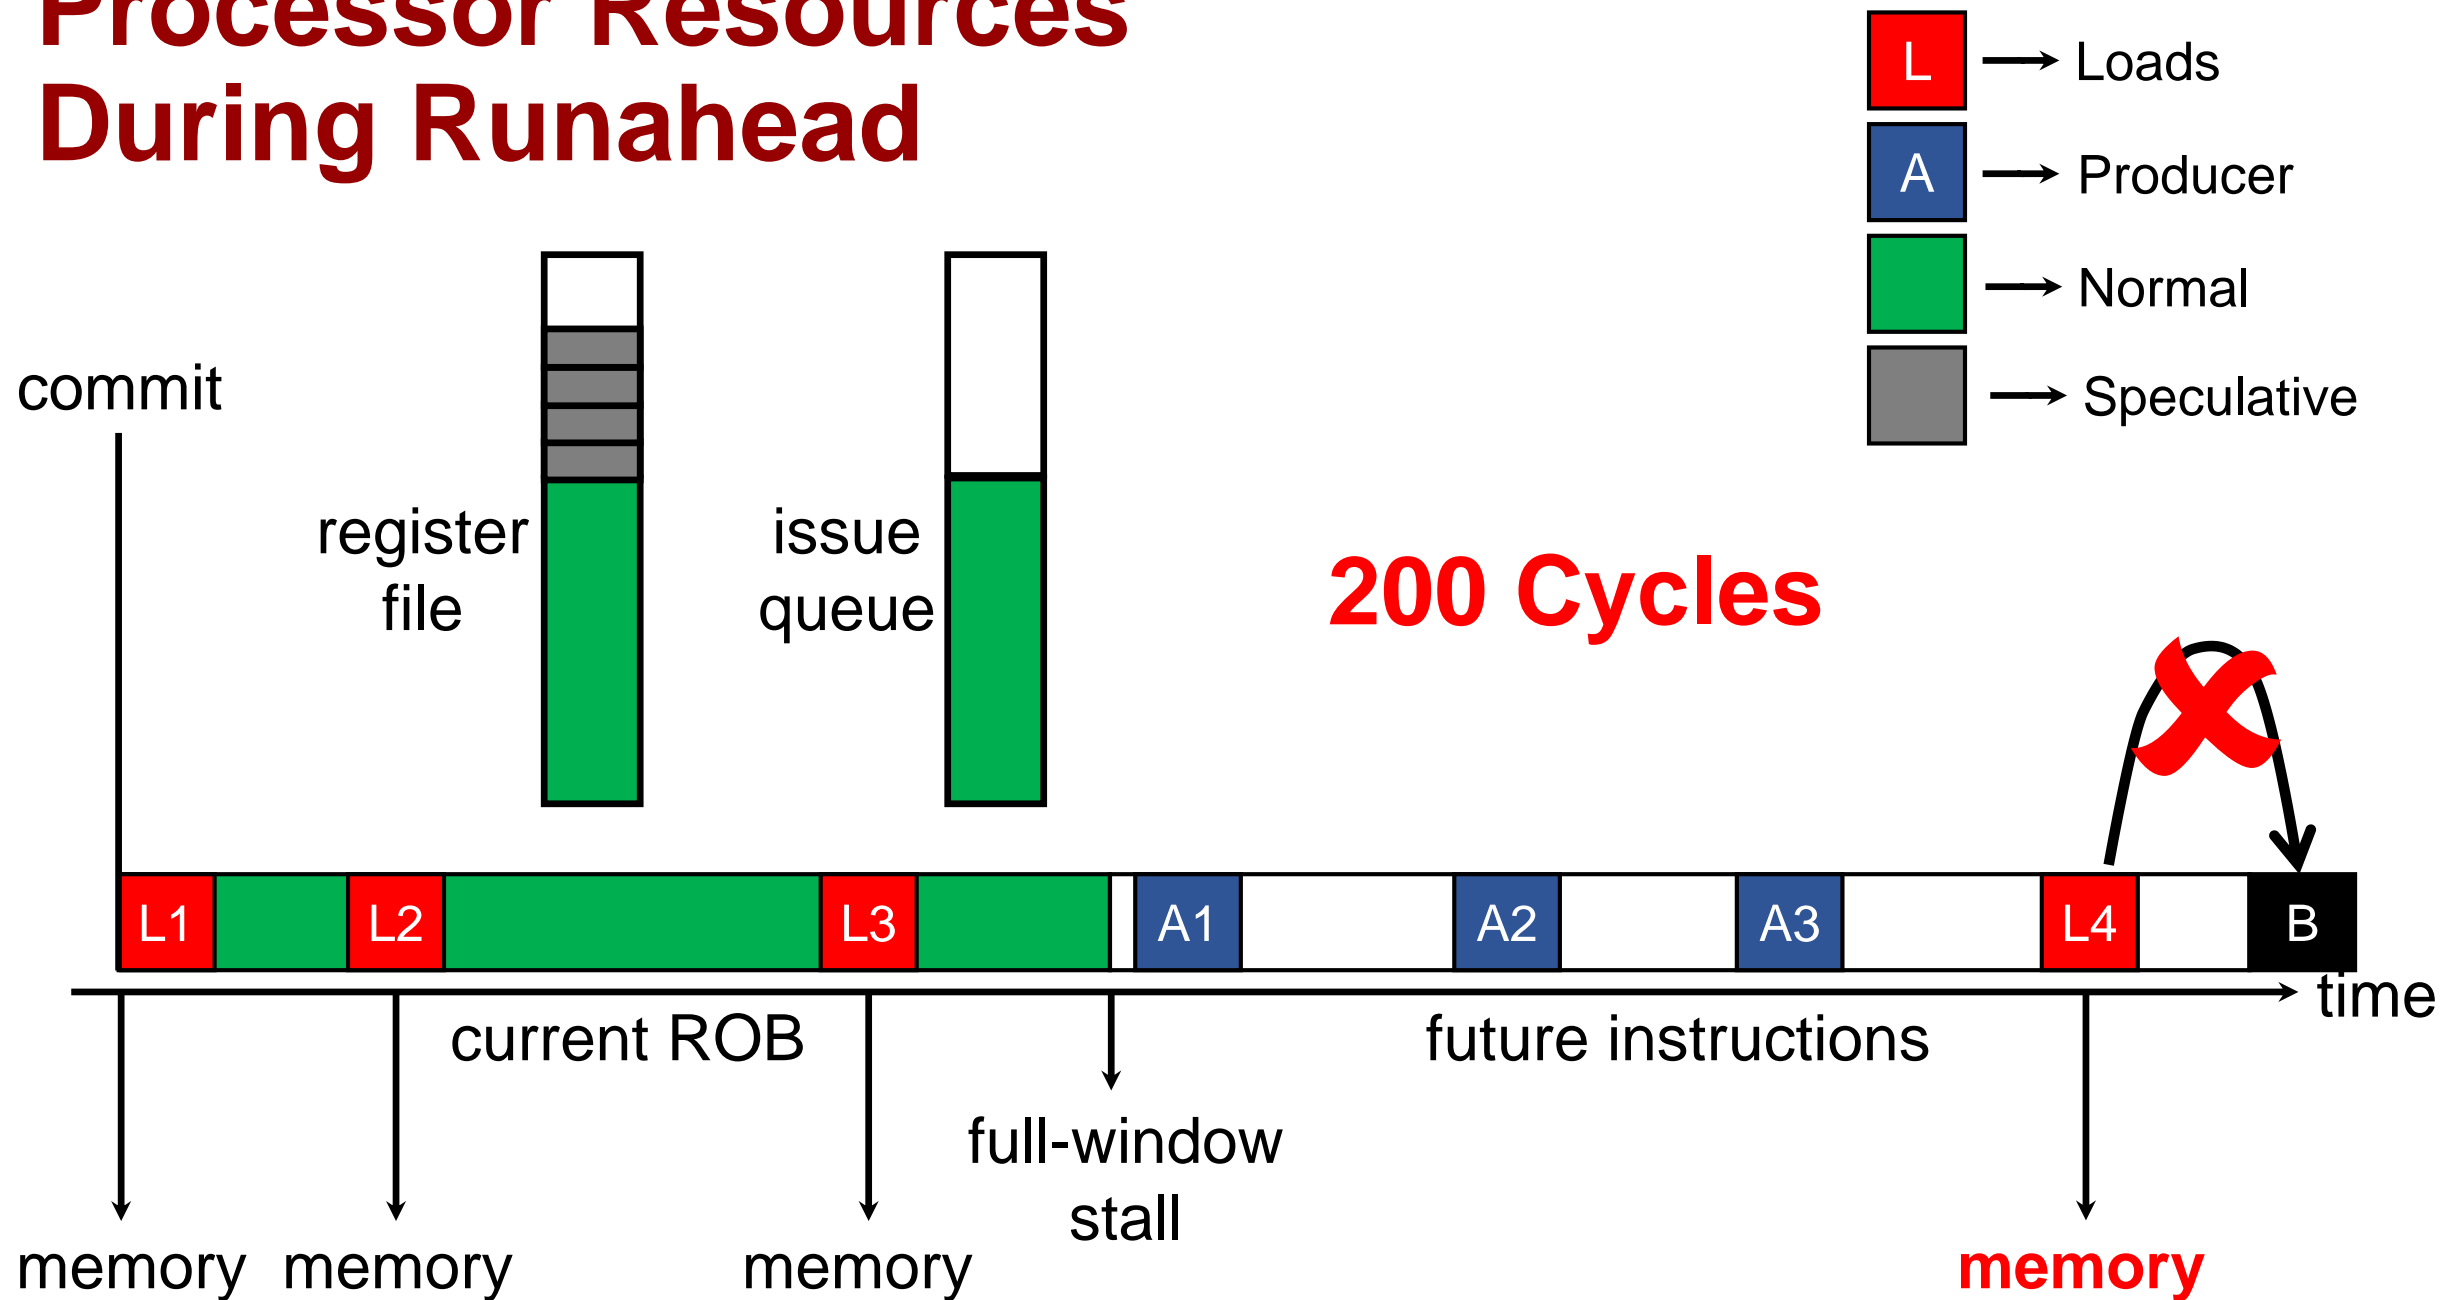
Processor Resources During Runahead



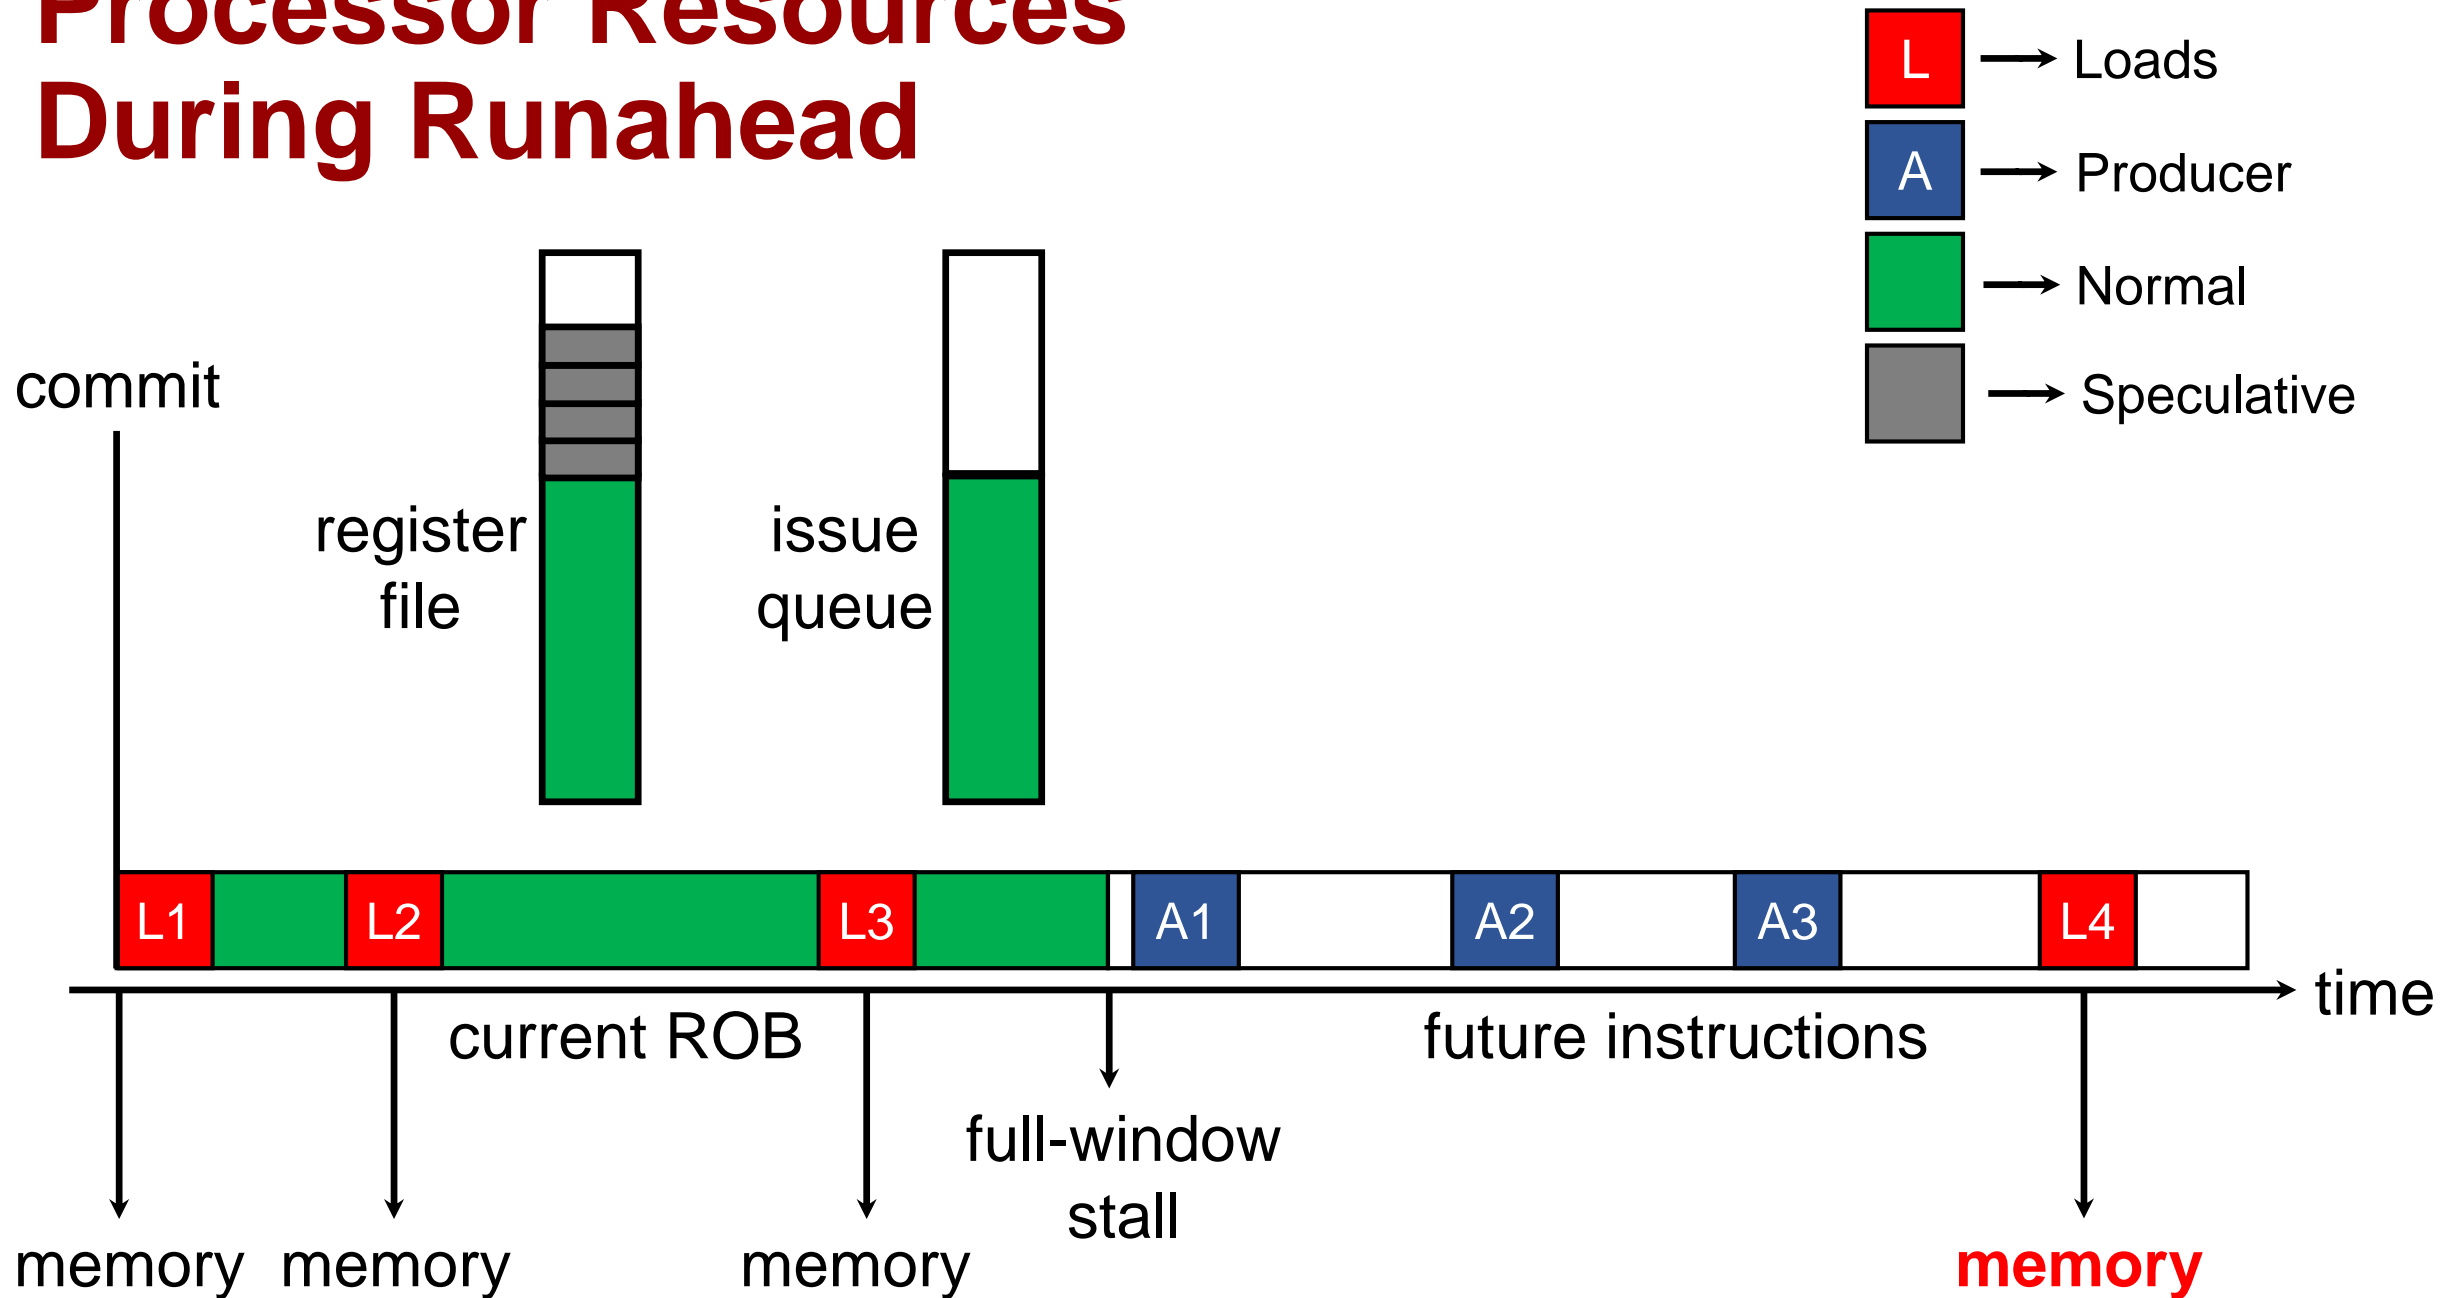
Processor Resources During Runahead



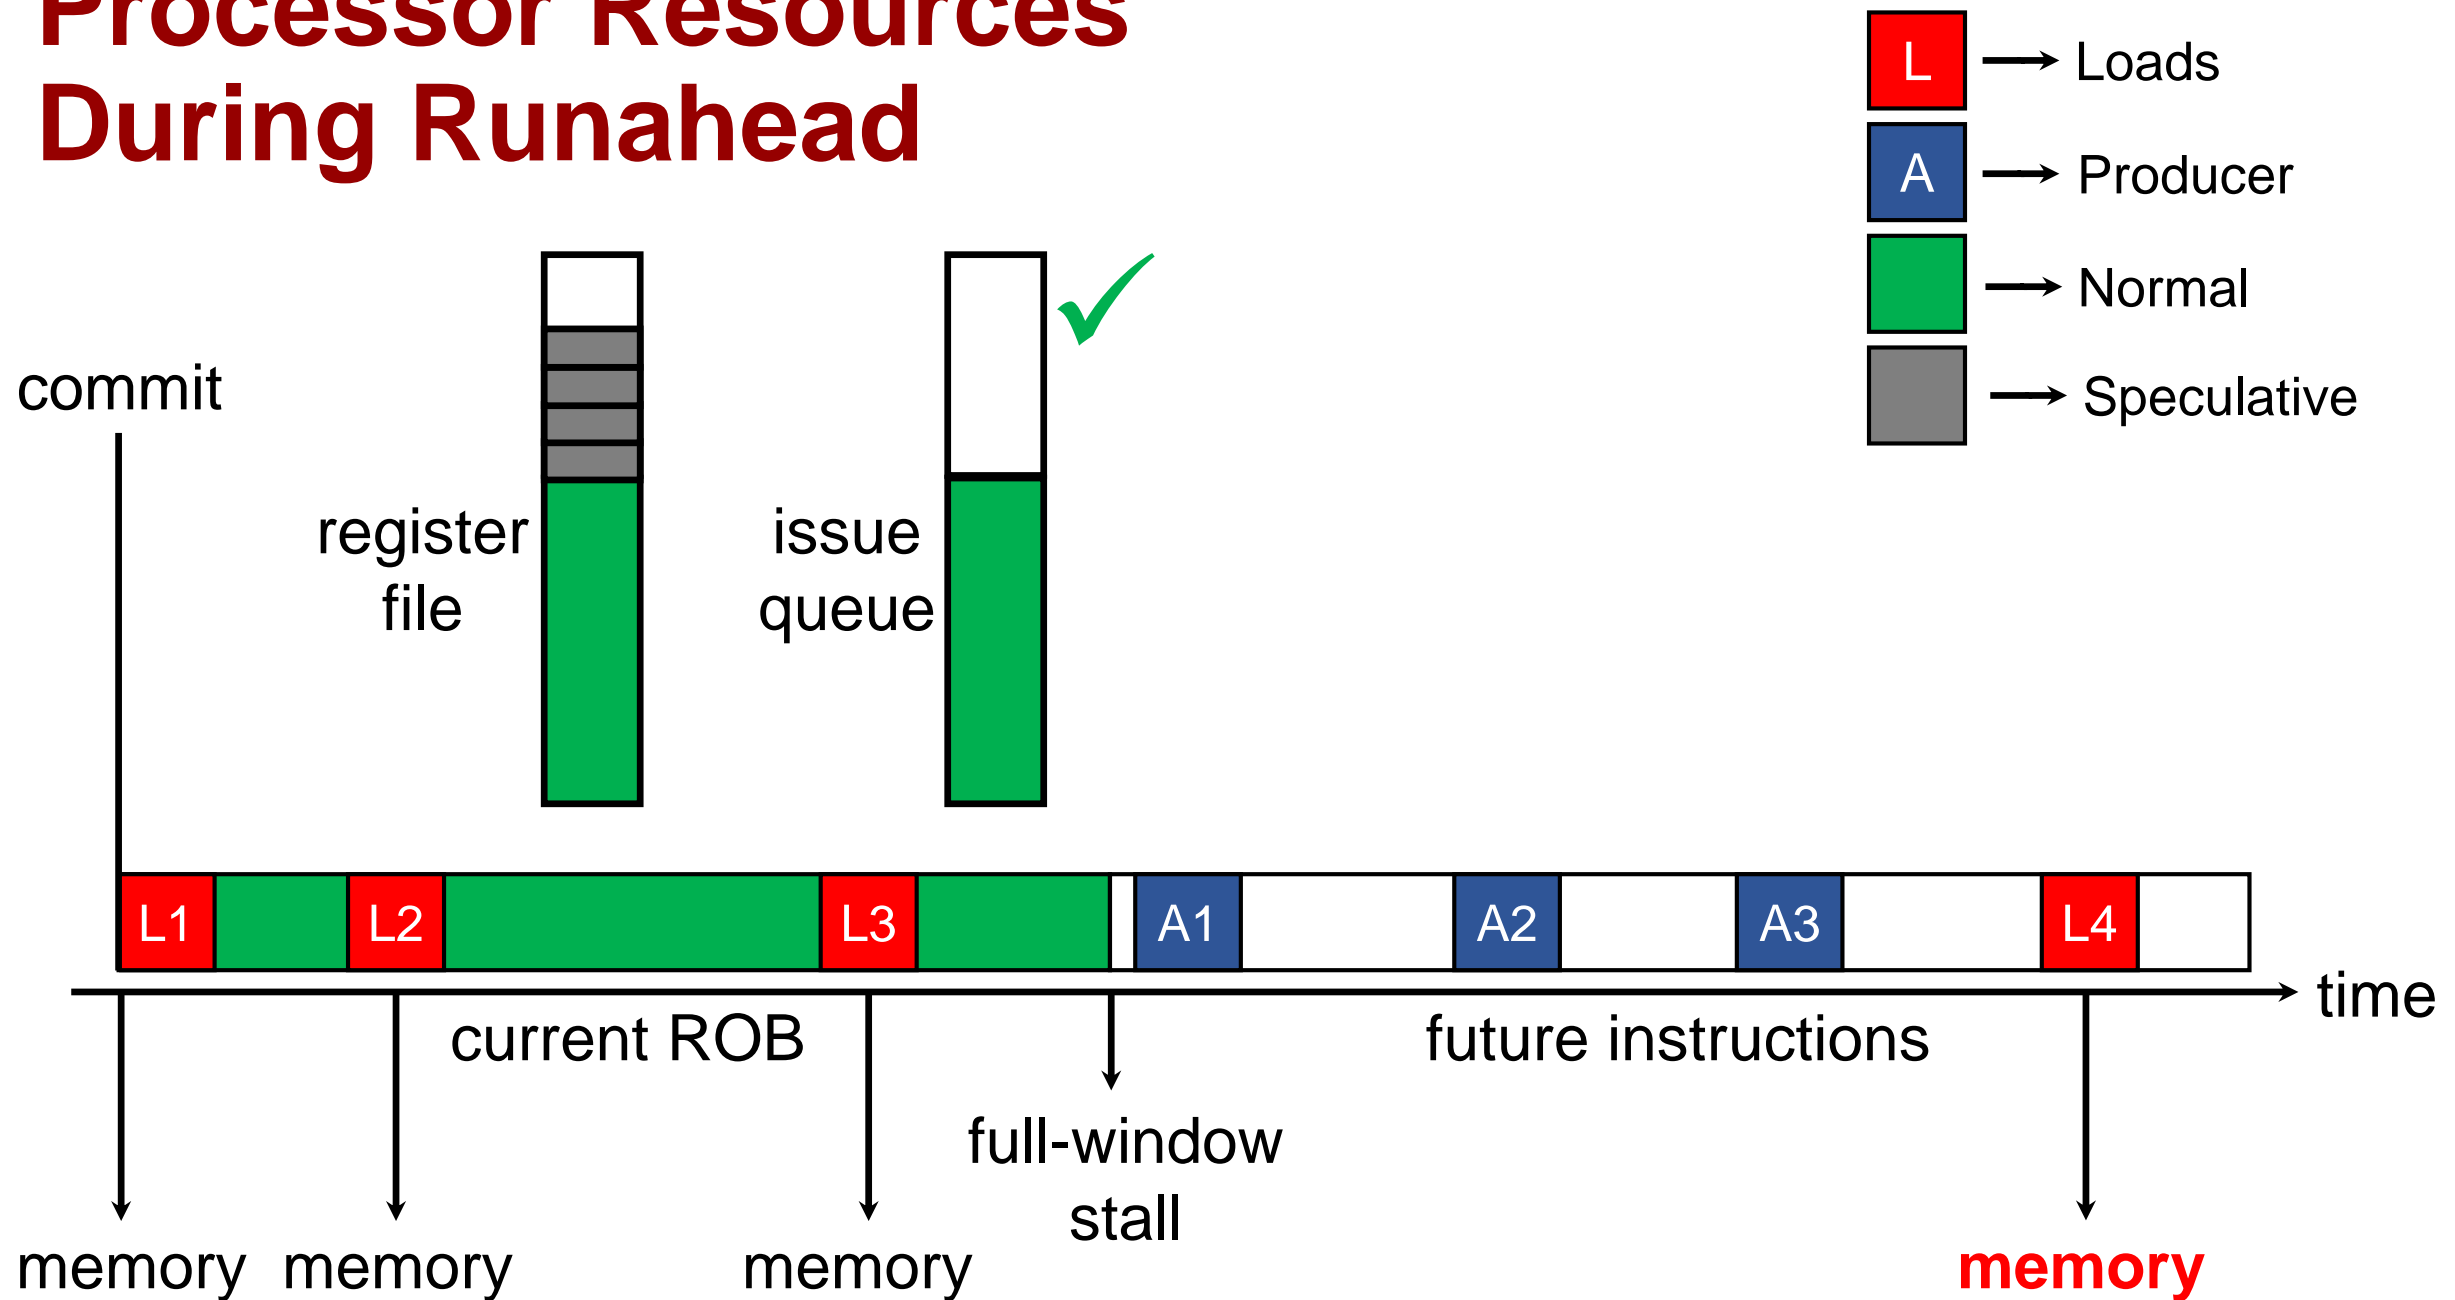
Processor Resources During Runahead



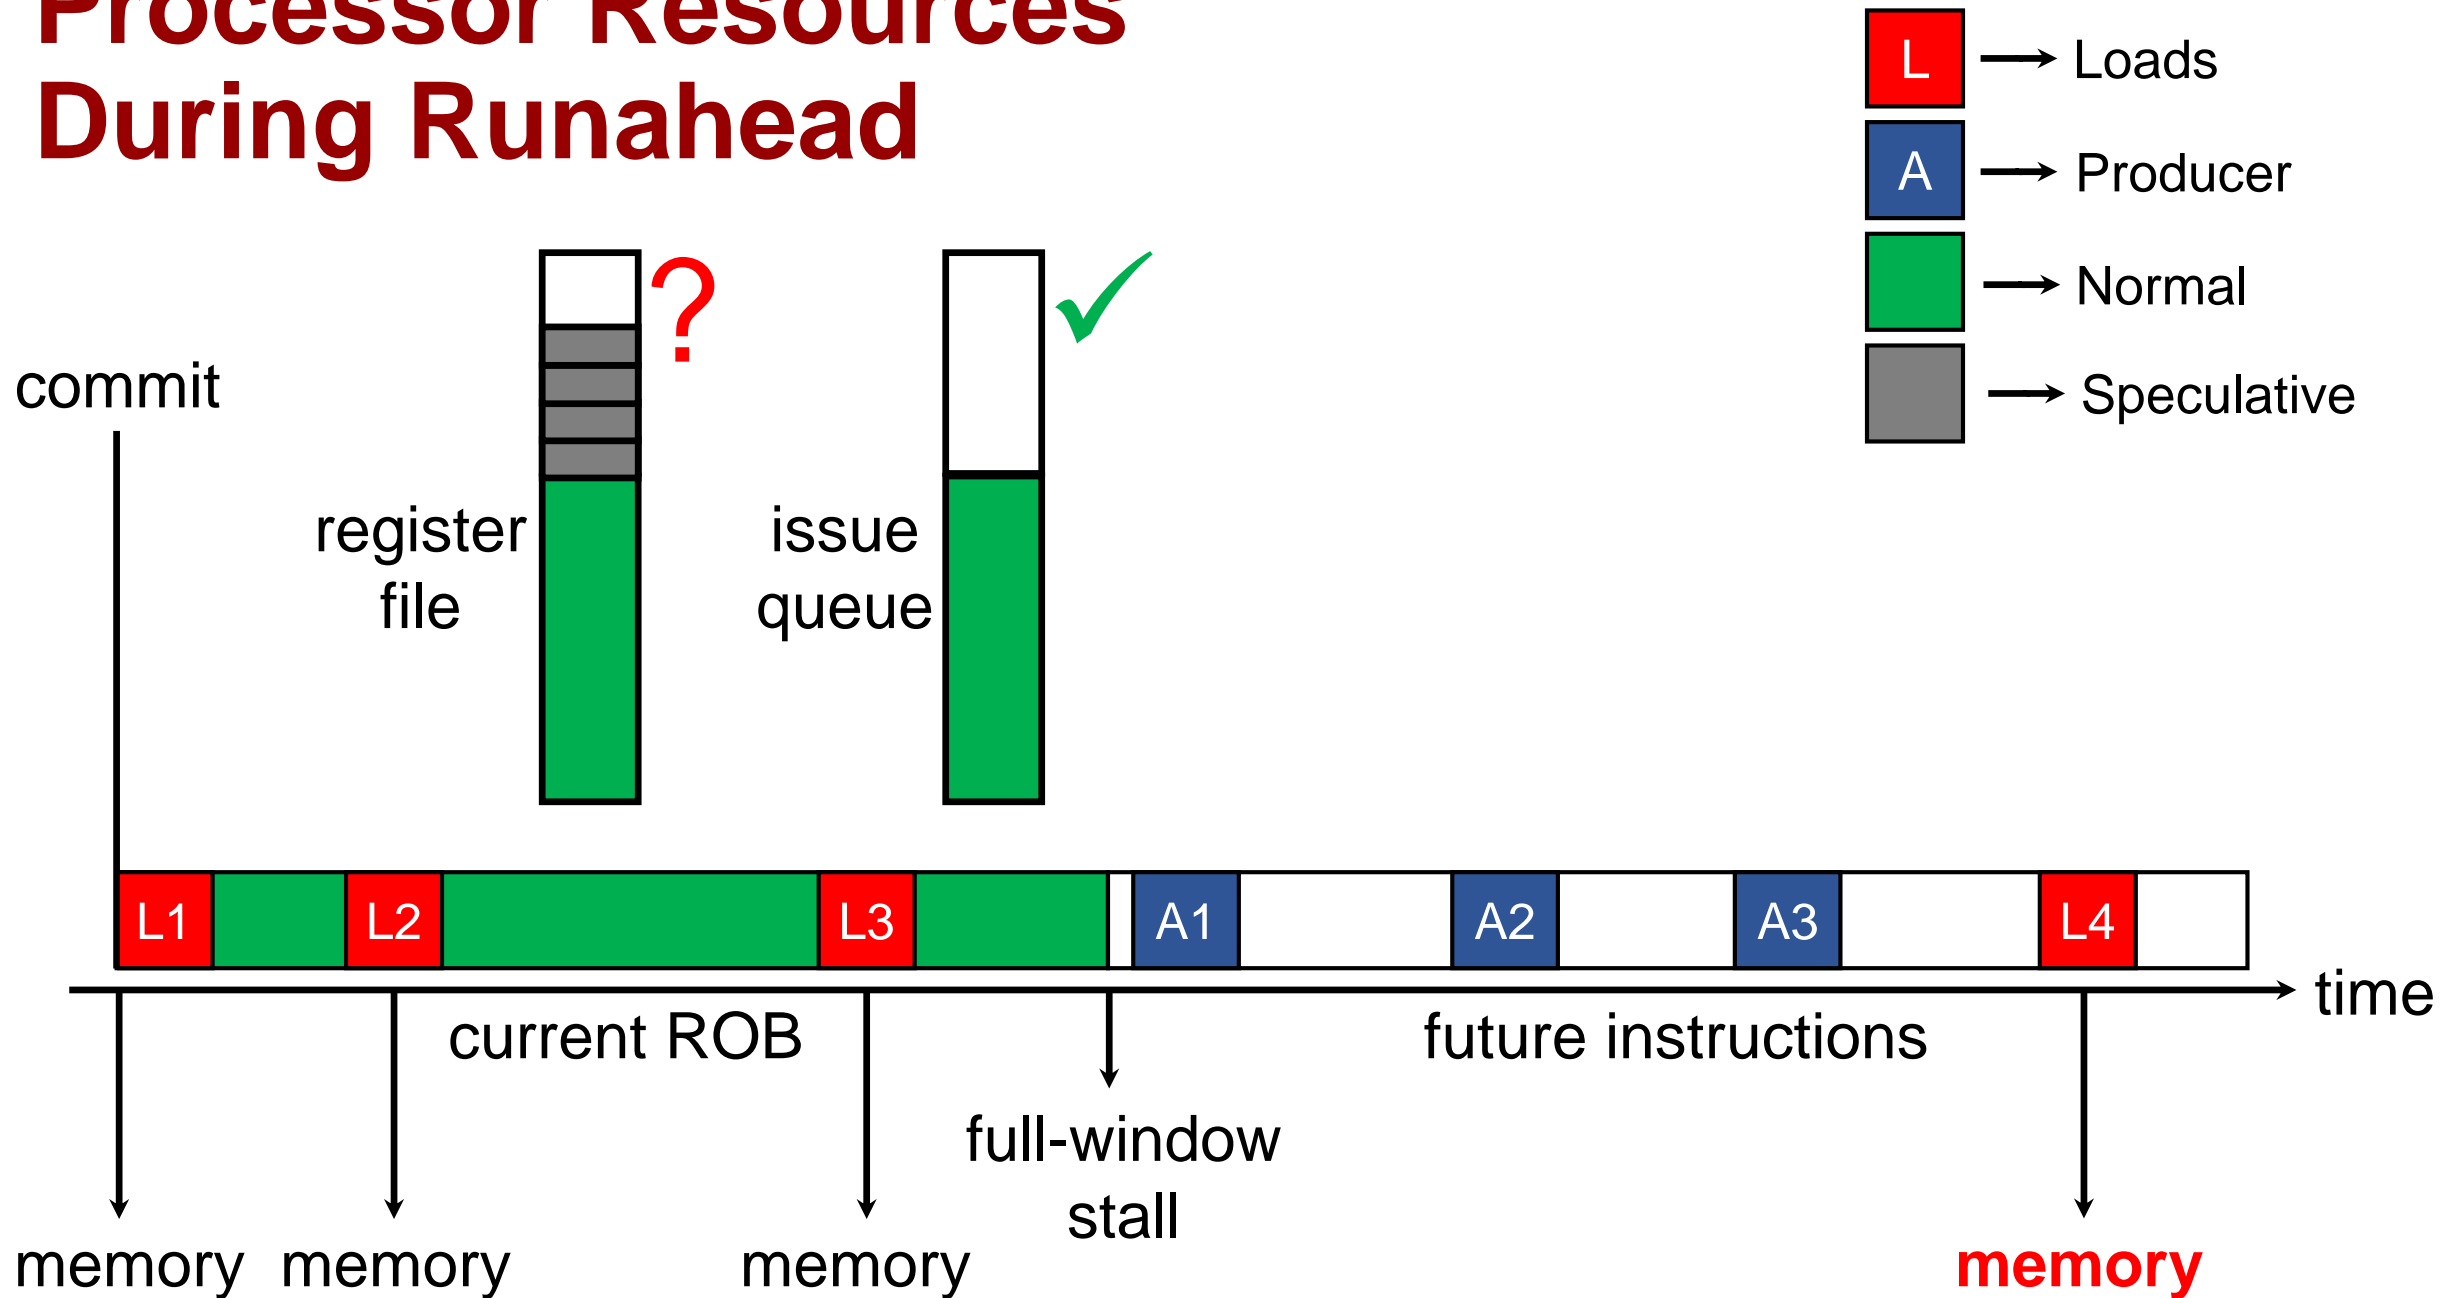
Processor Resources During Runahead



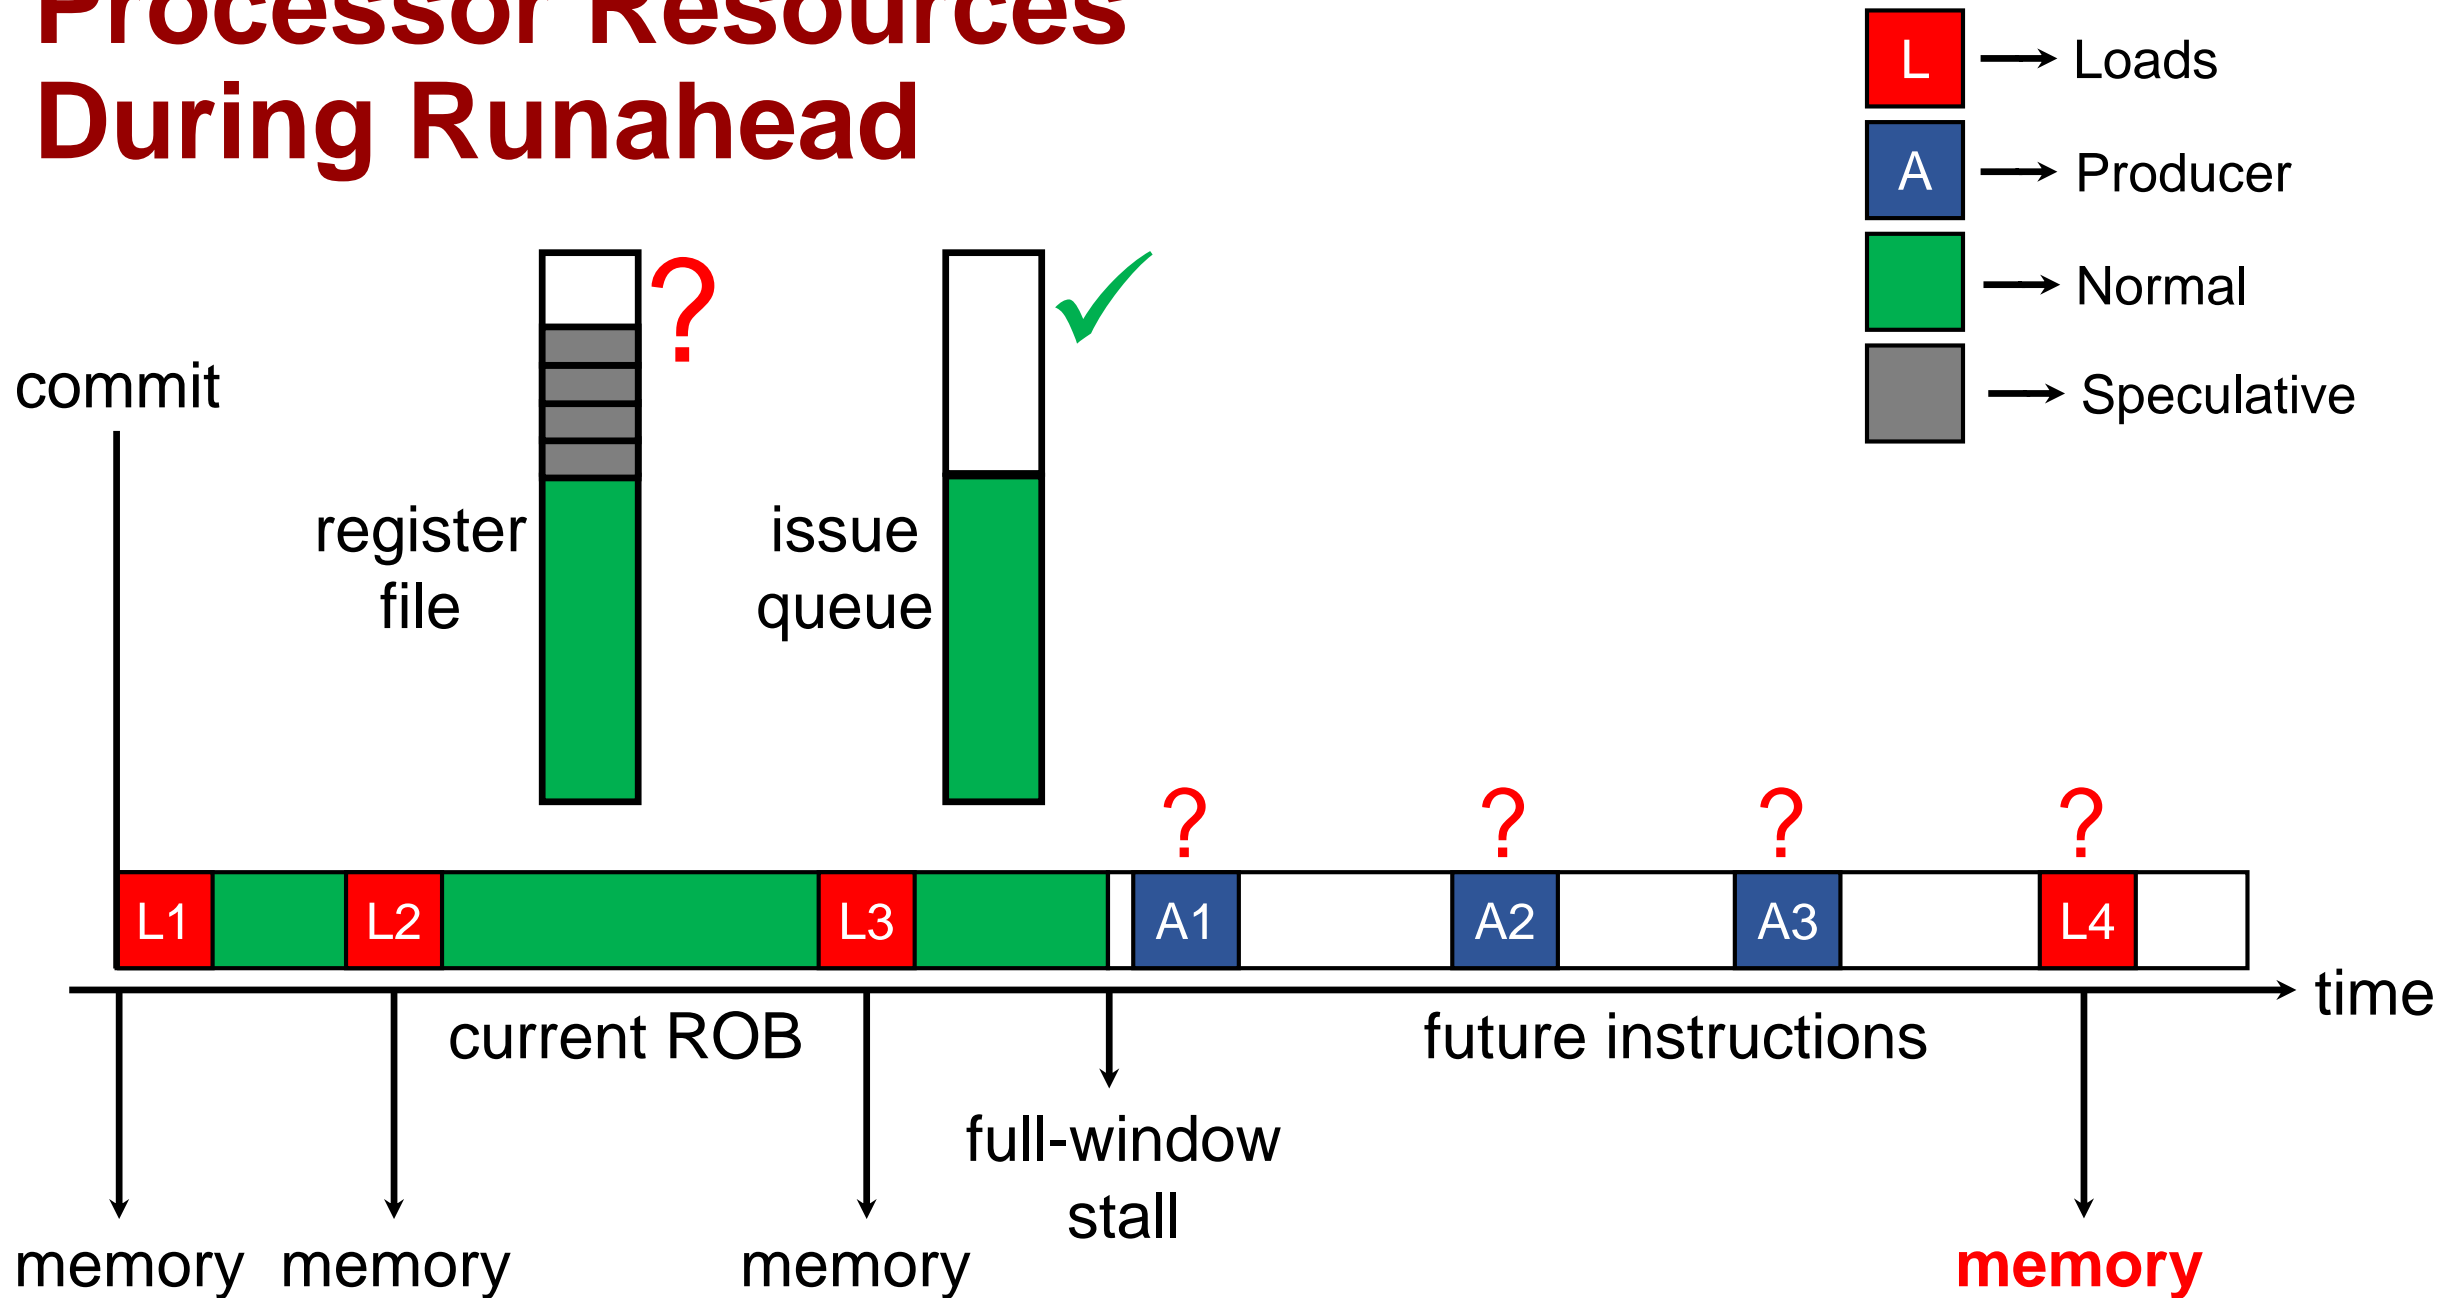
Processor Resources During Runahead



Processor Resources During Runahead



Processor Resources During Runahead



Two Key Questions

Two Key Questions


1. How to identify **only useful** instructions?

Two Key Questions

1. How to identify **only useful** instructions?
2. How to **recycle** (physical) **registers**?

Two Key Questions

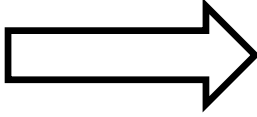
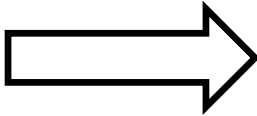
1. How to identify **only useful** instructions?



Iterative Backward
Dependency Analysis (IBDA)


2. How to **recycle**
(physical) **registers**?

Two Key Questions

1. How to identify **only useful** instructions?  Iterative Backward Dependency Analysis (IBDA)
2. How to **recycle** (physical) **registers**?  Runahead Register Reclamation

Iteratively Identifying the Stalling Slices

Iteratively Identifying the Stalling Slices



Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

A1	r2	←	r1
A2	r3	←	r2
A3	r4	←	r3
L4	r5	←	r4

Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch.
register

A1	r2	←	r1
A2	r3	←	r2
A3	r4	←	r3
L4	r5	←	r4

Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

	Arch. register		Phy. register
--	-------------------	--	------------------

A1	r2	←	r1
A2	r3	←	r2
A3	r4	←	r3
L4	r5	←	r4

Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch.
register Phy.
register


r1	P1
r2	P2
r3	P3
r4	P4
r5	P5

A1	r2	←	r1
A2	r3	←	r2
A3	r4	←	r3
L4	r5	←	r4

Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. Phy.
register register Last-writer
 instruction




r1	P1	
r2	P2	
r3	P3	
r4	P4	
r5	P5	

Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. Phy. Last-writer
register register instruction

r1	P1	A0
r2	P2	
r3	P3	
r4	P4	
r5	P5	




Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. Phy. Last-writer
register register instruction

r1	P1	A0
r2	P2	A1
r3	P3	
r4	P4	
r5	P5	




Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. Phy. Last-writer
register register instruction

r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	
r5	P5	




Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. Phy. Last-writer
register register instruction


r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	



Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4




Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4

Stalling Slice Table (SST)




Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4

Stalling Slice Table (SST)



Iteration-1:
L4 stalls the window

Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4

Stalling Slice Table (SST)

L4

Iteration-1:
L4 stalls the window


Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4

Stalling Slice Table (SST)

L4




Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4

Stalling Slice Table (SST)

L4



Iteration-2:
L4 hits in the SST


Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4


Stalling Slice Table (SST)

L4



Iteration-2:

L4 hits in the SST



While renaming source r4, read A3


Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4


Stalling Slice Table (SST)

L4
A3



Iteration-2:

L4 hits in the SST



While renaming source r4, read A3


Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4

Stalling Slice Table (SST)

L4
A3



Iteration-3:
A3 hits in the SST


Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4


Stalling Slice Table (SST)

L4
A3



Iteration-3:

A3 hits in the SST



While renaming source r3, read A2


Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4


Stalling Slice Table (SST)

L4
A3
A2



Iteration-3:

A3 hits in the SST



While renaming source r3, read A2


Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4

Stalling Slice Table (SST)

L4
A3
A2



Iteration-4:
A2 hits in the SST


Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4


Stalling Slice Table (SST)

L4
A3
A2



Iteration-4:

A2 hits in the SST



While renaming source r2, read A1


Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)

Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4


Stalling Slice Table (SST)

L4
A3
A2
A1



Iteration-4:

A2 hits in the SST




While renaming source r2, read A1

Iteratively Identifying the Stalling Slices

Register Allocation Table (RAT)


Arch. register	Phy. register	Last-writer instruction
r1	P1	A0
r2	P2	A1
r3	P3	A2
r4	P4	A3
r5	P5	L4

Stalling Slice Table (SST)



Iteration-4:

A2 hits in the SST



While renaming source r2, read A1

Runahead Register Reclamation

Runahead Register Reclamation

normal mode

Runahead Register Reclamation

normal mode

	instruction	dest	src1	src2	OldPhy register

Runahead Register Reclamation

normal mode

	instruction	dest	src1	src2	OldPhy register
I1	add r1 \leftarrow r2, r3	P1	P2	P3	P0

Runahead Register Reclamation

normal mode

	instruction	dest	src1	src2	OldPhy register
I1	add r1 \leftarrow r2, r3	P1	P2	P3	P0
I2	mul r2 \leftarrow r1, r4	P5	P1	P4	P2

Runahead Register Reclamation

normal mode

	instruction	dest	src1	src2	OldPhy register
I1	add r1 \leftarrow r2, r3	P1	P2	P3	P0
I2	mul r2 \leftarrow r1, r4	P5	P1	P4	P2
I3	ld r1 \leftarrow mem[x]	P6			P1

Runahead Register Reclamation

normal mode

	instruction	dest	src1	src2	OldPhy register
I1	add r1 \leftarrow r2, r3	P1	P2	P3	P0
I2	mul r2 \leftarrow r1, r4	P5	P1	P4	P2
I3	ld r1 \leftarrow mem[x]	P6			P1
I4	add r2 \leftarrow r1, r3	P7	P6	P3	P5

Runahead Register Reclamation

normal mode

	instruction	dest	src1	src2	OldPhy register
I1	add r1 \leftarrow r2, r3	P1	P2	P3	P0
I2	mul r2 \leftarrow r1, r4	P5	P1	P4	P2
I3	ld r1 \leftarrow mem[x]	P6			P1
I4	add r2 \leftarrow r1, r3	P7	P6	P3	P5
I5	add r2 \leftarrow r4, r5	P9	P4	P8	P7

Runahead Register Reclamation

normal mode

	instruction	dest	src1	src2	OldPhy register
I1	add r1 \leftarrow r2, r3	P1	P2	P3	P0
I2	mul r2 \leftarrow r1, r4	P5	P1	P4	P2
I3	ld r1 \leftarrow mem[x]	P6			P1
I4	add r2 \leftarrow r1, r3	P7	P6	P3	P5
I5	add r2 \leftarrow r4, r5	P9	P4	P8	P7
I6	sub r1 \leftarrow r2, r6	P11	P9	P10	P6

Runahead Register Reclamation

normal mode

runahead mode

	instruction	dest	src1	src2	OldPhy register
I1	add r1 \leftarrow r2, r3	P1	P2	P3	P0
I2	mul r2 \leftarrow r1, r4	P5	P1	P4	P2
I3	ld r1 \leftarrow mem[x]	P6			P1
I4	add r2 \leftarrow r1, r3	P7	P6	P3	P5
I5	add r2 \leftarrow r4, r5	P9	P4	P8	P7
I6	sub r1 \leftarrow r2, r6	P11	P9	P10	P6

Runahead Register Reclamation

normal mode

	instruction	dest	src1	src2	OldPhy register
I1	add r1 \leftarrow r2, r3	P1	P2	P3	P0
I2	mul r2 \leftarrow r1, r4	P5	P1	P4	P2
I3	ld r1 \leftarrow mem[x]	P6			P1
I4	add r2 \leftarrow r1, r3	P7	P6	P3	P5
I5	add r2 \leftarrow r4, r5	P9	P4	P8	P7
I6	sub r1 \leftarrow r2, r6	P11	P9	P10	P6

runahead mode

	OldPhy register
I1	P0
I2	P2
I3	P1
I4	P5
I5	P7
I6	P6

Runahead Register Reclamation

runahead mode

	OldPhy register
I1	P0
I2	P2
I3	P1
I4	P5
I5	P7
I6	P6

Runahead Register Reclamation

runahead mode

	OldPhy register
I1	P0
I2	P2
I3	P1
I4	P5
I5	P7
I6	P6

Precise Register Deallocation Queue (PRDQ)

Runahead Register Reclamation

runahead mode

	OldPhy register
I1	P0
I2	P2
I3	P1
I4	P5
I5	P7
I6	P6

dispatch 

Precise Register Deallocation Queue (PRDQ)

Runahead Register Reclamation

runahead mode

	OldPhy register	Executed ?
I1	P0	
I2	P2	
I3	P1	
I4	P5	
I5	P7	
I6	P6	

dispatch →

Precise Register Deallocation Queue (PRDQ)

Runahead Register Reclamation

runahead mode

	OldPhy register	Executed ?
I1	P0	0
I2	P2	0
I3	P1	0
I4	P5	0
I5	P7	0
I6	P6	0

dispatch →

Precise Register Deallocation Queue (PRDQ)

Runahead Register Reclamation

runahead mode

	OldPhy register	Executed ?
I1	P0	0
I2	P2	0
I3	P1	0
I4	P5	0
I5	P7	0
I6	P6	0

dispatch →

← execute

Precise Register Deallocation Queue (PRDQ)

Runahead Register Reclamation

runahead mode

	OldPhy register	Executed ?
I1	P0	1
I2	P2	0
I3	P1	0
I4	P5	0
I5	P7	0
I6	P6	0

dispatch →

← execute

Precise Register Deallocation Queue (PRDQ)

Runahead Register Reclamation

runahead mode

	OldPhy register	Executed ?
I1	P0	1
I2	P2	1
I3	P1	0
I4	P5	0
I5	P7	0
I6	P6	0

dispatch →

← execute

Precise Register Deallocation Queue (PRDQ)

Runahead Register Reclamation

runahead mode

	OldPhy register	Executed ?
I1	P0	1
I2	P2	1
I3	P1	0
I4	P5	0
I5	P7	1
I6	P6	0

dispatch →

← execute

Precise Register Deallocation Queue (PRDQ)

Runahead Register Reclamation

runahead mode

		OldPhy register	Executed ?	
	I1	P0 ✓	1	
	I2	P2 ✓	1	
dispatch →	I3	P1	0	← execute
	I4	P5	0	
	I5	P7	1	
	I6	P6	0	

Precise Register Deallocation Queue (PRDQ)

Runahead Register Reclamation

runahead mode

		OldPhy register	Executed ?	
	I1	P0 ✓	1	
	I2	P2 ✓	1	
dispatch →	I3	P1	0	← execute
	I4	P5	0	
	I5	P7 ✗	1	
	I6	P6	0	

Precise Register Deallocation Queue (PRDQ)

Runahead Register Reclamation

runahead mode

	OldPhy register	Executed ?
I1	P0 ✓	1
I2	P2 ✓	1
I3	P1	0
I4	P5	0
I5	P7 ✗	1
I6	P6	0

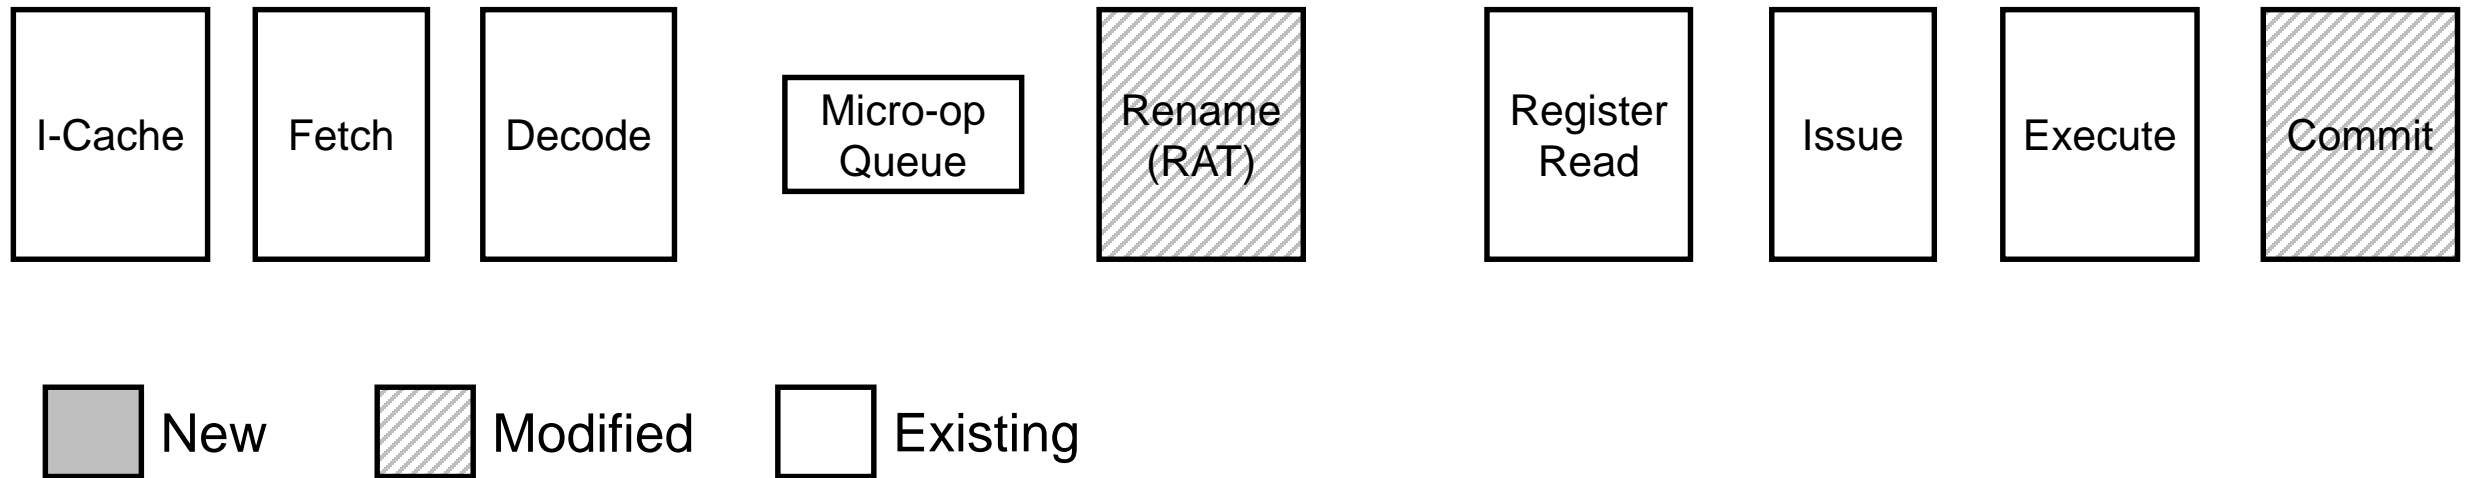
dispatch →

← execute

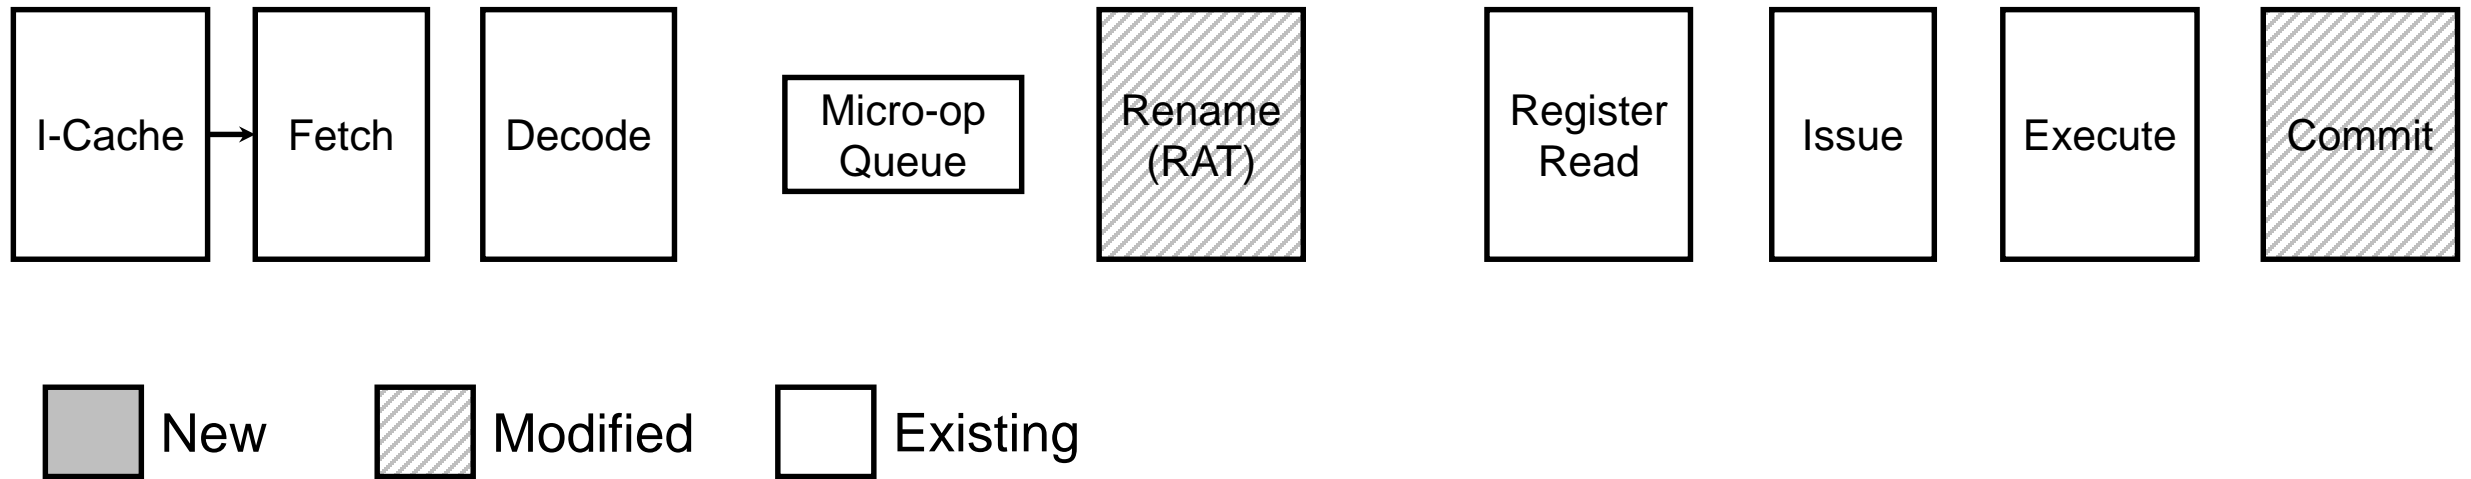
Precise Register Deallocation Queue (PRDQ)

Putting it All Together

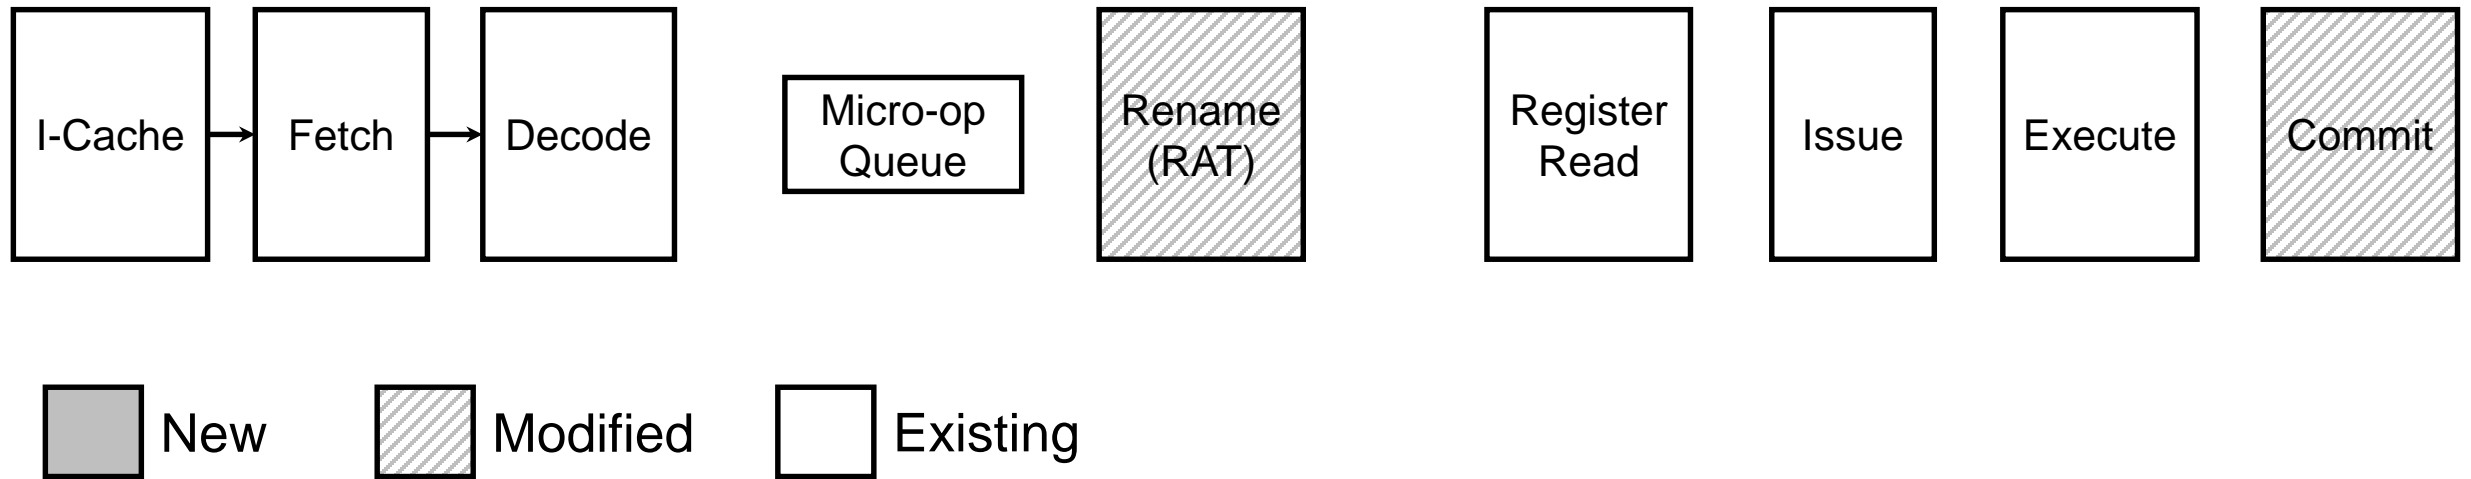
Putting it All Together



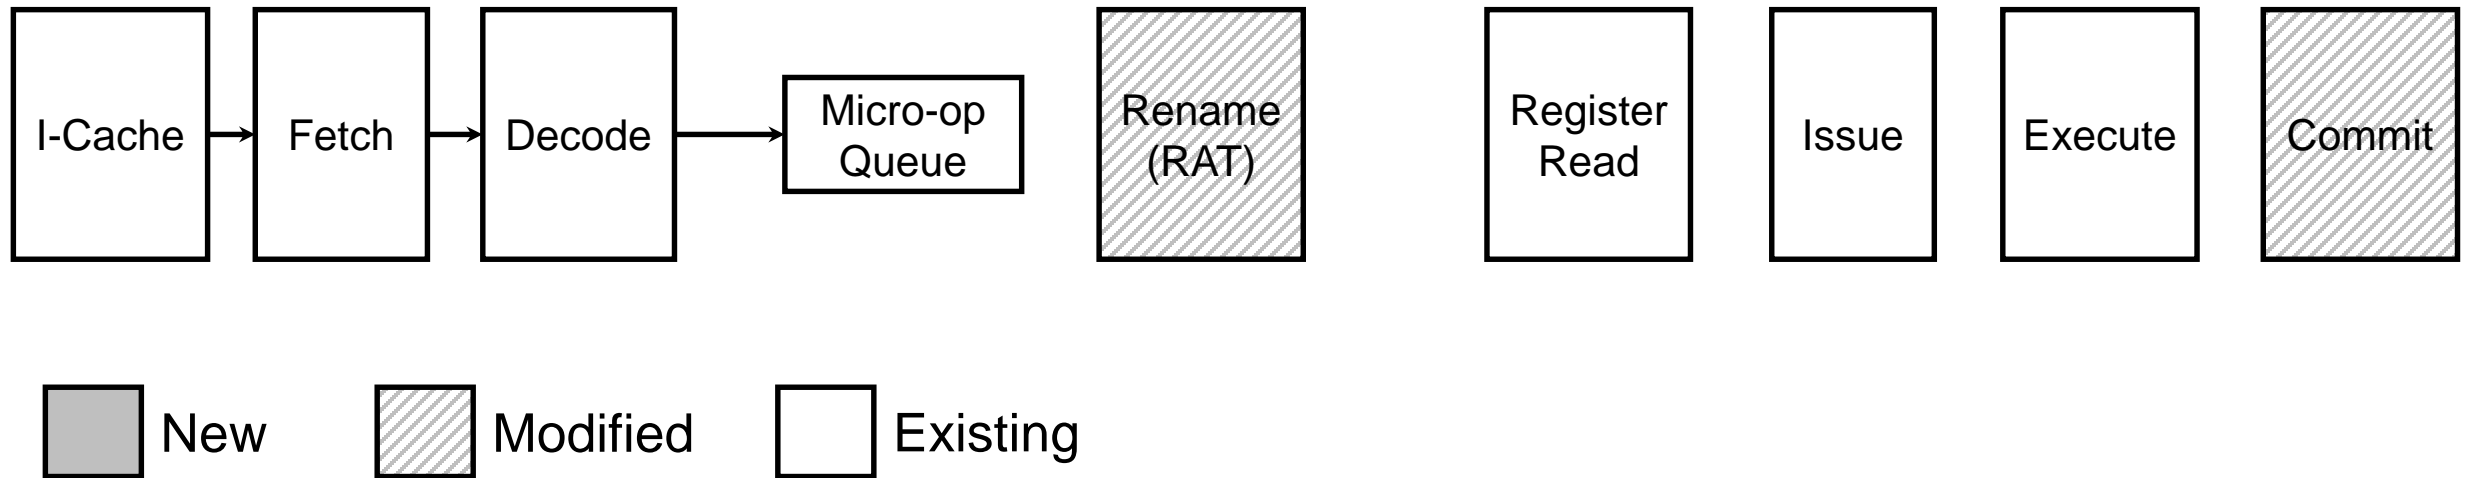
Putting it All Together



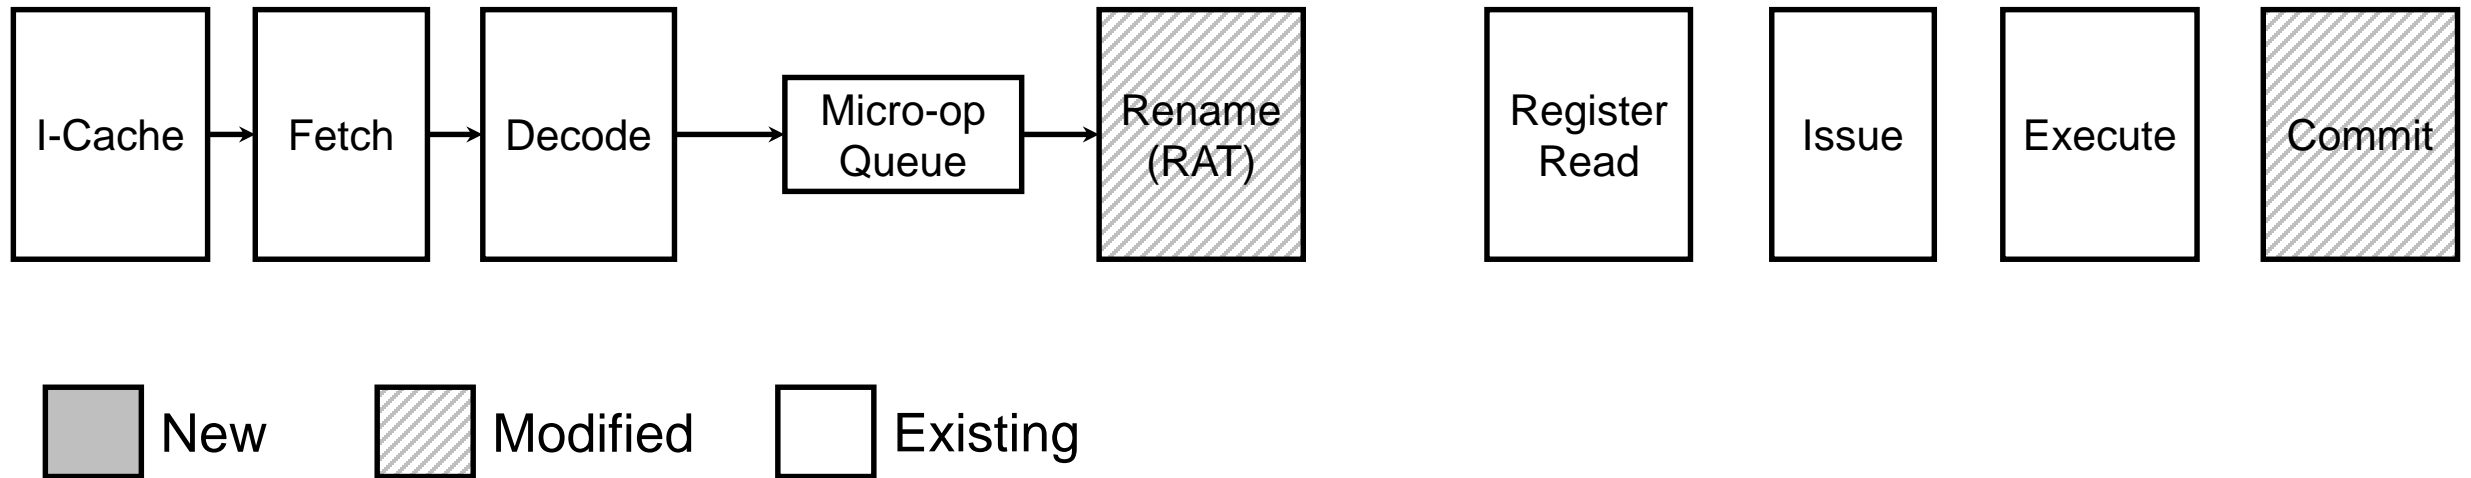
Putting it All Together



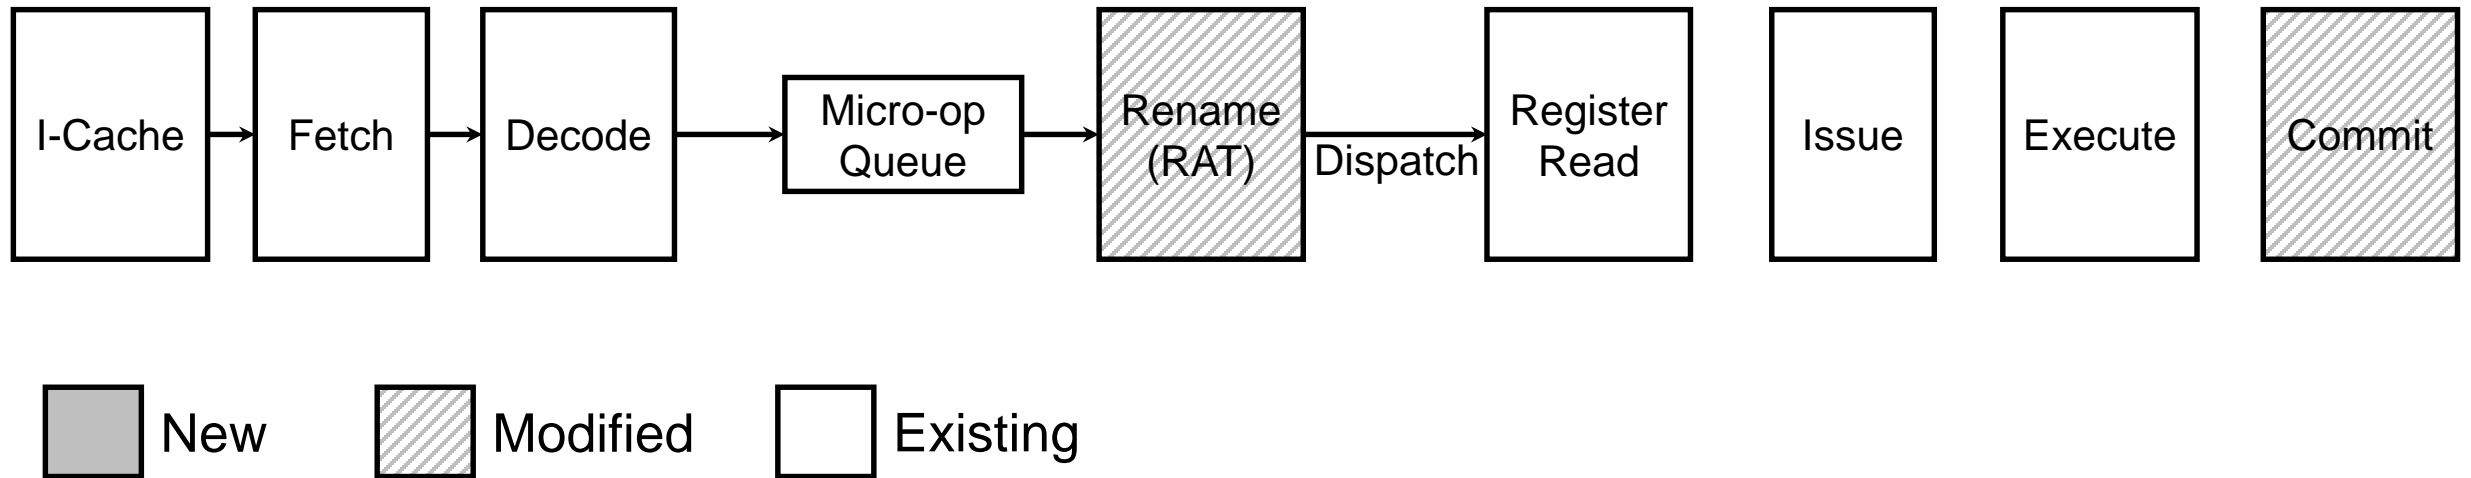
Putting it All Together



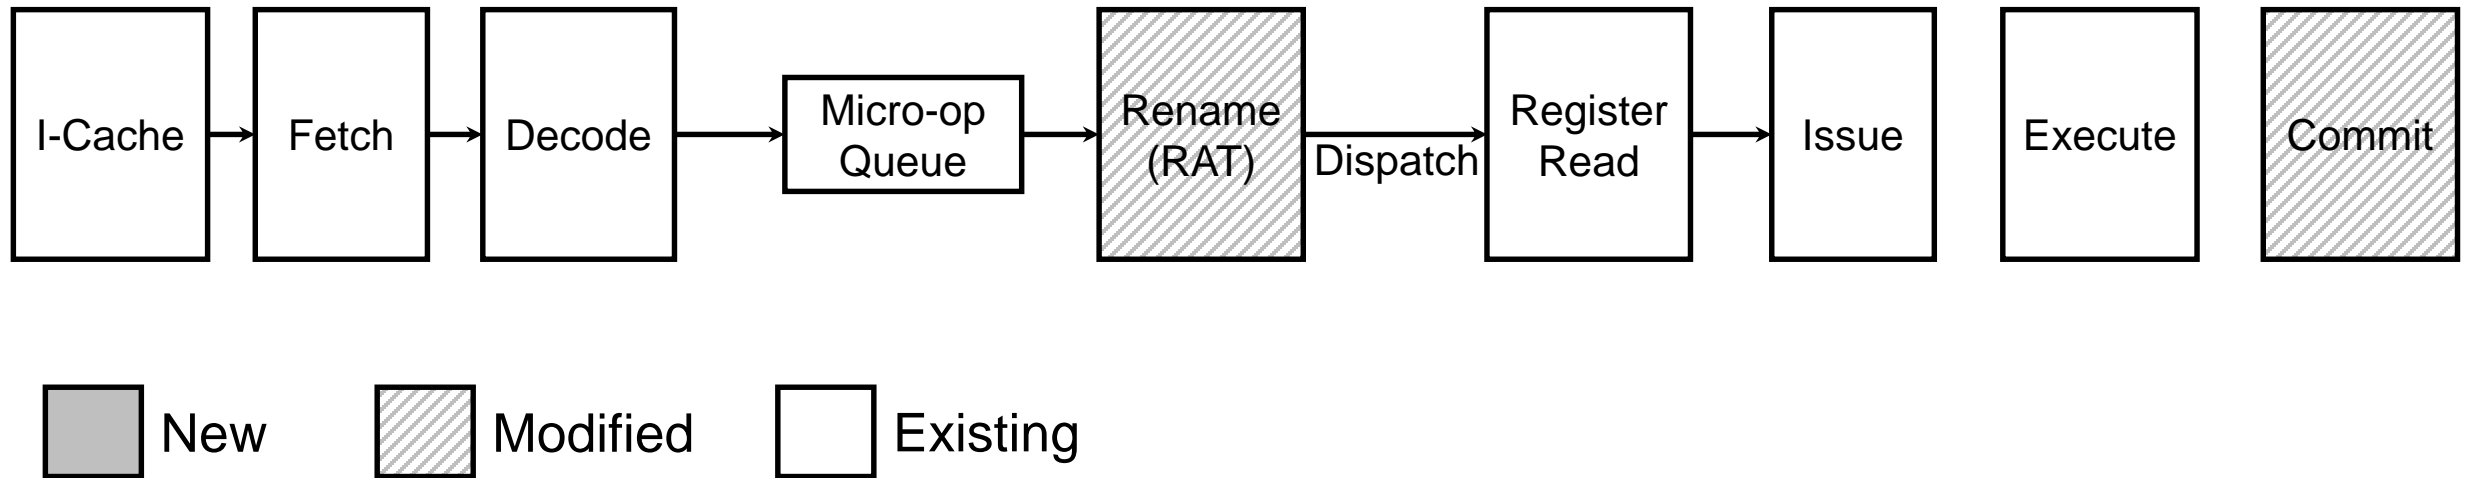
Putting it All Together



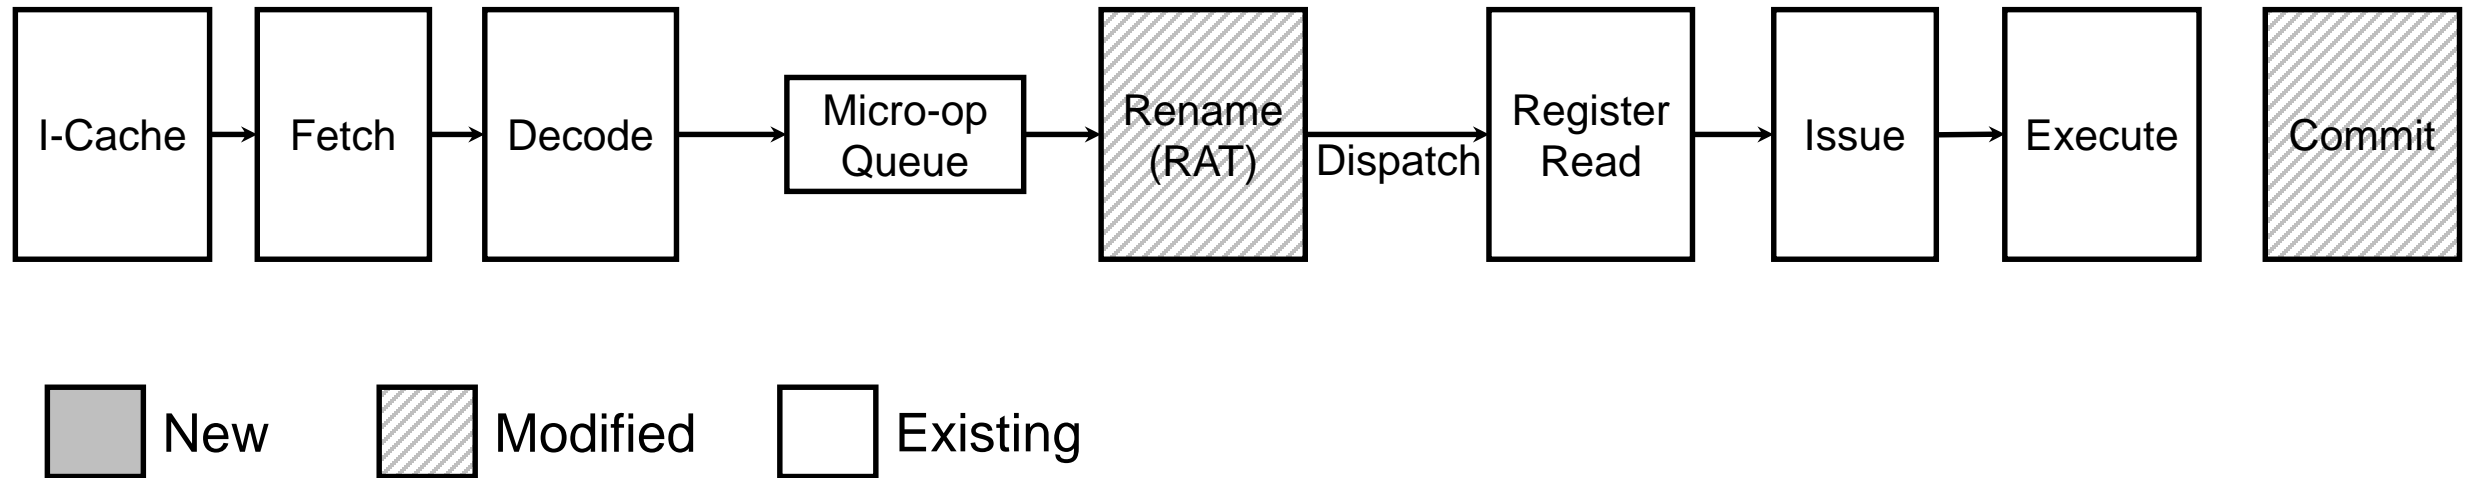
Putting it All Together



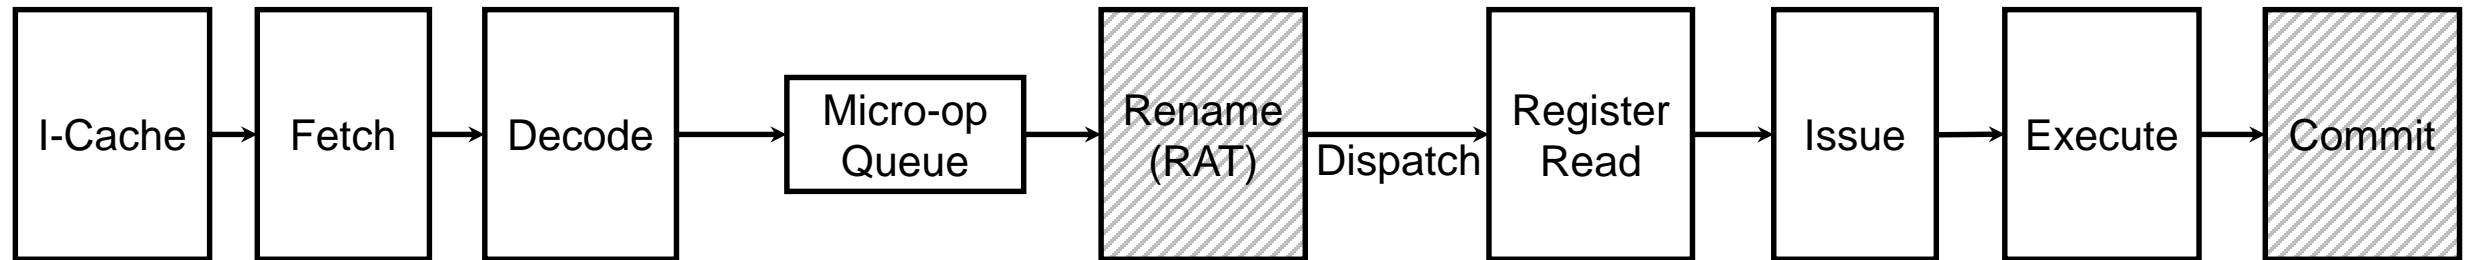
Putting it All Together



Putting it All Together



Putting it All Together



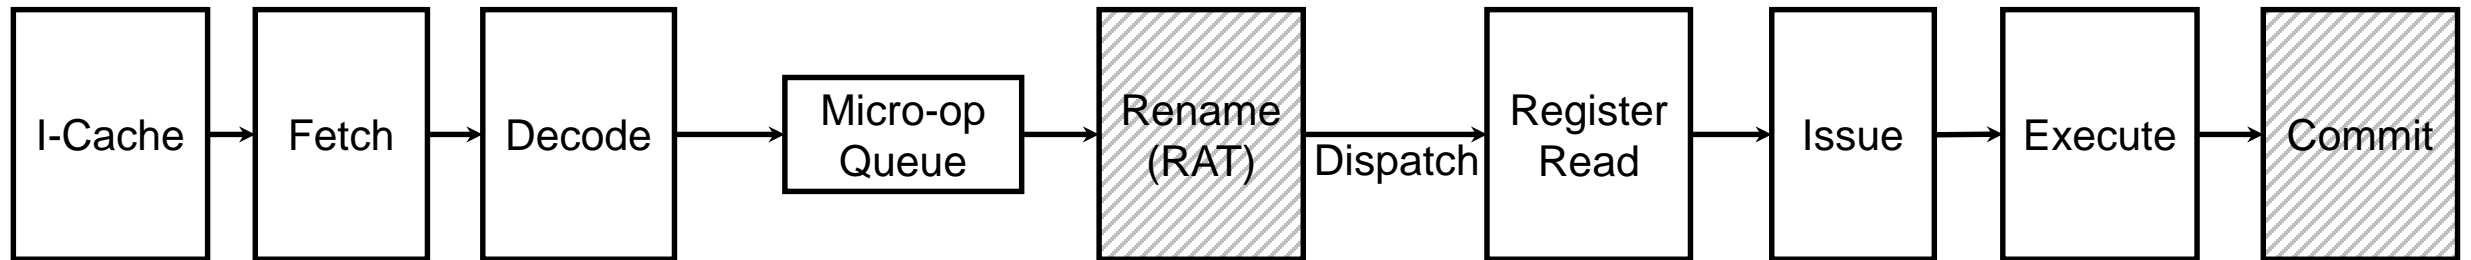
 New  Modified  Existing

→ Normal Mode

Putting it All Together

Stalling
Slice
Table

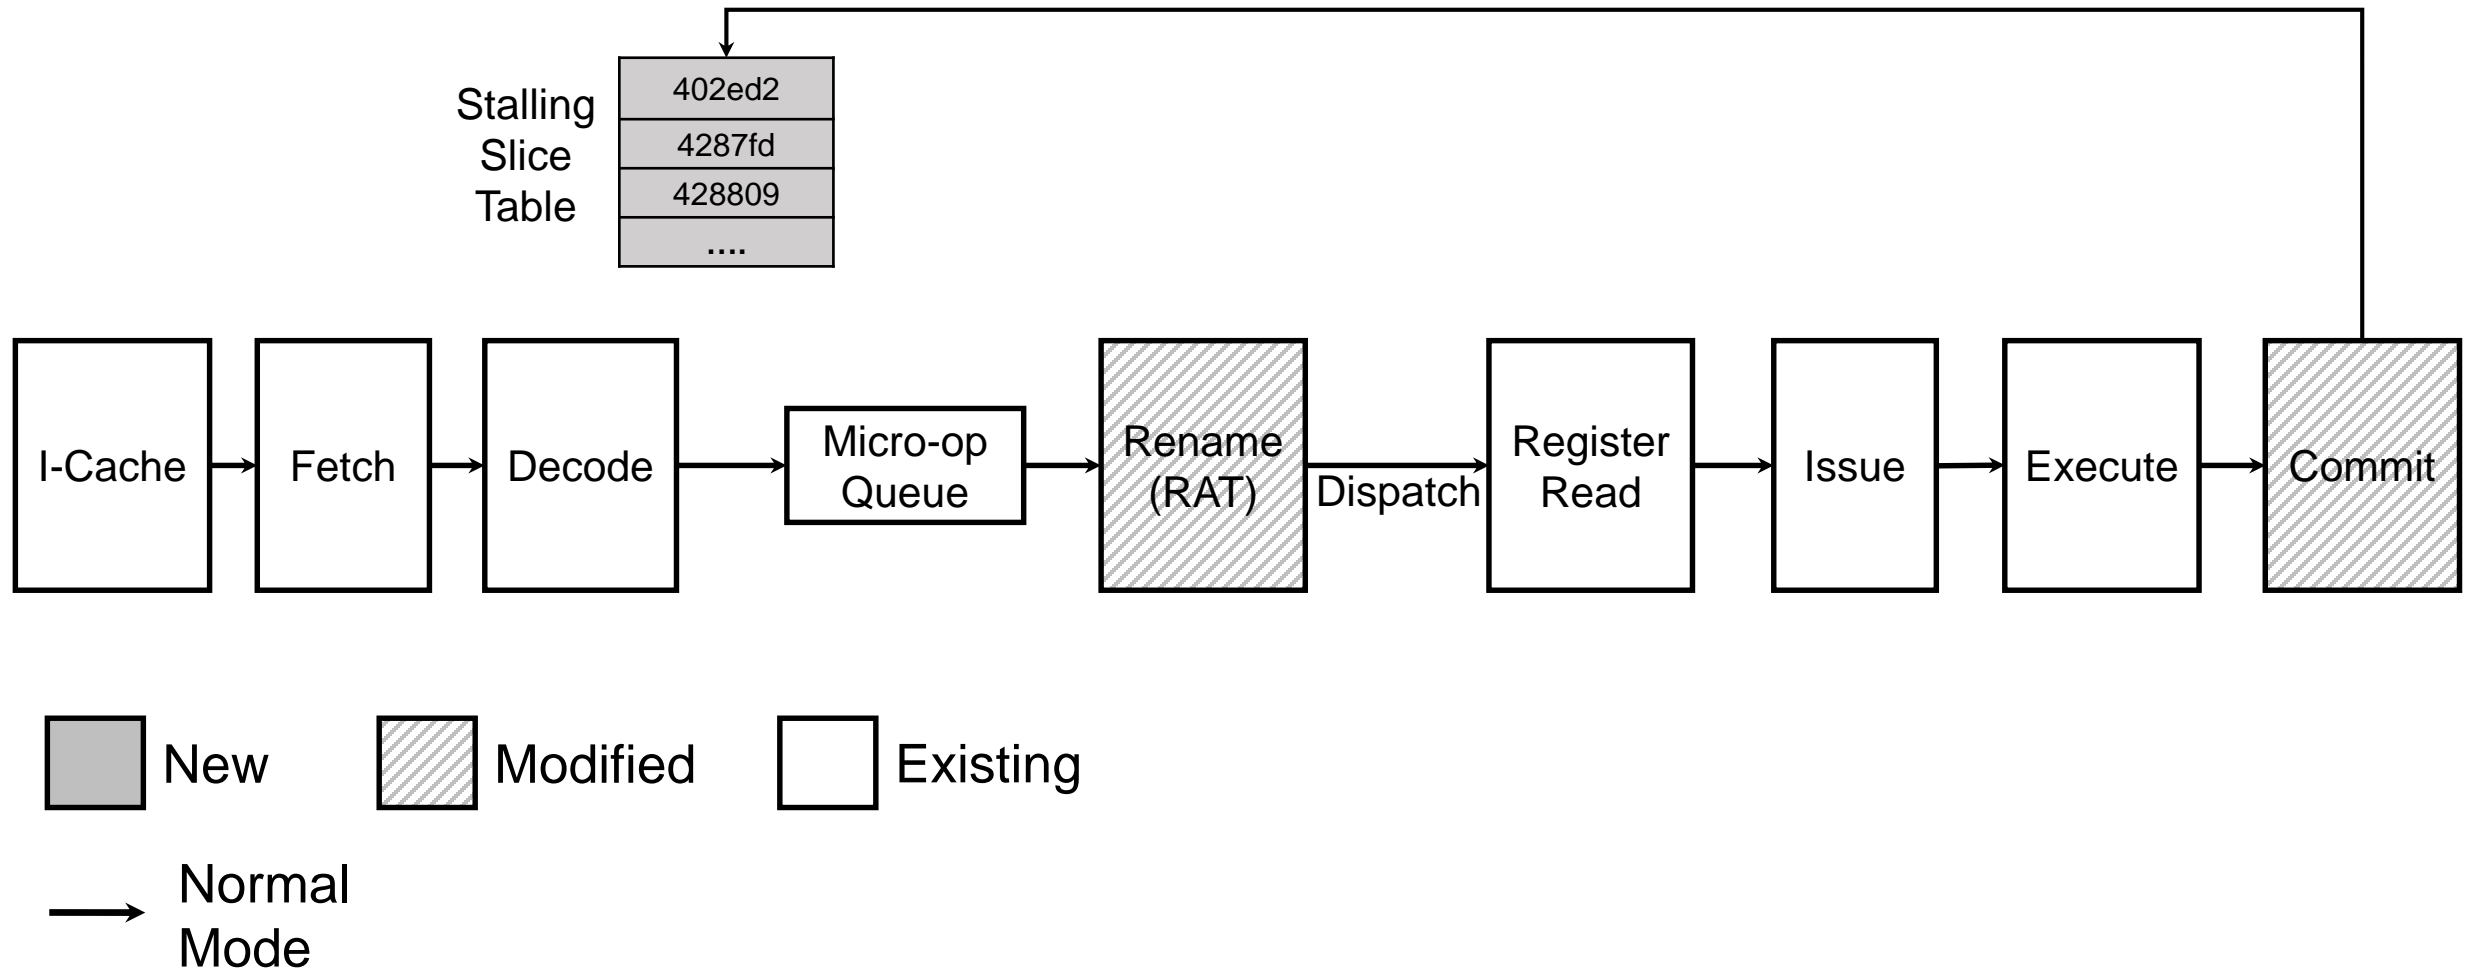
402ed2
4287fd
428809
....



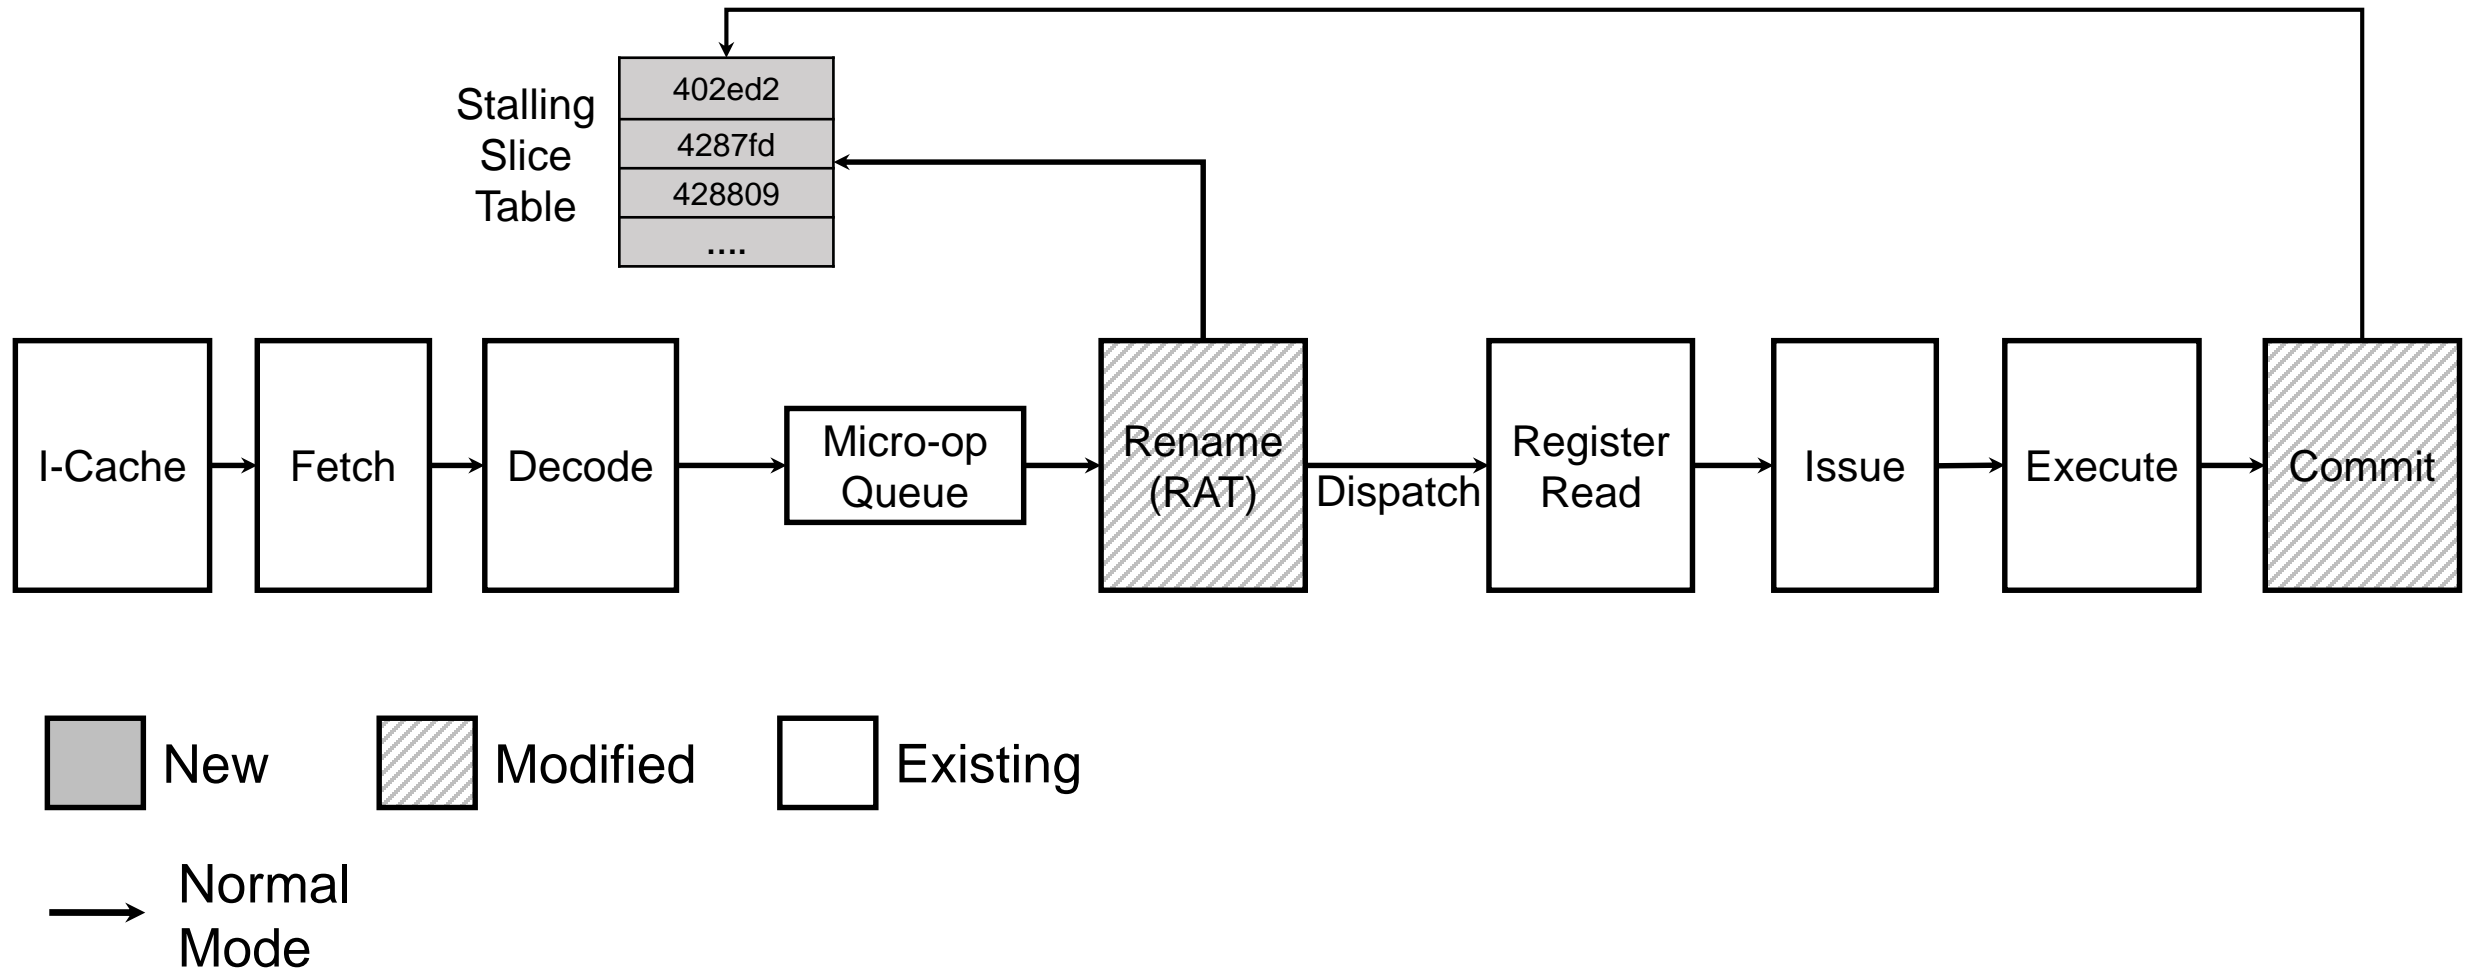
 New  Modified  Existing

→ Normal
Mode

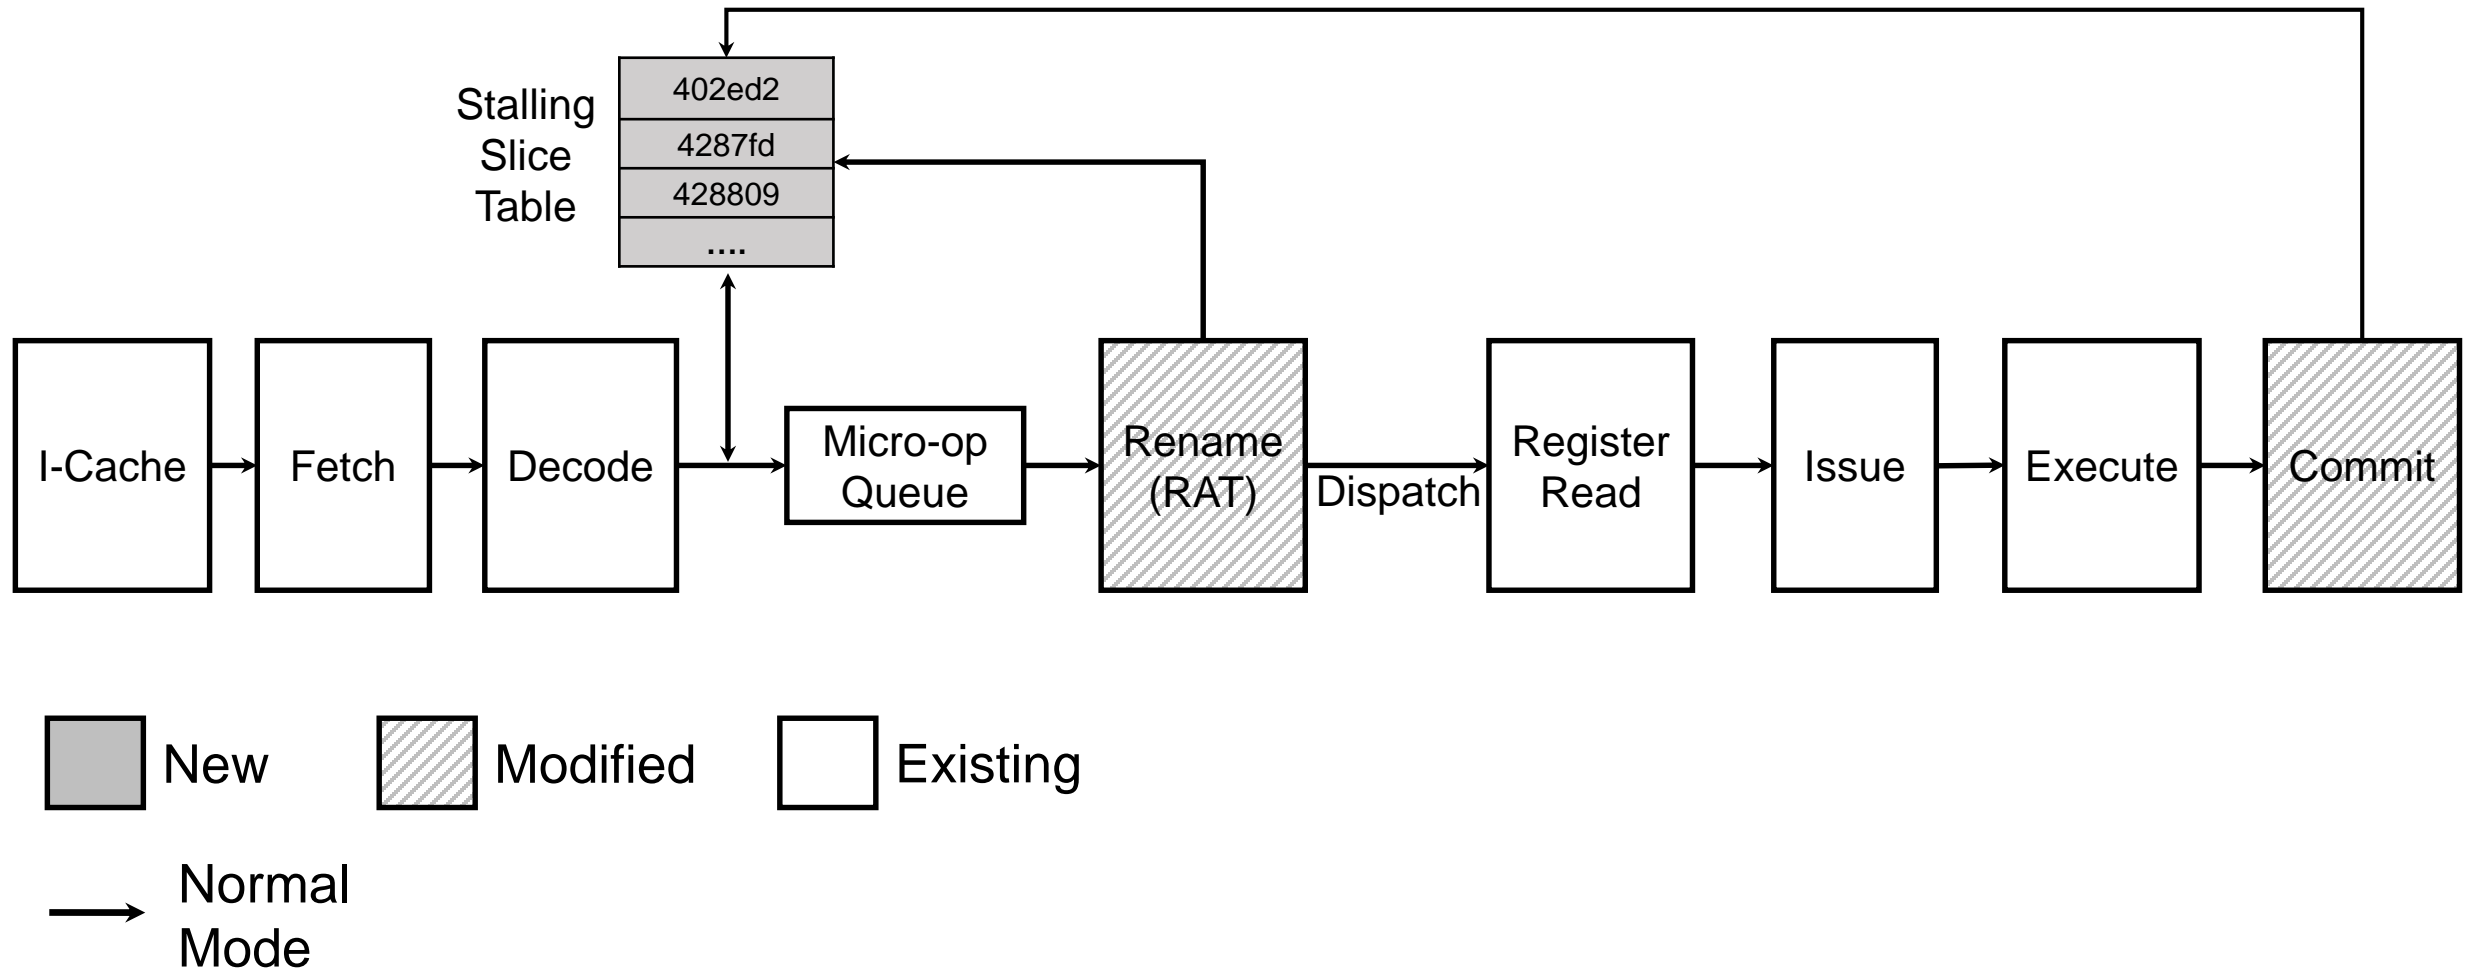
Putting it All Together



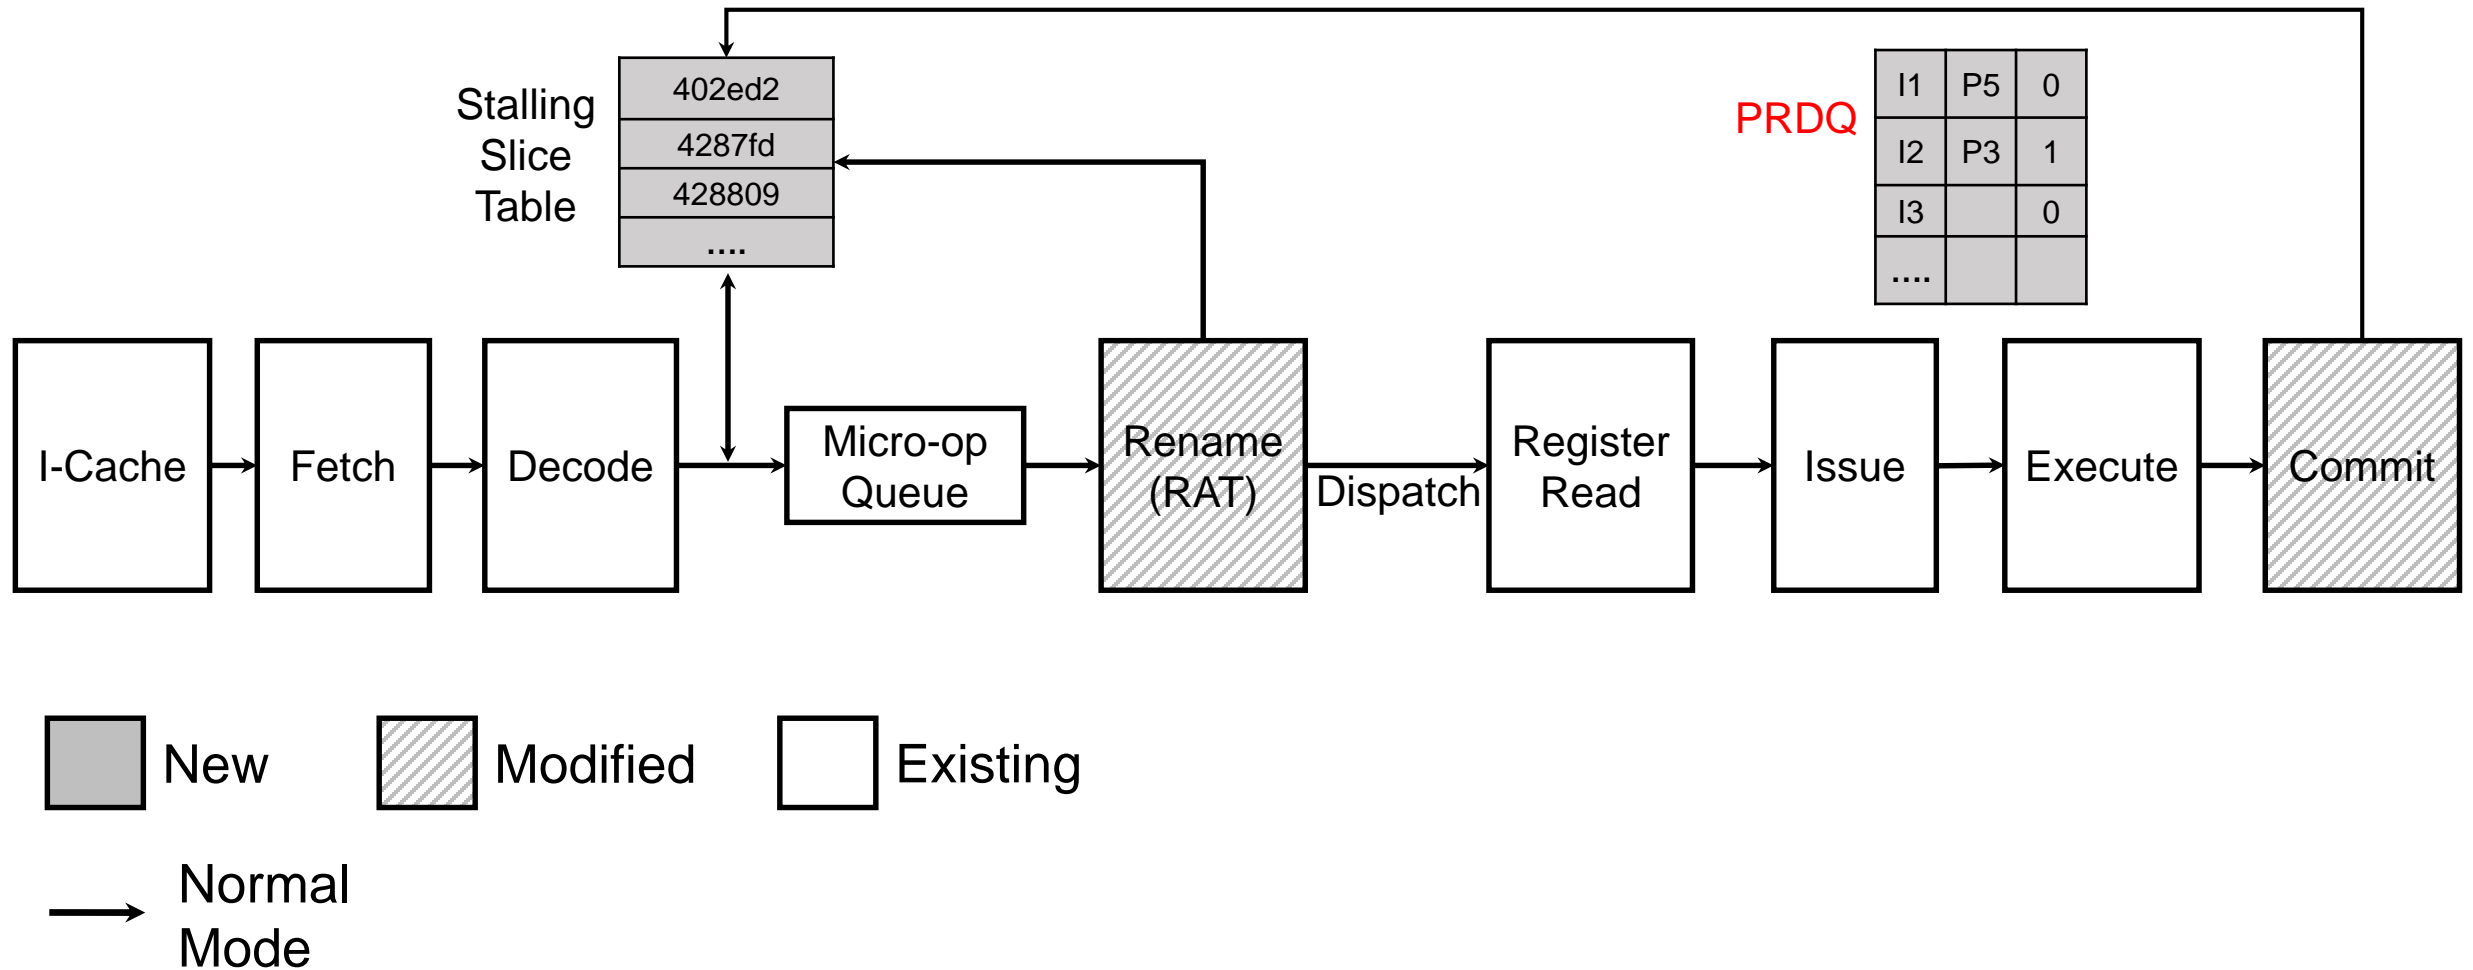
Putting it All Together



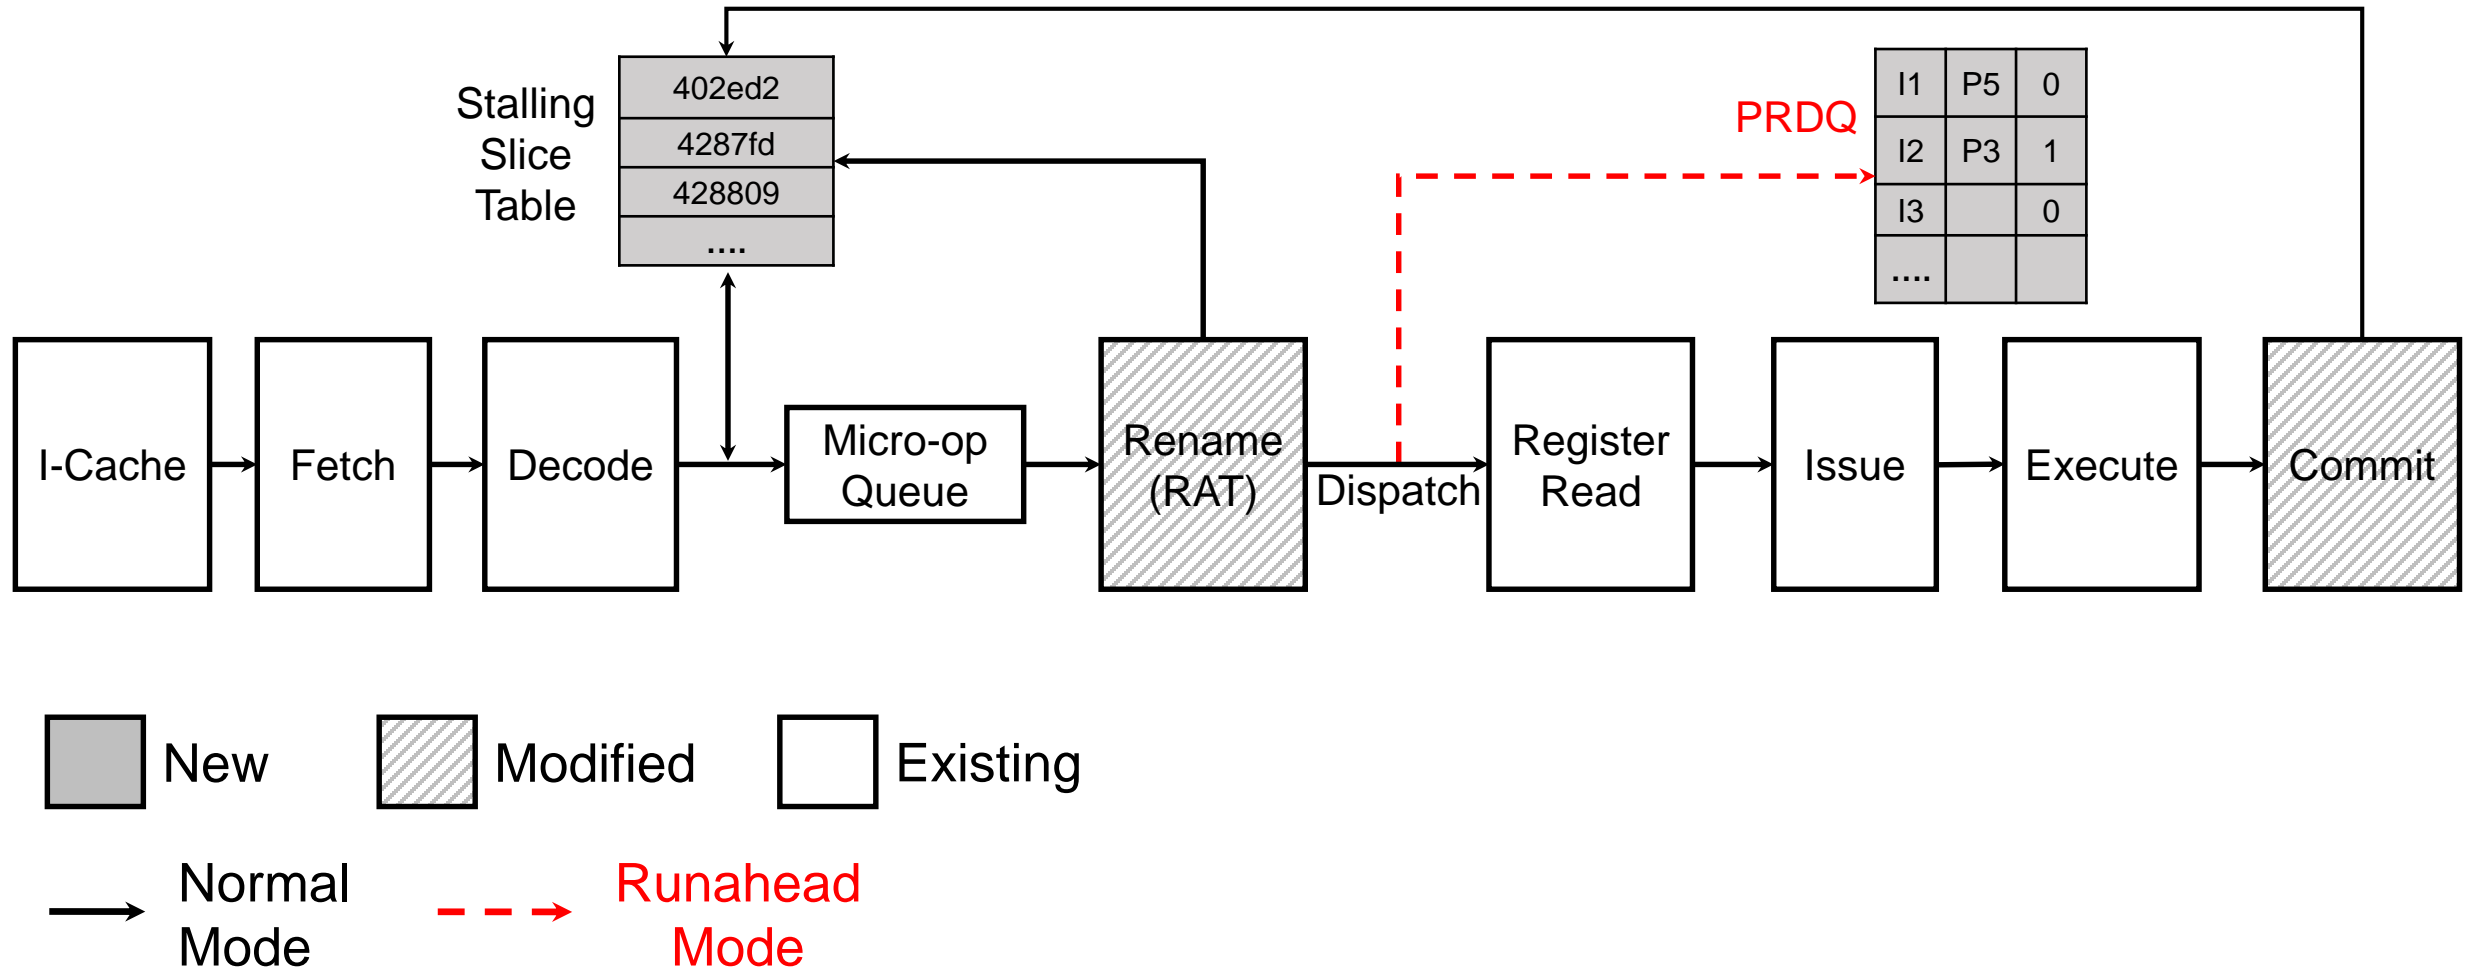
Putting it All Together



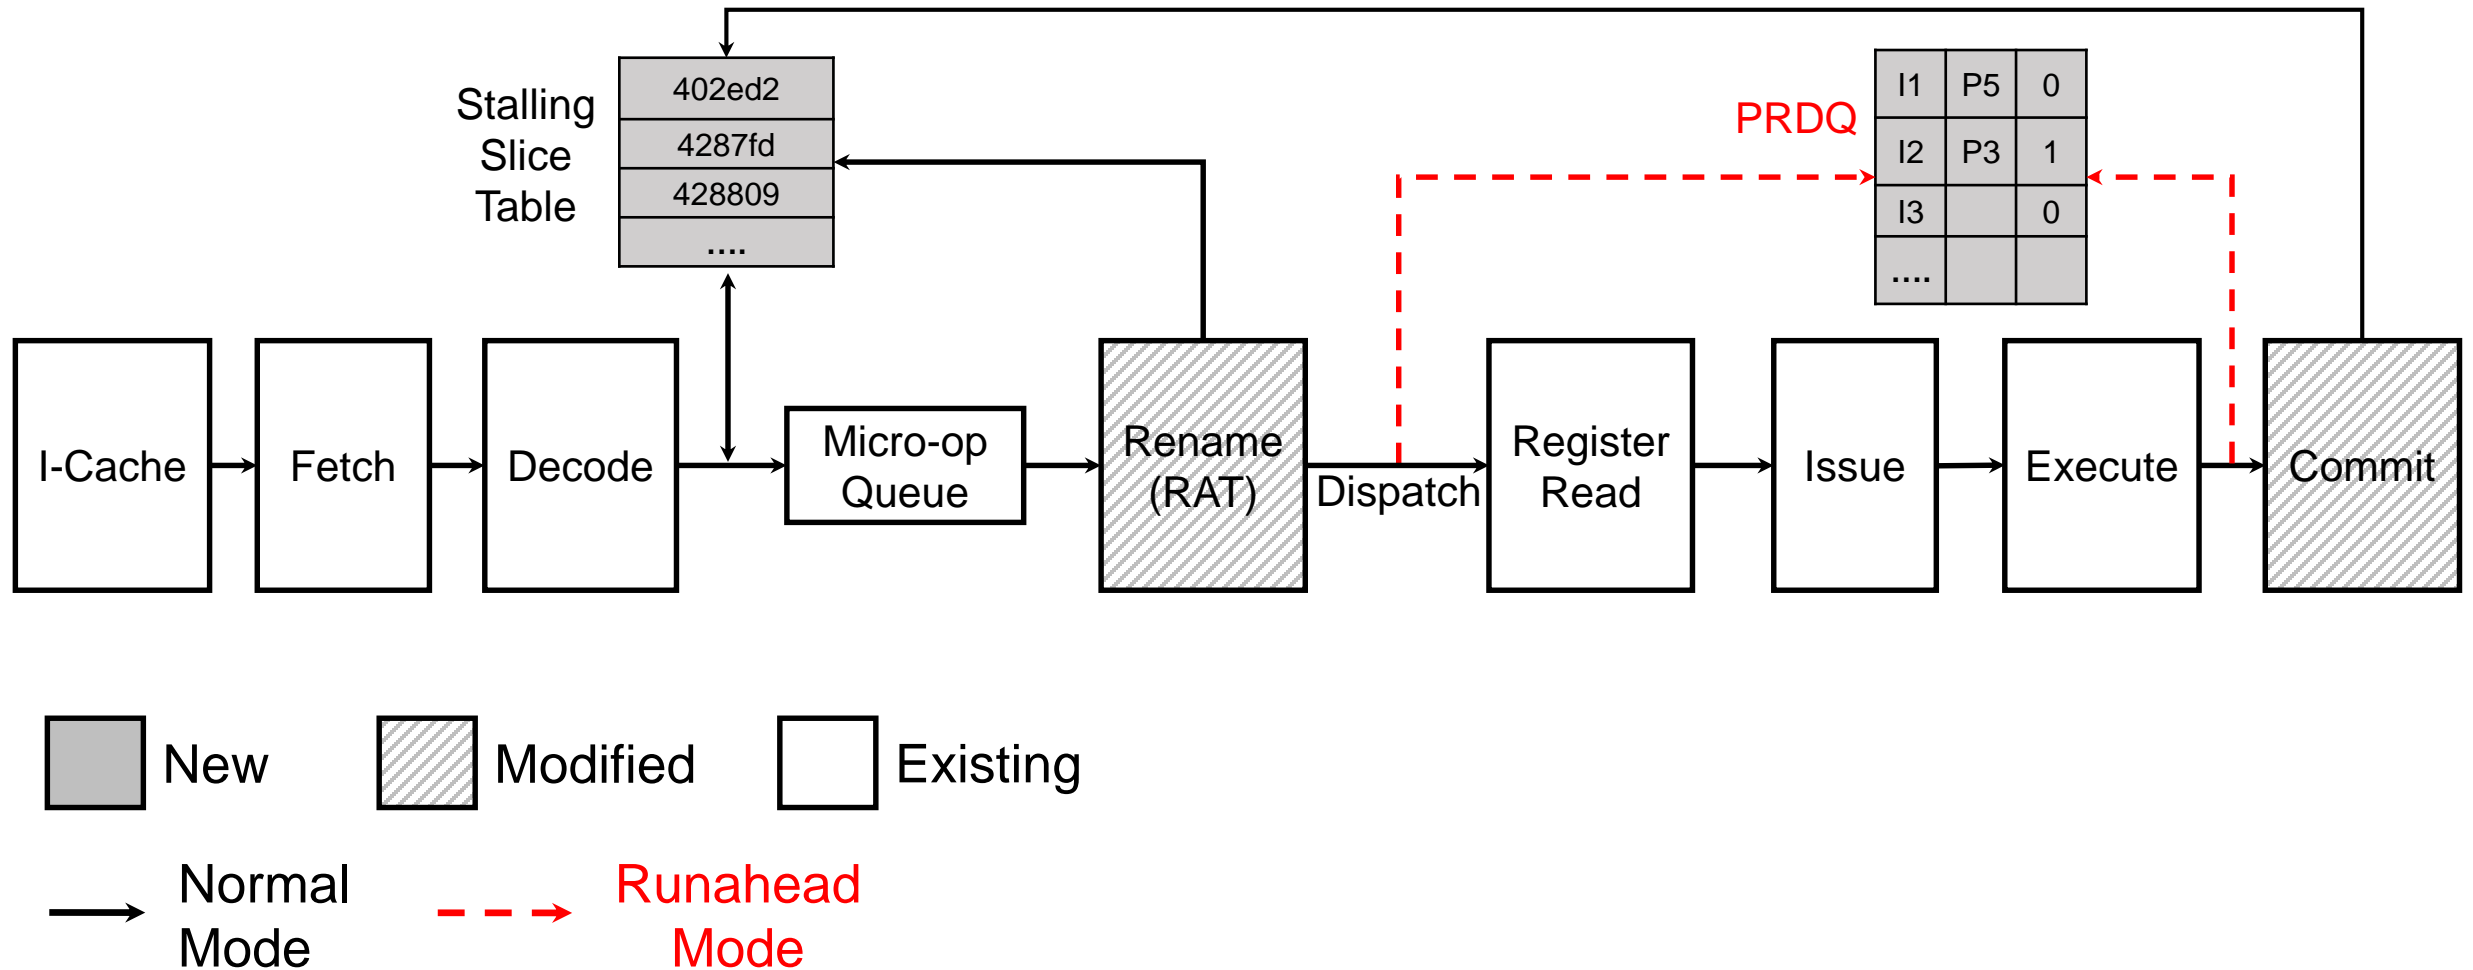
Putting it All Together



Putting it All Together



Putting it All Together



Evaluation

Evaluation

Simulator: Sniper 6.0, McPAT



Evaluation

Simulator: Sniper 6.0, McPAT



Workloads: SPEC CPU2006/CPU2017, 1B SimPoints

Evaluation

Simulator: Sniper 6.0, McPAT



Workloads: SPEC CPU2006/CPU2017, 1B SimPoints

Baseline: ROB=192, issue queue=92, load/store queue=64, register file=168/168

Evaluation

OoO: Baseline out-of-order core

Evaluation

OoO: Baseline out-of-order core

RA: Runahead execution*
-- No short runahead intervals

*[Mutlu et al. ISCA'05]

Evaluation

OoO: Baseline out-of-order core

RA: Runahead execution*

- No short runahead intervals
- No overlapping intervals

*[Mutlu et al. ISCA'05]

Evaluation

OoO: Baseline out-of-order core

RA: Runahead execution*

- No short runahead intervals

- No overlapping intervals

RA-buffer: Runahead buffer**

*[Mutlu et al. ISCA'05]

Evaluation

OoO: Baseline out-of-order core

RA: Runahead execution*

- No short runahead intervals

- No overlapping intervals

RA-buffer: Runahead buffer**

RA-hybrid: Better performing mechanism between RA-buffer and RA

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]

Evaluation

OoO: Baseline out-of-order core

RA: Runahead execution*
-- No short runahead intervals
-- No overlapping intervals

PRE:
Precise runahead
execution***

RA-buffer: Runahead buffer**


RA-hybrid: Better performing mechanism between RA-buffer and RA

*[Mutlu et al. ISCA'05]


**[Hashemi et al. MICRO'15]

***[Naithani et al. HPCA'20]

Evaluation – Performance




Evaluation – Performance




RA: 15.9%

Evaluation – Performance



RA: 15.9% RA-buffer: 13.3%

Evaluation – Performance




RA: 15.9%

RA-buffer: 13.3%

RA-hybrid: 20%

Evaluation – Performance




RA: 15.9%

RA-buffer: 13.3%

RA-hybrid: 20%

PRE: 38.2%

Evaluation – Performance




RA: 15.9%

RA-buffer: 13.3%

RA-hybrid: 20%

PRE: 38.2%

Evaluation – Performance




RA: 15.9%

RA-buffer: 13.3%

RA-hybrid: 20%

PRE: 38.2%

Evaluation – Performance




RA: 15.9%

RA-buffer: 13.3%

RA-hybrid: 20%

PRE: 38.2%

Evaluation – Performance




RA: 15.9%

RA-buffer: 13.3%

RA-hybrid: 20%

PRE: 38.2%

Evaluation – Performance




RA: 15.9%

RA-buffer: 13.3%

RA-hybrid: 20%

PRE: 38.2%

Evaluation – Performance




RA: 15.9%

RA-buffer: 13.3%


RA-hybrid: 20%

PRE: 38.2%

Evaluation – Memory-Level Parallelism




Evaluation – Memory-Level Parallelism



RA: 1.5X


Evaluation – Memory-Level Parallelism



RA: 1.5X

RA-buffer: 1.3X

Evaluation – Memory-Level Parallelism




RA: 1.5X

RA-buffer: 1.3X

RA-hybrid: 1.6X

Evaluation – Memory-Level Parallelism




RA: 1.5X

RA-buffer: 1.3X


RA-hybrid: 1.6X

PRE: 2.0X

Evaluation – LLC Miss Count Reduction




Evaluation – LLC Miss Count Reduction




RA: 26.4%

Evaluation – LLC Miss Count Reduction




RA: 26.4% RA-buffer: 27.7%

Evaluation – LLC Miss Count Reduction




RA: 26.4% RA-buffer: 27.7% RA-hybrid: 31%

Evaluation – LLC Miss Count Reduction




RA: 26.4% RA-buffer: 27.7% RA-hybrid: 31% PRE: 50.2%

Evaluation – Energy




Evaluation – Energy



RA: +2.4%


Evaluation – Energy



RA: +2.4%

RA-buffer: Same

Evaluation – Energy




RA: +2.4%

RA-buffer: Same

RA-hybrid: Same

Evaluation – Energy



RA: +2.4%

RA-buffer: Same

RA-hybrid: Same

PRE: -6.2%

Conclusions

Conclusions

1. Never flushes the ROB

Conclusions

1. Never flushes the ROB
2. Executes only useful instructions in runahead mode

Conclusions

1. Never flushes the ROB
2. Executes only useful instructions in runahead mode
3. Efficiently manages microarchitectural resources

Conclusions

1. Never flushes the ROB
2. Executes only useful instructions in runahead mode
3. Efficiently manages microarchitectural resources

18.2% better performance

Conclusions

1. Never flushes the ROB
2. Executes only useful instructions in runahead mode
3. Efficiently manages microarchitectural resources

18.2% better performance

6.2% better energy

Precise Runahead Execution



Ajeya Naithani
Josue Feliu
Almutaz Adileh
Lieven Eeckhout