

Improving Soft Error Reliability in Modern Processors

Ajeya Naithani

Advisor: Prof. Lieven Eeckhout

Doctoral Thesis Defense

Reliability in Computer Science

The term reliability refers to the ability of a computer **hardware** or **software** component to **consistently perform** according to its specifications.

What if the Reliability is Compromised?

What if the Reliability is Compromised?

A problem has been detected and Windows has been shut down to prevent damage to your computer.

The problem seems to be caused by the following file: kbdhid.sys

MANUALLY_INITIATED_CRASH

If this is the first time you've seen this stop error screen, restart your computer. If this screen appears again, follow these steps:

Check to make sure any new hardware or software is properly installed. If this is a new installation, ask your hardware or software manufacturer for any Windows updates you might need.

If problems continue, disable or remove any newly installed hardware or software. Disable BIOS memory options such as caching or shadowing. If you need to use safe mode to remove or disable components, restart your computer, press F8 to select Advanced Startup Options, and then select Safe Mode.

Technical Information:

*** STOP: 0x000000e2 (0x00000000, 0x00000000, 0x00000000, 0x00000000)

*** kbdhid.sys - Address 0x94efd1aa base at 0x94efb000 DateStamp 0x4a5bc705

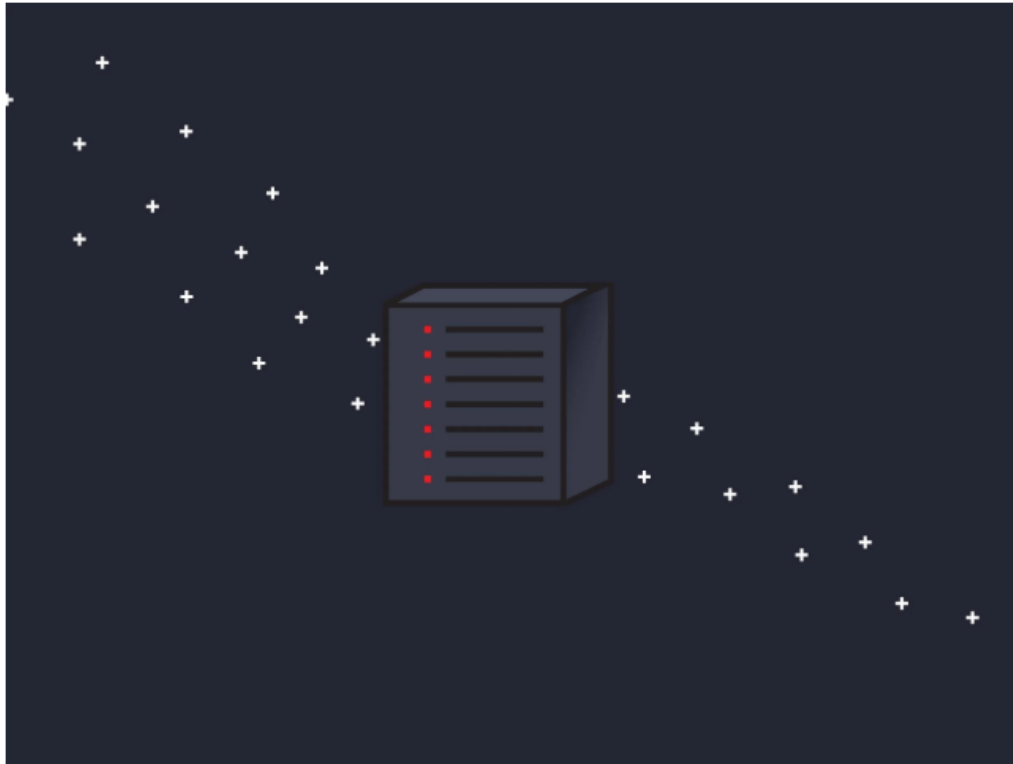
What if the Reliability is Compromised?

SARAH SCOTTS

SCIENCE 06.04.2018 07:00 AM

Cosmic Ray Showers Crash Supercomputers. Here's What to Do About It

What happens when a national laboratory's supercomputers start glitching?



EMILY WAITE

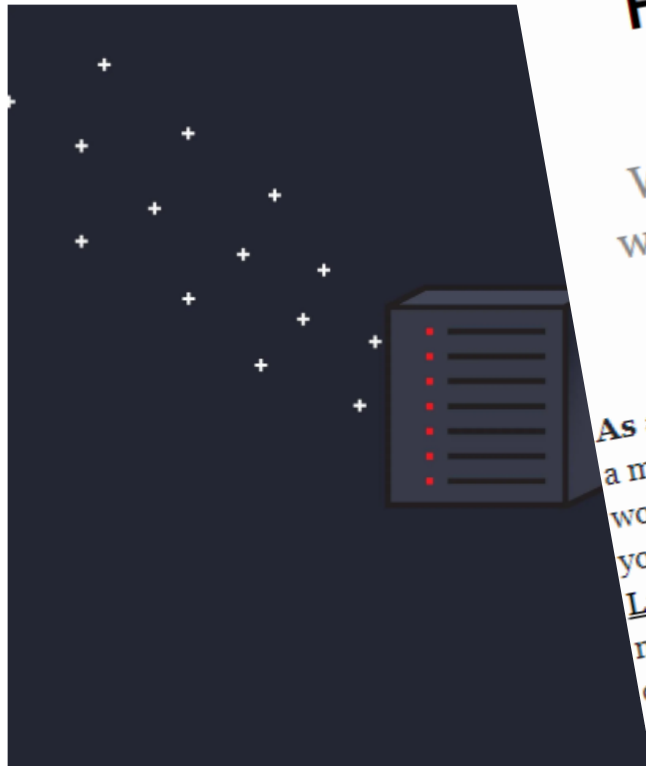
What if the Reliability is Compromised?

SARAH SCOLES

SCIENCE 06.04.2018 07:00 AM

Cosmic Ray Showers Crash Supercomputers. Here's About It

What happens when a national laboratory's superc



EMILY WAITE

How To Kill A Supercomputer: Dirty Power, Cosmic Rays, and Bad Solder

Will future exascale supercomputers be able to withstand the steady onslaught of routine faults?

By Al Geist

As a child, were you ever afraid that a monster lurking in your bedroom would leap out of the dark and get you? My job at Oak Ridge National Laboratory is to worry about a similar monster, hiding in the steel cabinets of the supercomputers and threatening to crash the largest computing machines on the planet.

The monster is something
supercomputer specialists call



What if the Reliability is Compromised?

The Uber Crash Won't Be the Last Shocking Self-Driving Death

These systems will fail, and when they do, the result will look nothing like human-powered crashes.



KIEN HONG LE/BLOOMBERG/GETTY IMAGES

Did you Read your Car Insurance Carefully?

Assicurato ha diritto all'indennizzo
invalidità permanente a
zione che la stessa si manifesti
due anni dall'Infortunio.

valutazione dell'invalidità
nente sarà effettuata in base alla
a che segue nella pagina
ssiva. Se la lesione comporta
minorazione anziché la perdita
anatomica o funzionale di
i o arti, le percentuali della
la vengono ridotte in
orzione alla funzionalità
ata.

perdita totale anatomica o
onale di più organi odarti
orta l'applicazione di una
ntuale di invalidità pari alla
a delle singole percentuali
osciute per ciascuna lesione con
ssimo del 100%. Per i casi non

ART. C.3 - ESCLUSIONI

Sono esclusi dall'assicurazione i sinistri determinati da:

- a) partecipazione a corse o gare e relative prove ufficiali e verifiche preliminari e finali previste nel regolamento particolare di gara;
- b) tumulti popolari, atti di terrorismo, vandalismo, attentati ai quali l'Assicurato abbia partecipato attivamente;
- c) guerra, insurrezioni, terremoti, eruzioni vulcaniche;
- d) trasmutazione del nucleo dell'atomo come pure dovuti ad esposizione a radiazioni ionizzanti;

ART. C.4 - LIQUIDAZIONE

Liquidazione incaricato dalla
Società
giustifi
oppur
interv
Sanita
quota
prede
anticip
sanita
presen
incaric
danno

The insurance does not cover those accidents caused by:

[...]

exposure to ionizing radiation*

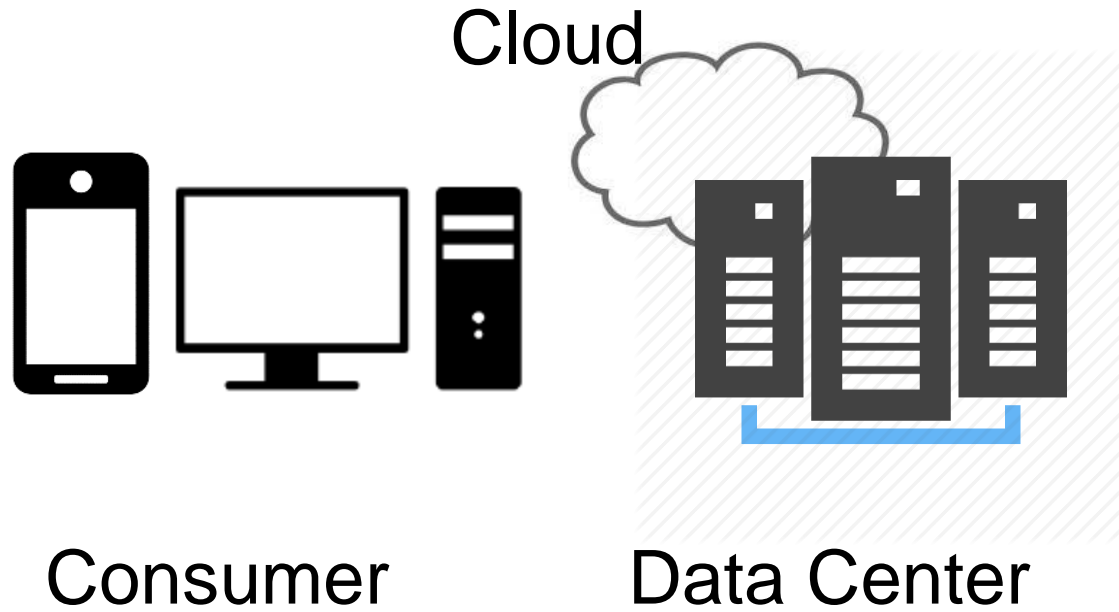
Different Types of Computing Systems

Different Types of Computing Systems

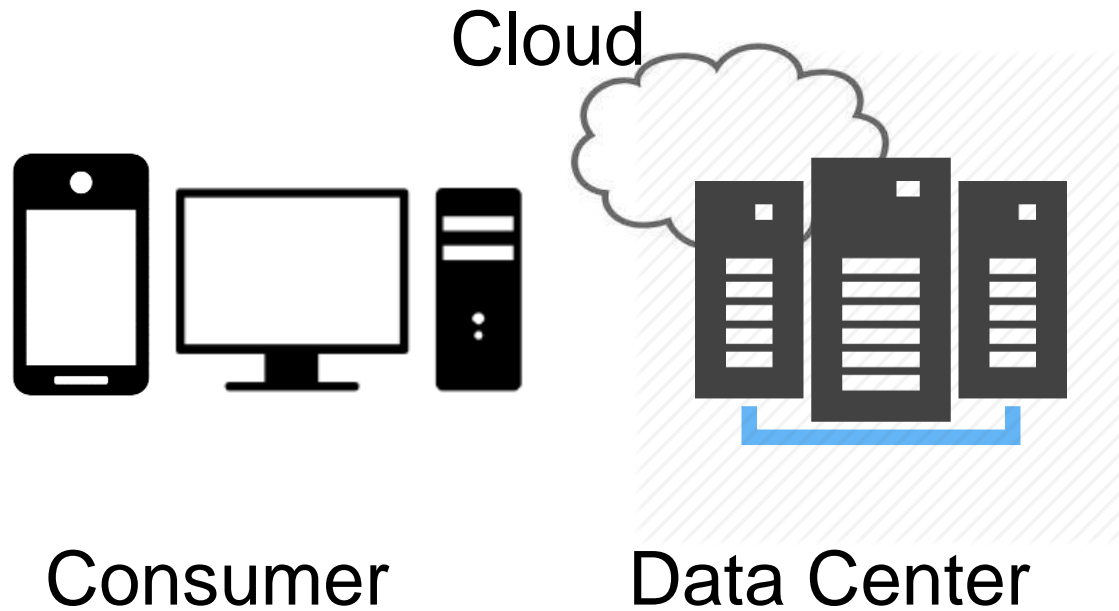


Consumer

Different Types of Computing Systems

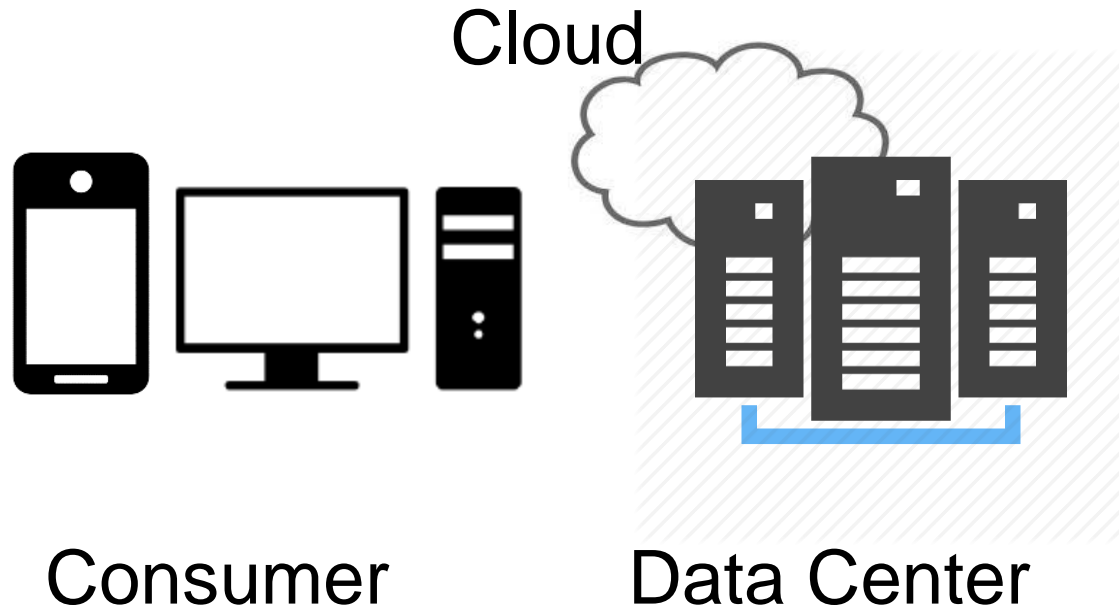


Different Types of Computing Systems



Scientific Computing

Different Types of Computing Systems

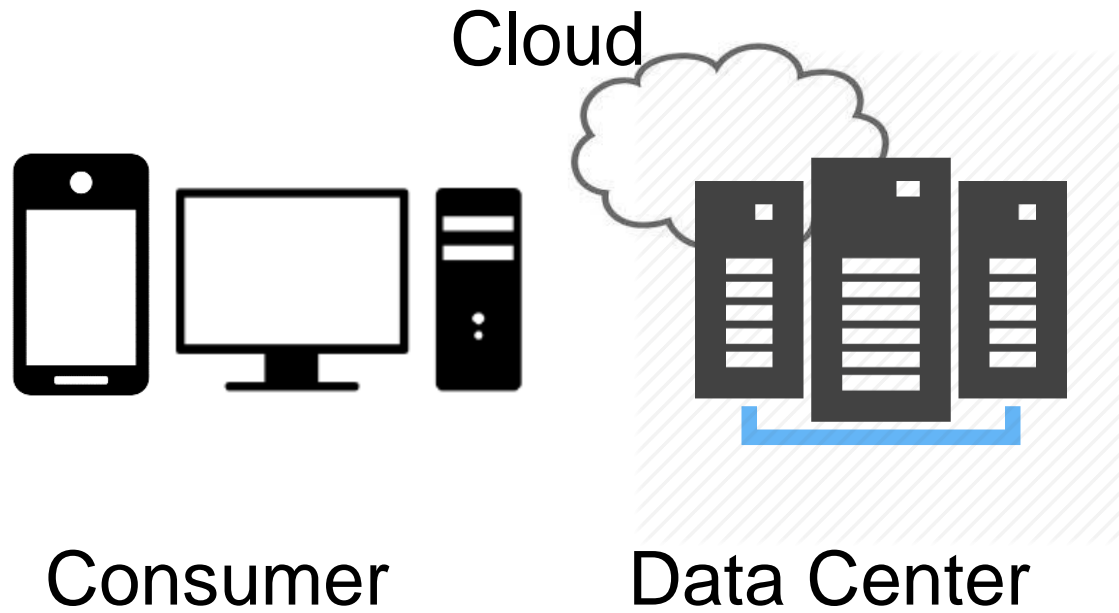


Scientific Computing



Automotive

Different Reliability Requirements

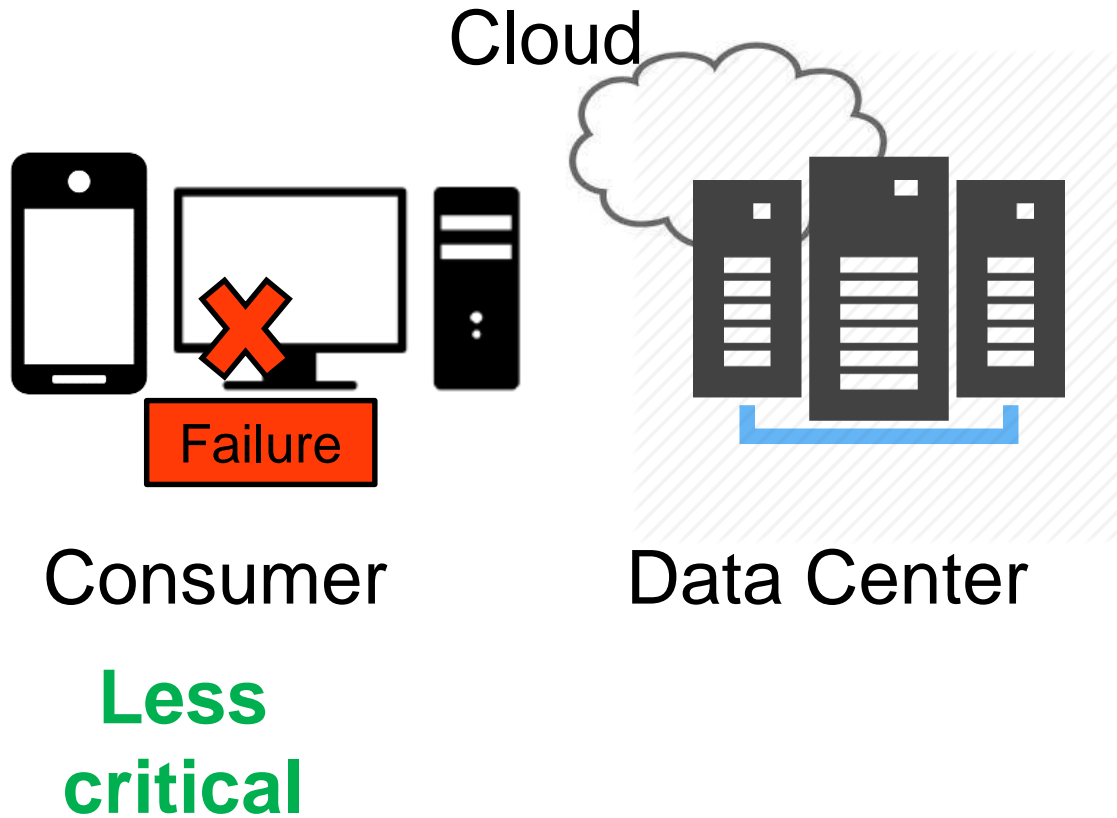


Scientific Computing



Automotive

Different Reliability Requirements

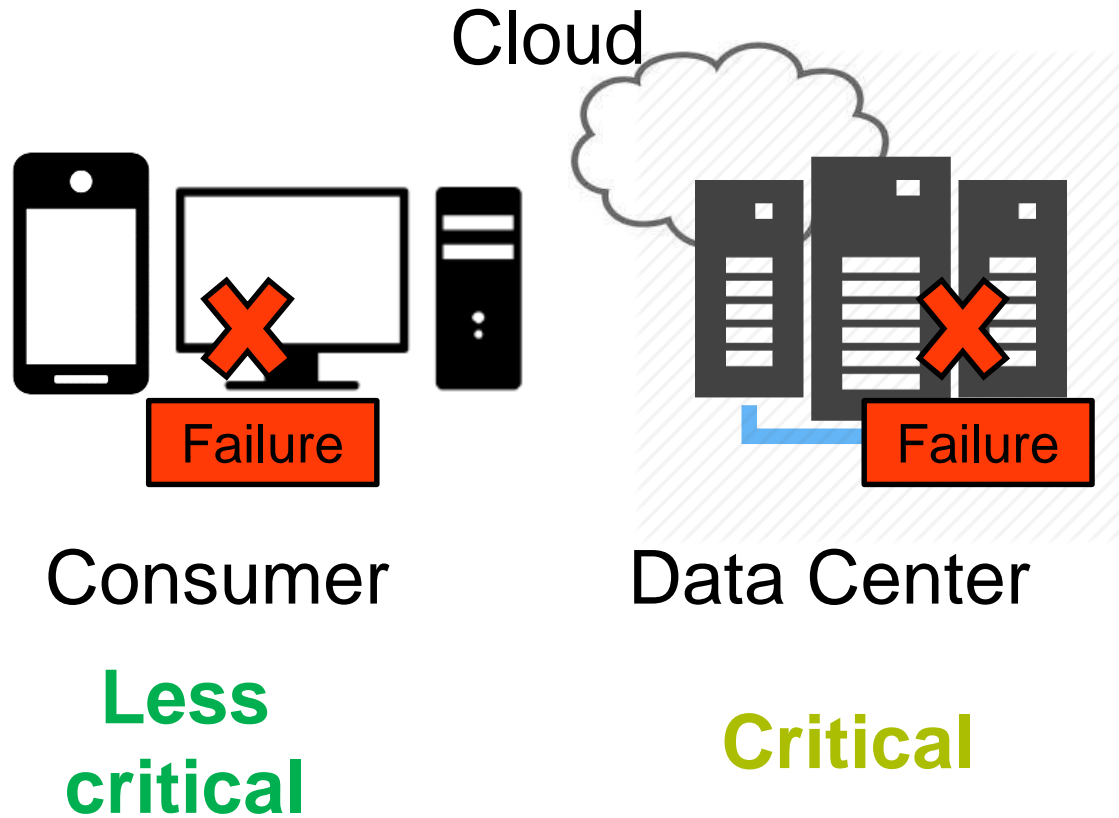


Scientific Computing



Automotive

Different Reliability Requirements

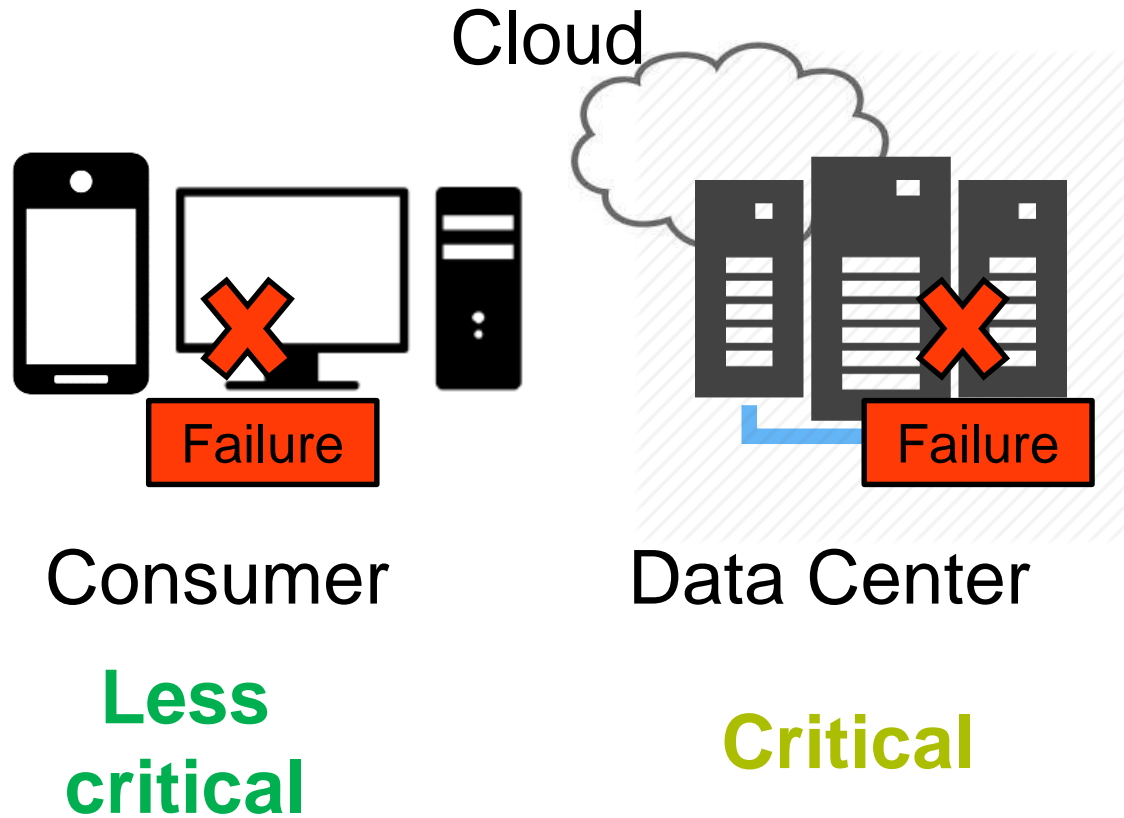


Scientific Computing



Automotive

Different Reliability Requirements



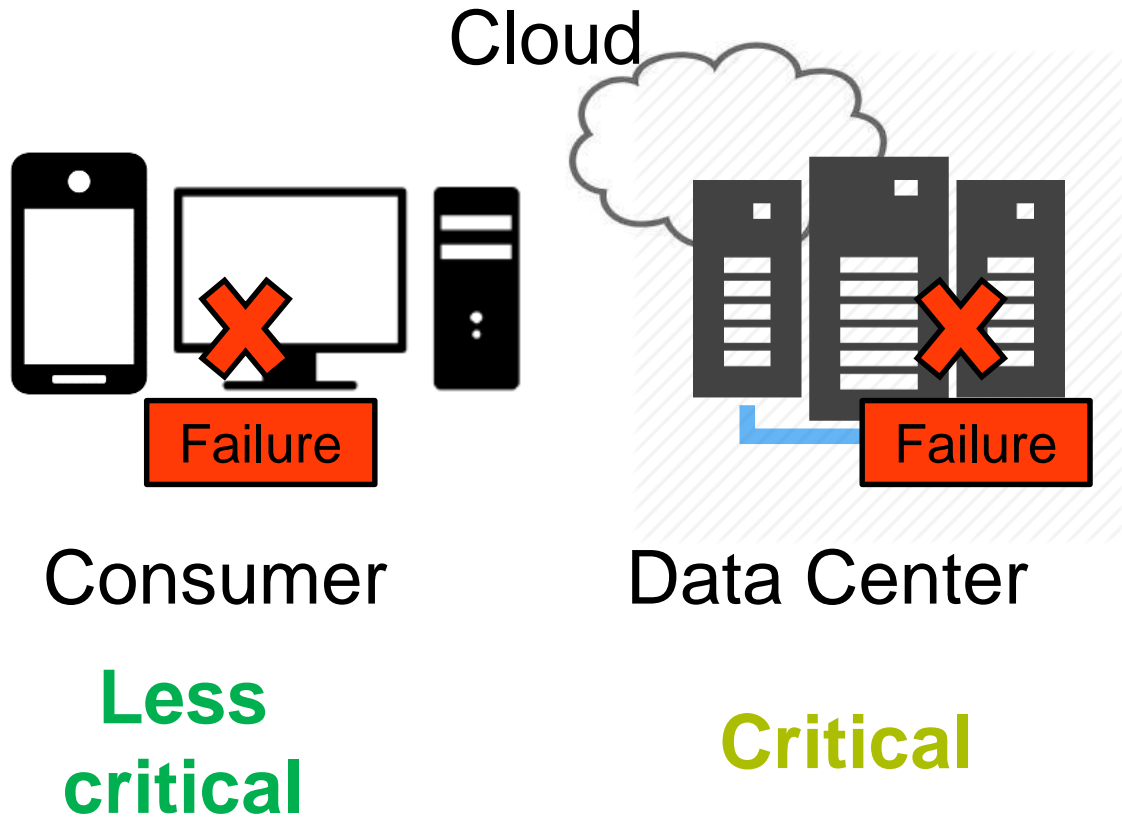
Scientific Computing

More
Critical



Automotive

Different Reliability Requirements



Scientific Computing

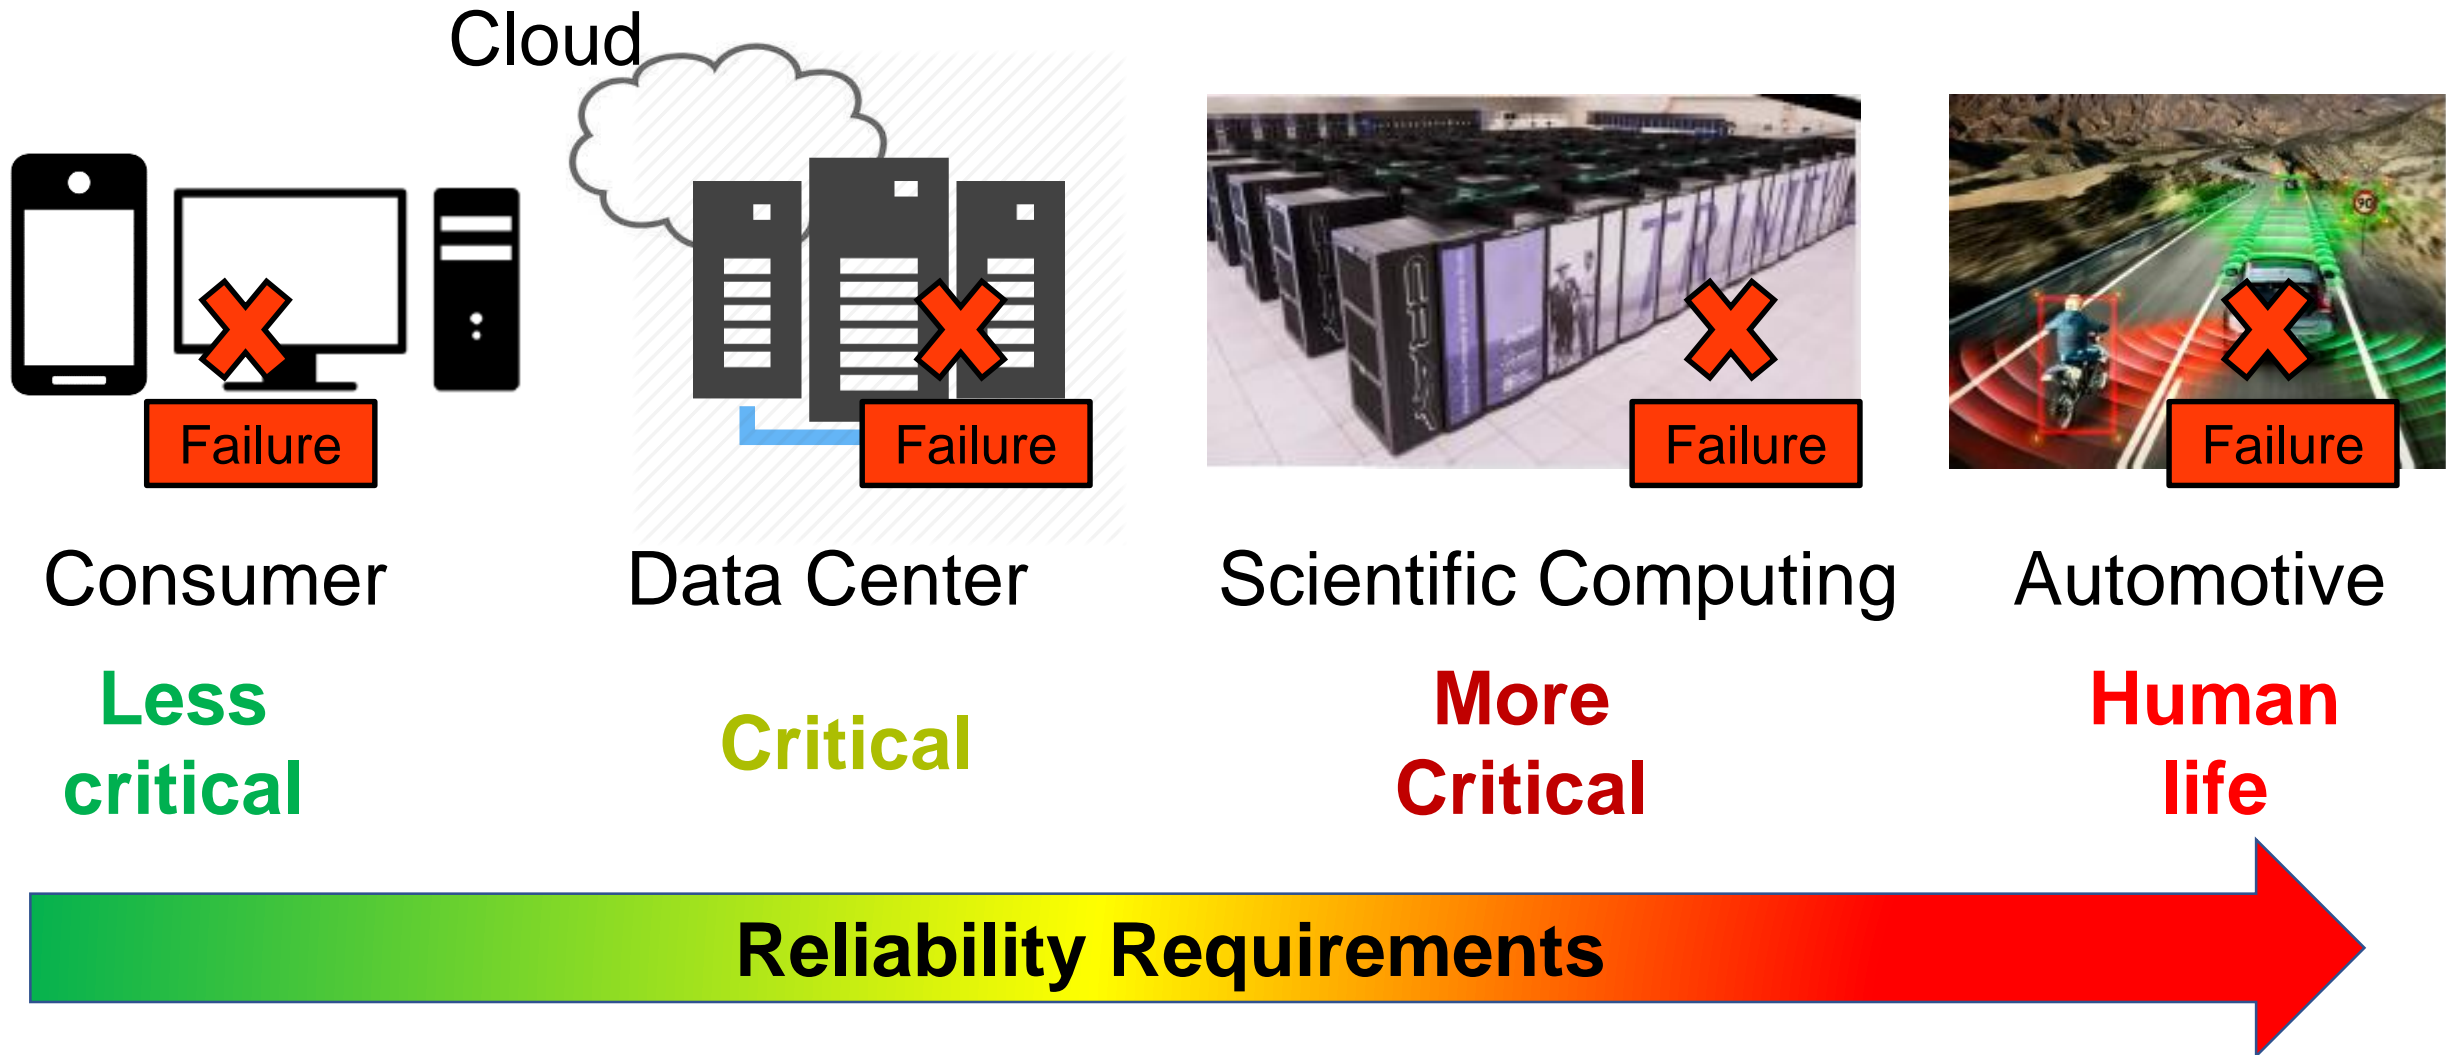
**More
Critical**



Automotive

**Human
life**

Different Reliability Requirements



Different Devices Running These Systems

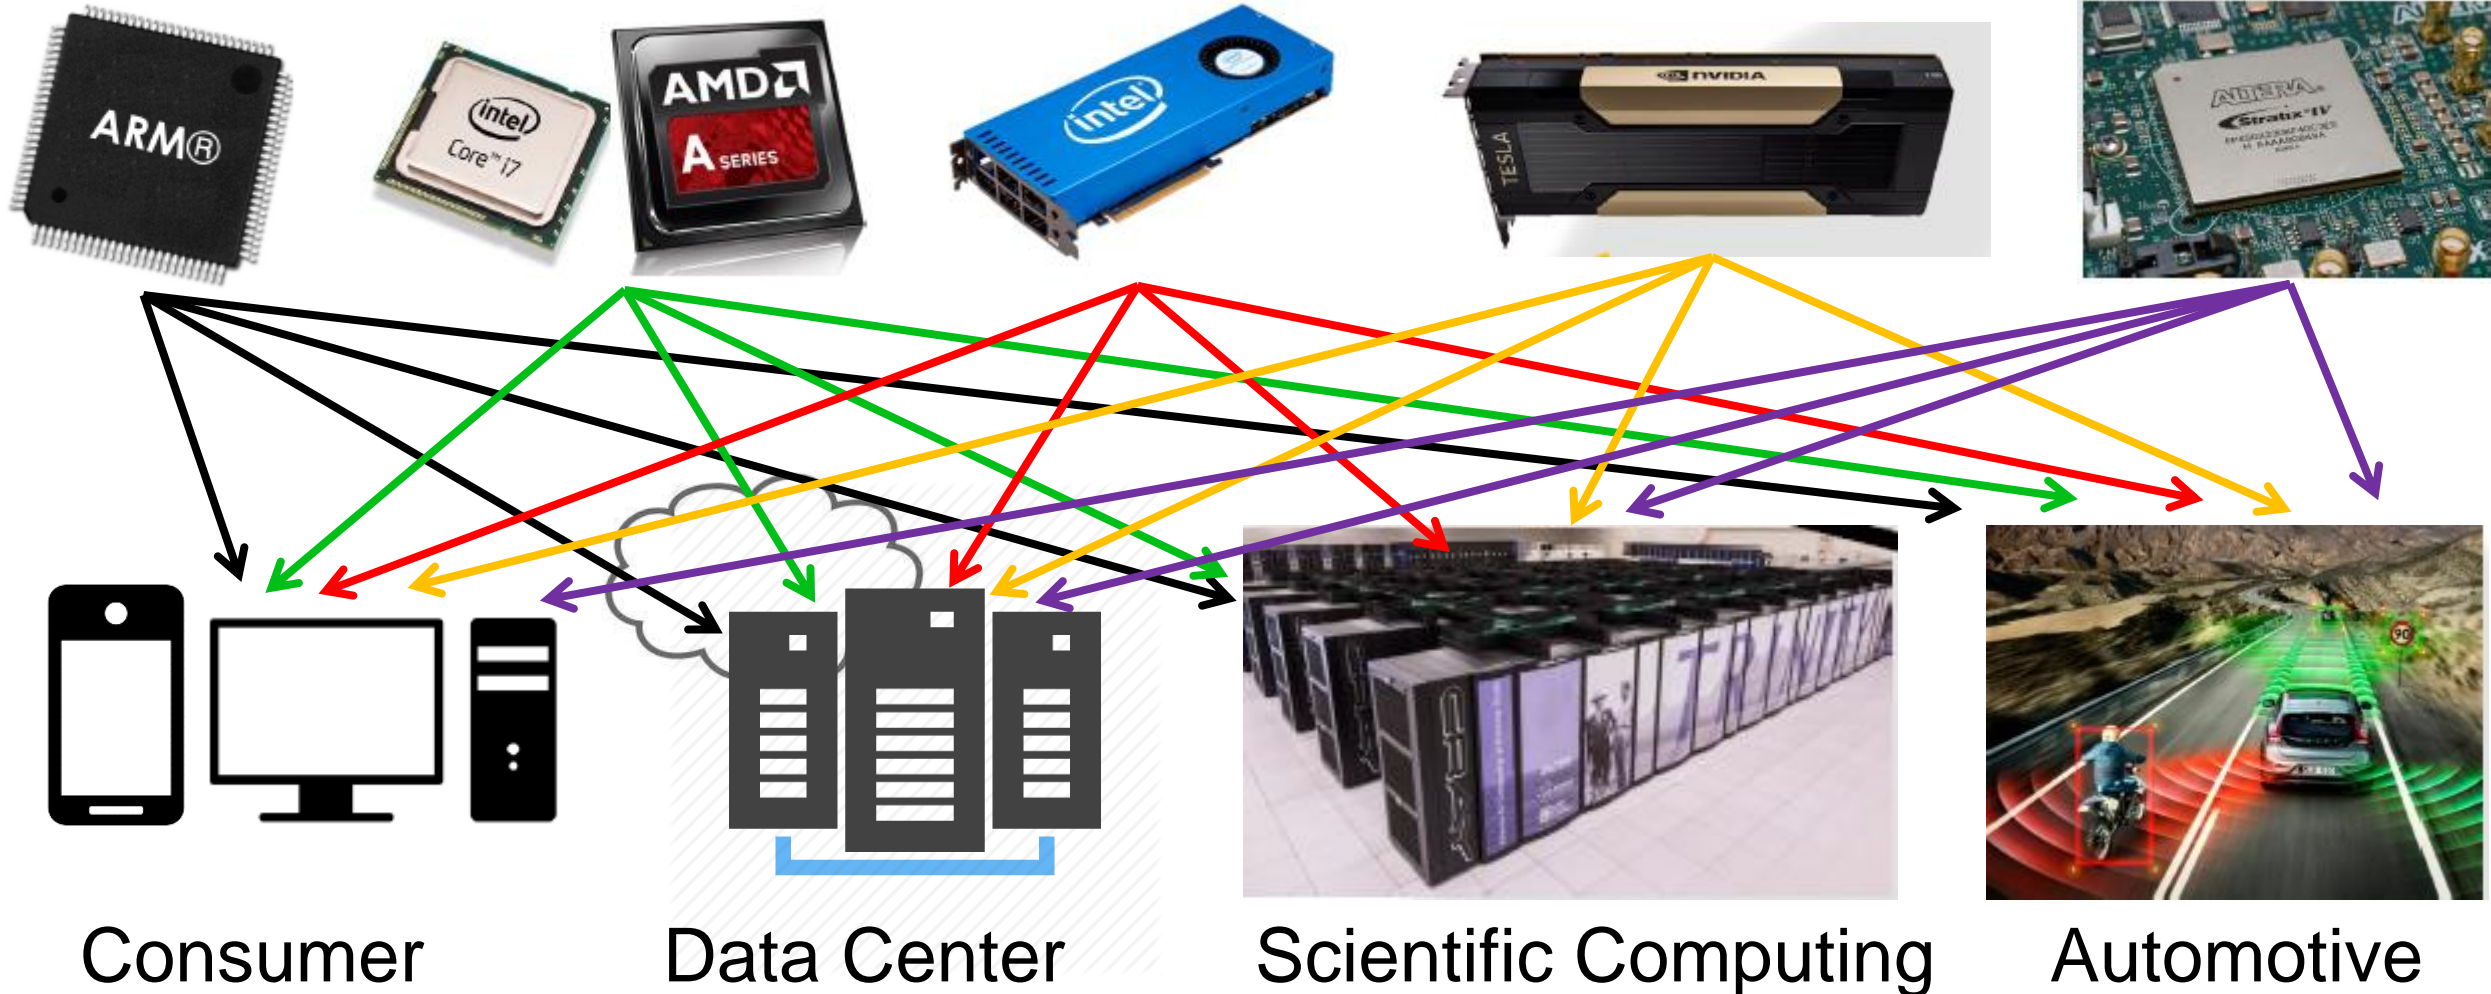
ARM

CPU

Coprocessor

GPU

FPGA/SoC



Two Key Components: Processor and Memory

Two Key Components: Processor and Memory

Processor

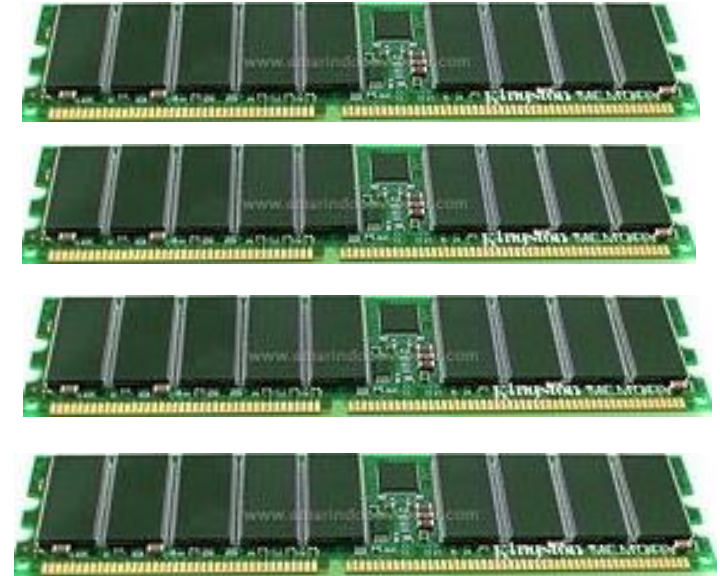


Two Key Components: Processor and Memory

Processor



Memory



Two Key Components: Processor and Memory

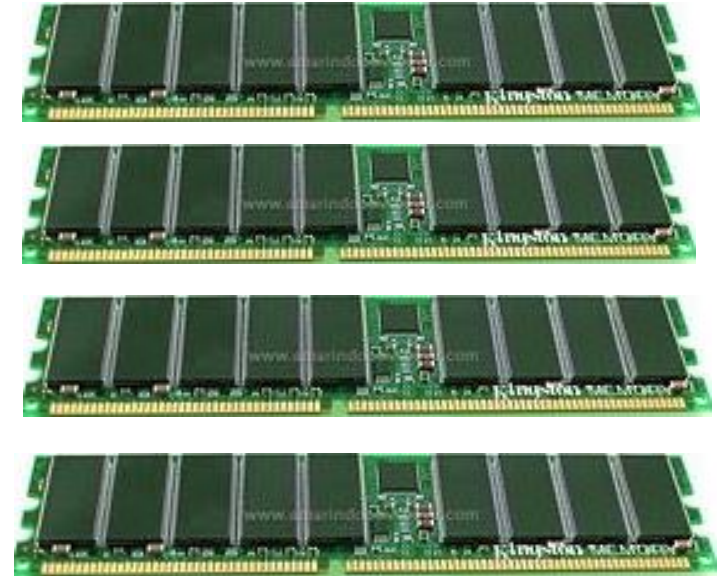
Processor



Read



Memory



Two Key Components: Processor and Memory

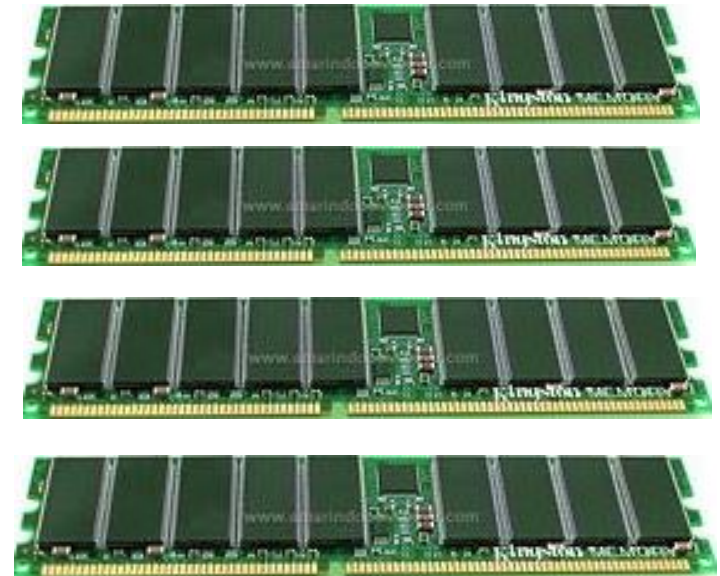
Processor

Memory

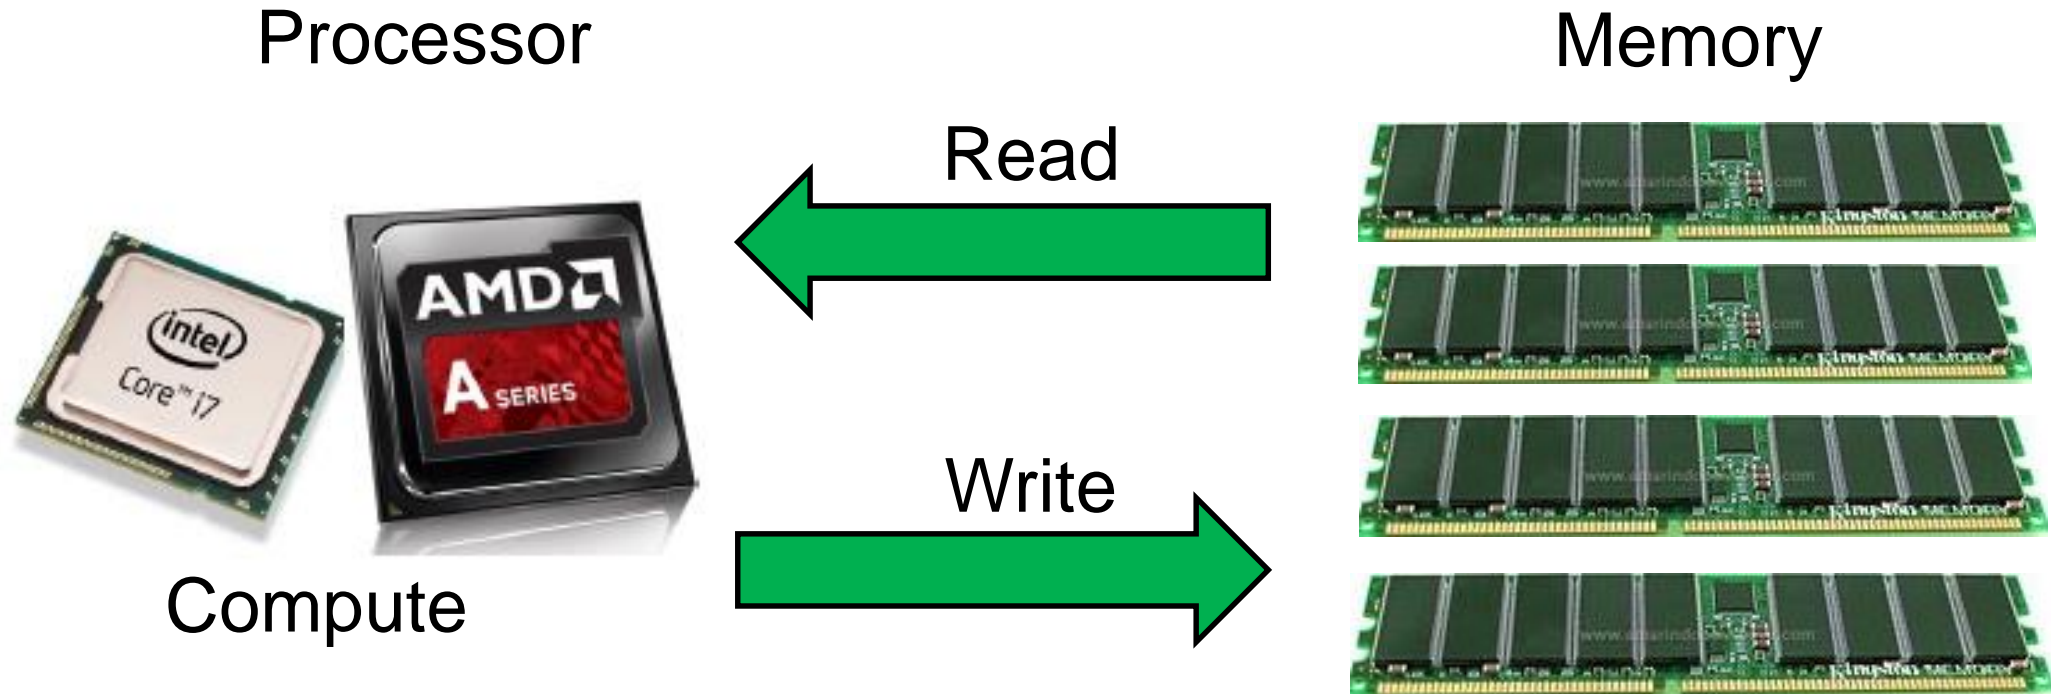
Read



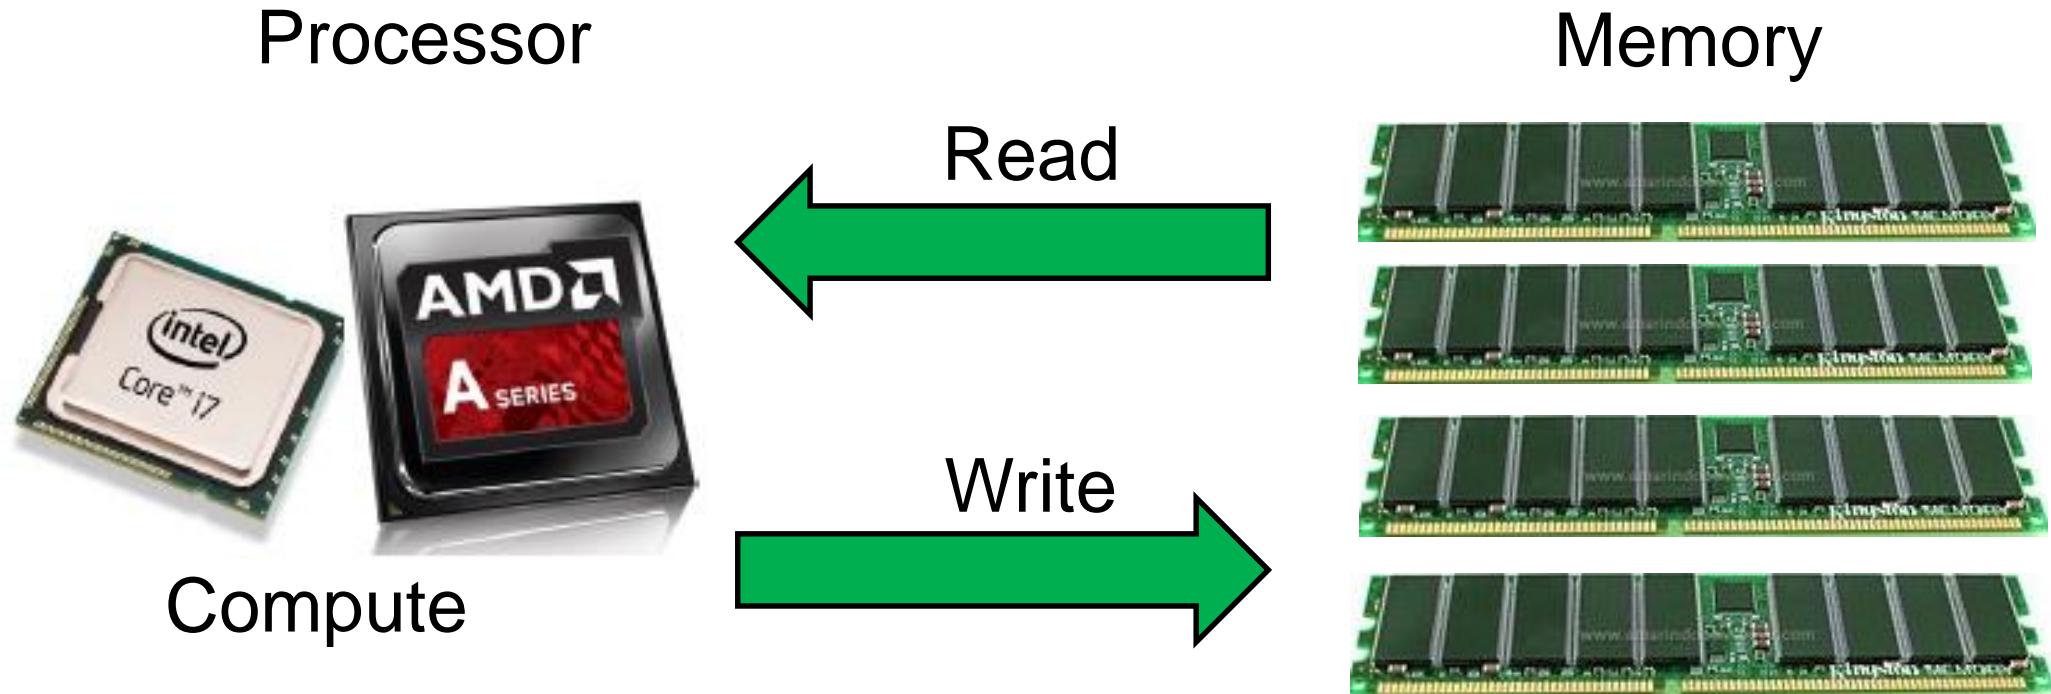
Compute



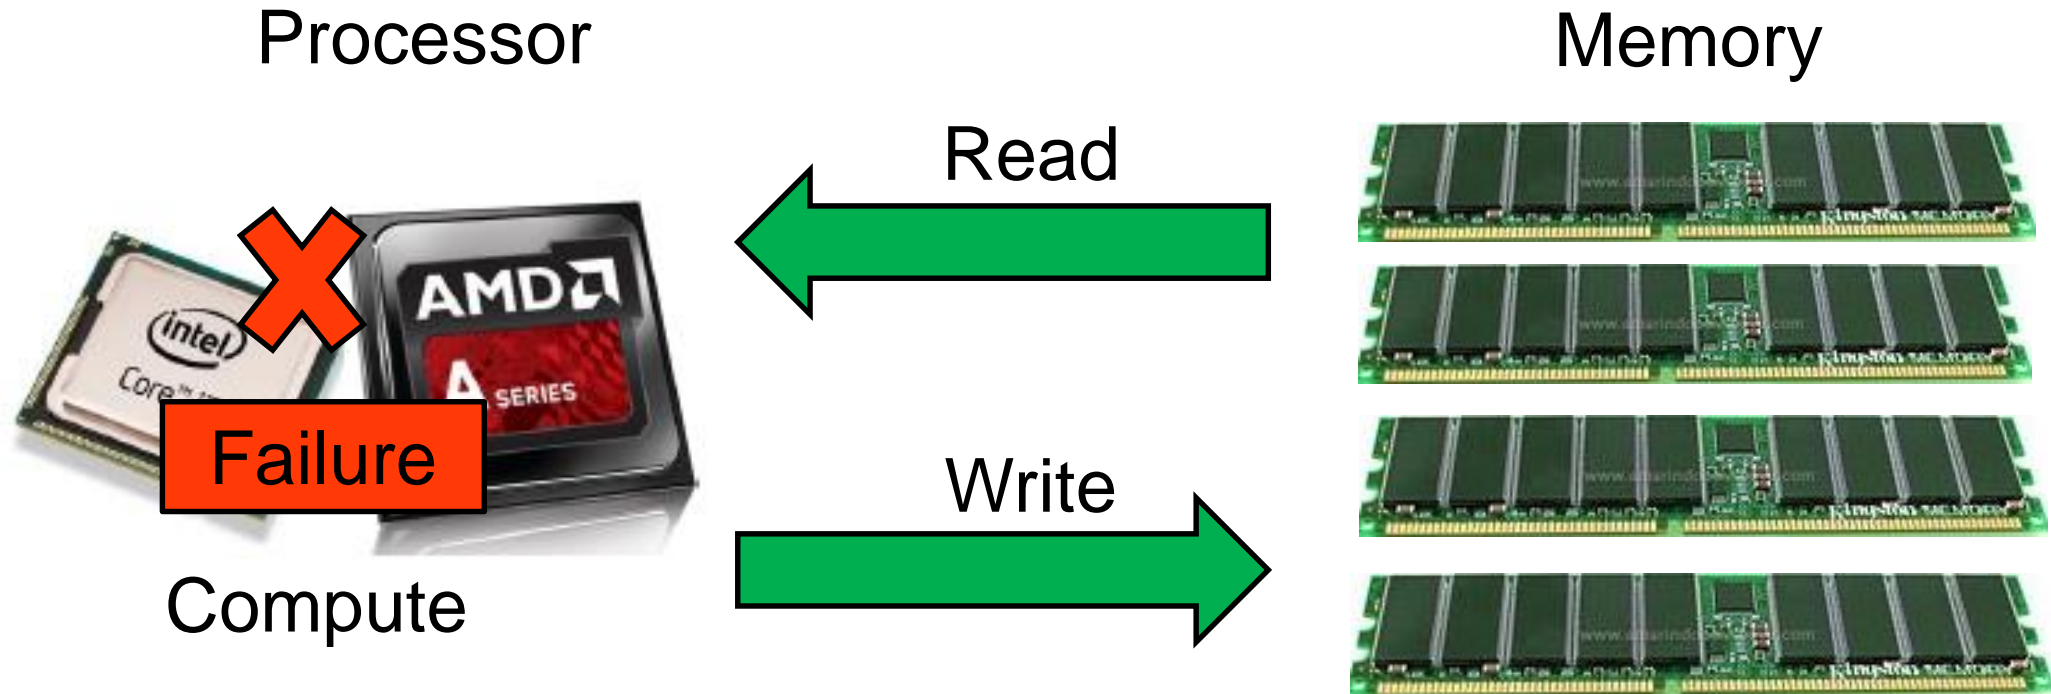
Two Key Components: Processor and Memory



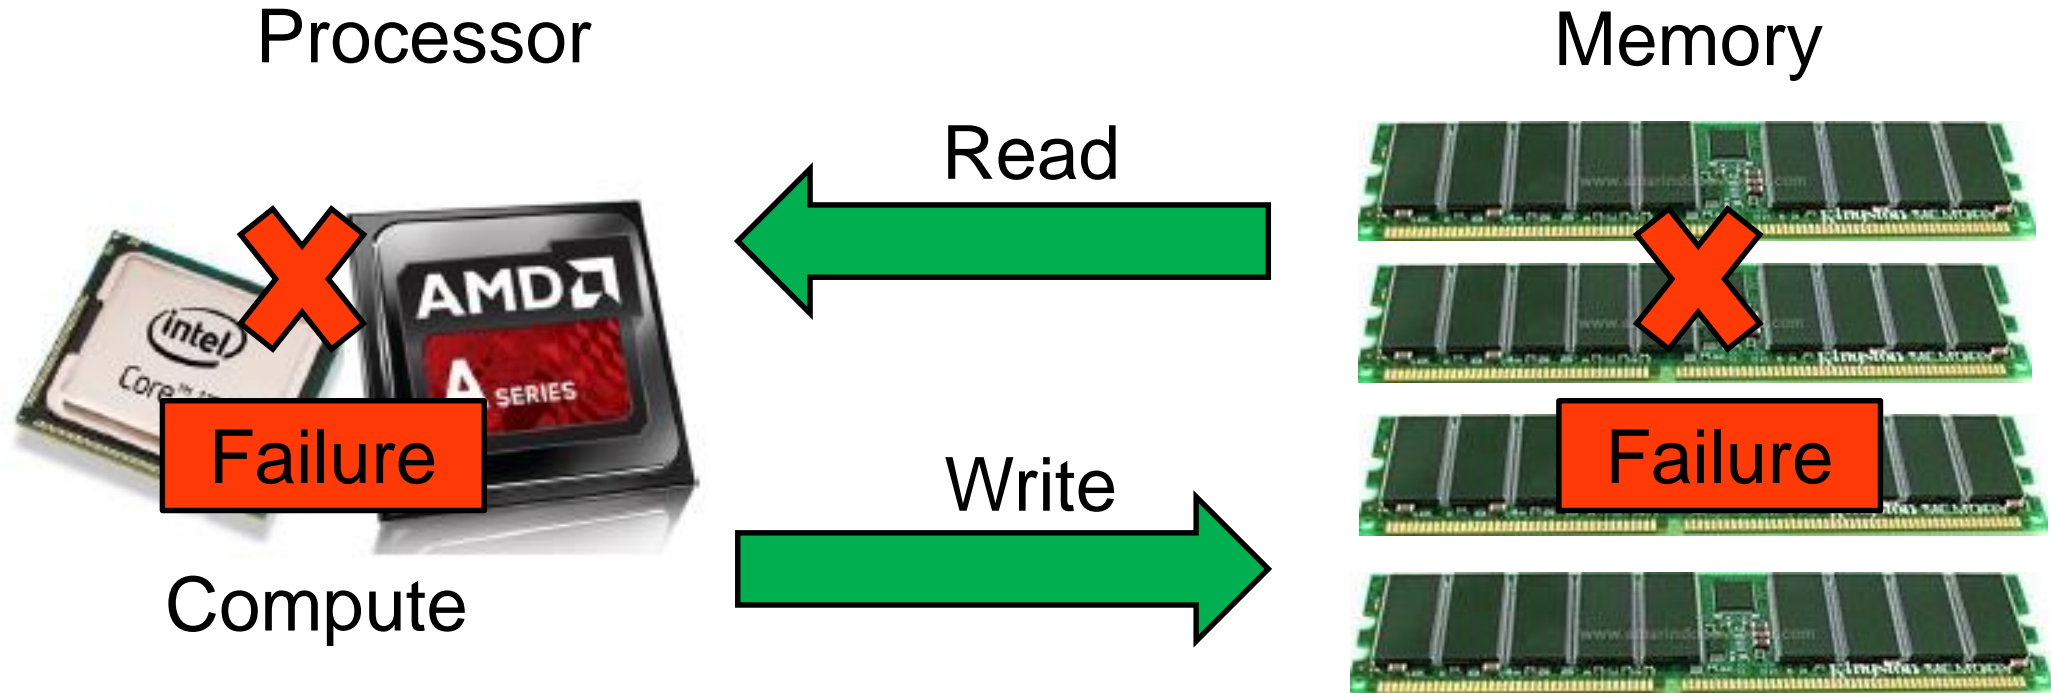
Two Key Sources of Failures: Processor and Memory



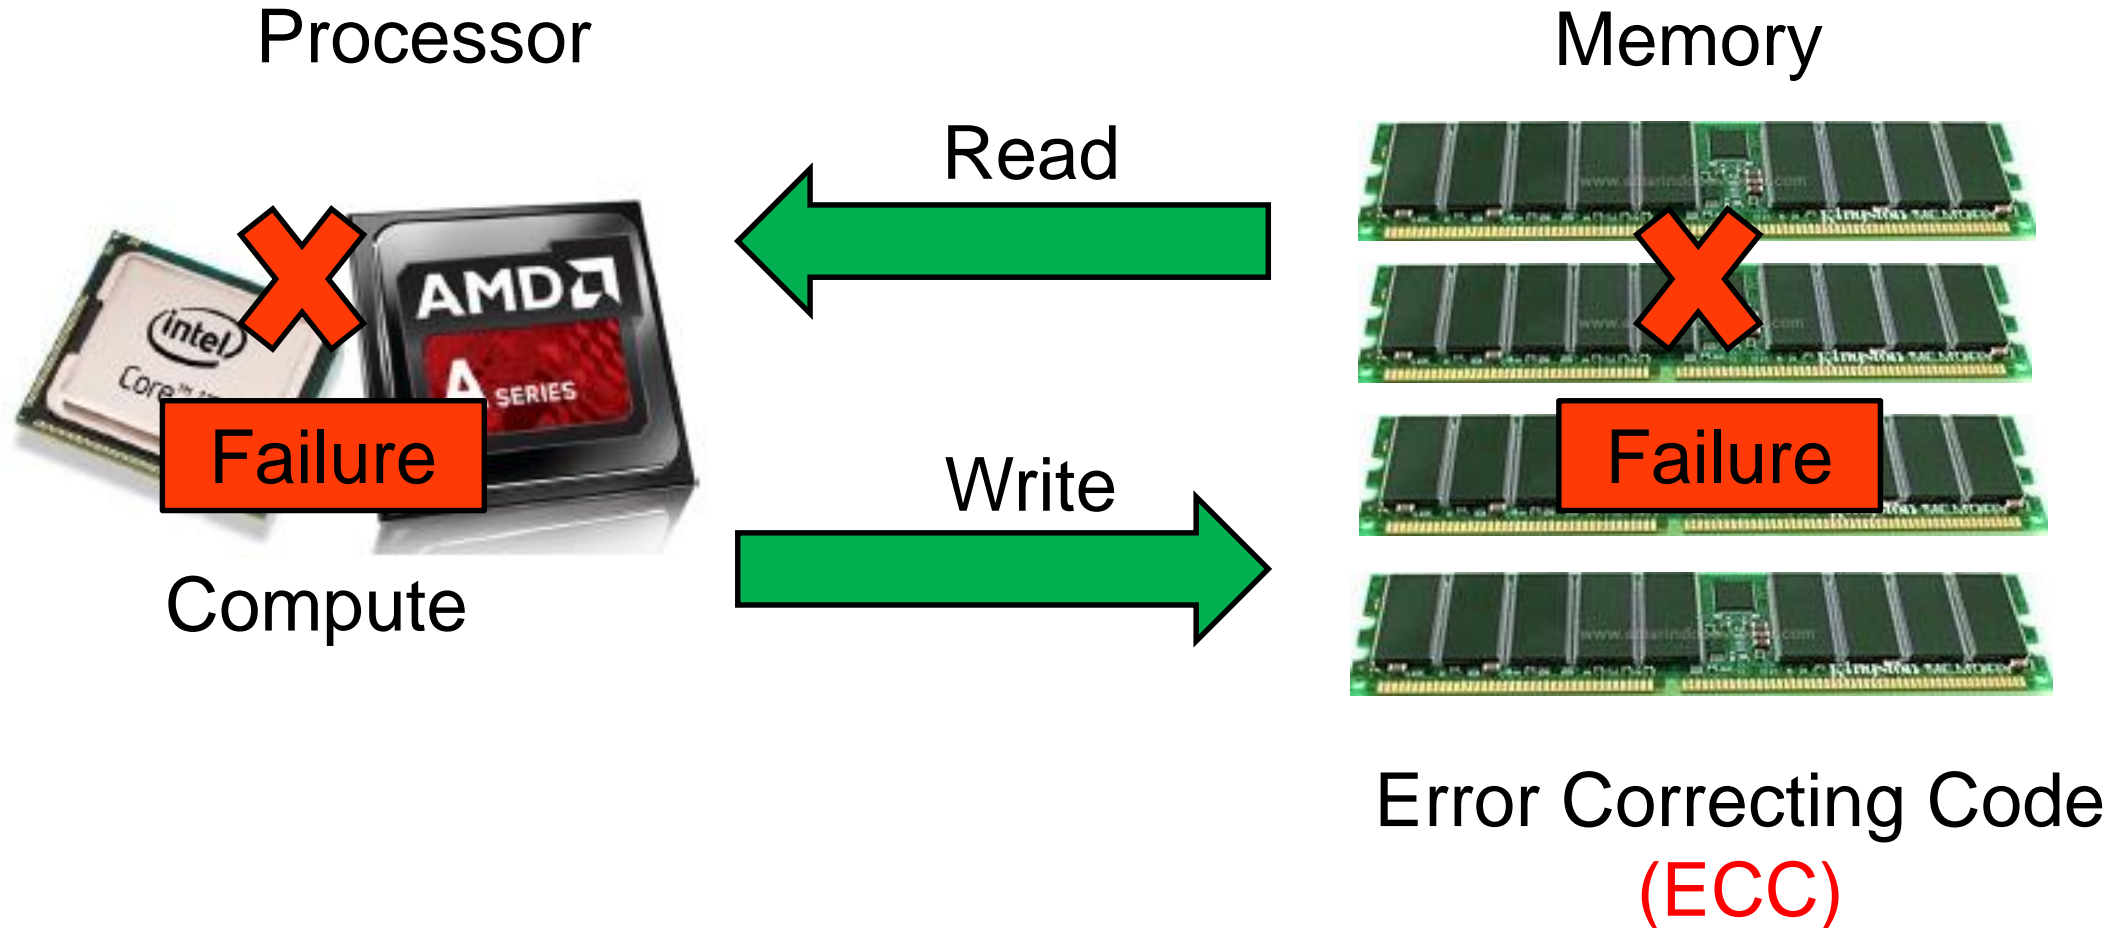
Two Key Sources of Failures: Processor and Memory



Two Key Sources of Failures: Processor and Memory



Two Key Sources of Failures: Processor and Memory



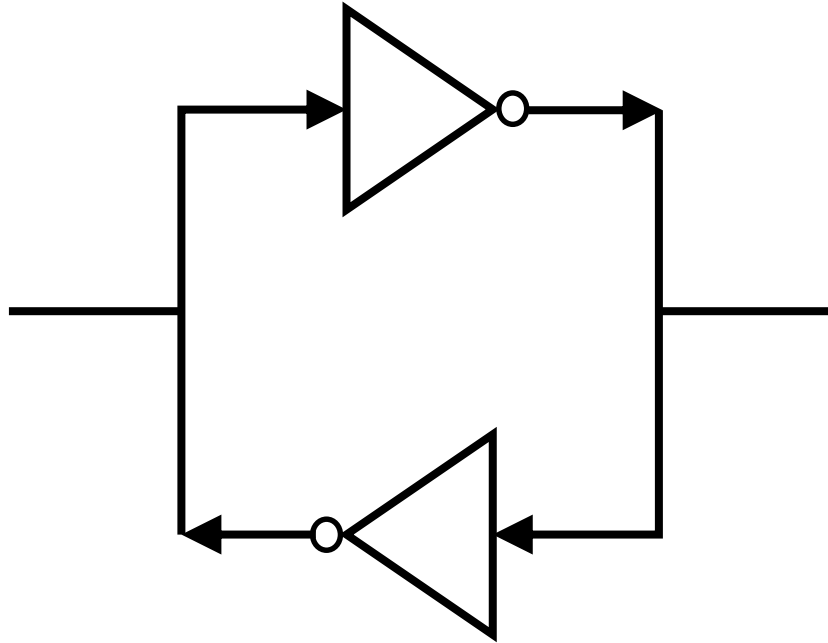
Processor Failures are a Result of Underlying Faults

Processor



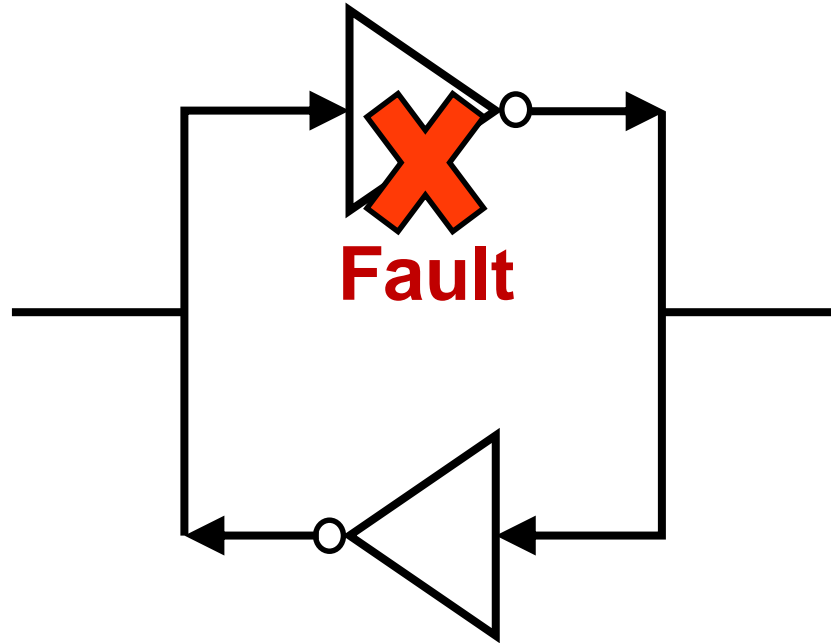
Processor Failures are a Result of Underlying Faults

Processor



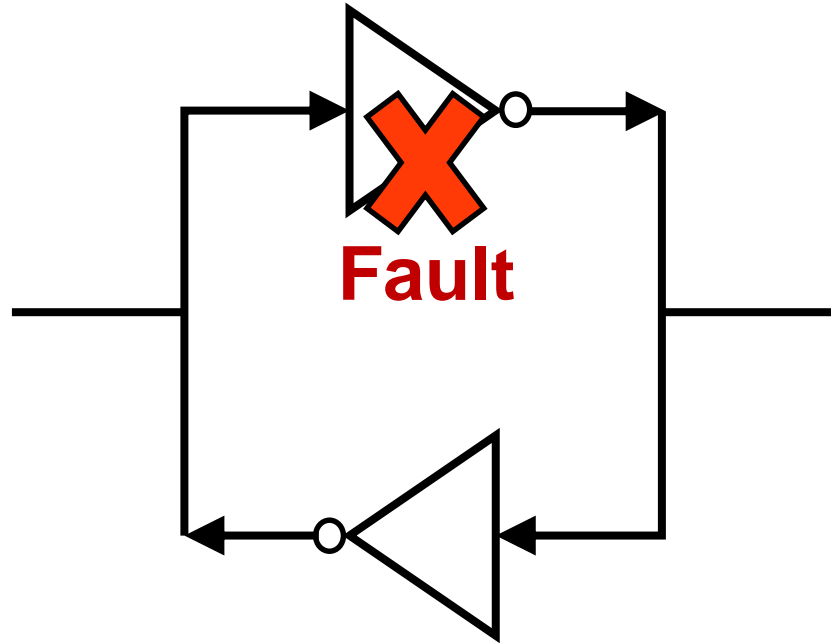
Processor Failures are a Result of Underlying Faults

Processor



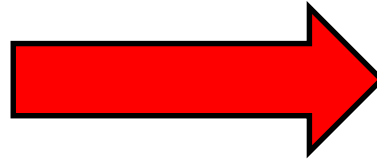
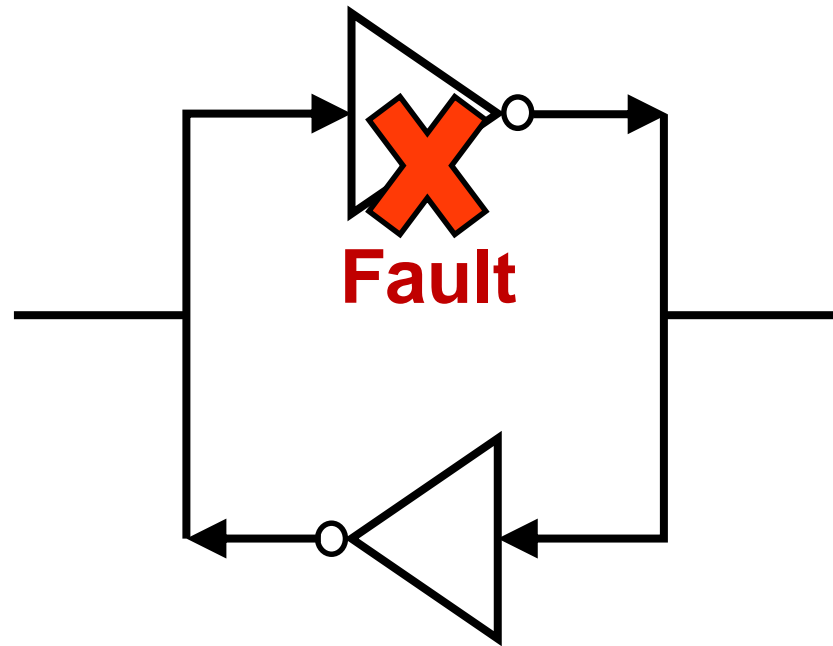
Processor Failures are a Result of Underlying Faults

Processor



Example: **non-functioning** or **defective** transistor

Processor Failures are a Result of Underlying Faults



Processor



Example: **non-functioning** or **defective** transistor

Three Types of Faults

Three Types of Faults

1. **Permanent** – to be fixed at hardware level or cannot be fixed

Three Types of Faults

1. **Permanent** – to be fixed at hardware level or cannot be fixed
-- Wearout, Electromigration

Three Types of Faults

1. **Permanent** – to be fixed at hardware level or cannot be fixed
-- Wearout, Electromigration
2. **Intermittent** – under specific conditions

Three Types of Faults

1. **Permanent** – to be fixed at hardware level or cannot be fixed
 - Wearout, Electromigration
2. **Intermittent** – under specific conditions
 - Elevated temperature

Three Types of Faults

1. **Permanent** – to be fixed at hardware level or cannot be fixed
 - Wearout, Electromigration
2. **Intermittent** – under specific conditions
 - Elevated temperature
3. **Transient** – random, temporary, non-reproducible

Three Types of Faults

1. **Permanent** – to be fixed at hardware level or cannot be fixed
 - Wearout, Electromigration
2. **Intermittent** – under specific conditions
 - Elevated temperature
3. **Transient** – random, temporary, non-reproducible
 - Energy particle strikes or **radiations**

What are Radiation-Induced Faults?

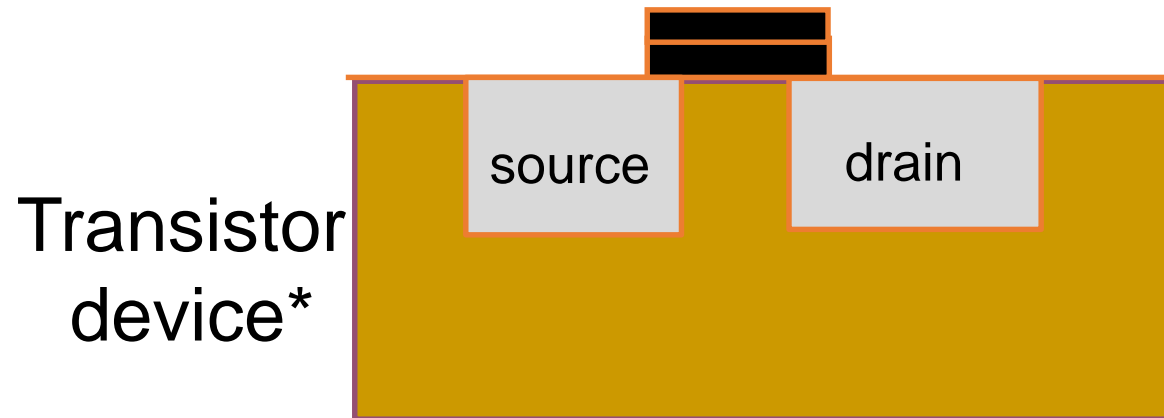
What are Radiation-Induced Faults?

- Main source: **neutrons** from deep space



What are Radiation-Induced Faults?

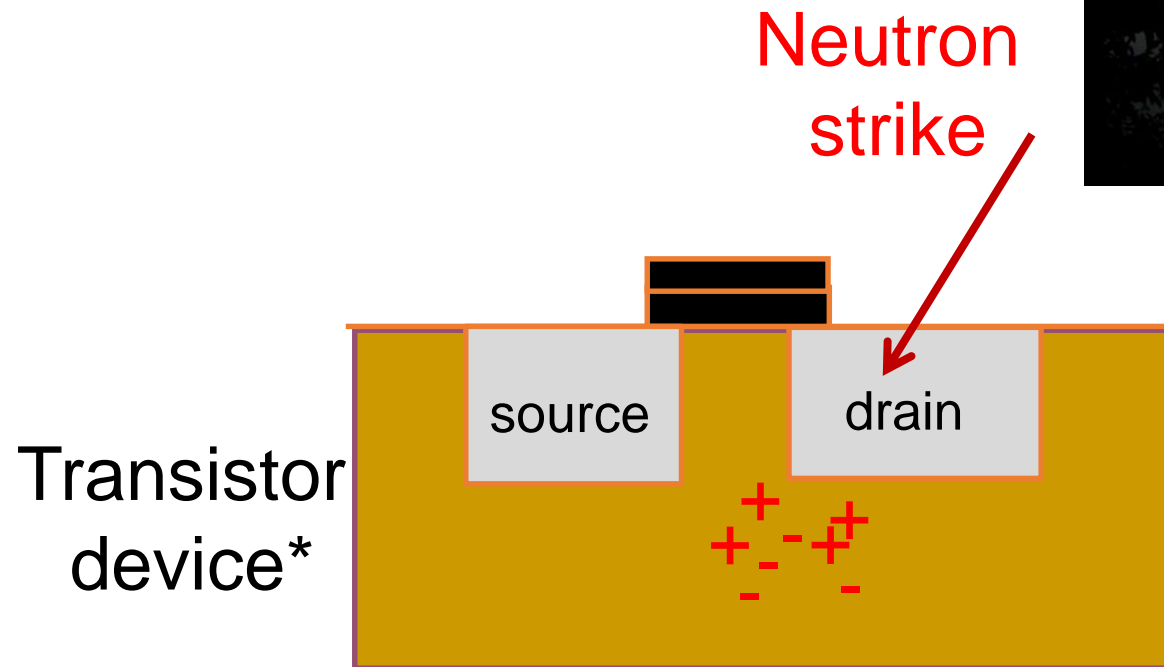
- Main source: **neutrons** from deep space



*[Soft Error Problem, Shubhu Mukherjee, HPCA2005]

What are Radiation-Induced Faults?

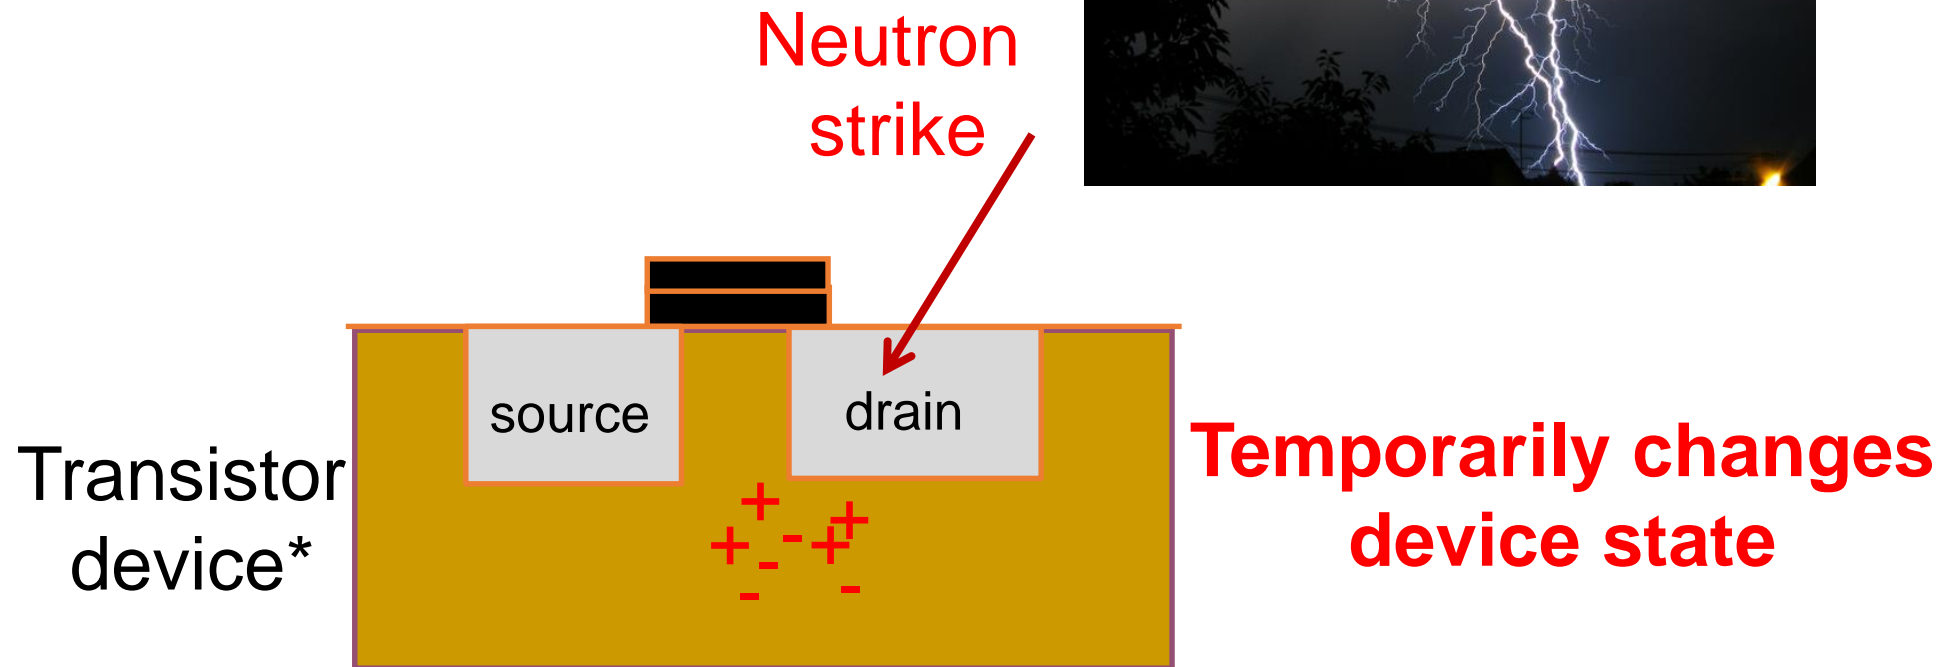
- Main source: **neutrons** from deep space



*[Soft Error Problem, Shubhu Mukherjee, HPCA2005]

What are Radiation-Induced Faults?

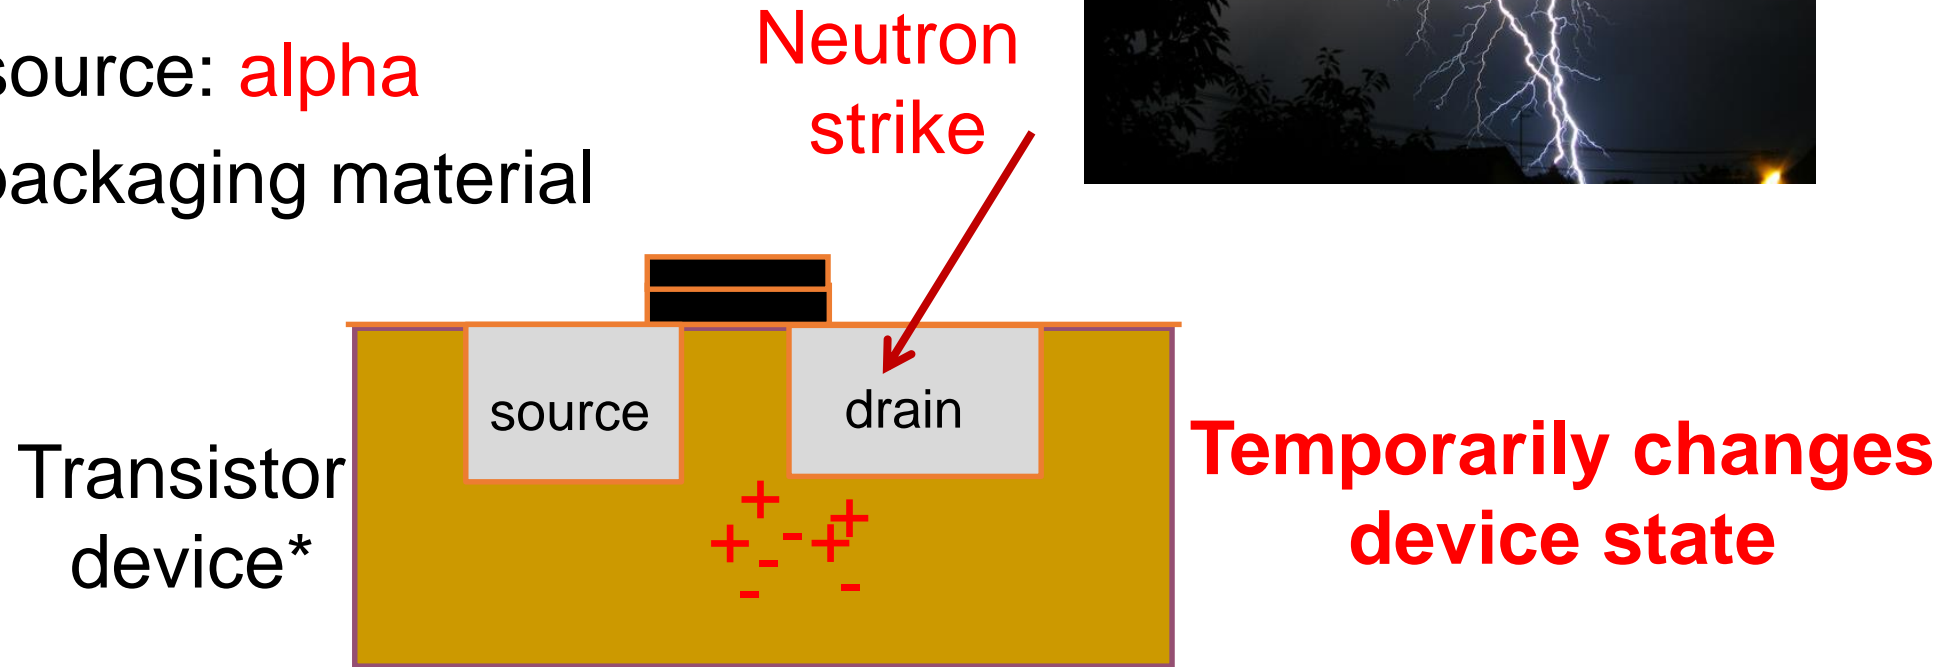
- Main source: **neutrons** from deep space



*[Soft Error Problem, Shubhu Mukherjee, HPCA2005]

What are Radiation-Induced Faults?

- Main source: **neutrons** from deep space
- Secondary source: **alpha particles** from packaging material



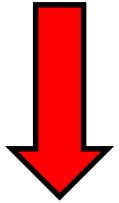
Faults → Errors → Failures

Faults → Errors → Failures

Fault

Faults → Errors → Failures

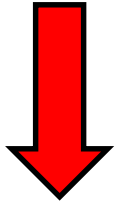
Fault



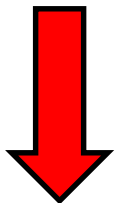
Error

Faults → Errors → Failures

Fault



Error

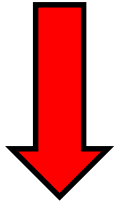


Failure

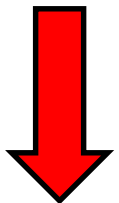
Faults → Errors → Failures

Fault

Radiations



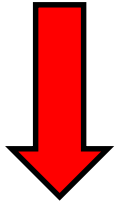
Error



Failure

Faults → Errors → Failures

Fault



Error



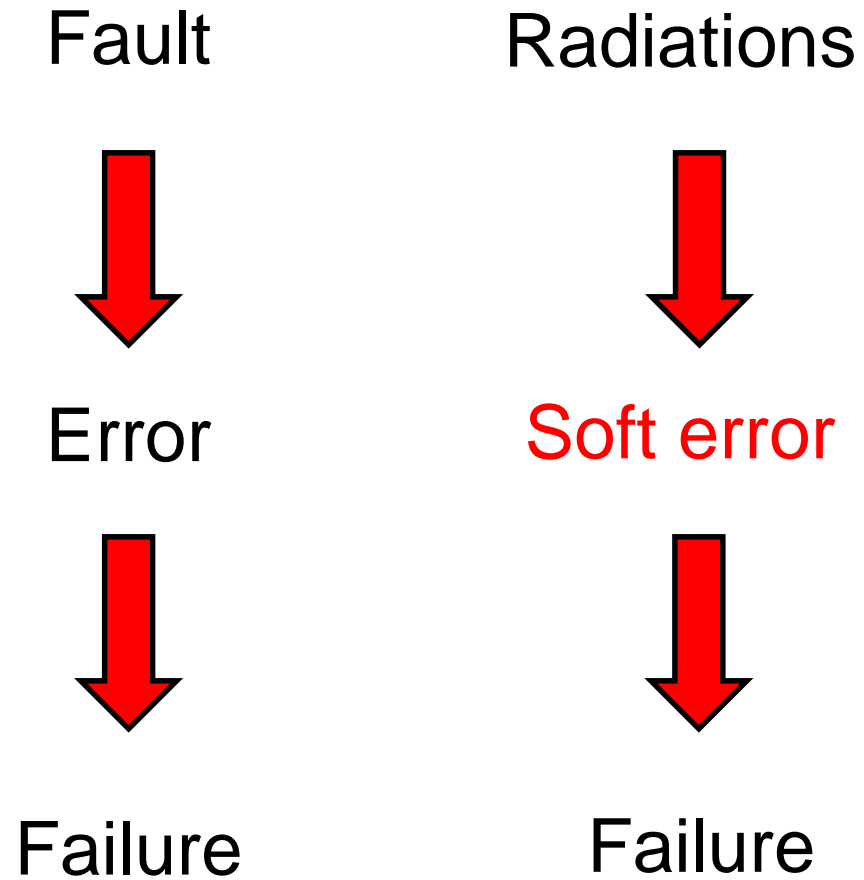
Failure

Radiations



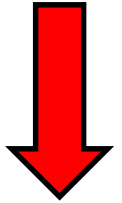
Soft error

Faults → Errors → Failures

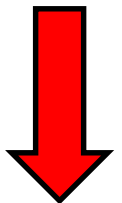


Faults → Errors → Failures

Fault



Error

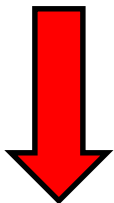


Failure

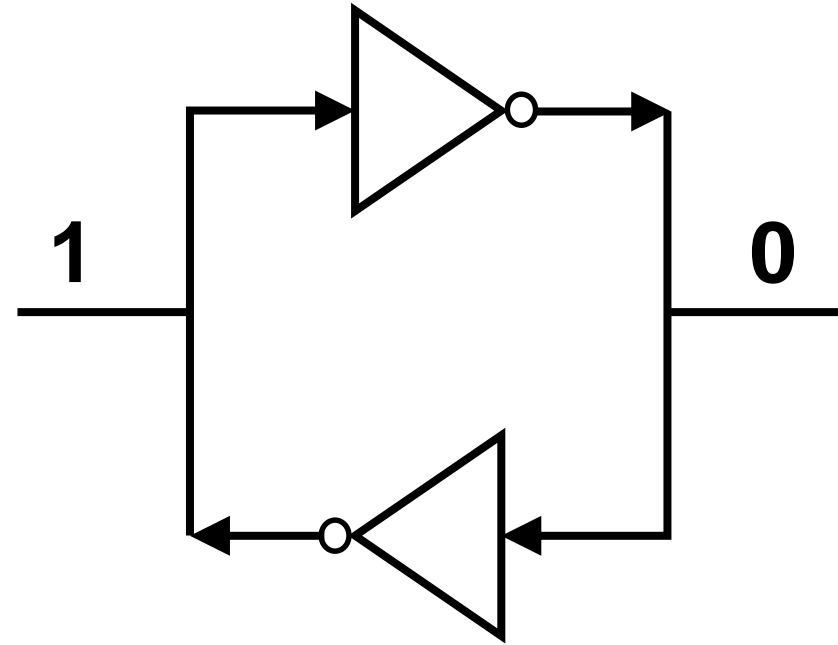
Radiations



Soft error

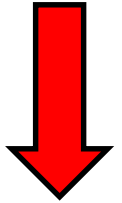


Failure

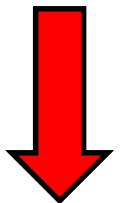


Faults → Errors → Failures

Fault



Error

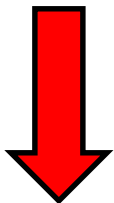


Failure

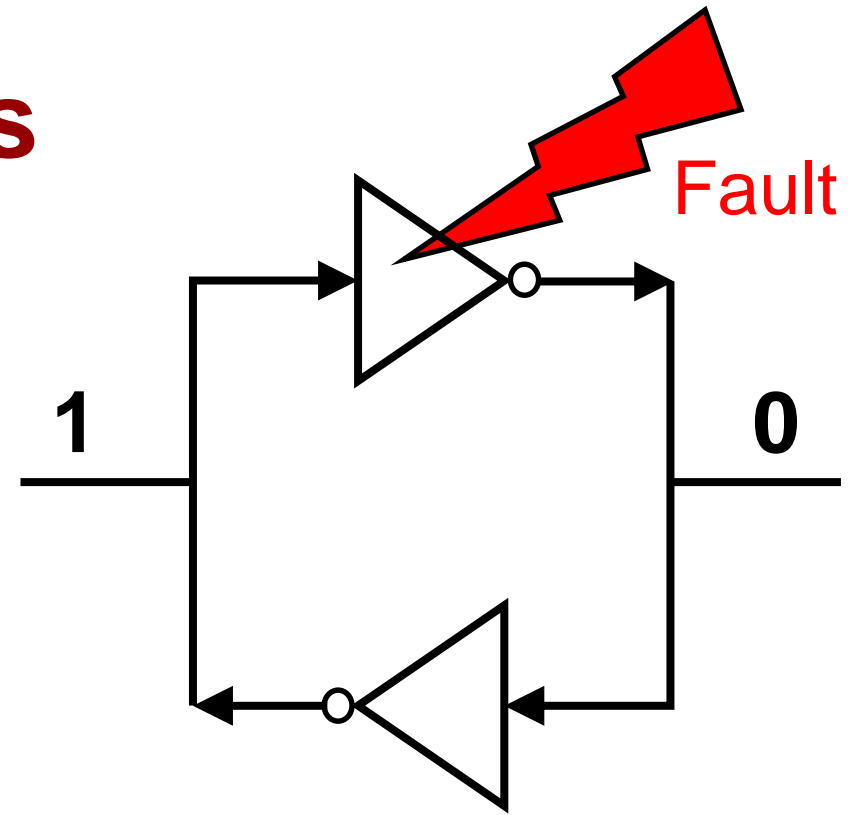
Radiations



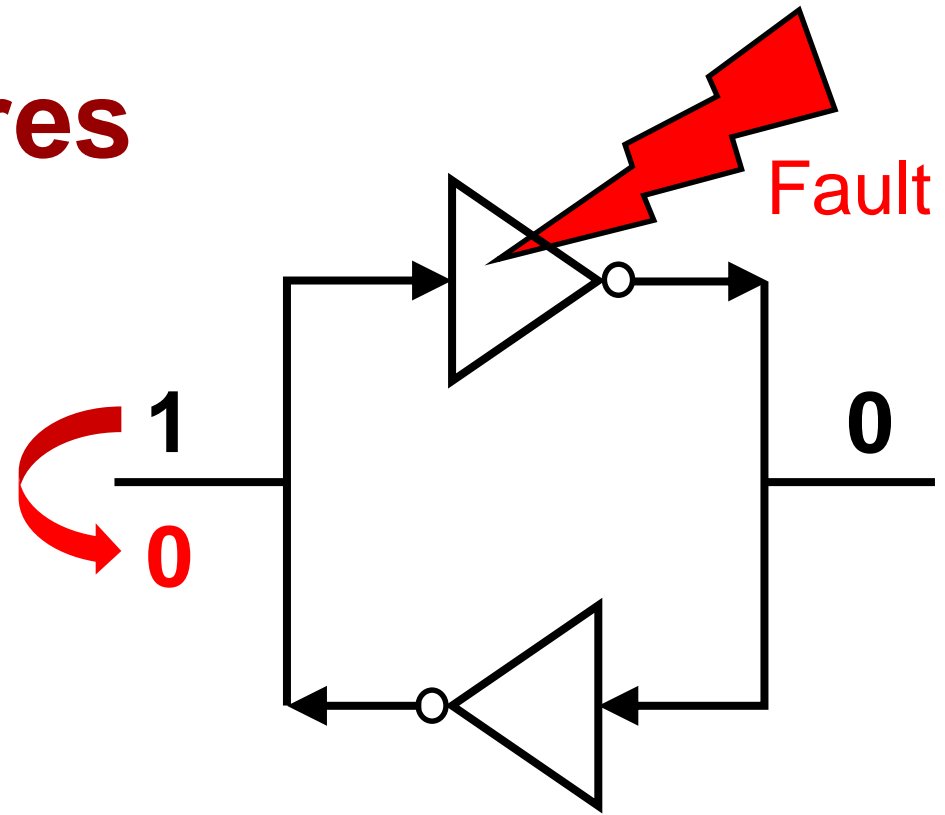
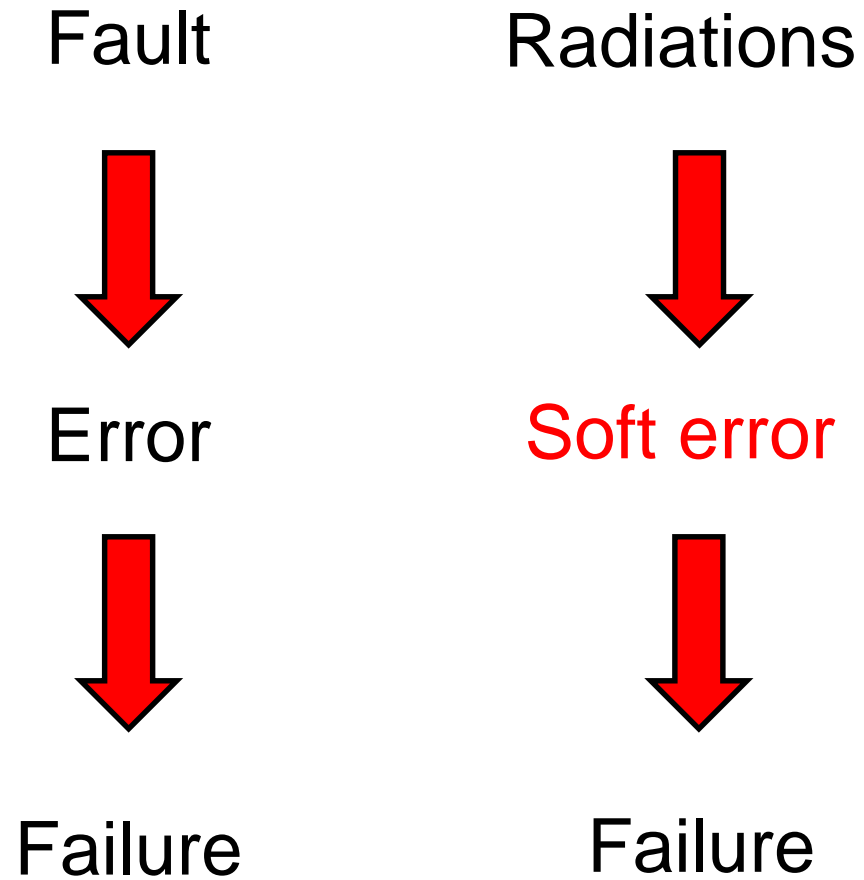
Soft error



Failure

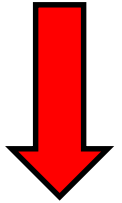


Faults → Errors → Failures

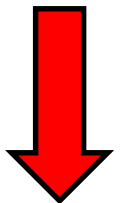


Faults → Errors → Failures

Fault



Error

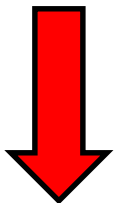


Failure

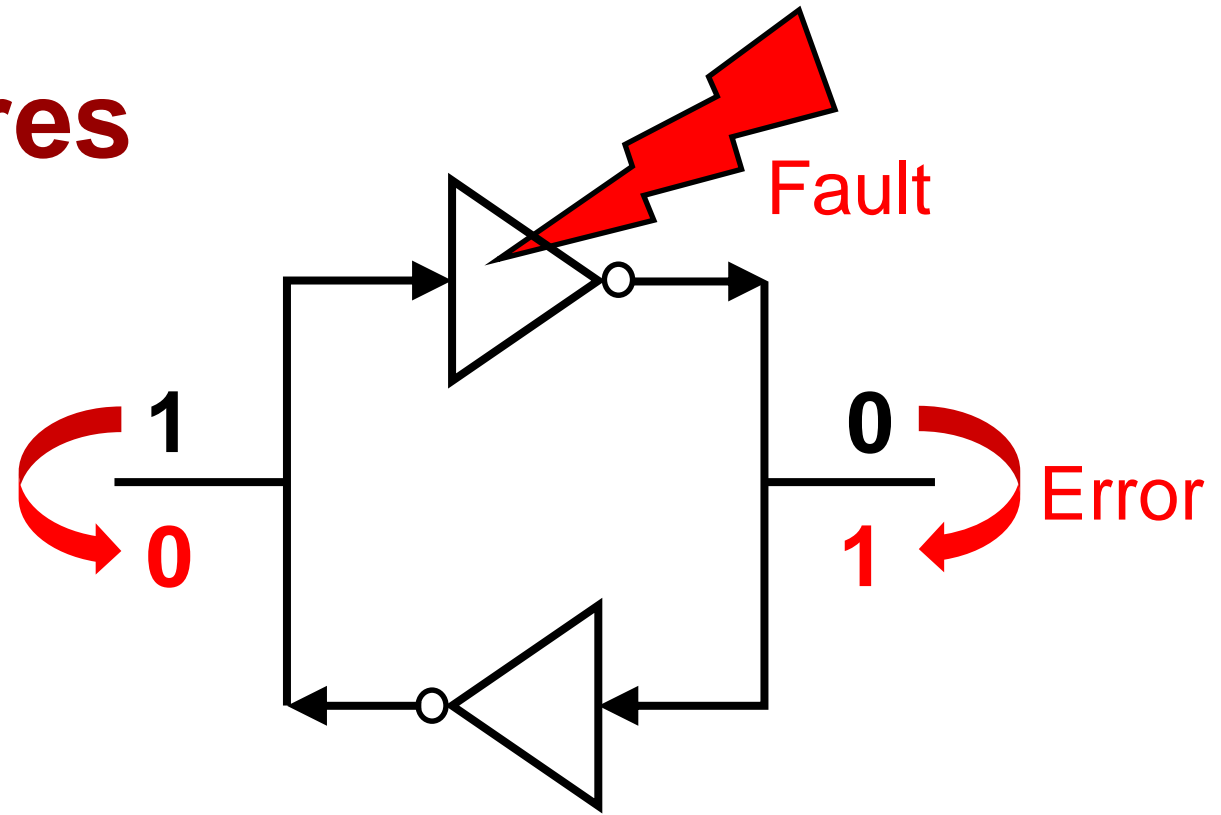
Radiations



Soft error

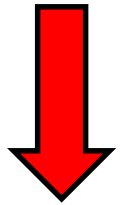


Failure

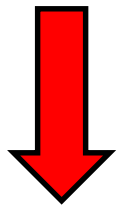


Faults → Errors → Failures

Fault



Error

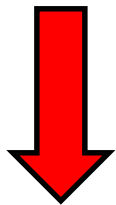


Failure

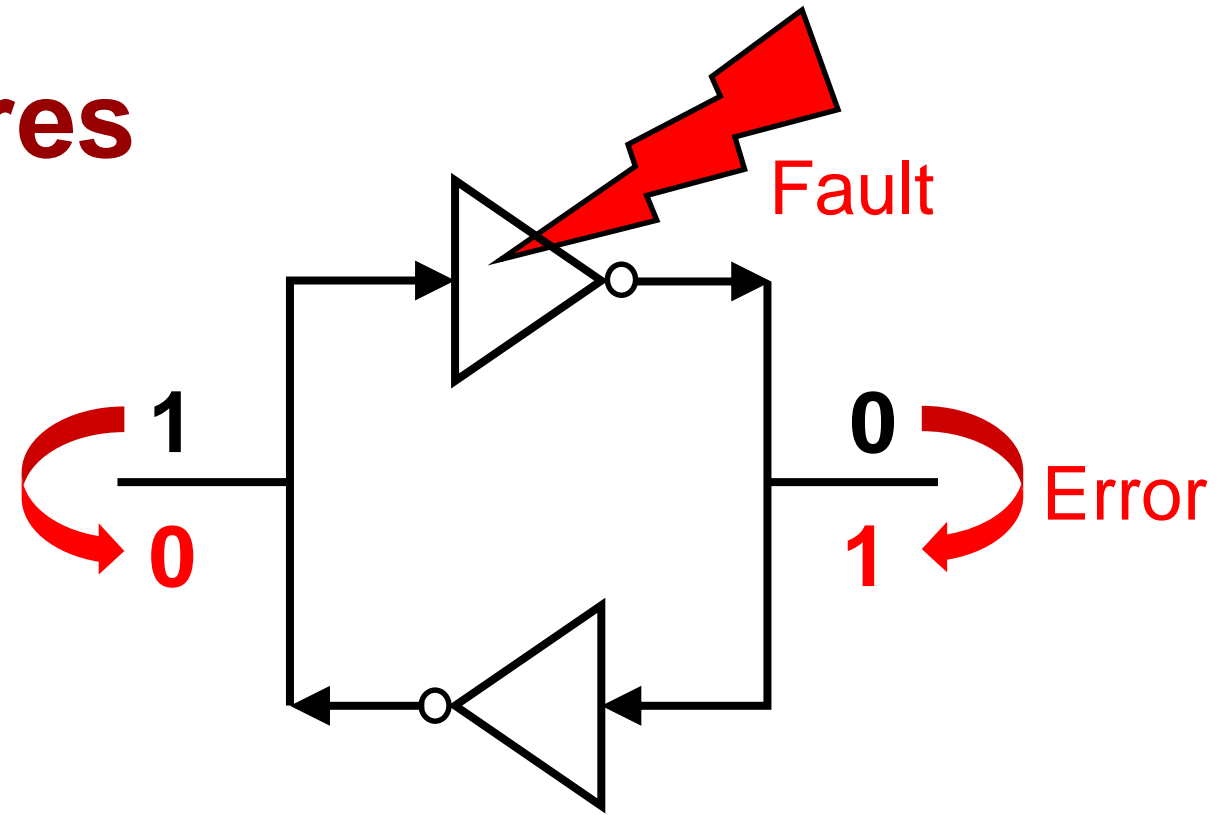
Radiations



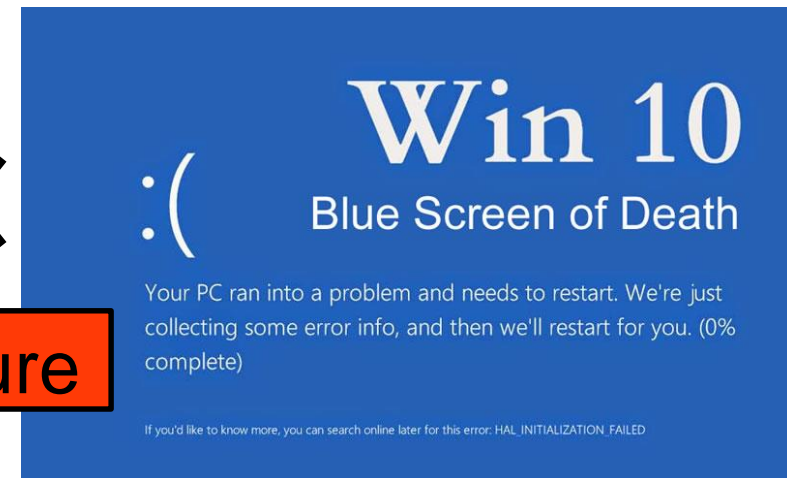
Soft error



Failure



Failure



Why have Soft Errors Become a Critical Issue?

Why have Soft Errors Become a Critical Issue?

High performance

Why have Soft Errors Become a Critical Issue?

High performance

- 1988 → 2017 more than 1000 X increase

Why have Soft Errors Become a Critical Issue?

High performance

- 1988 → 2017 more than 1000 X increase

Two key contributors

Why have Soft Errors Become a Critical Issue?

High performance

- 1988 → 2017 more than 1000 X increase

Two key contributors

1. **Technology scaling** – smaller transistors → better performance

Why have Soft Errors Become a Critical Issue?

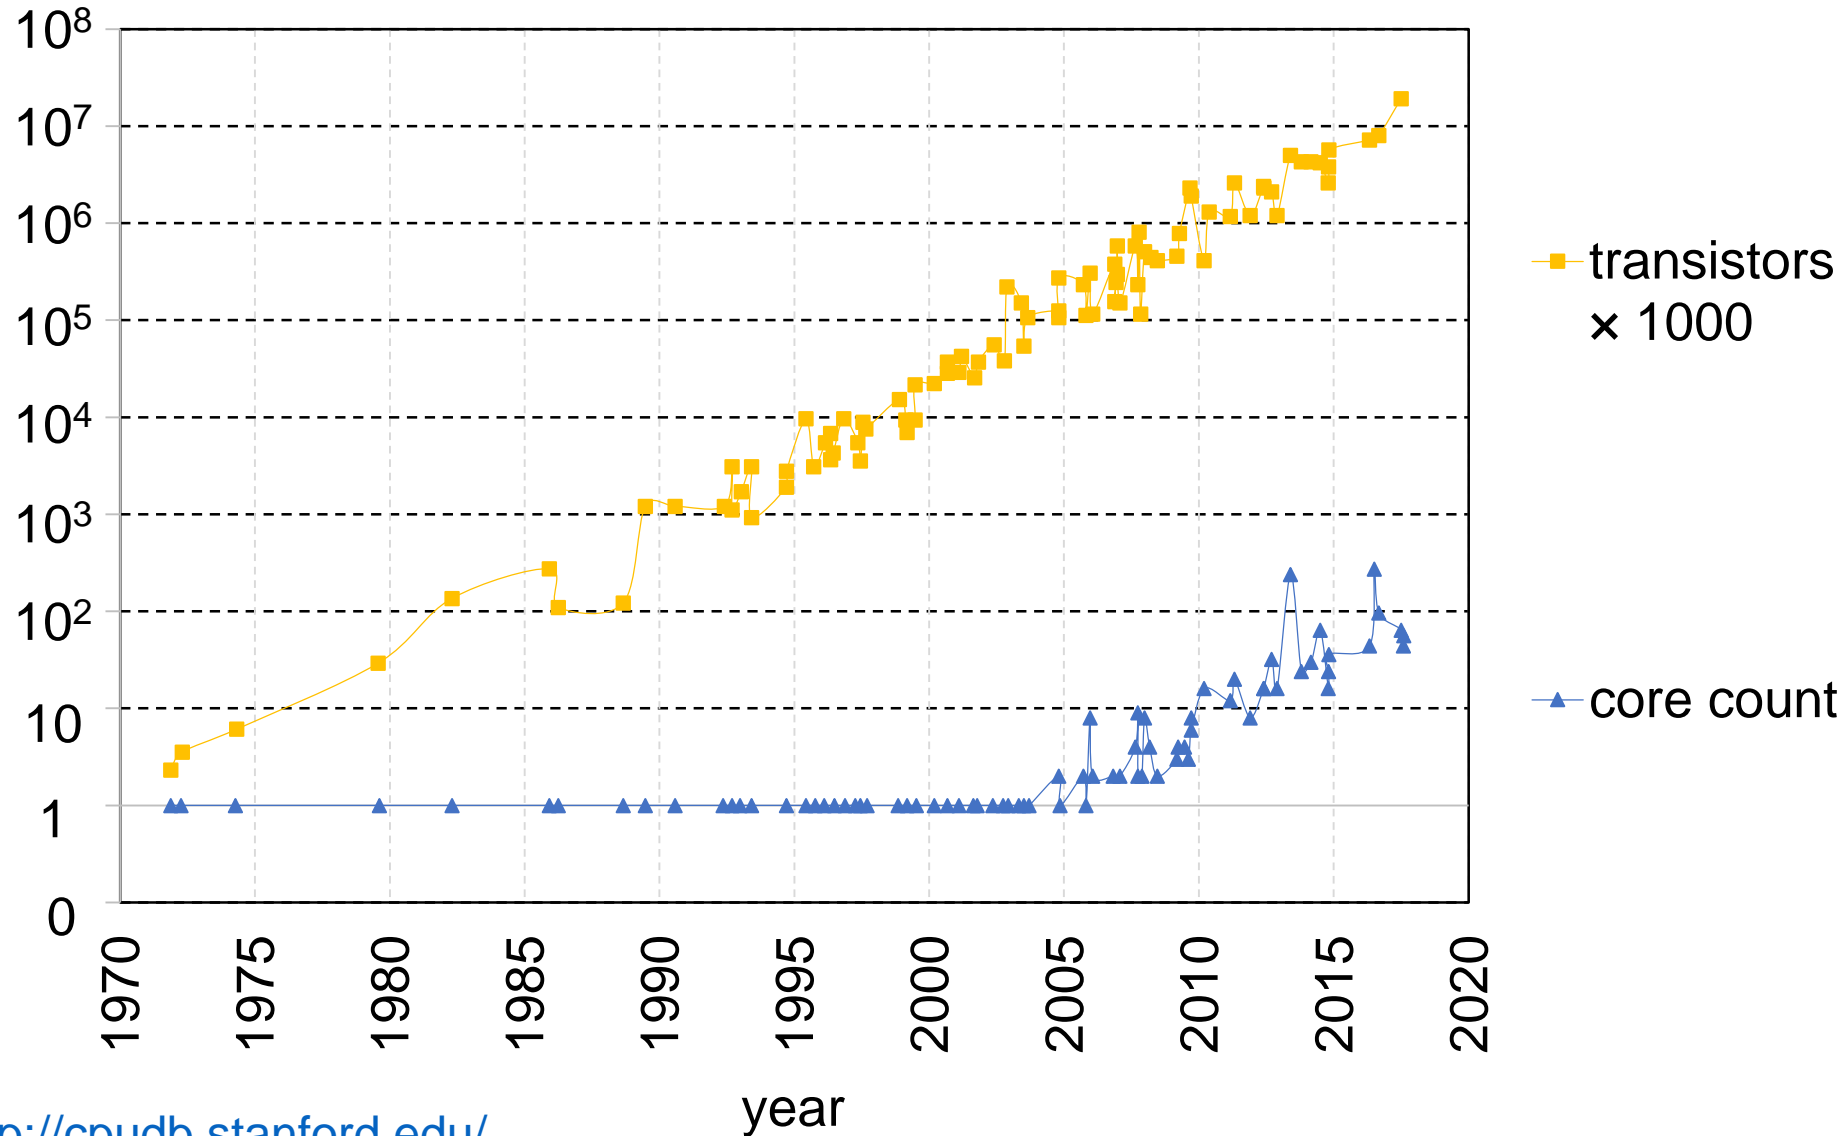
High performance

- 1988 → 2017 more than 1000 X increase

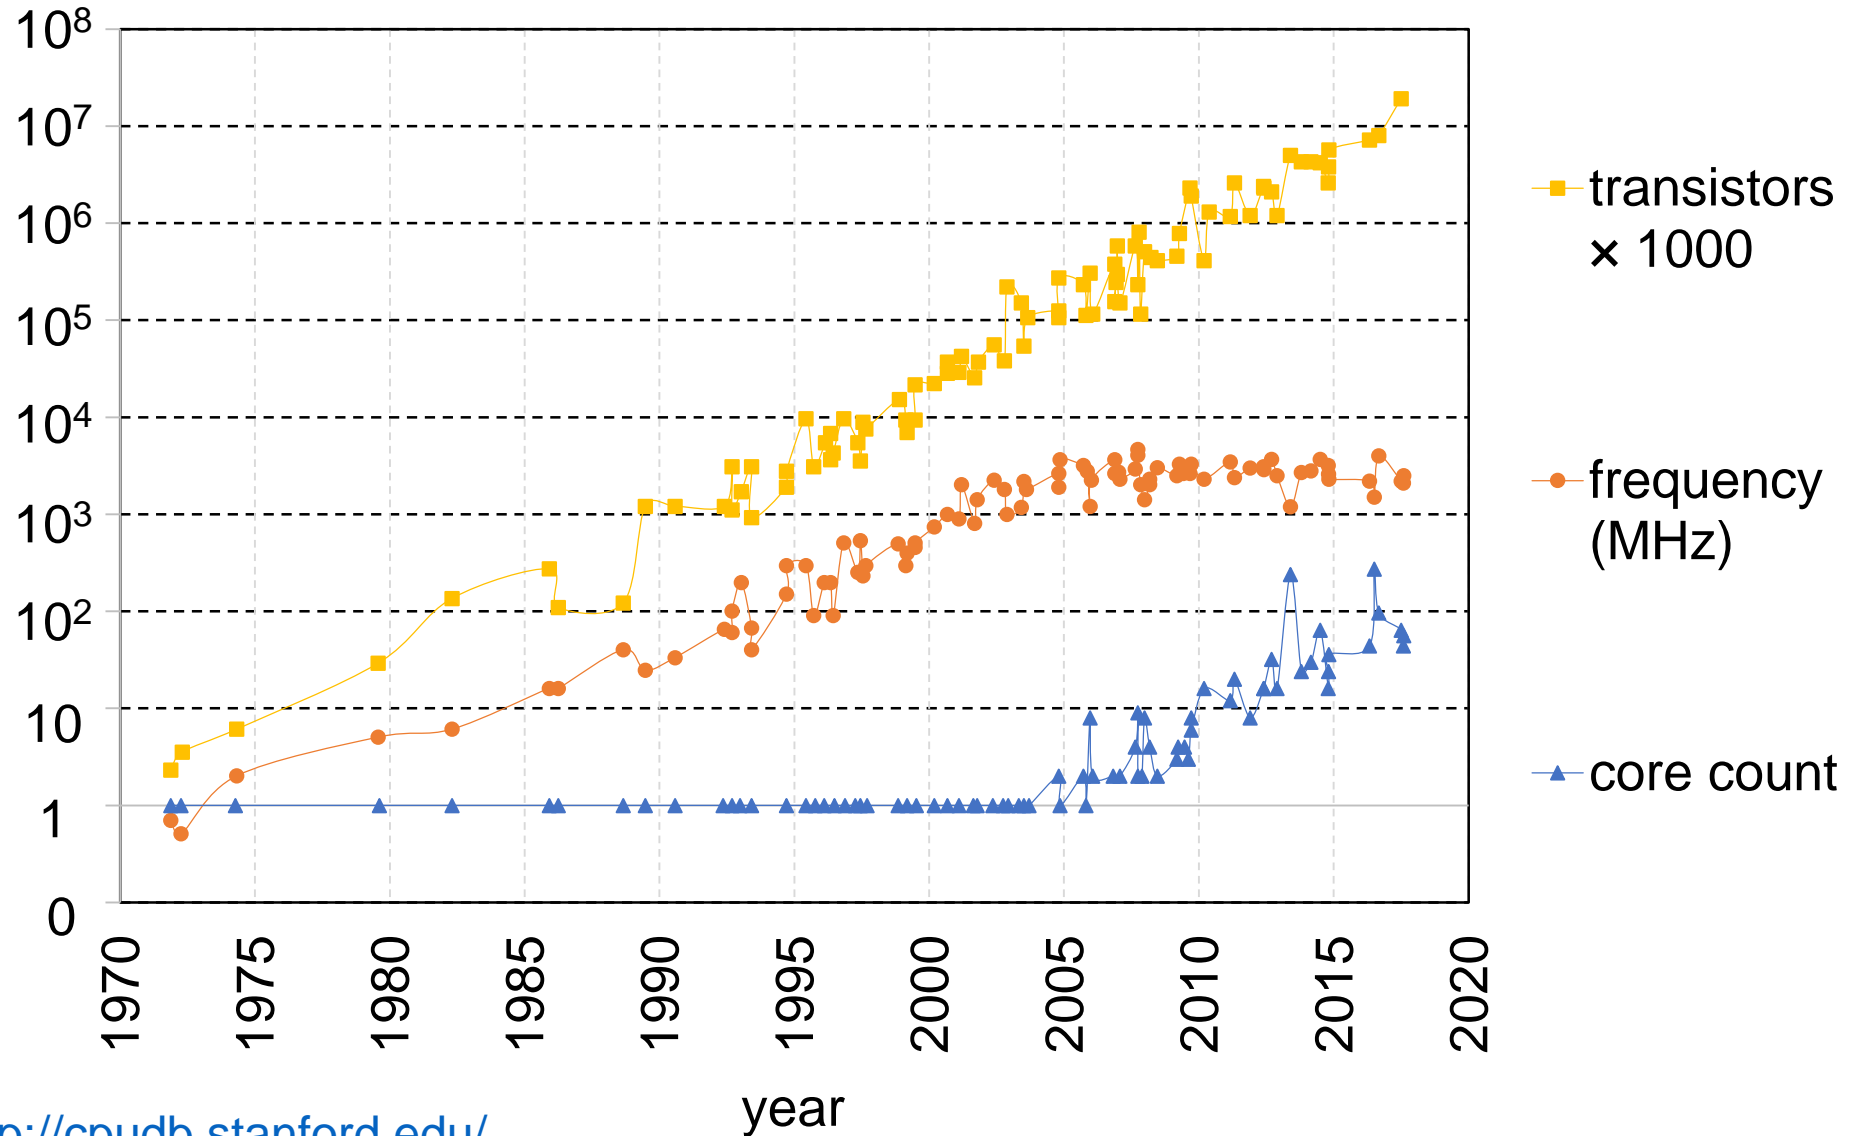
Two key contributors

1. **Technology scaling** – smaller transistors → better performance
2. **Microarchitectural innovations** – pipelining, out-of-order execution, prefetching, branch prediction

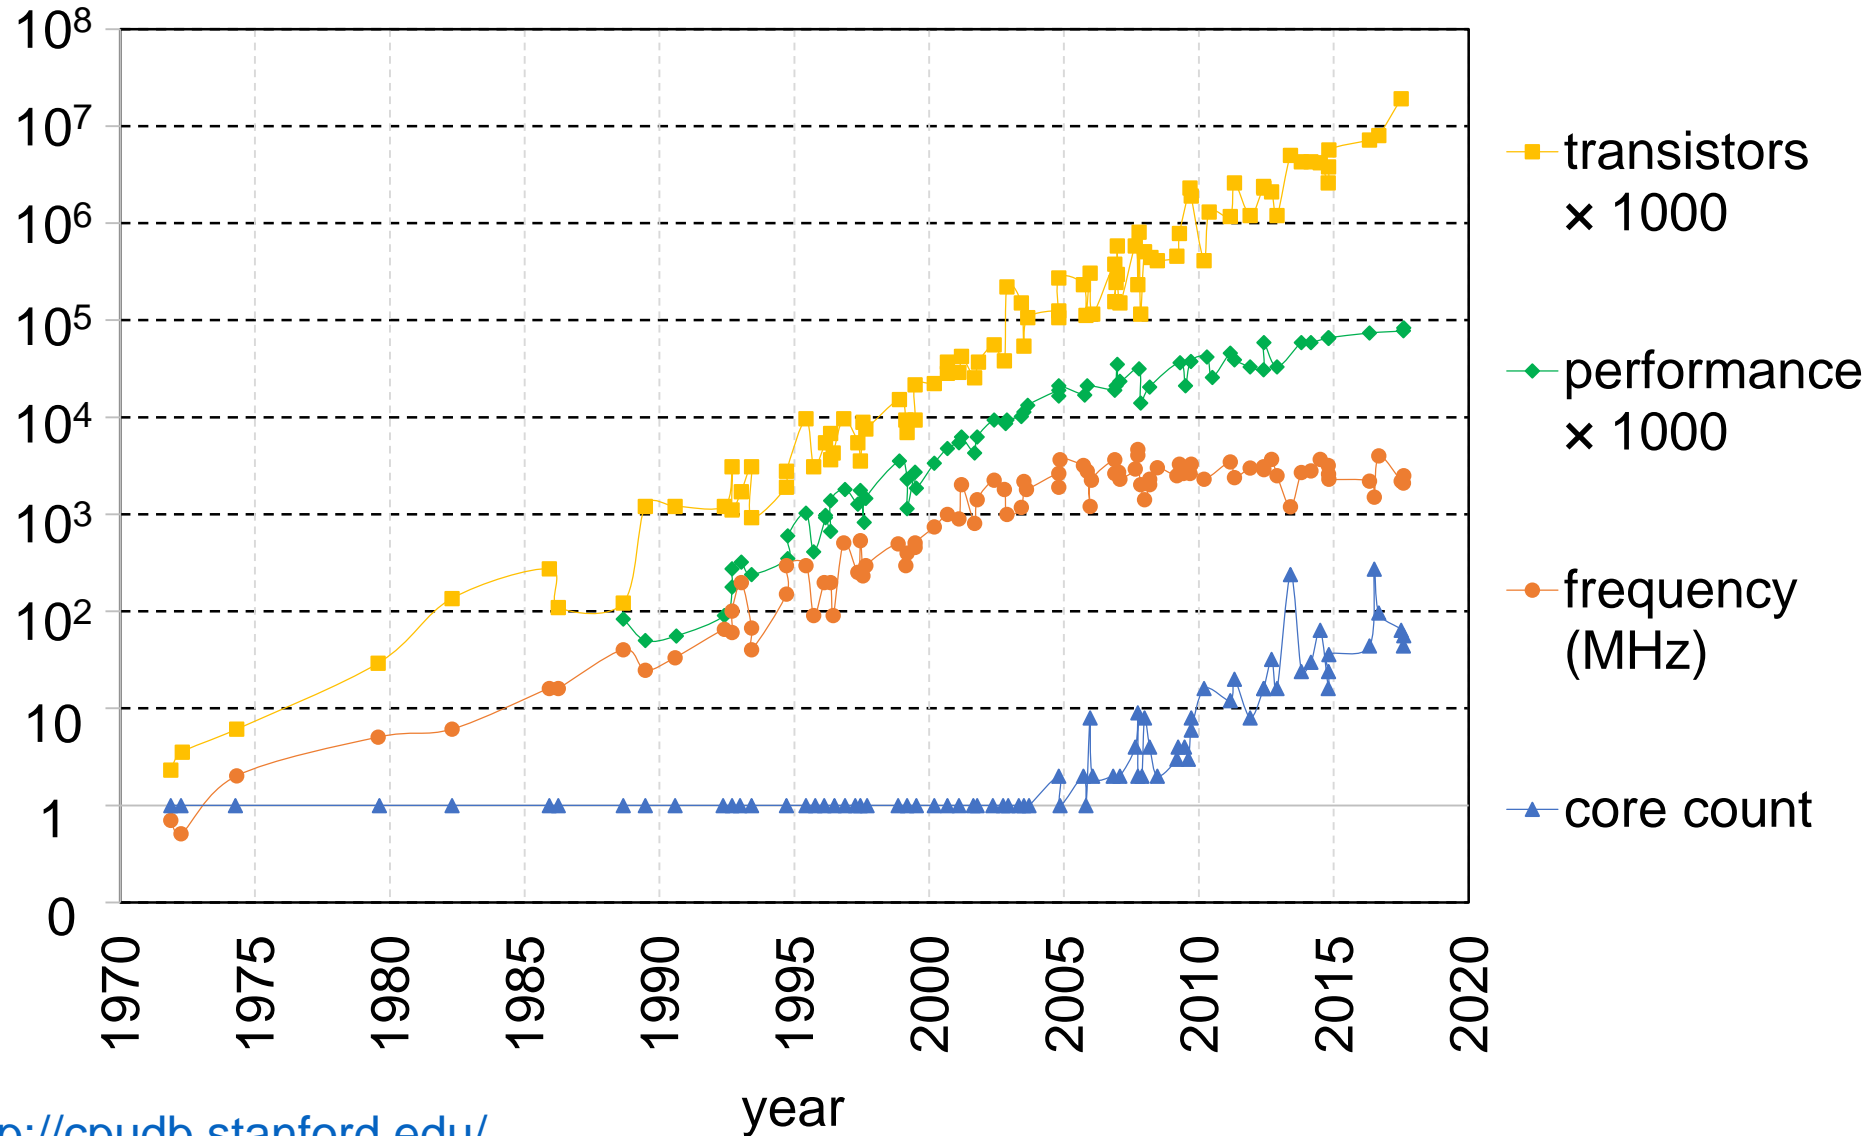
Consequences of Technology Scaling



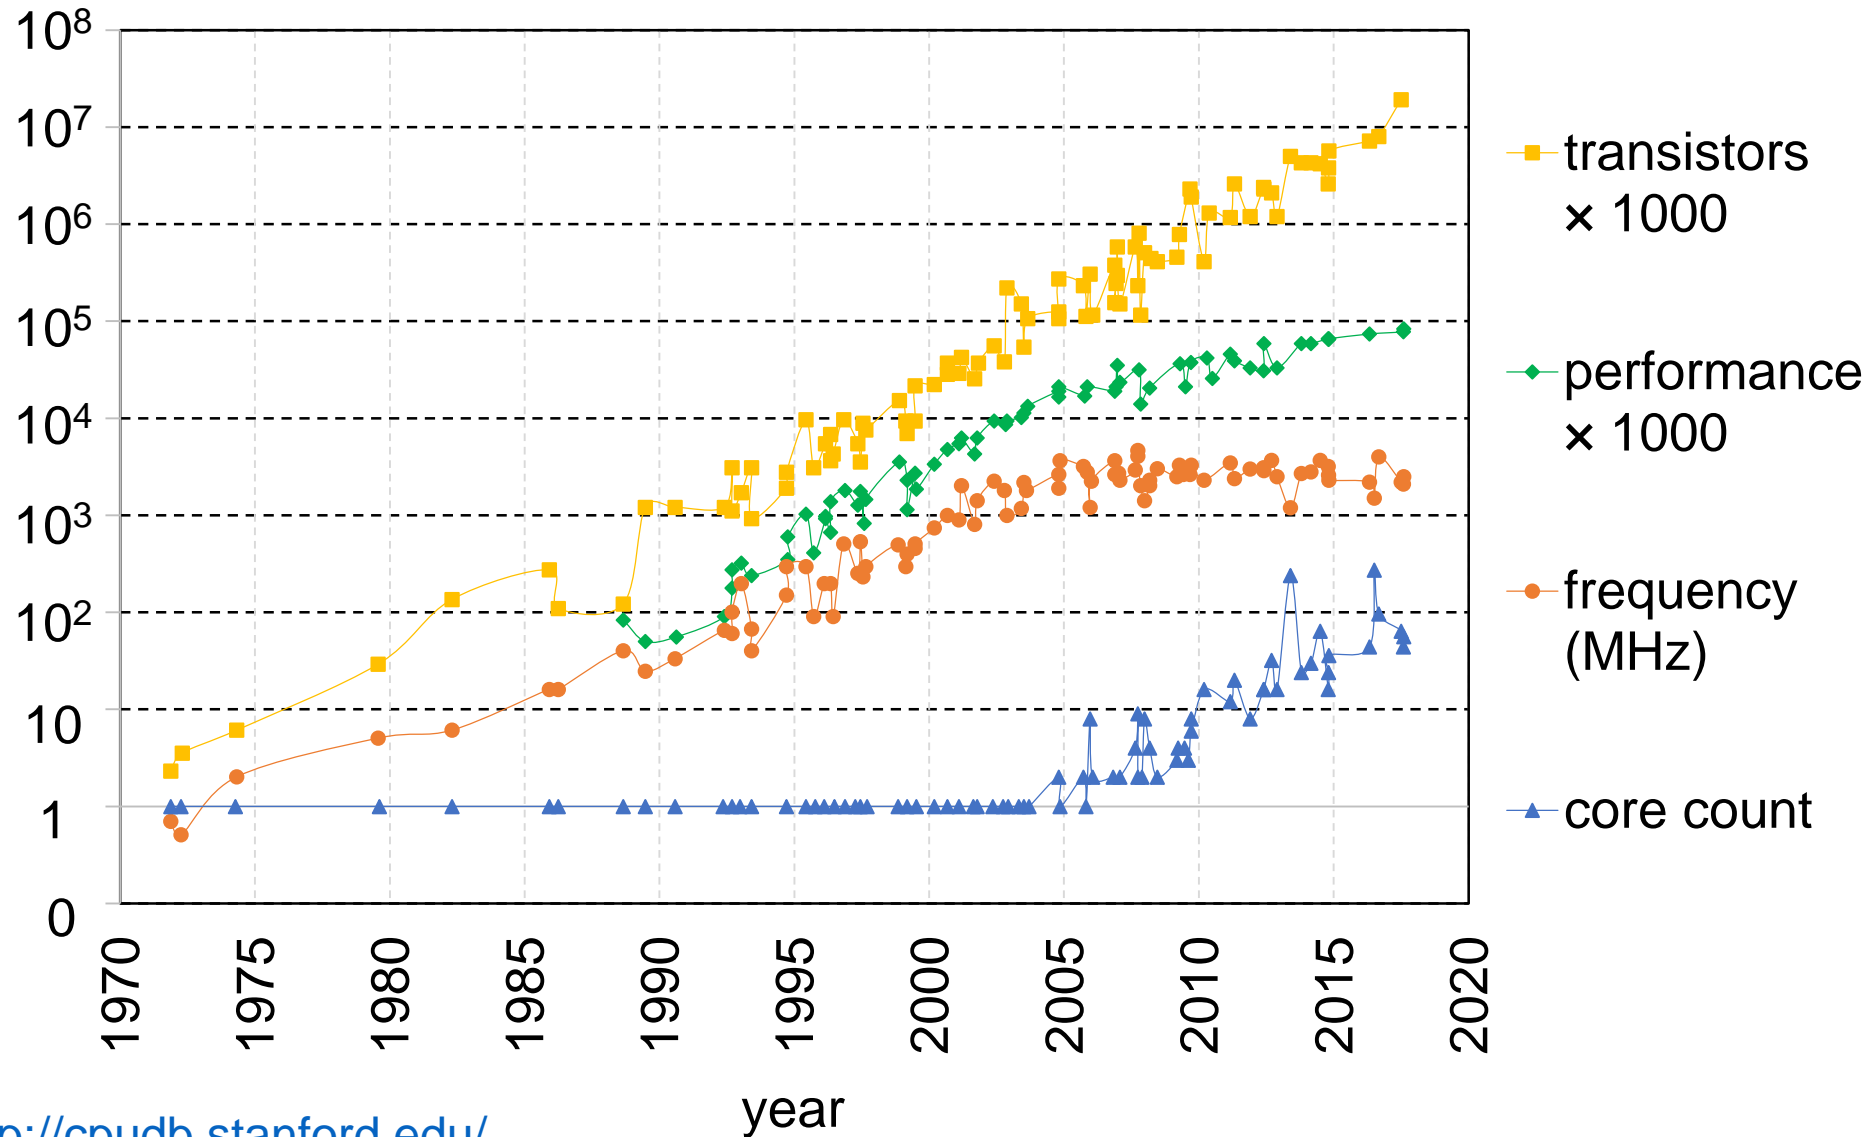
Consequences of Technology Scaling



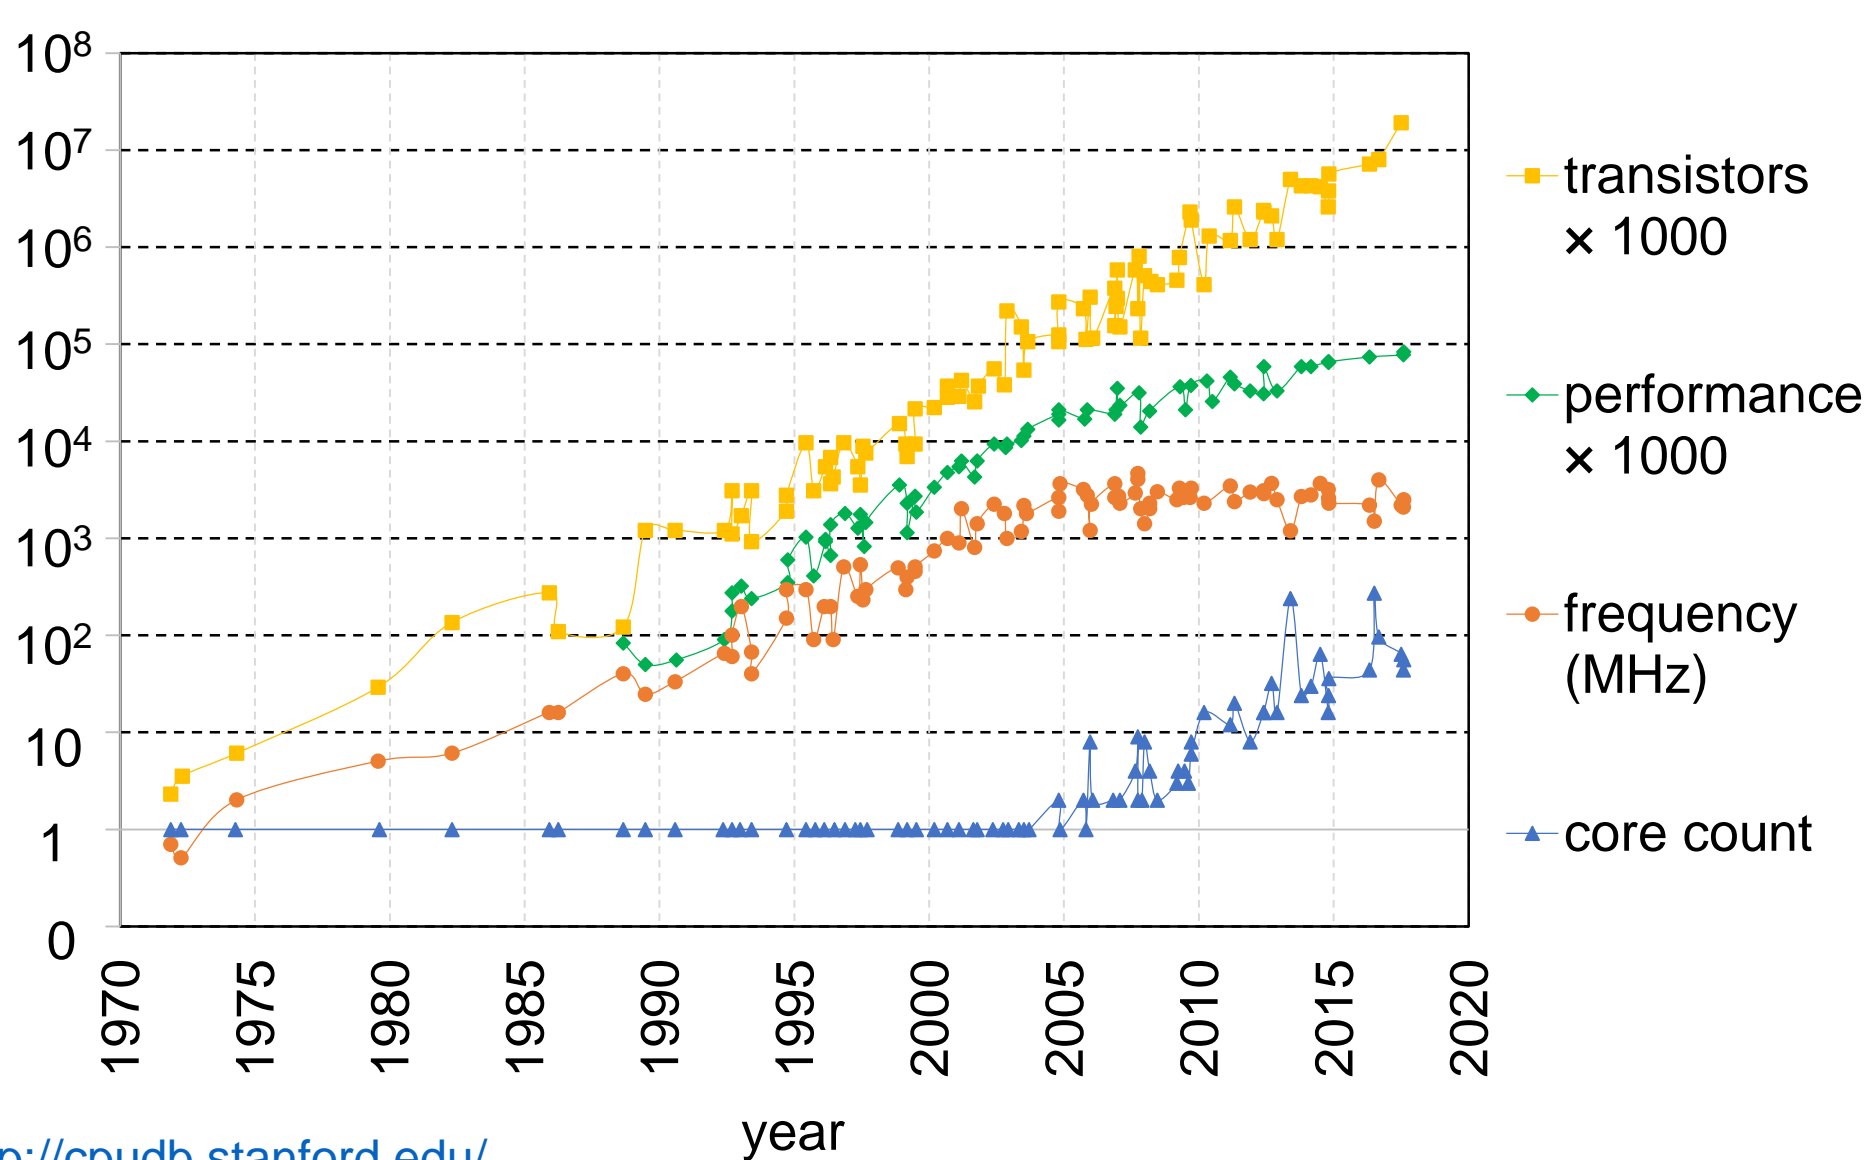
Consequences of Technology Scaling



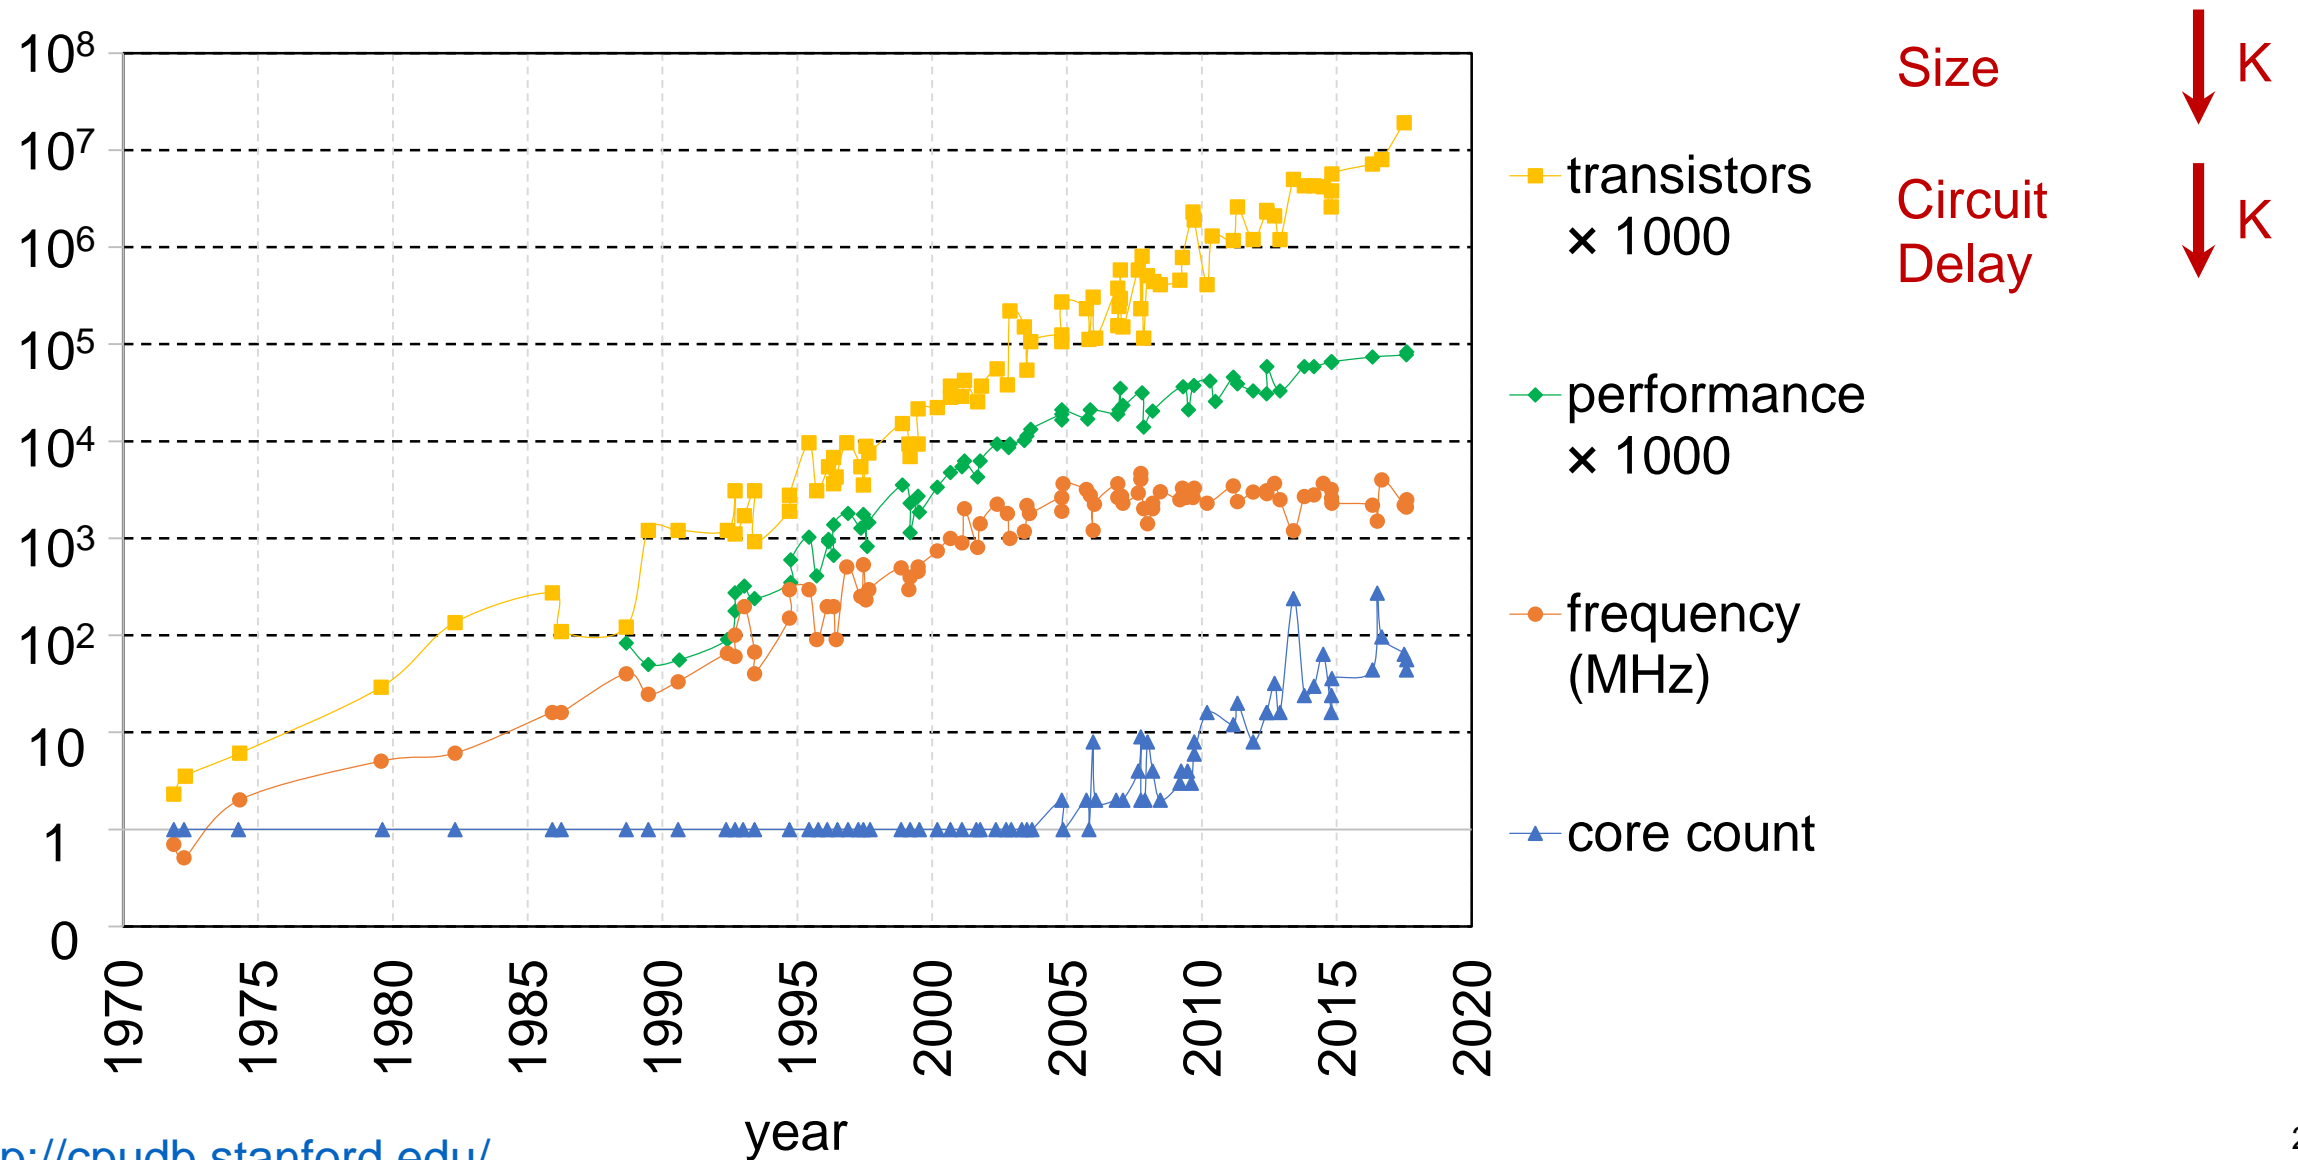
Consequences of Technology Scaling



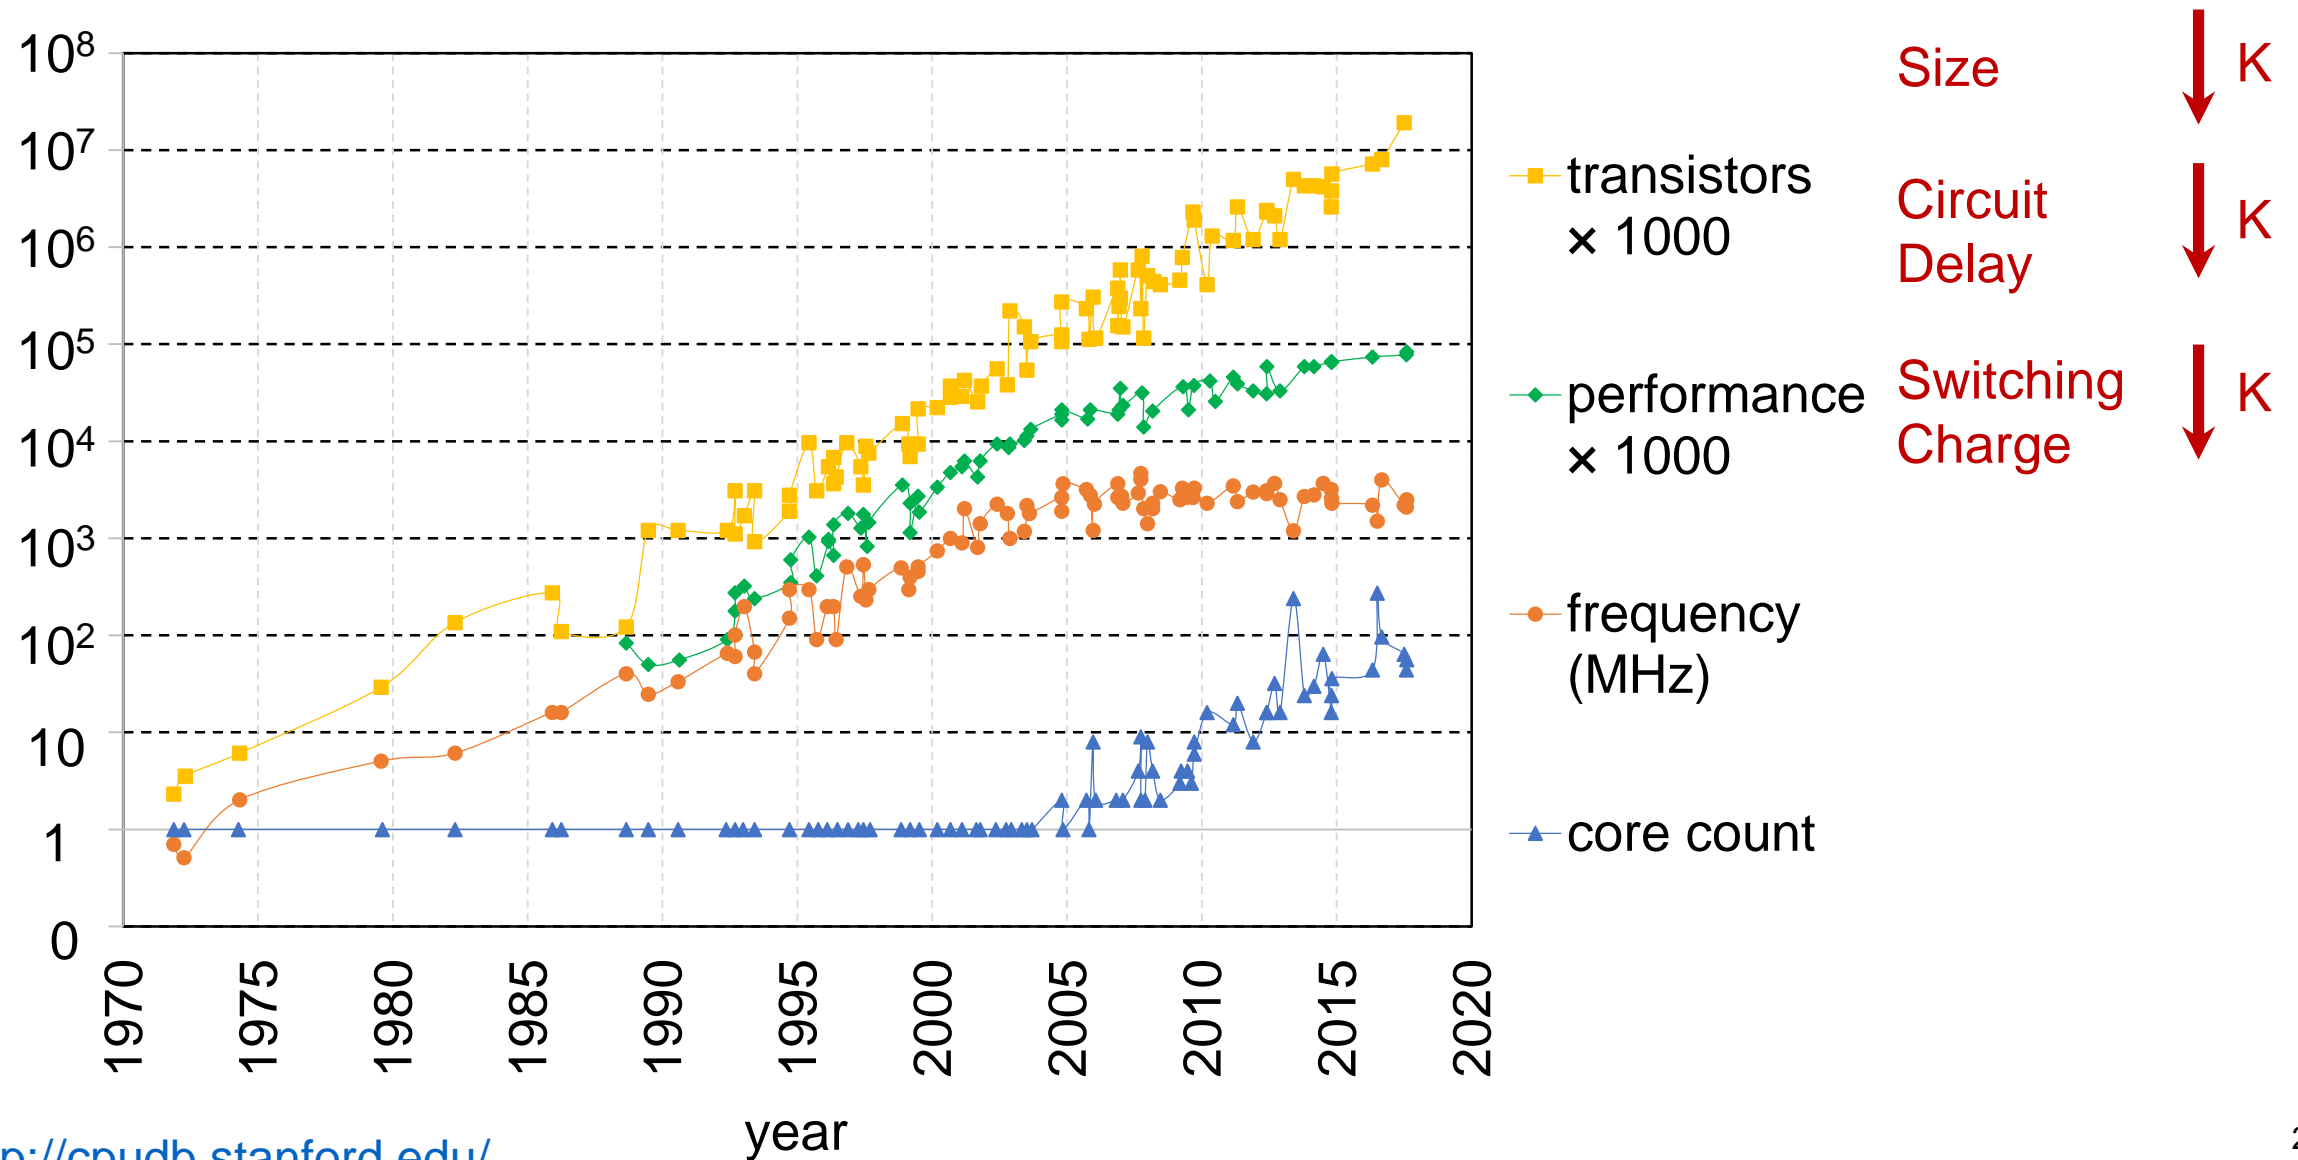
Consequences of Technology Scaling



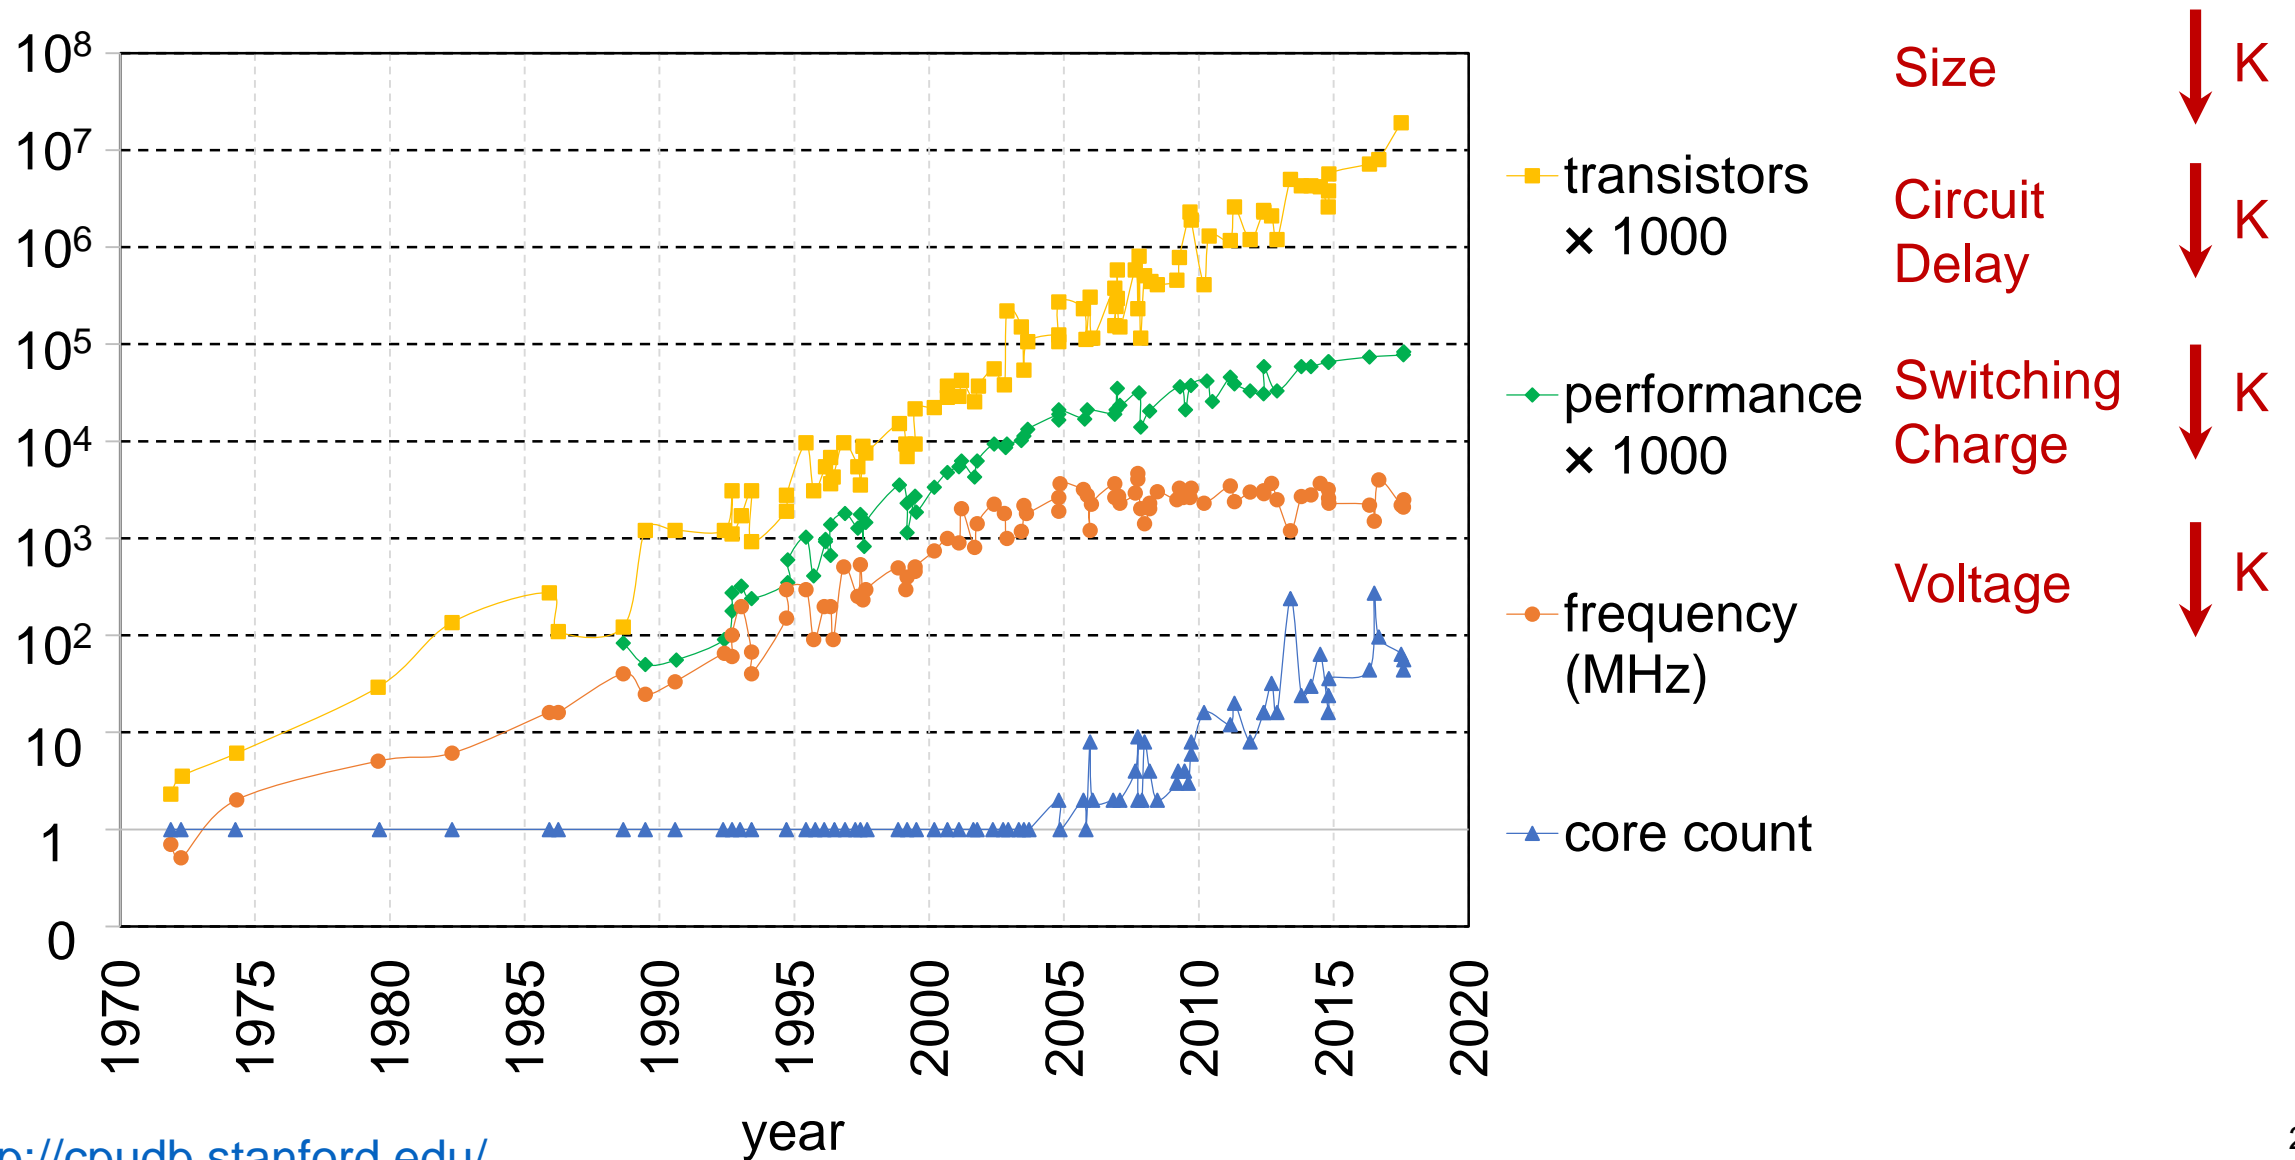
Consequences of Technology Scaling



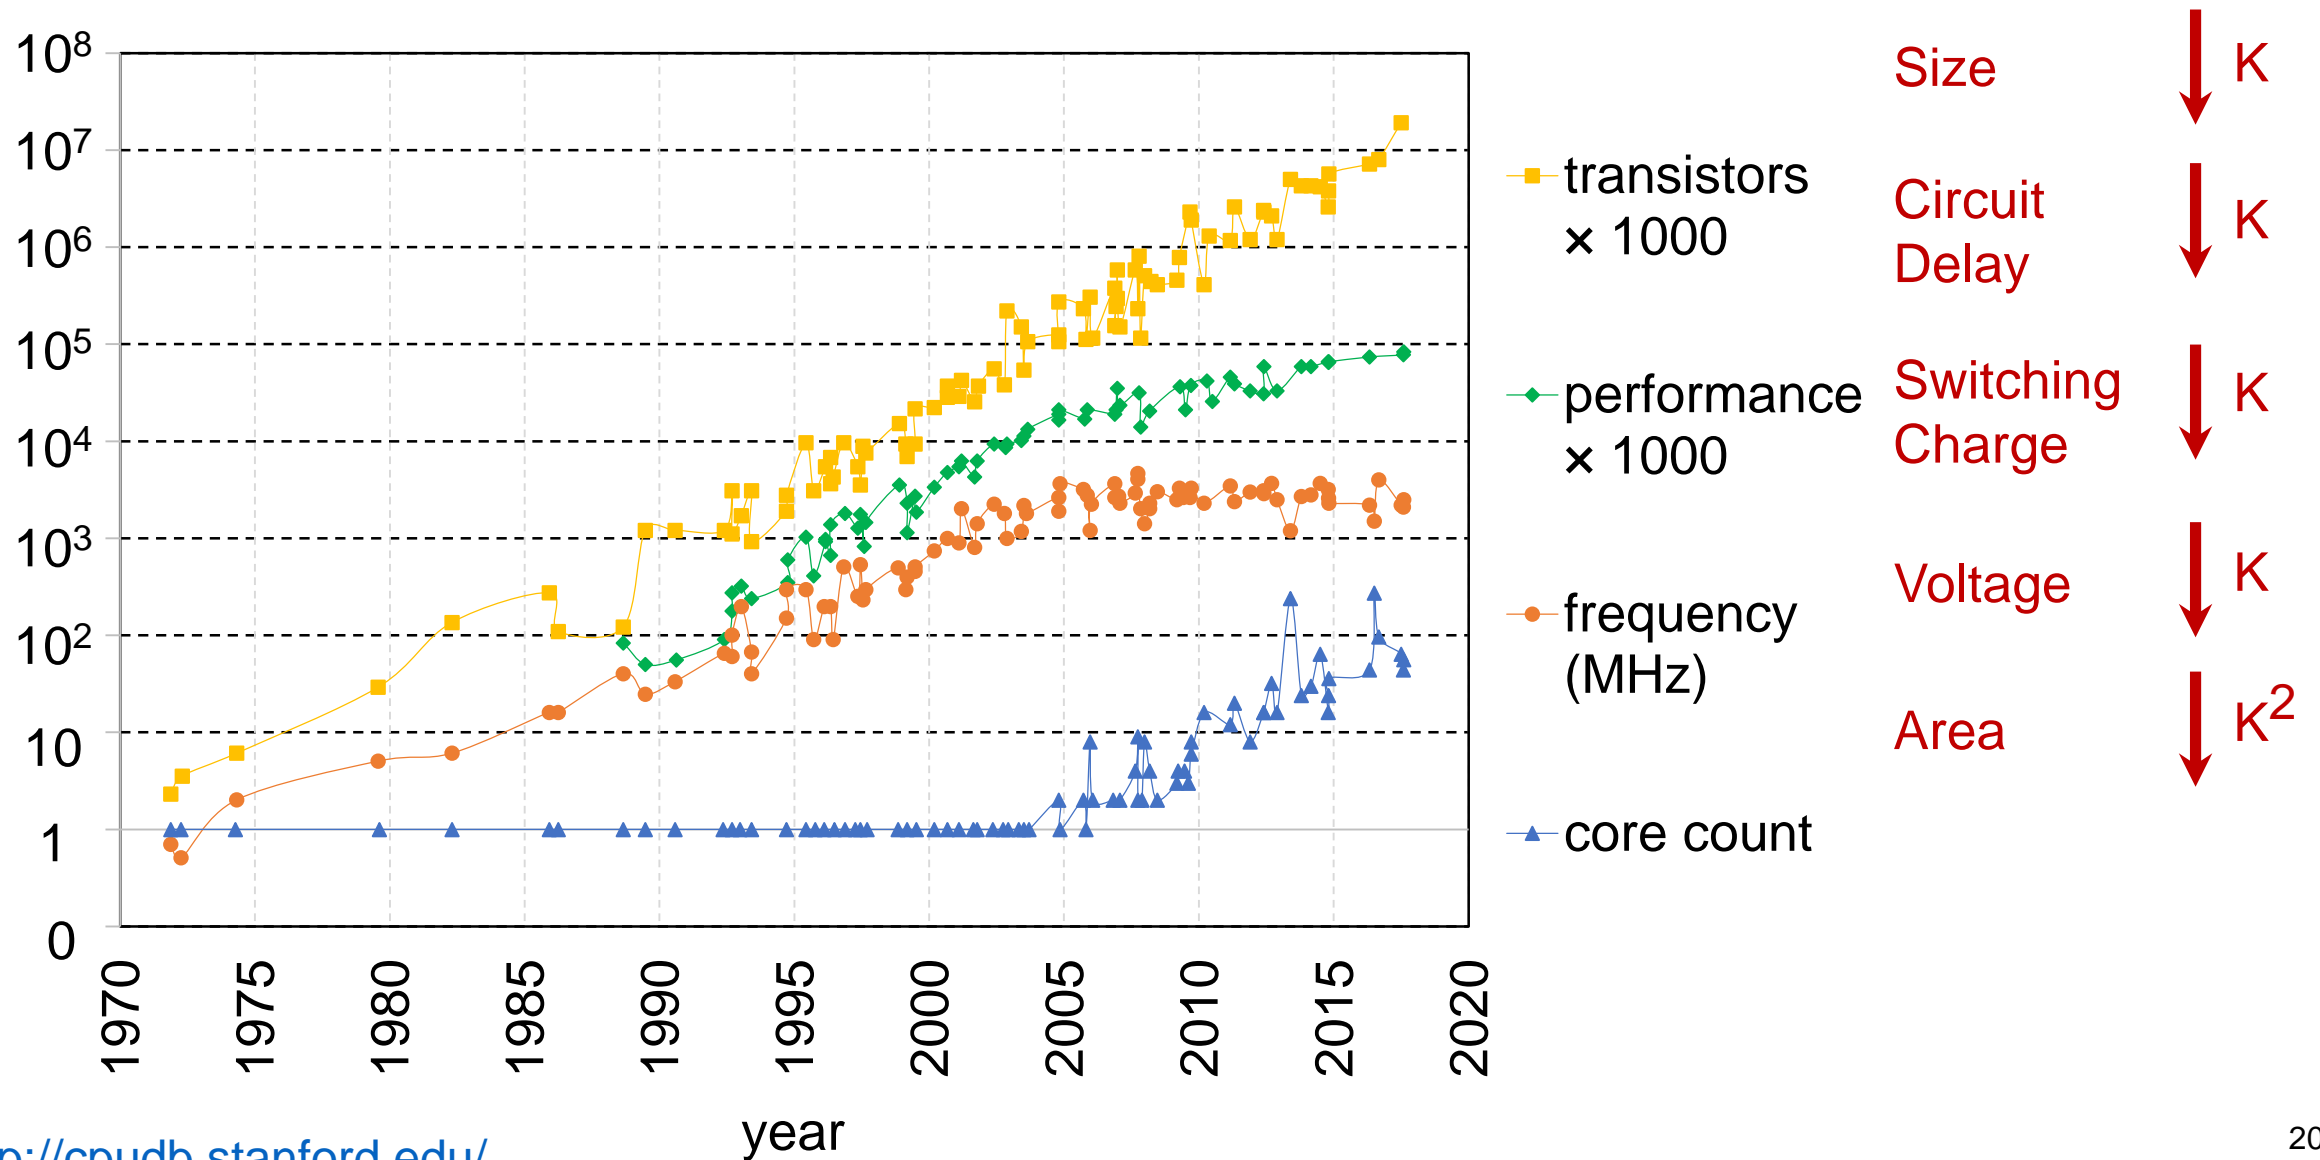
Consequences of Technology Scaling



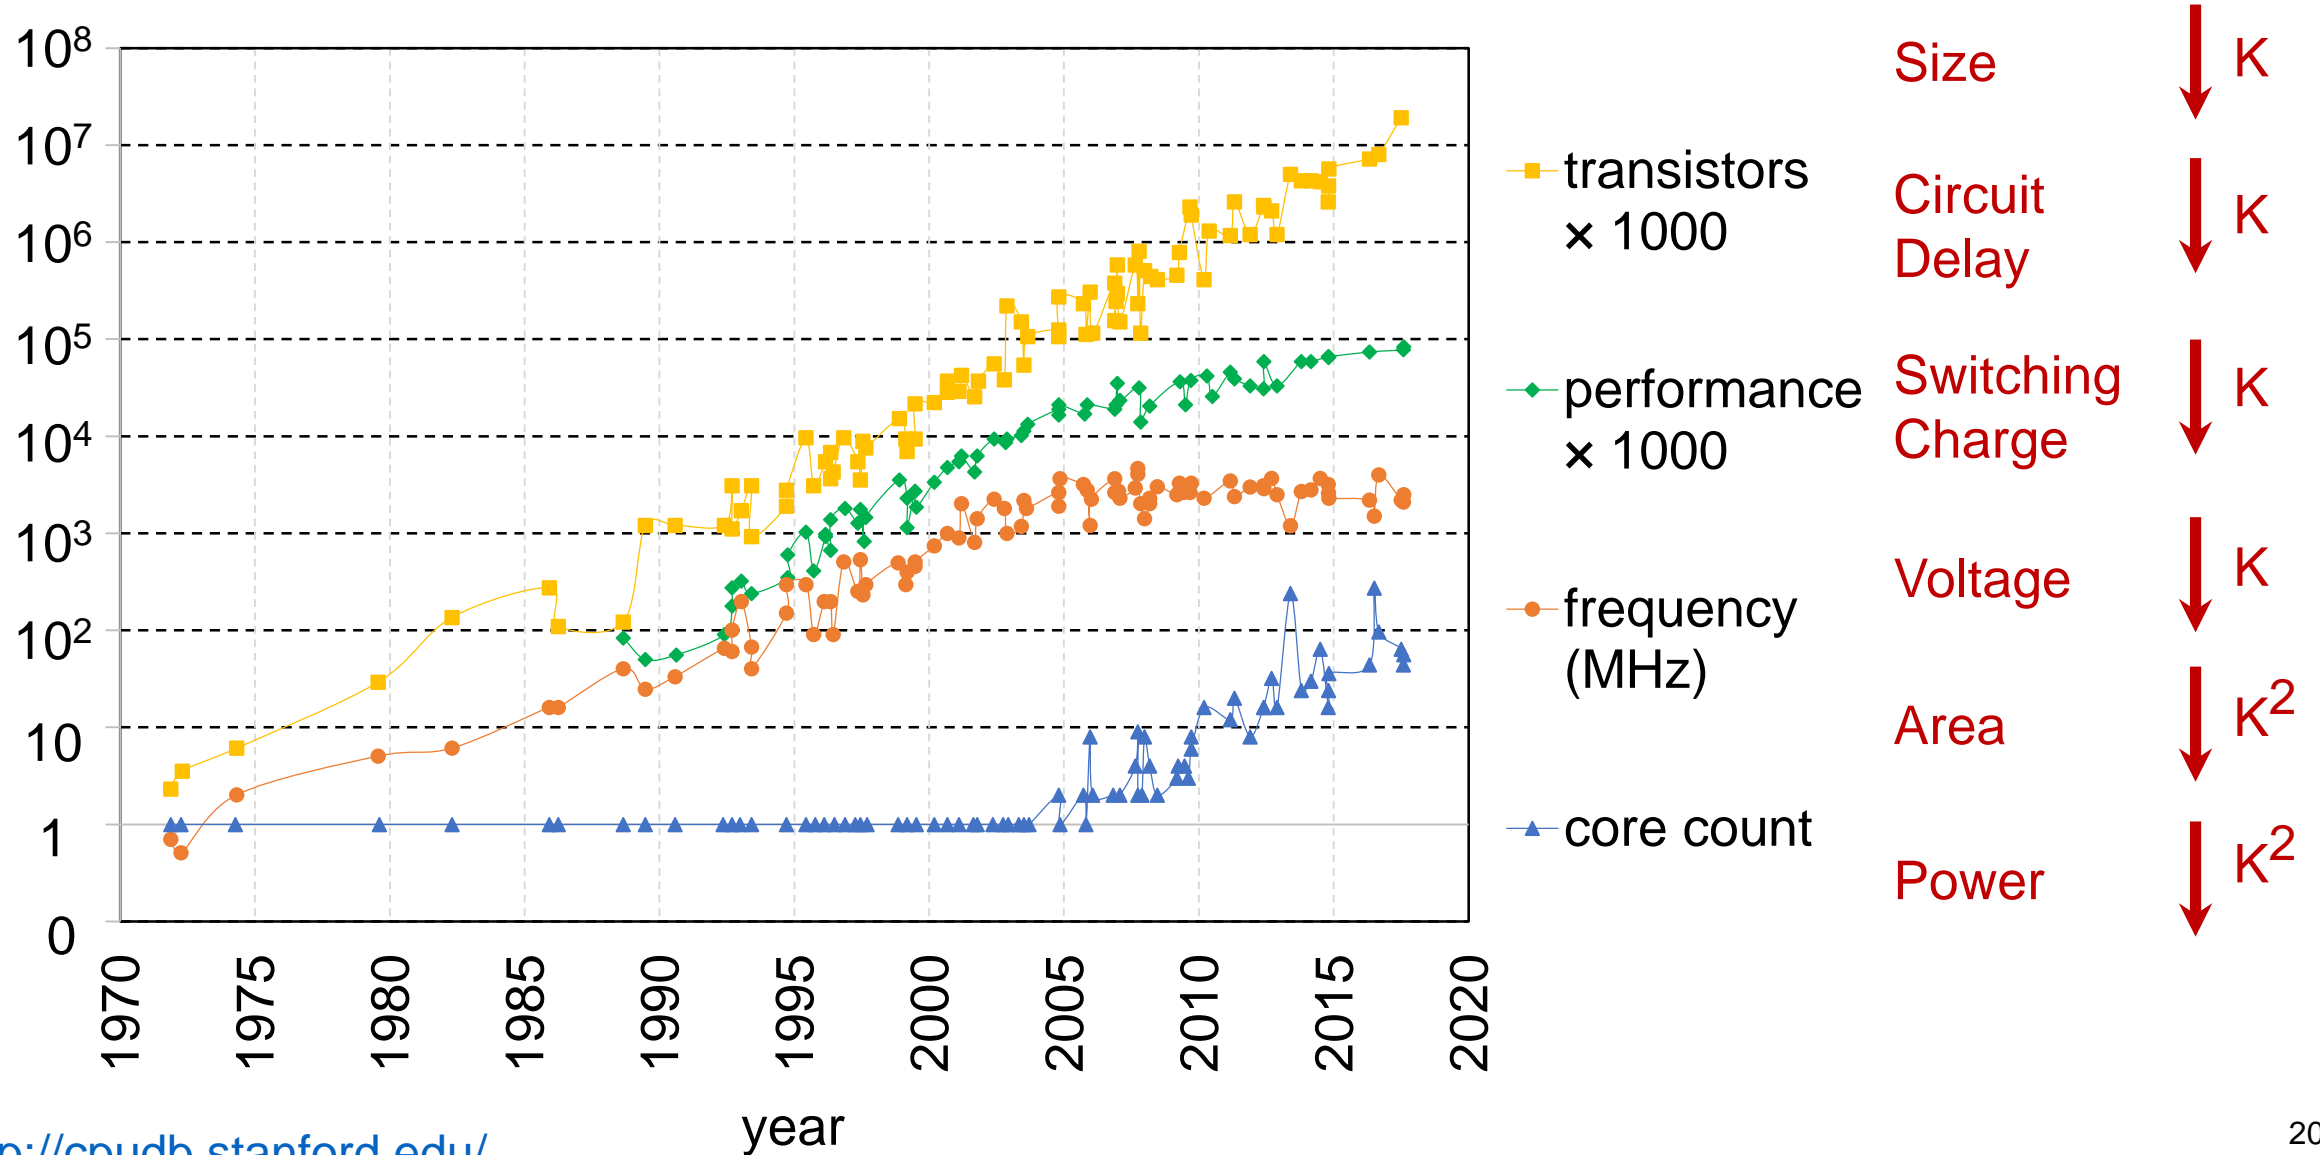
Consequences of Technology Scaling



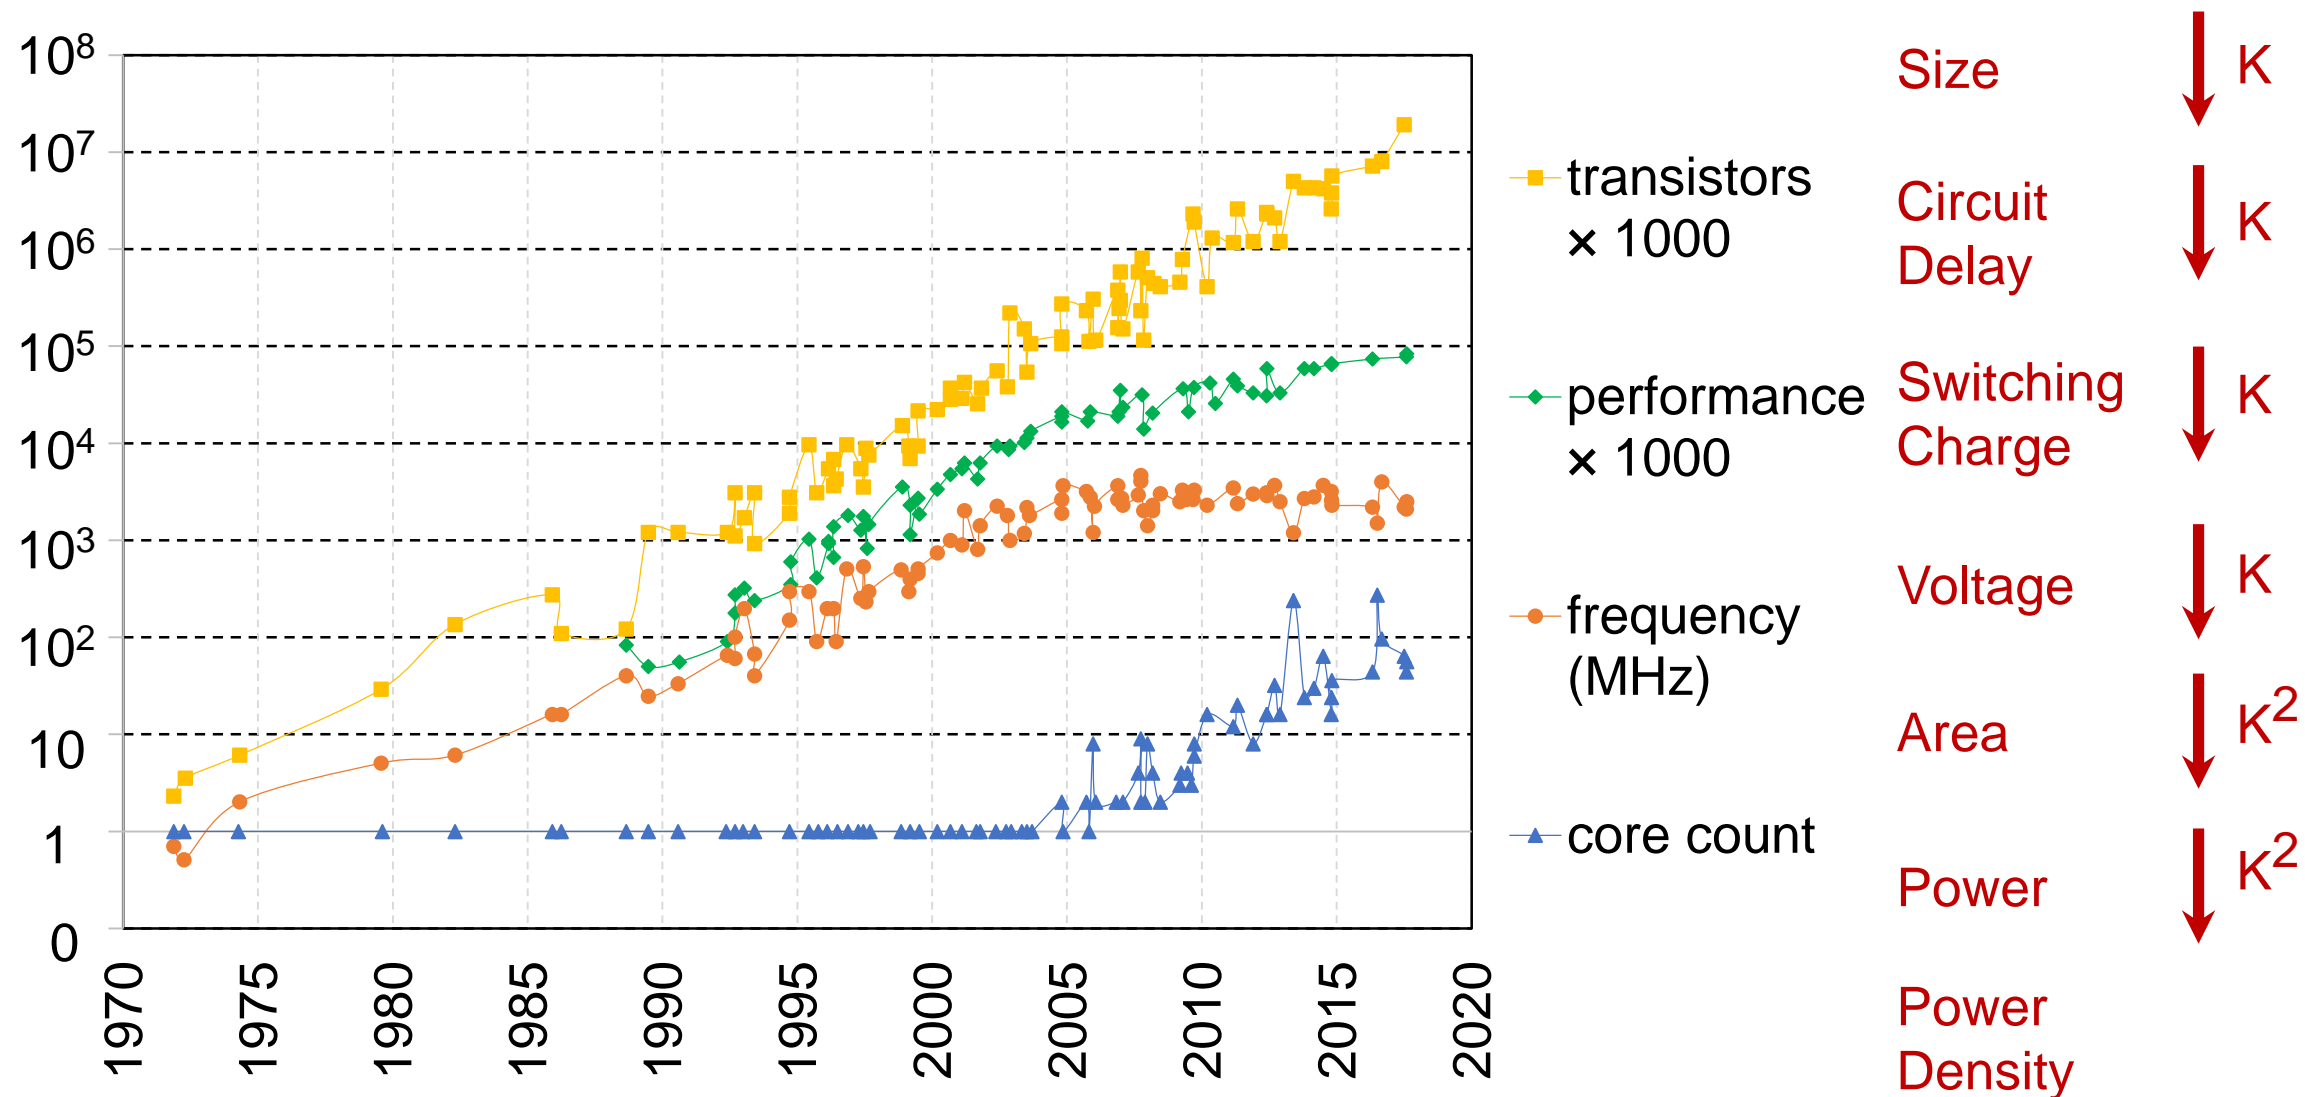
Consequences of Technology Scaling



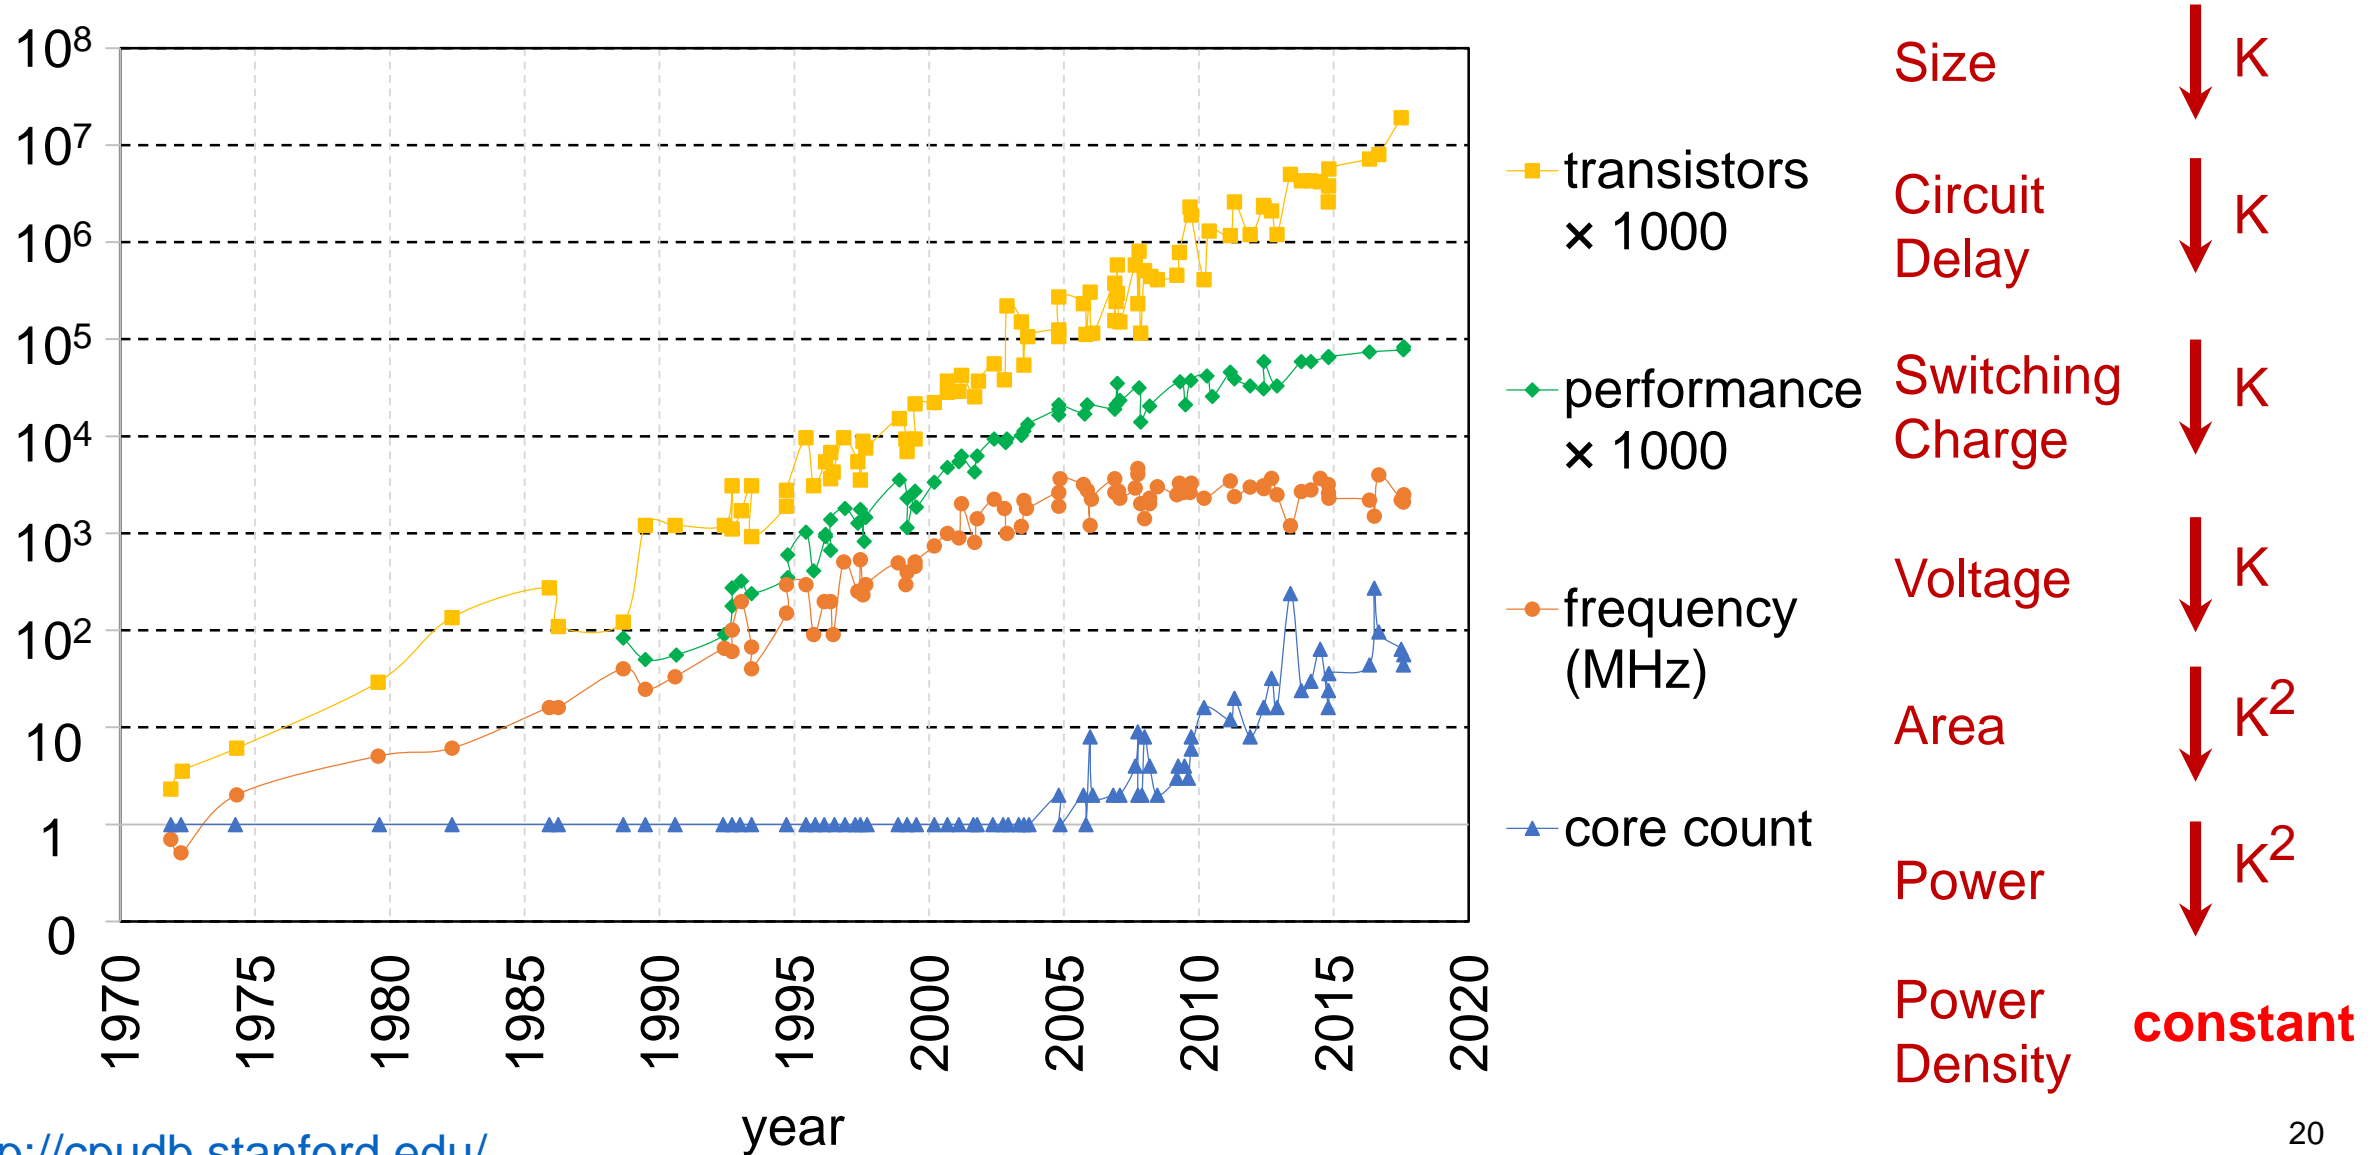
Consequences of Technology Scaling



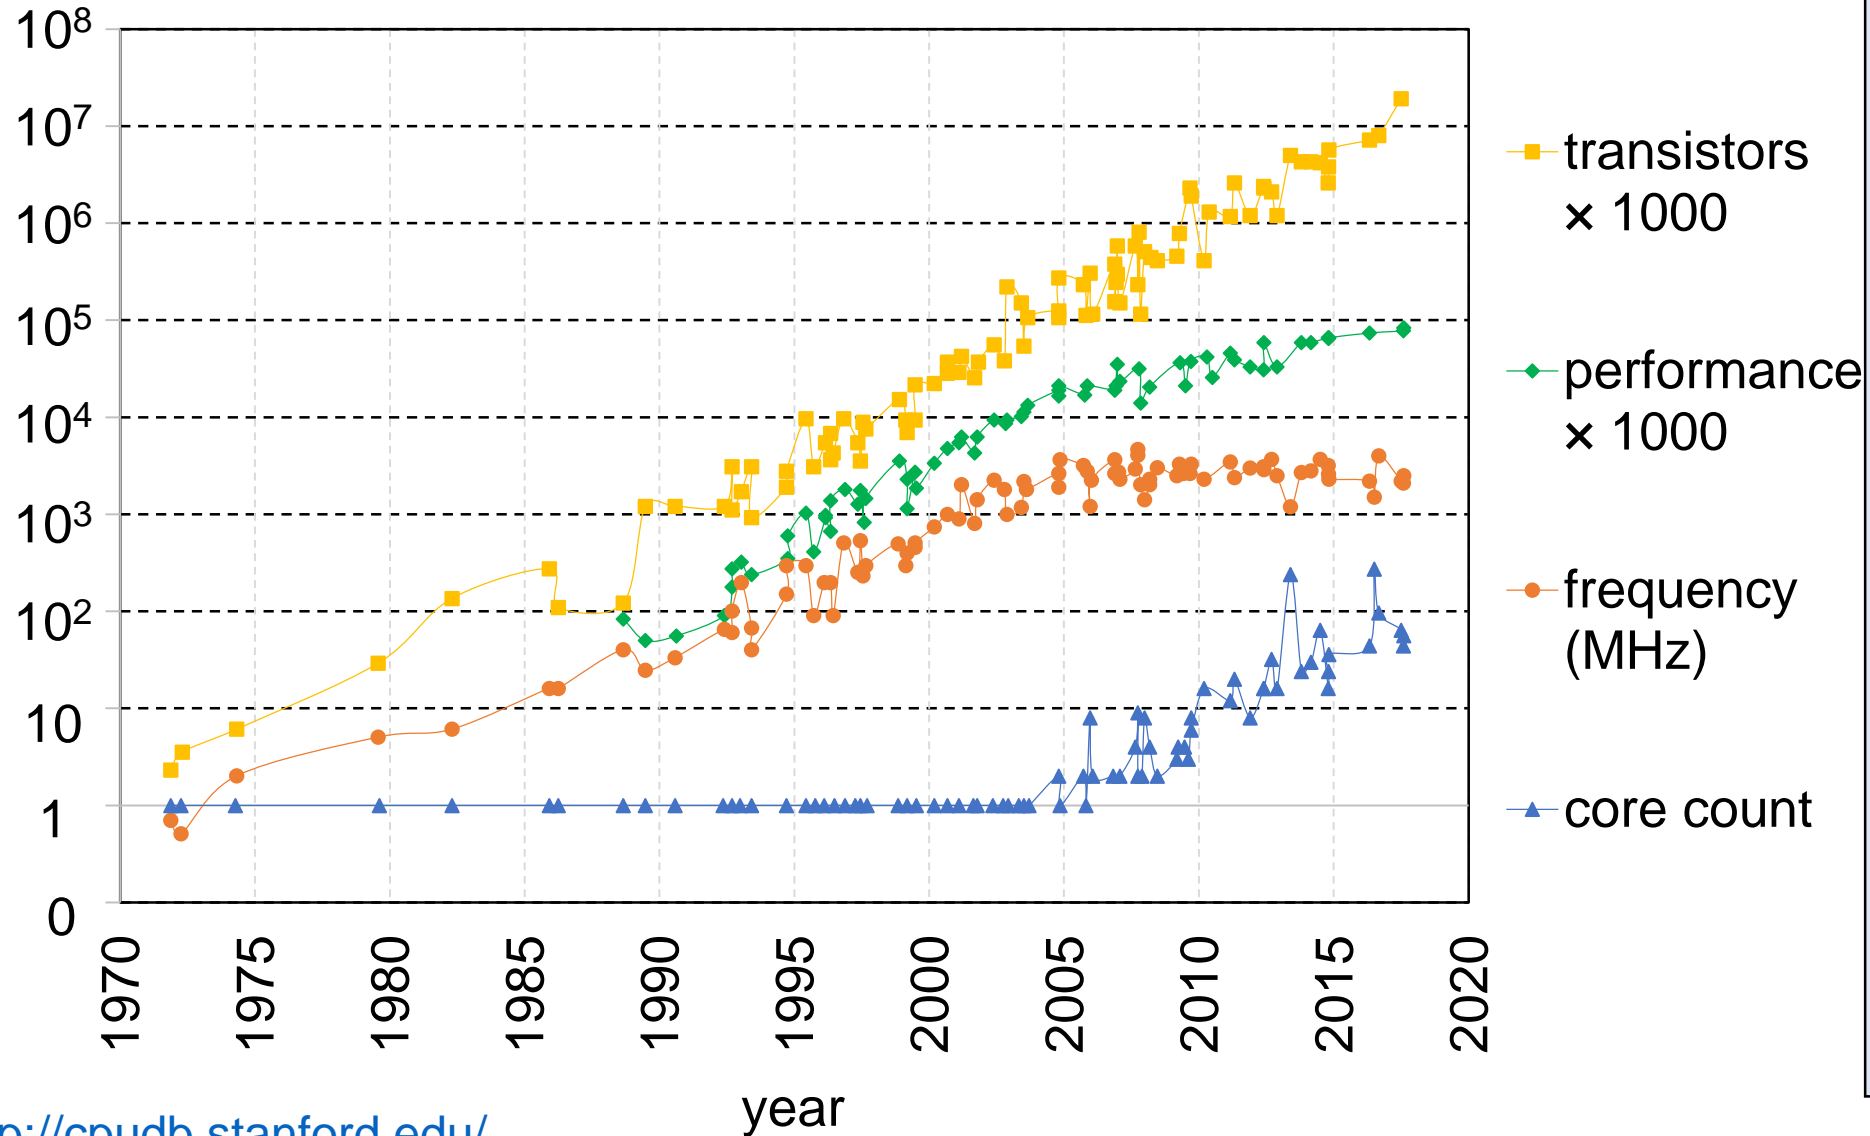
Consequences of Technology Scaling



Consequences of Technology Scaling



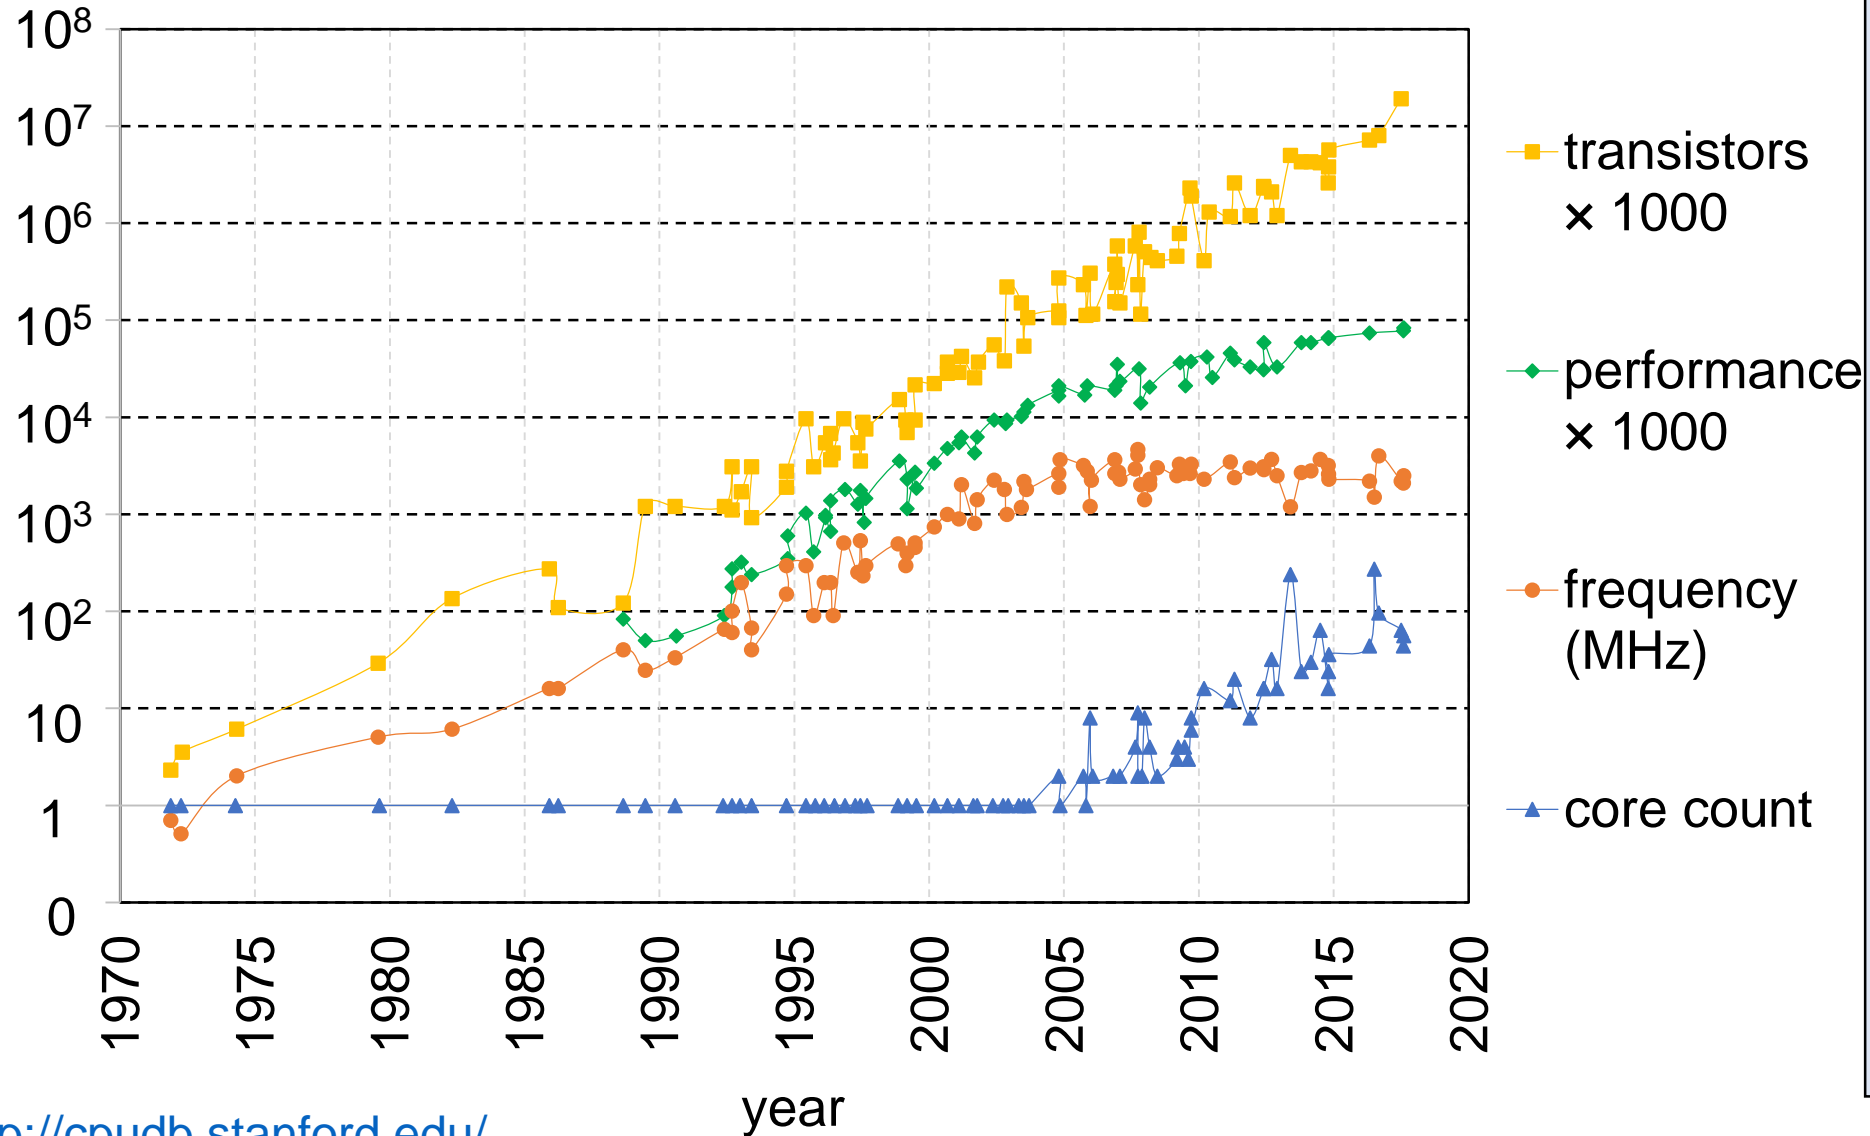
Consequences of Technology Scaling



Dennard Scaling

Size	$\downarrow K$
Circuit Delay	$\downarrow K$
Switching Charge	$\downarrow K$
Voltage	$\downarrow K$
Area	$\downarrow K^2$
Power	$\downarrow K^2$
Power Density	constant

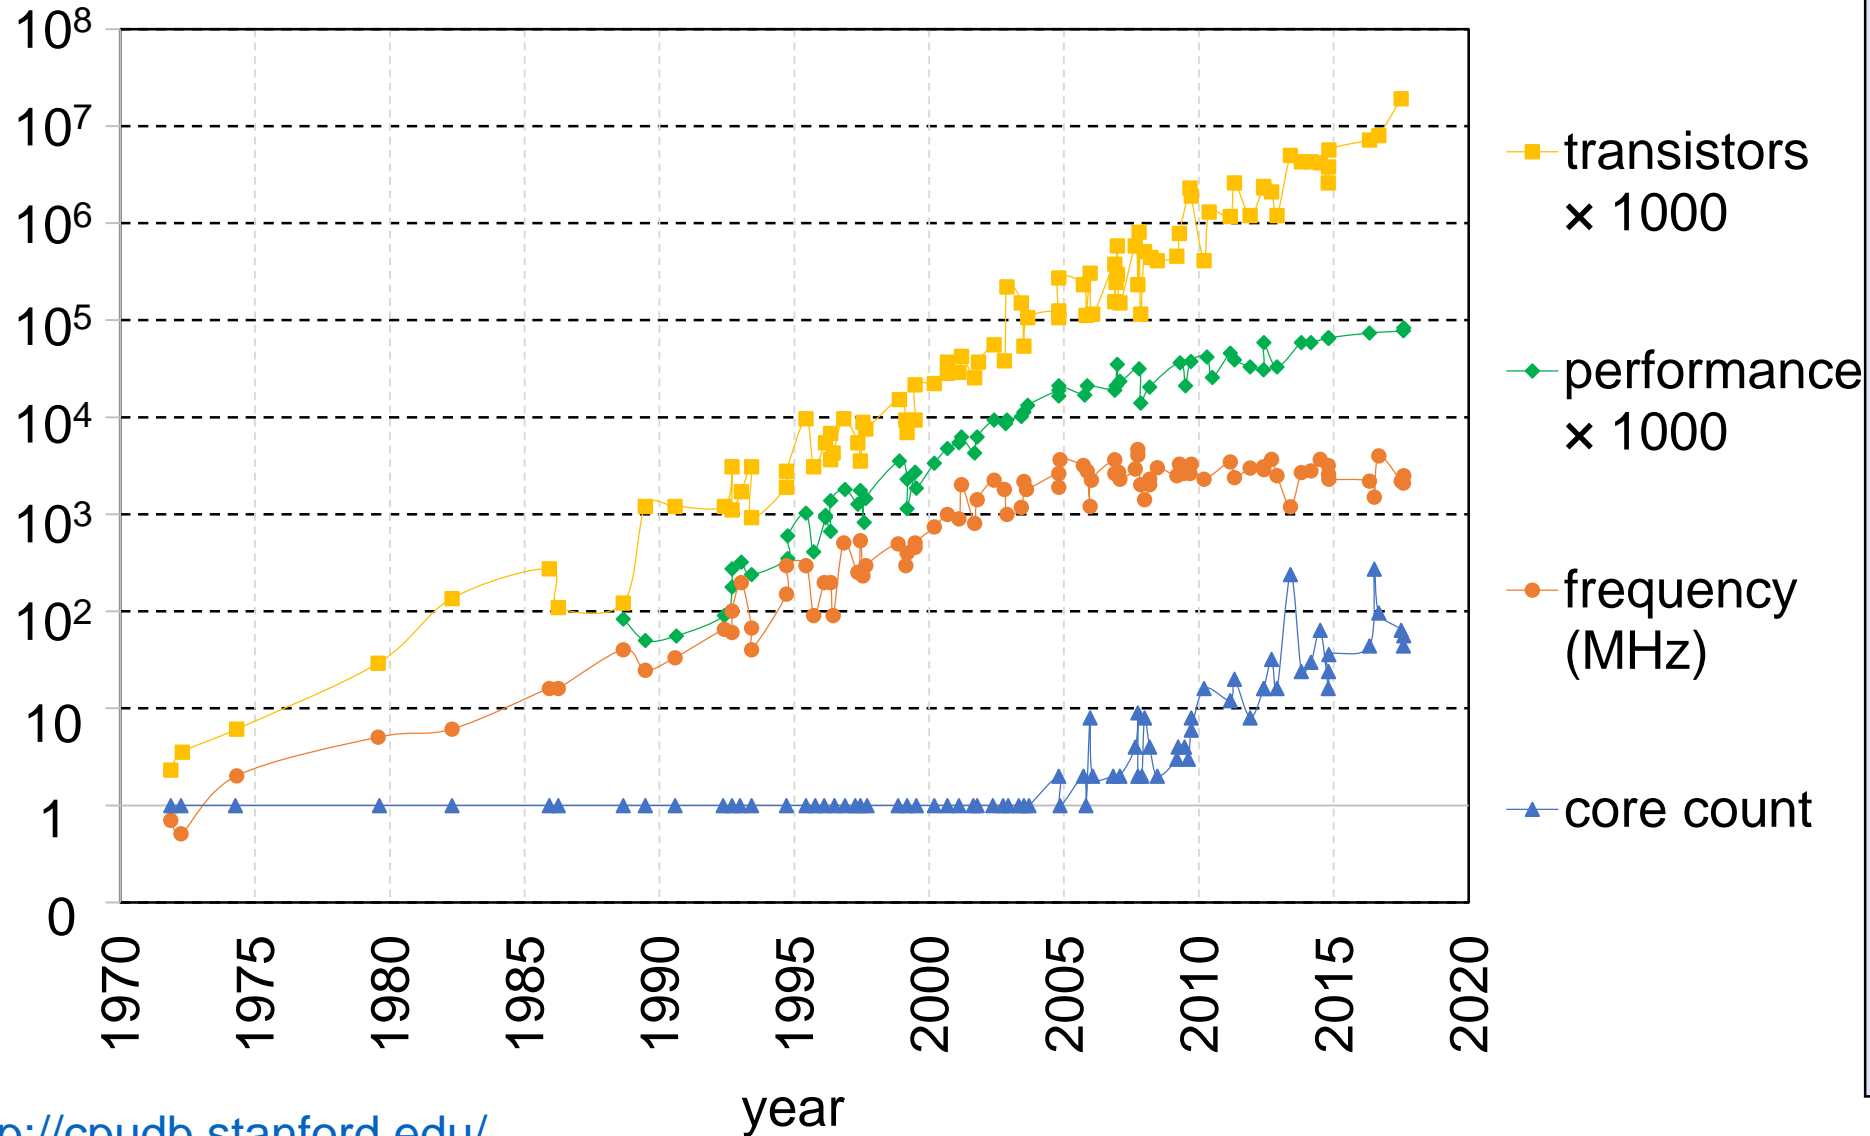
Consequences of Technology Scaling



Dennard Scaling

Size	$\downarrow K$
Circuit Delay	$\downarrow K$
Switching Charge	$\downarrow K$
Voltage	$\downarrow K$
Area	$\downarrow K^2$
Power	$\downarrow K^2$
Power Density	constant

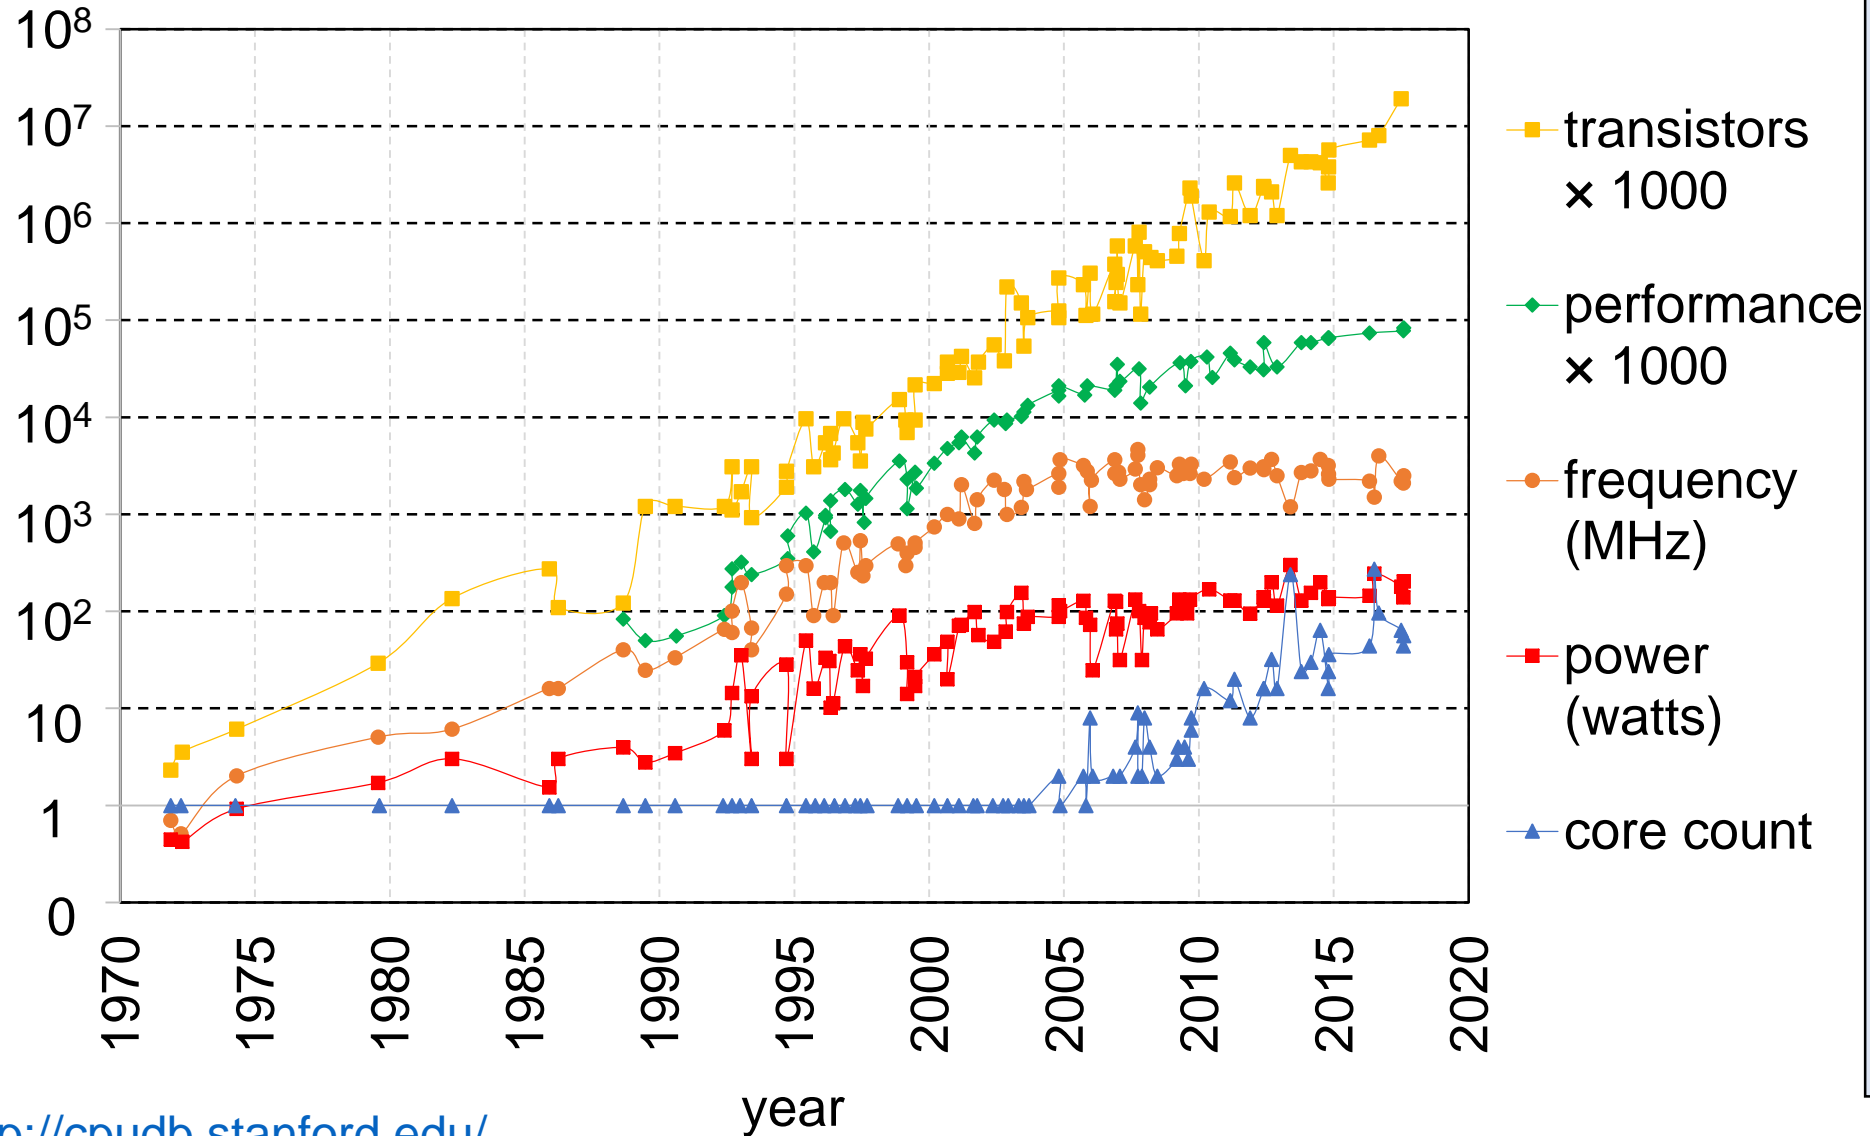
Consequences of Technology Scaling



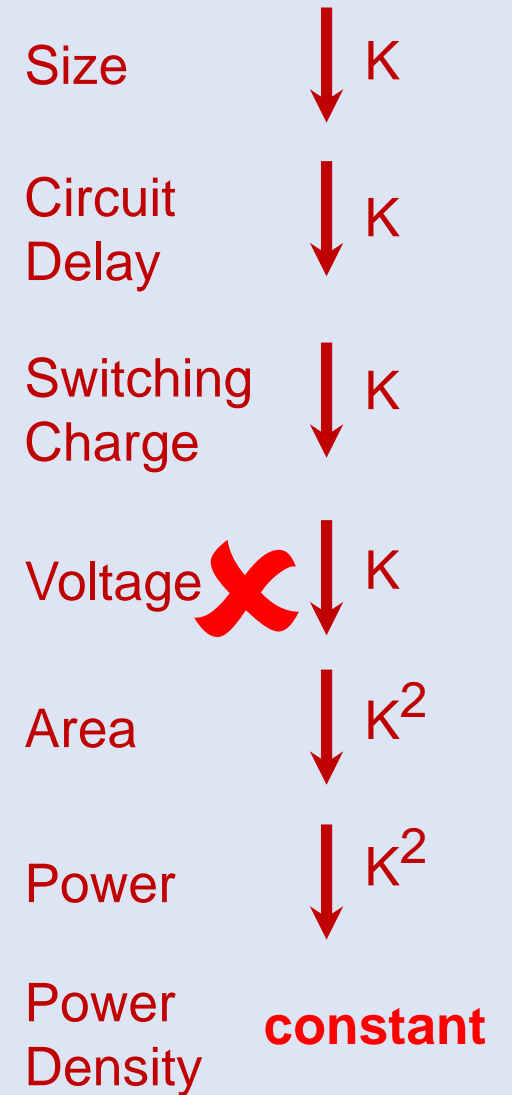
Dennard Scaling

Size	$\downarrow K$
Circuit Delay	$\downarrow K$
Switching Charge	$\downarrow K$
Voltage	$\downarrow K$
Area	$\downarrow K^2$
Power	$\downarrow K^2$
Power Density	constant

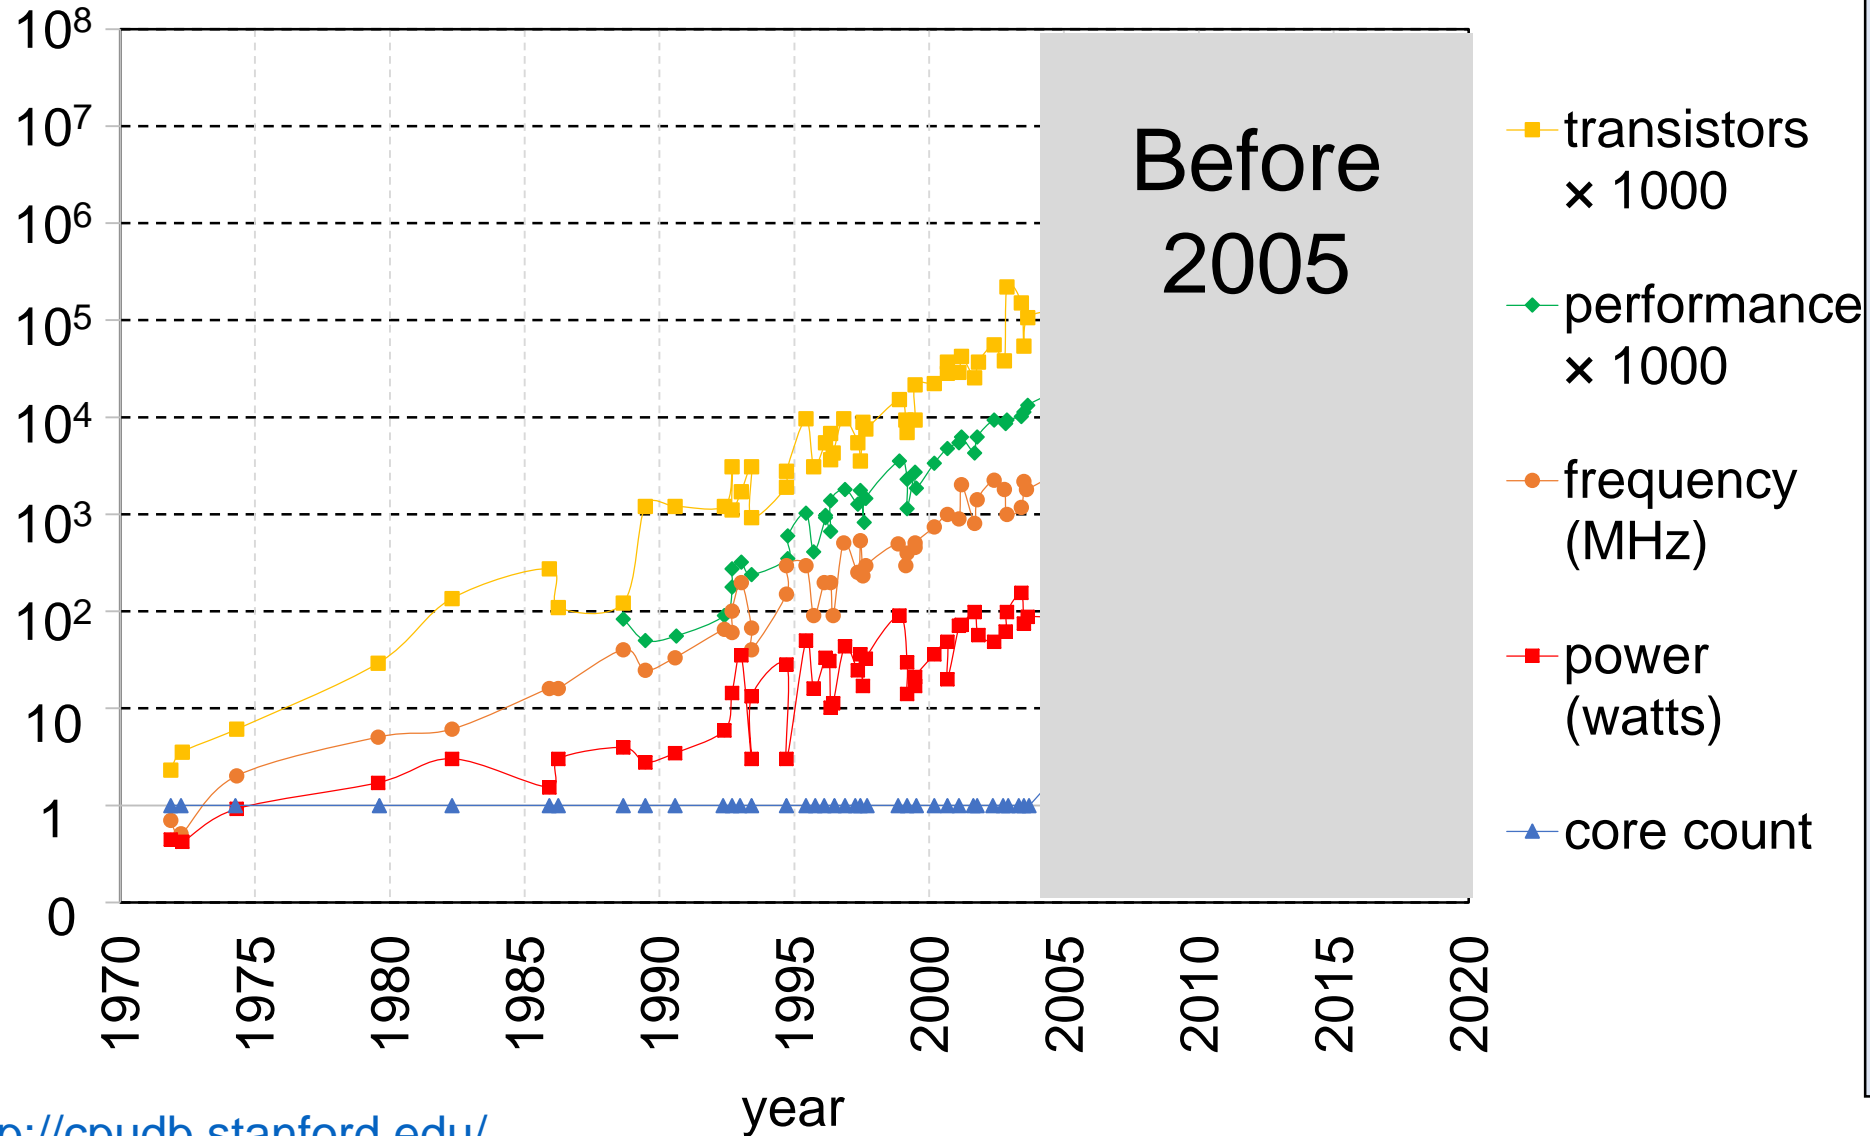
Consequences of Technology Scaling



Dennard Scaling



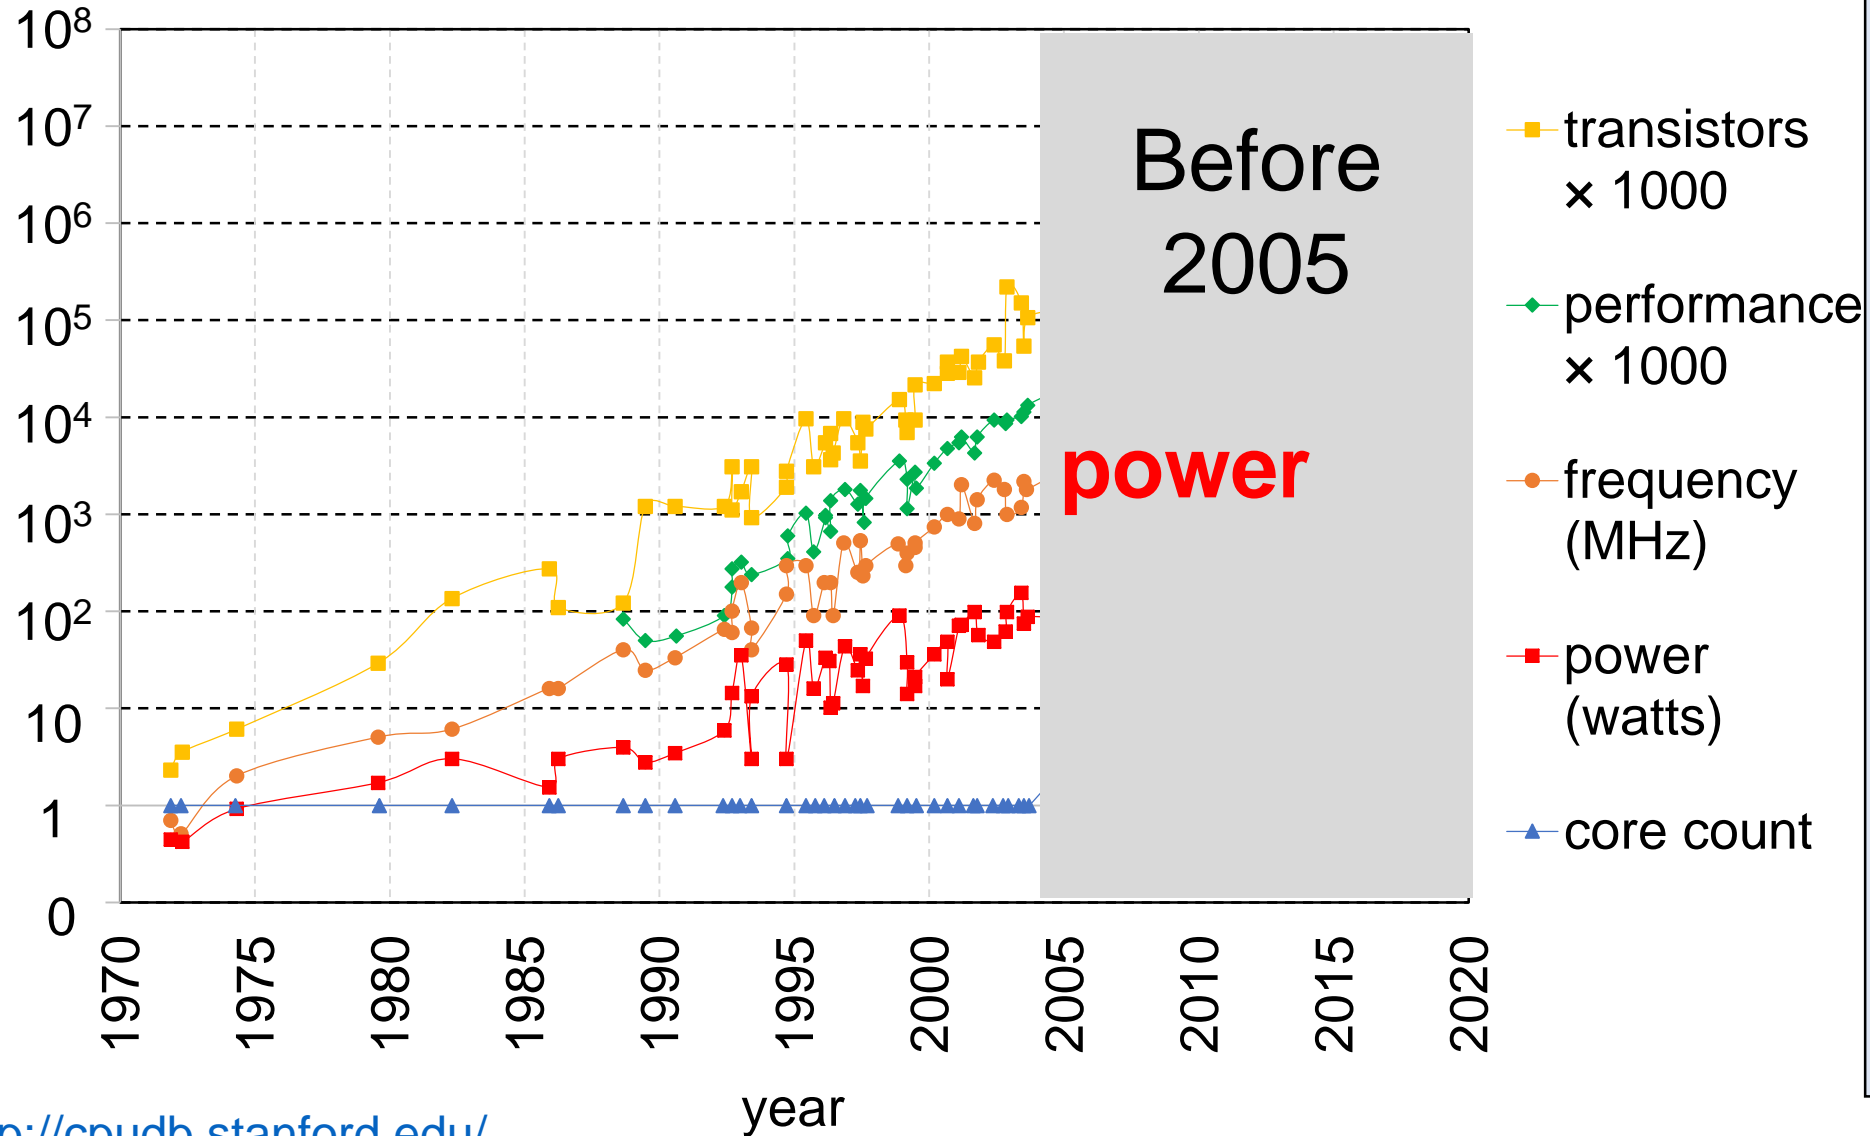
Consequences of Technology Scaling



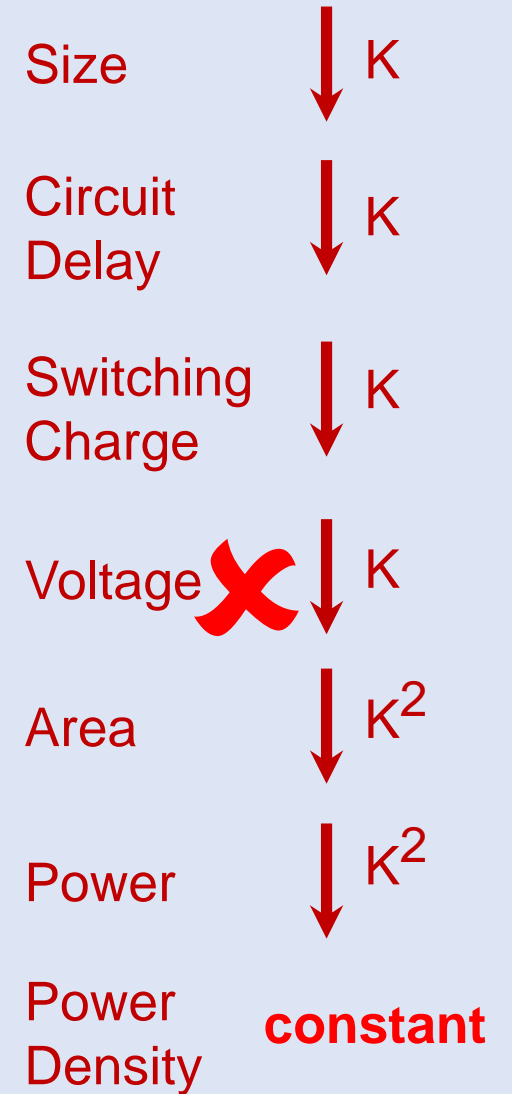
Dennard Scaling

Size	$\downarrow K$
Circuit Delay	$\downarrow K$
Switching Charge	$\downarrow K$
Voltage	$\downarrow K$
Area	$\downarrow K^2$
Power	$\downarrow K^2$
Power Density	constant

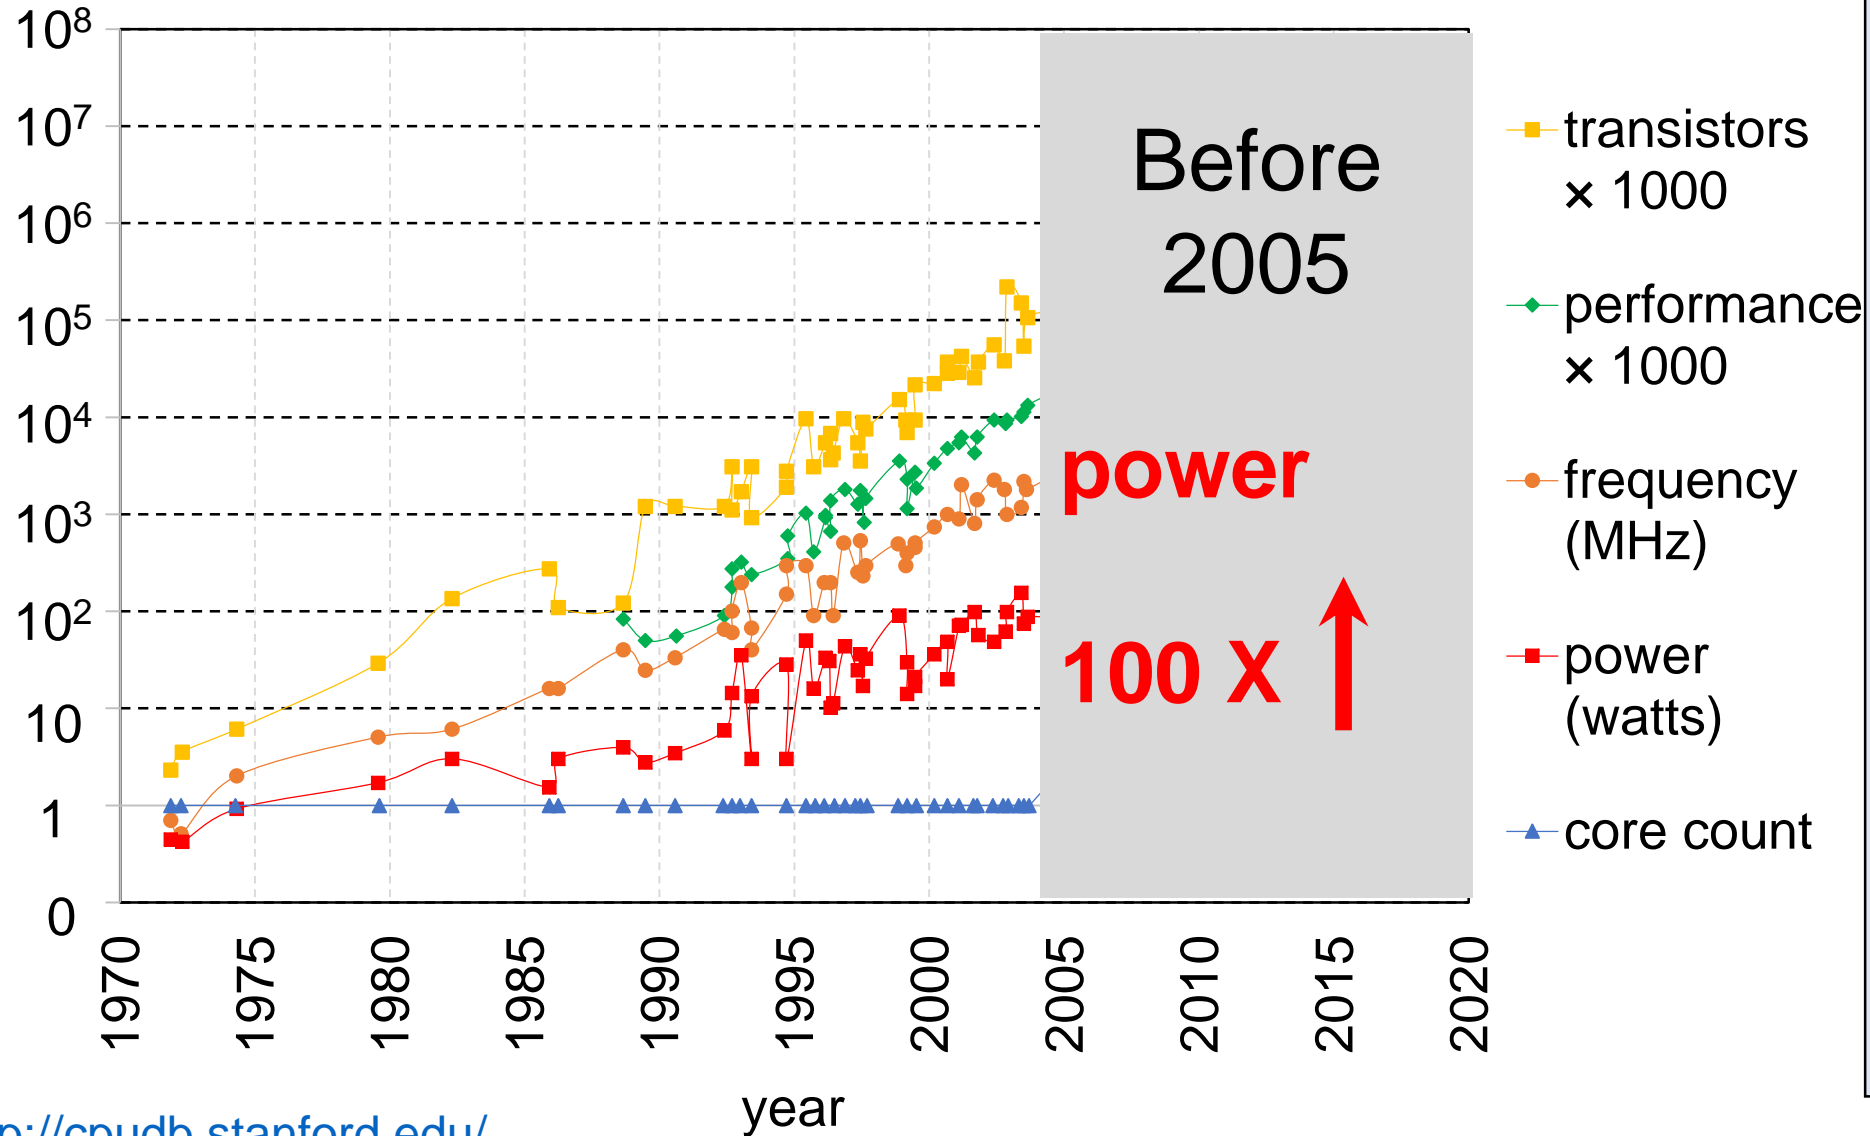
Consequences of Technology Scaling



Dennard Scaling



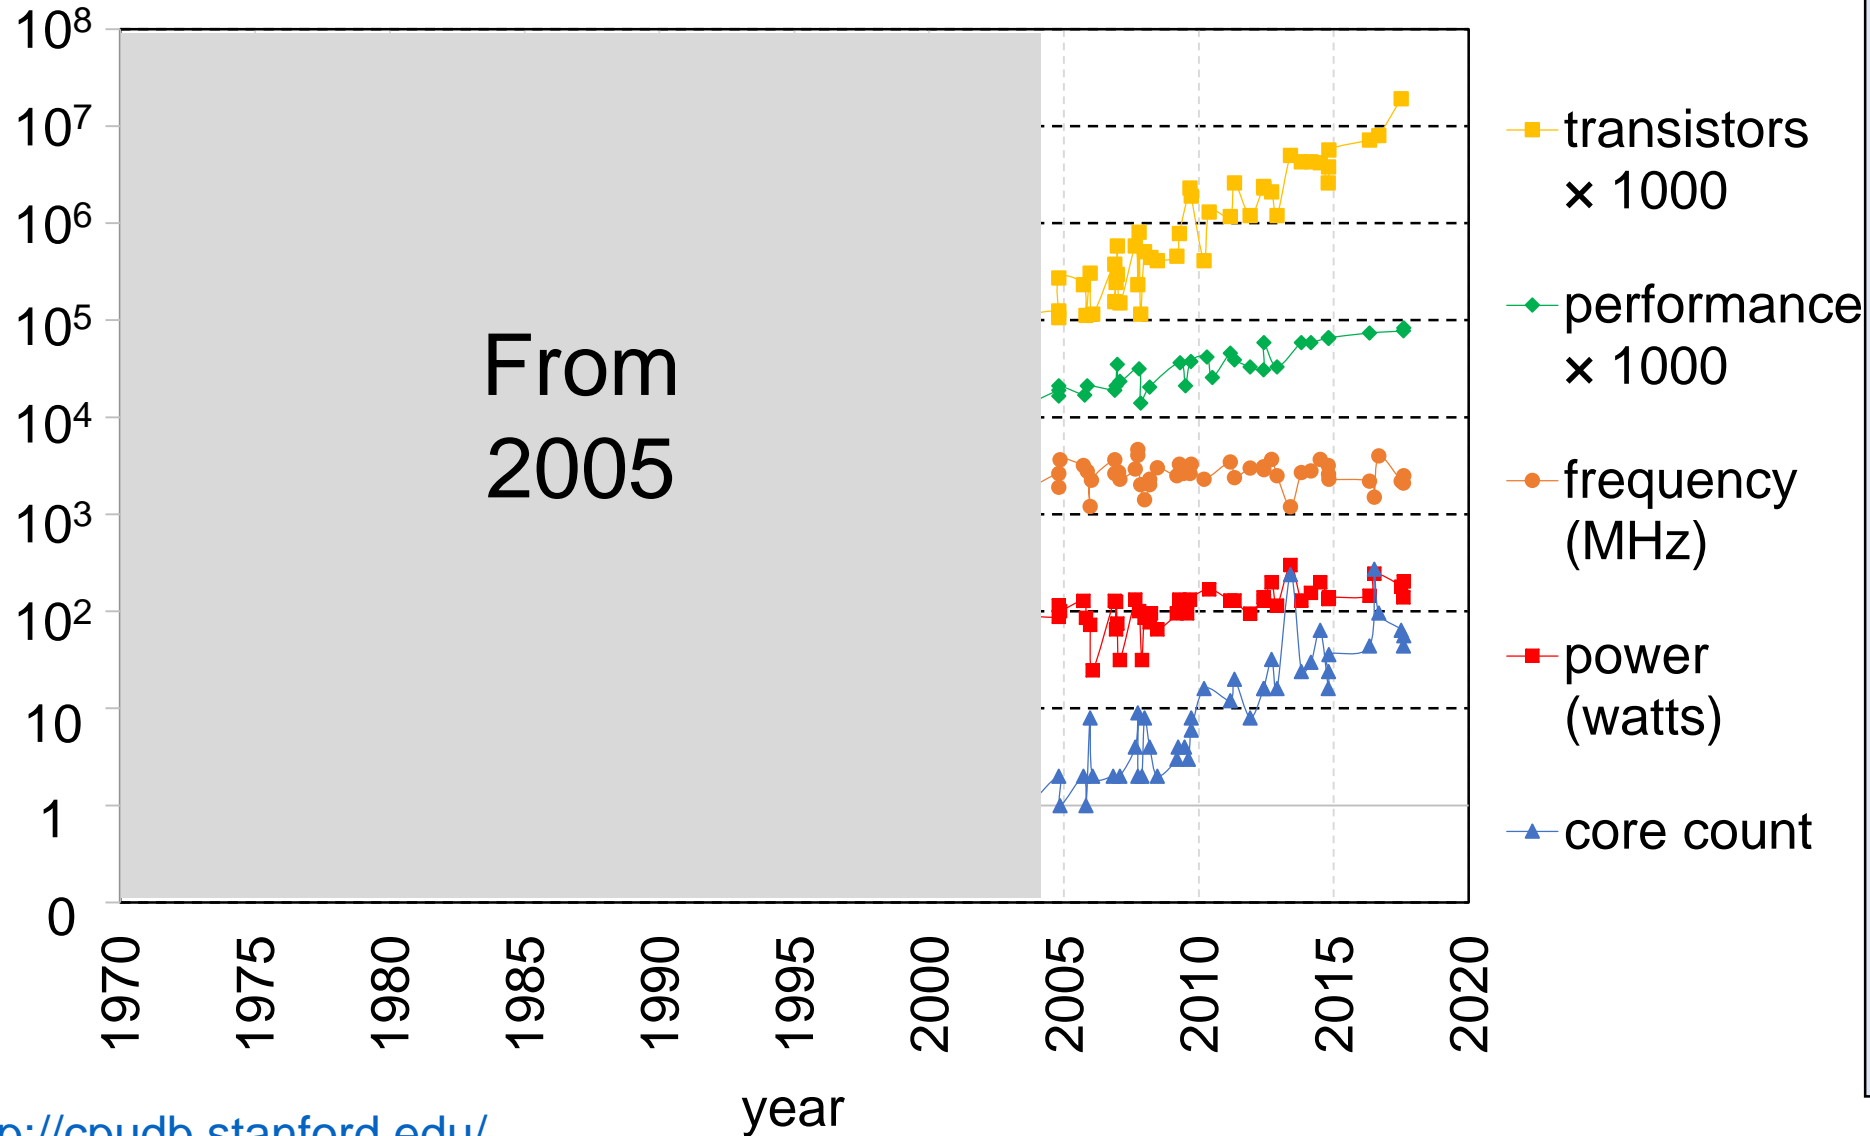
Consequences of Technology Scaling



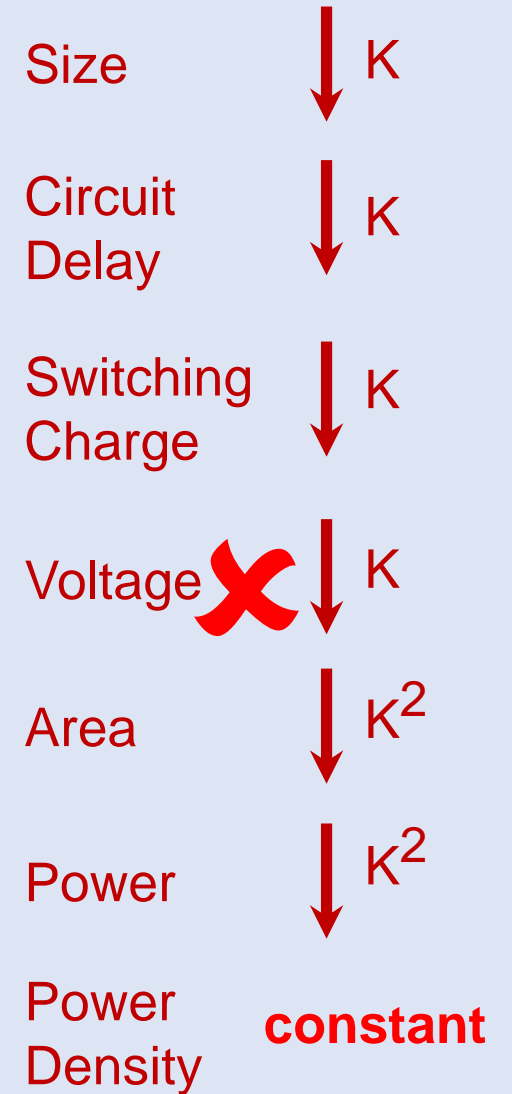
Dennard Scaling

Size	↓ K
Circuit Delay	↓ K
Switching Charge	↓ K
Voltage	↓ K
Area	↓ K ²
Power	↓ K ²
Power Density	constant

Consequences of Technology Scaling



Dennard Scaling

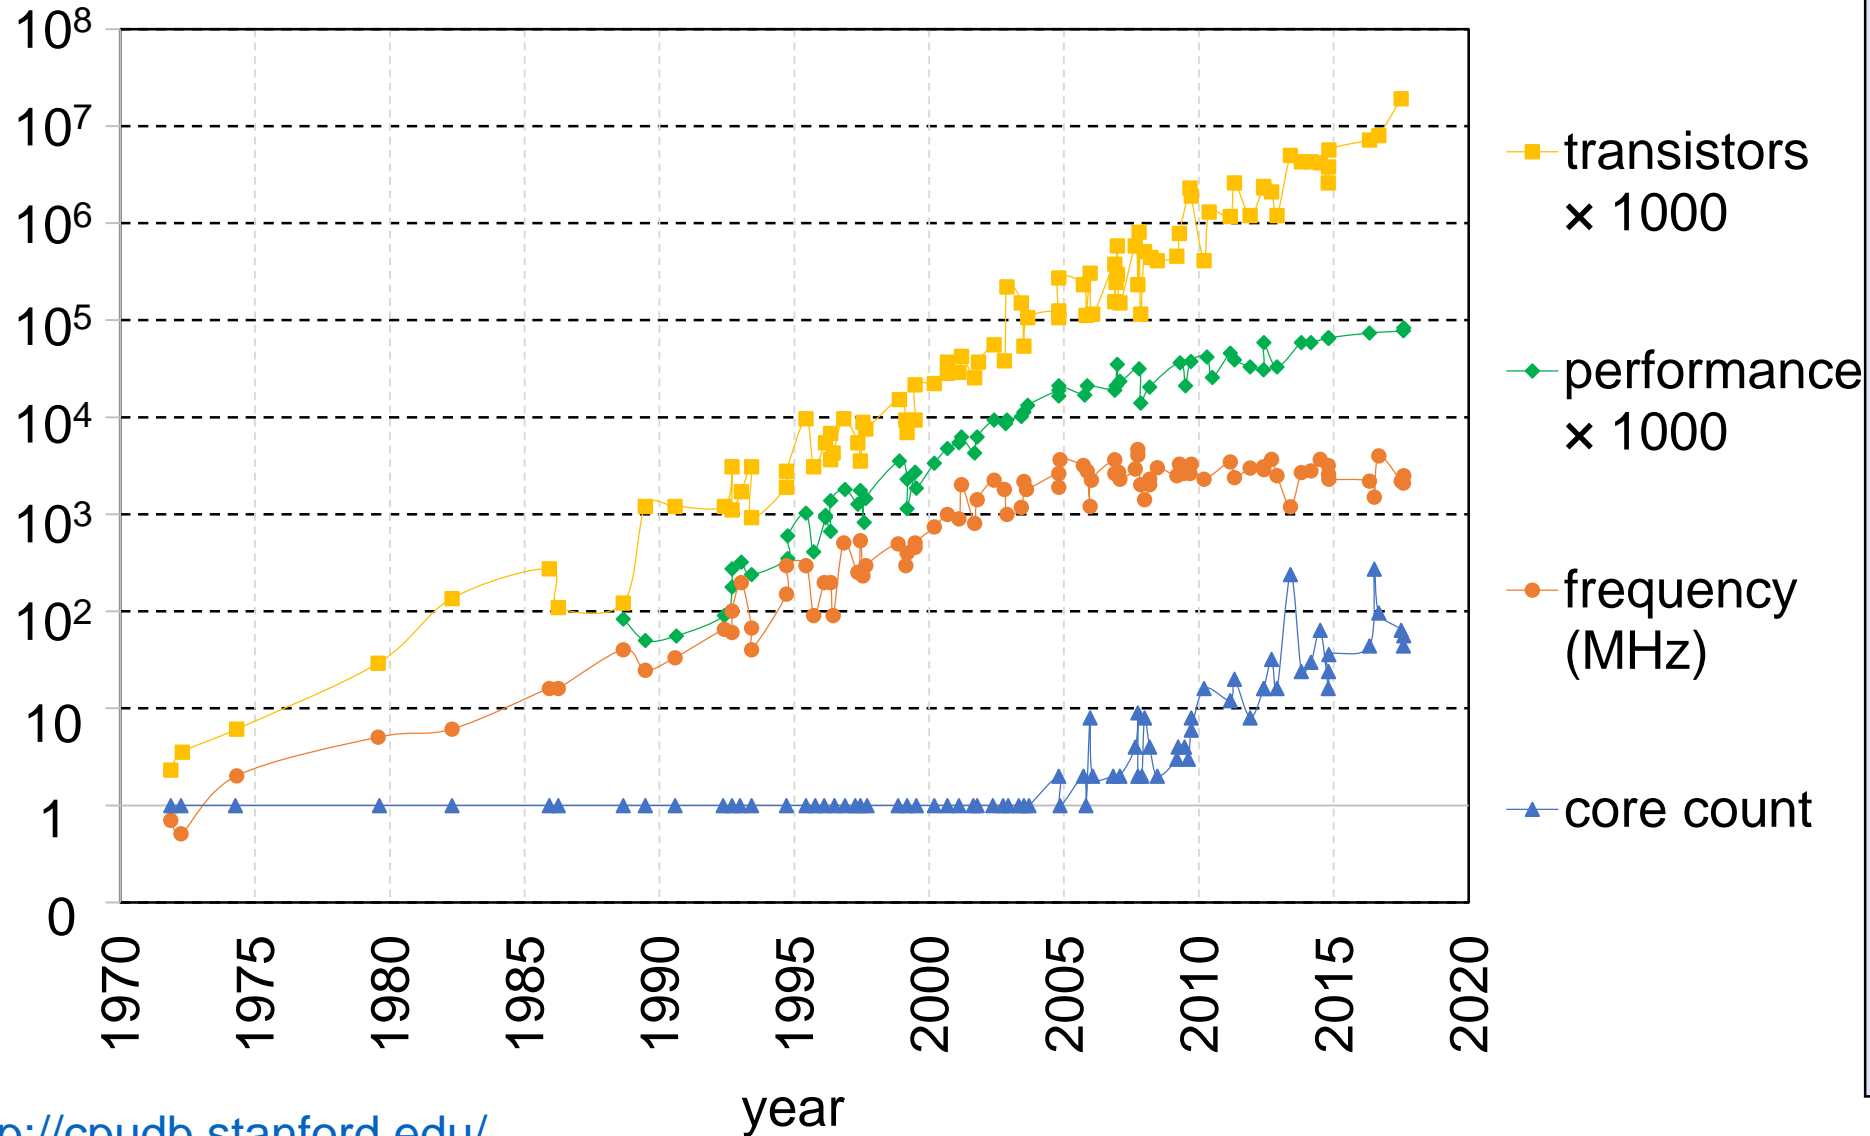


Consequences of Technology Scaling on Soft Error Reliability

Consequences of Technology Scaling on Soft Error Reliability

Smaller transistors → Reduced charge

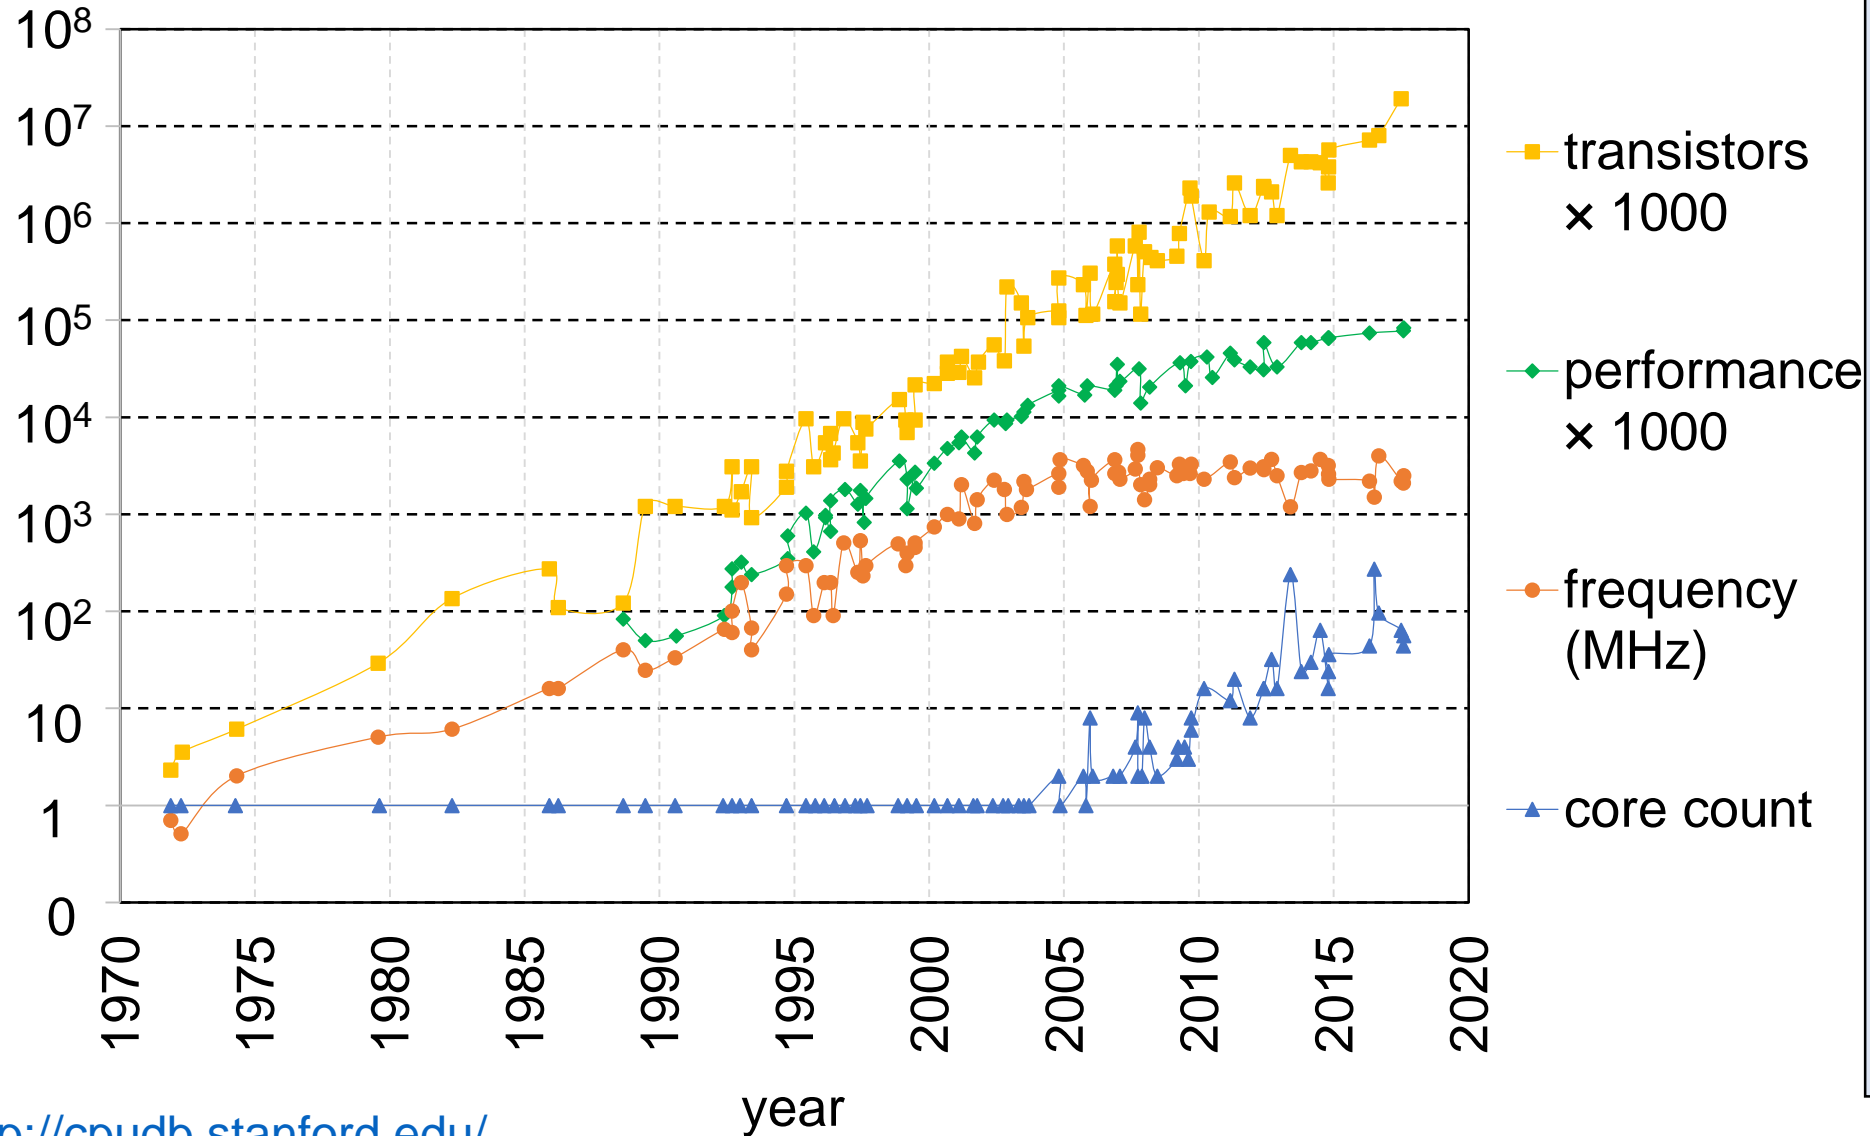
Consequences of Technology Scaling



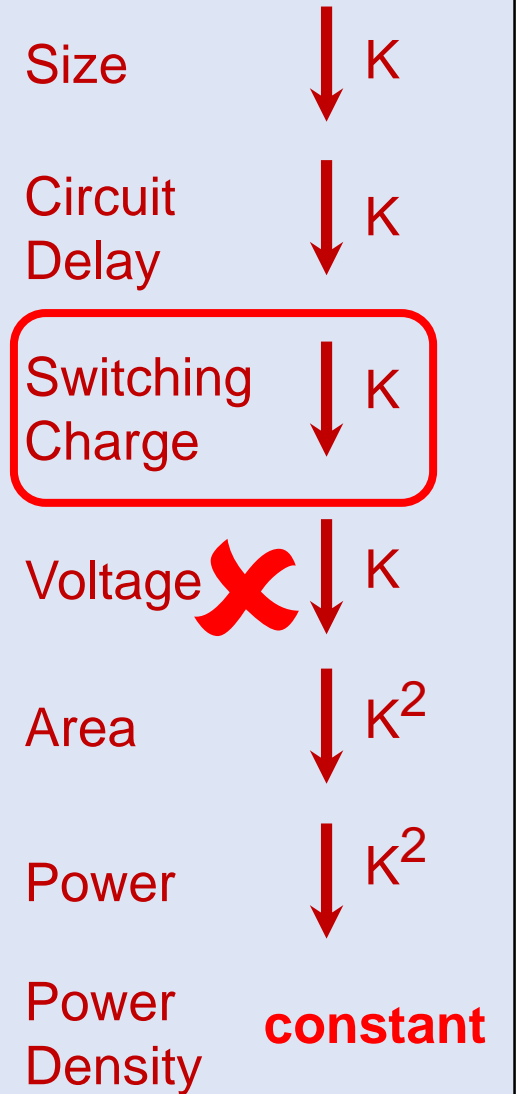
Dennard Scaling

Size	$\downarrow K$
Circuit Delay	$\downarrow K$
Switching Charge	$\downarrow K$
Voltage	$\downarrow K$
Area	$\downarrow K^2$
Power	$\downarrow K^2$
Power Density	constant

Consequences of Technology Scaling



Dennard Scaling



Consequences of Technology Scaling on Soft Error Reliability

Smaller transistors → Reduced charge

Consequences of Technology Scaling on Soft Error Reliability

Smaller transistors → Reduced charge
– Easily surpassed by neutrons

Consequences of Technology Scaling on Soft Error Reliability

Smaller transistors → Reduced charge

- Easily surpassed by neutrons

More transistors → More faults → More errors/failures

Consequences of Technology Scaling on Soft Error Reliability

Smaller transistors → Reduced charge

– Easily surpassed by neutrons

More transistors → More faults → More errors/failures

Technology
scaling

Consequences of Technology Scaling on Soft Error Reliability

Smaller transistors → Reduced charge

– Easily surpassed by neutrons

More transistors → More faults → More errors/failures

Technology
scaling



Increased
rate of soft errors

Modern Processors are Increasingly Vulnerable to Soft Errors

Modern Processors are Increasingly Vulnerable to Soft Errors

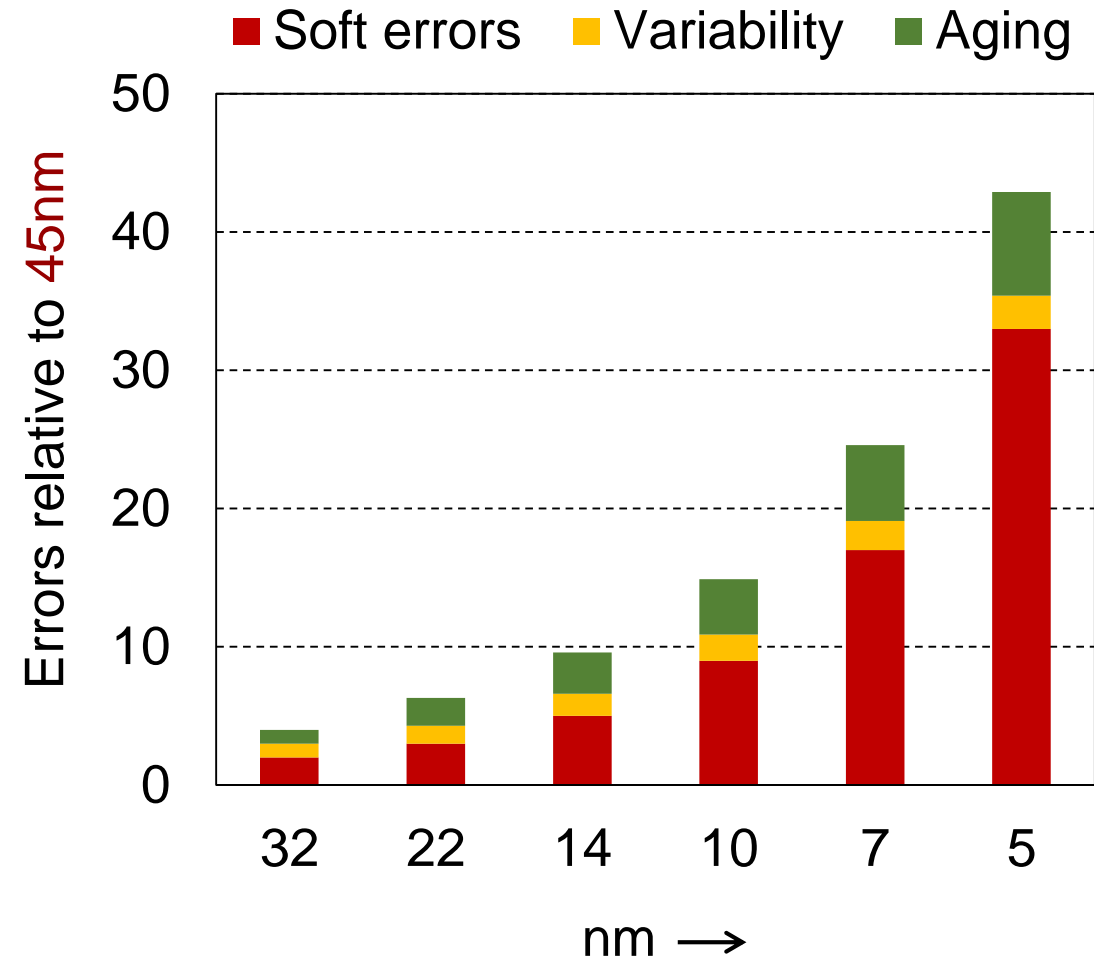
- Aggressive technology scaling

Modern Processors are Increasingly Vulnerable to Soft Errors

- Aggressive technology scaling
- Large microarchitectural state

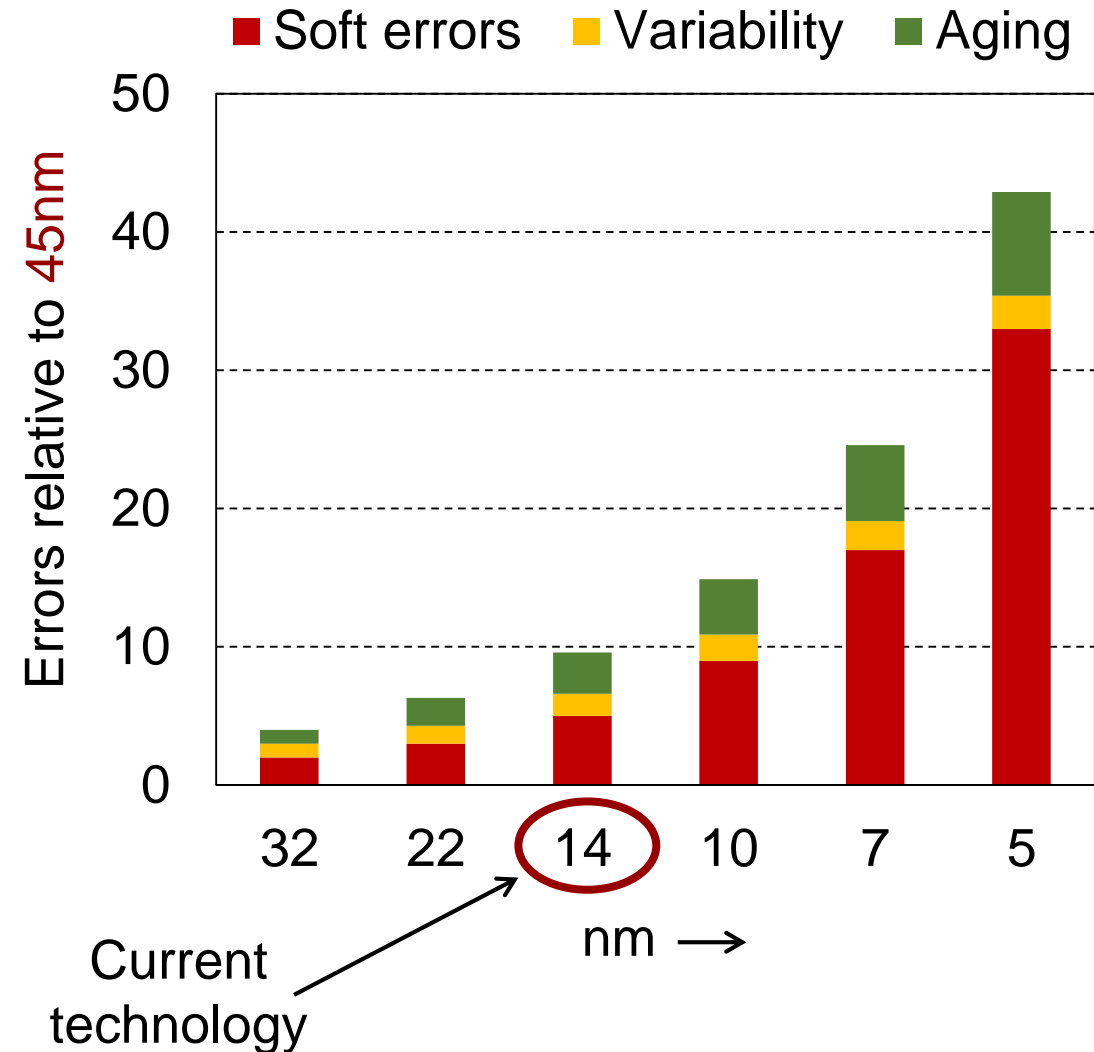
Modern Processors are Increasingly Vulnerable to Soft Errors

- Aggressive technology scaling
- Large microarchitectural state



Modern Processors are Increasingly Vulnerable to Soft Errors

- Aggressive technology scaling
- Large microarchitectural state



This Dissertation: Three New Contributions

This Dissertation: Three New Contributions

1. Exploiting core heterogeneity to improve reliability
[HPCA 2017, TC 2018]

This Dissertation: Three New Contributions

1. Exploiting core heterogeneity to improve reliability
[HPCA 2017, TC 2018]
2. Dispatch Halting: Improving out-of-order core reliability
[DSN 2020, Under review]

This Dissertation: Three New Contributions

1. Exploiting core heterogeneity to improve reliability
[HPCA 2017, TC 2018]
2. Dispatch Halting: Improving out-of-order core reliability
[DSN 2020, Under review]
3. Precise Runahead Execution: Improving performance and evaluating reliability
[CAL 2019, HPCA 2020]

Contribution #1 [HPCA 2017, TC 2018]

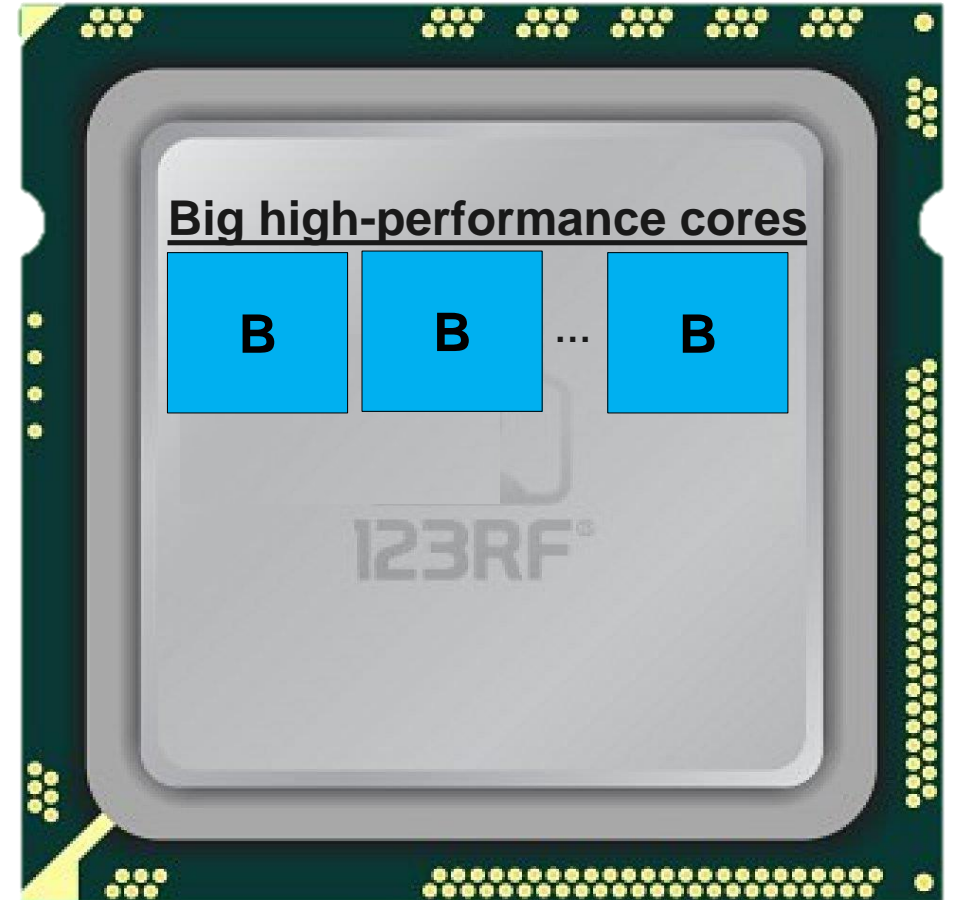
**Exploiting Core Heterogeneity
to Improve Reliability**

Benefits of Heterogeneous Multicores

- Multiple core types

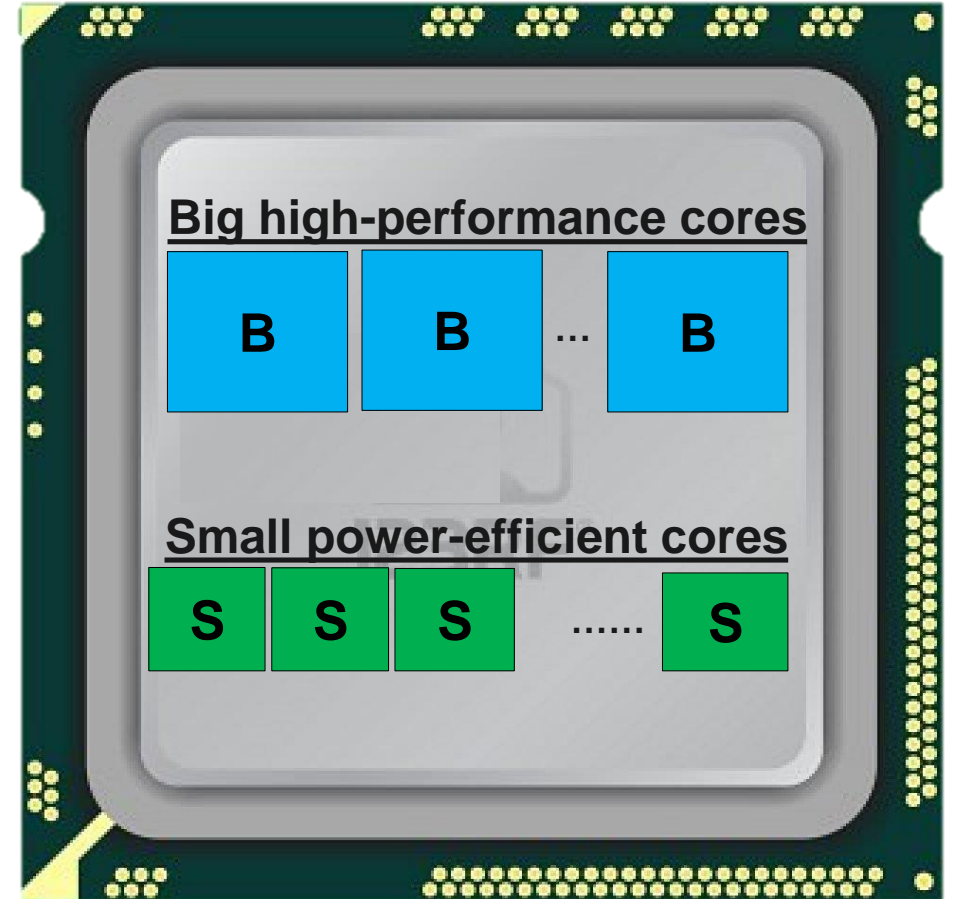
Benefits of Heterogeneous Multicores

- Multiple core types



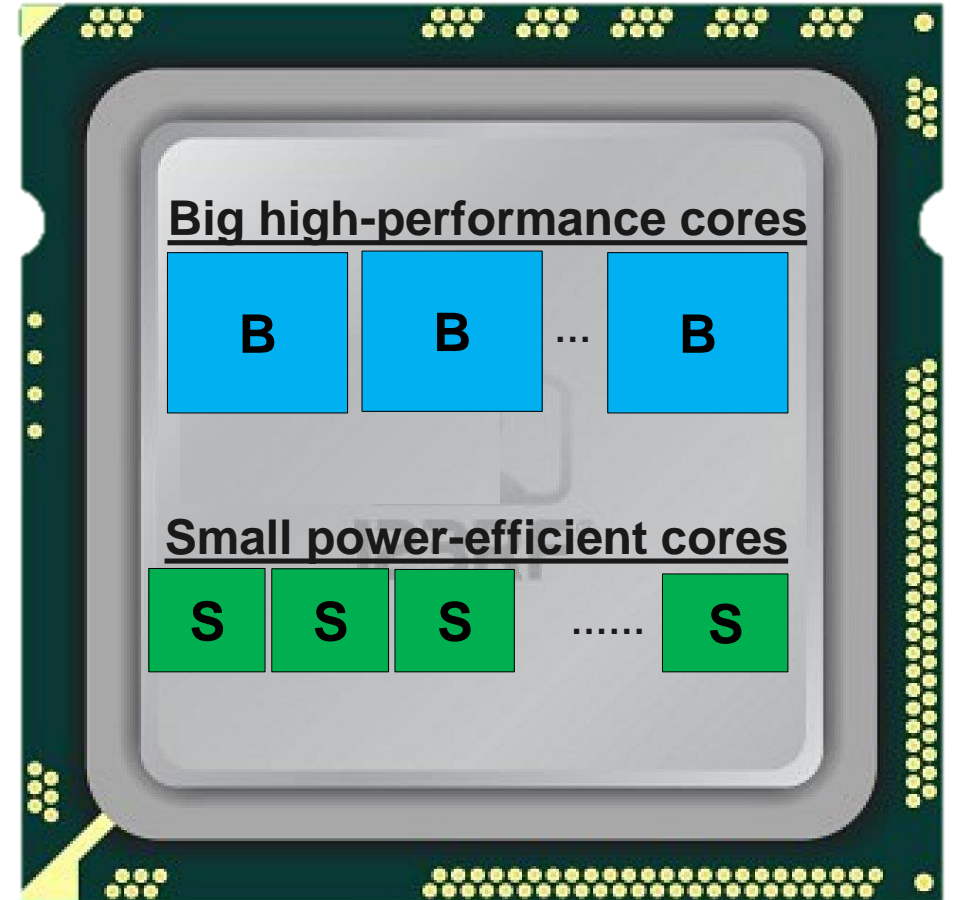
Benefits of Heterogeneous Multicores

- Multiple core types



Benefits of Heterogeneous Multicores

- Multiple core types
- Well-established **power benefits***

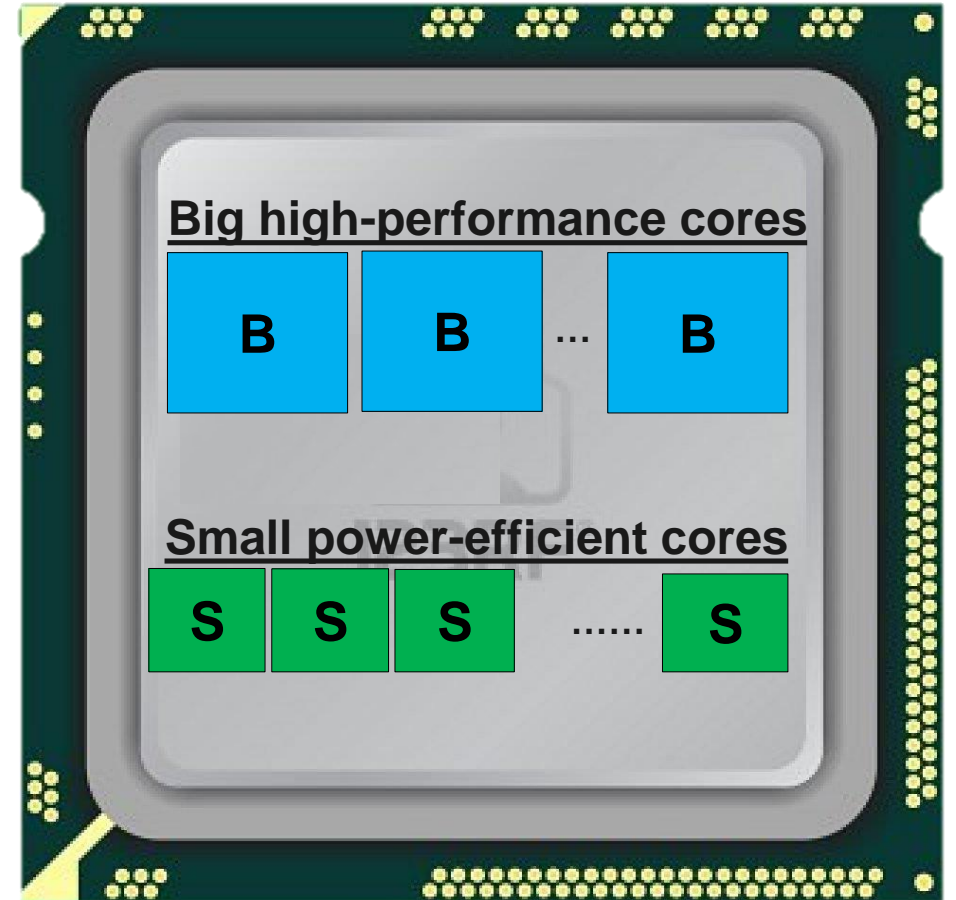


*[Kumar et al. MICRO'03, ISCA'04]

Benefits of Heterogeneous Multicores

- Multiple core types
- Well-established **power benefits***
- Well-established **scheduling techniques****

e.g. Big.LITTLE and Kal-El



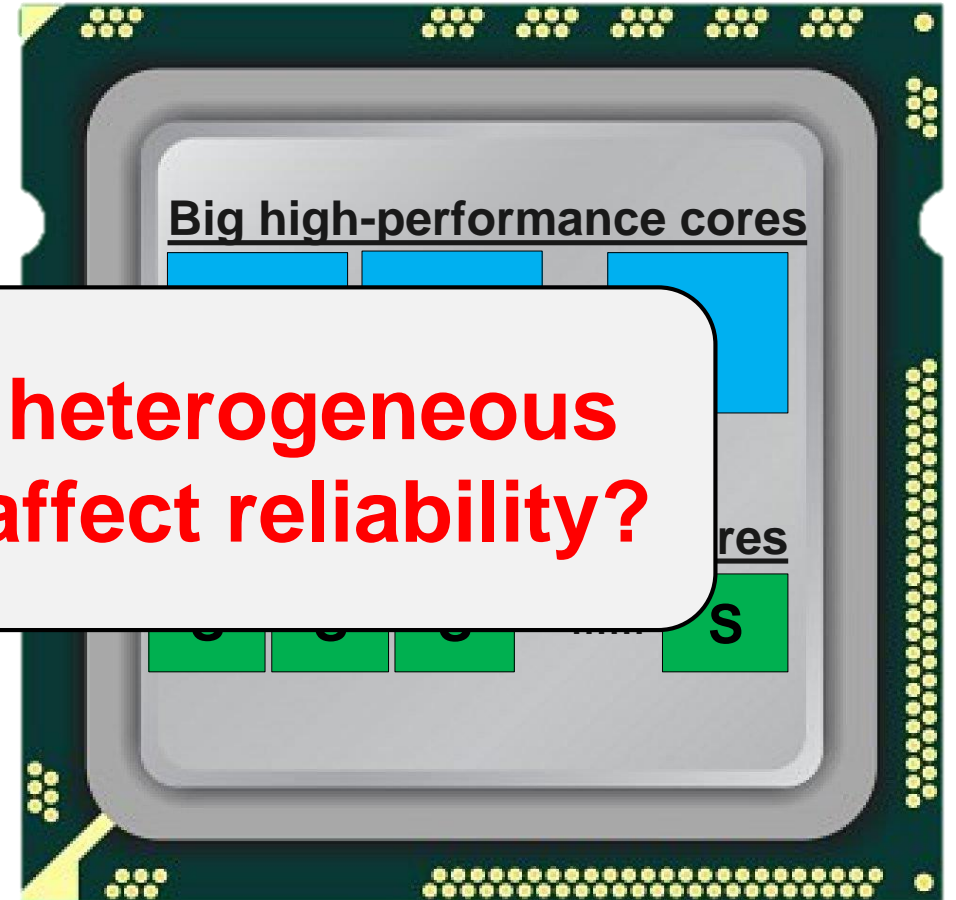
*[Kumar et al. MICRO'03, ISCA'04] **[Van Craeynest et al. ISCA'12]

Benefits of Heterogeneous Multicores

- Multiple core types
- Well-established techniques
- Well-established techniques**

No prior work → How heterogeneous chip-multiprocessors affect reliability?

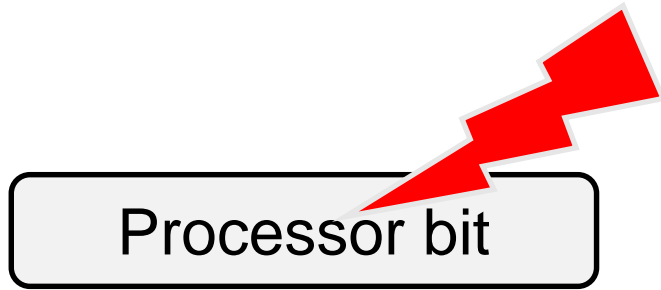
e.g. Big.LITTLE and Kal-El



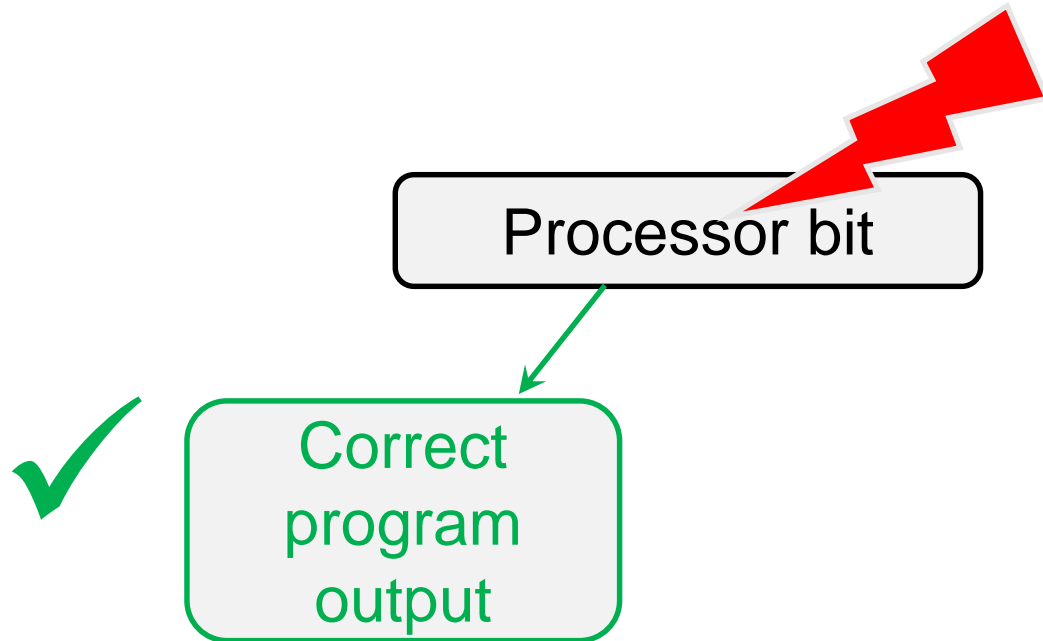
ACE Analysis for Estimating Soft Errors

Processor bit

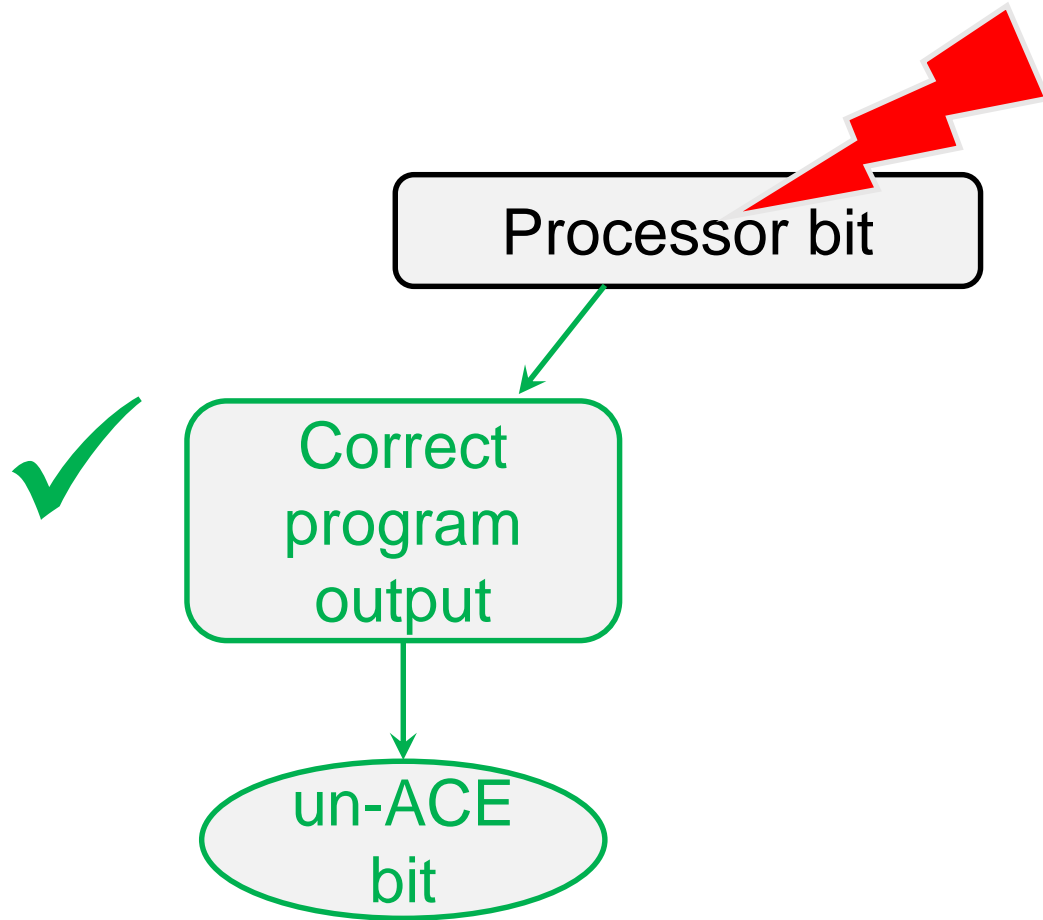
ACE Analysis for Estimating Soft Errors



ACE Analysis for Estimating Soft Errors

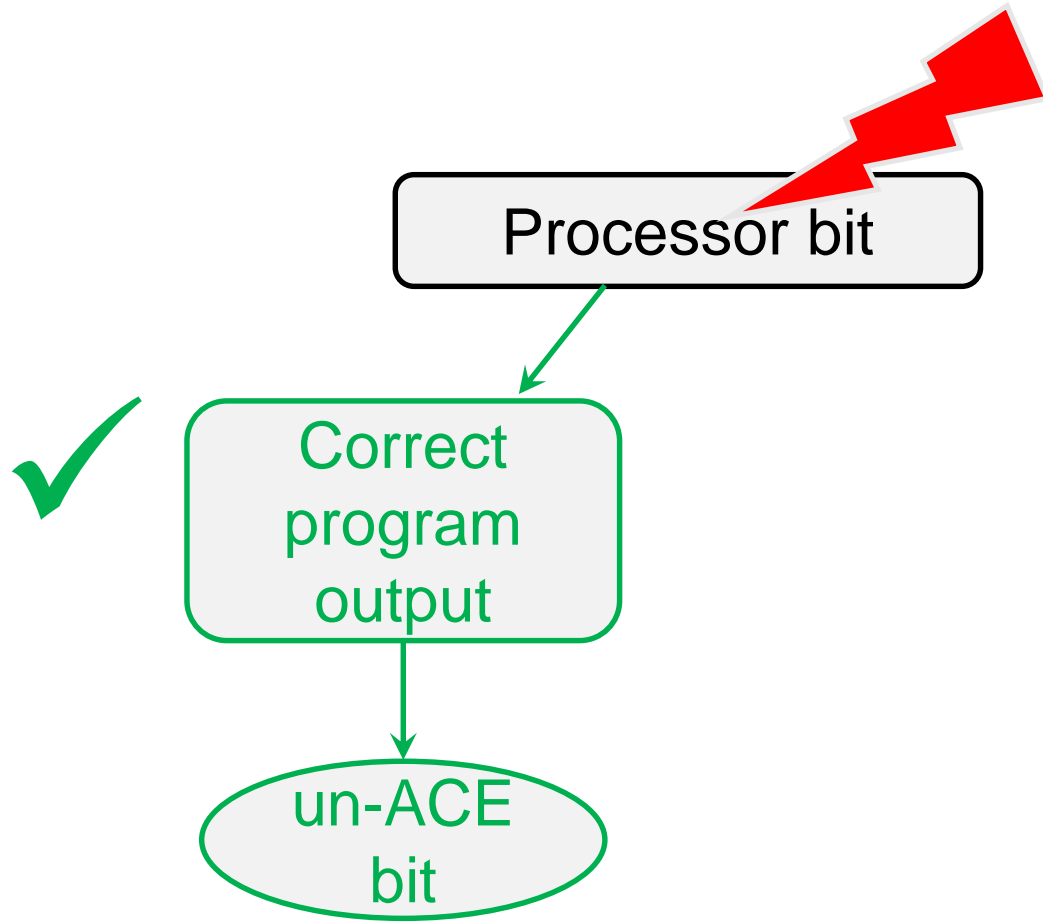


ACE Analysis for Estimating Soft Errors



ACE →
Architecturally
Correct
Execution

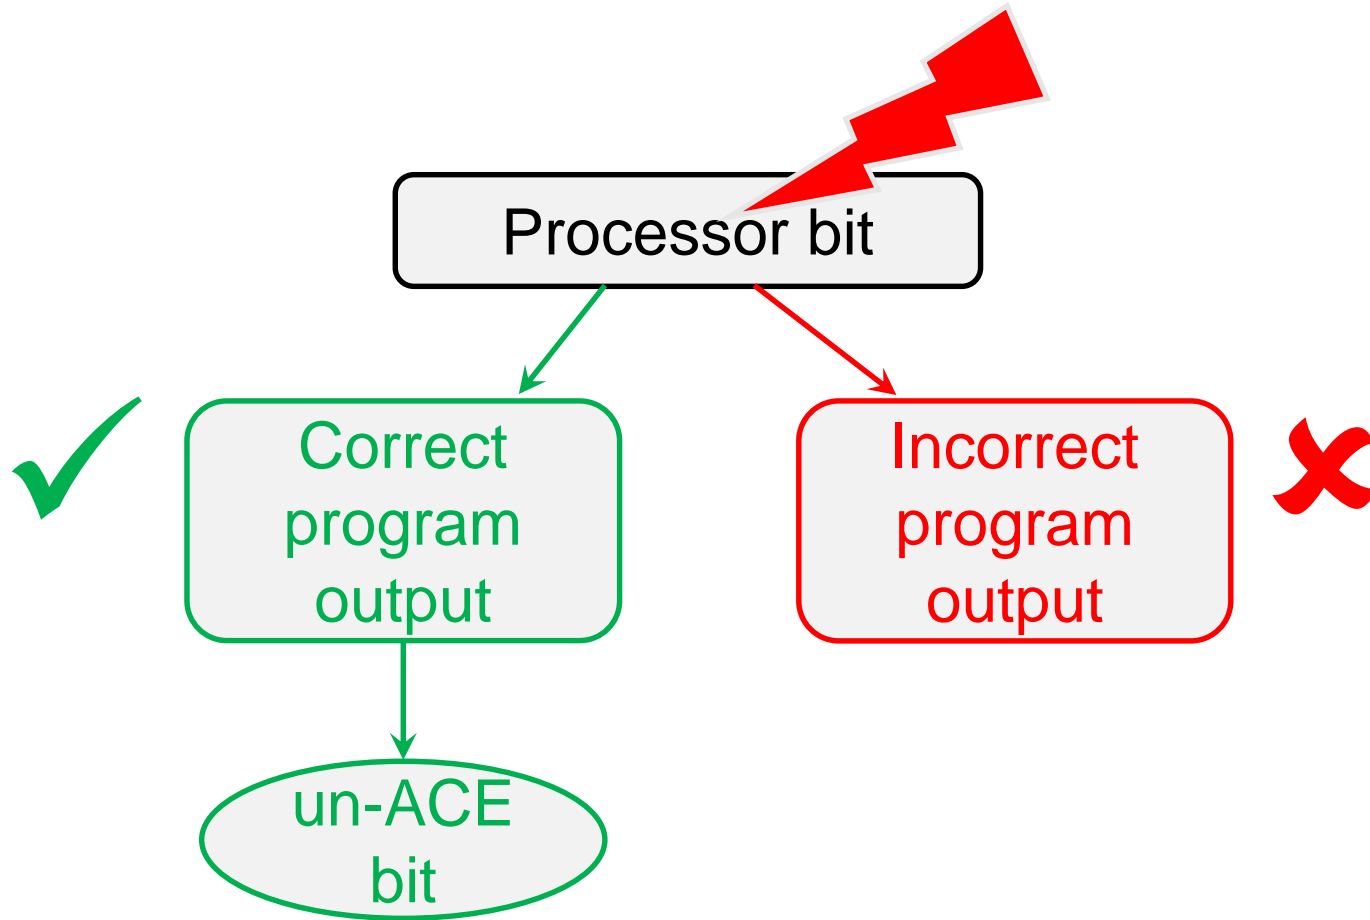
ACE Analysis for Estimating Soft Errors



ACE →
Architecturally
Correct
Execution

Branch predictor,
wrong-path instructions

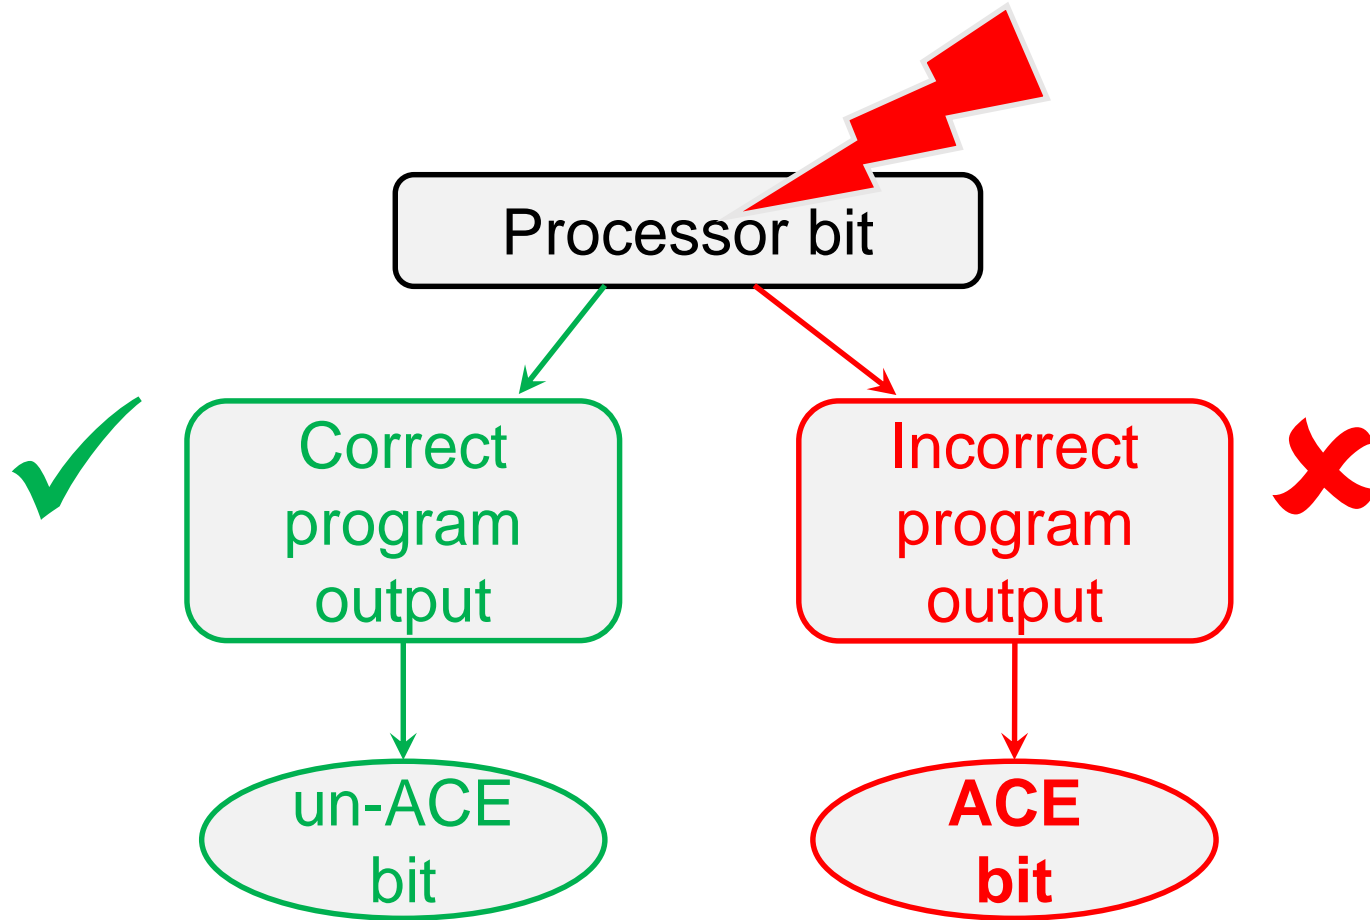
ACE Analysis for Estimating Soft Errors



Branch predictor,
wrong-path instructions

ACE →
Architecturally
Correct
Execution

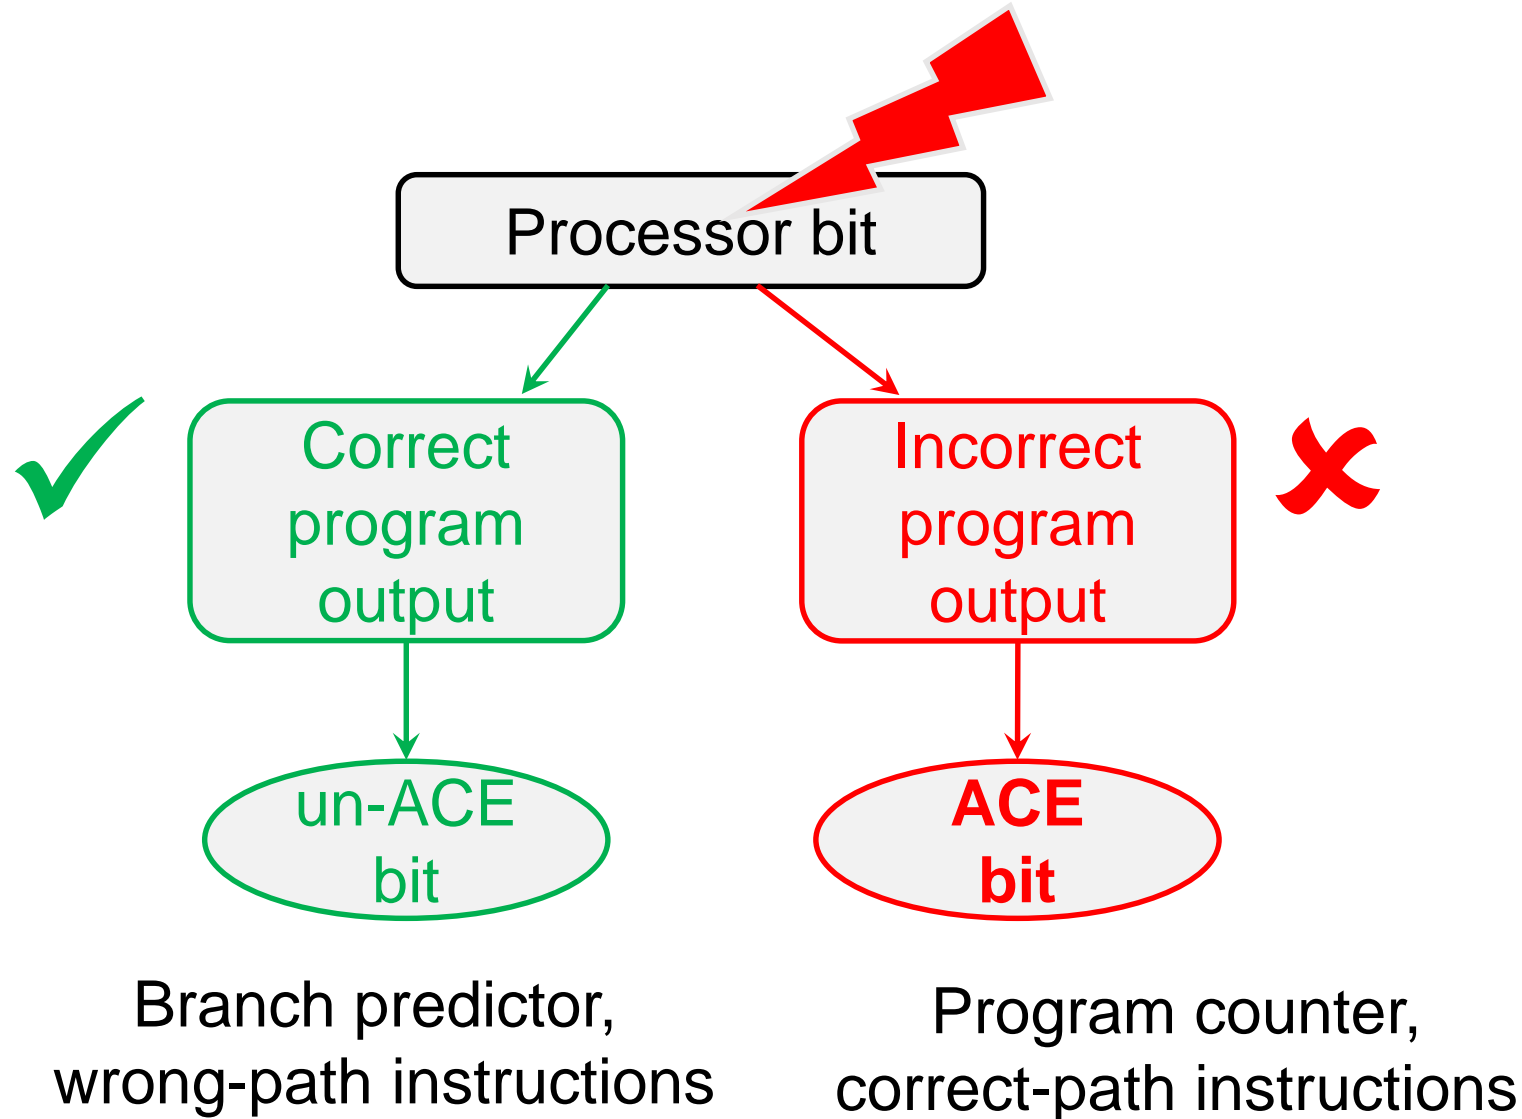
ACE Analysis for Estimating Soft Errors



ACE →
Architecturally
Correct
Execution

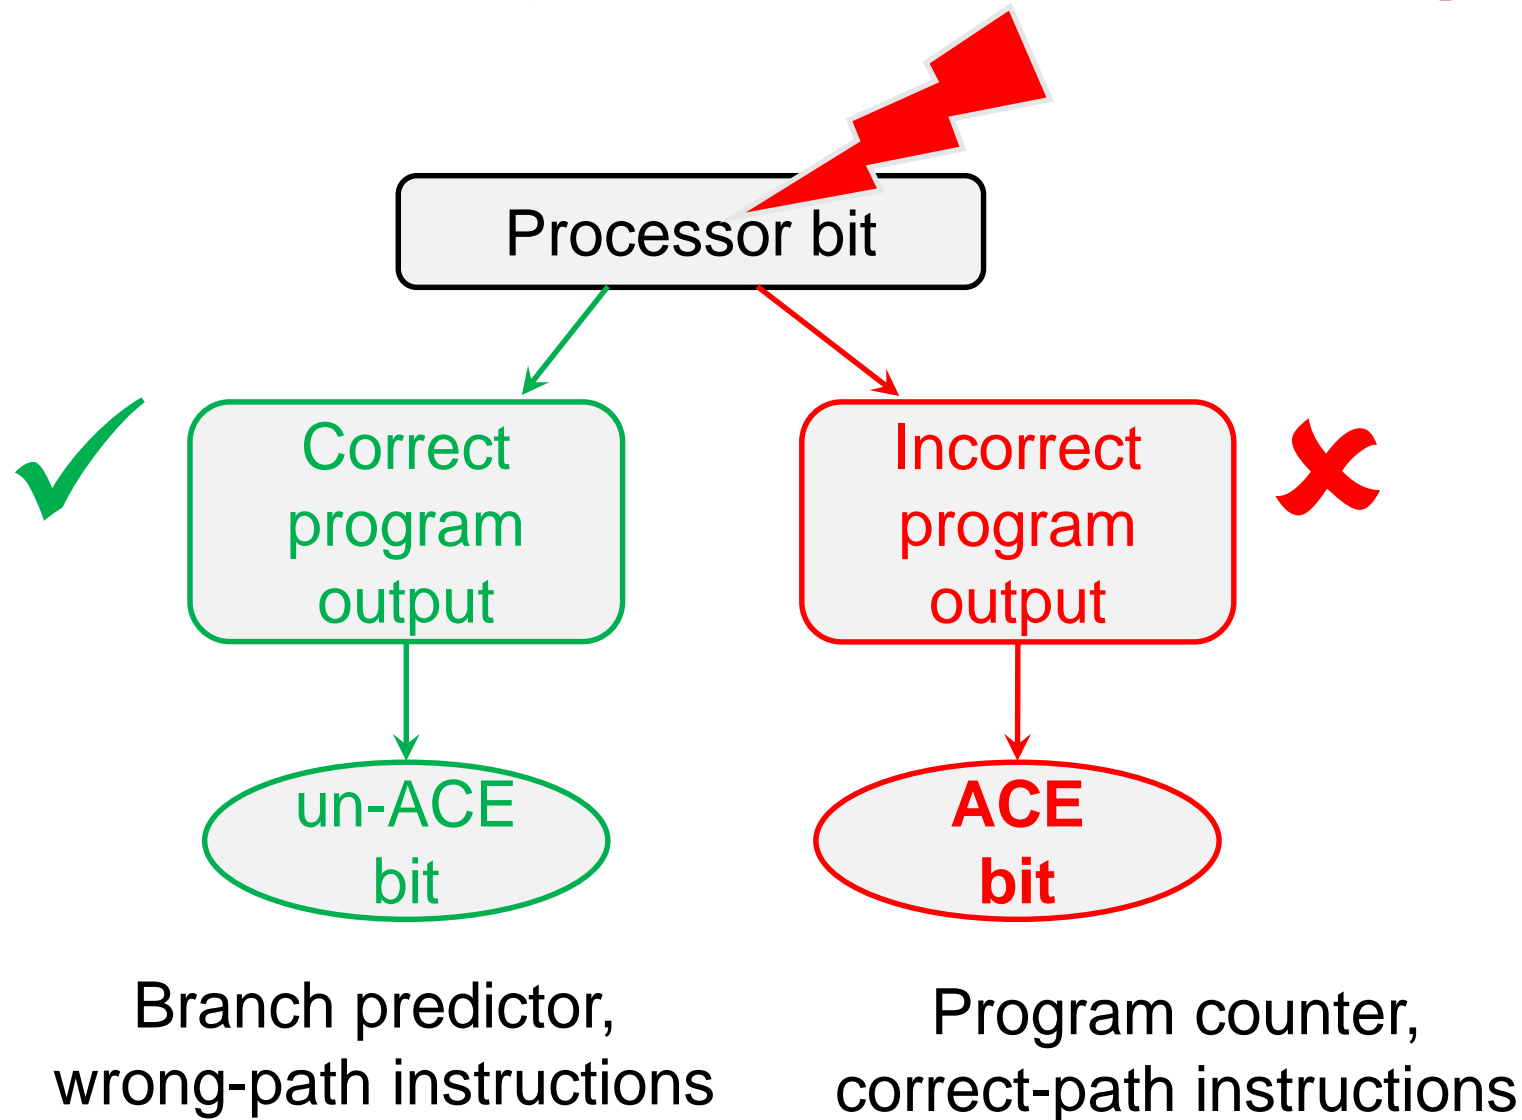
Branch predictor,
wrong-path instructions

ACE Analysis for Estimating Soft Errors



ACE →
Architecturally
Correct
Execution

ACE Analysis for Estimating Soft Errors



ACE →
Architecturally
Correct
Execution

**For correct program
execution, all ACE
bits must be correct**

ACE Analysis for Estimating Soft Errors

ACE Analysis for Estimating Soft Errors

ACE Cycles = Total cycles a bit was ACE

ACE Analysis for Estimating Soft Errors

ACE Cycles = Total cycles a bit was ACE

ACE Bit Count (**ABC**) = Sum of ACE cycles for all committed instructions

ACE Analysis for Estimating Soft Errors

ACE Cycles = Total cycles a bit was ACE

ACE Bit Count (**ABC**) = Sum of ACE cycles for all committed instructions

Soft Error Rate (**SER**) = $\frac{ABC}{\text{Total cycles}} \times \text{IFR}$ IFR = Probability of a soft error

ACE Analysis for Estimating Soft Errors

ACE Cycles = Total cycles a bit was ACE

ACE Bit Count (**ABC**) = Sum of ACE cycles for all committed instructions

Soft Error Rate (**SER**) = $\frac{ABC}{\text{Total cycles}} \times \text{IFR}$ IFR = Probability of a soft error

For the same
execution time

ACE Analysis for Estimating Soft Errors

ACE Cycles = Total cycles a bit was ACE

ACE Bit Count (**ABC**) = Sum of ACE cycles for all committed instructions

Soft Error Rate (**SER**) = $\frac{ABC}{\text{Total cycles}} \times \text{IFR}$ IFR = Probability of a soft error

For the same
execution time

ABC

SER

ACE Analysis for Estimating Soft Errors

ACE Cycles = Total cycles a bit was ACE

ACE Bit Count (**ABC**) = Sum of ACE cycles for all committed instructions

Soft Error Rate (**SER**) = $\frac{ABC}{\text{Total cycles}} \times \text{IFR}$ IFR = Probability of a soft error

For the same
execution time

ABC ↑

SER

ACE Analysis for Estimating Soft Errors

ACE Cycles = Total cycles a bit was ACE

ACE Bit Count (**ABC**) = Sum of ACE cycles for all committed instructions

Soft Error Rate (**SER**) = $\frac{ABC}{\text{Total cycles}} \times \text{IFR}$ IFR = Probability of a soft error

For the same
execution time

ABC ↑

SER ↑

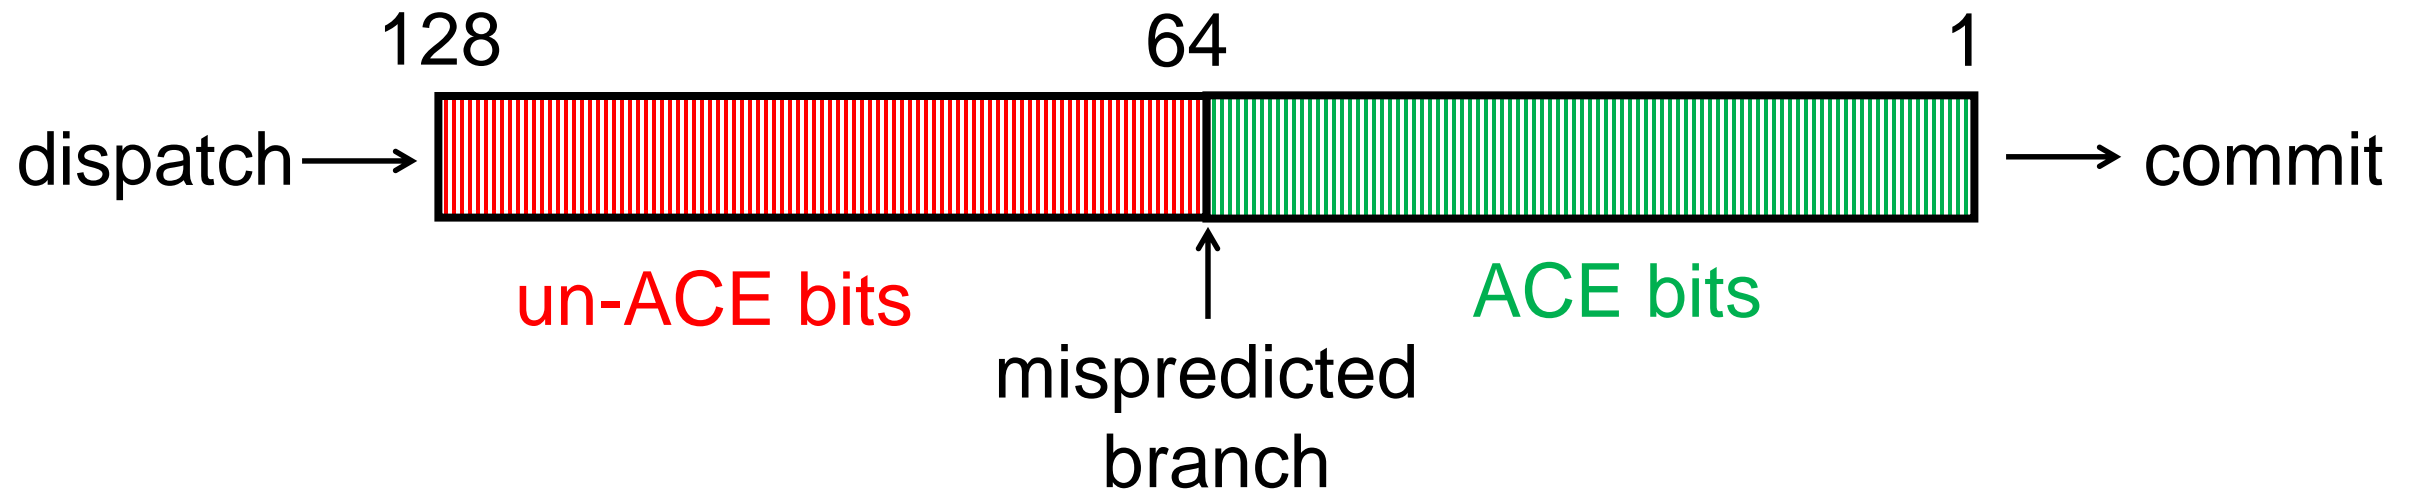
Example: Reorder Buffer (ROB)

Example: Reorder Buffer (ROB)

Size = 128, each entry = 100 bits, total bits = 12800

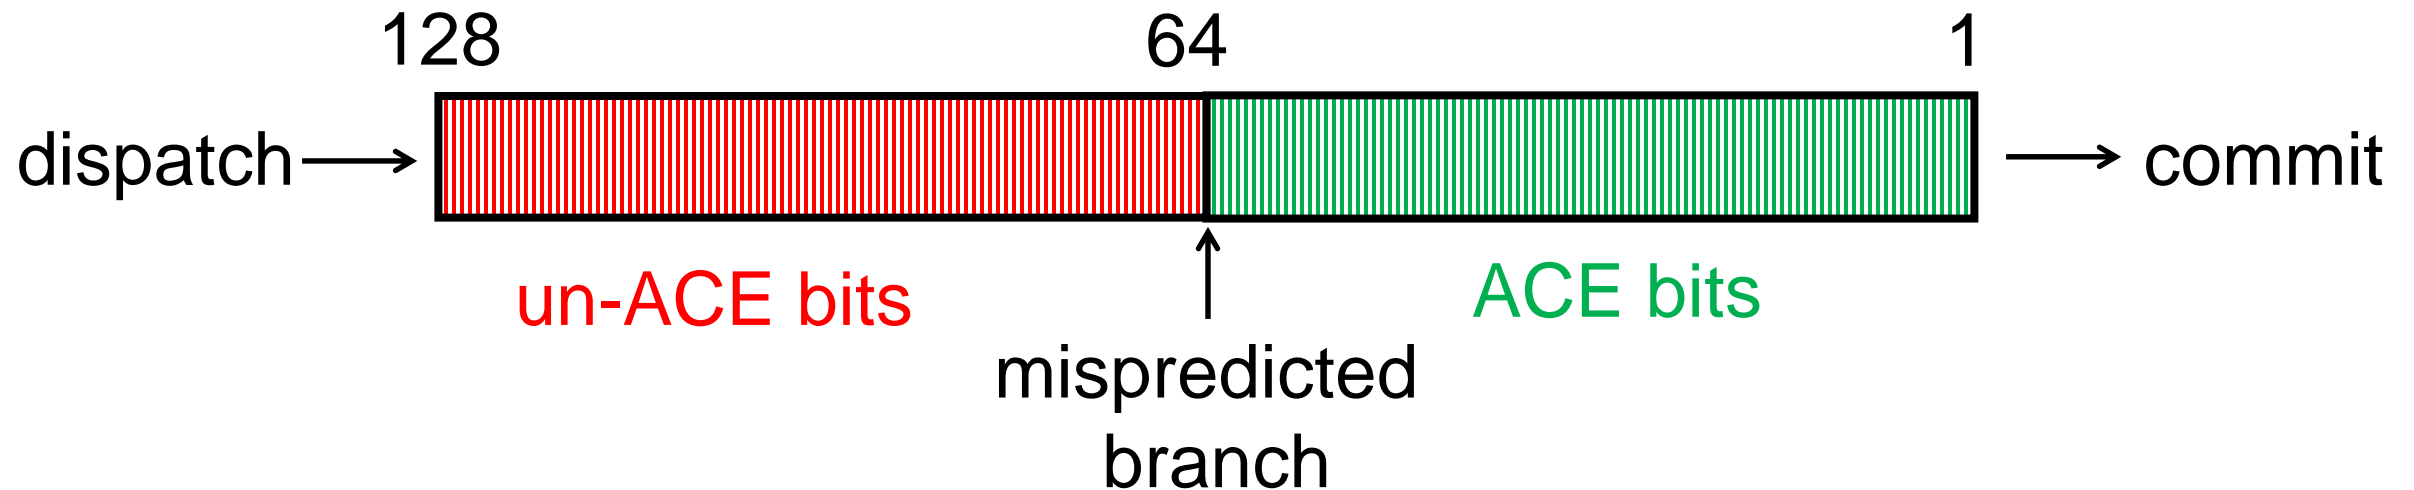
Example: Reorder Buffer (ROB)

Size = 128, each entry = 100 bits, total bits = 12800



Example: Reorder Buffer (ROB)

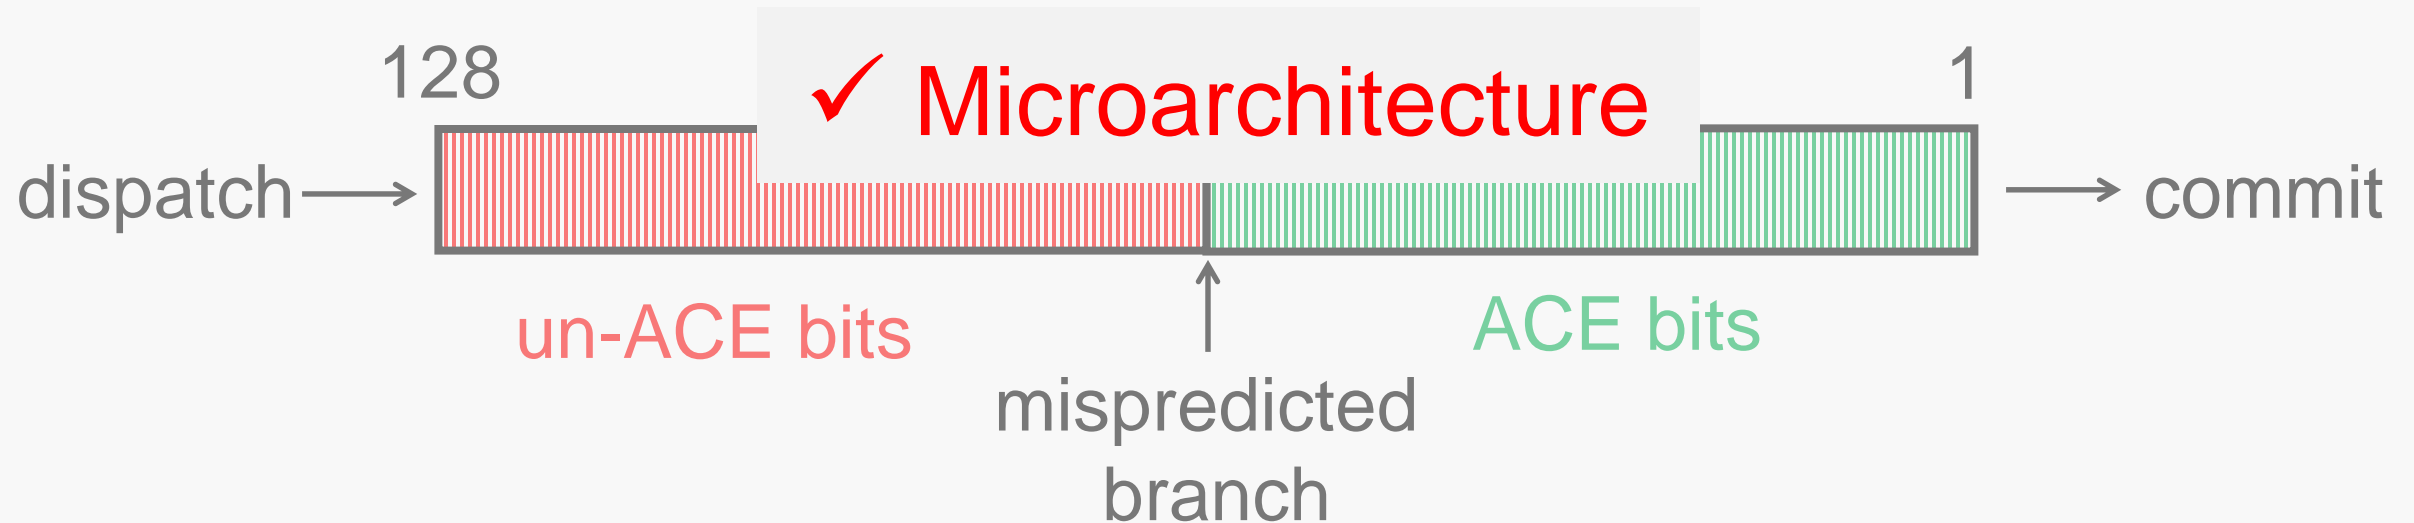
Size = 128, each entry = 100 bits, total bits = 12800



$$\text{ROB ACE bits} = 64 \times 100 = 6400$$

Example: Reorder Buffer (ROB)

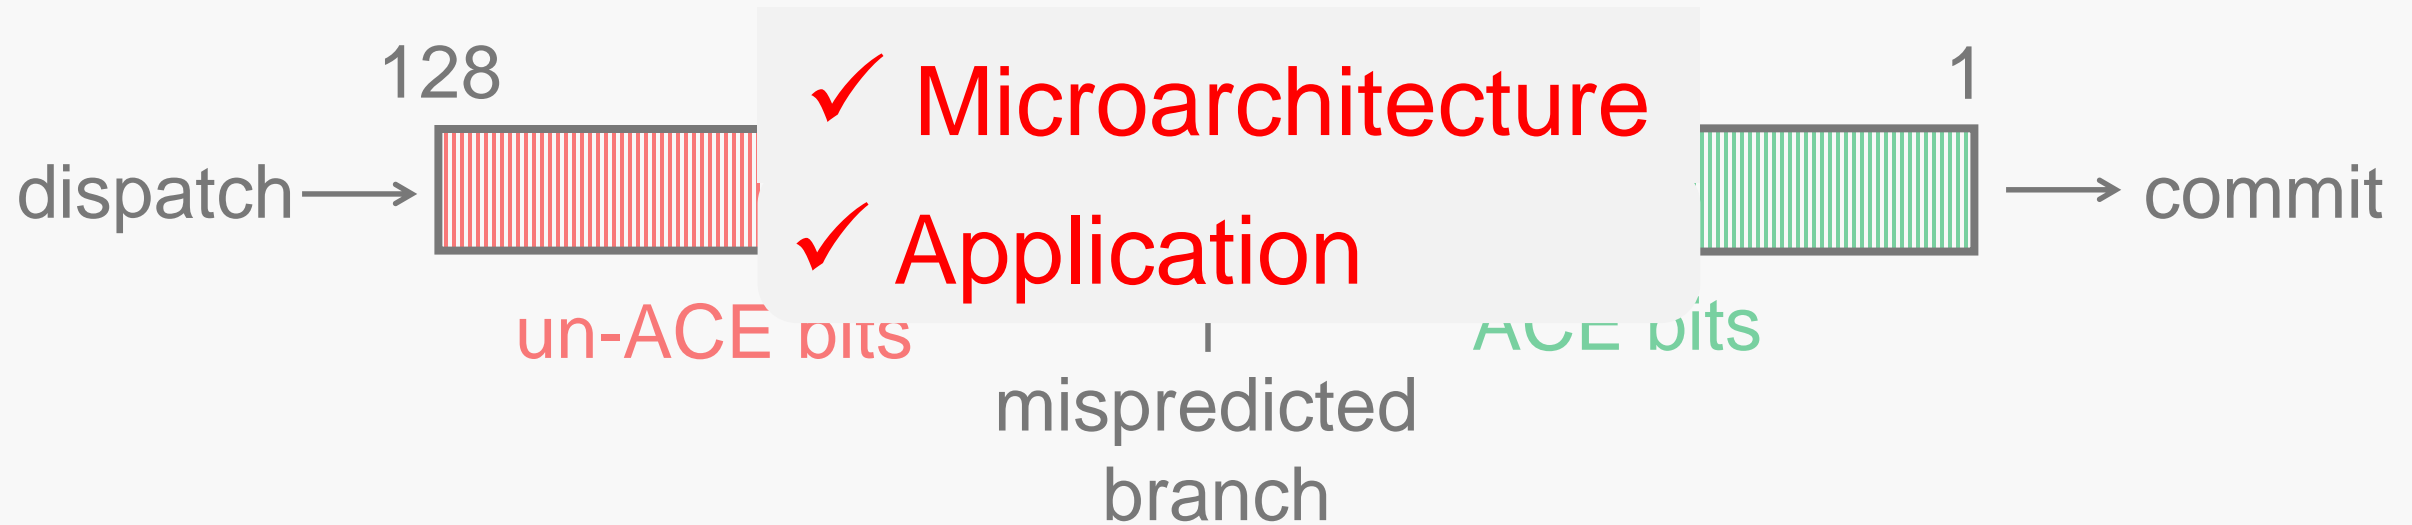
Size = 128, each entry = 100 bits, total bits = 12800



ROB ACE bits = $64 \times 100 = 6400$

Example: Reorder Buffer (ROB)

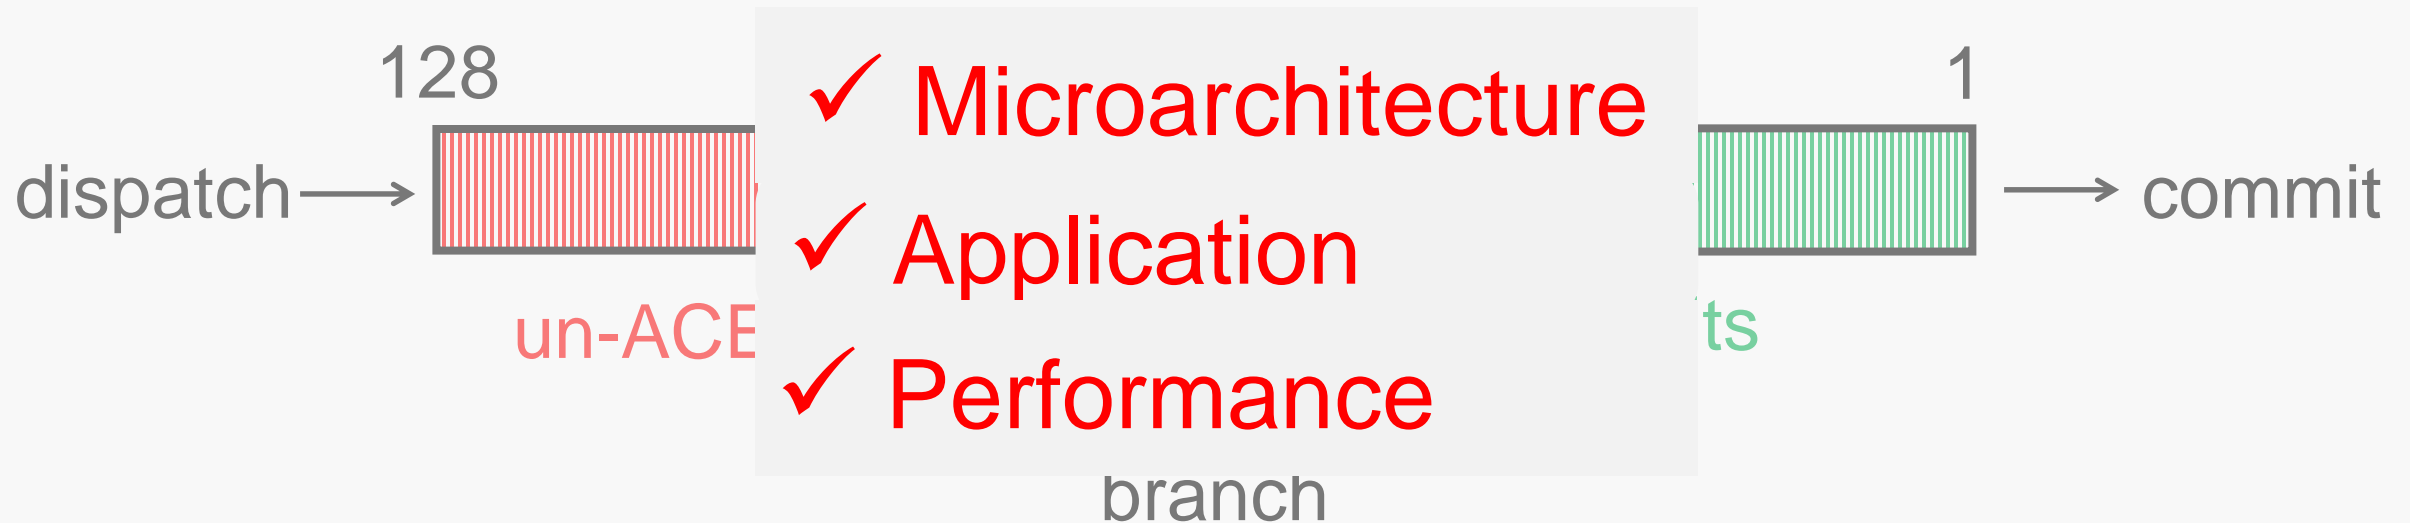
Size = 128, each entry = 100 bits, total bits = 12800



ROB ACE bits = $64 \times 100 = 6400$

Example: Reorder Buffer (ROB)

Size = 128, each entry = 100 bits, total bits = 12800



ROB ACE bits = $64 \times 100 = 6400$

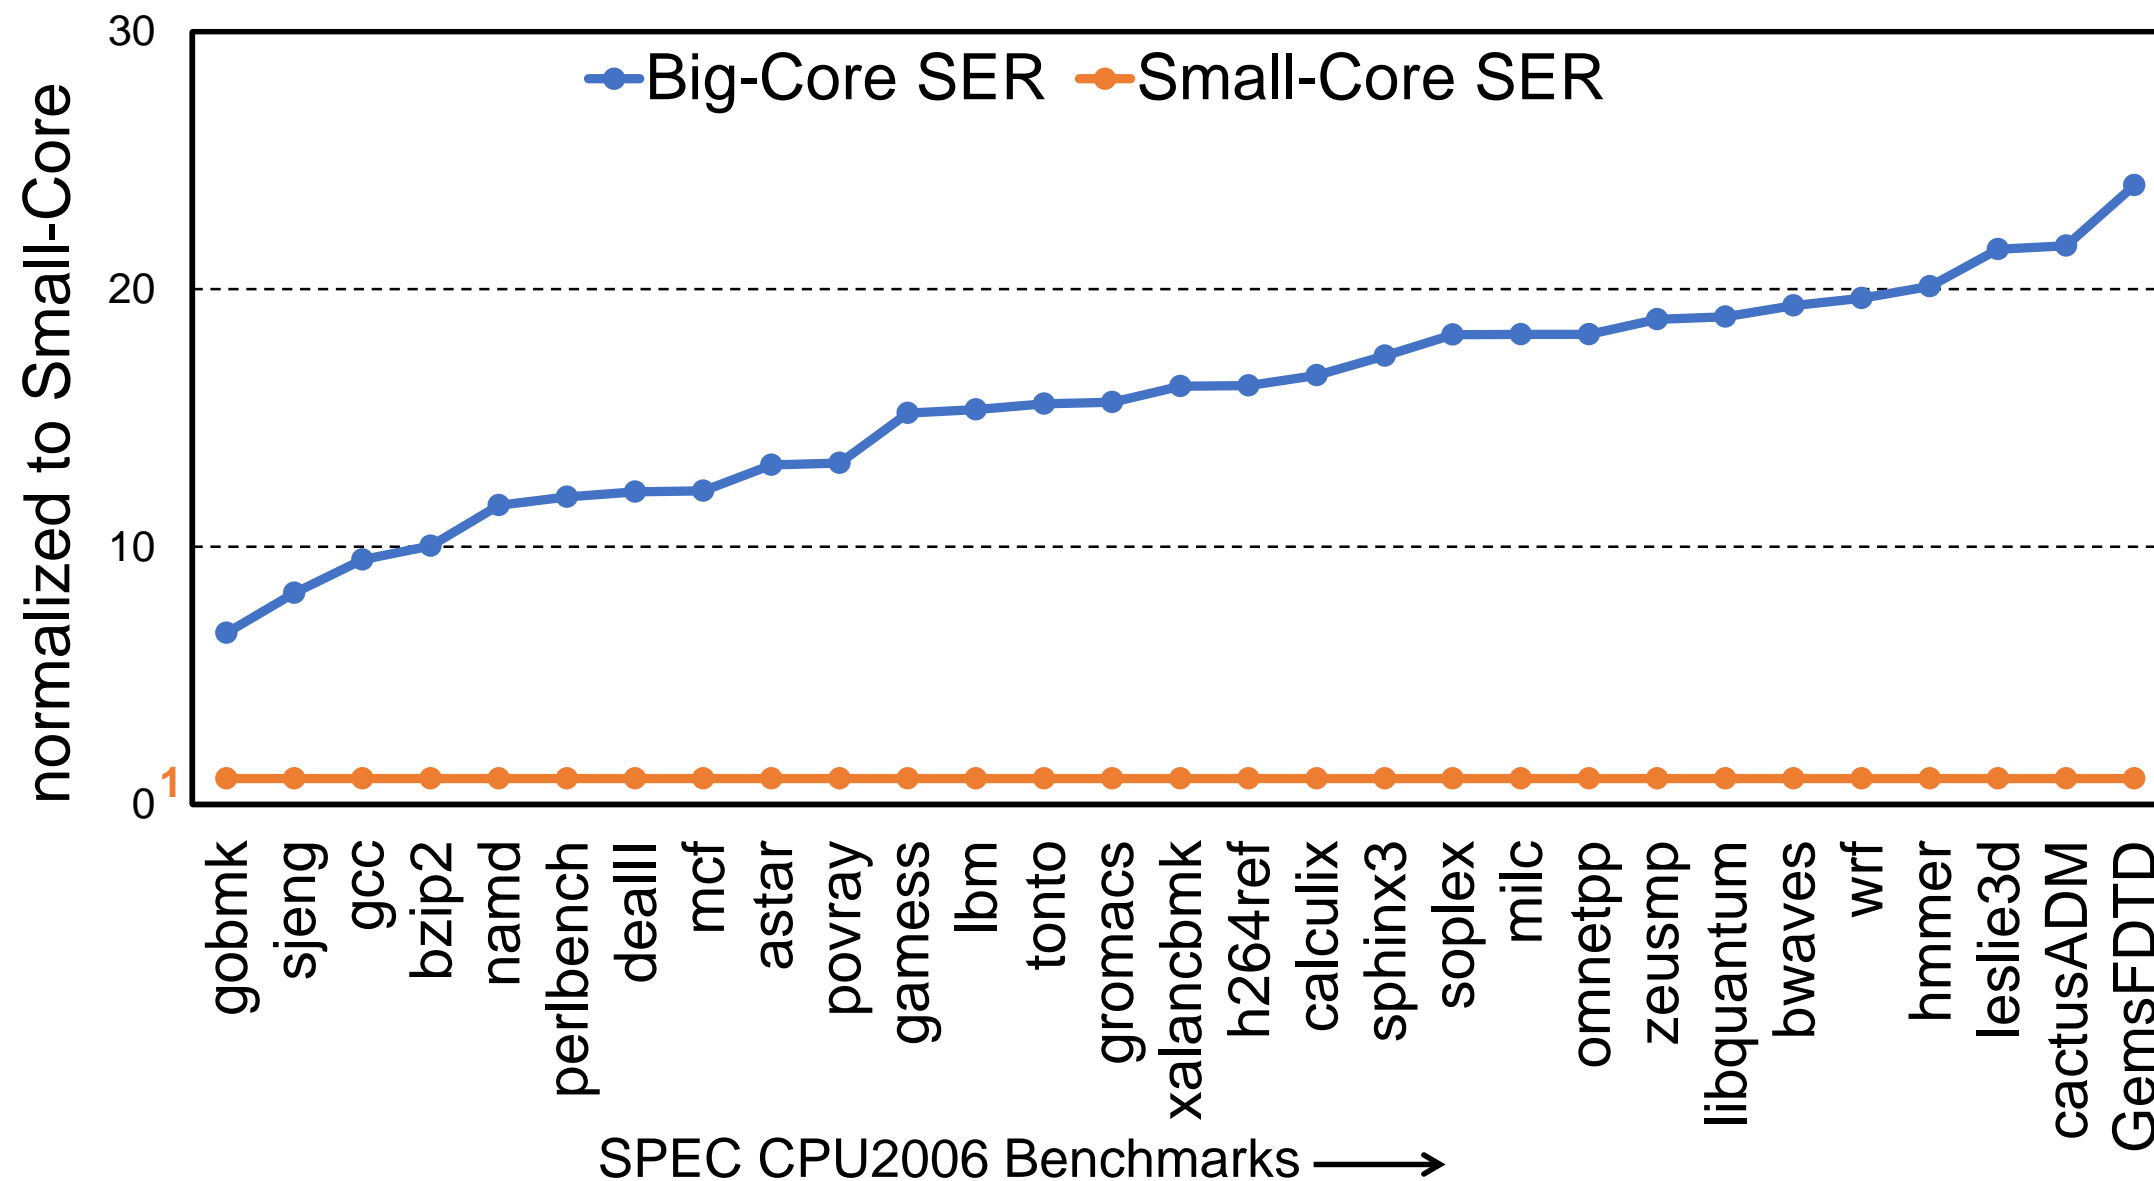
Example: Reorder Buffer (ROB)

Size = 128, each entry = 100 bits, total bits = 12800

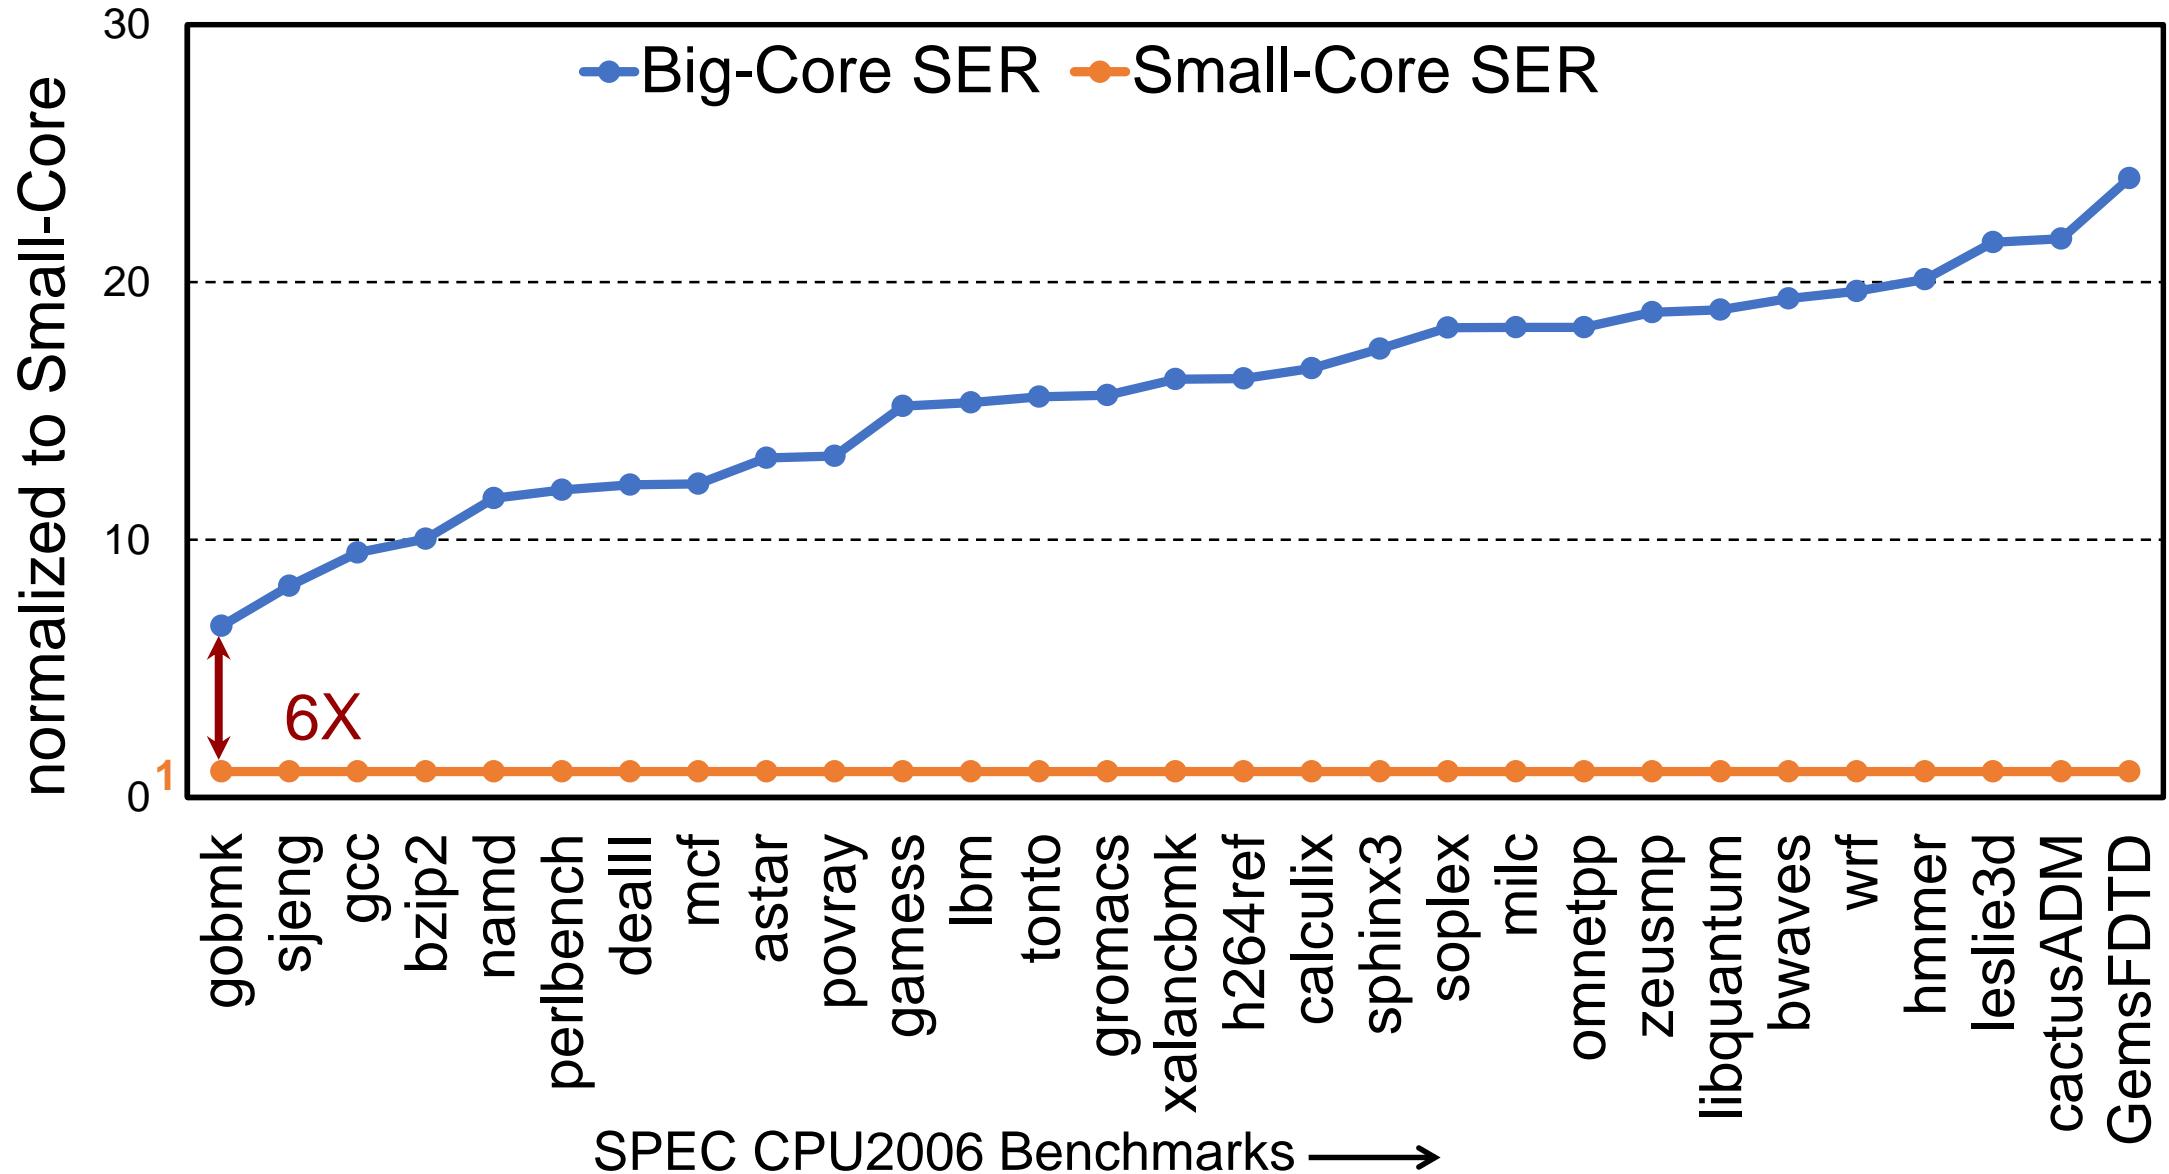


ROB ACE bits = $64 \times 100 = 6400$

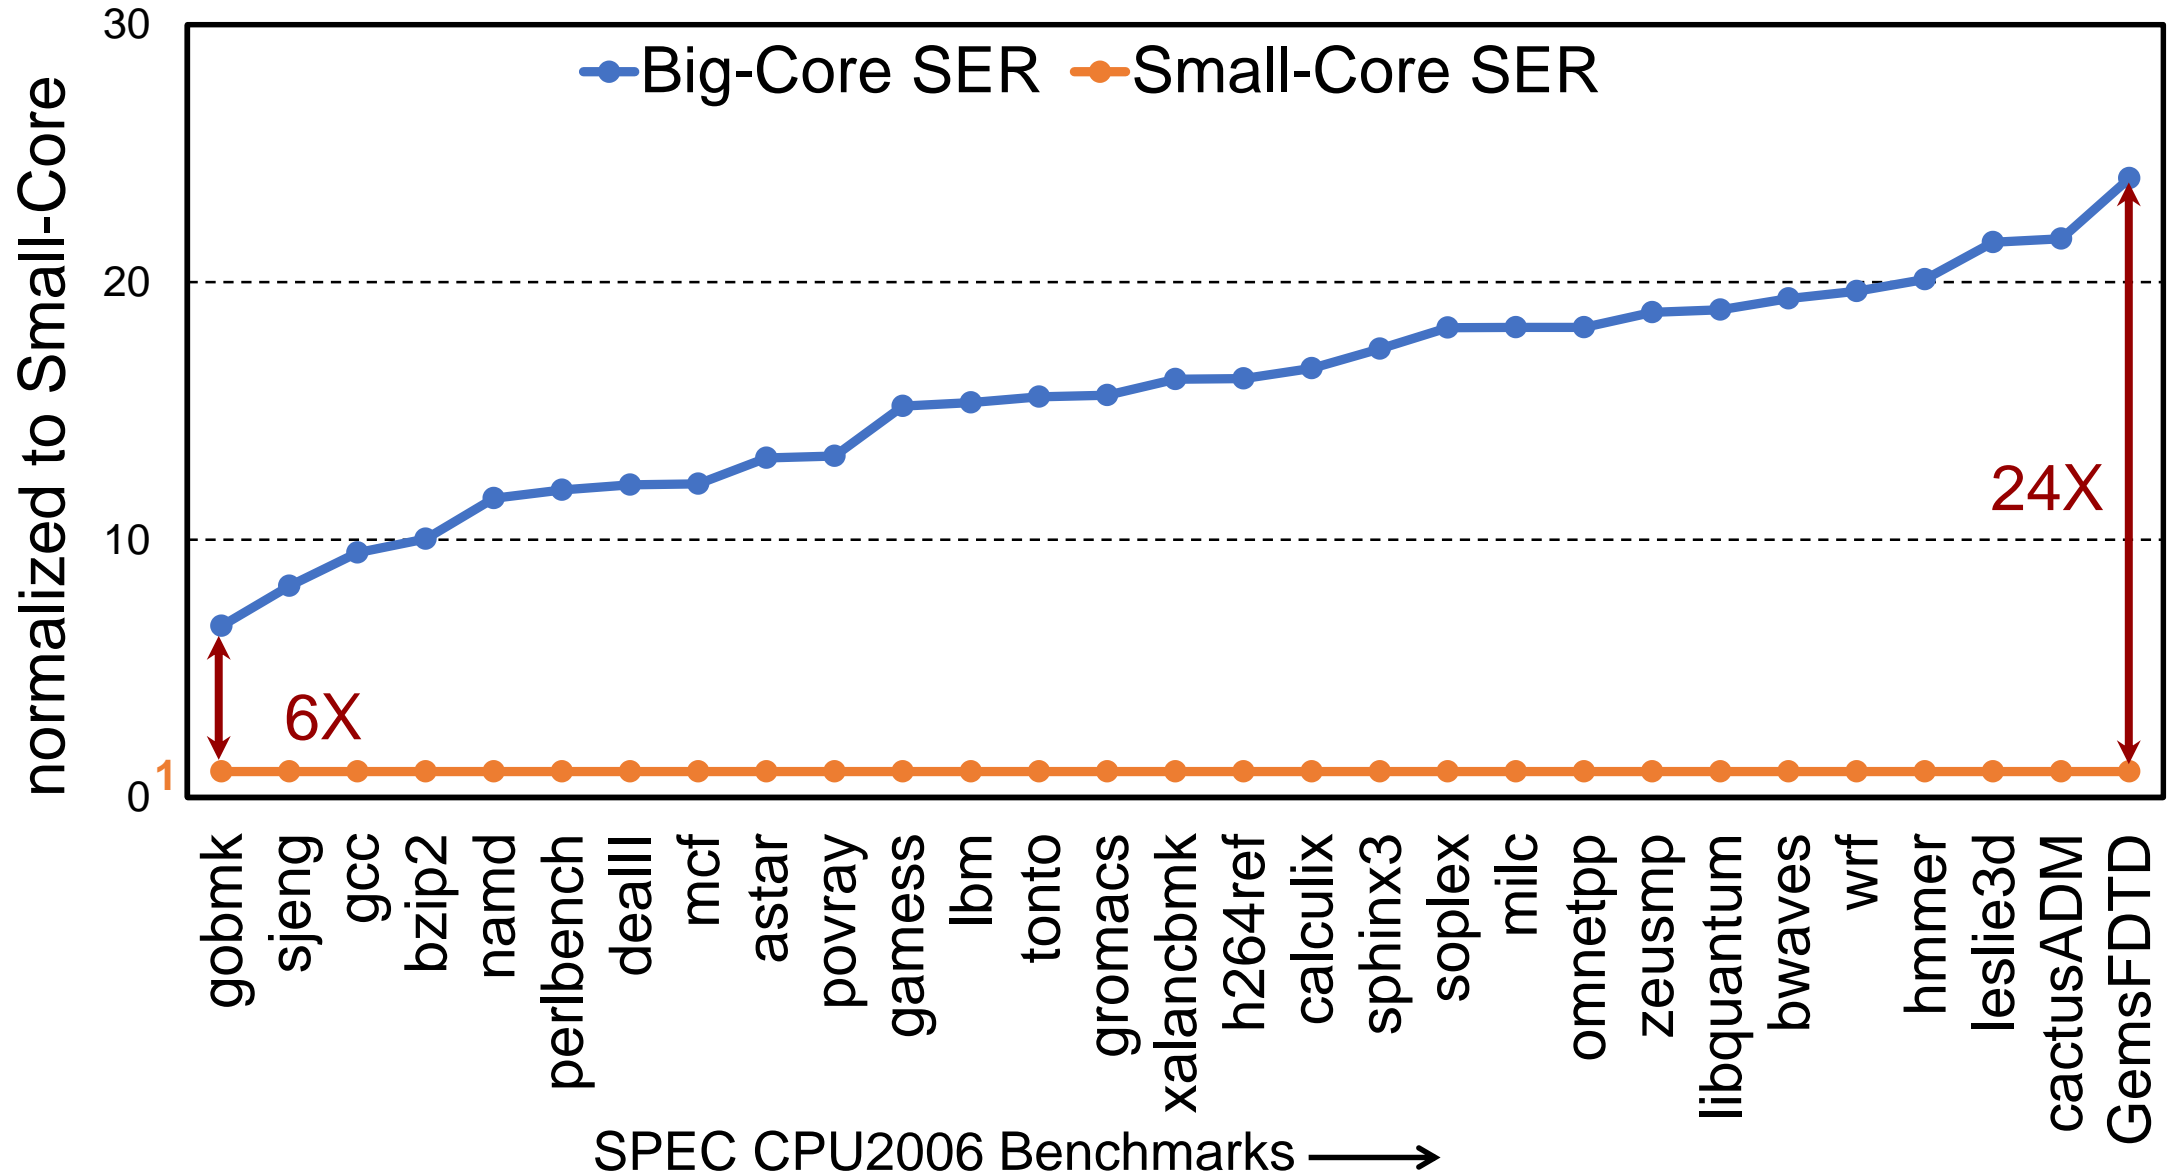
Vulnerability vs. Core Type



Vulnerability vs. Core Type

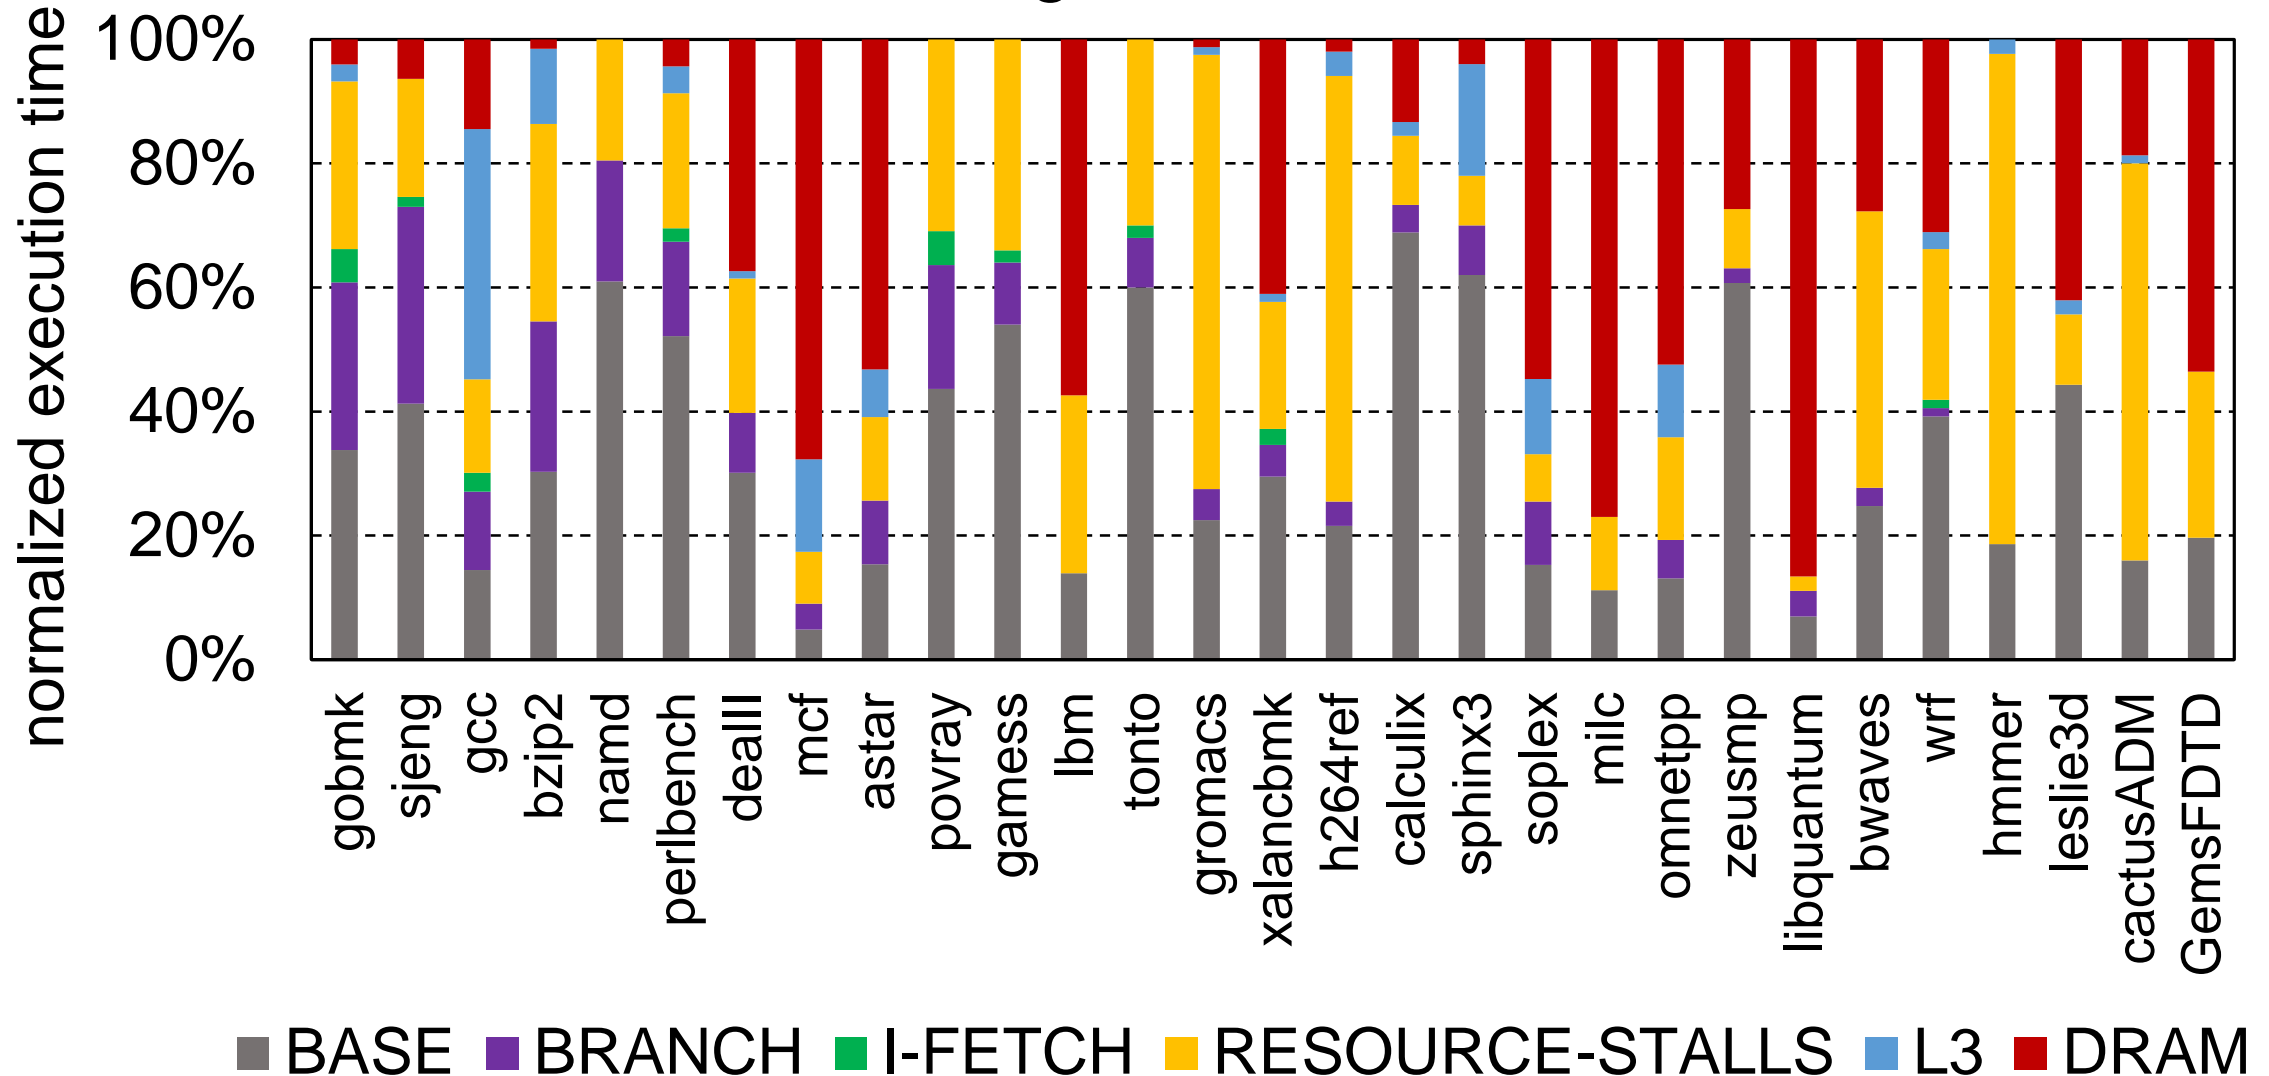


Vulnerability vs. Core Type



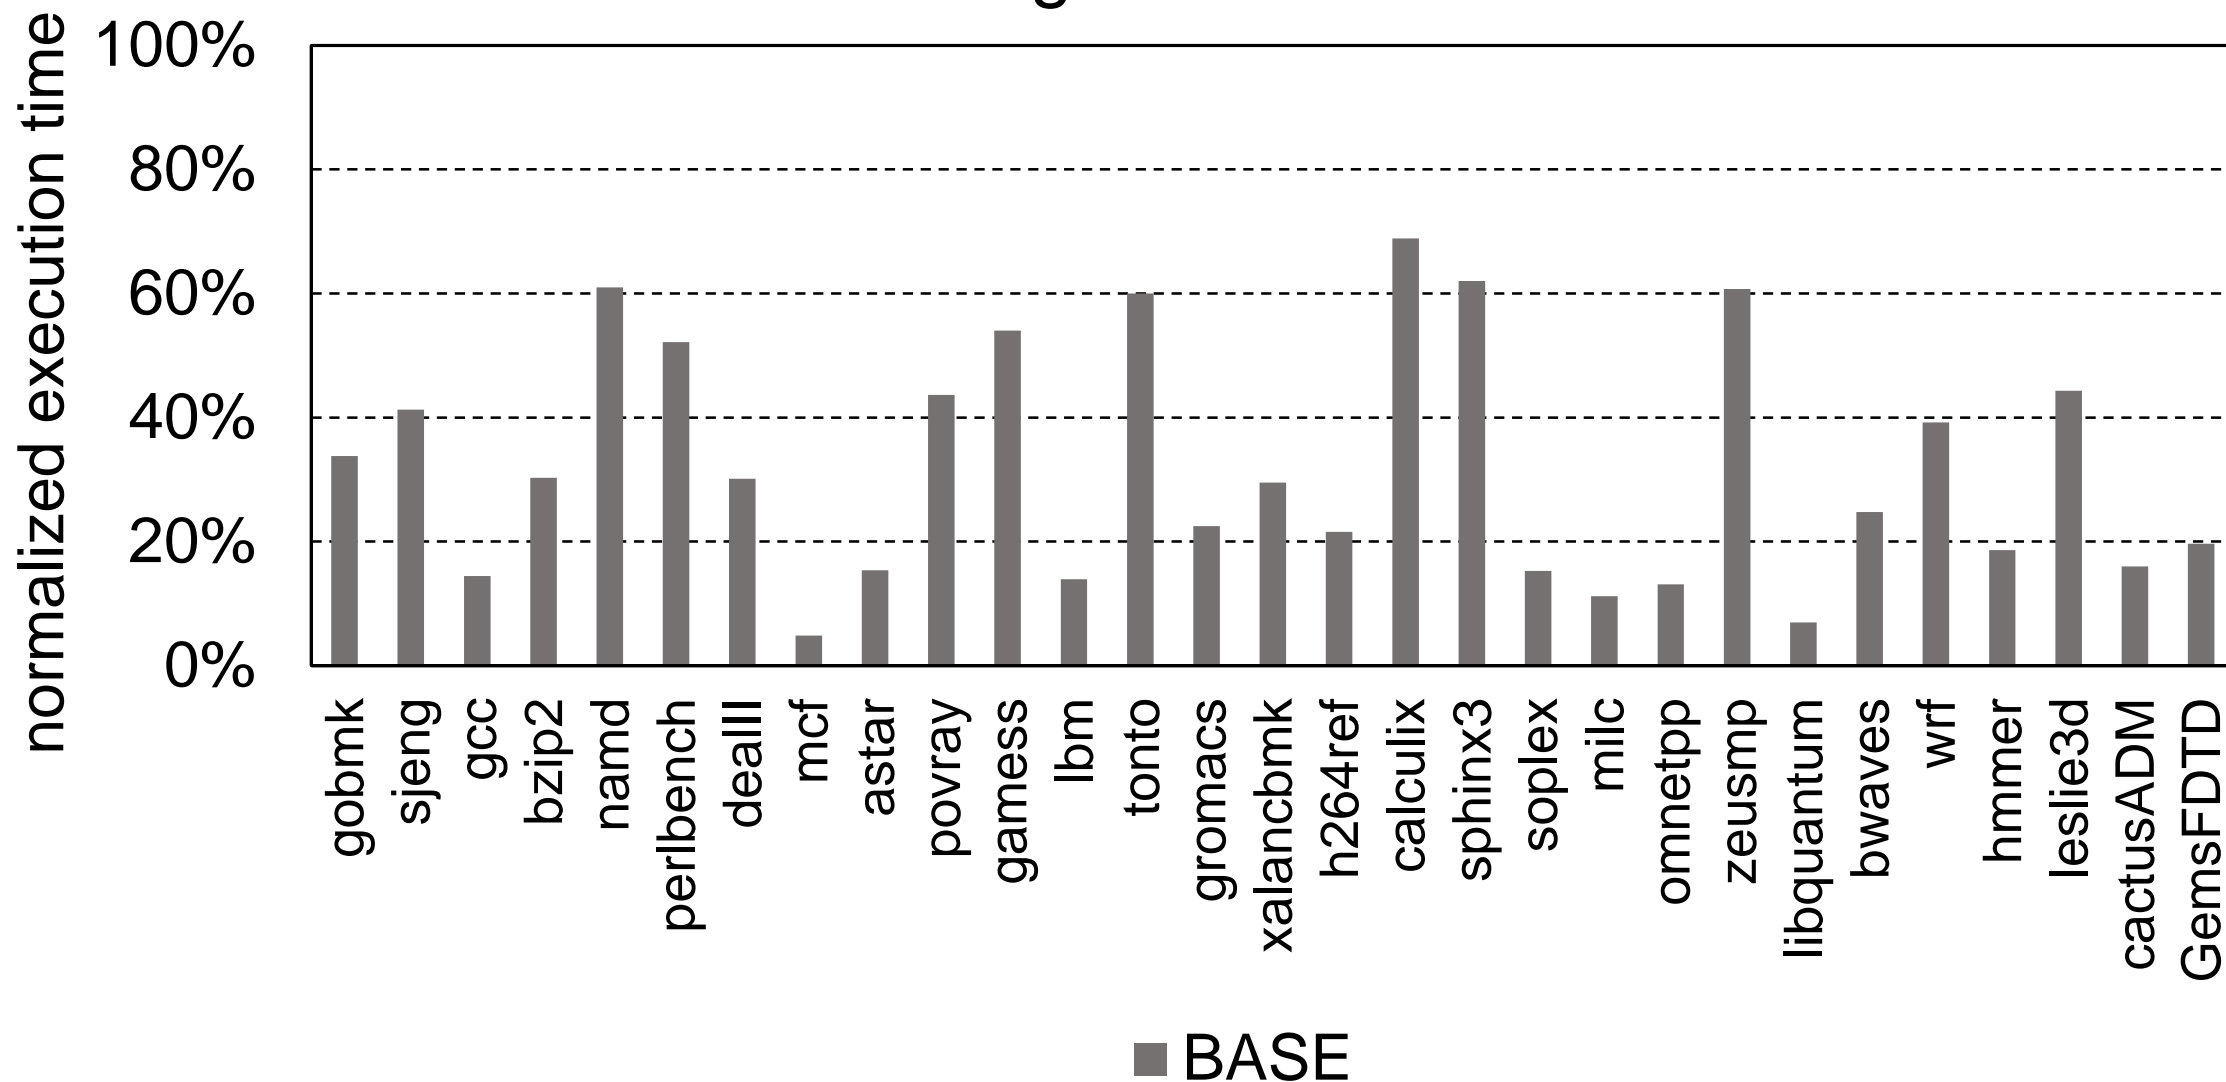
Vulnerability vs. CPI Stacks

Increasing SER →



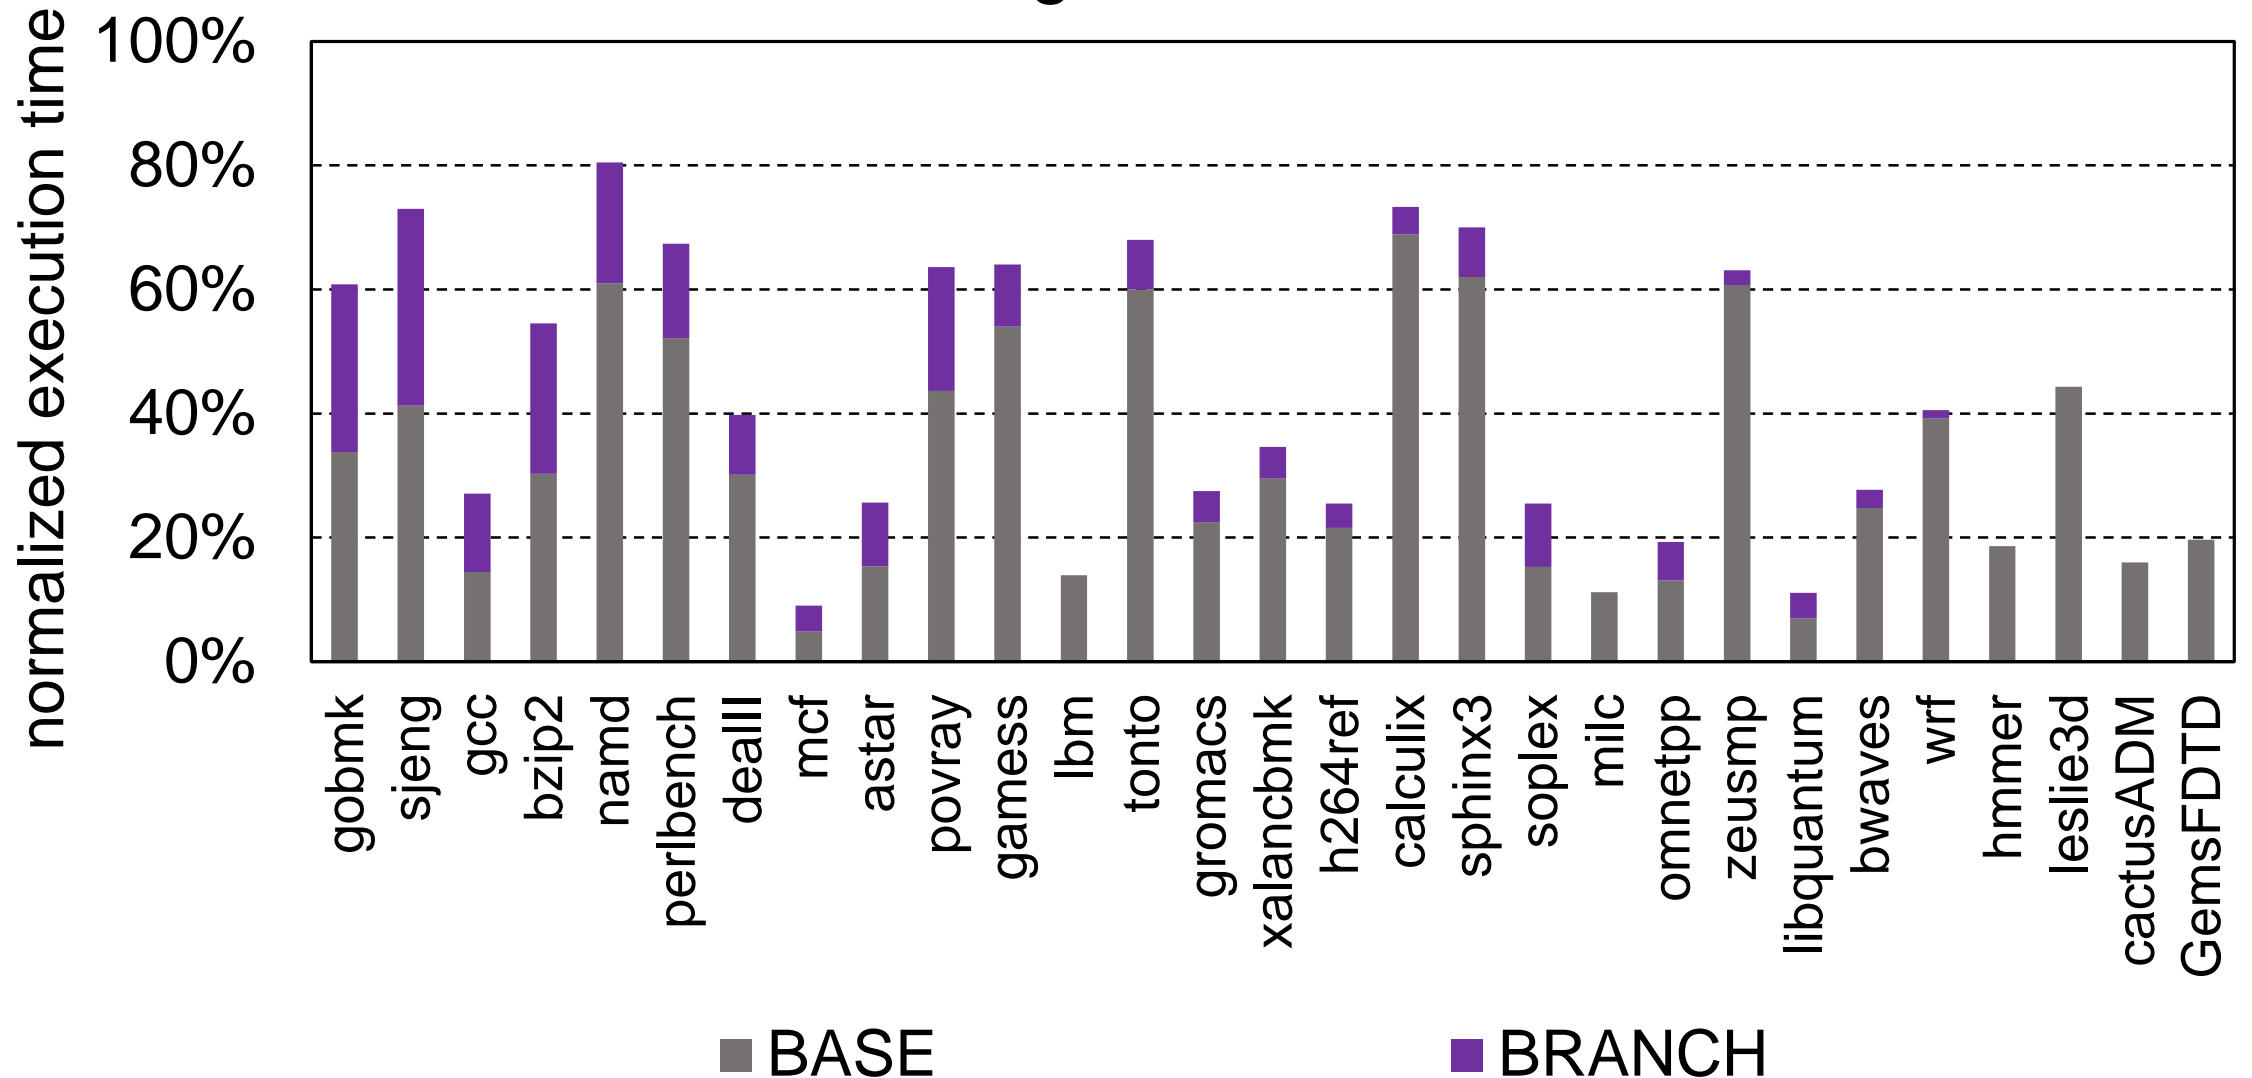
Vulnerability vs. CPI Stacks

Increasing SER →



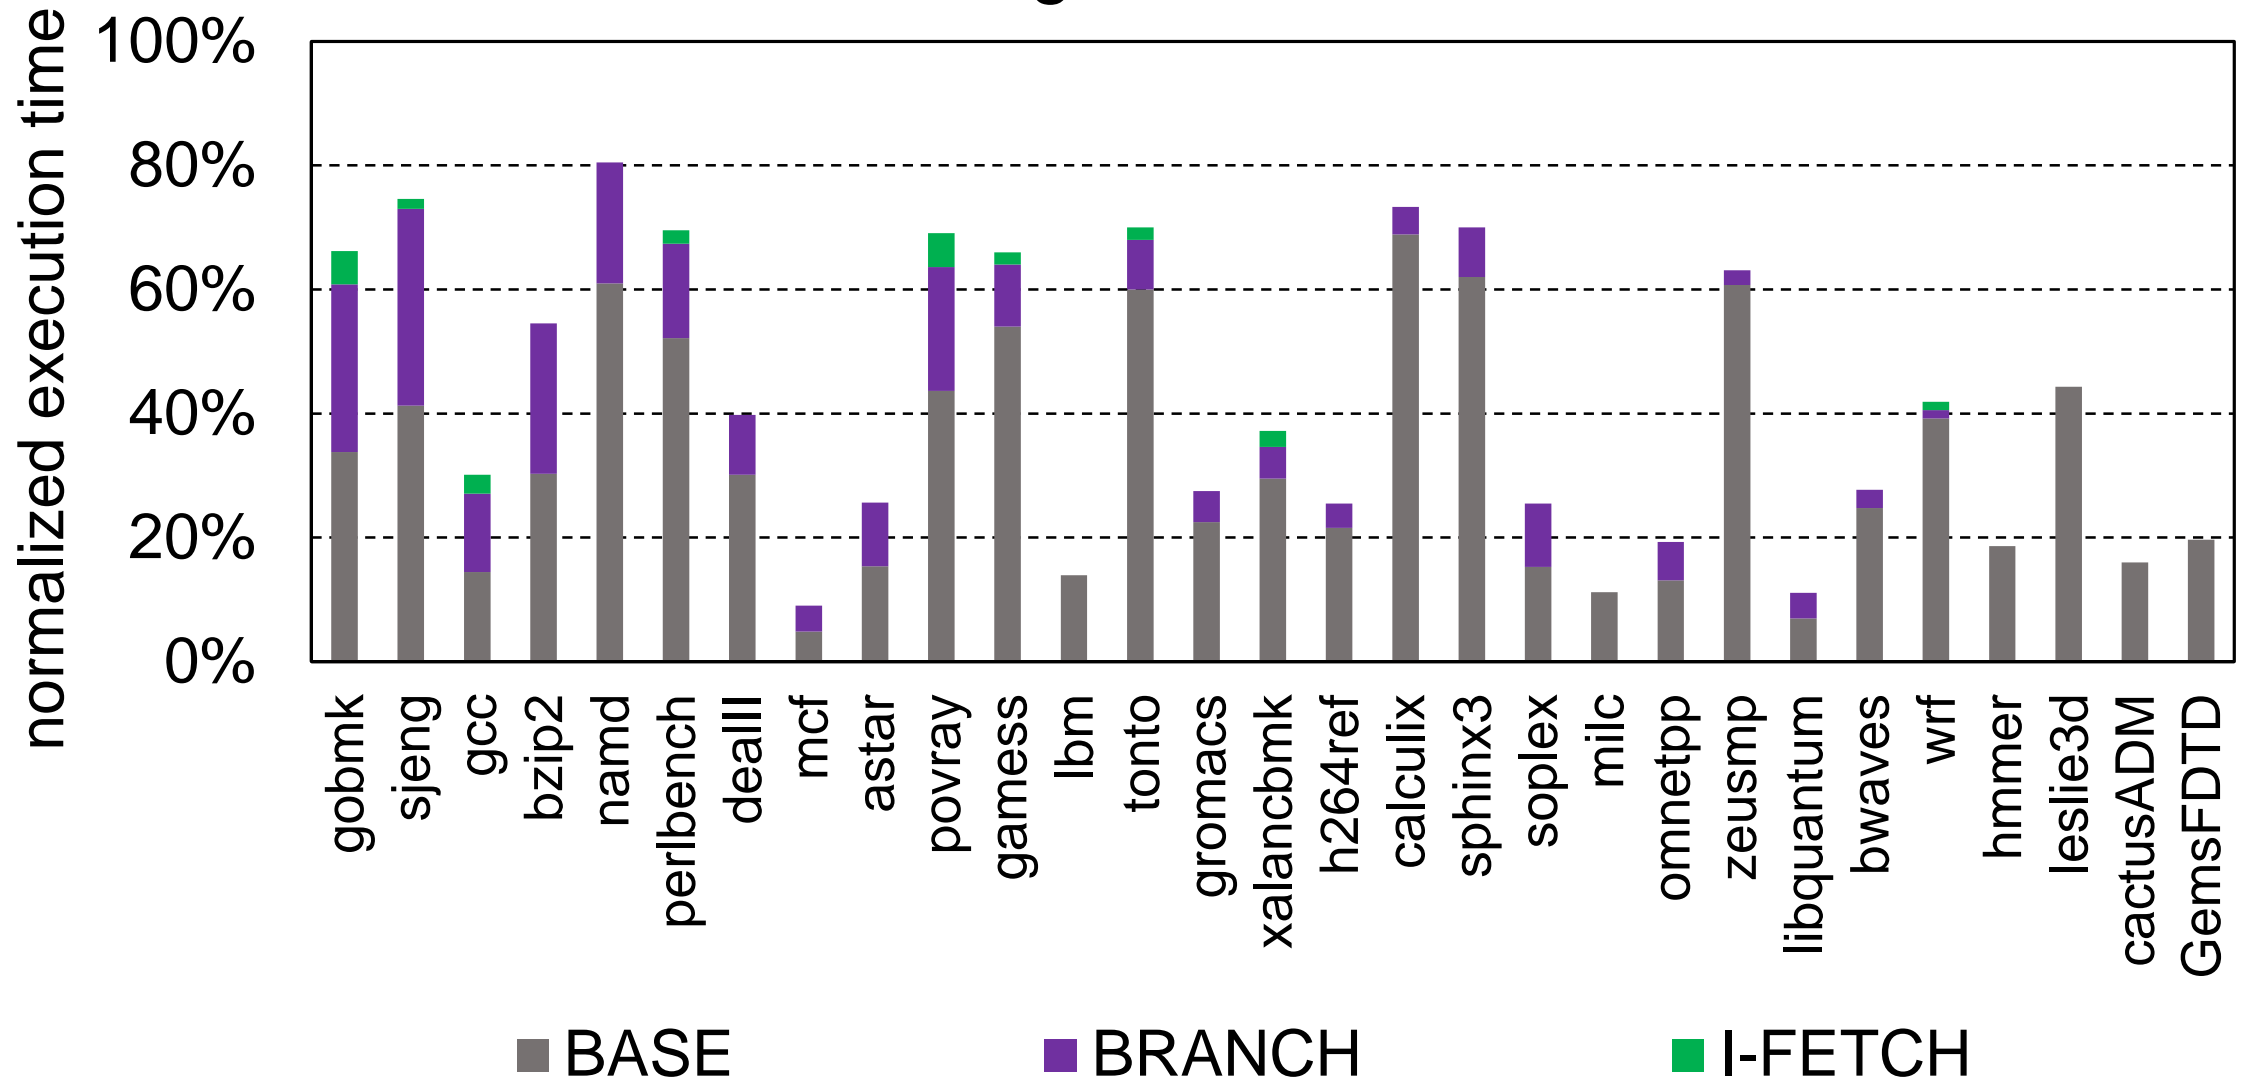
Vulnerability vs. CPI Stacks

Increasing SER →



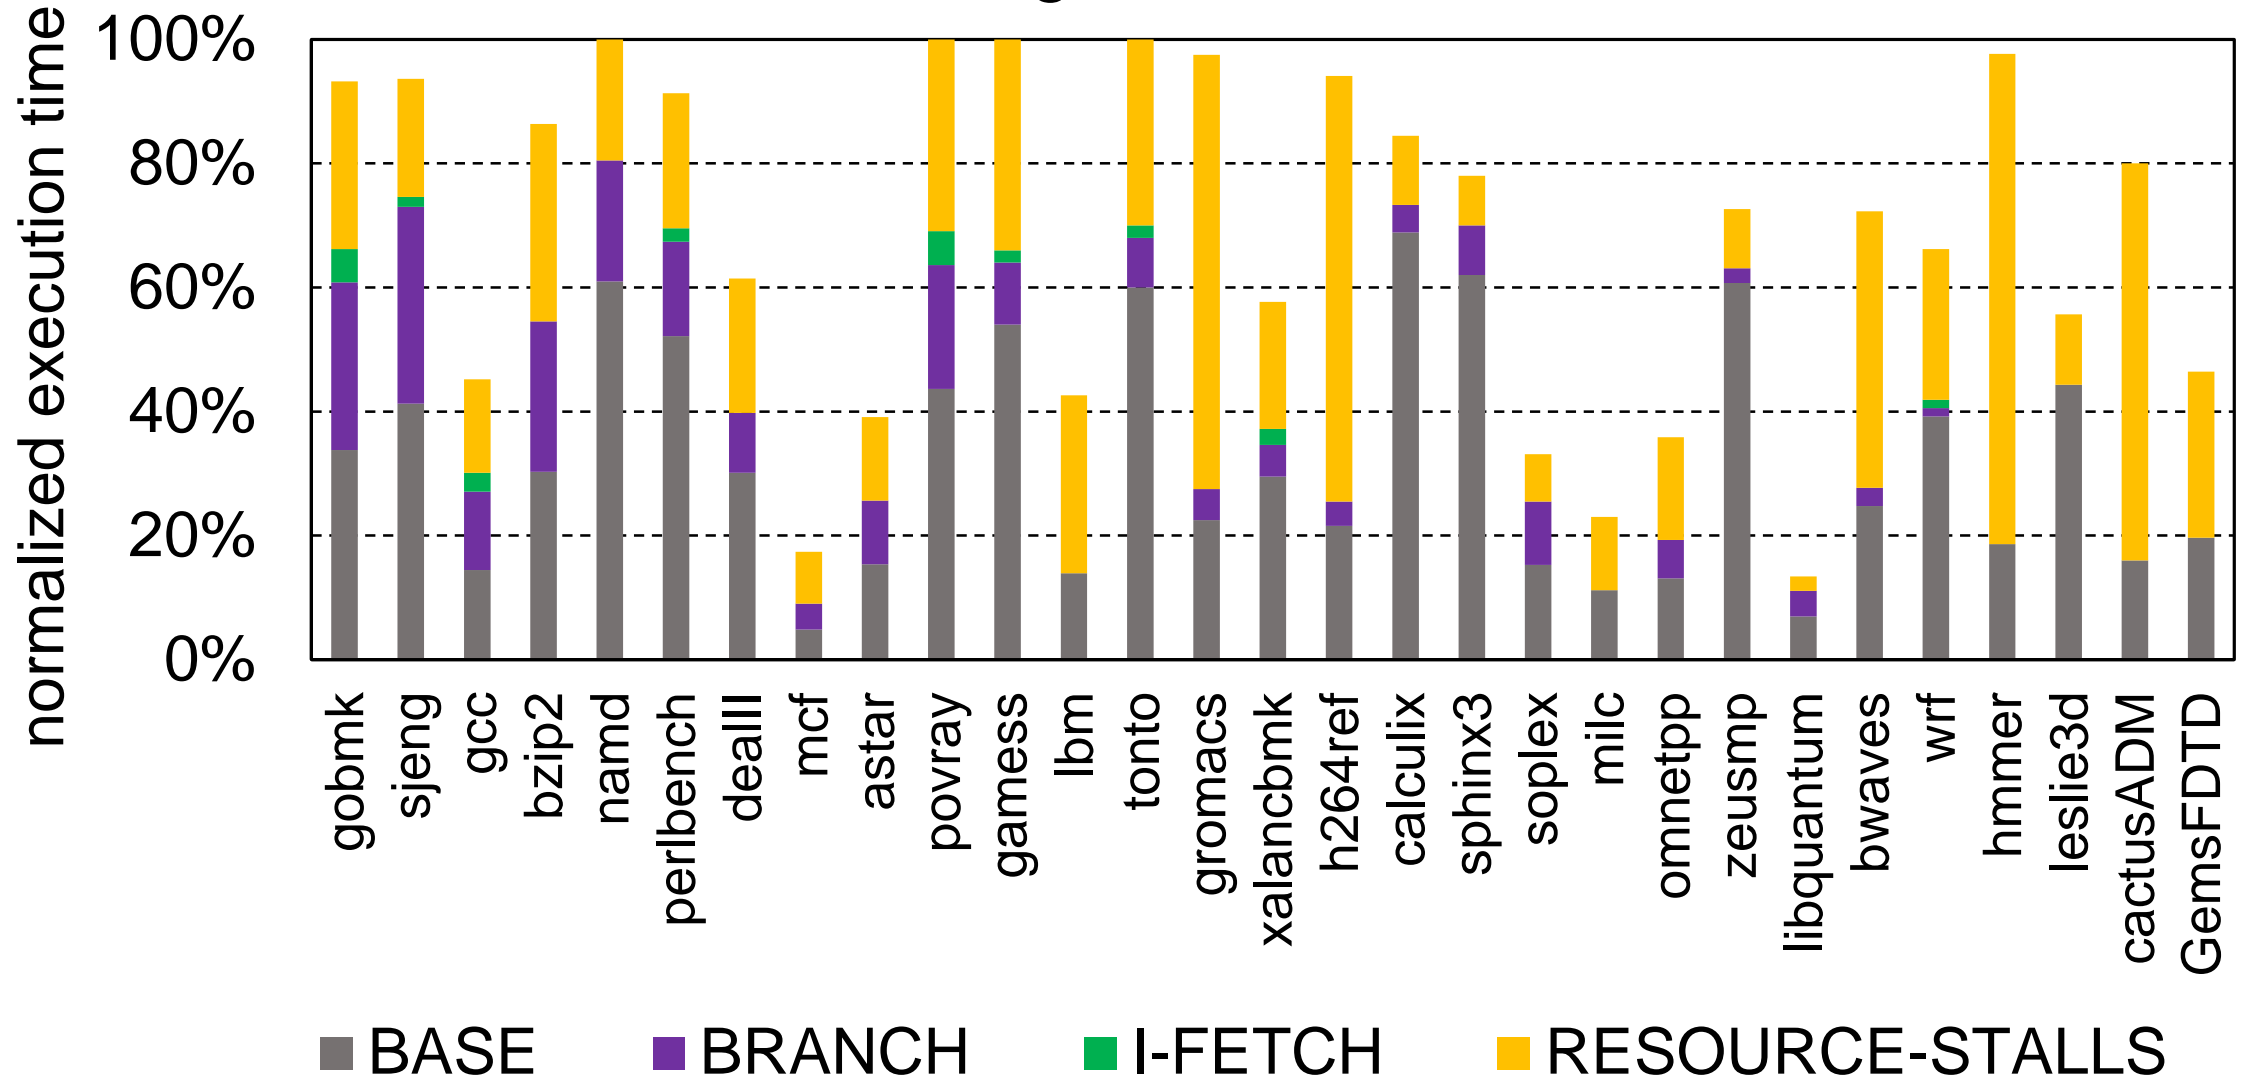
Vulnerability vs. CPI Stacks

Increasing SER →



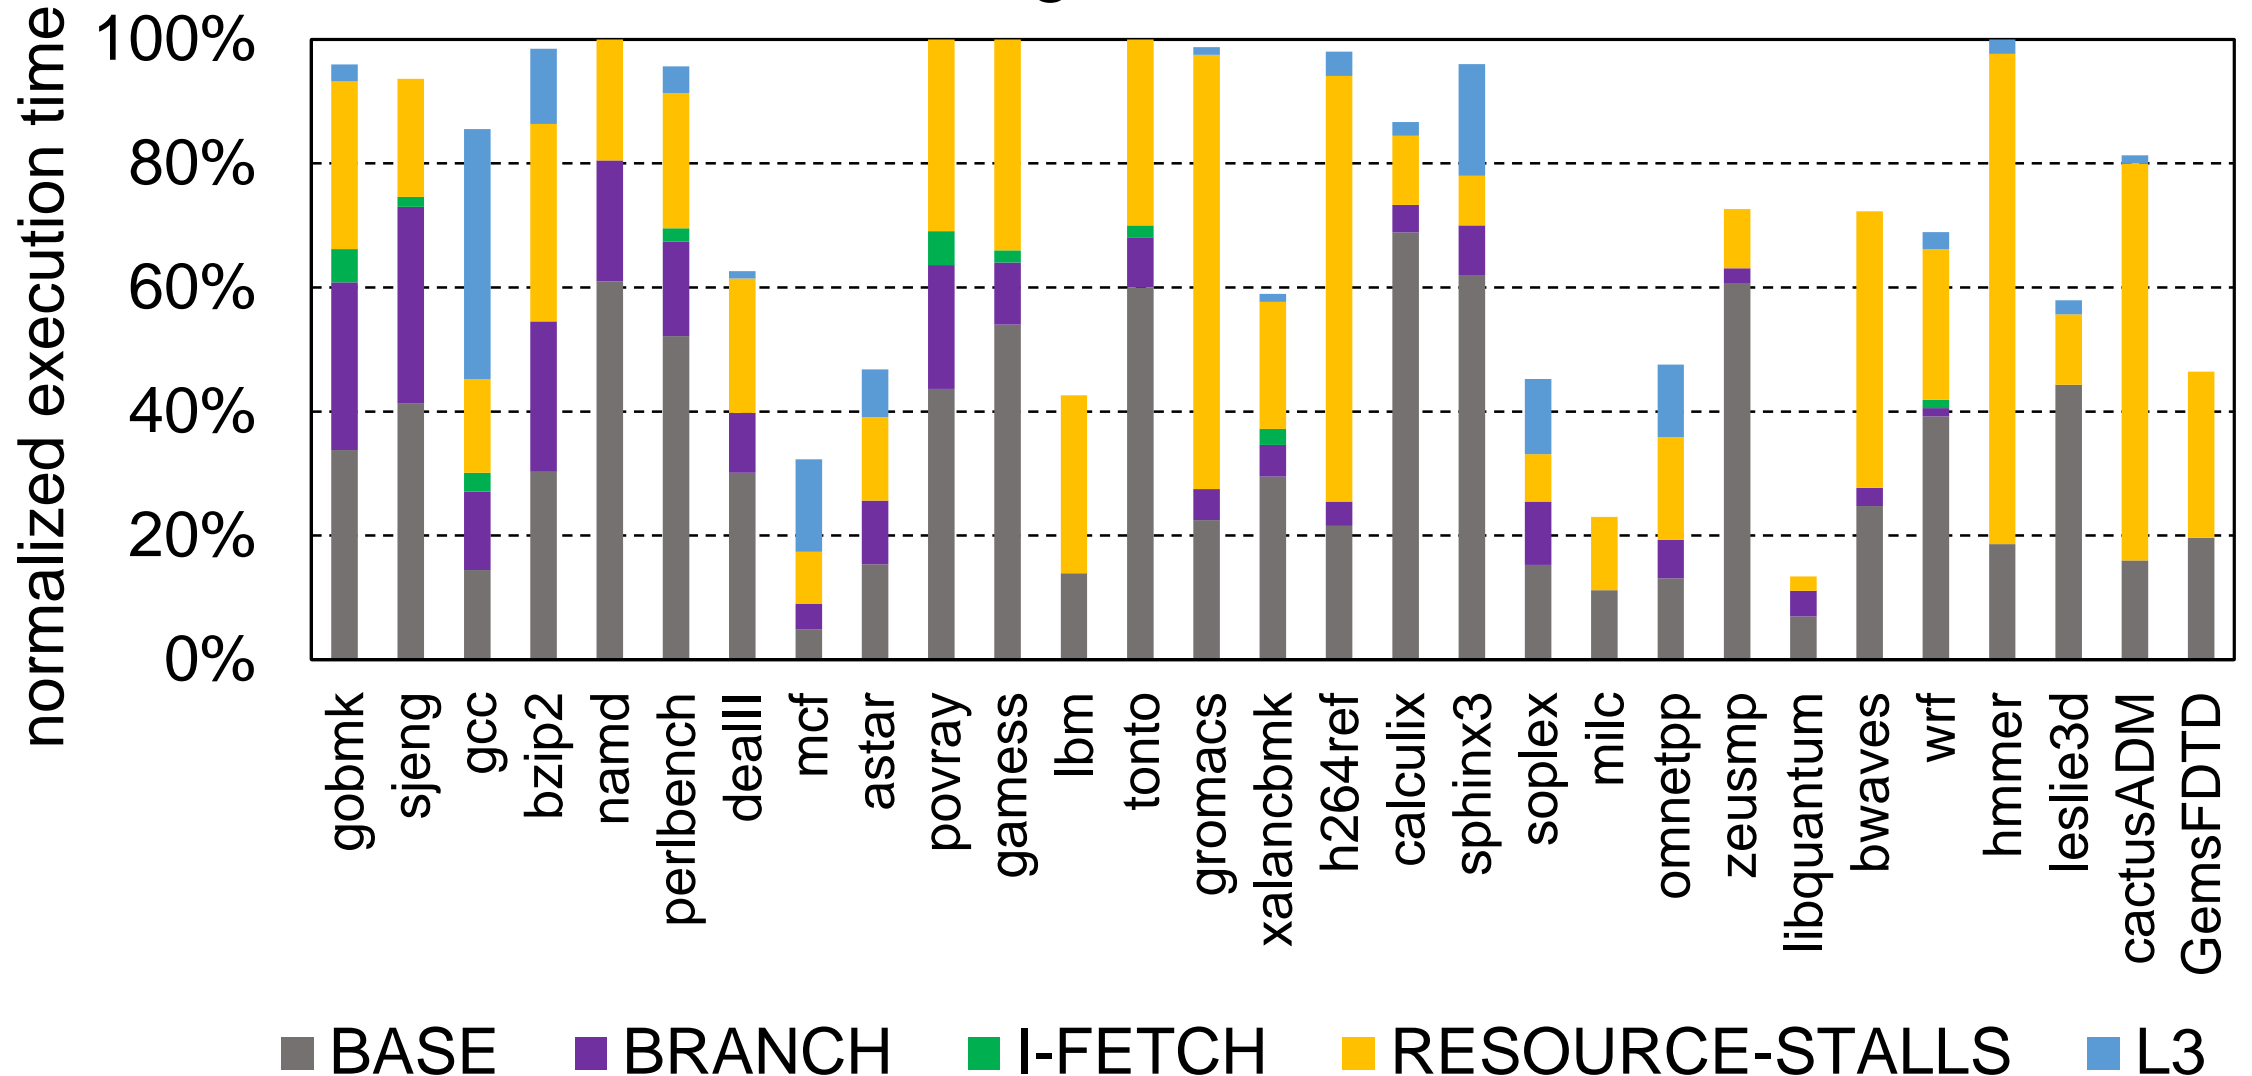
Vulnerability vs. CPI Stacks

Increasing SER →



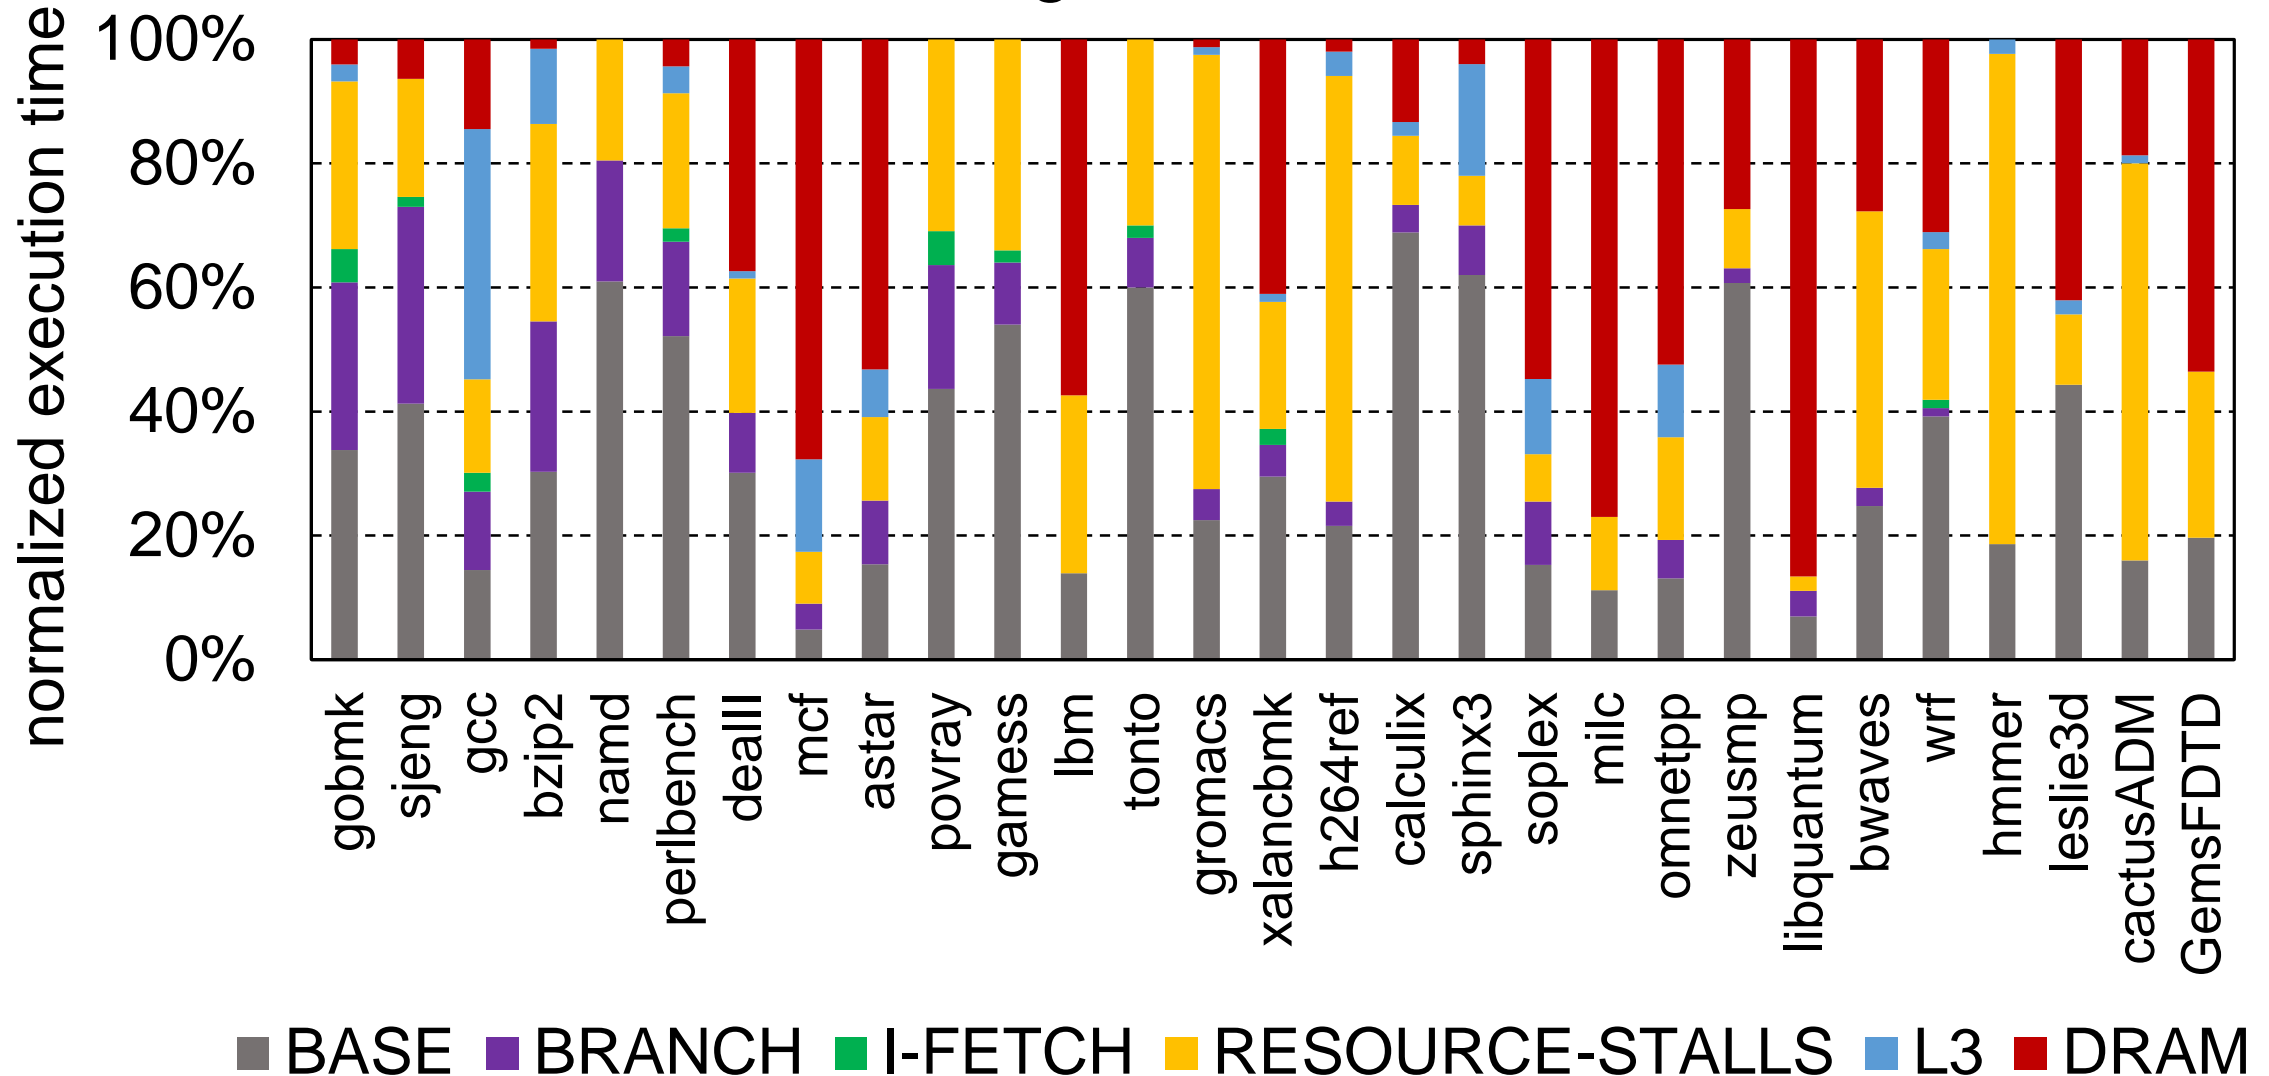
Vulnerability vs. CPI Stacks

Increasing SER →



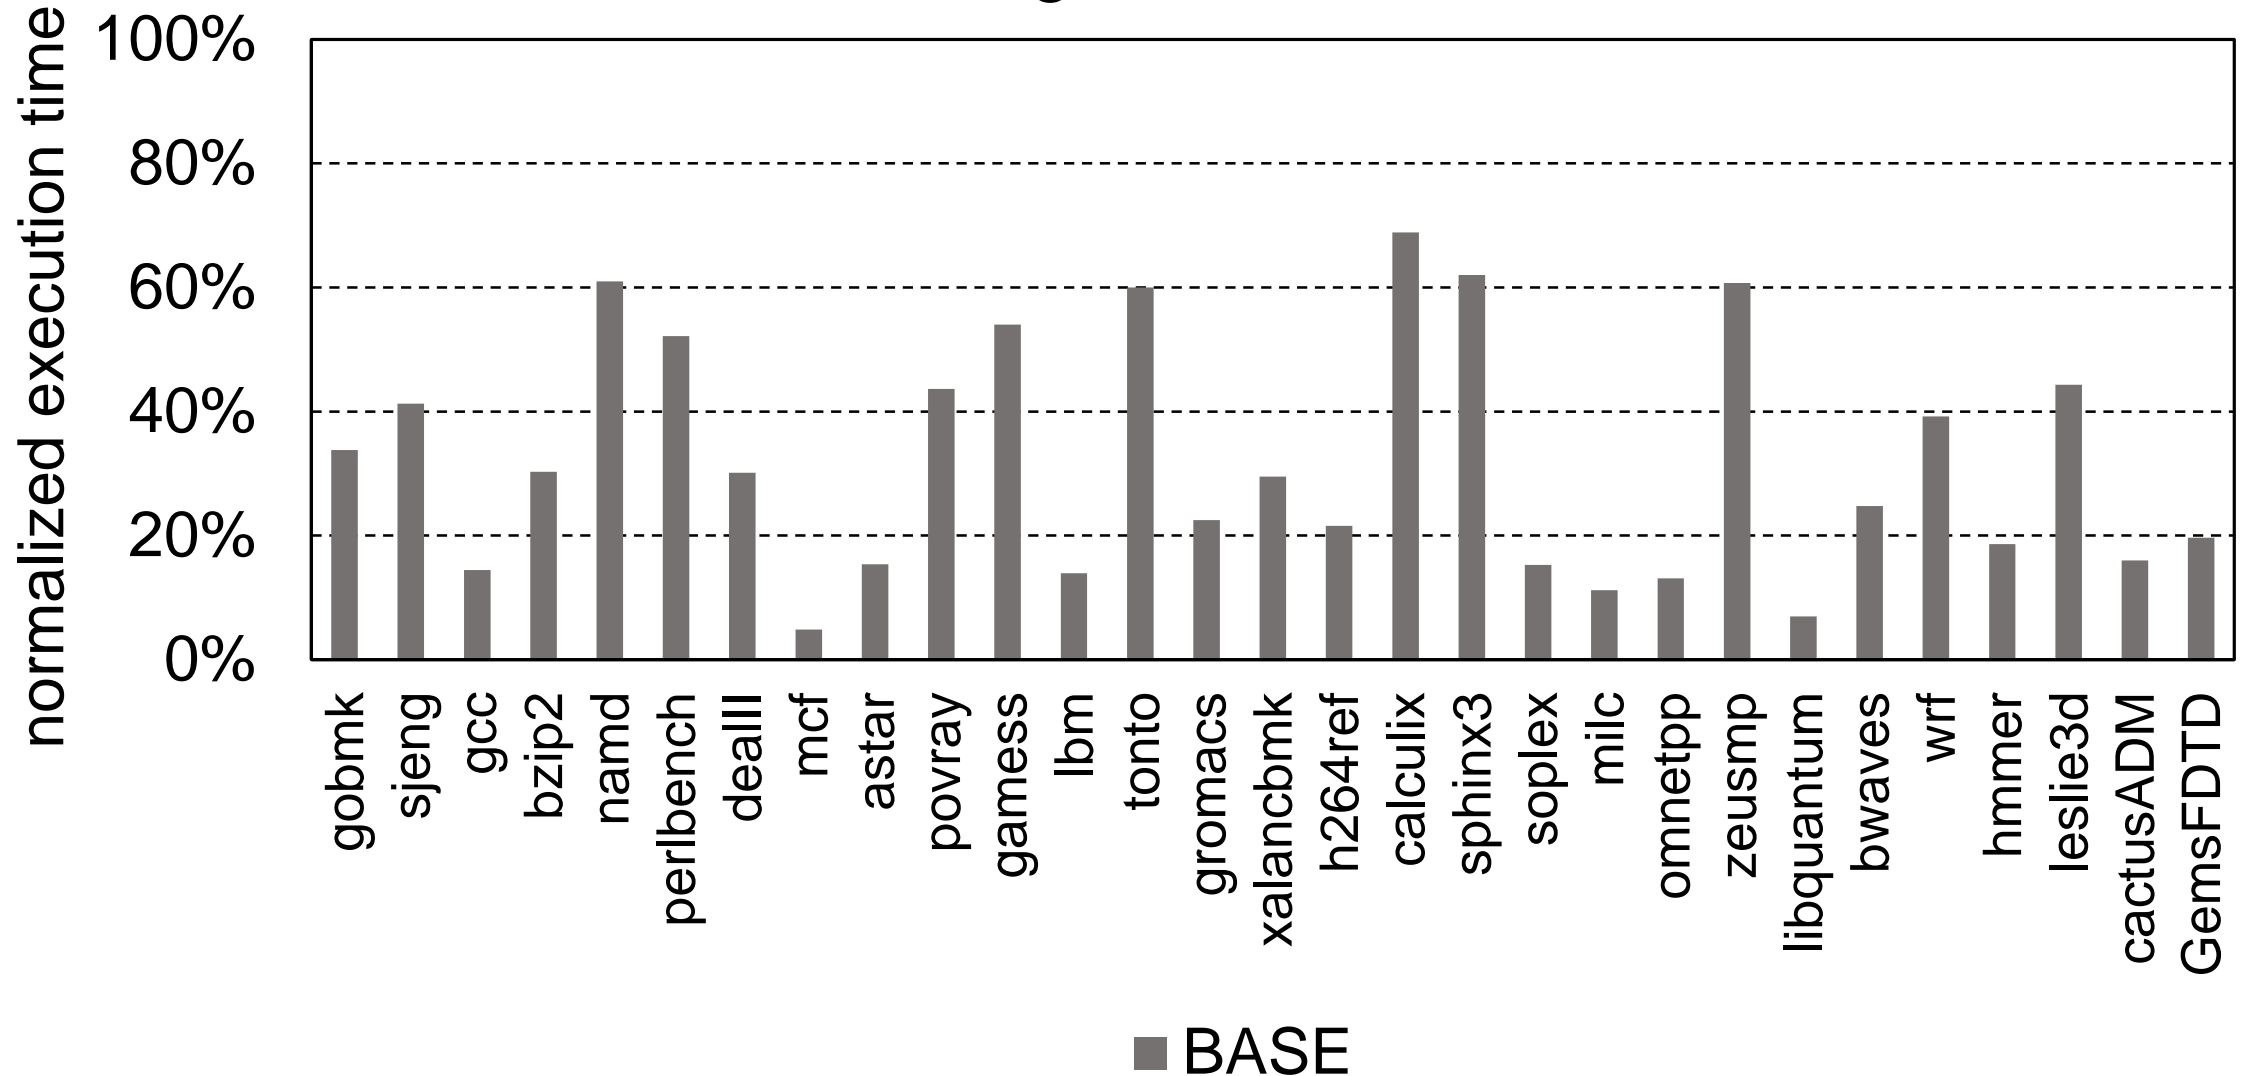
Vulnerability vs. CPI Stacks

Increasing SER →



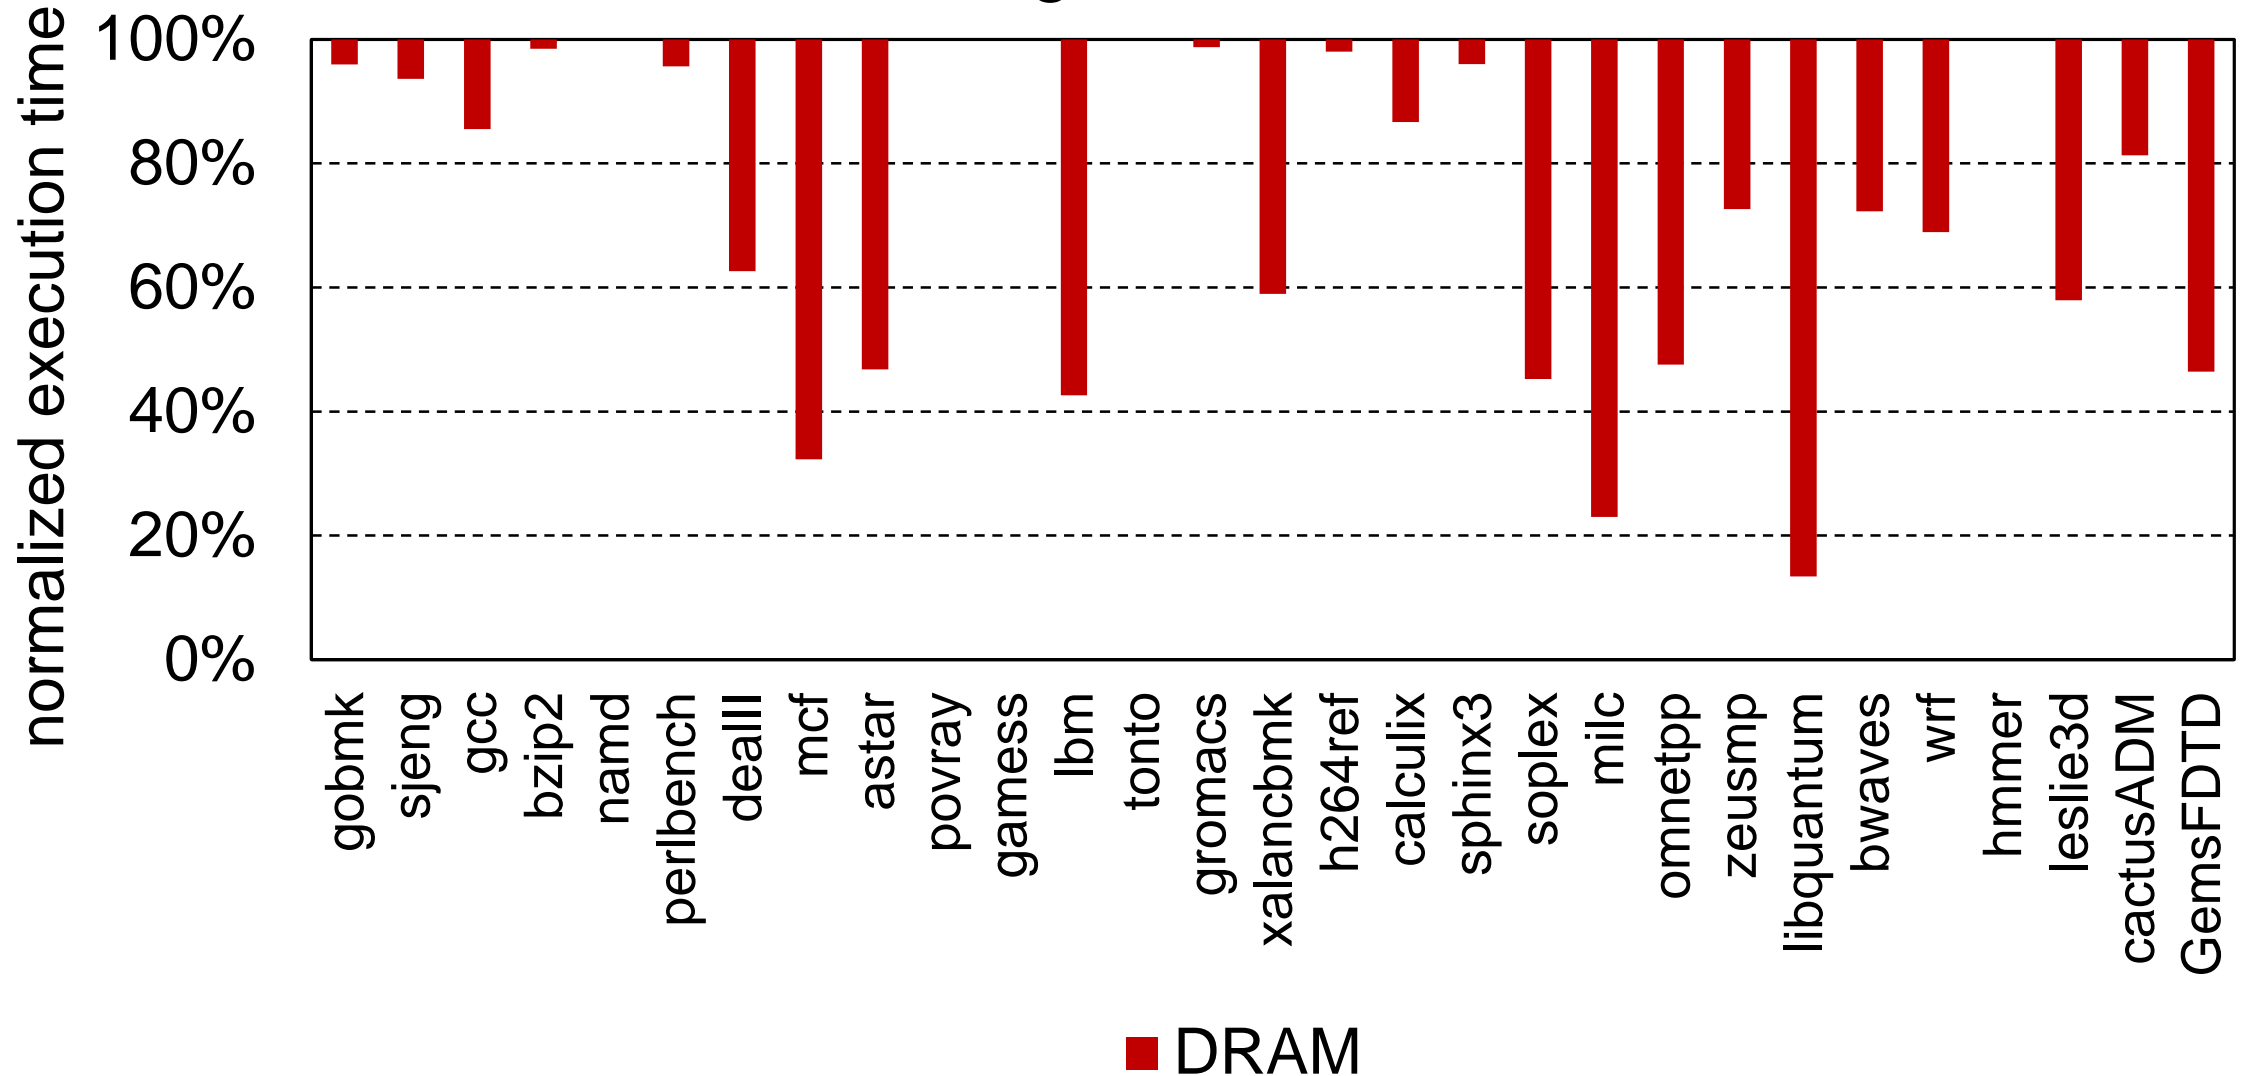
Vulnerability vs. Compute-Intensity

Increasing SER →



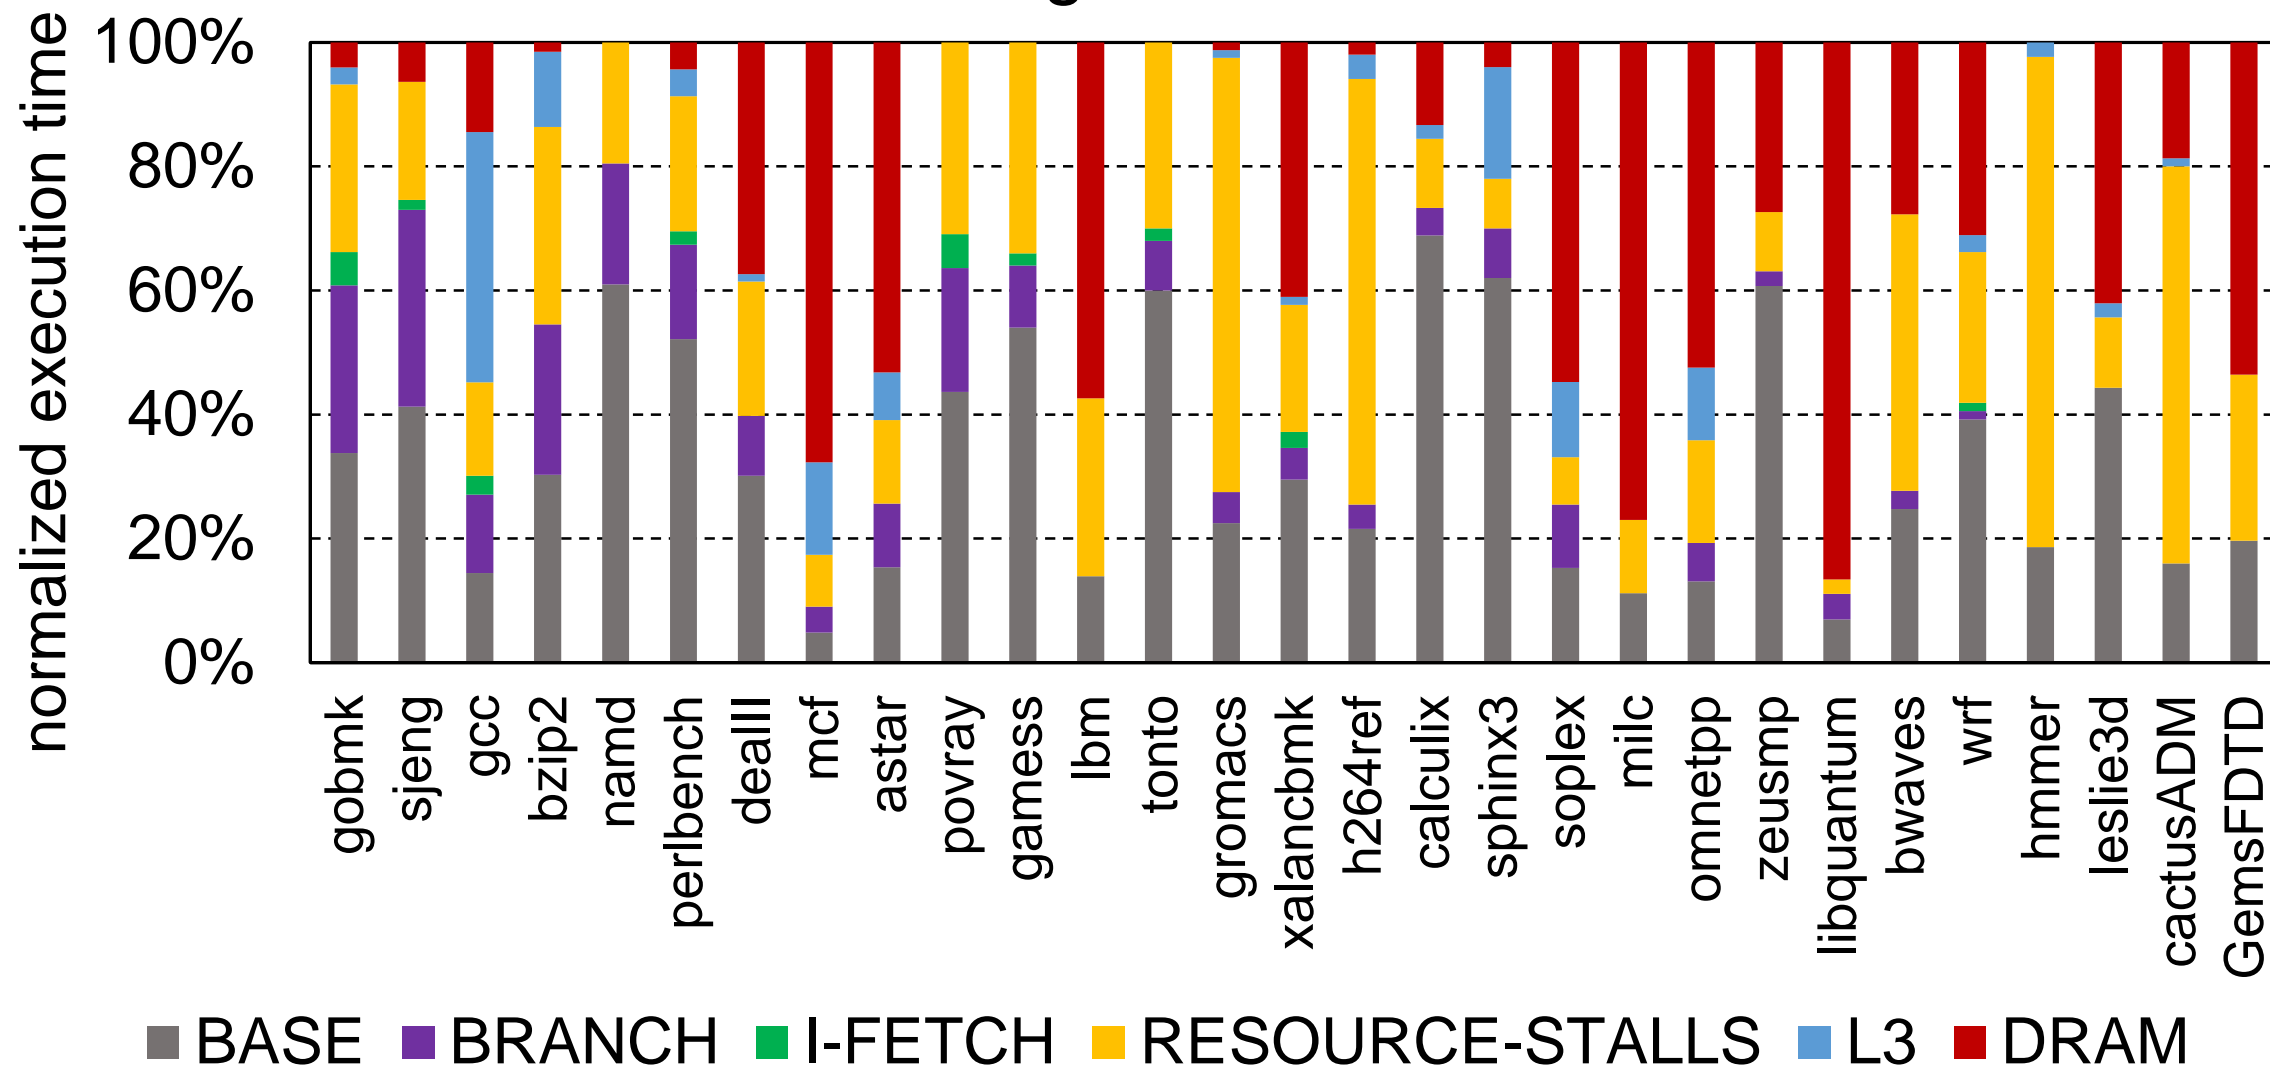
Vulnerability vs. Memory-Intensity

Increasing SER →



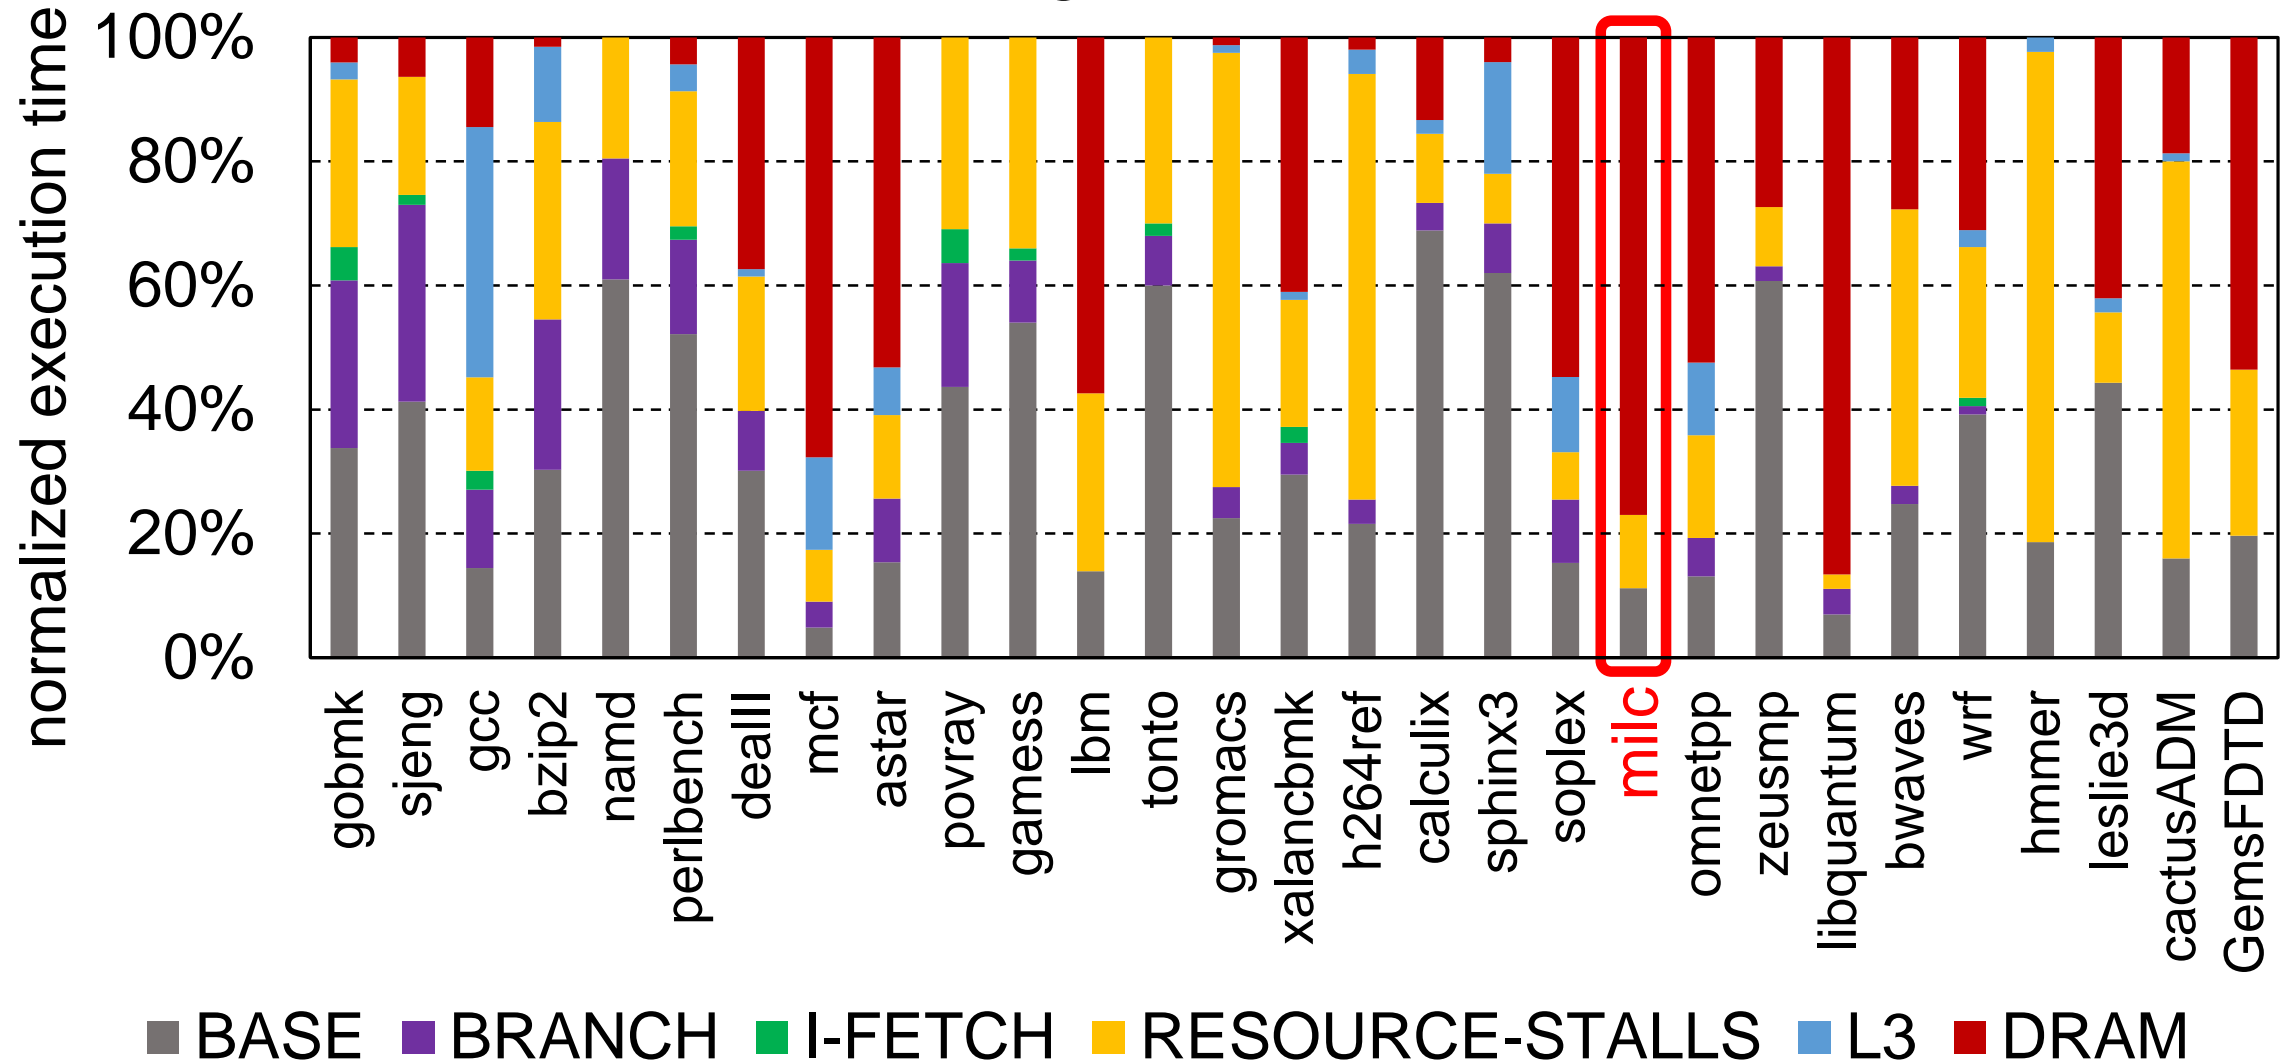
Vulnerability vs. CPI Stacks

Increasing SER →



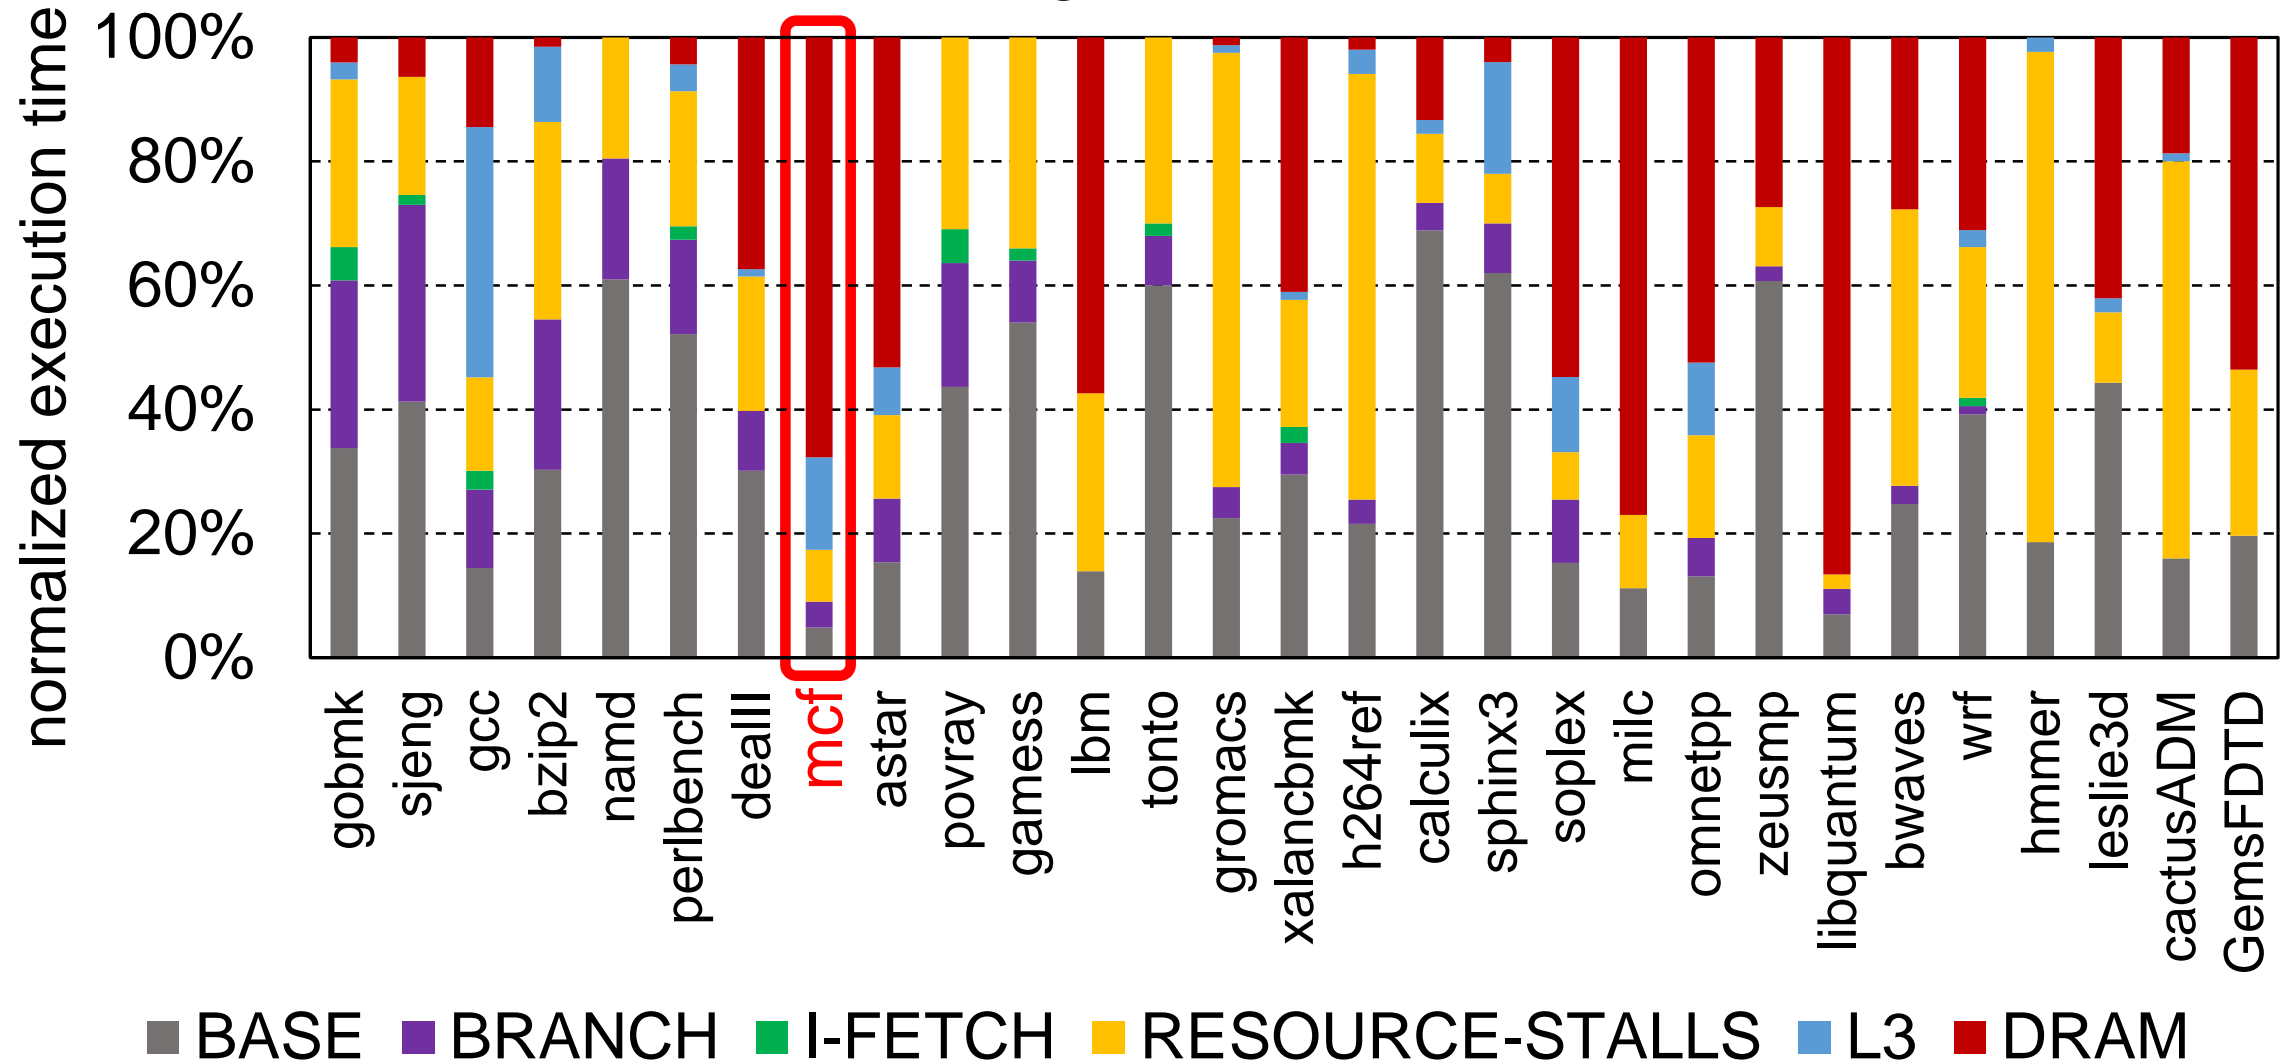
Vulnerability vs. CPI Stacks

Increasing SER →



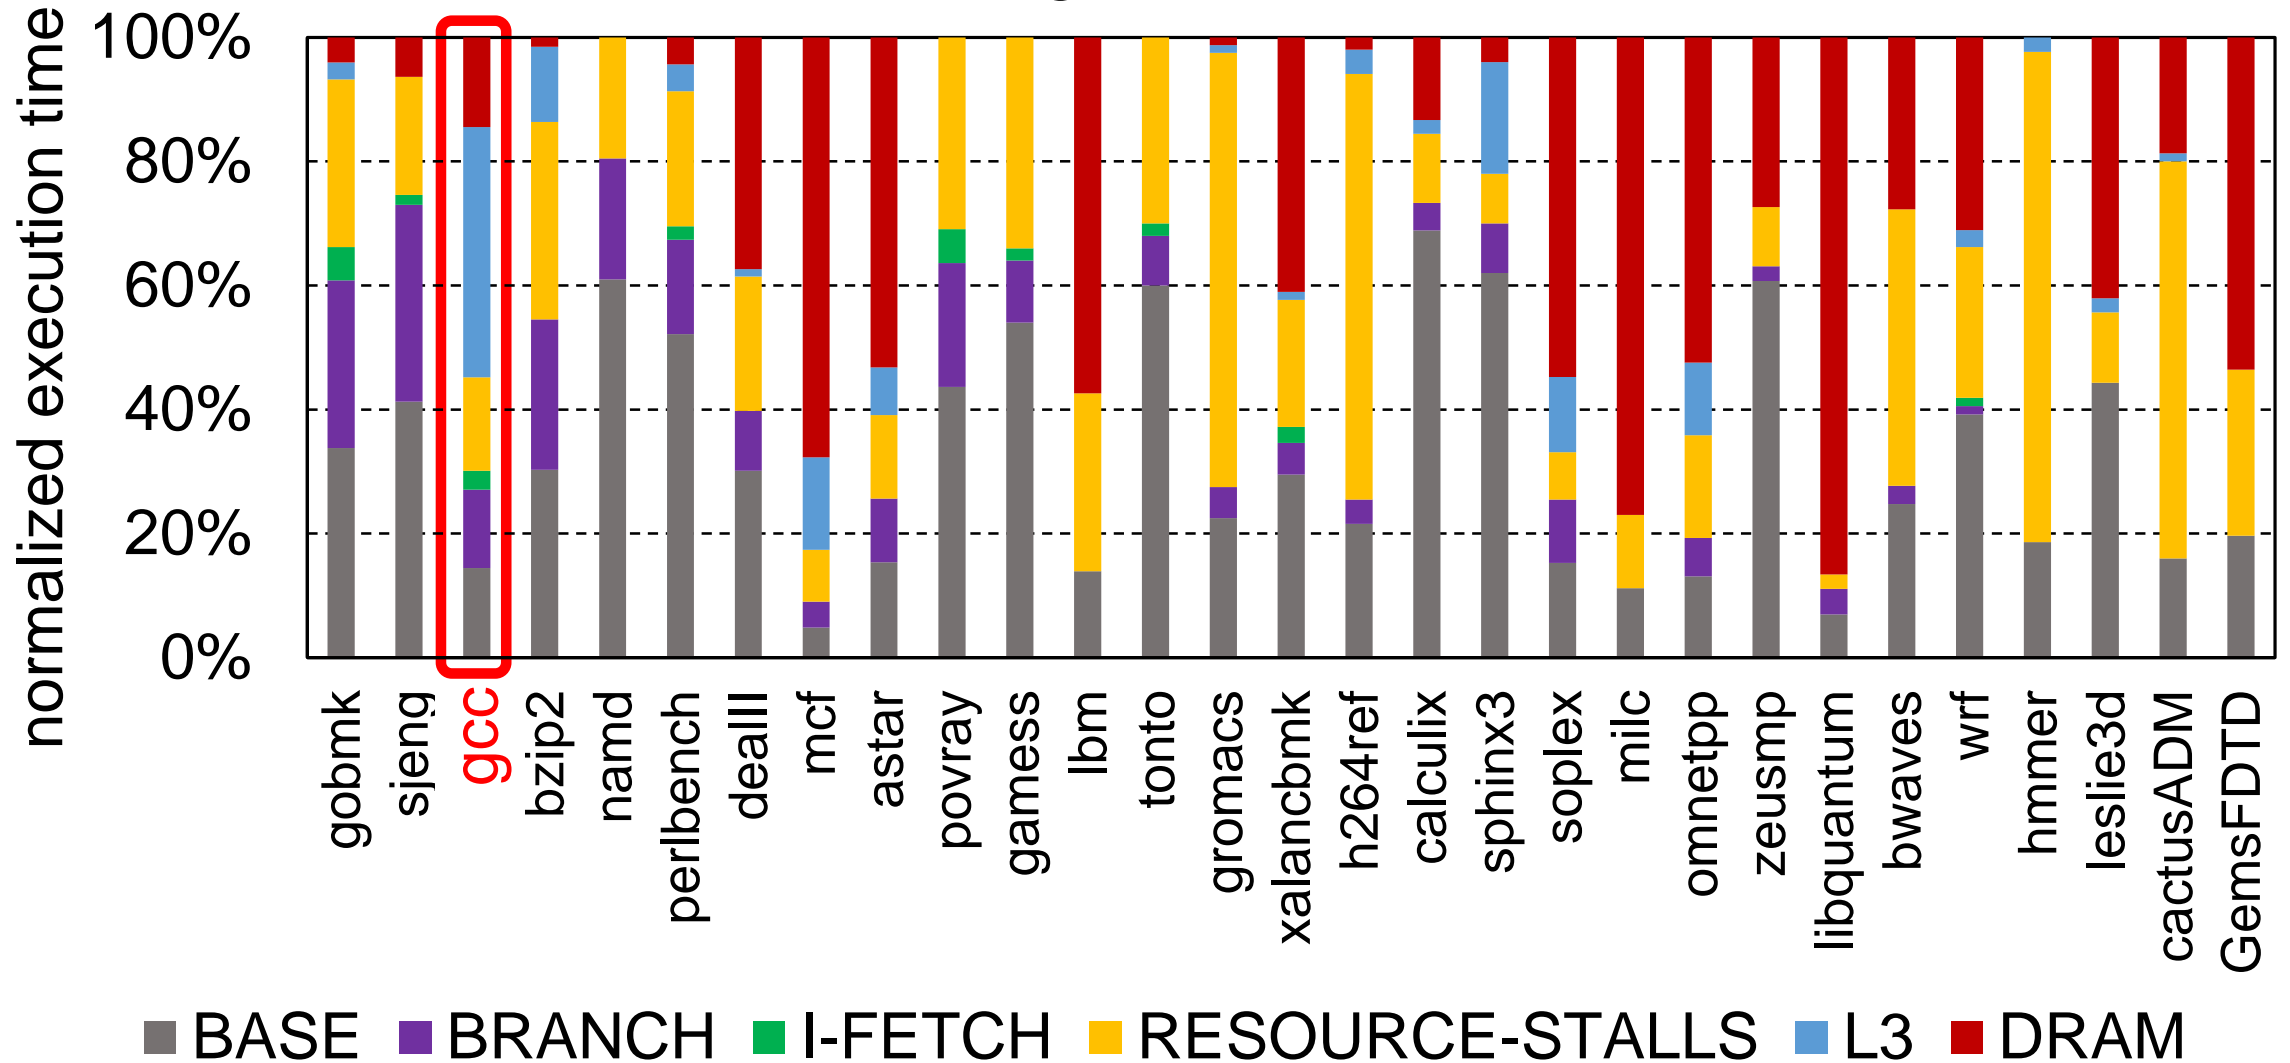
Vulnerability vs. CPI Stacks

Increasing SER →



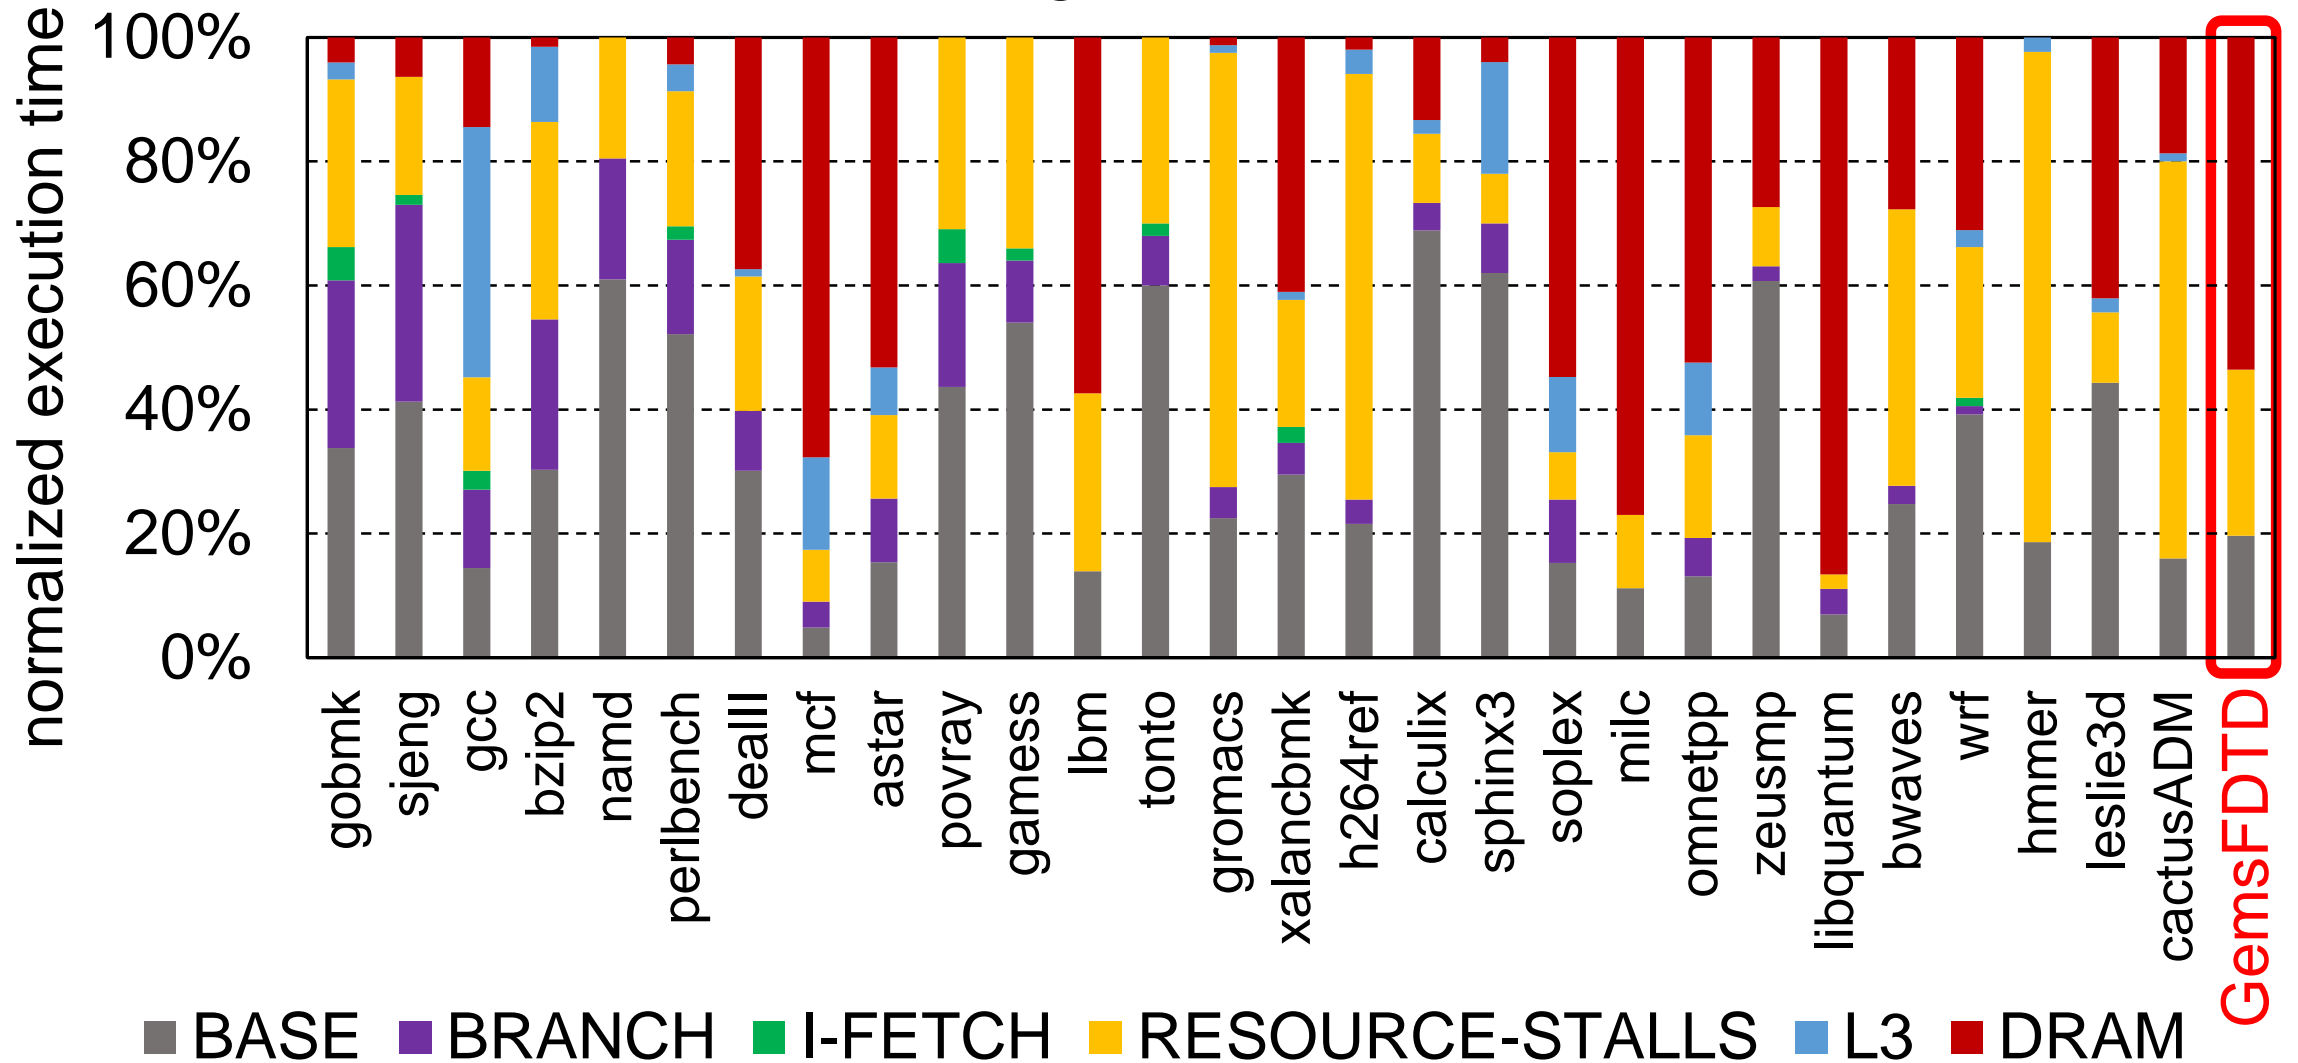
Vulnerability vs. CPI Stacks

Increasing SER →



Vulnerability vs. CPI Stacks

Increasing SER →



Reliability vs. Application Type

Reliability vs. Application Type

- No simple workload characteristic
 - For example, **memory-intensity** or **compute-intensity**

Reliability vs. Application Type

- No simple workload characteristic
 - For example, **memory-intensity** or **compute-intensity**
- Predicting SER is not straightforward*
 - **Complex interaction** of various workload characteristics

*[Nair et al. ISCA'12]

Reliability vs. Application Type

- No simple workload characteristic
 - For example, **memory-intensity** or **compute-intensity**
- Predicting SER is not straightforward*
 - **Complex interaction** of various workload characteristics
- A **dynamic mechanism** required
 - To monitor reliability on either core type

*[Nair et al. ISCA'12]

Reliability vs. Application Type

- No simple workload characteristic
 - For example, **memory-intensity** or **compute-intensity**
- Predicting **Reliability-Aware Scheduling** characteristics
 - **Complex**
- A **dynamic**
 - To monitor reliability on either core type

Reliability Metric for Multiprogram Workloads

Reliability Metric for Multiprogram Workloads

In **isolation** on one core:

Reliability Metric for Multiprogram Workloads

In **isolation** on one core:

A Time = 1

SER = 4

Reliability Metric for Multiprogram Workloads

In **isolation** on one core:

A Time = 1

SER = 4

B Time = 1

SER = 1

Reliability Metric for Multiprogram Workloads

In **isolation** on one core:

A Time = 1

SER = 4

B Time = 1

SER = 1

As a **2-program workload** on a **CMP** with two cores:

A Time = 1

SER = 4

Reliability Metric for Multiprogram Workloads

In **isolation** on one core:

A Time = 1

SER = 4

B Time = 1

SER = 1

As a **2-program workload** on a **CMP** with two cores:

A Time = 1

SER = 4

B Time = 2

SER = 1

Reliability Metric for Multiprogram Workloads

In **isolation** on one core:

A Time = 1

SER = 4

B Time = 1

SER = 1

As a **2-program workload** on a **CMP** with two cores:

A Time = 1

SER = 4

+

B Time = 2

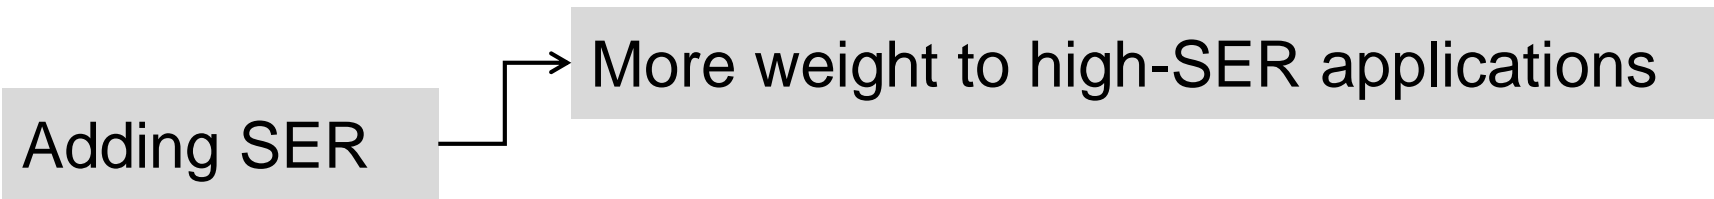
SER = 1

Total SER = 5

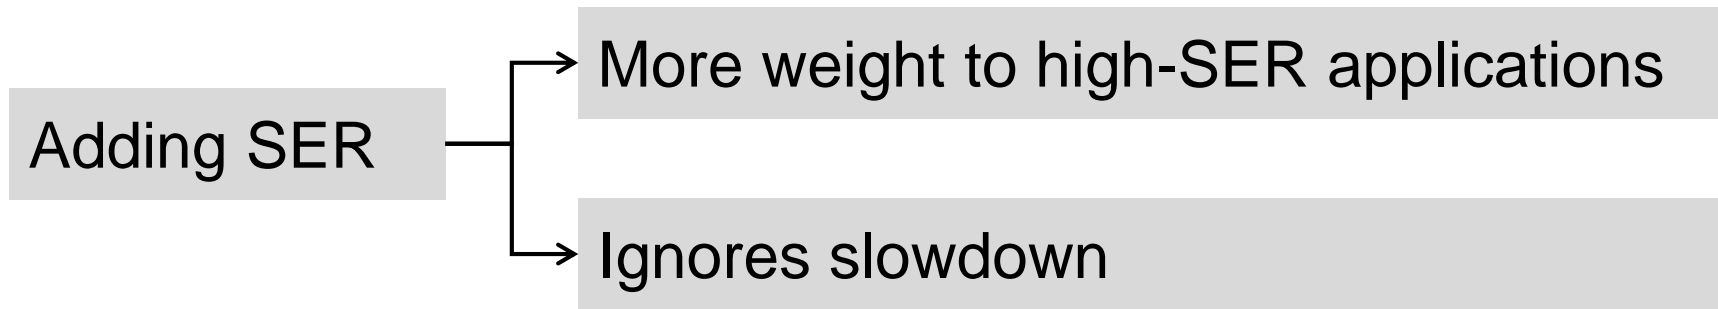
Reliability Metric for Multiprogram Workloads

Adding SER

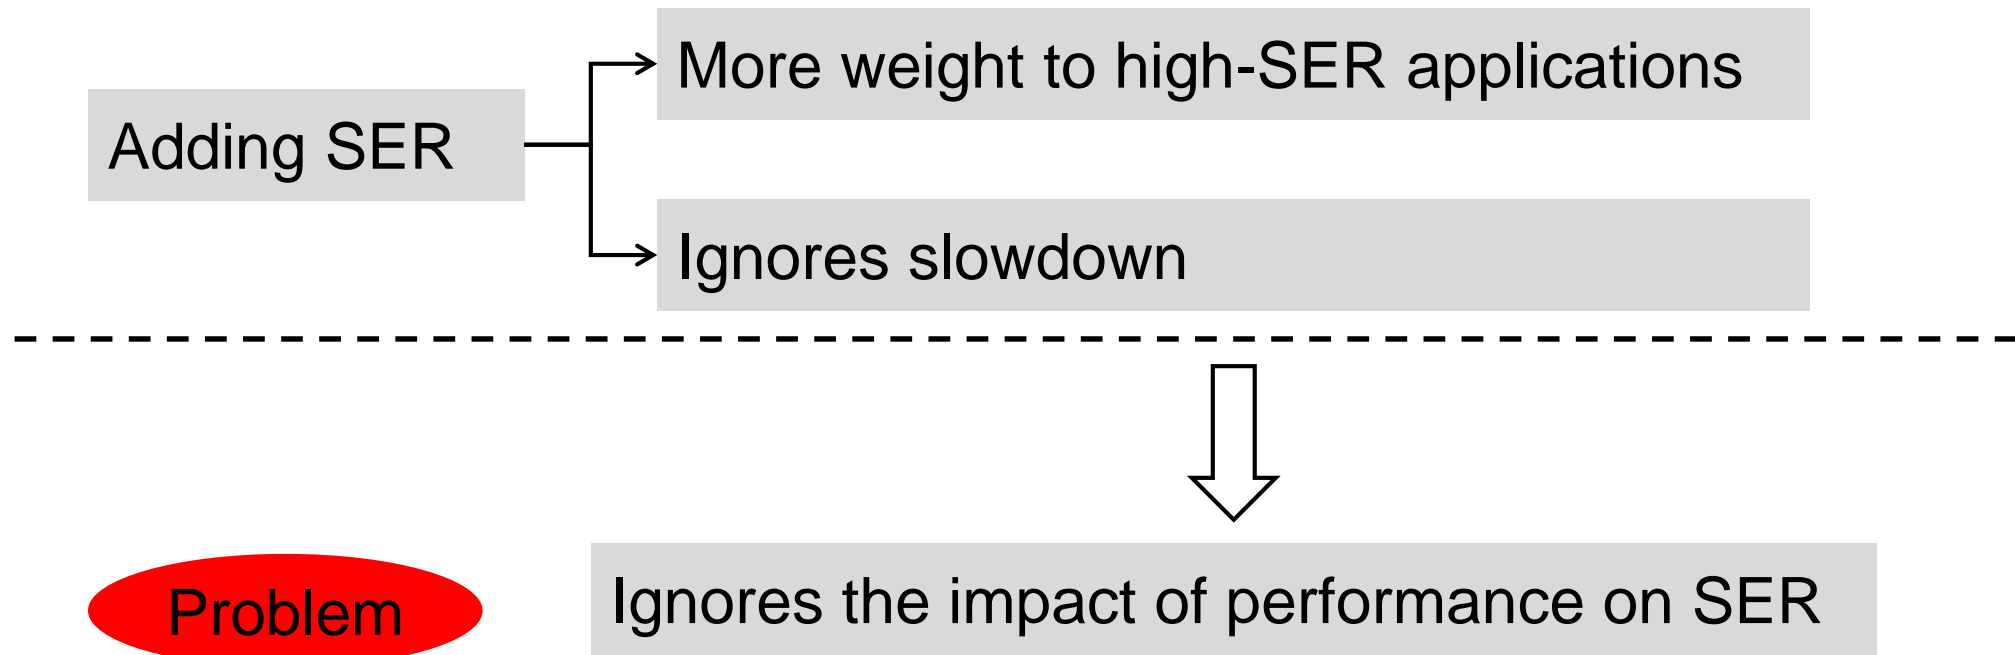
Reliability Metric for Multiprogram Workloads



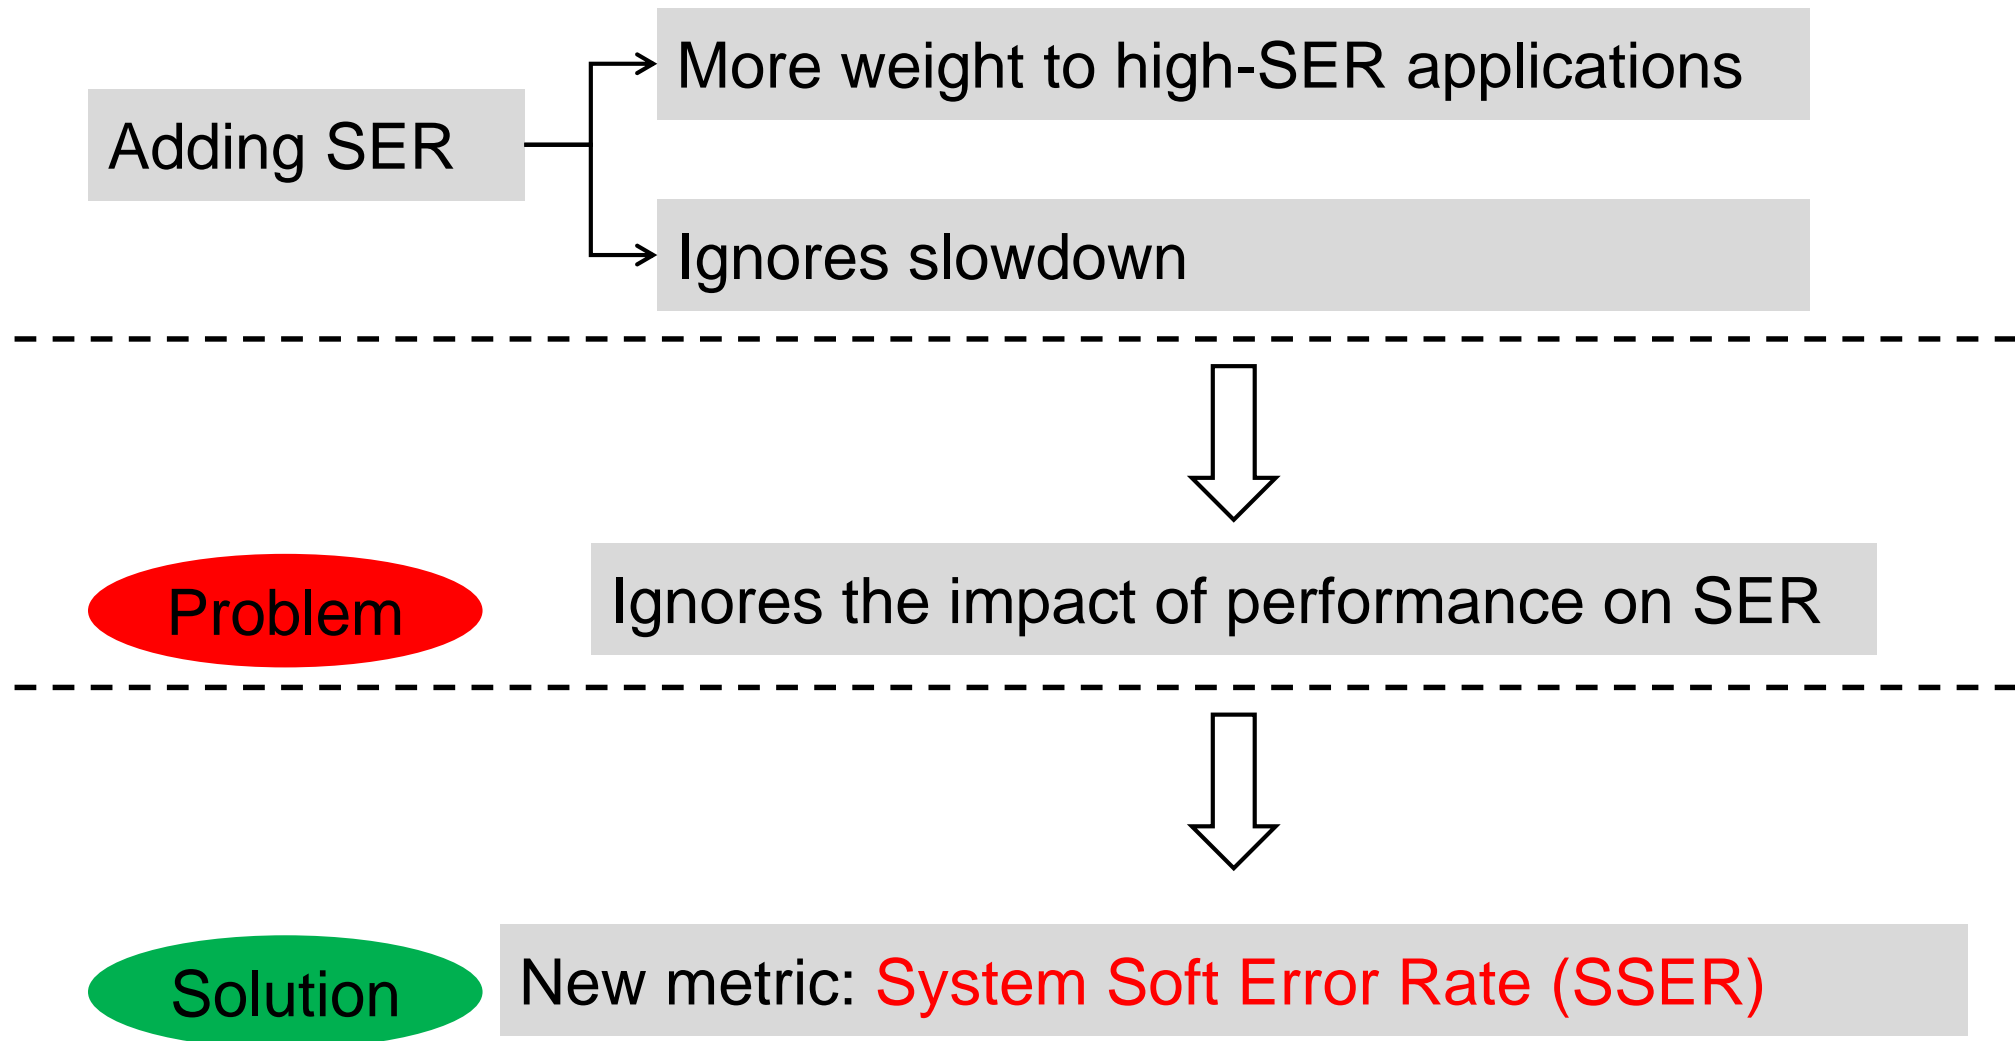
Reliability Metric for Multiprogram Workloads



Reliability Metric for Multiprogram Workloads



Reliability Metric for Multiprogram Workloads



System Soft Error Rate (SSER) Correctly Accounts for the Slowdown

In isolation on one core:

A Time = 1

SER = 4

B Time = 1

SER = 1

As a 2-program workload on a CMP with two cores:

A Time = 1

SER = 4

+

B Time = 2

SER = 1

Total SER = 5

System Soft Error Rate (SSER) Correctly Accounts for the Slowdown

In isolation on one core:

A Time = 1

SER = 4

B Time = 1

SER = 1

As a 2-program workload on a CMP with two cores:

A Time = 1

SER = 4

1×4

+

B Time = 2

SER = 1

Total SER = 5

System Soft Error Rate (SSER) Correctly Accounts for the Slowdown

In isolation on one core:

A Time = 1

SER = 4

B Time = 1

SER = 1

As a 2-program workload on a CMP with two cores:

A Time = 1

SER = 4

1×4 Weighted SER = 4

B Time = 2

SER = 1

+

Total SER = 5

System Soft Error Rate (SSER) Correctly Accounts for the Slowdown

In isolation on one core:

A Time = 1

SER = 4

B Time = 1

SER = 1

As a 2-program workload on a CMP with two cores:

A Time = 1

SER = 4

1×4 Weighted SER = 4

B Time = 2

SER = 1

2×1

+

Total SER = 5

System Soft Error Rate (SSER) Correctly Accounts for the Slowdown

In isolation on one core:

A Time = 1

SER = 4

B Time = 1

SER = 1

As a 2-program workload on a CMP with two cores:

A Time = 1

SER = 4

1×4 Weighted SER = 4

B Time = 2

SER = 1

2×1 Weighted SER = 2

Total SER = 5

System Soft Error Rate (SSER) Correctly Accounts for the Slowdown

In isolation on one core:

A Time = 1

SER = 4

B Time = 1

SER = 1

As a 2-program workload on a CMP with two cores:

A Time = 1

SER = 4

Weighted SER = 4

B Time = 2

SER = 1

Weighted SER = 2

Total SER = 5

SSER = 6

System Soft Error Rate (SSER) Correctly Accounts for the Slowdown

In isolation on one core:

A Time = 1

SER = 4

B Time = 1

SER = 1

As a 2-program workload on a CMP with two cores:

A Time = 1

SER = 4

Weighted SER = 4

+

B Time = 2

SER = 1

Weighted SER = 2

Total SER =

5



SSER =

6

System Soft Error Rate (SSER) Correctly Accounts for the Slowdown

In isolation on one core:

A Time = 1

SER = 4

B Time = 1

SER = 1

As a 2-program workload on a CMP with two cores:

A Time = 1

SER = 4

Weighted SER = 4

+

B Time = 2

SER = 1

Weighted SER = 2

Total SER =

5



SSER =

6



Scheduling Algorithm

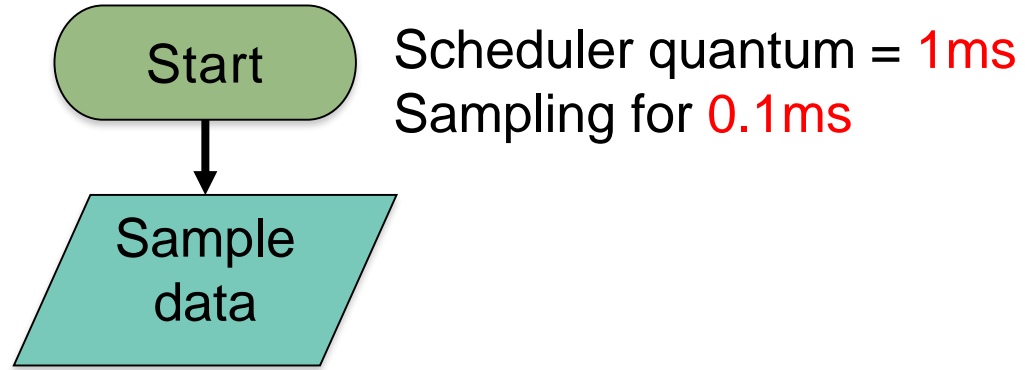
Scheduling Algorithm



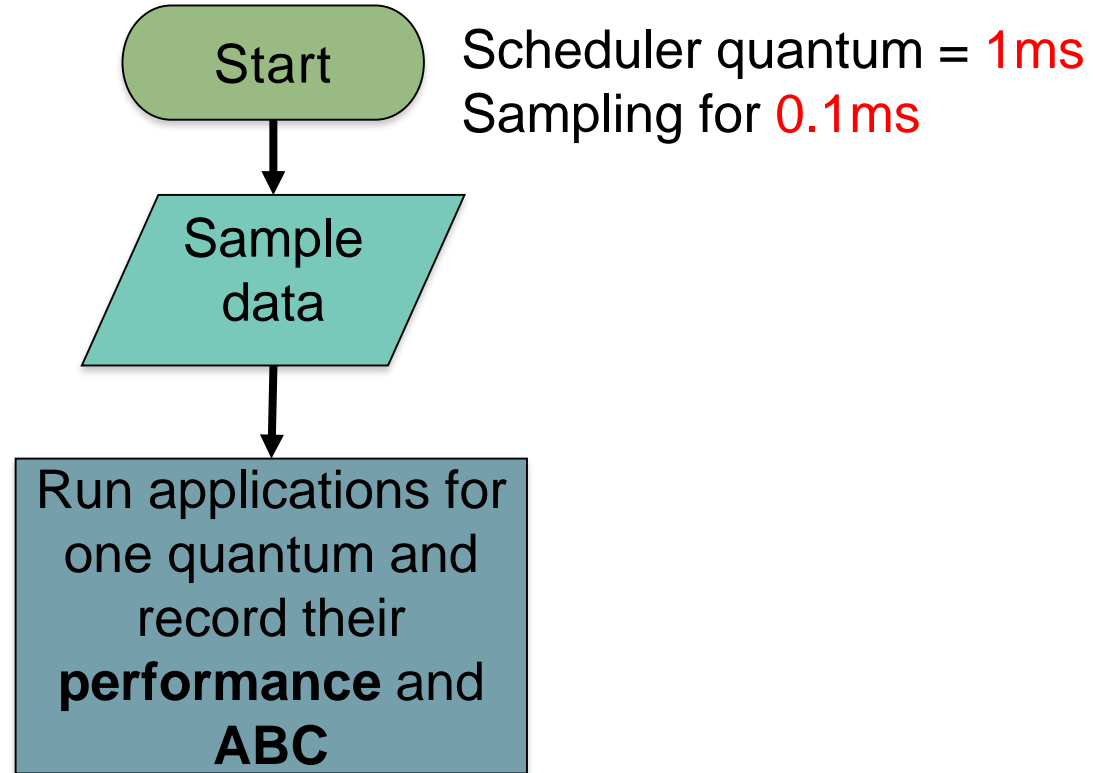
```
graph TD; Start([Start]);
```

Start

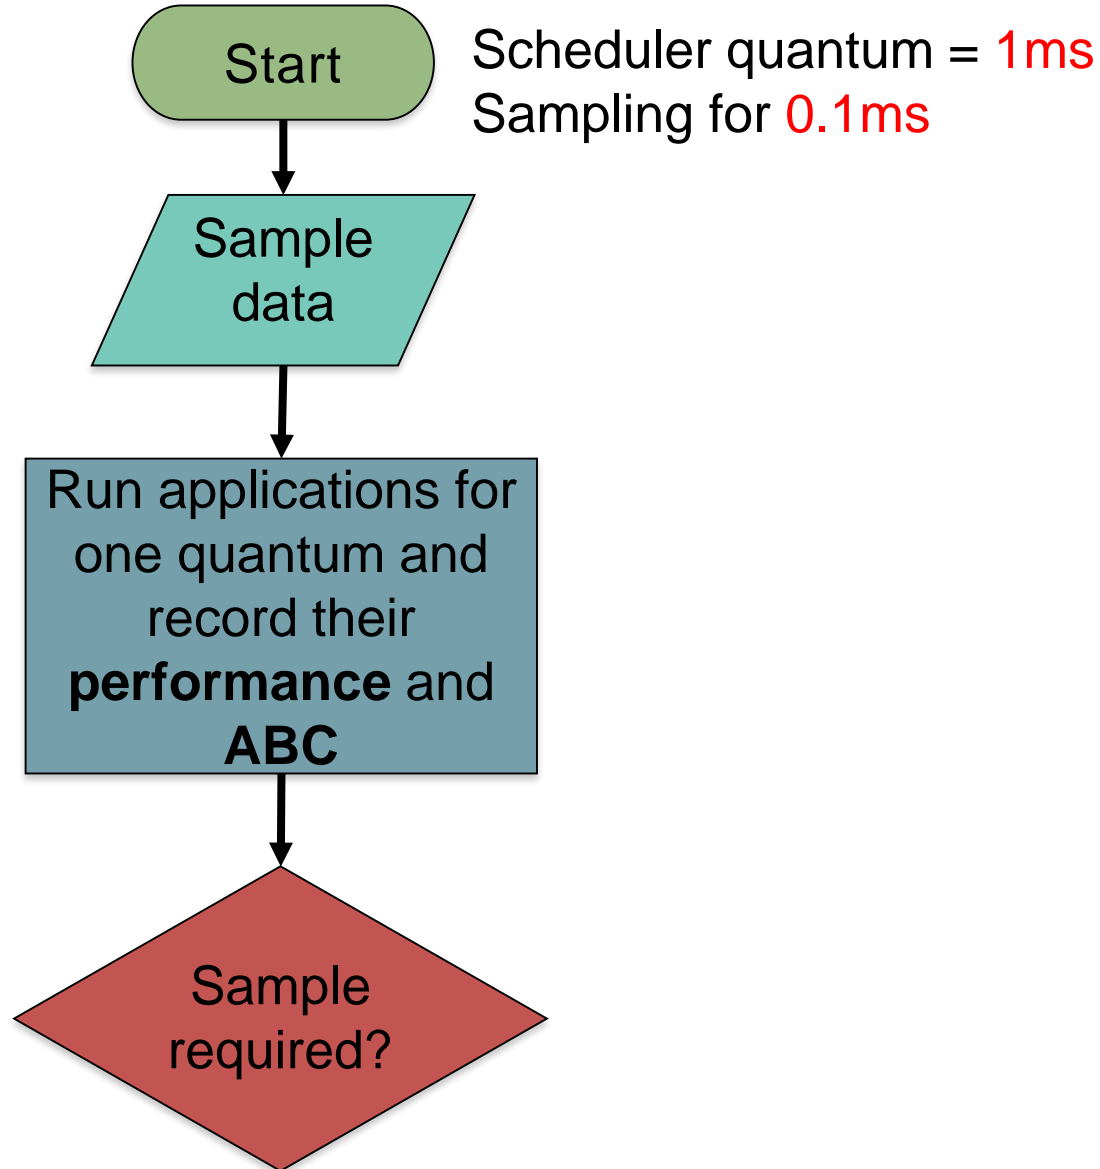
Scheduling Algorithm



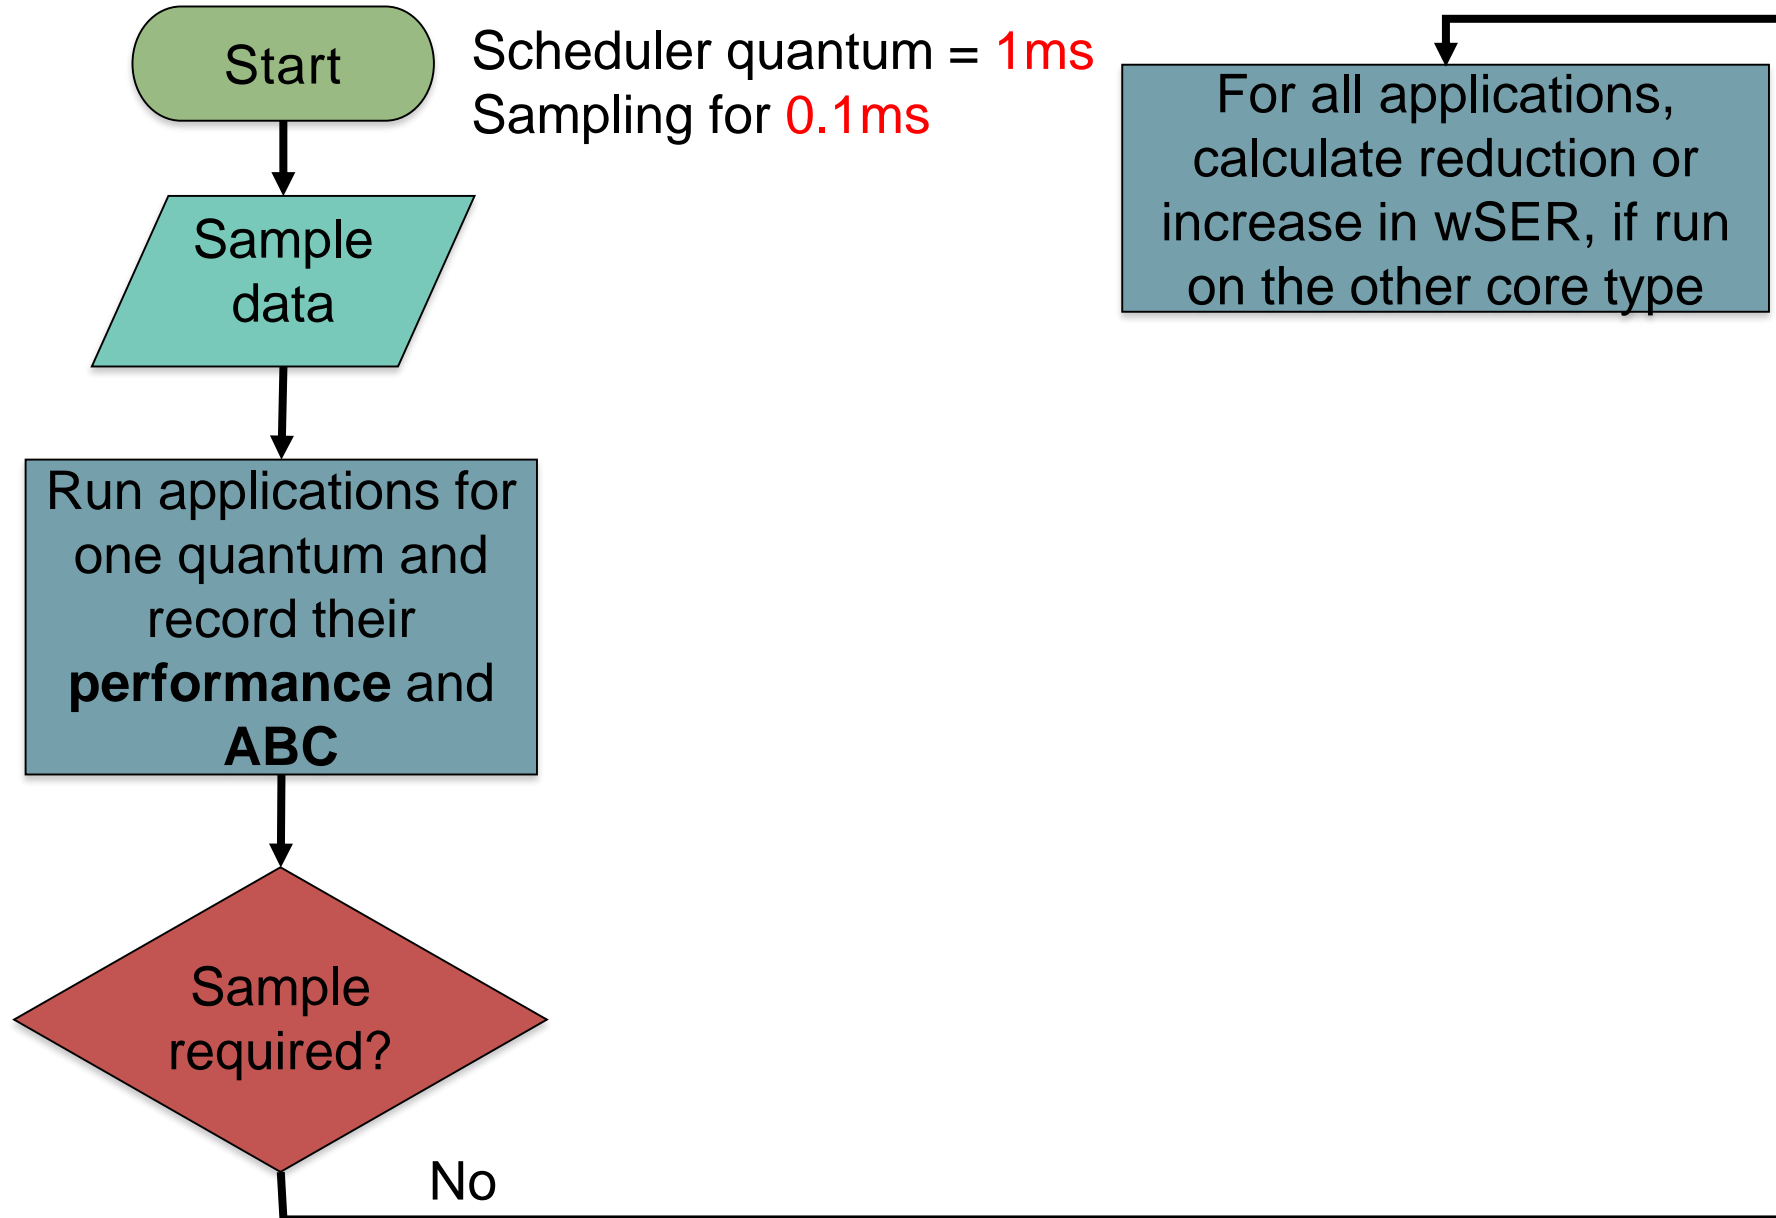
Scheduling Algorithm



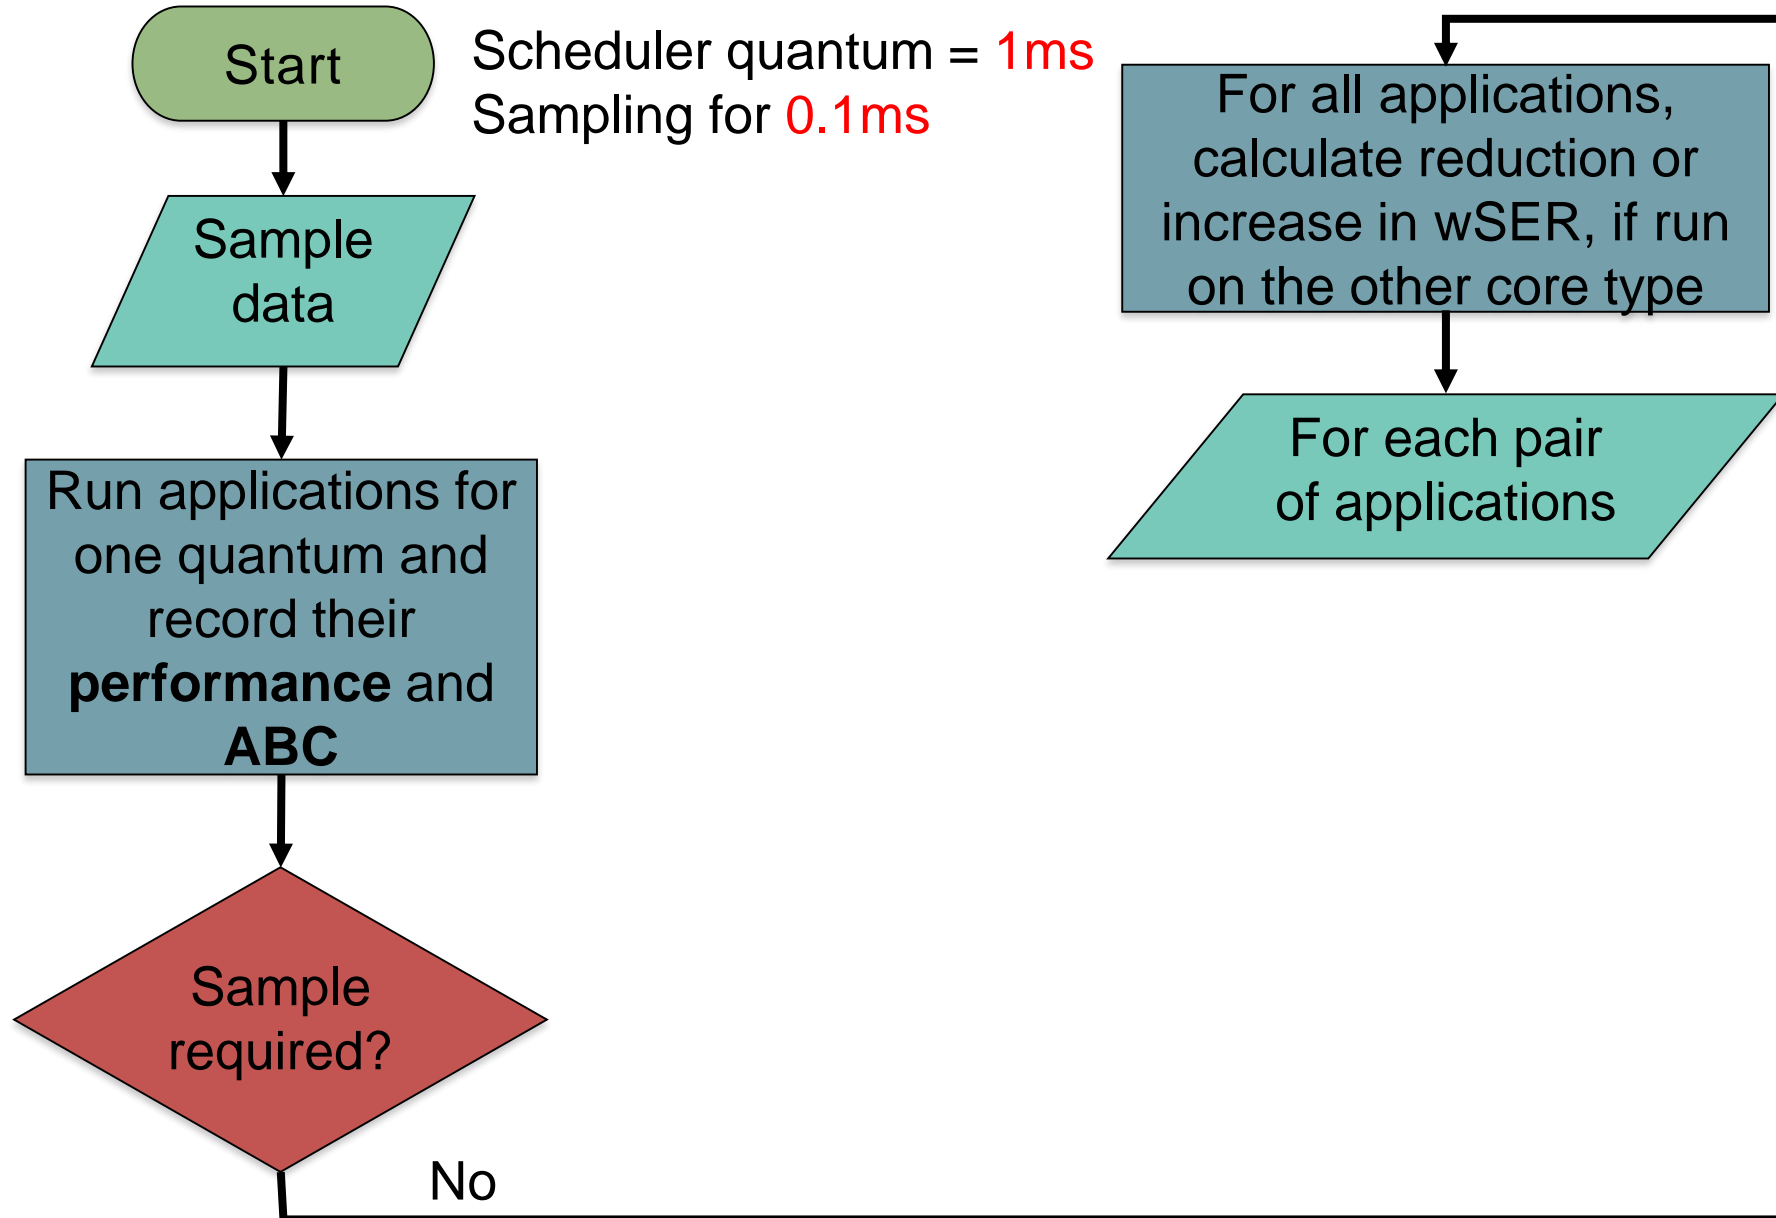
Scheduling Algorithm



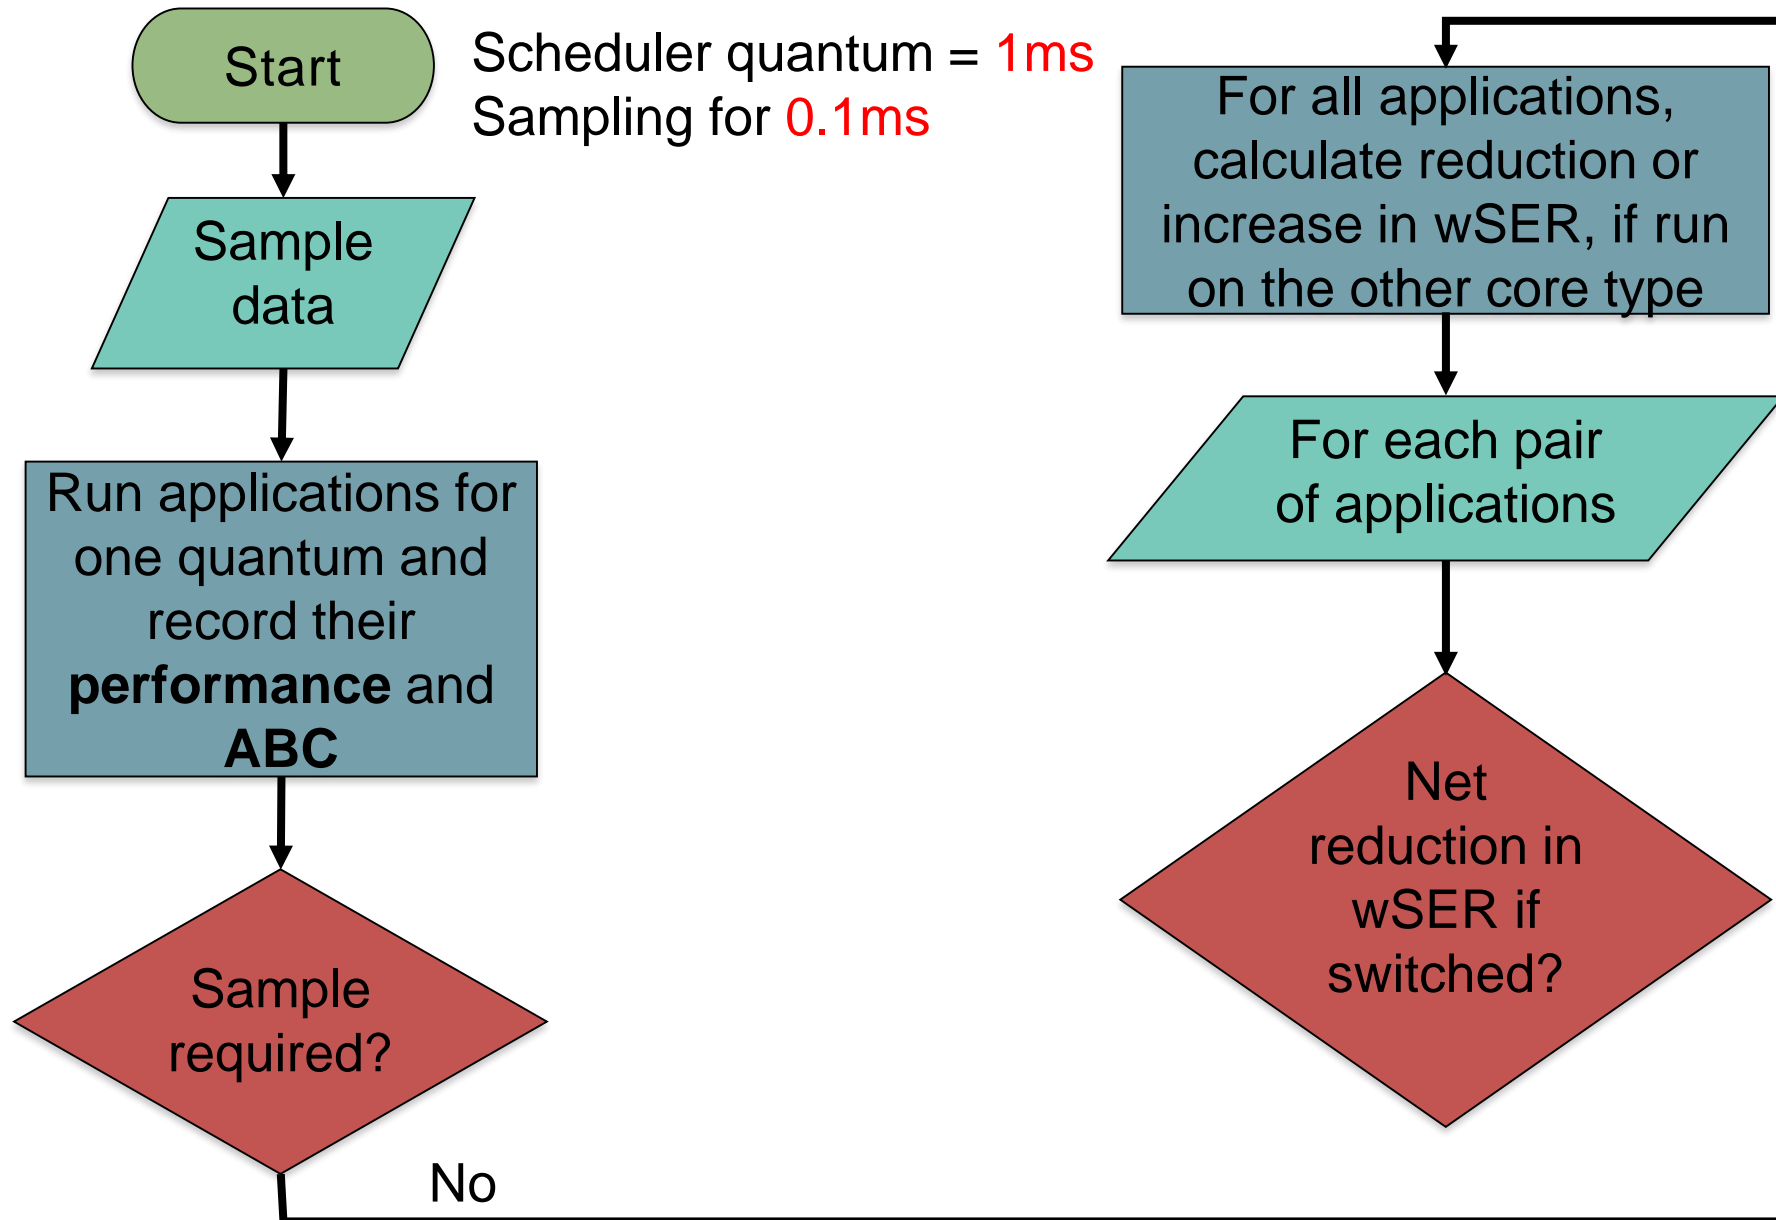
Scheduling Algorithm



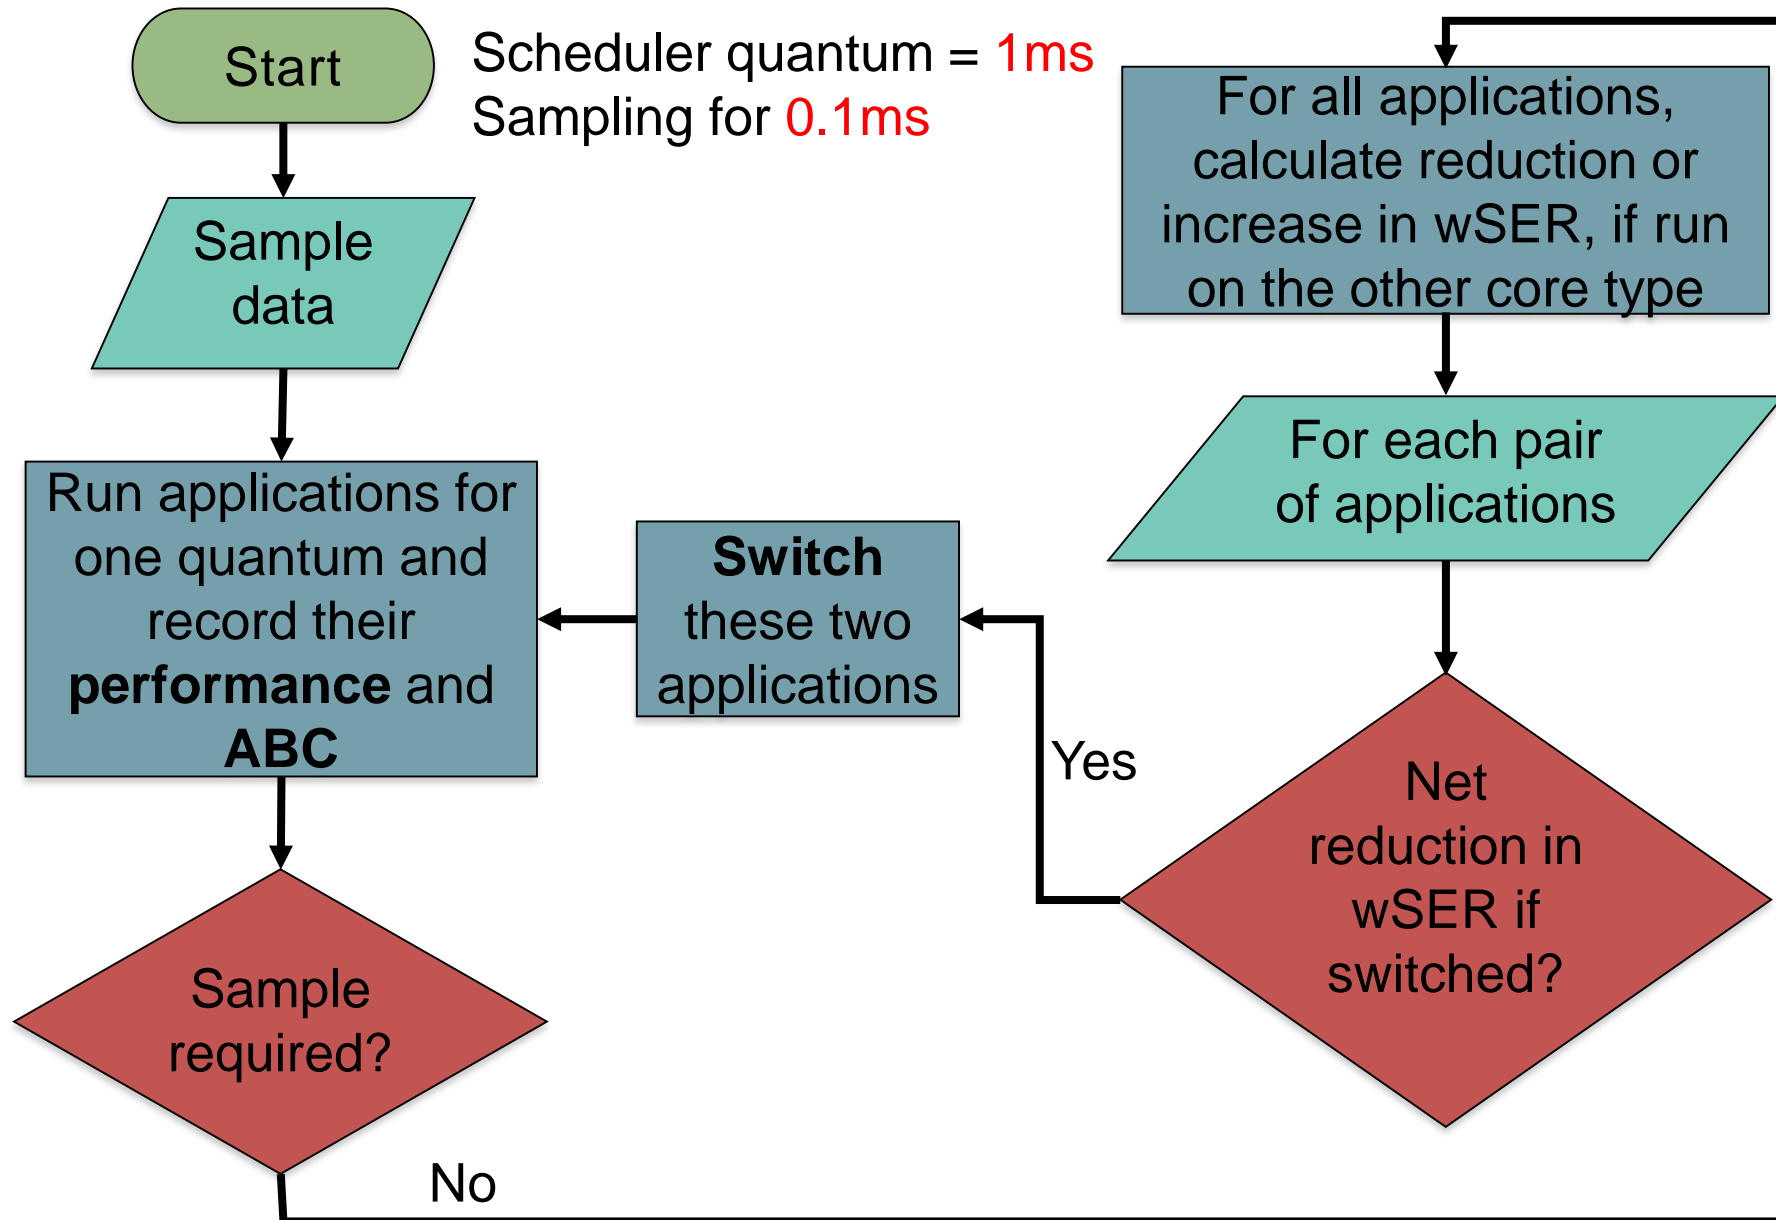
Scheduling Algorithm



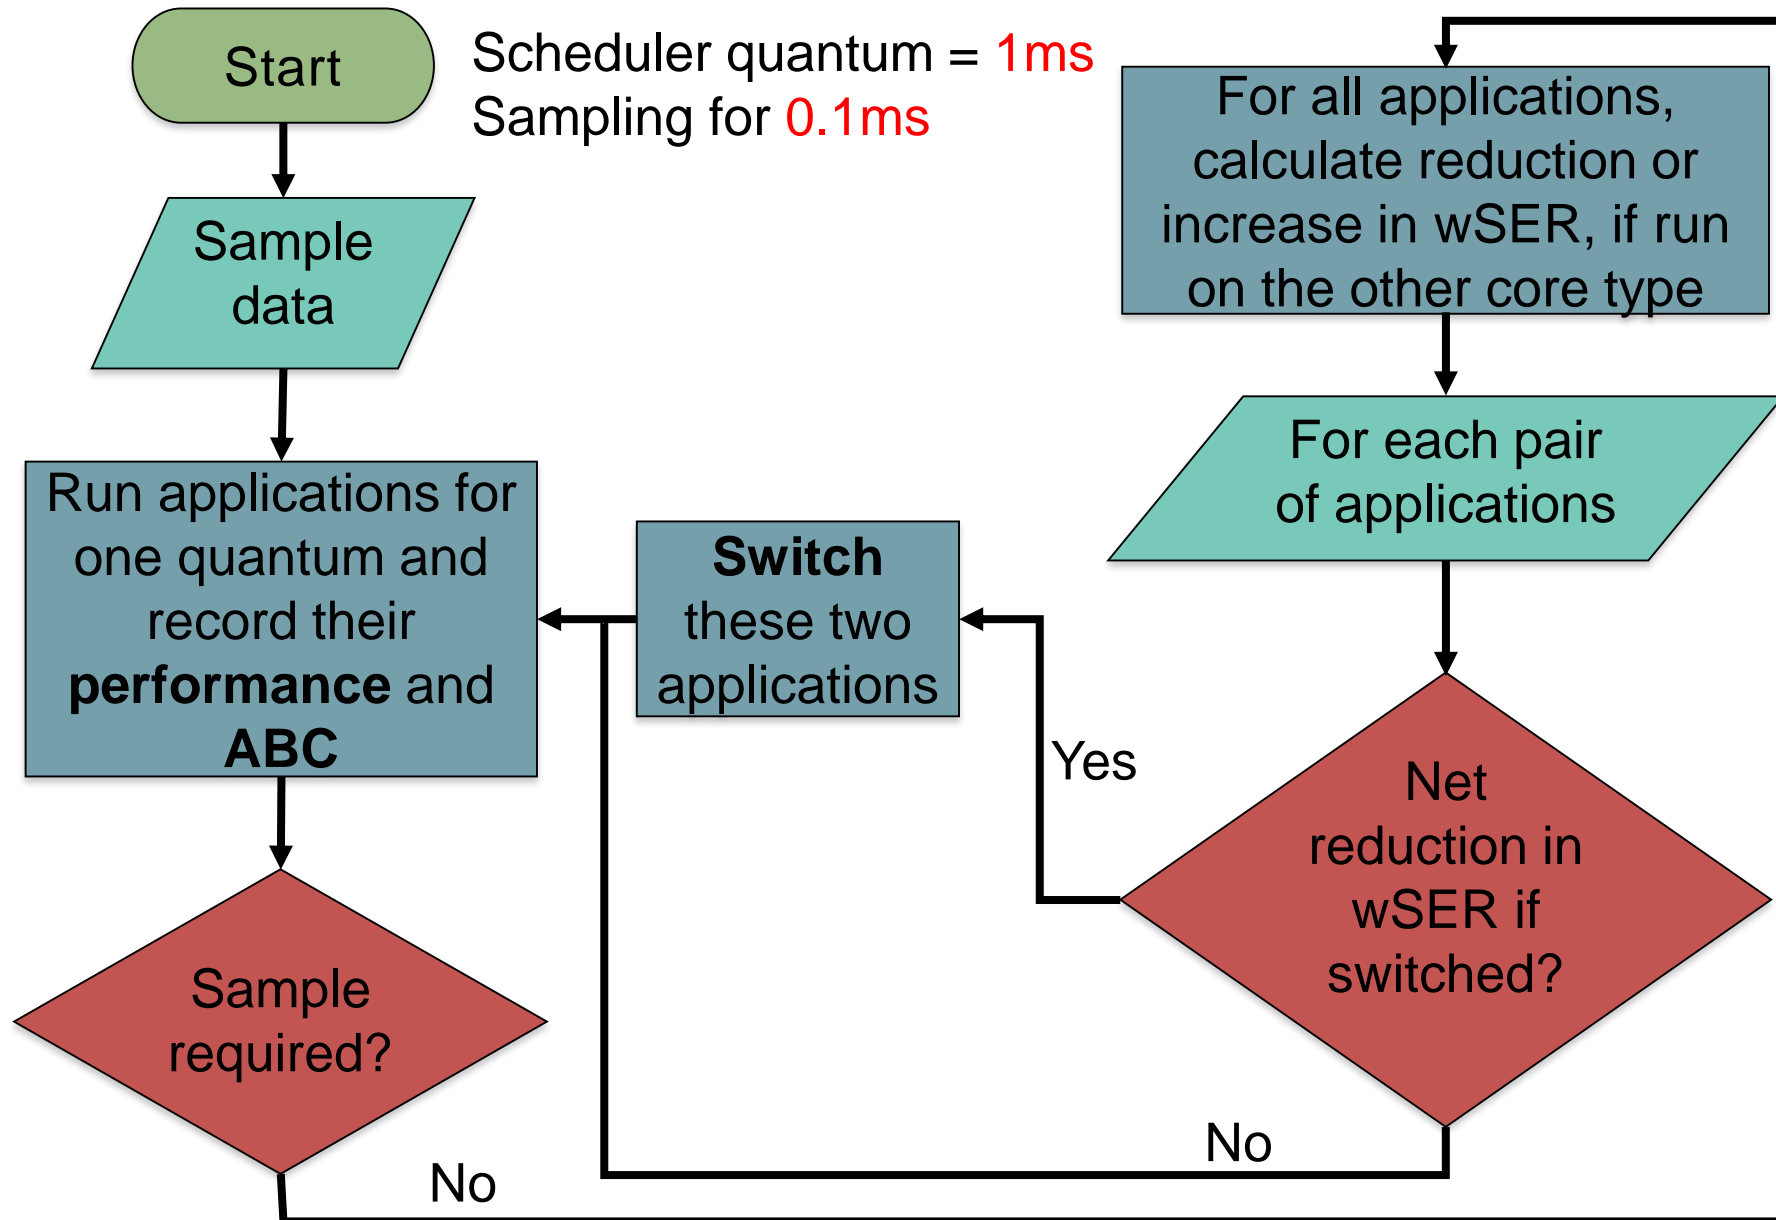
Scheduling Algorithm



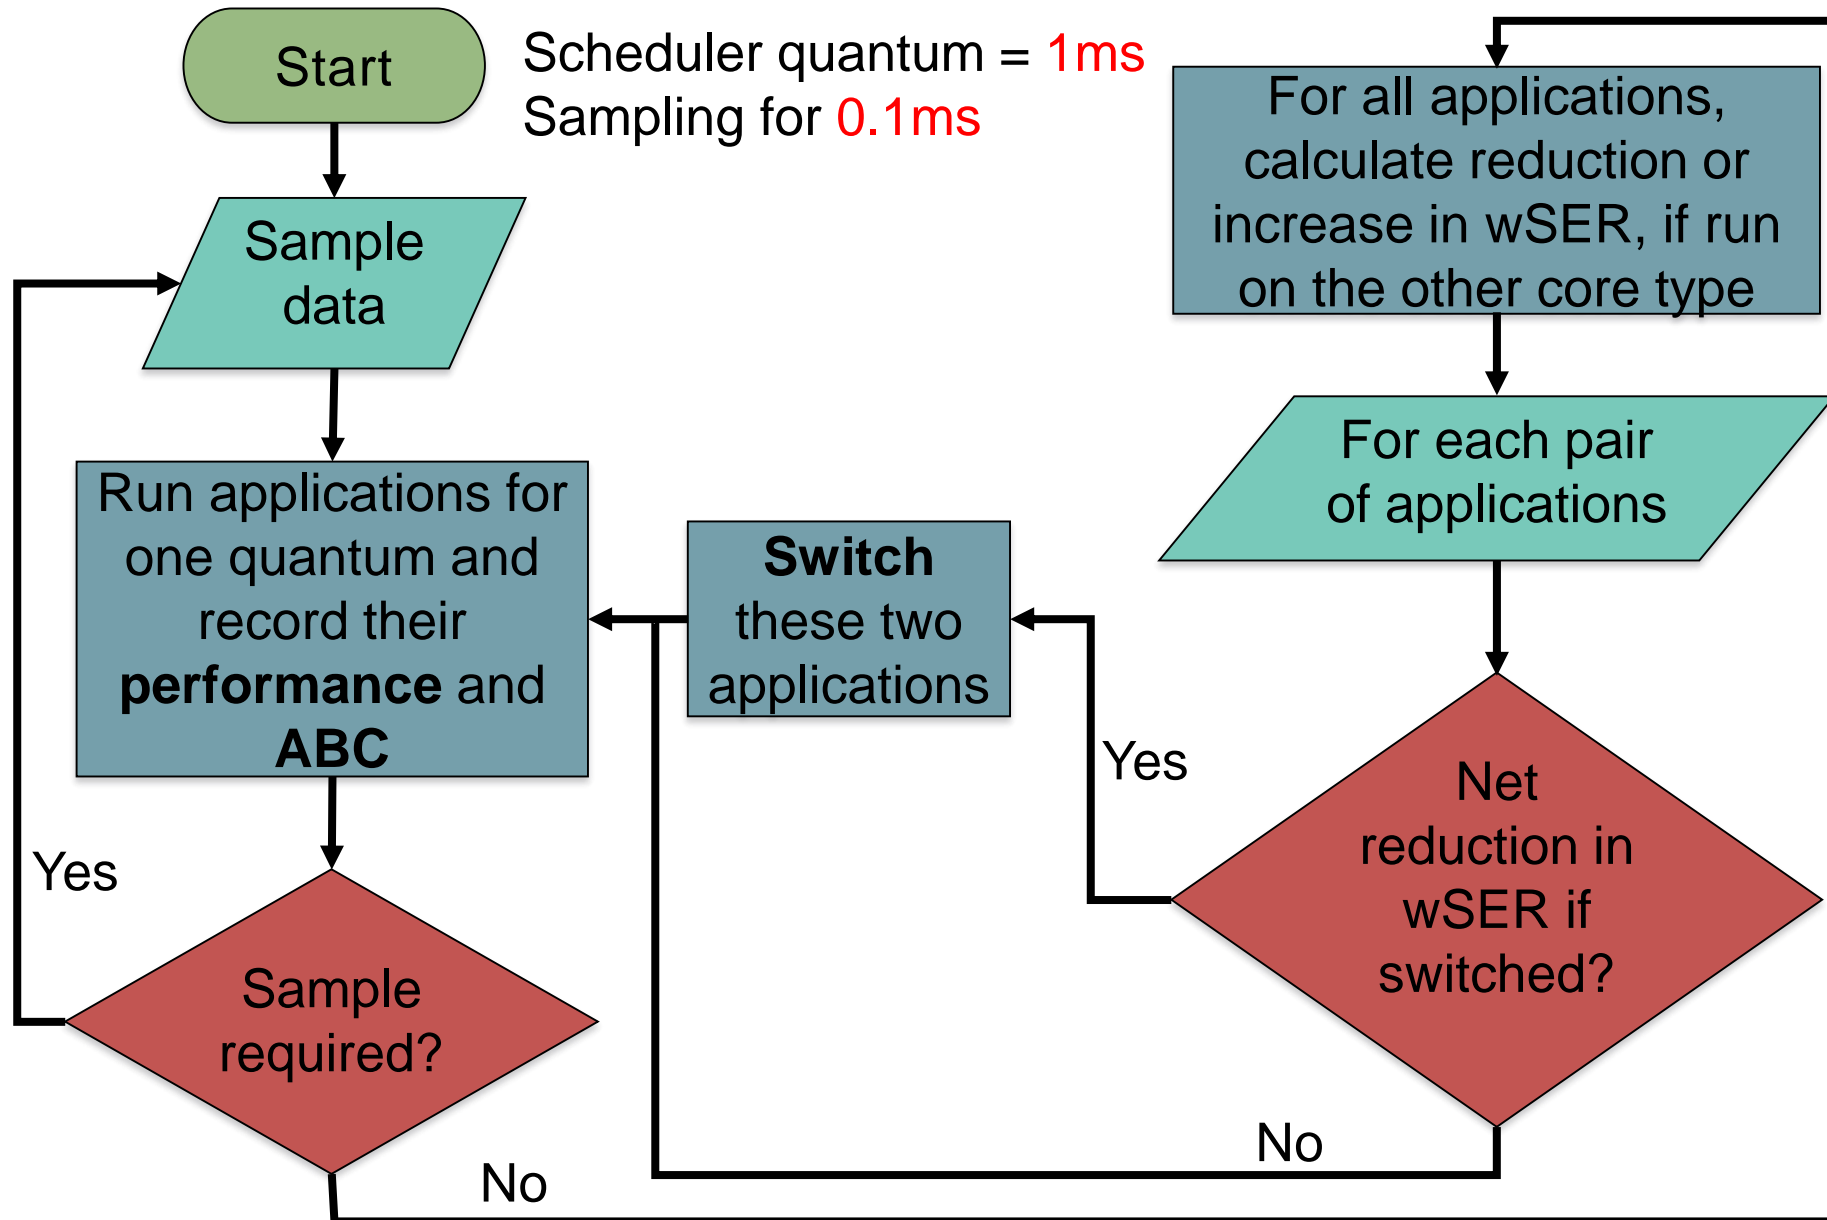
Scheduling Algorithm



Scheduling Algorithm



Scheduling Algorithm



Experimental Setup

Experimental Setup

Sniper 6.0

- Augmented with ACE Bit Counters

Experimental Setup

Sniper 6.0

-- Augmented with ACE Bit Counters

SPEC CPU2006

-- 1B instruction SimPoints

-- Based on **big-core SER**

- ✓ L → Low SER
- ✓ M → Medium SER
- ✓ H → High SER

Experimental Setup

Sniper 6.0

-- Augmented with ACE Bit Counters

- ✓ L → Low SER
- ✓ M → Medium SER
- ✓ H → High SER

SPEC CPU2006

-- 1B instruction SimPoints
-- Based on **big-core SER**

4-program workloads

-- 6 categories → HHHH, HHMM, HHLL, MMMM, MMLL, LLLL

Experimental Setup

Sniper 6.0

-- Augmented with ACE Bit Counters

SPEC CPU2006

-- 1B instruction SimPoints
-- Based on **big-core SER**

4-program workloads

-- 6 categories → HHHH, HHMM, HHLL, MMMM, MMLL, LLLL

- ✓ L → Low SER
- ✓ M → Medium SER
- ✓ H → High SER

HCMP → **two big** and
two small cores (2B2S)

Evaluation

Reliability-Optimized Scheduler

- Optimizes **SSER** using our scheduling algorithm

Evaluation

Reliability-Optimized Scheduler

- Optimizes **SSER** using our scheduling algorithm

Performance-Optimized Scheduler

- Optimizes System Throughput (**STP**)

Evaluation

Reliability-Optimized Scheduler

- Optimizes **SSER** using our scheduling algorithm

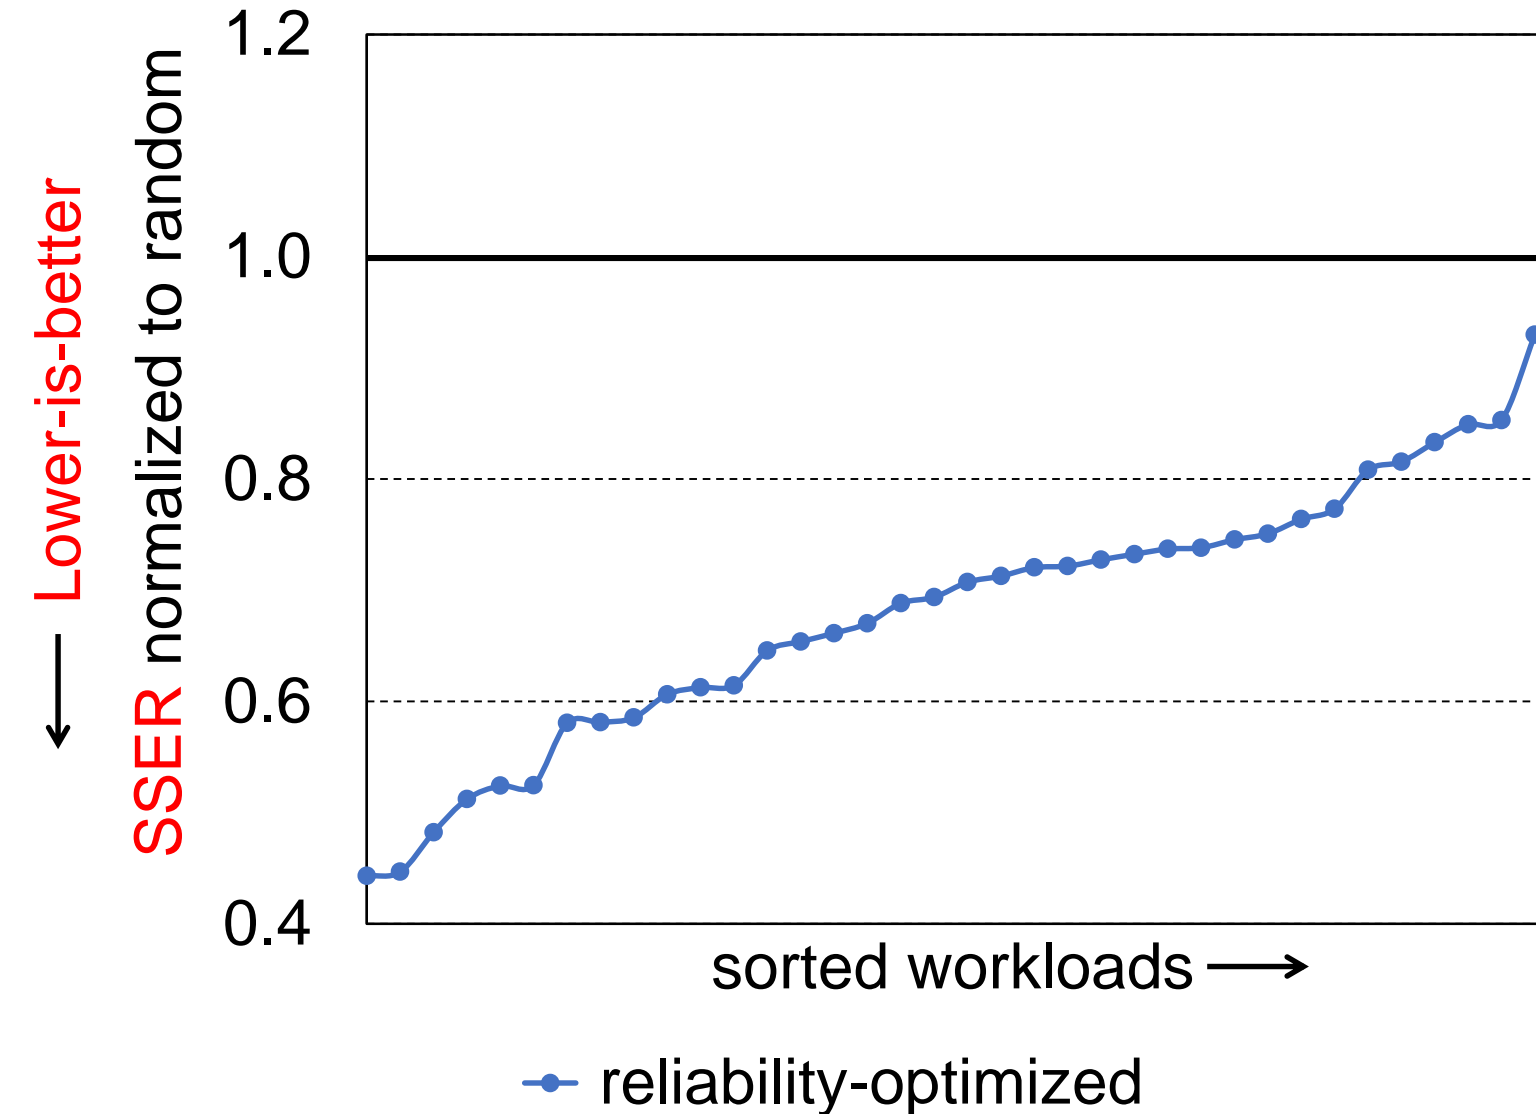
Performance-Optimized Scheduler

- Optimizes System Throughput (**STP**)

Random Scheduler

- Randomly selects applications to run on the **big core(s)**

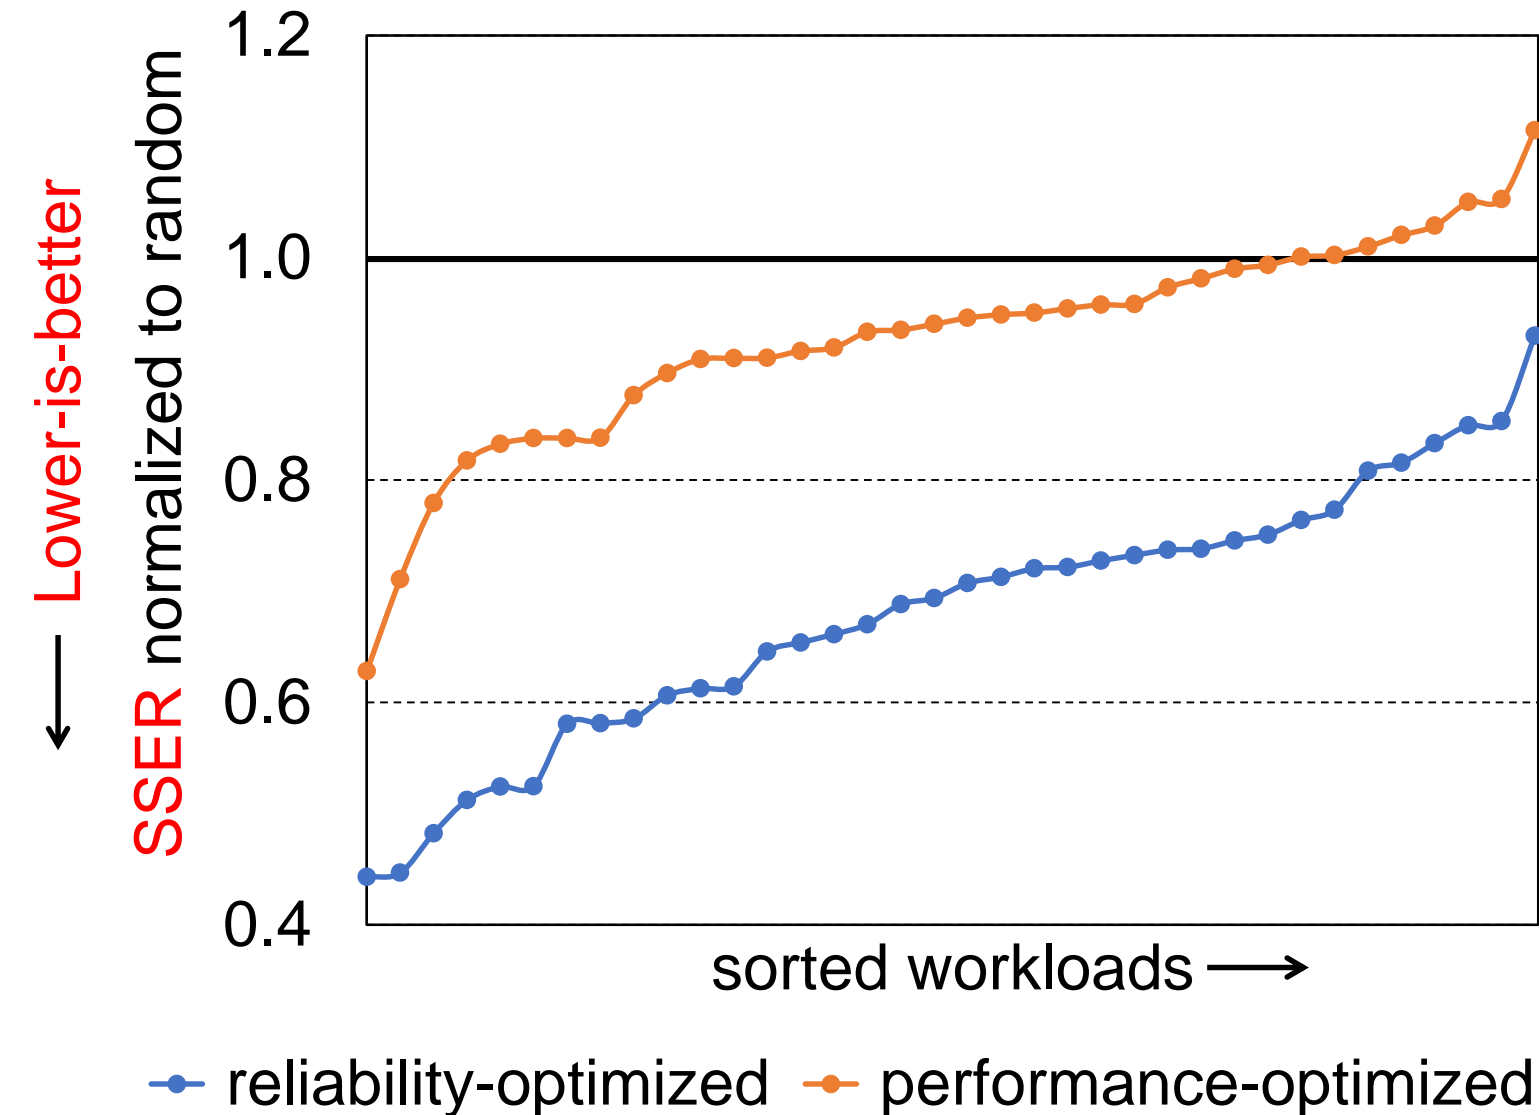
2B2S Results – Reliability



Compared to random

Reliability-optimized
-- improves SSER by 32%
-- up to 55.6%

2B2S Results – Reliability

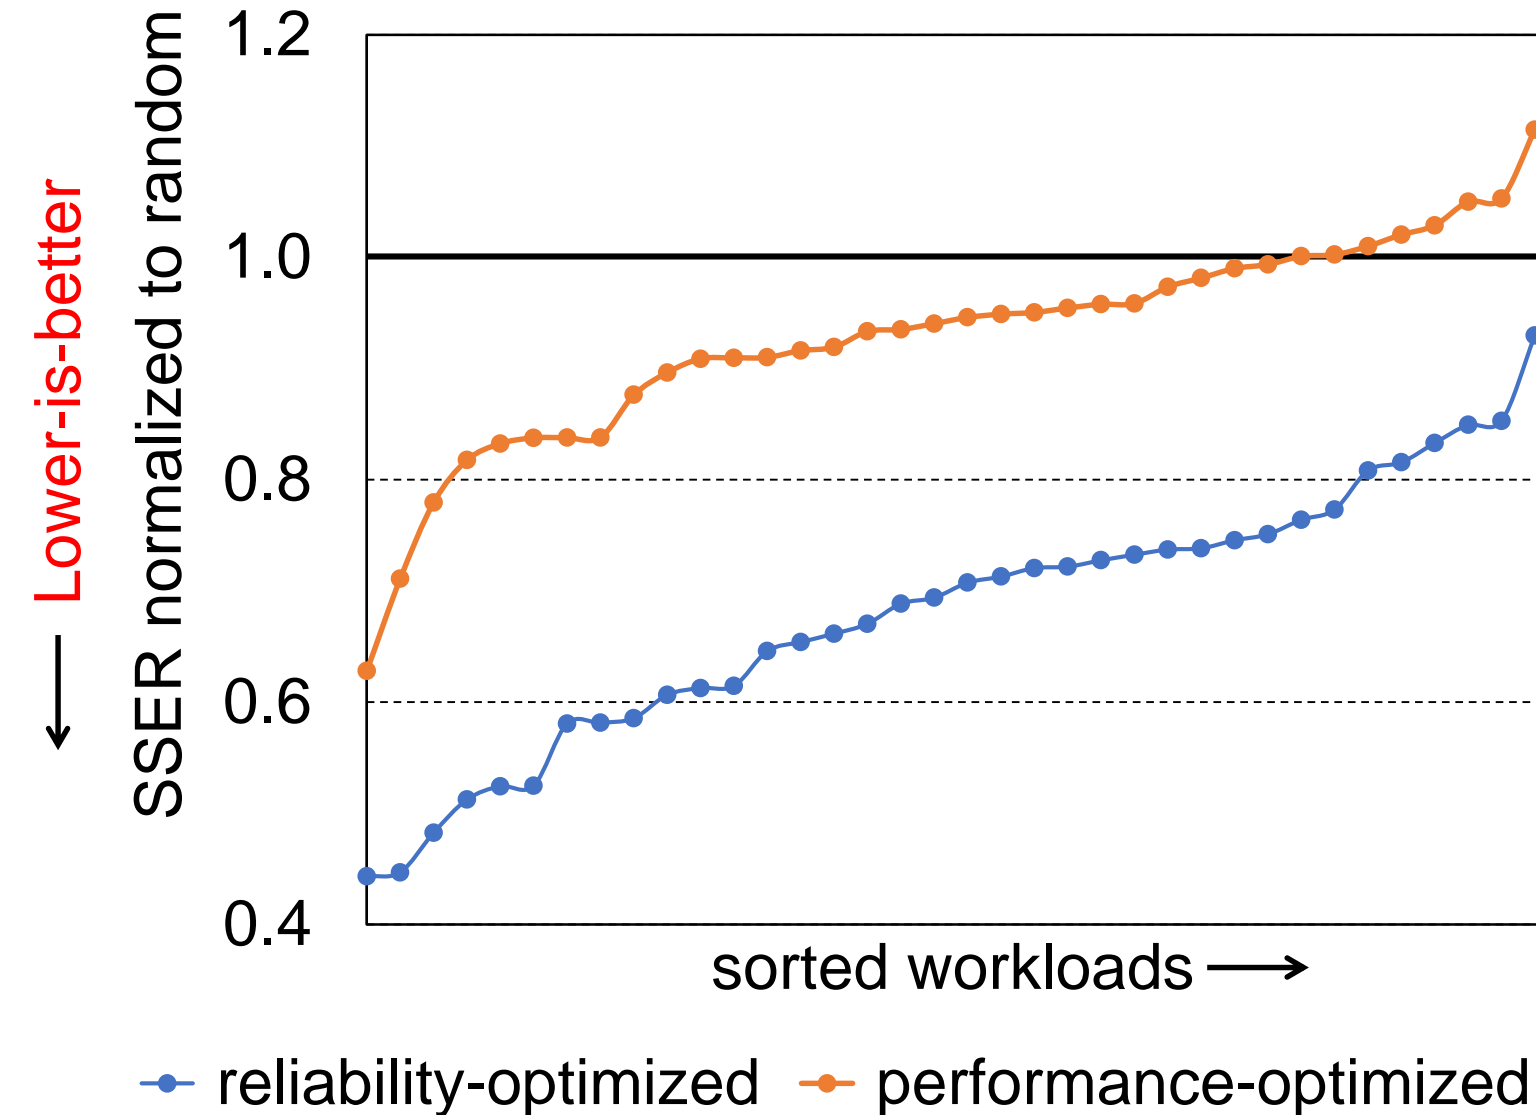


Compared to random

Reliability-optimized
-- improves SSER by 32%
-- up to 55.6%

Performance-optimized
-- improves SSER by 7.3%

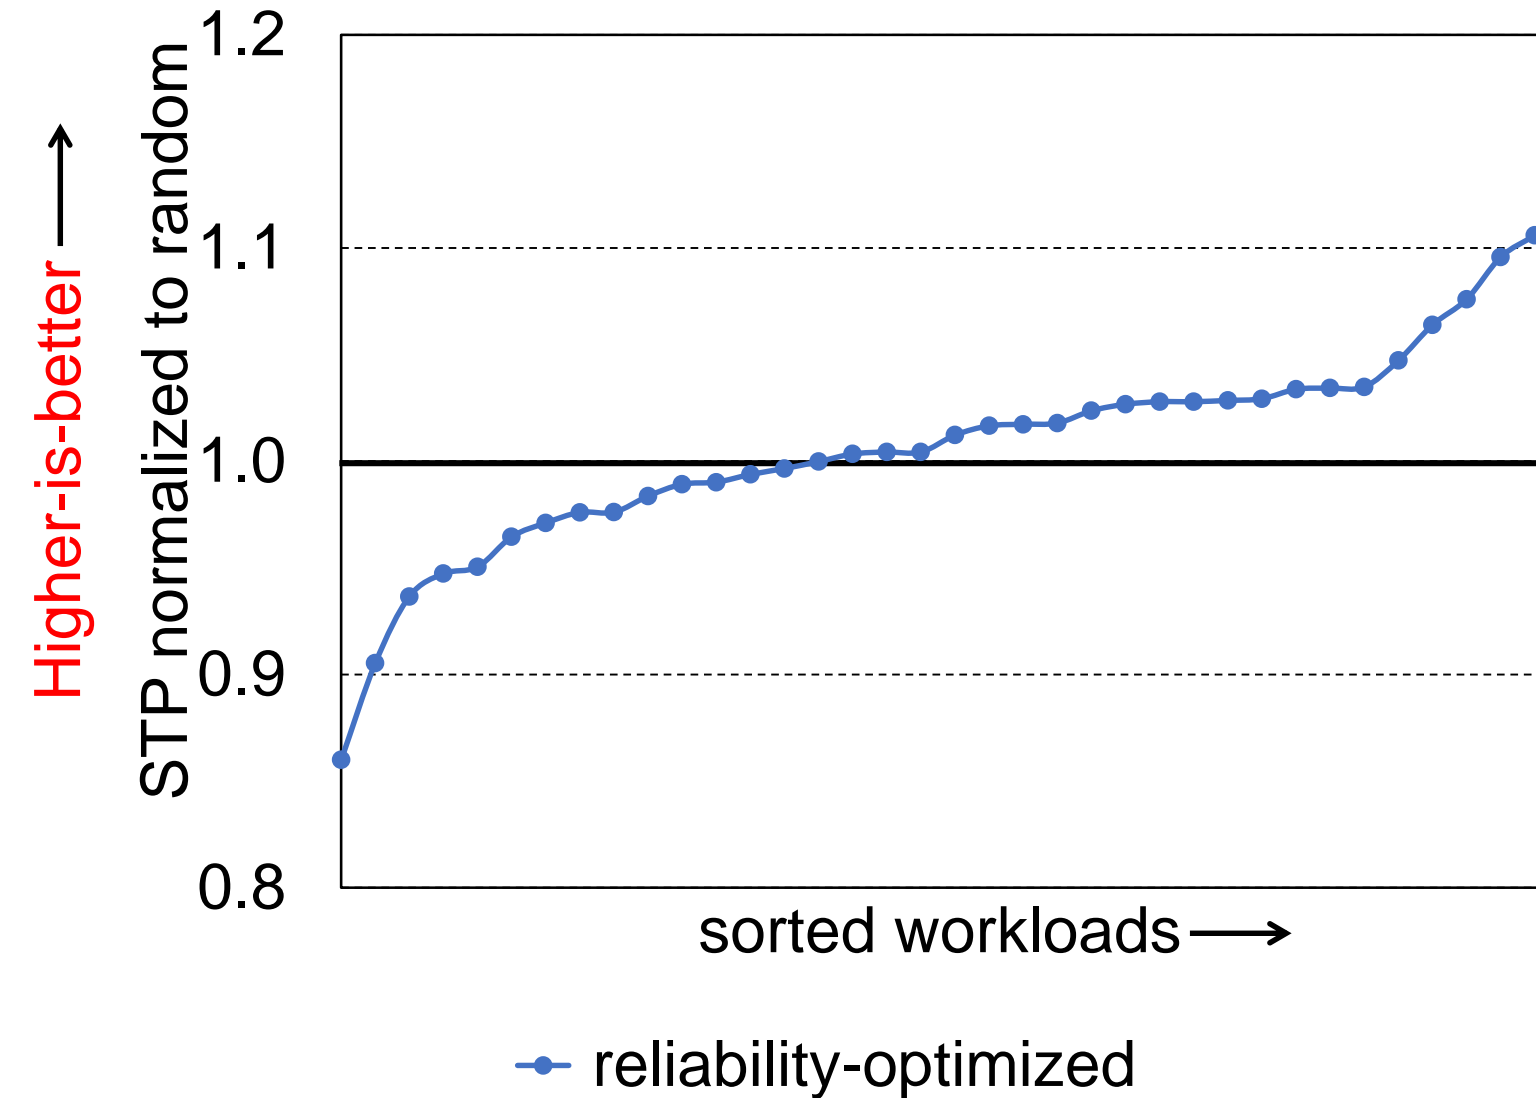
2B2S Results – Reliability



Compared to
performance-optimized

Reliability-optimized
-- improves SSER by 25.4%
-- up to 60.2%

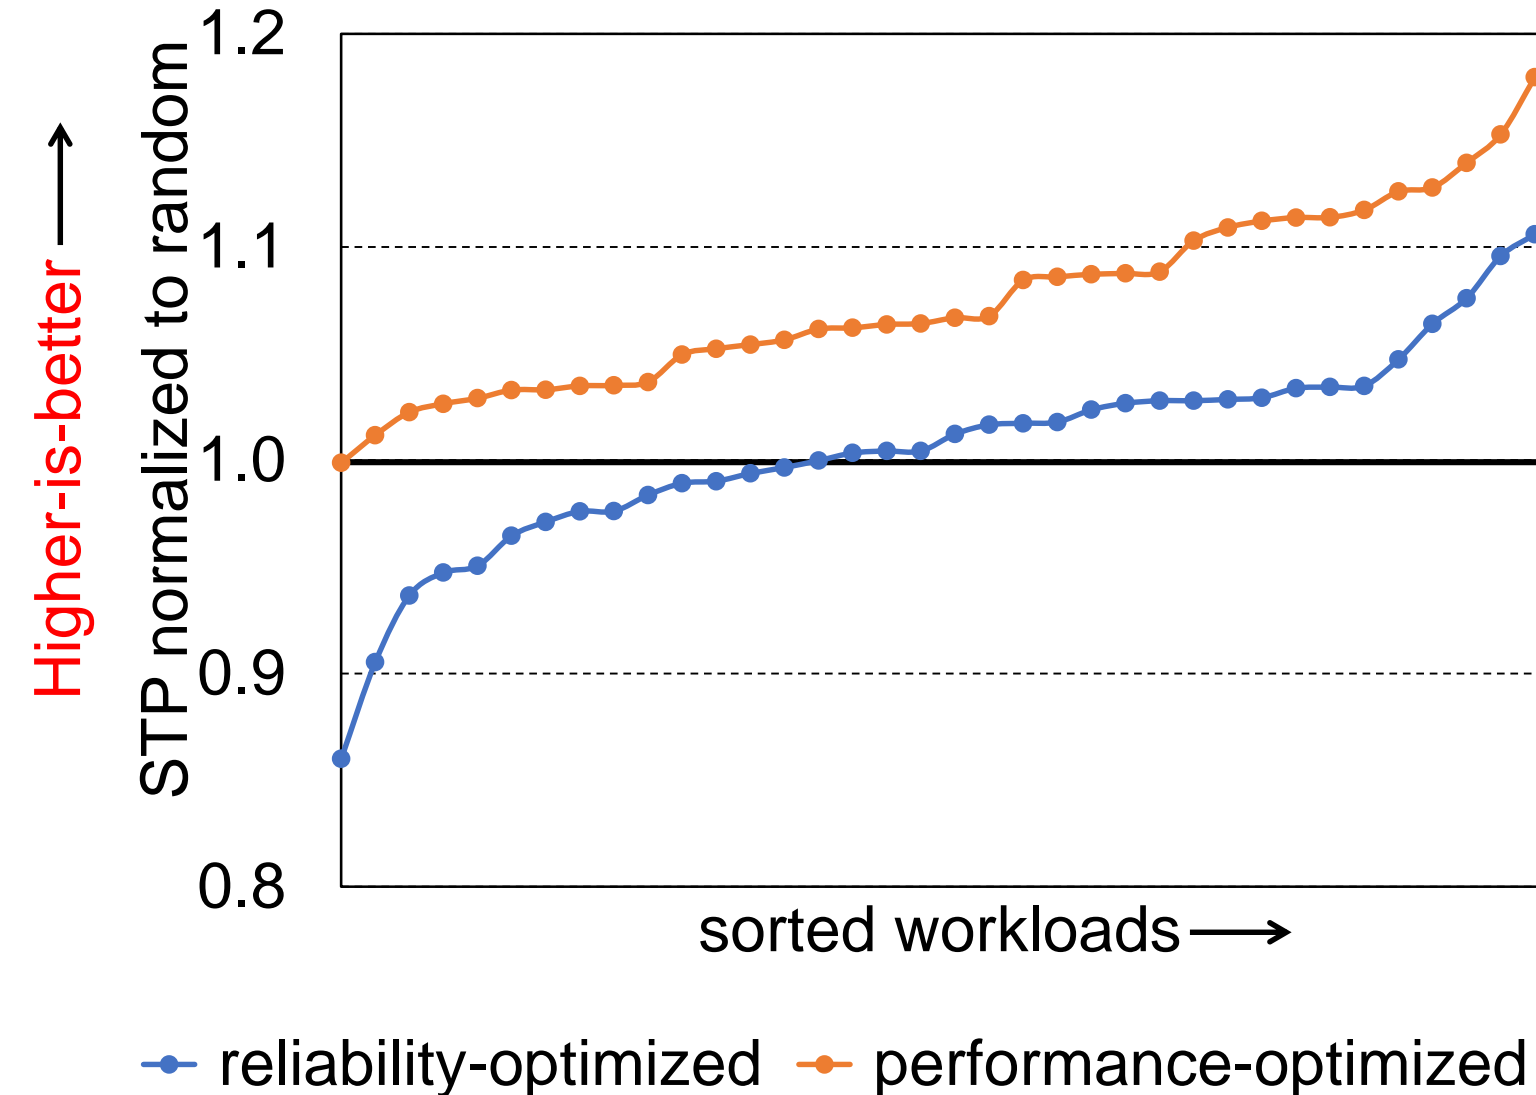
2B2S Results – Performance



Compared to random

Reliability-optimized
-- does not degrade STP

2B2S Results – Performance

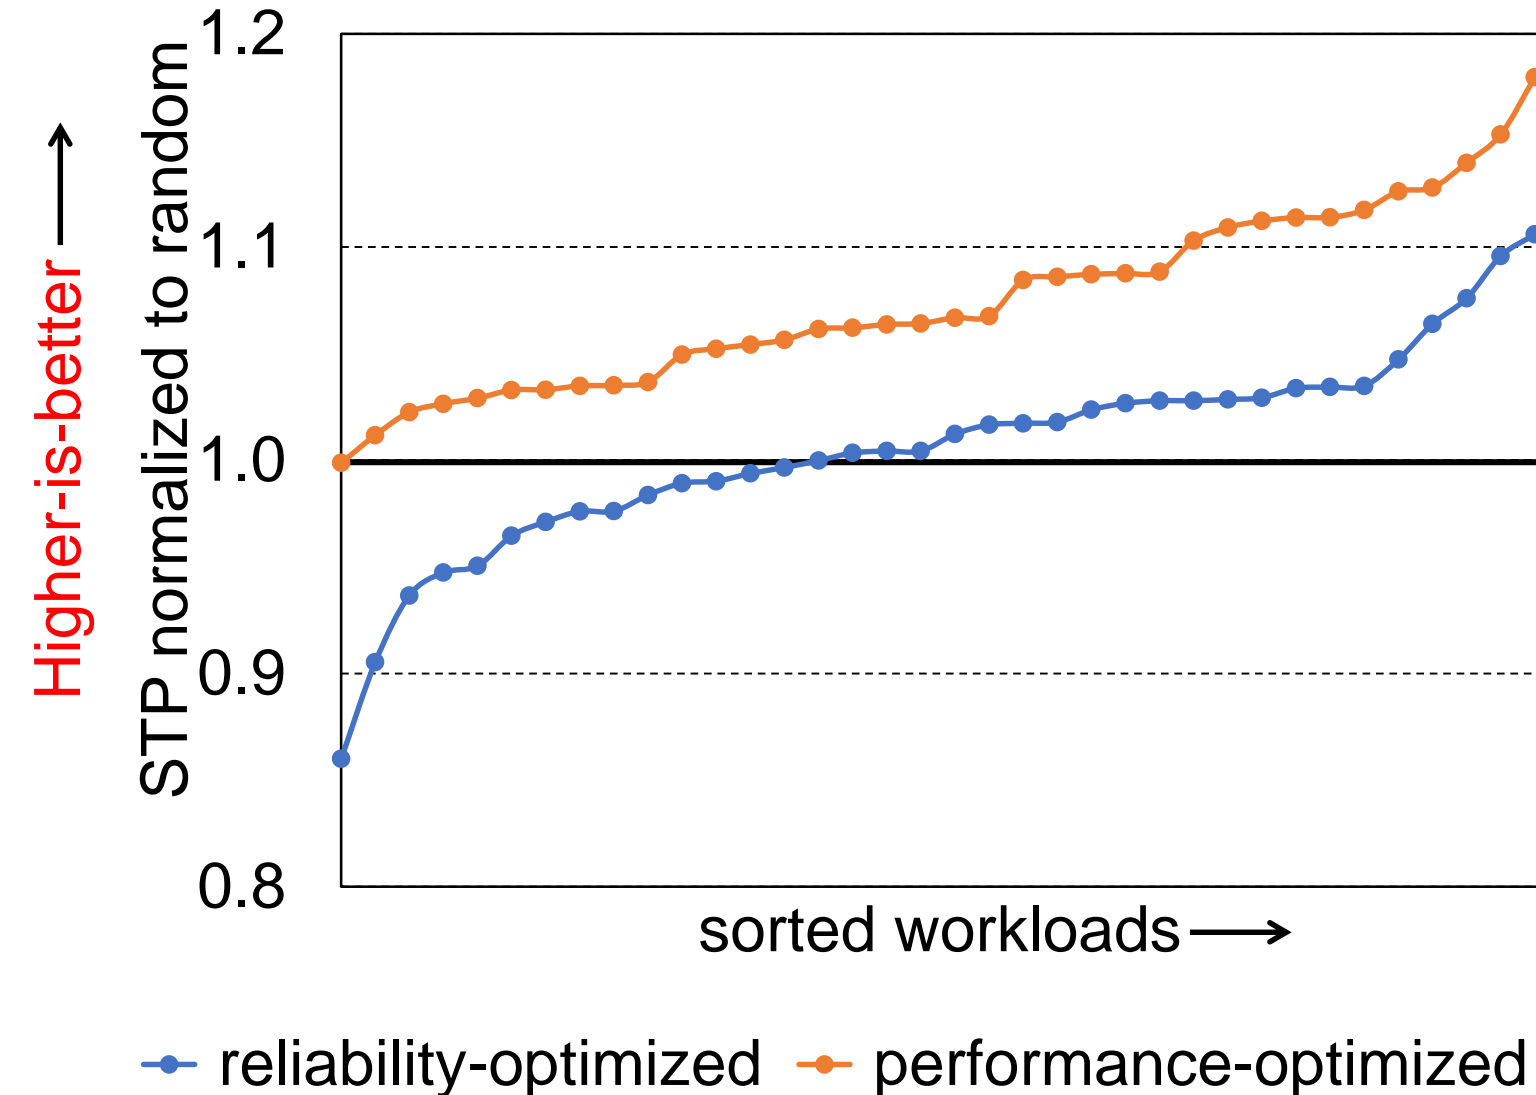


Compared to random

Reliability-optimized
-- does not degrade STP

Performance-optimized
-- improves STP by 6.3%
-- up to 18.7%

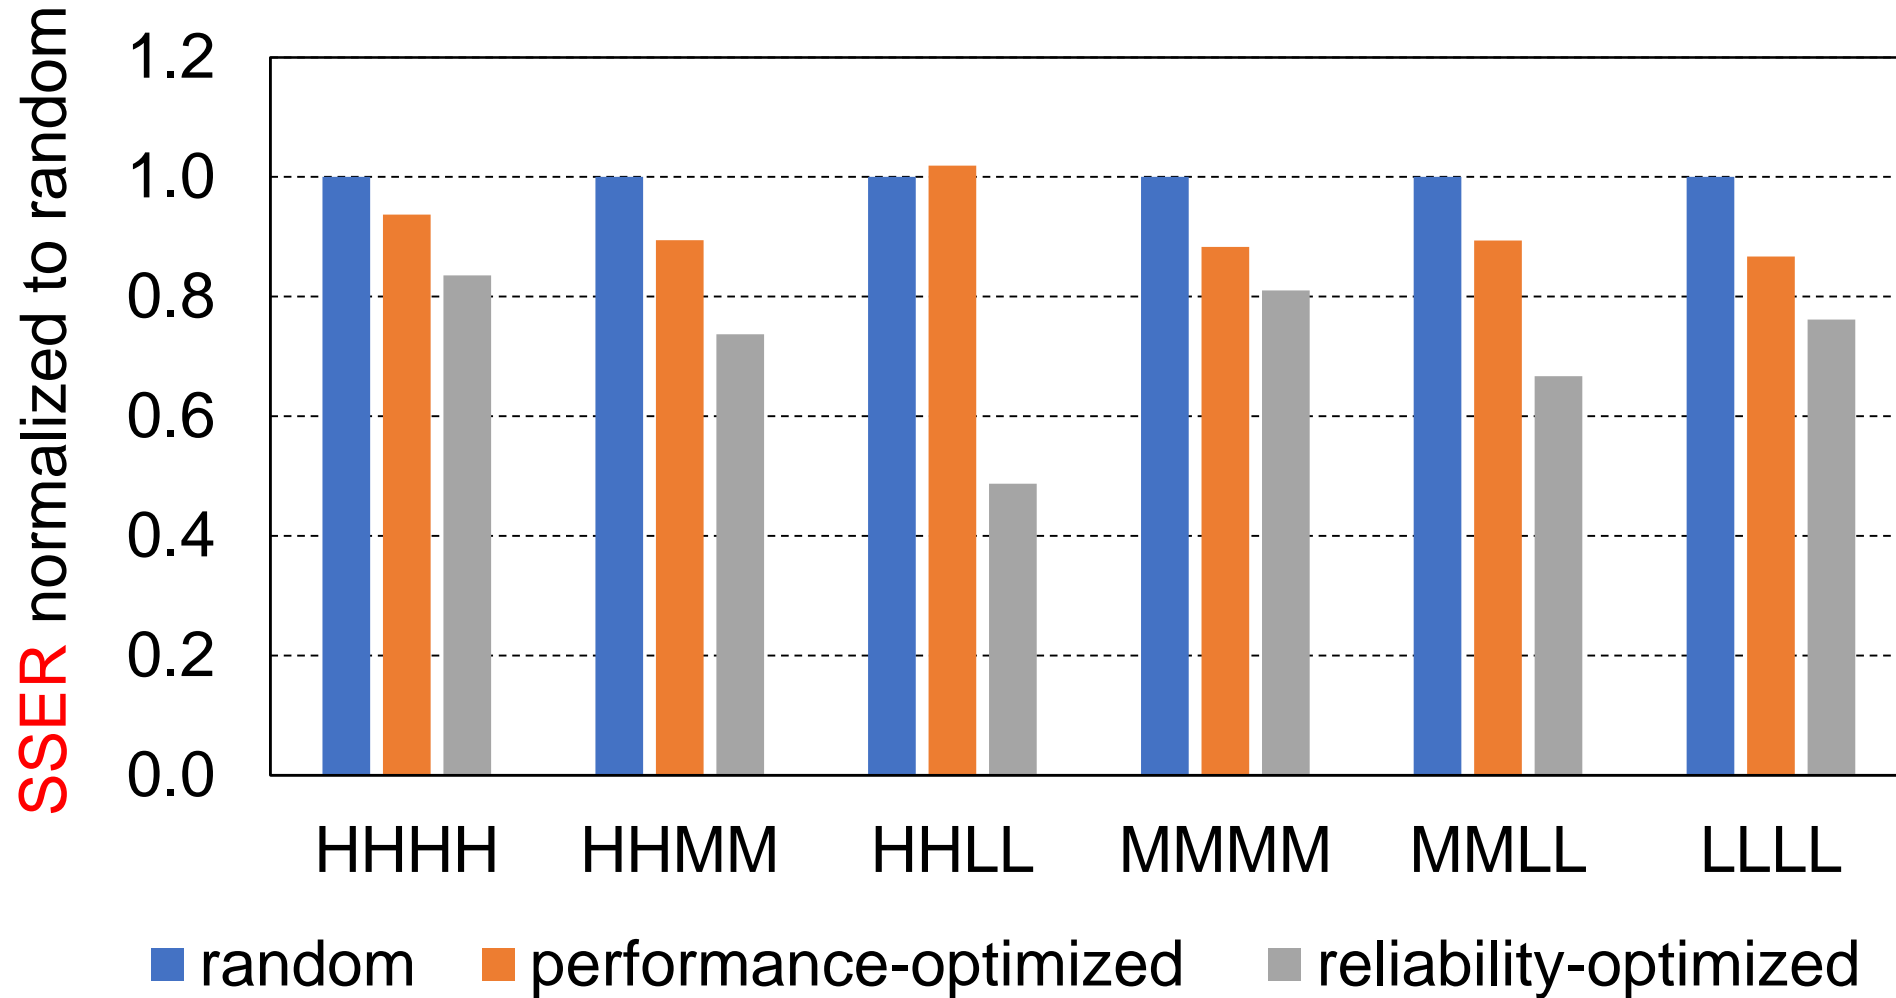
2B2S Results – Performance



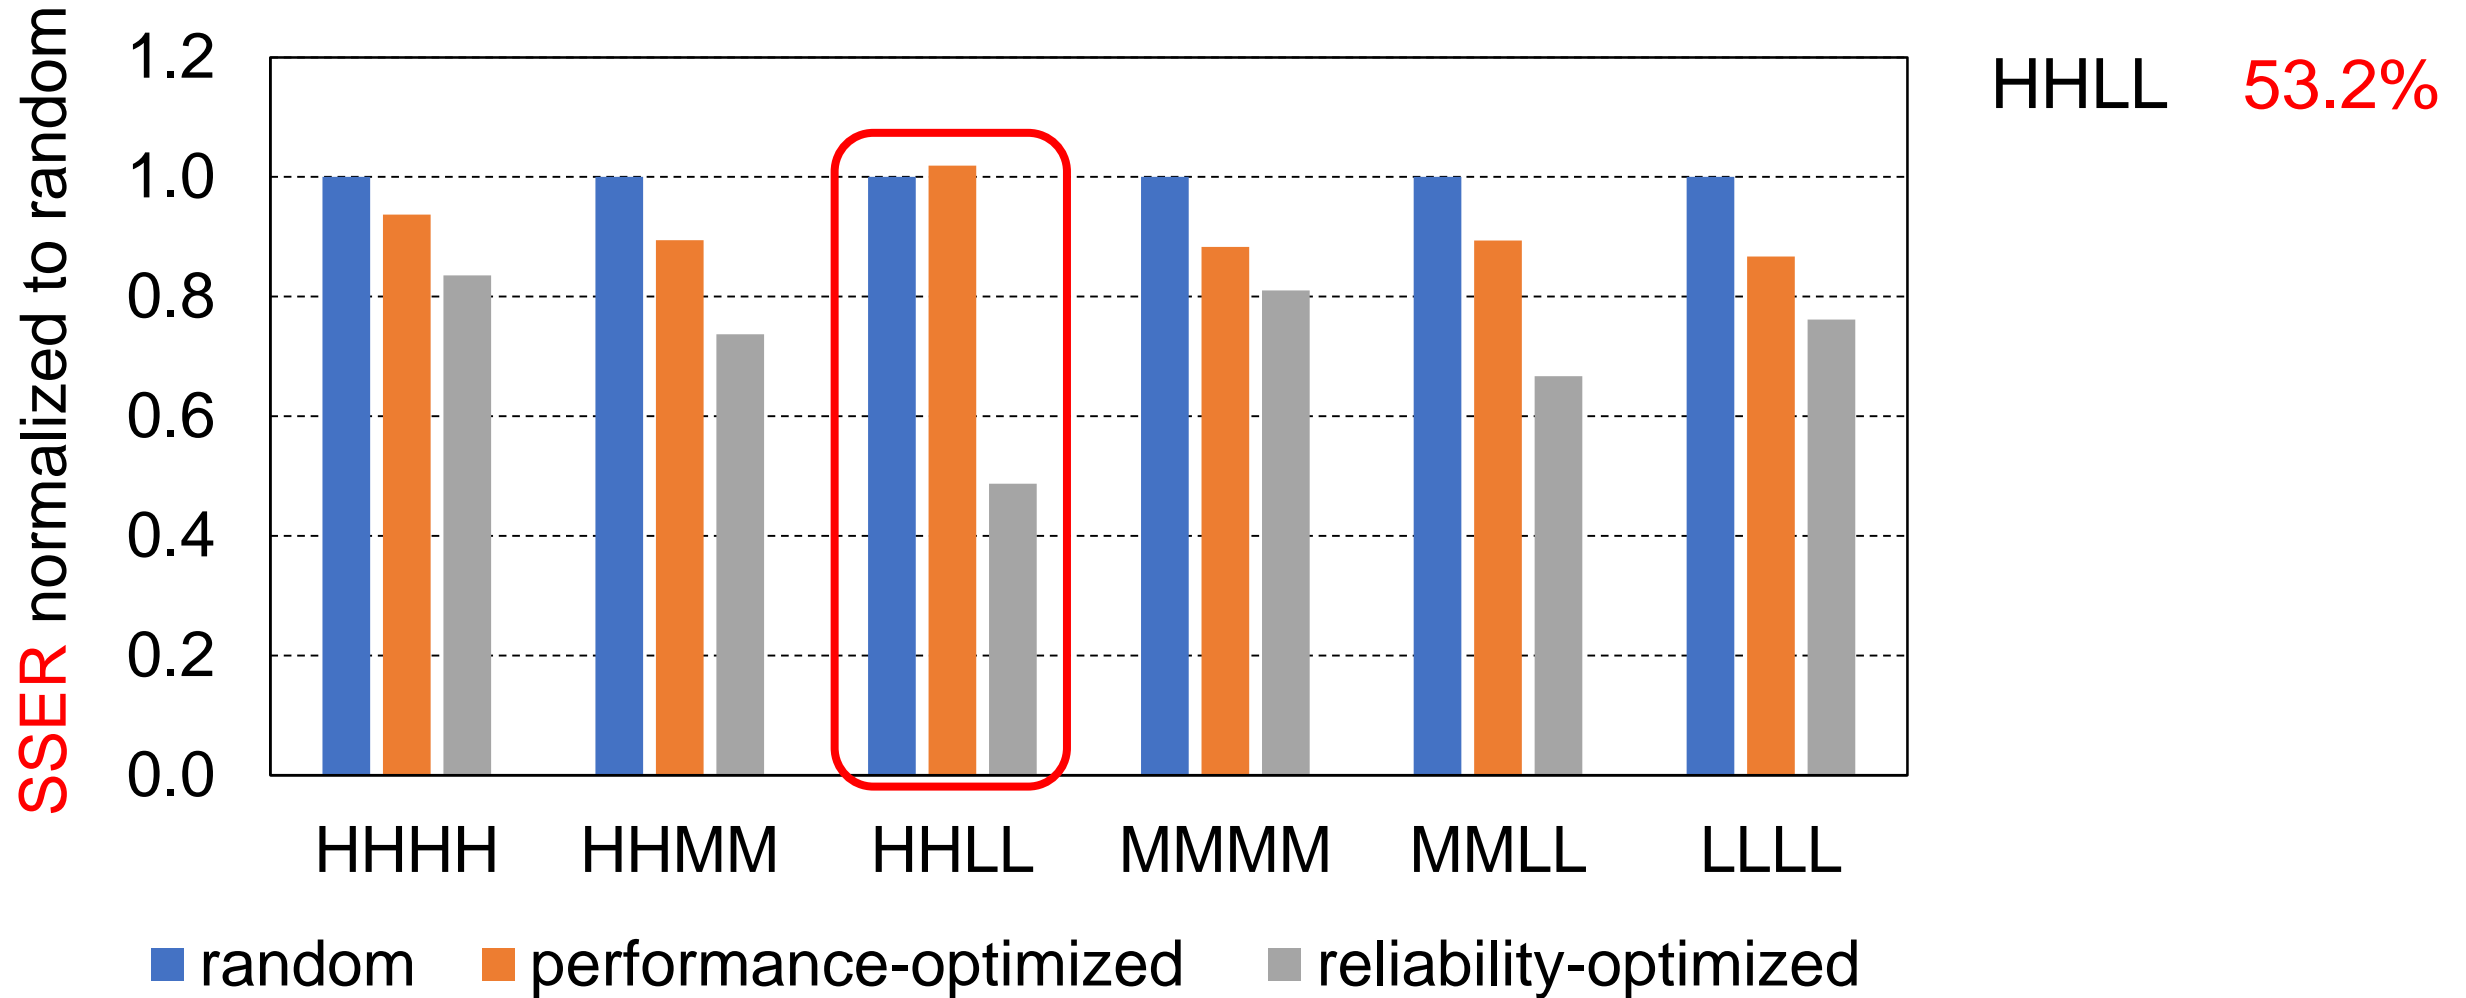
Compared to
performance-optimized

Reliability-optimized
-- degrades STP by 6.3%
-- up to 18.7%

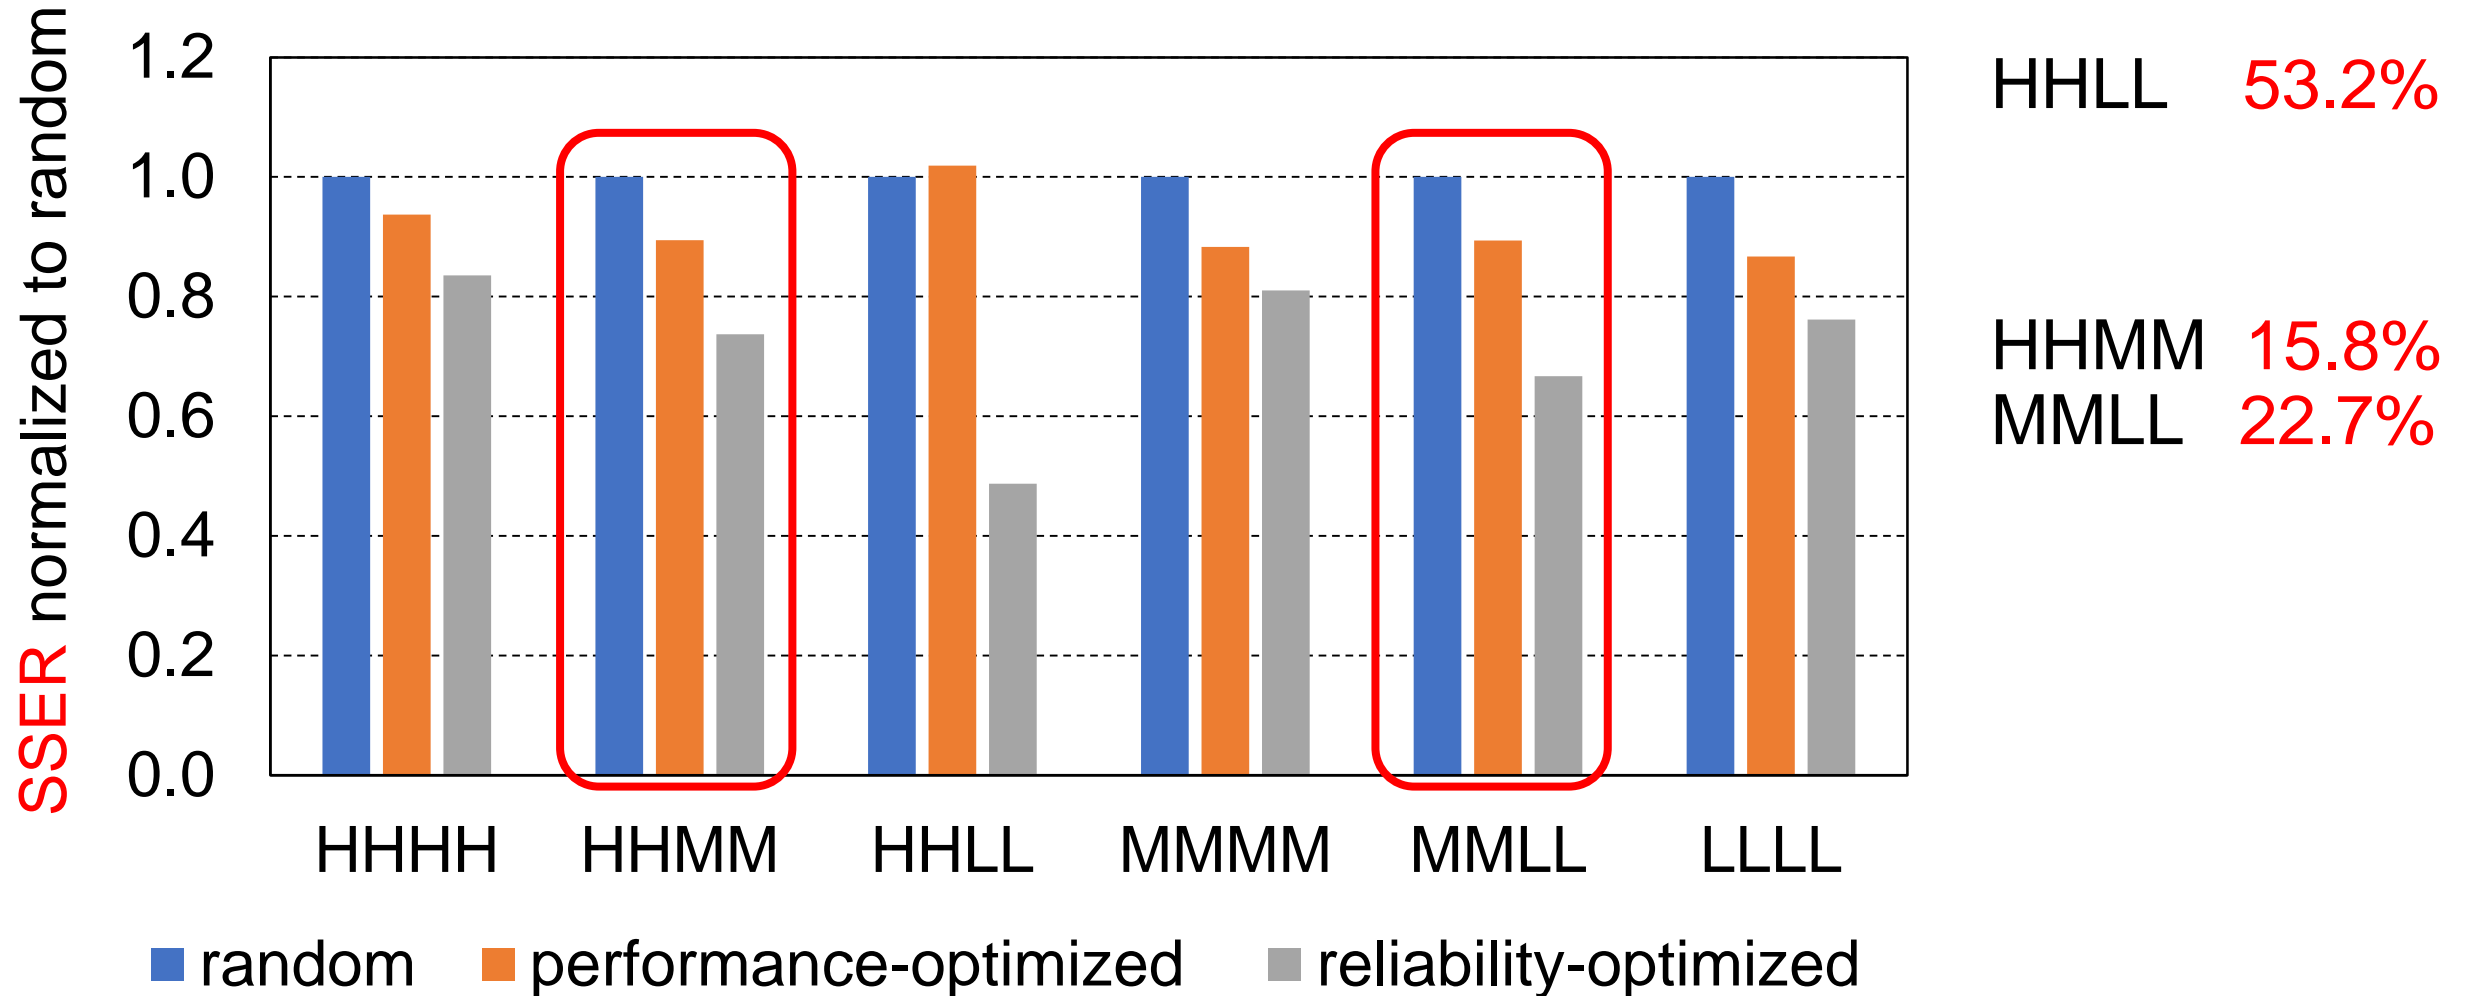
2B2S Results – by Workload Category



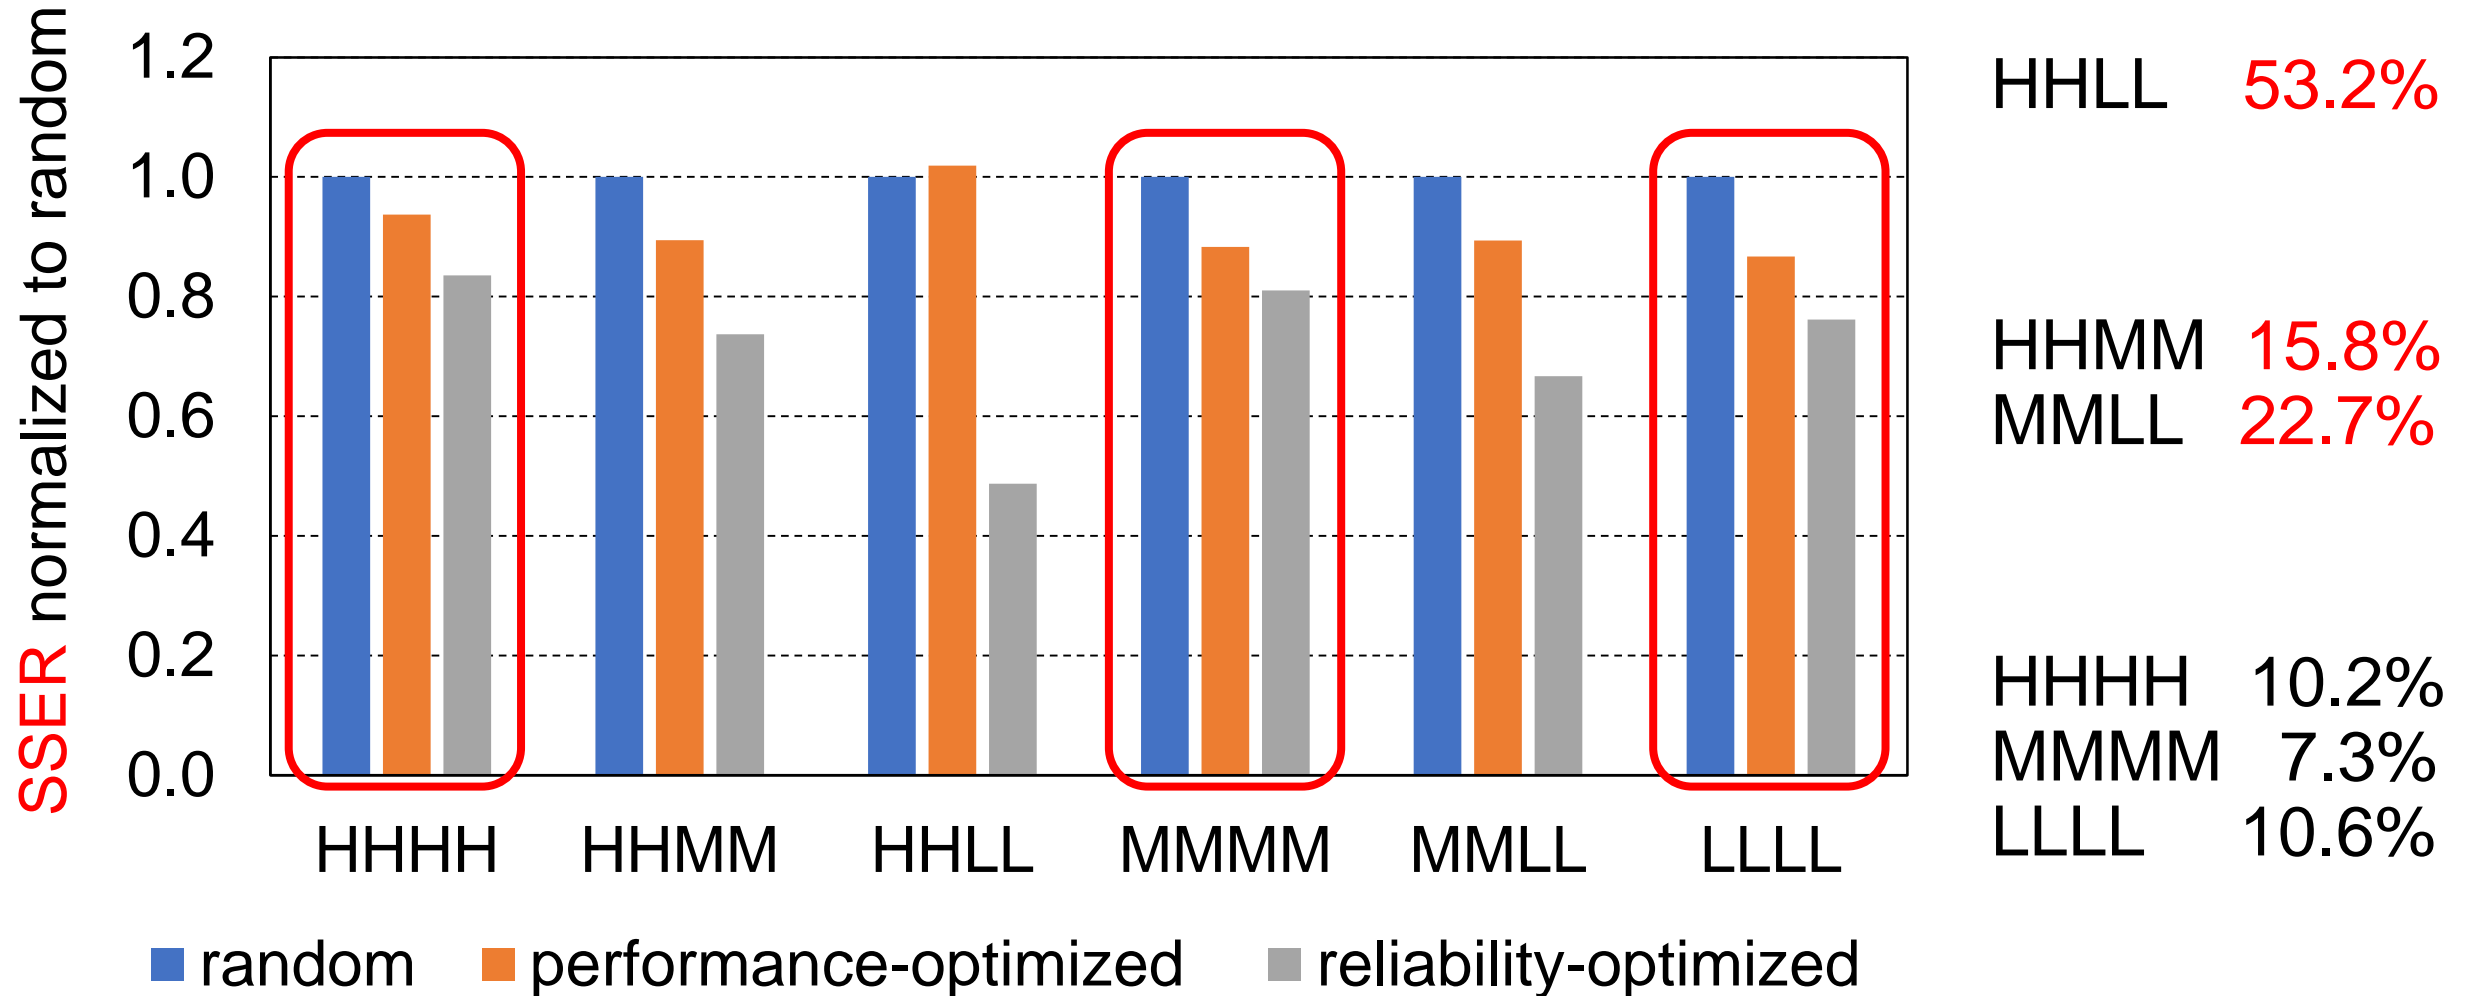
2B2S Results – by Workload Category



2B2S Results – by Workload Category



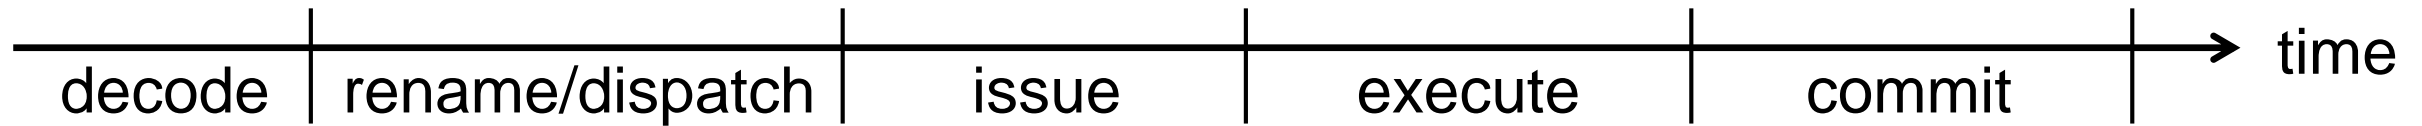
2B2S Results – by Workload Category



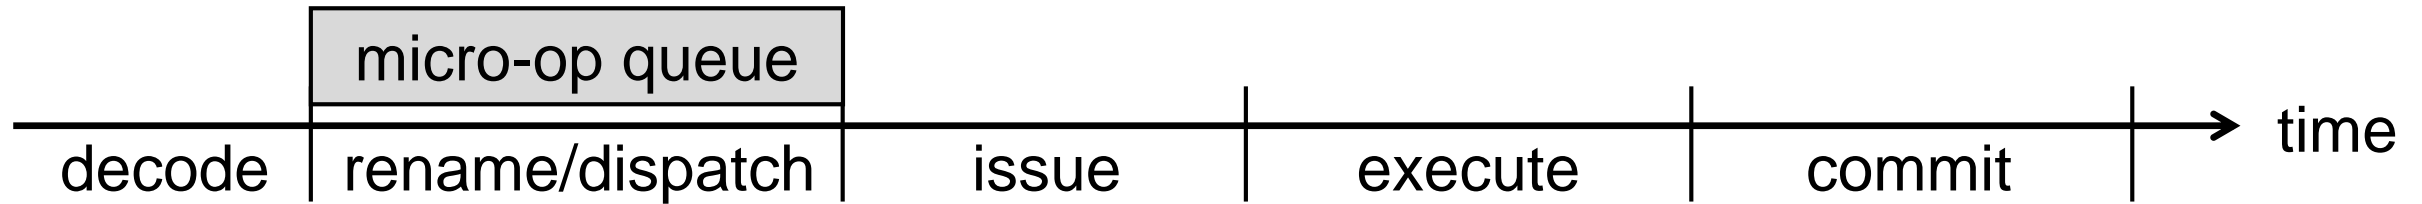
Contribution #2 [DSN 2020, Under review]

Dispatch Halting to Improve Out-of-Order Core Reliability

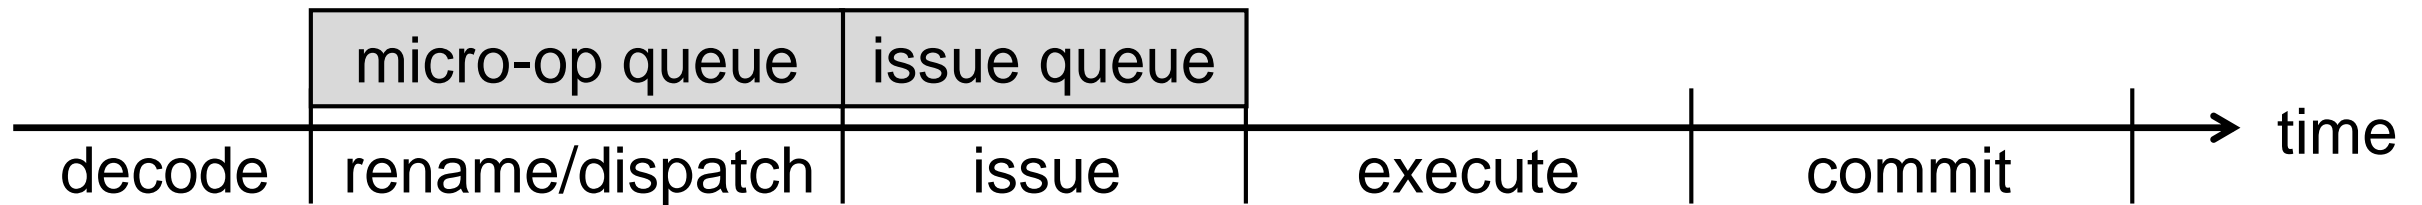
Vulnerability of Out-of-Order Core Resources



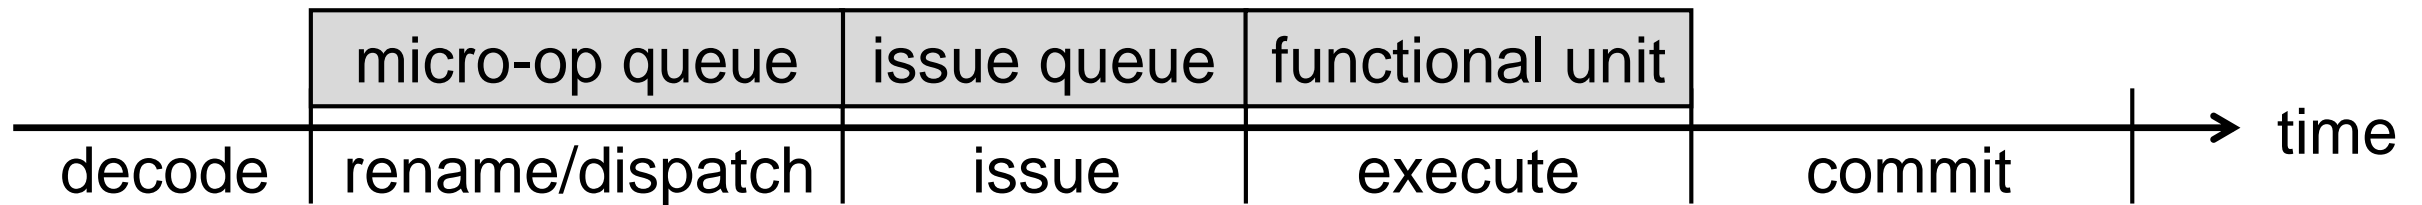
Vulnerability of Out-of-Order Core Resources



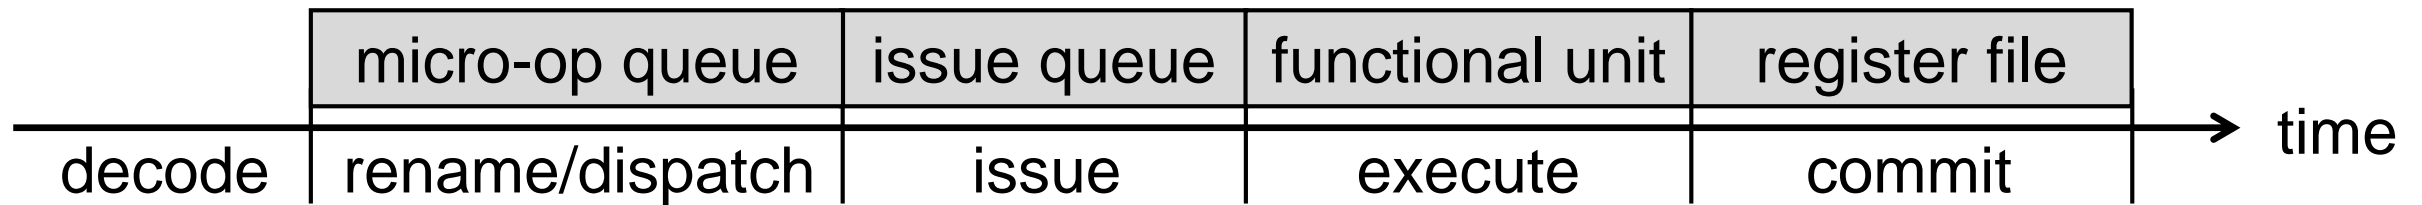
Vulnerability of Out-of-Order Core Resources



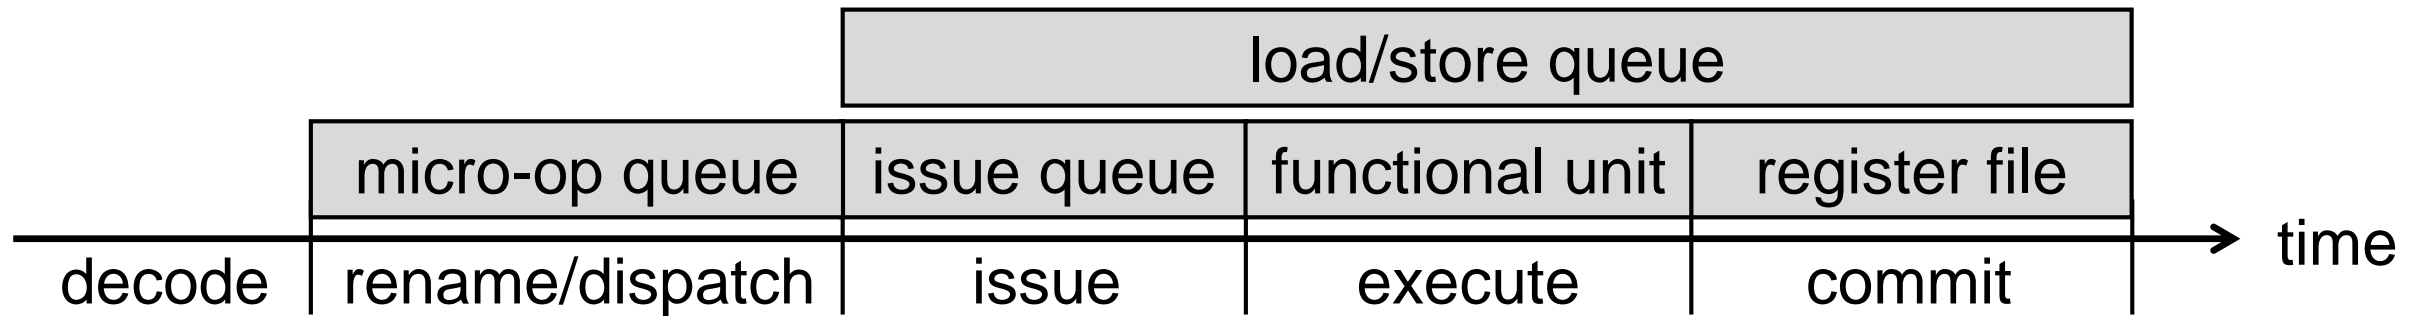
Vulnerability of Out-of-Order Core Resources



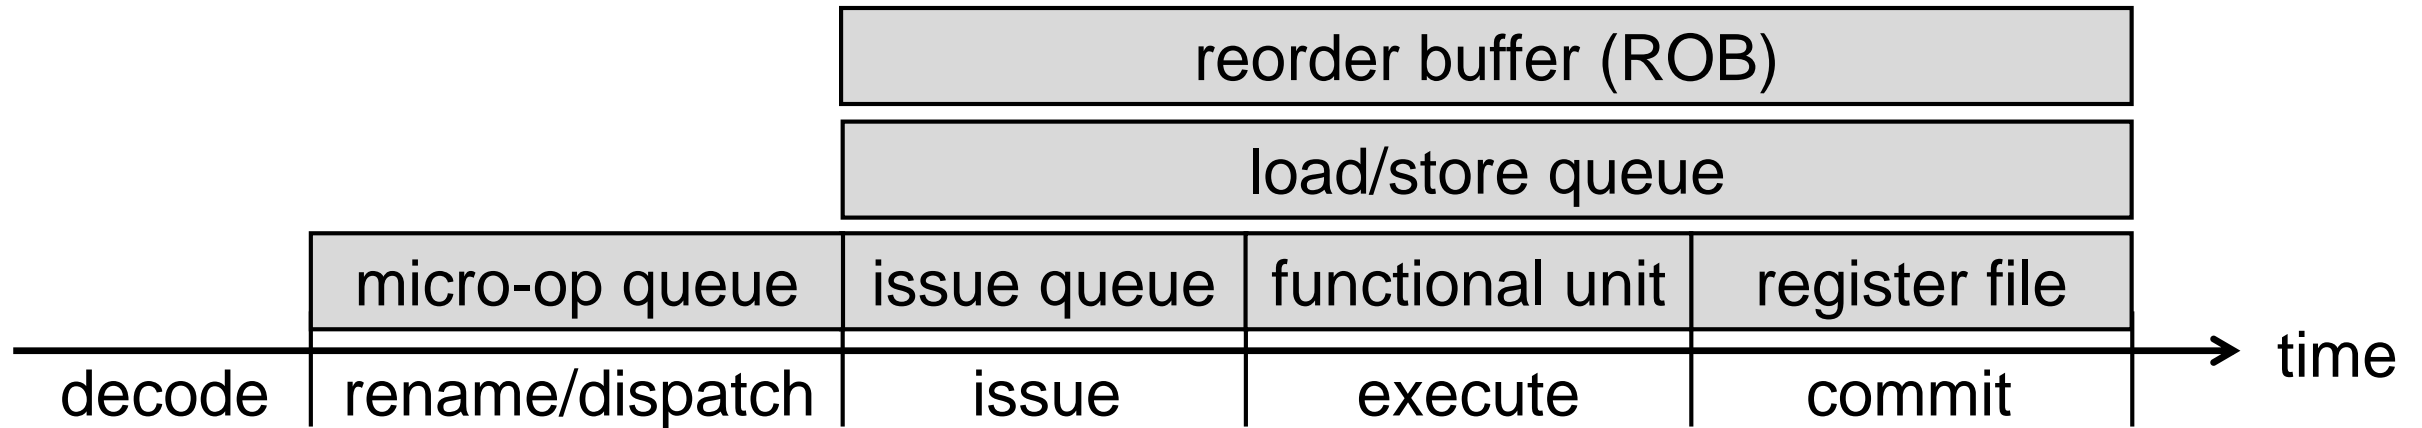
Vulnerability of Out-of-Order Core Resources



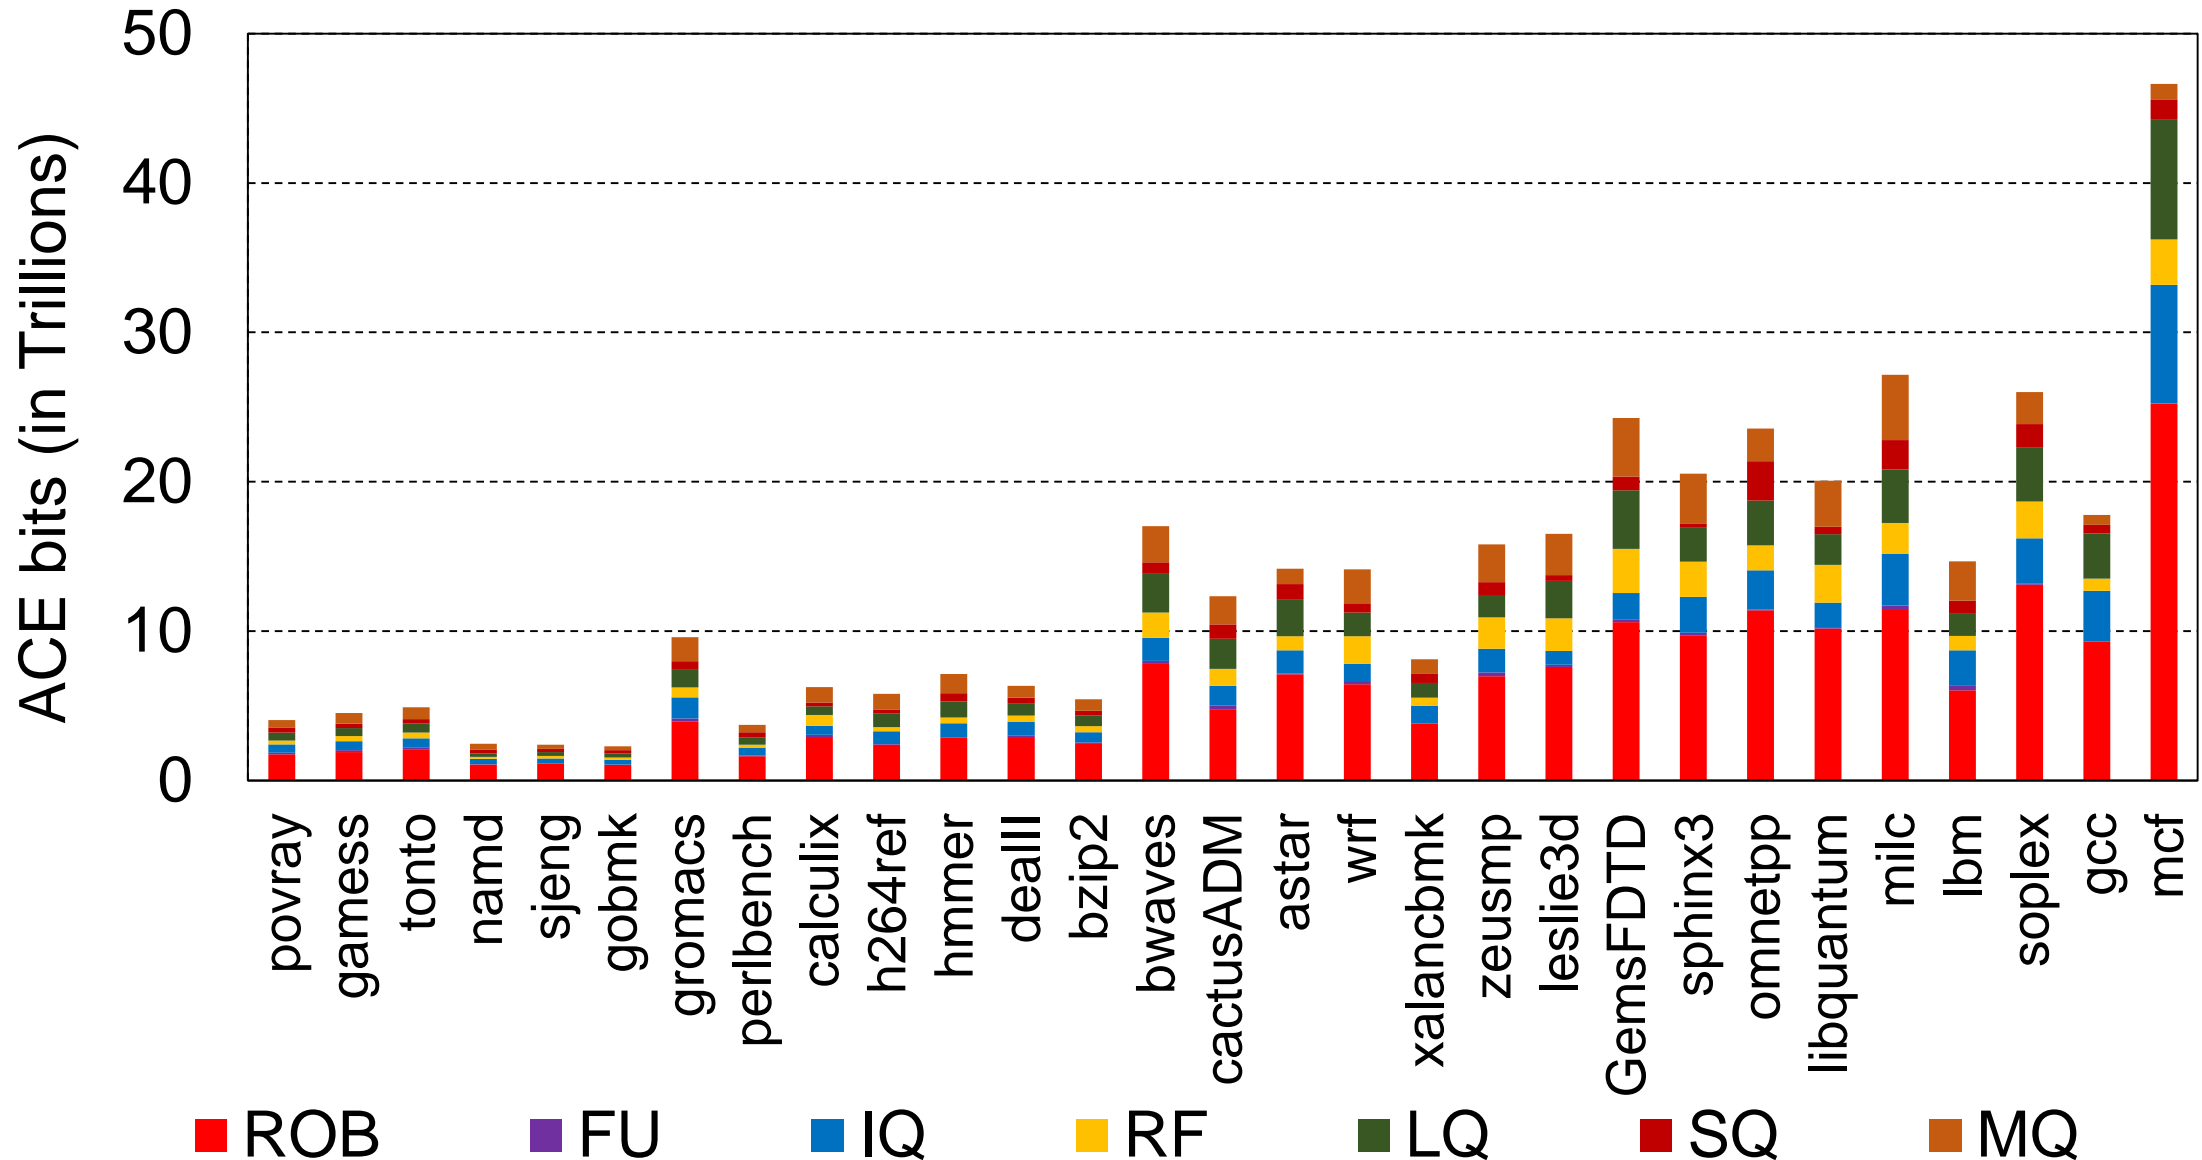
Vulnerability of Out-of-Order Core Resources



Vulnerability of Out-of-Order Core Resources



ACE Bit Count for OoO Resources



Vulnerable State under Full-ROB Stalls

Vulnerable State under Full-ROB Stalls

—————→ time

Vulnerable State under Full-ROB Stalls



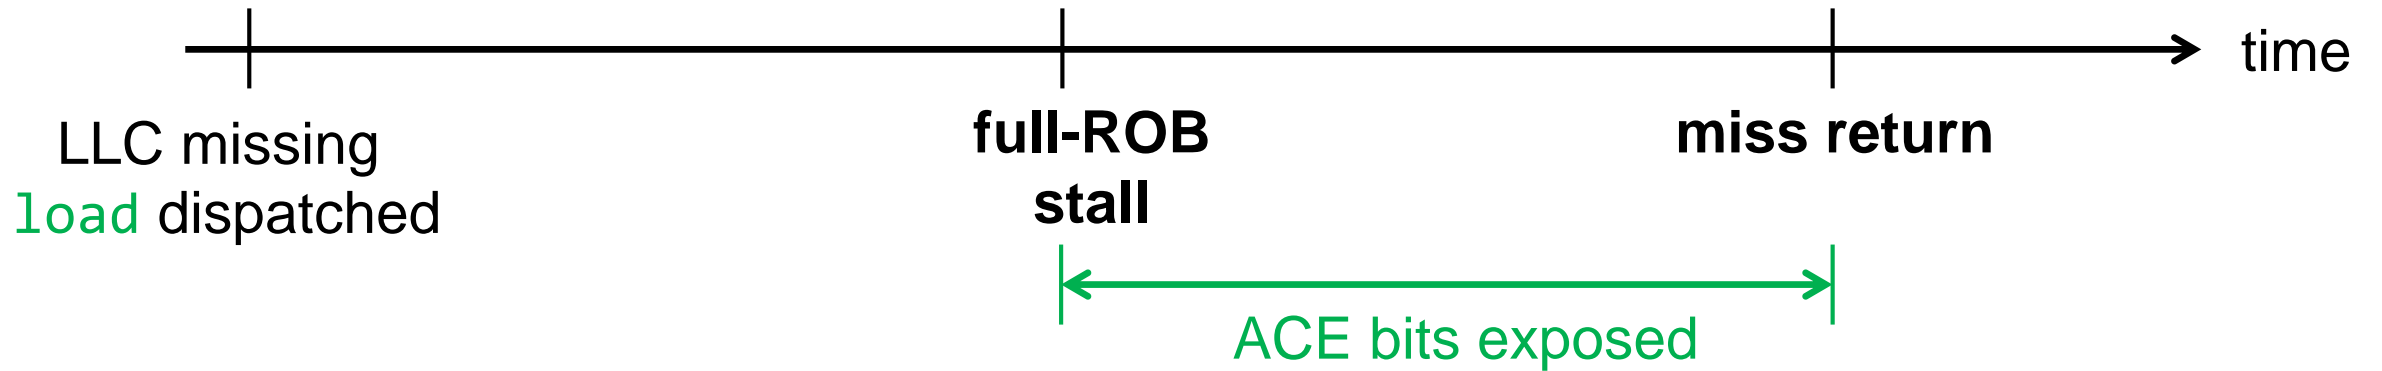
Vulnerable State under Full-ROB Stalls



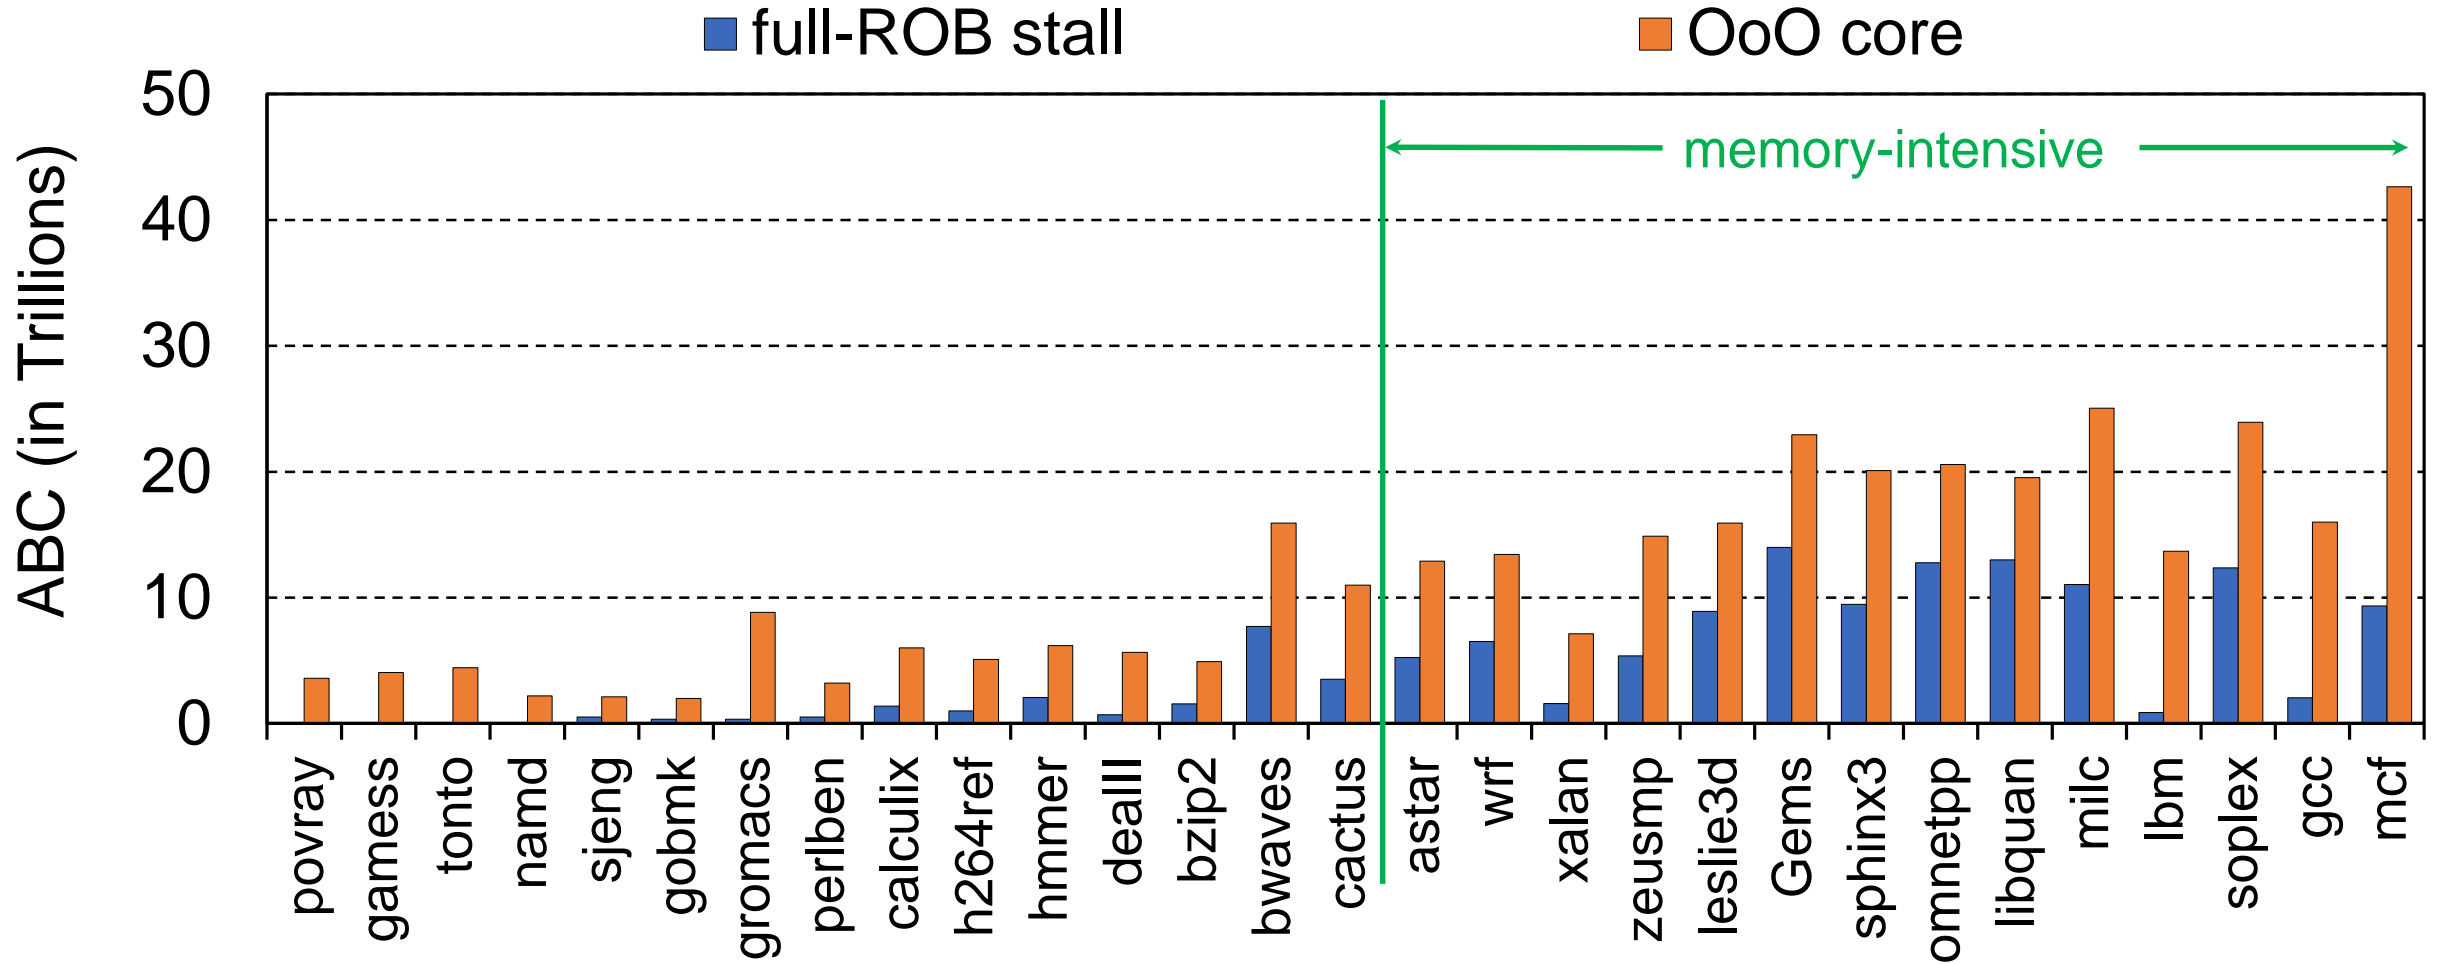
Vulnerable State under Full-ROB Stalls



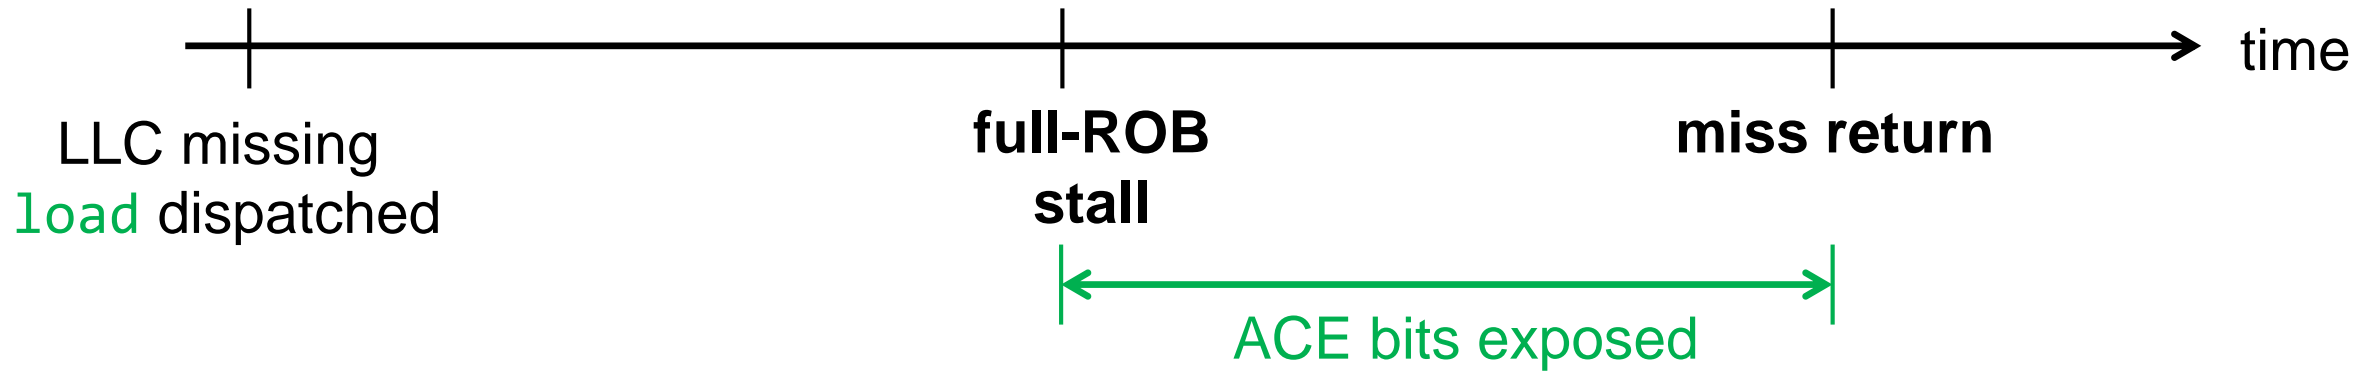
Vulnerable State under Full-ROB Stalls



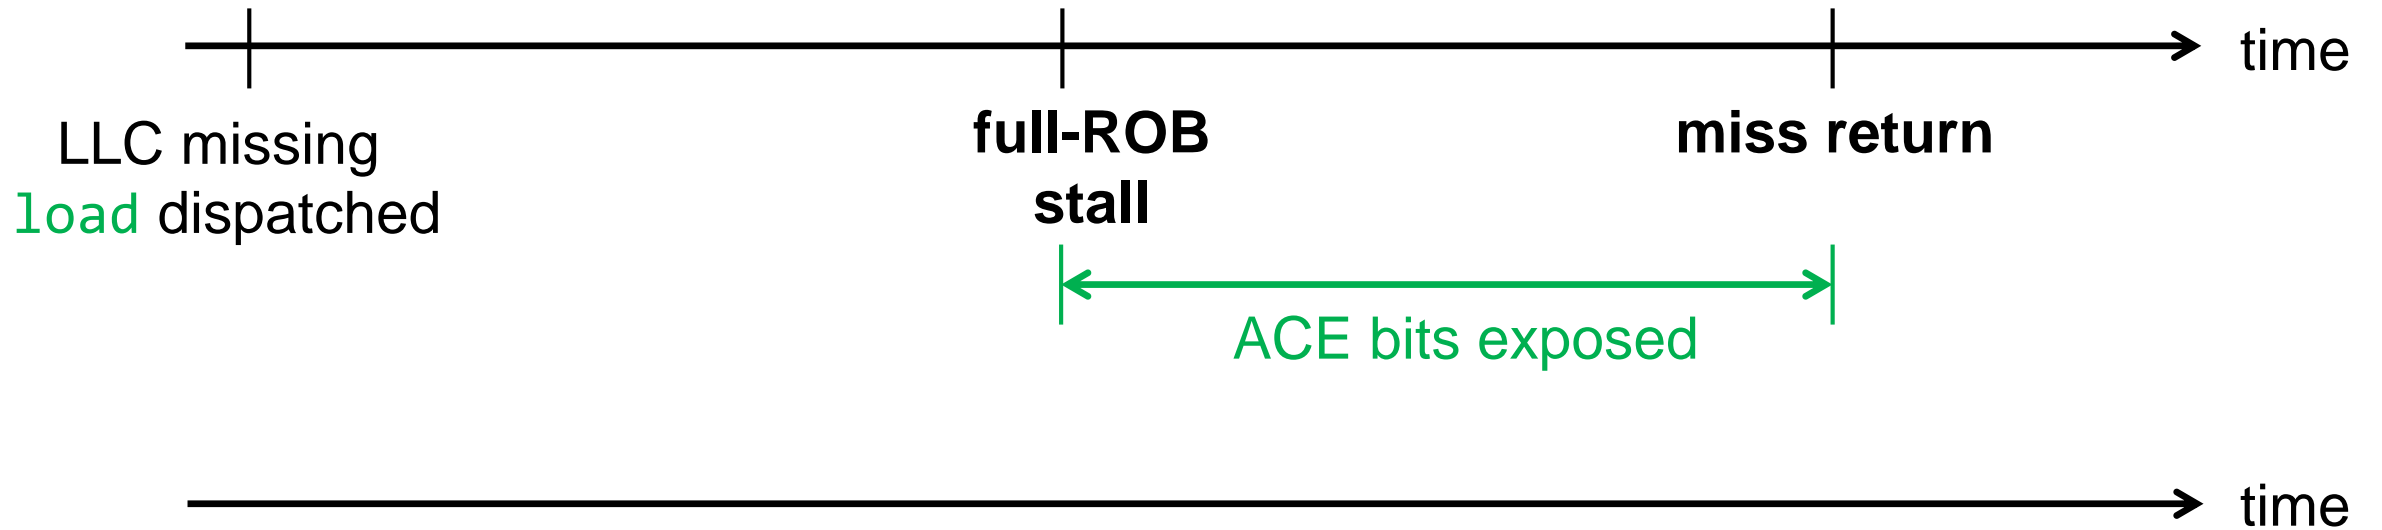
Vulnerable State under Full-ROB Stalls



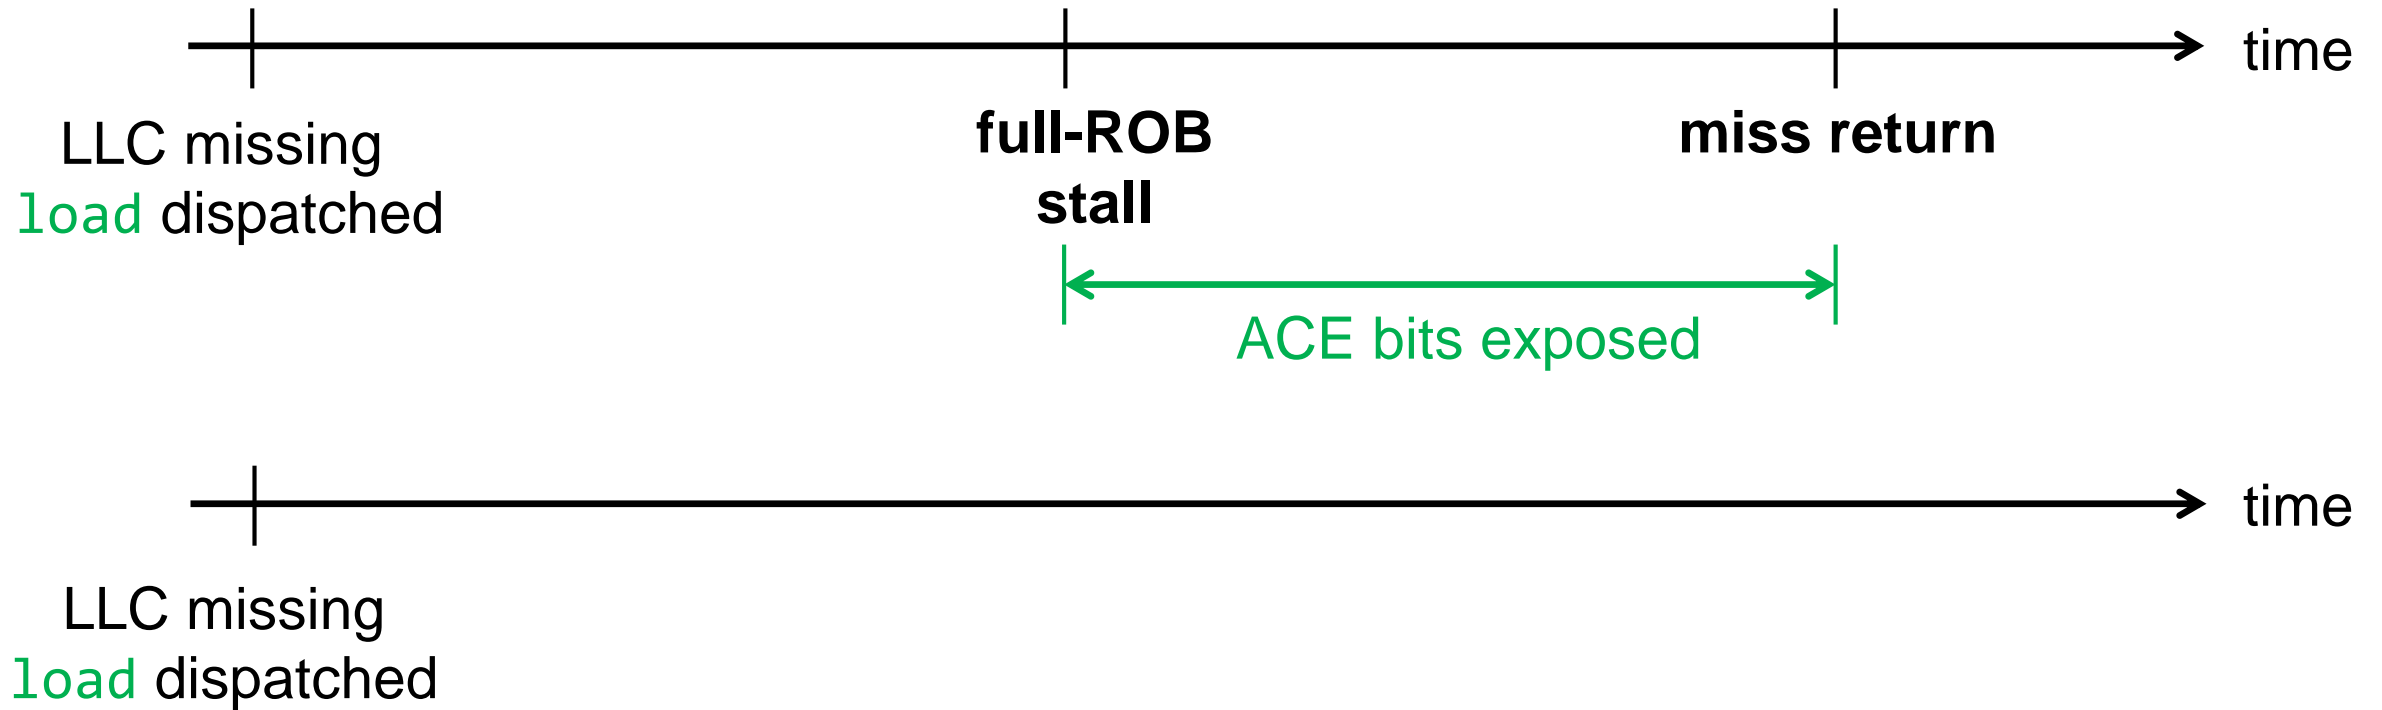
Vulnerable State under Blocked ROB Head



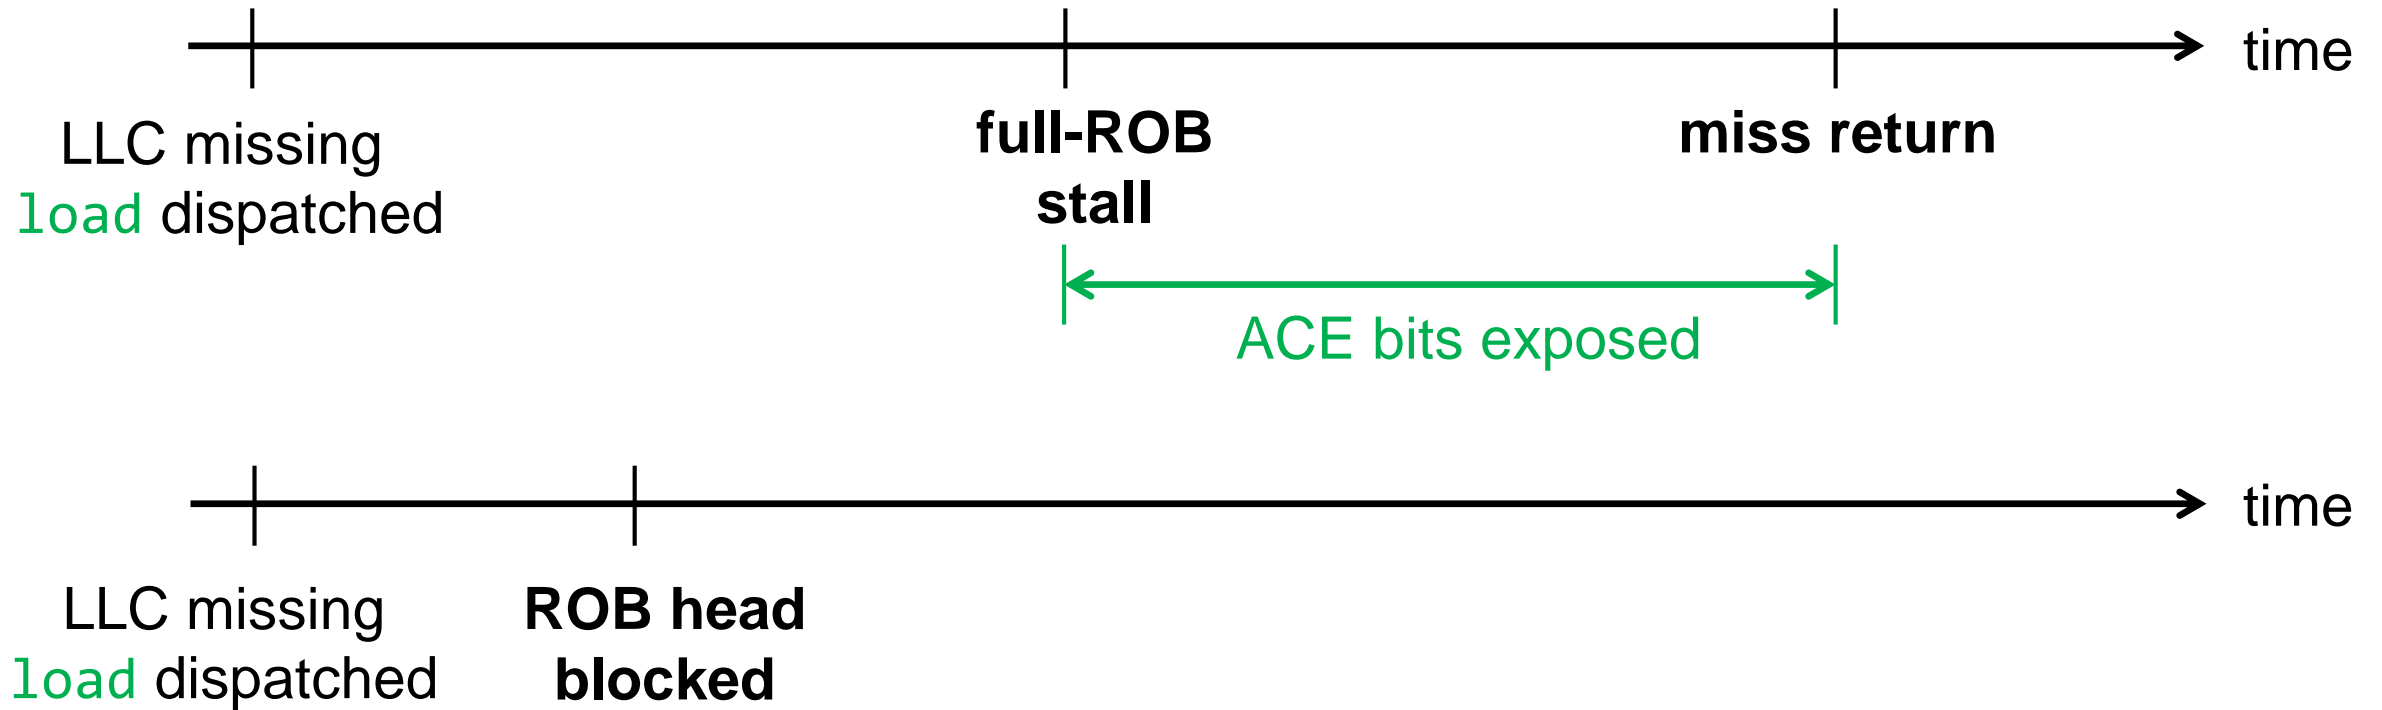
Vulnerable State under Blocked ROB Head



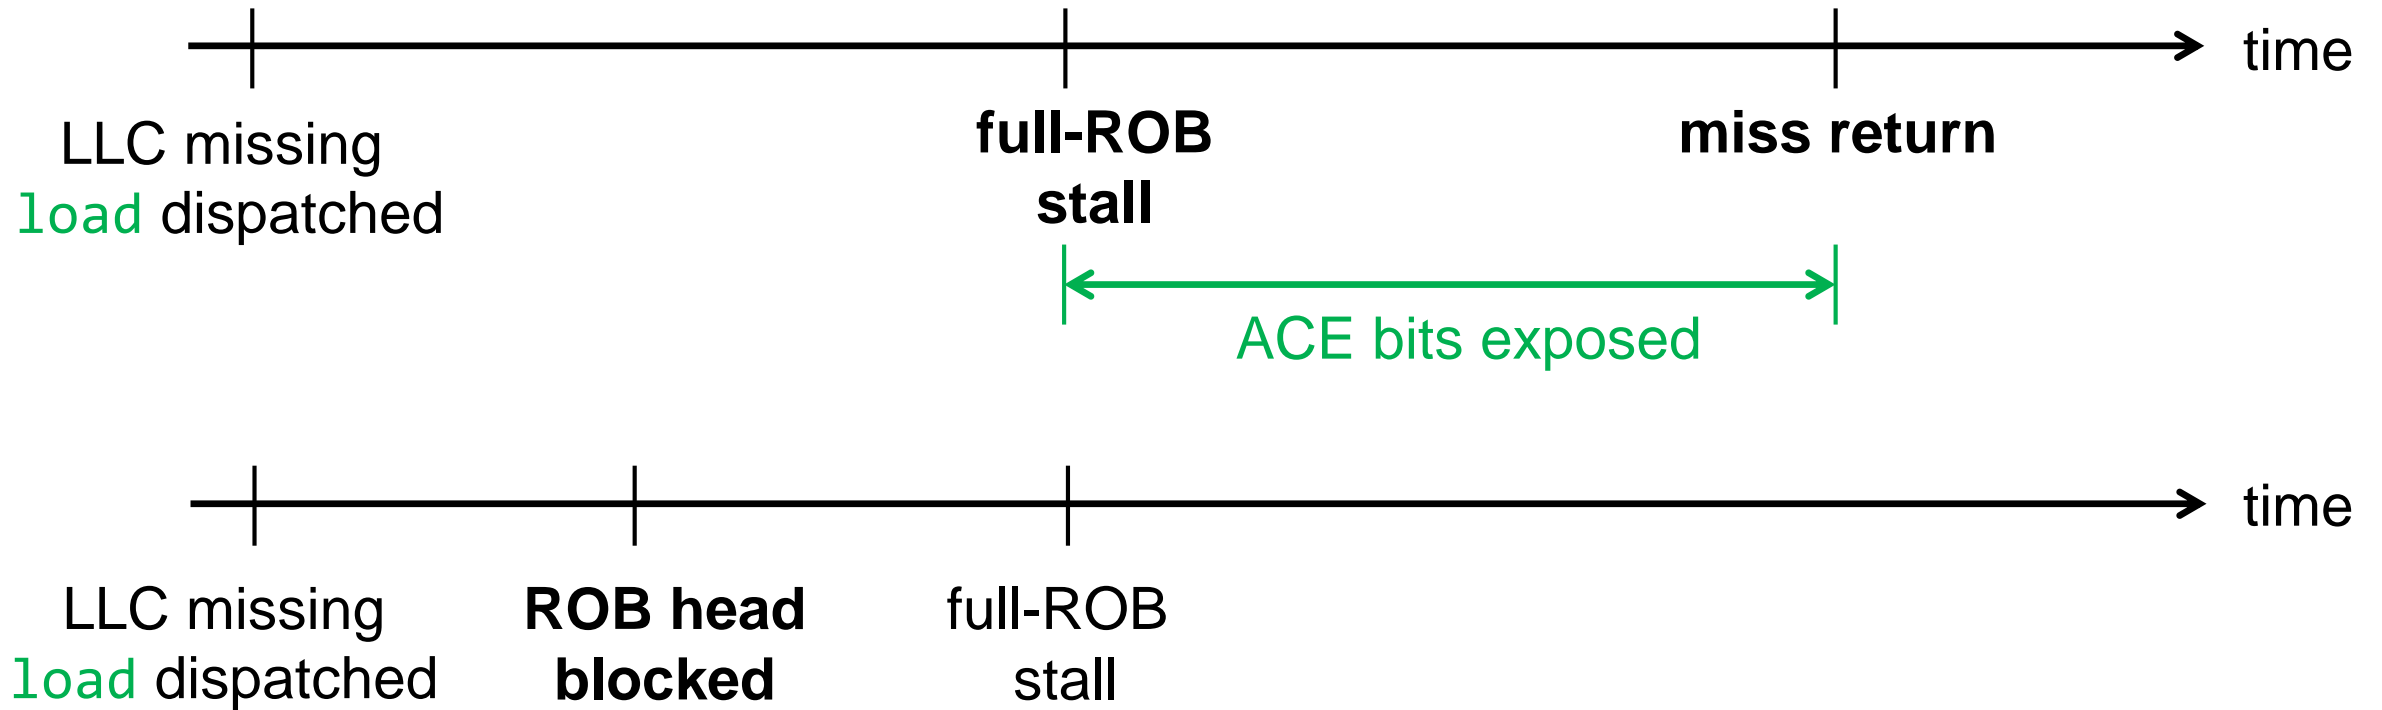
Vulnerable State under Blocked ROB Head



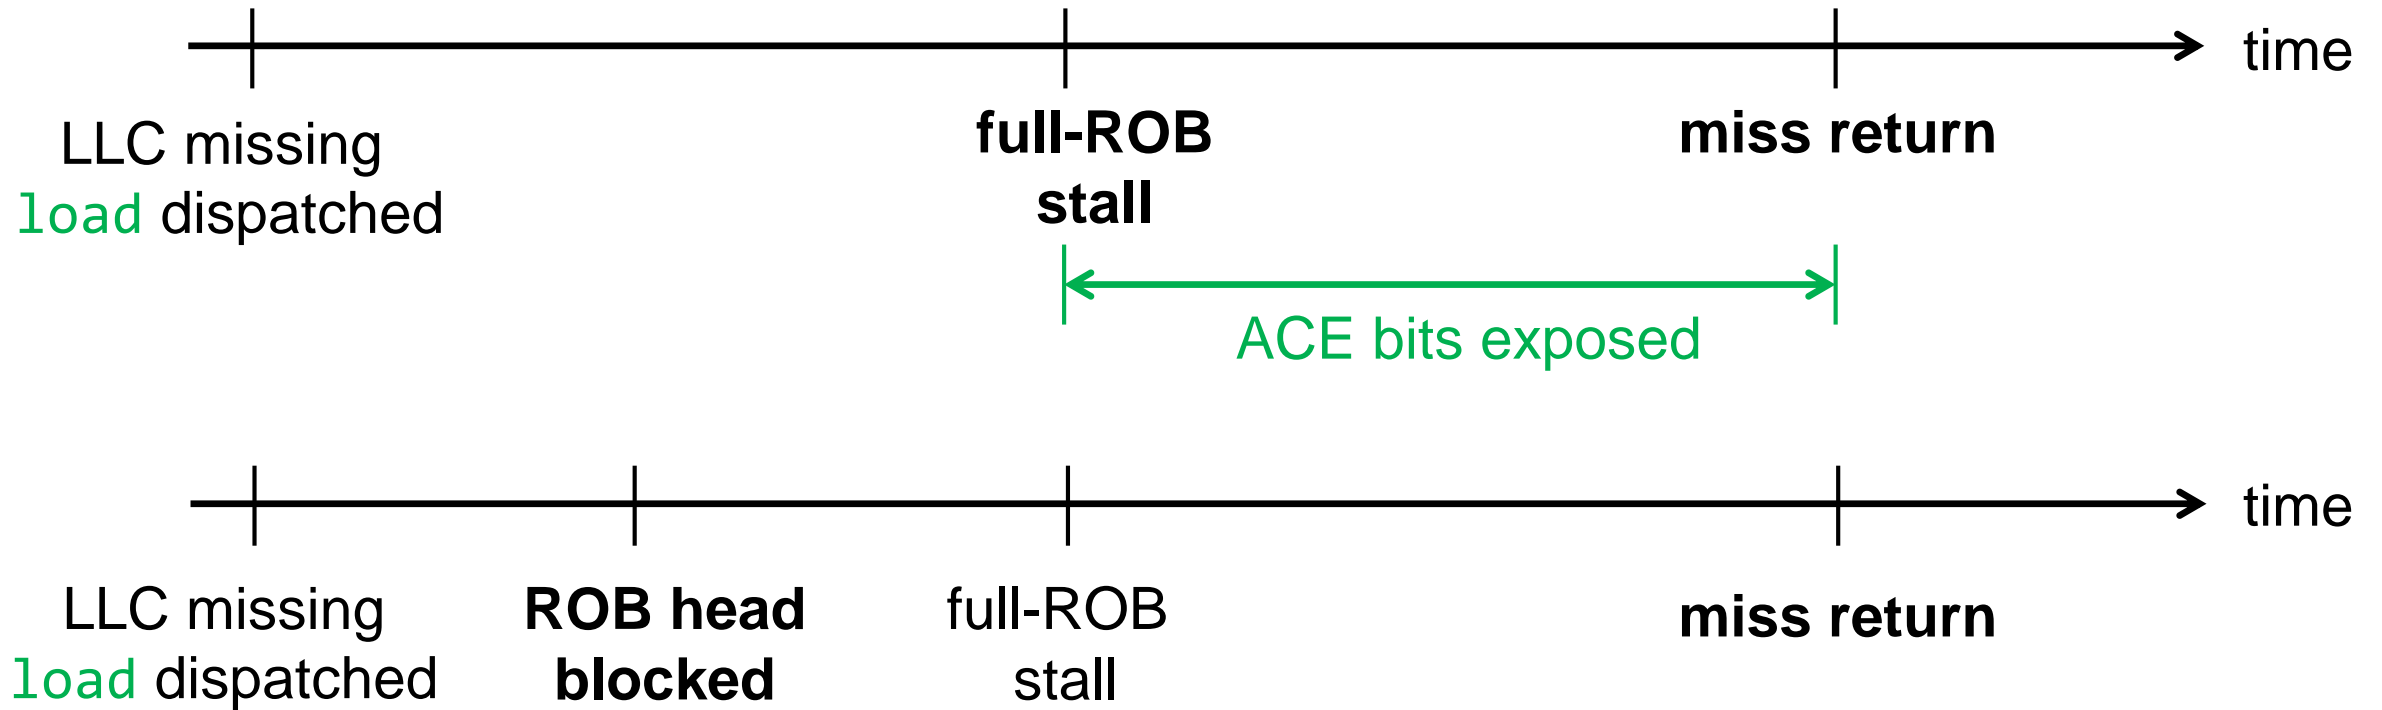
Vulnerable State under Blocked ROB Head



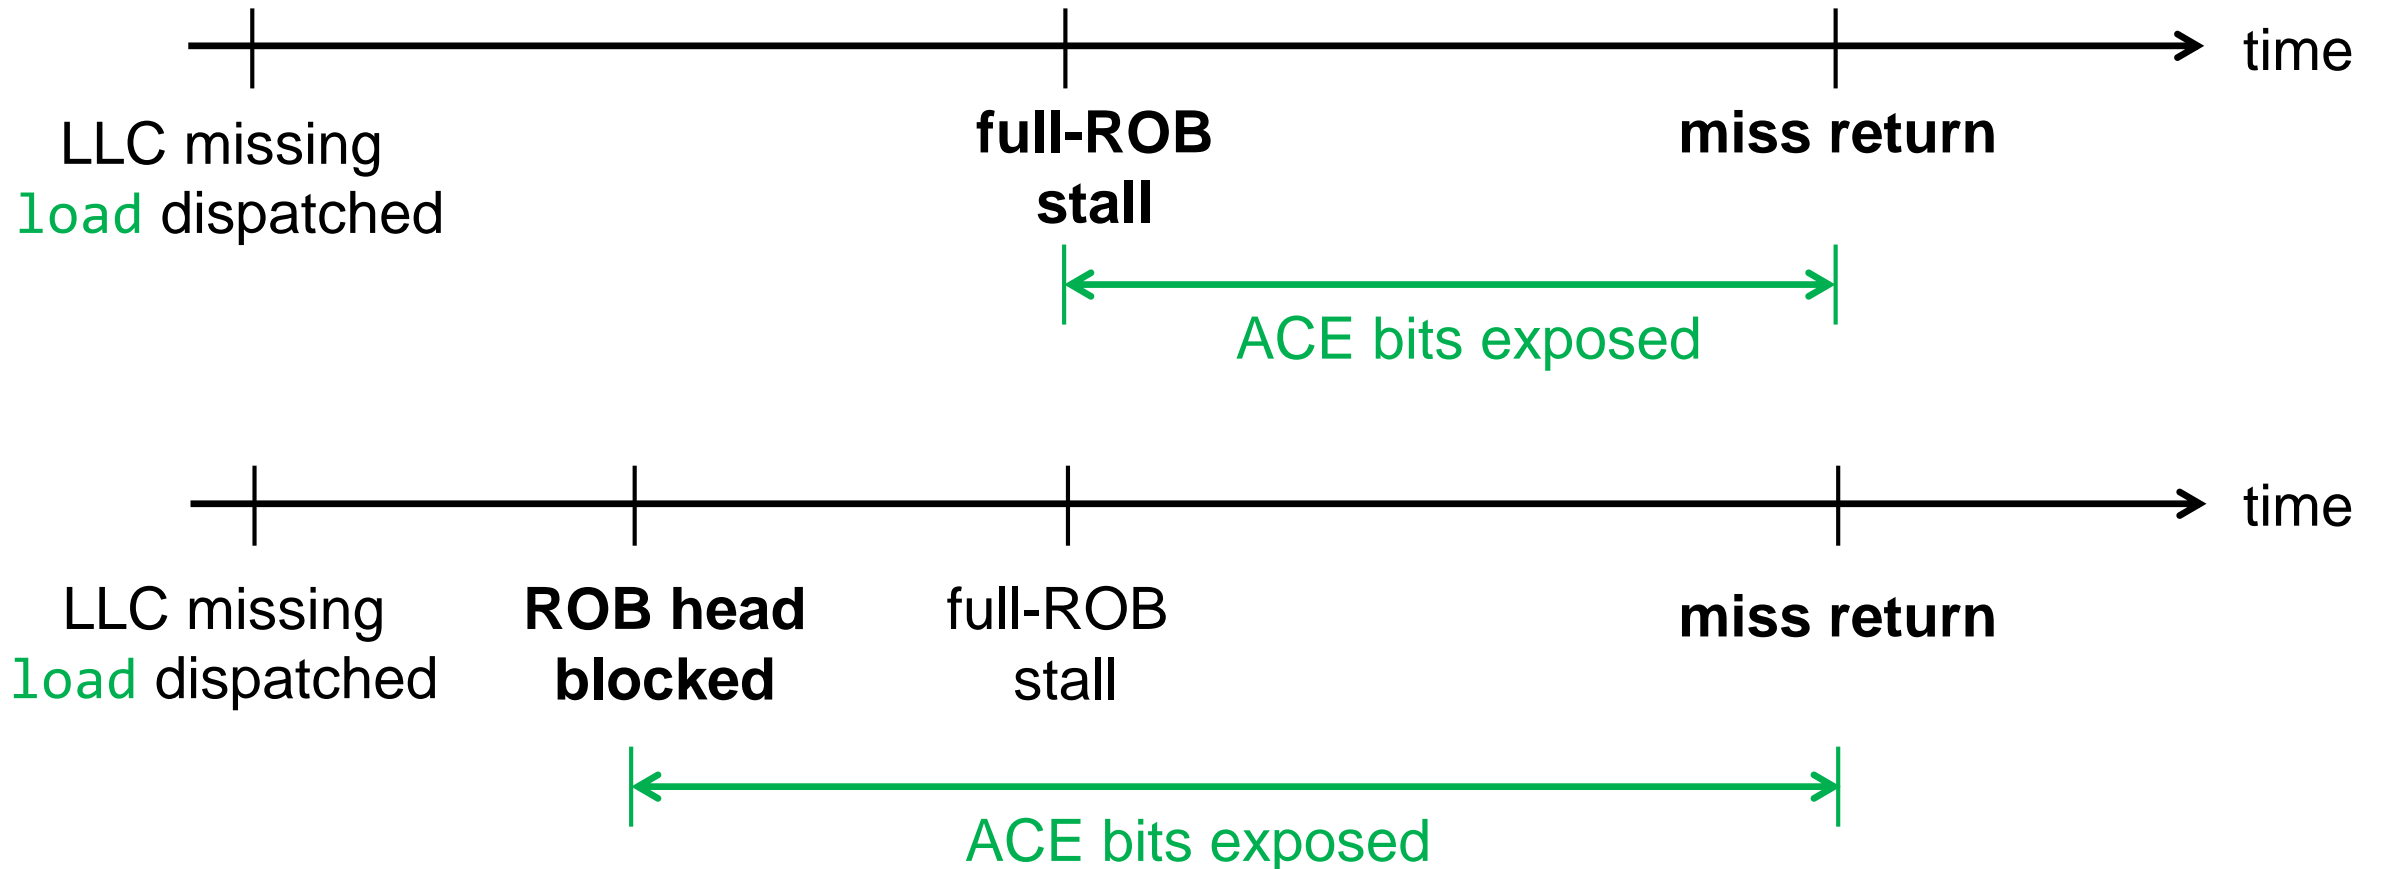
Vulnerable State under Blocked ROB Head



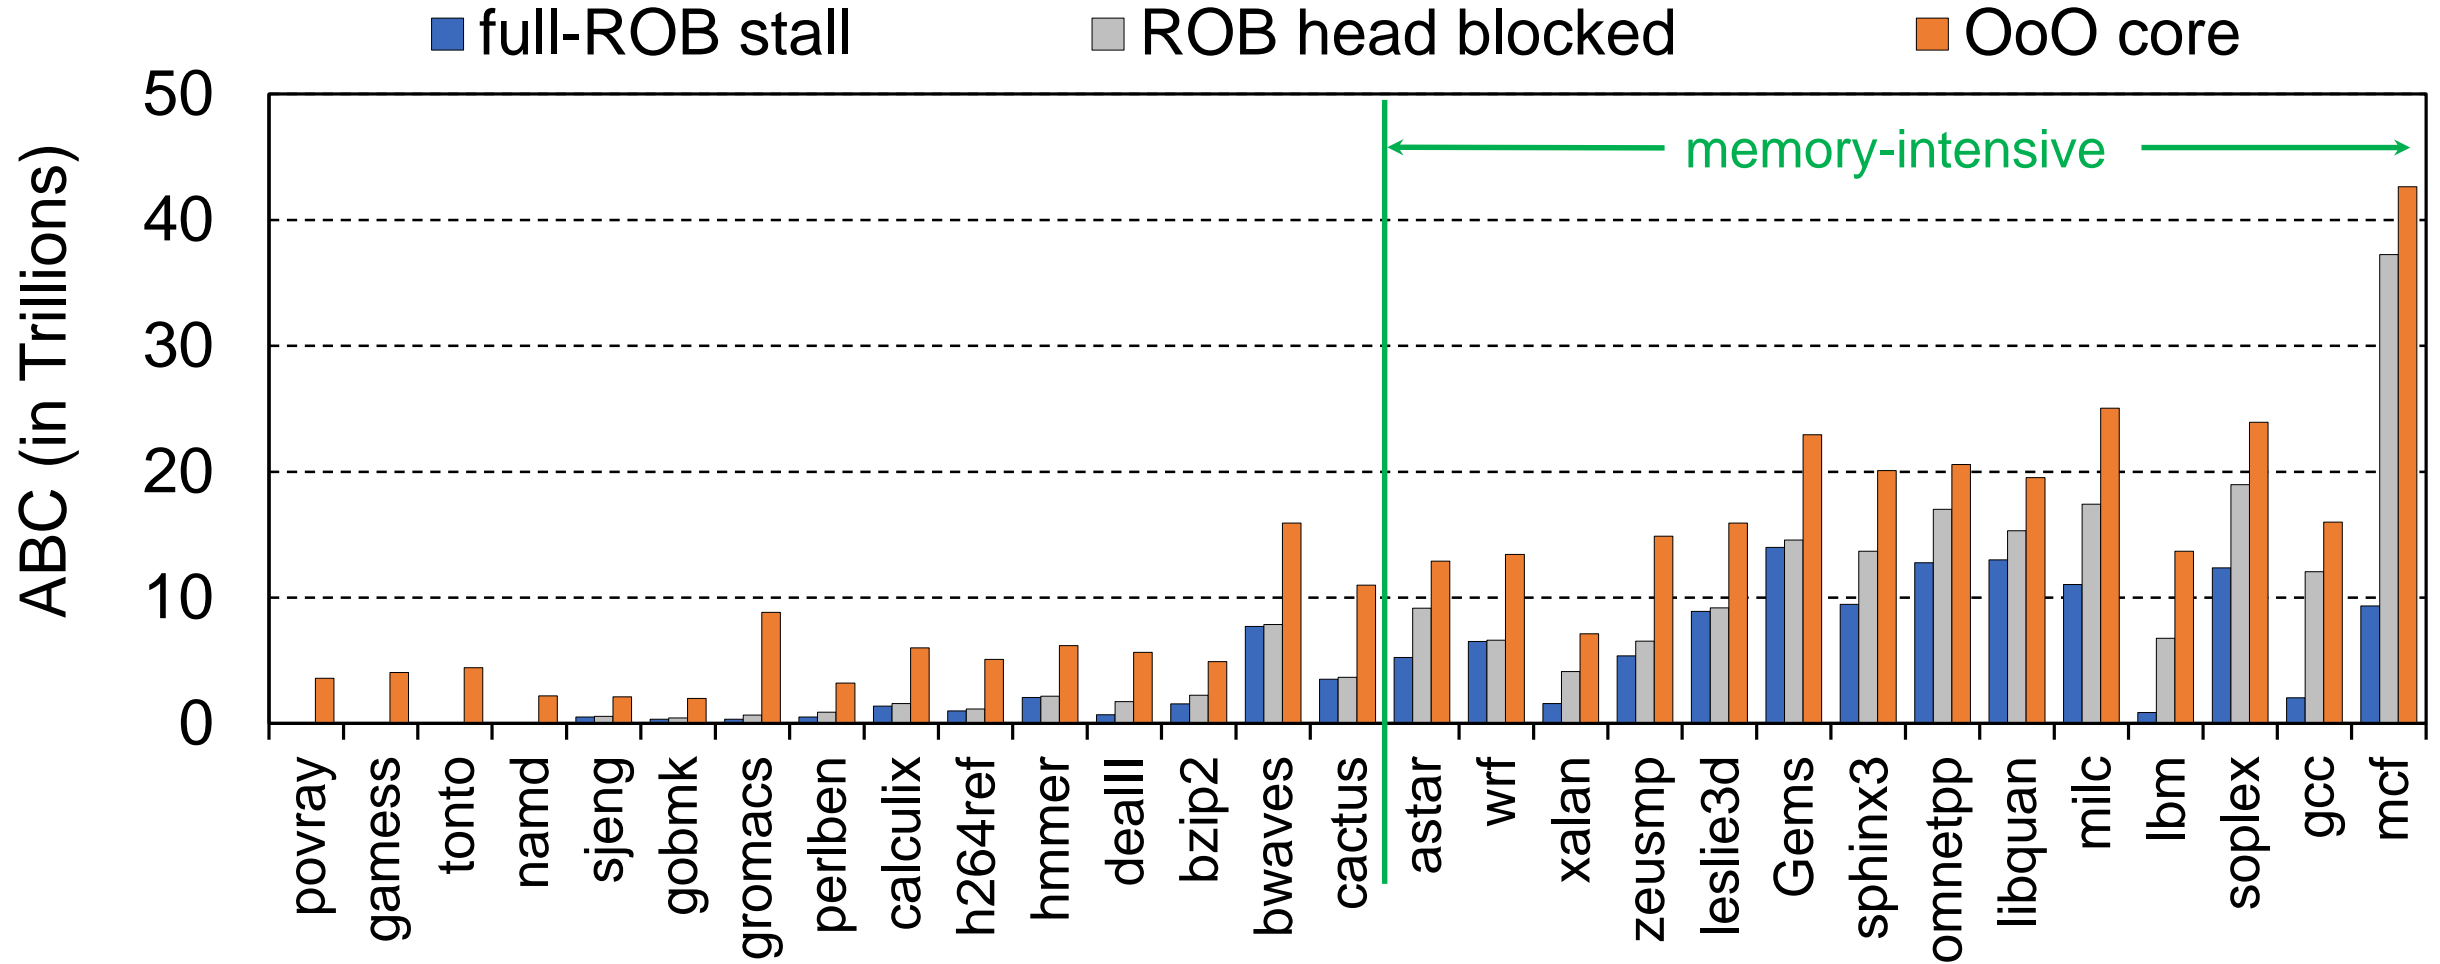
Vulnerable State under Blocked ROB Head



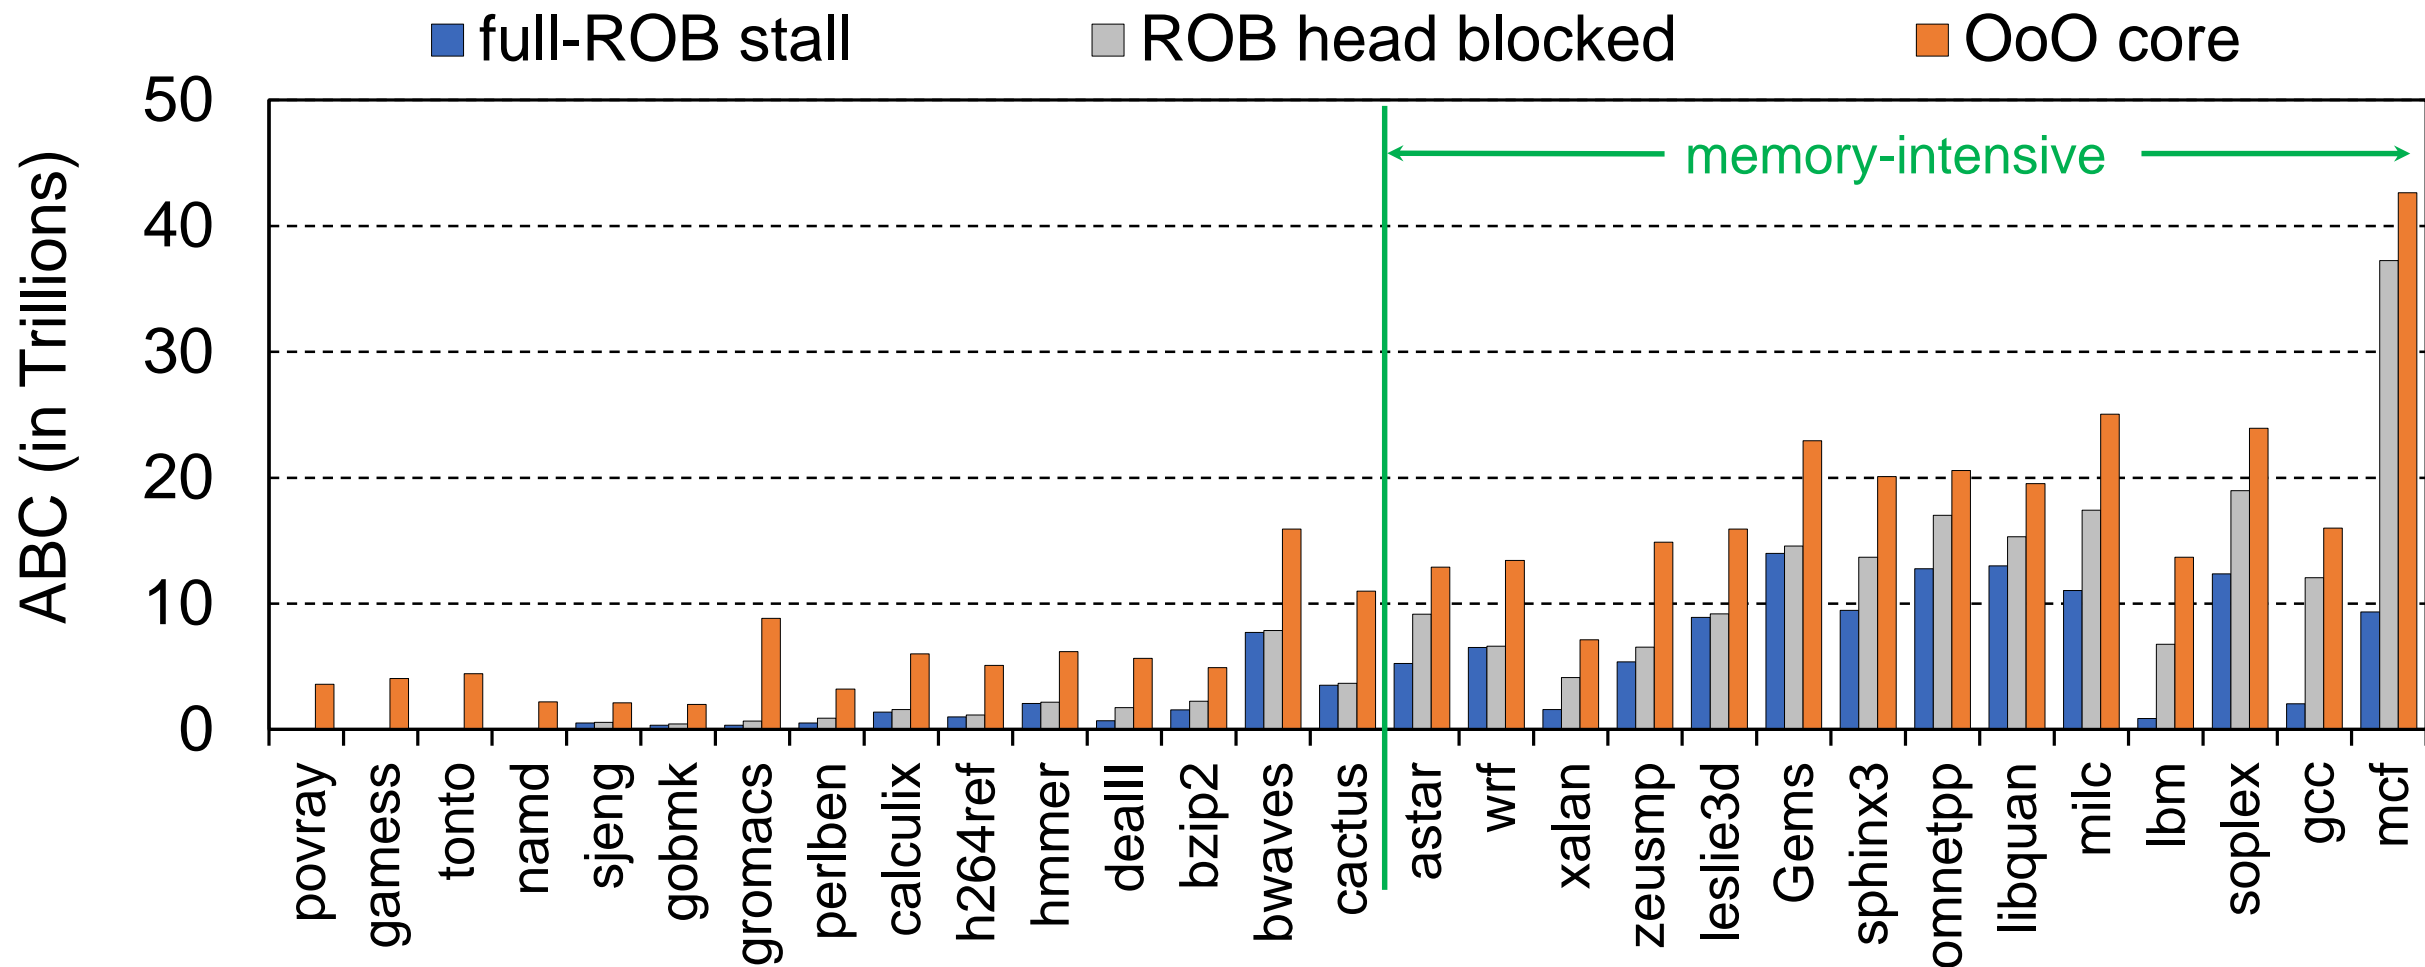
Vulnerable State under Blocked ROB Head



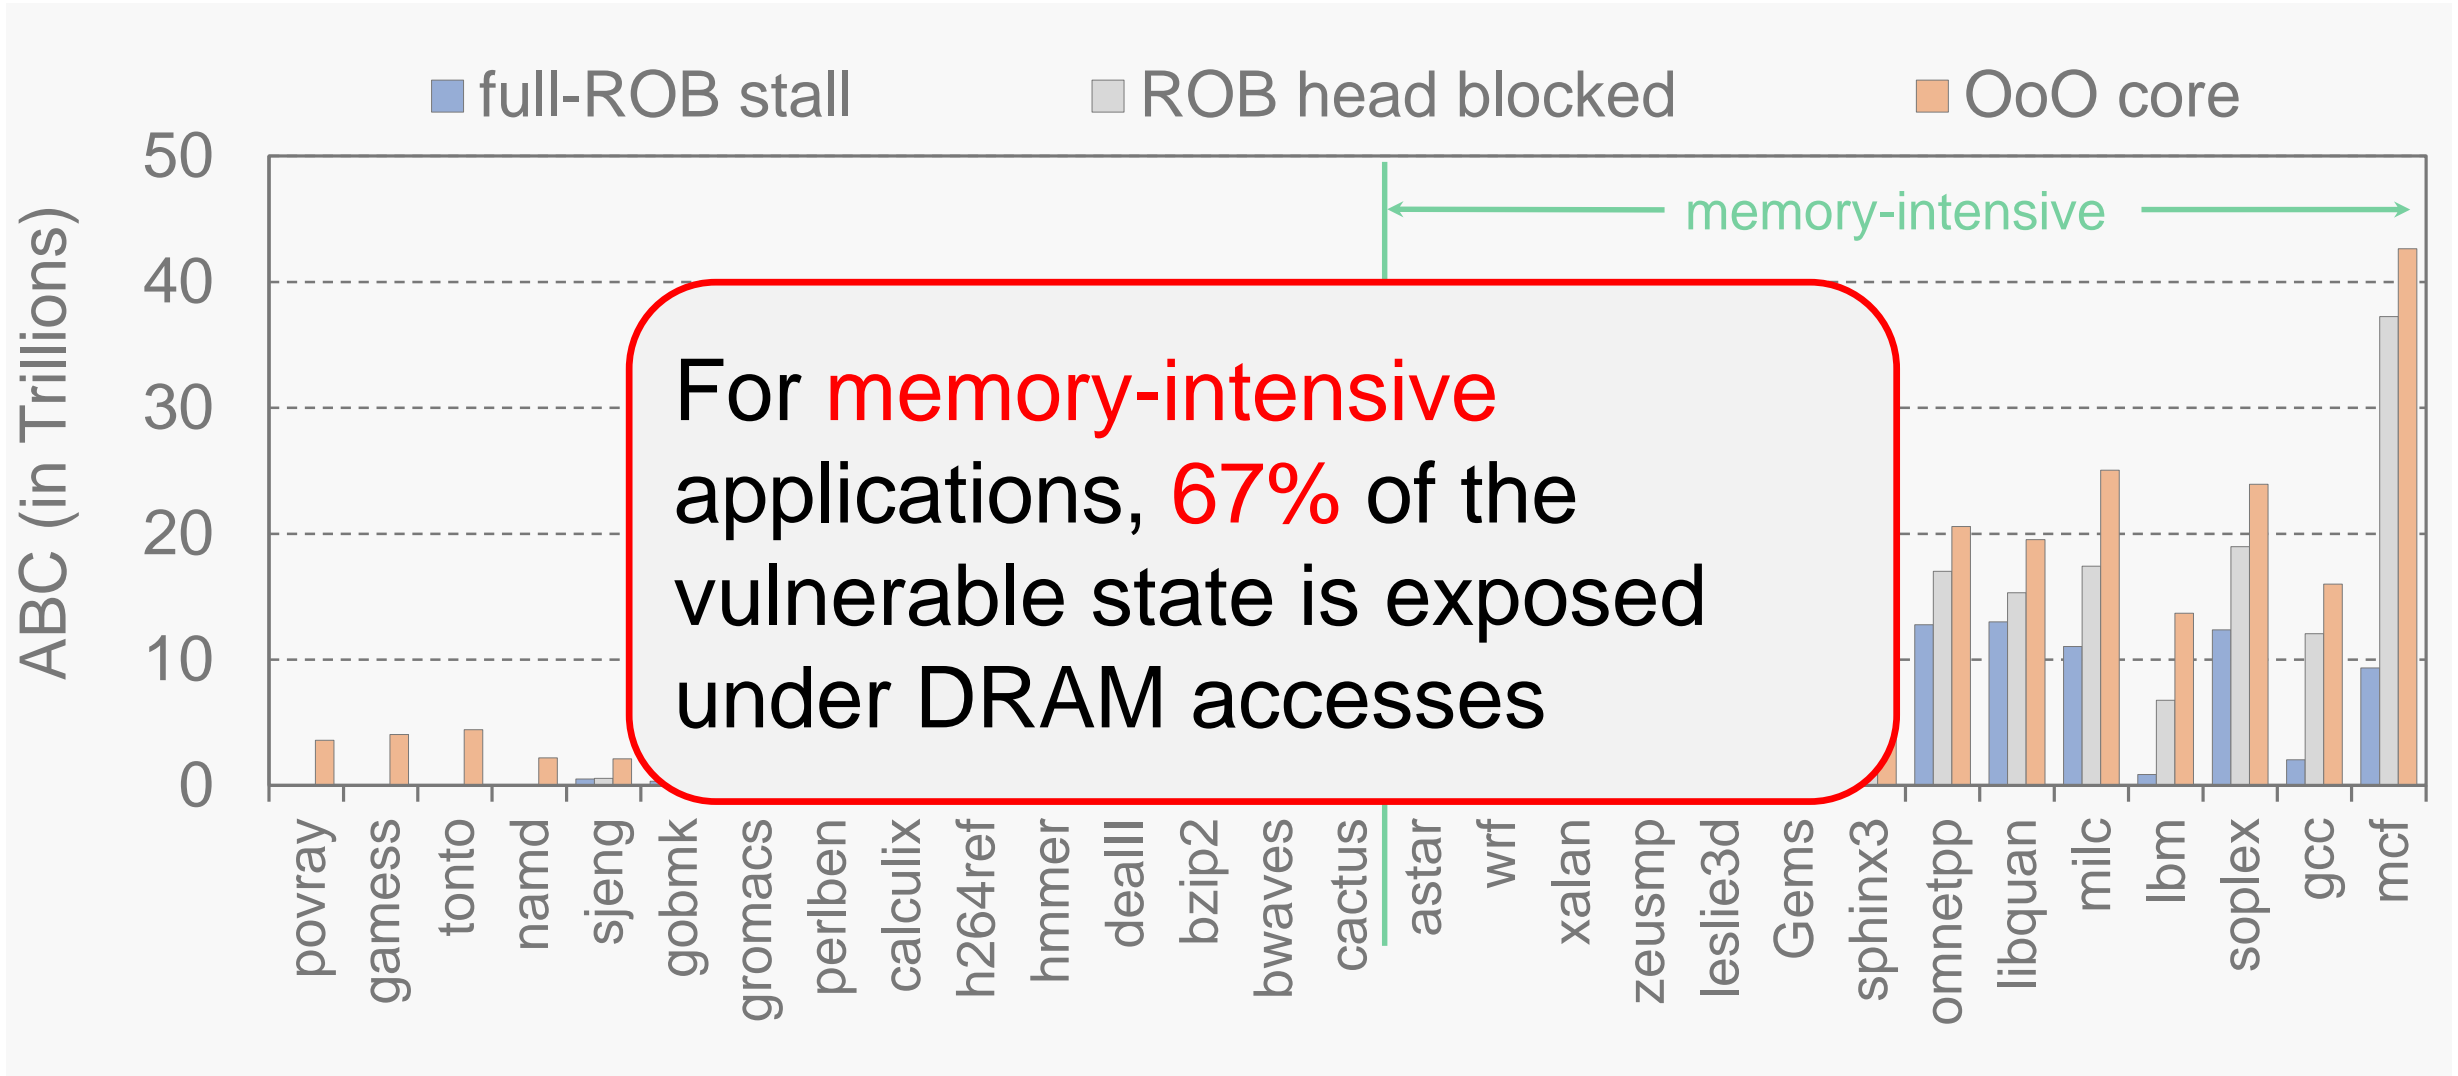
Vulnerable State under Blocked ROB Head



Quantifying the Potential



Quantifying the Potential



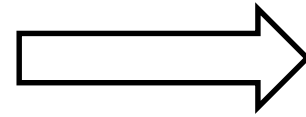
Dispatch Halting to Improve Reliability

Dispatch Halting to Improve Reliability

Proactively prevent instructions following an LLC missing load to allocate entries in the ROB (or back-end structures)

Dispatch Halting to Improve Reliability

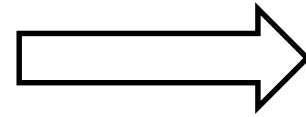
Proactively prevent instructions following an LLC missing load to allocate entries in the ROB (or back-end structures)



Proactive Dispatch Halting

Dispatch Halting to Improve Reliability

Proactively prevent instructions following an LLC missing load to allocate entries in the ROB (or back-end structures)

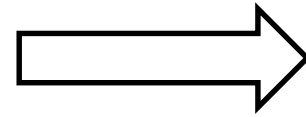


Proactive Dispatch Halting

Reactively mark the back-end speculative after a blocked ROB head

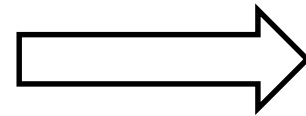
Dispatch Halting to Improve Reliability

Proactively prevent instructions following an LLC missing load to allocate entries in the ROB (or back-end structures)



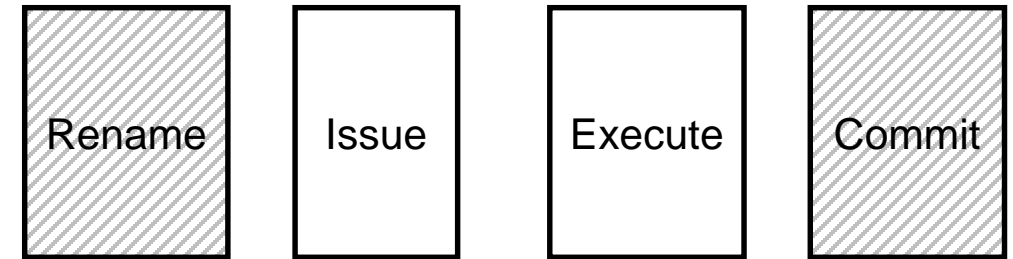
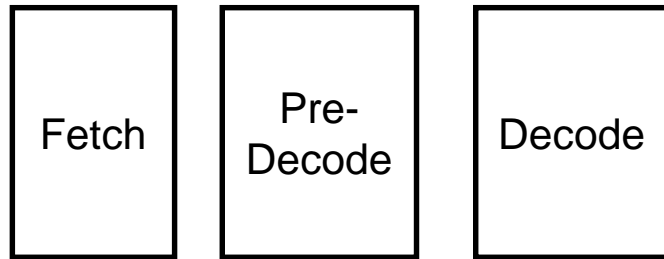
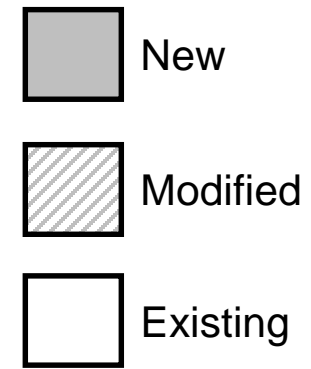
Proactive Dispatch Halting

Reactively mark the back-end speculative after a blocked ROB head

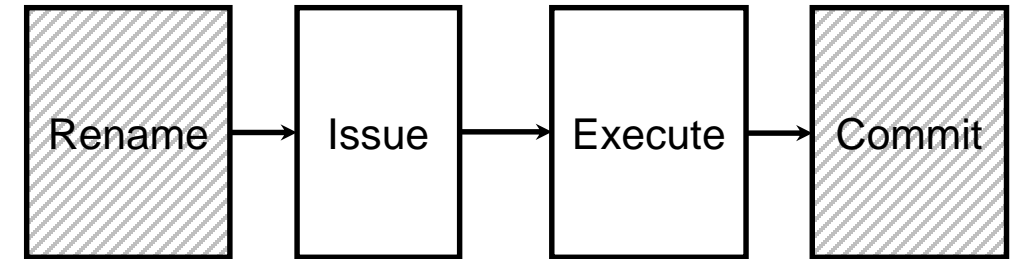
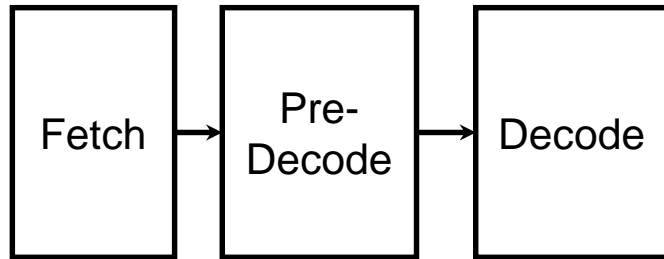
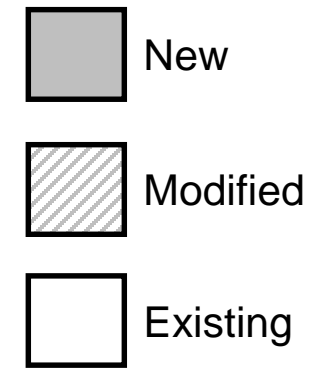


Reactive Dispatch Halting

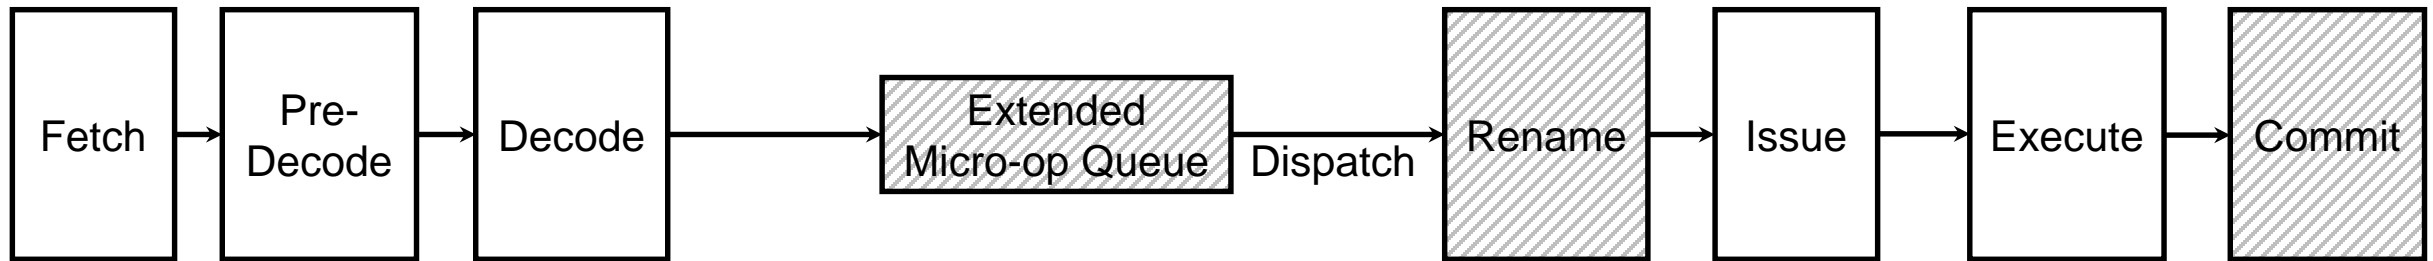
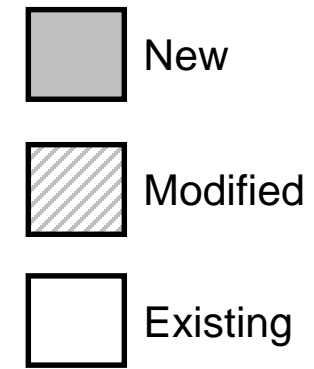
Proactive Dispatch Halting



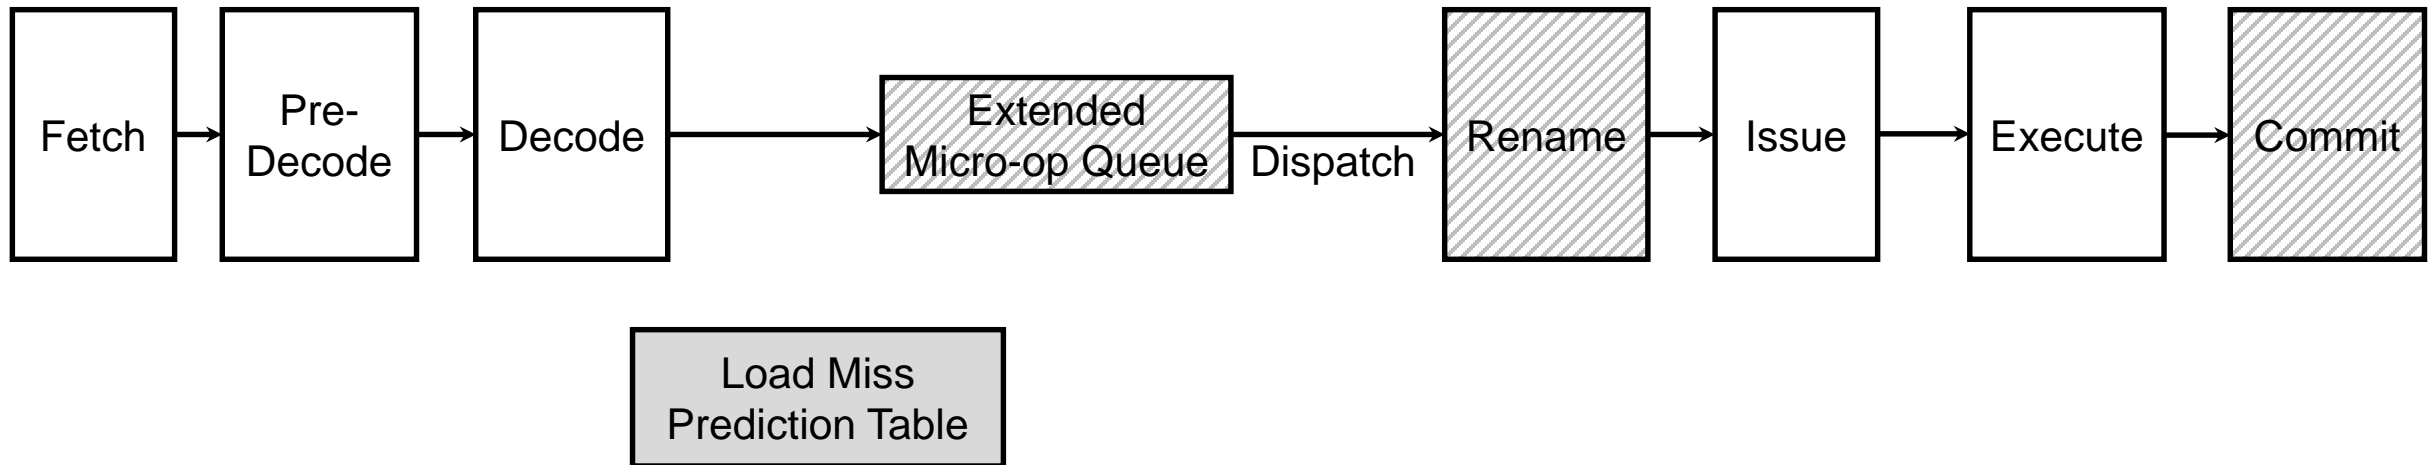
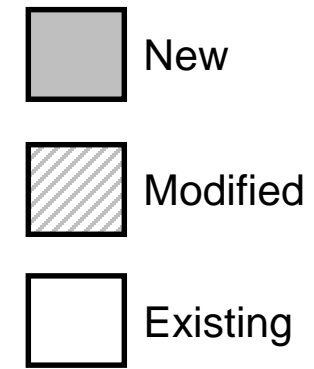
Proactive Dispatch Halting



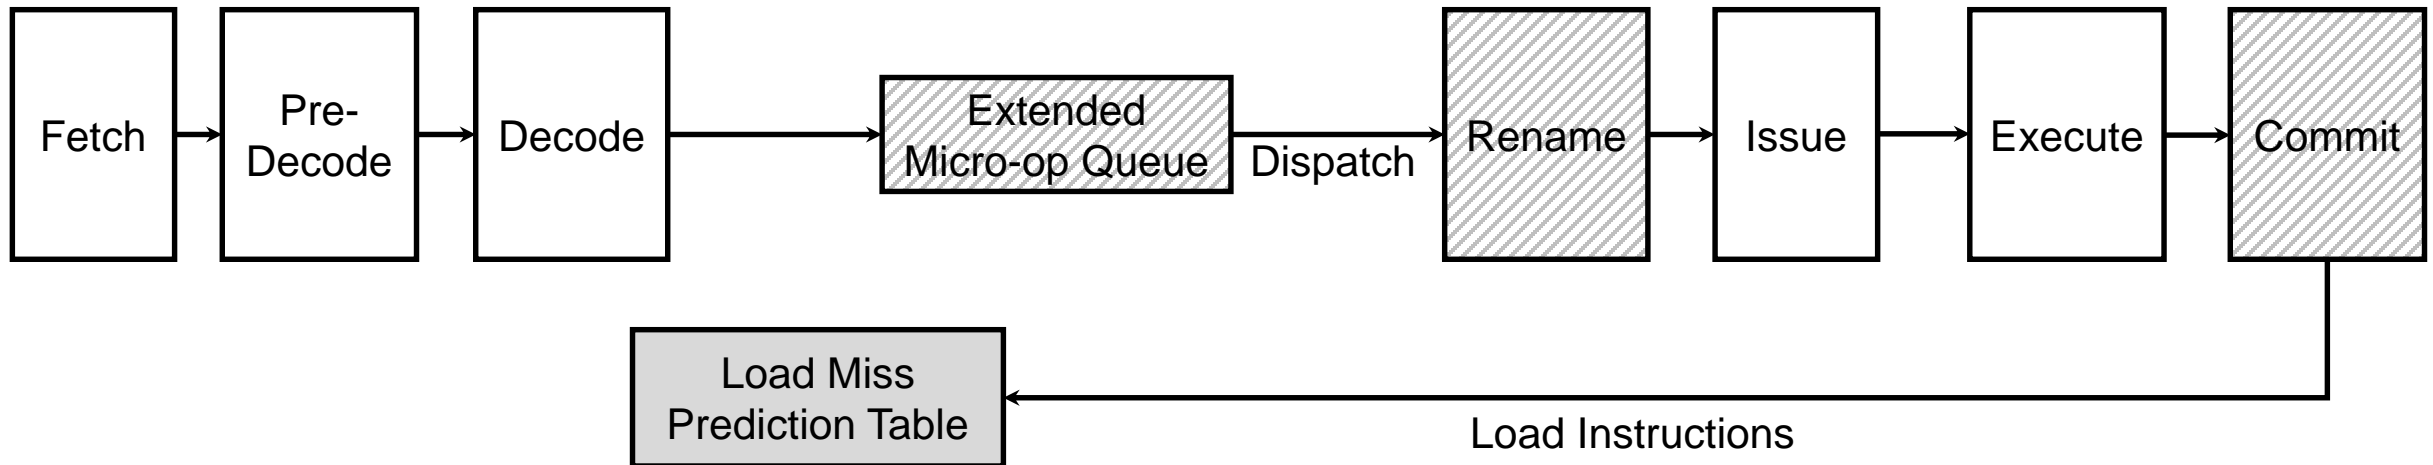
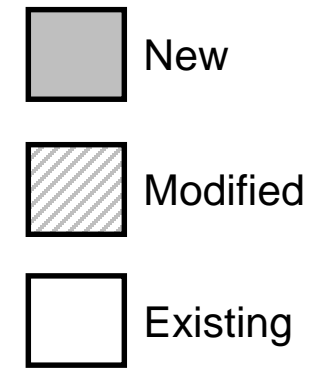
Proactive Dispatch Halting



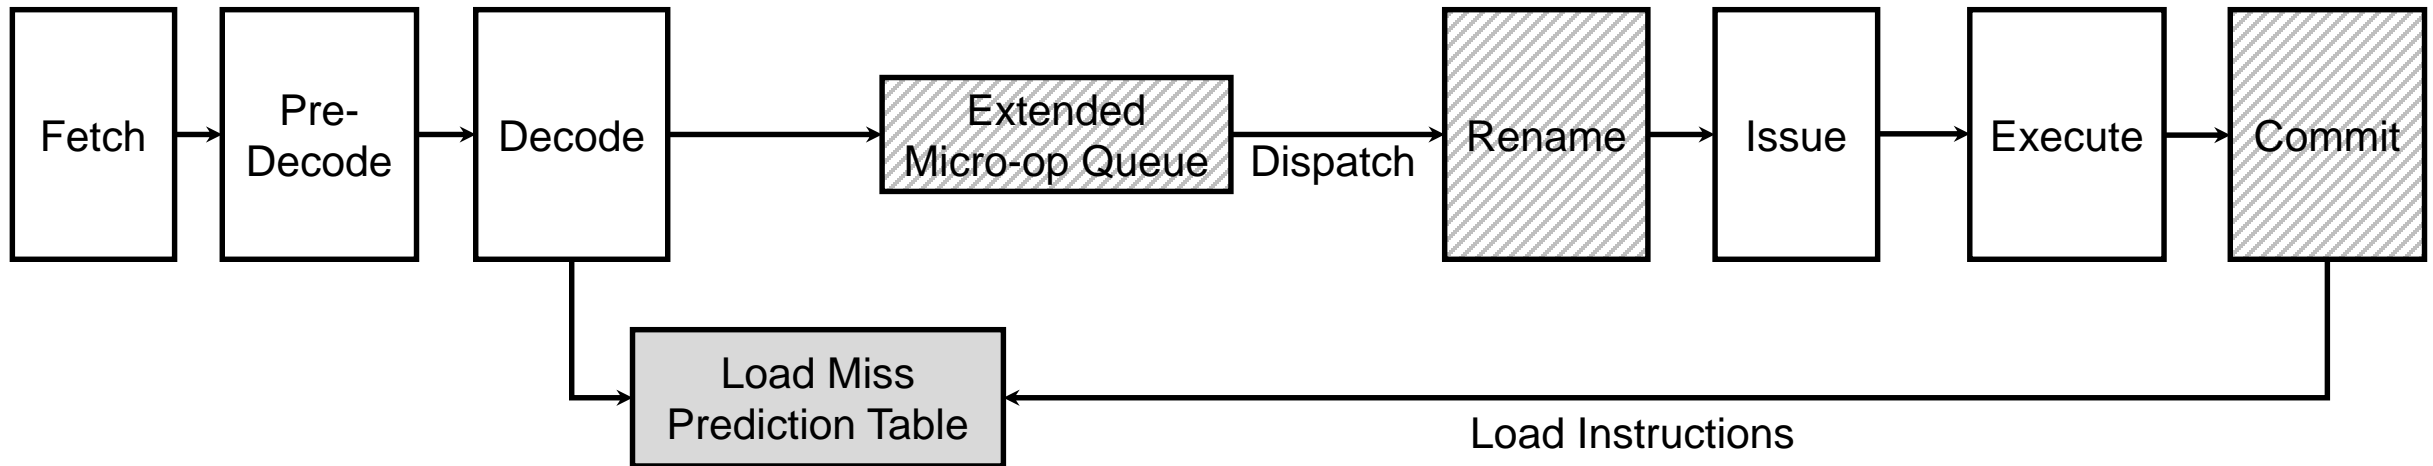
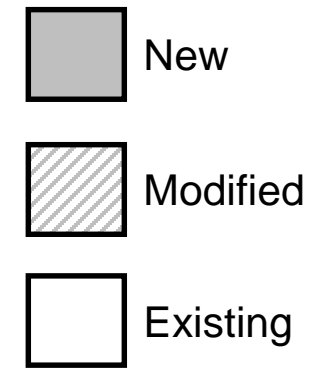
Proactive Dispatch Halting



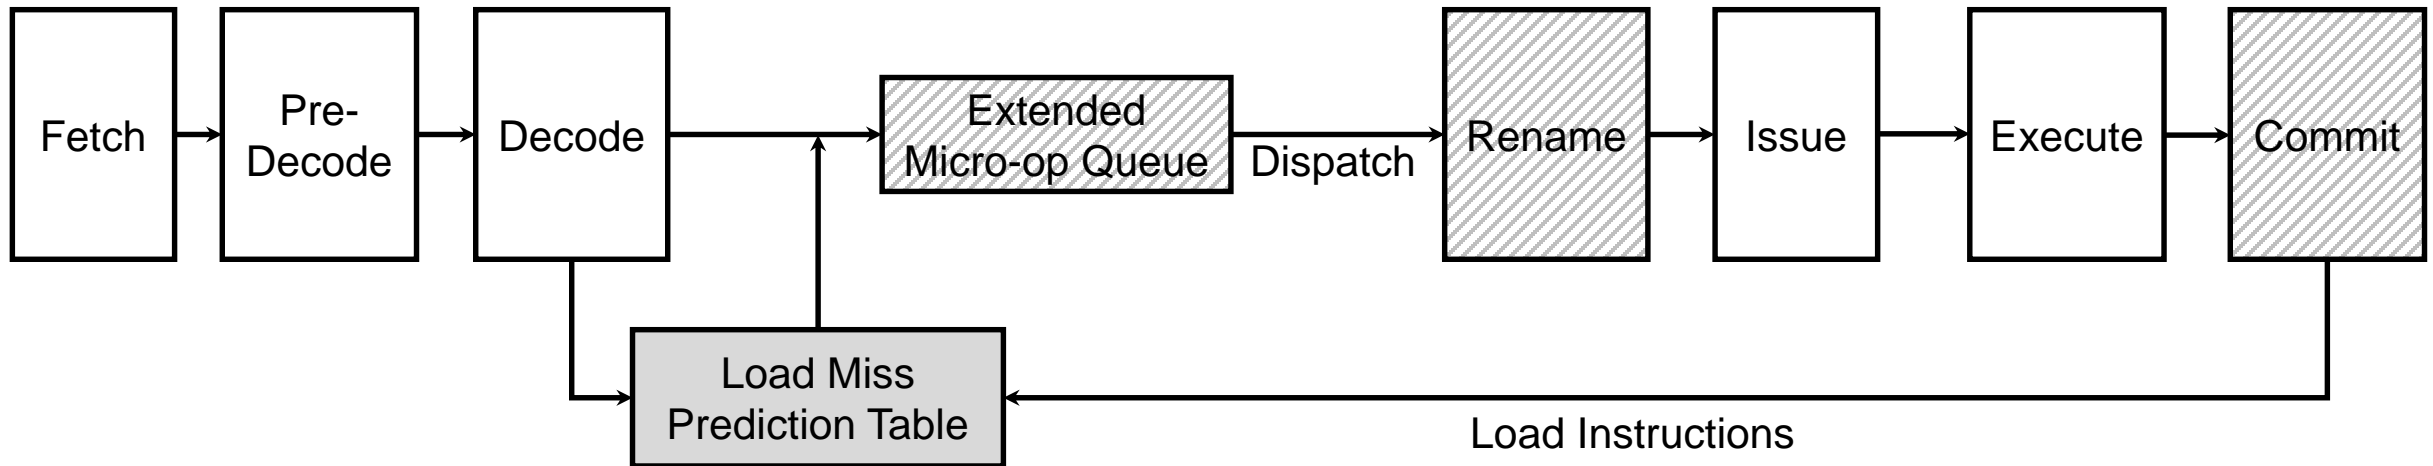
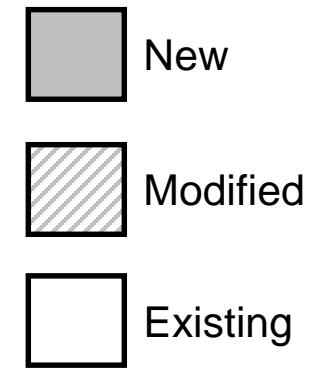
Proactive Dispatch Halting



Proactive Dispatch Halting

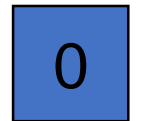


Proactive Dispatch Halting



Performance Bottlenecks

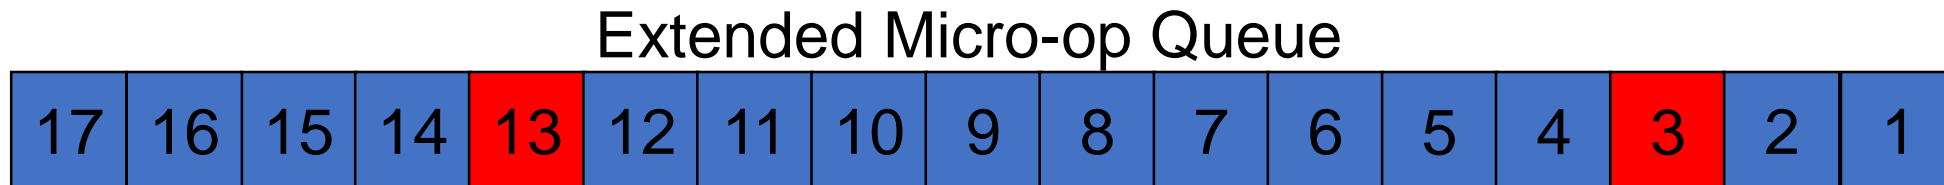
Extended Micro-op Queue




LLC missing
load
dispatched

Performance Bottlenecks

 → Load instructions




↓
LLC missing
load
dispatched

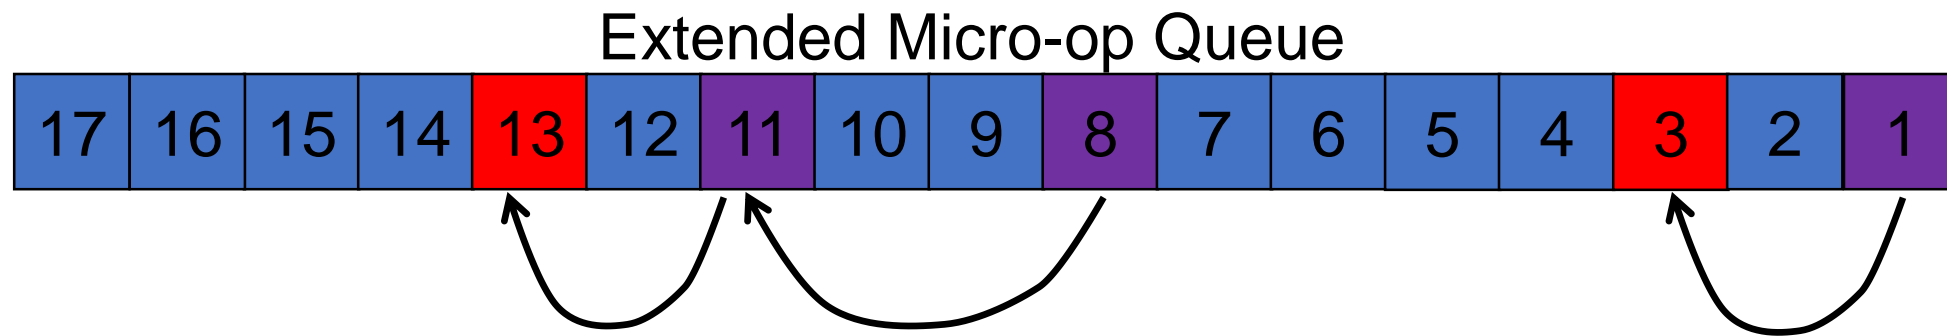
Performance Bottlenecks



→ Load instructions



→ Address Generating Instructions or Producer Instructions



LLC missing
load
dispatched

Performance Bottlenecks



→ Load instructions

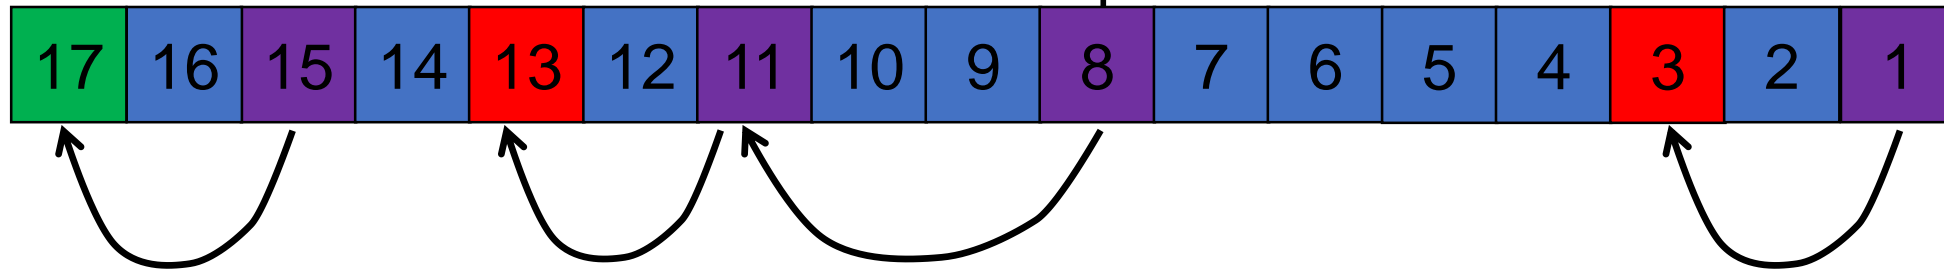


→ Address Generating Instructions or Producer Instructions



→ (Mispredicted) Branch Instructions

Extended Micro-op Queue



↓
LLC missing
load
dispatched

Handling Performance Bottlenecks

Handling Performance Bottlenecks

- Execute load/branch instructions and their producer instructions **speculatively**

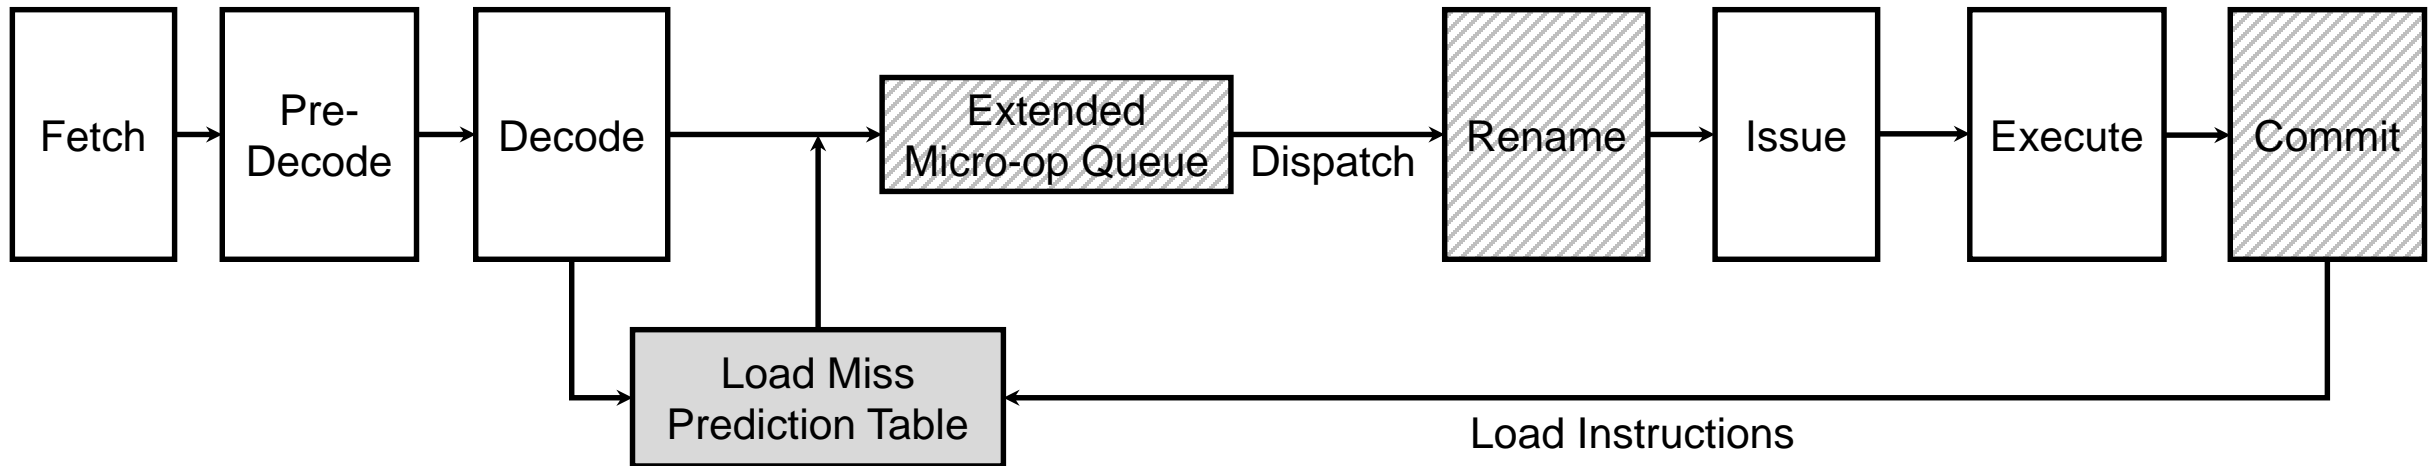
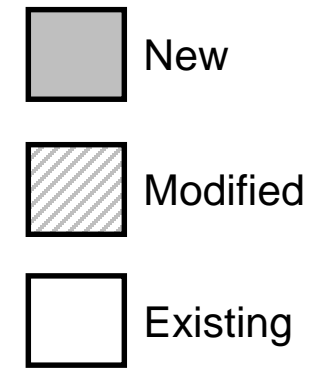
Handling Performance Bottlenecks

- Execute load/branch instructions and their producer instructions **speculatively**
- Find producer instructions **iteratively**
 - Modify front-end Register Allocation Table (RAT)

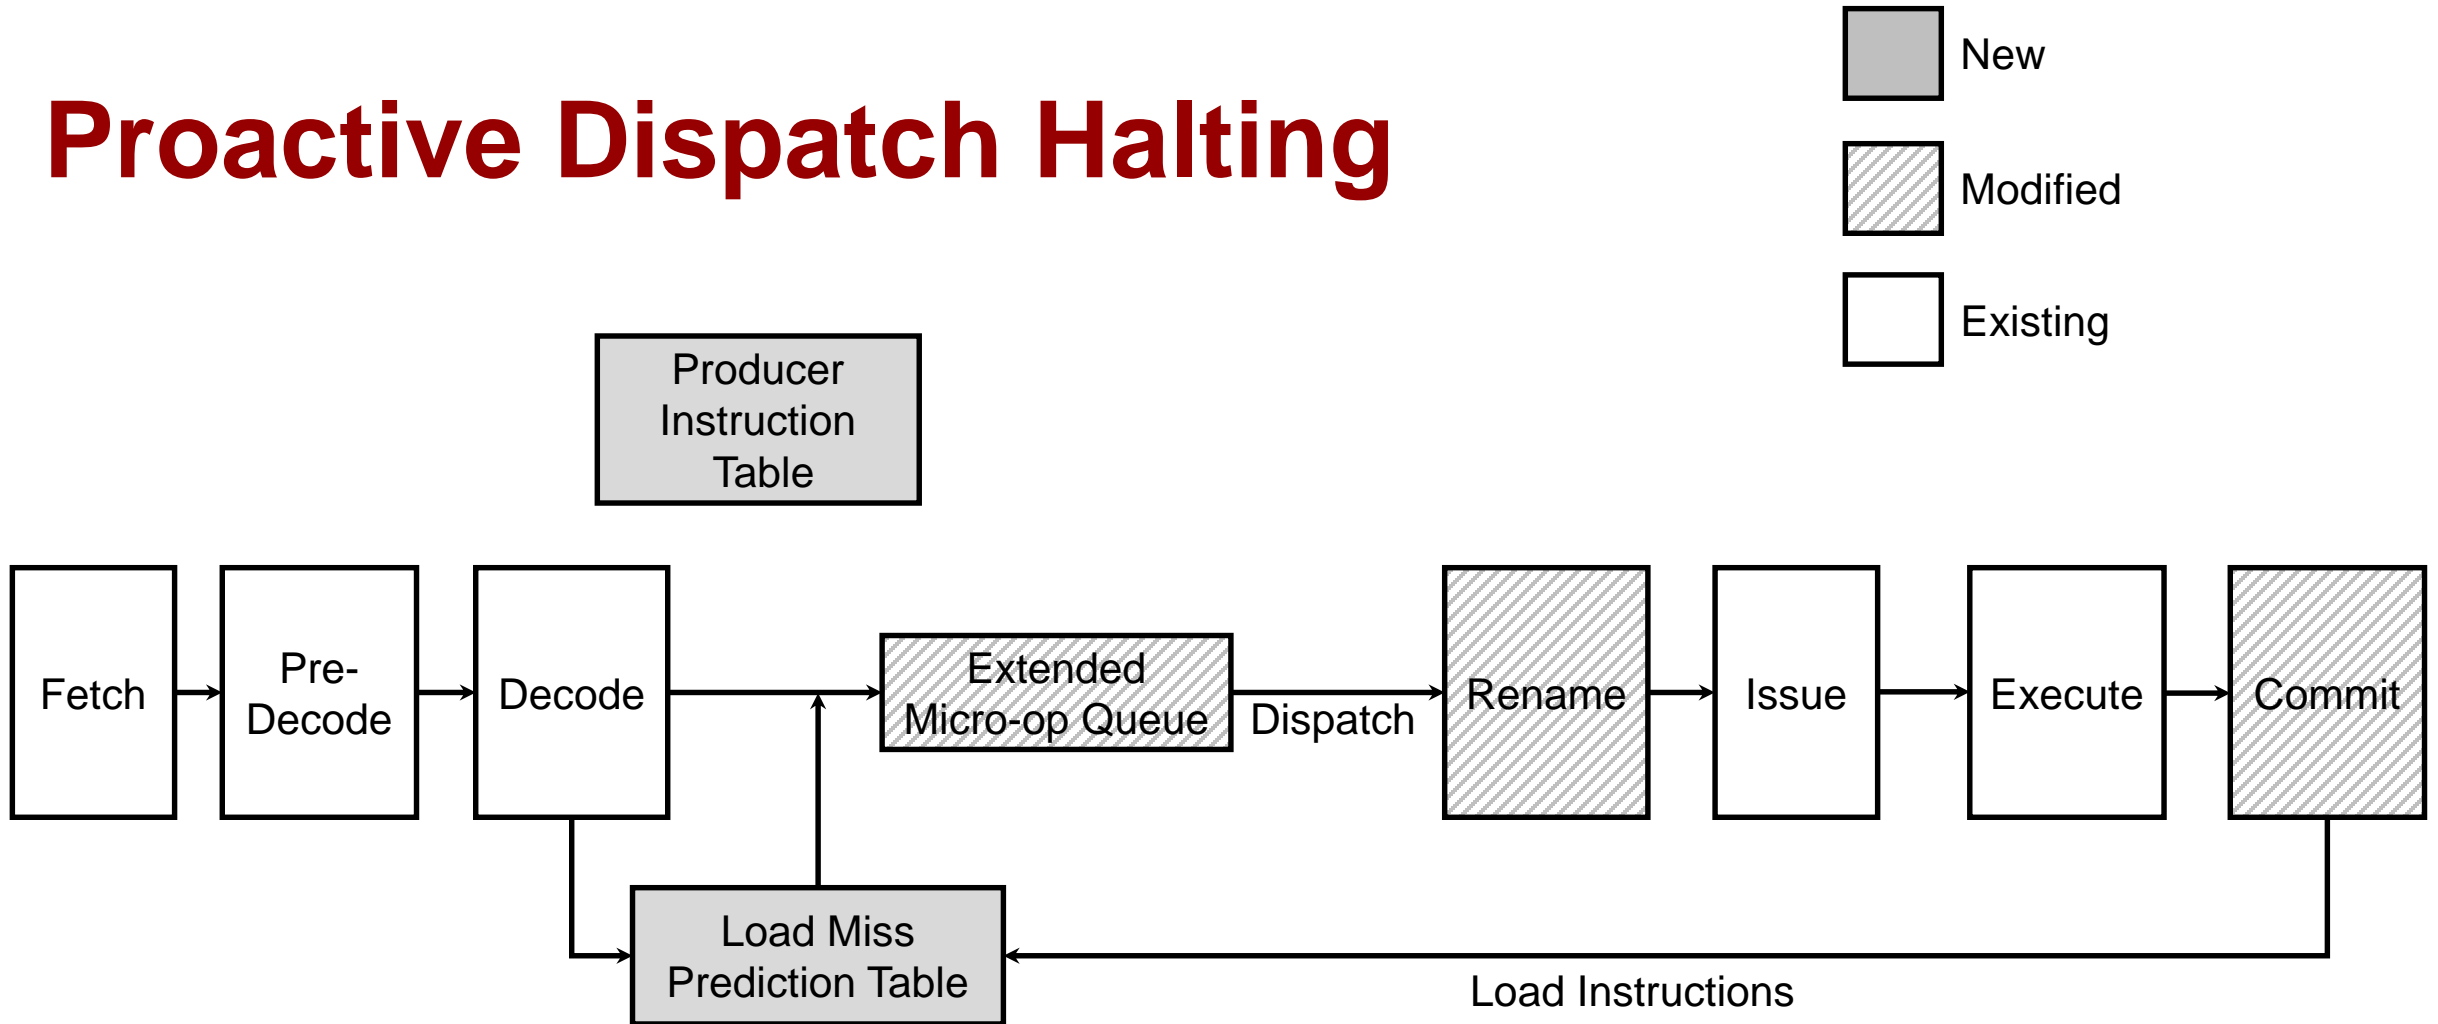
Handling Performance Bottlenecks

- Execute load/branch instructions and their producer instructions **speculatively**
- Find producer instructions **iteratively**
 - Modify front-end Register Allocation Table (RAT)
- One additional **RAT checkpoint**

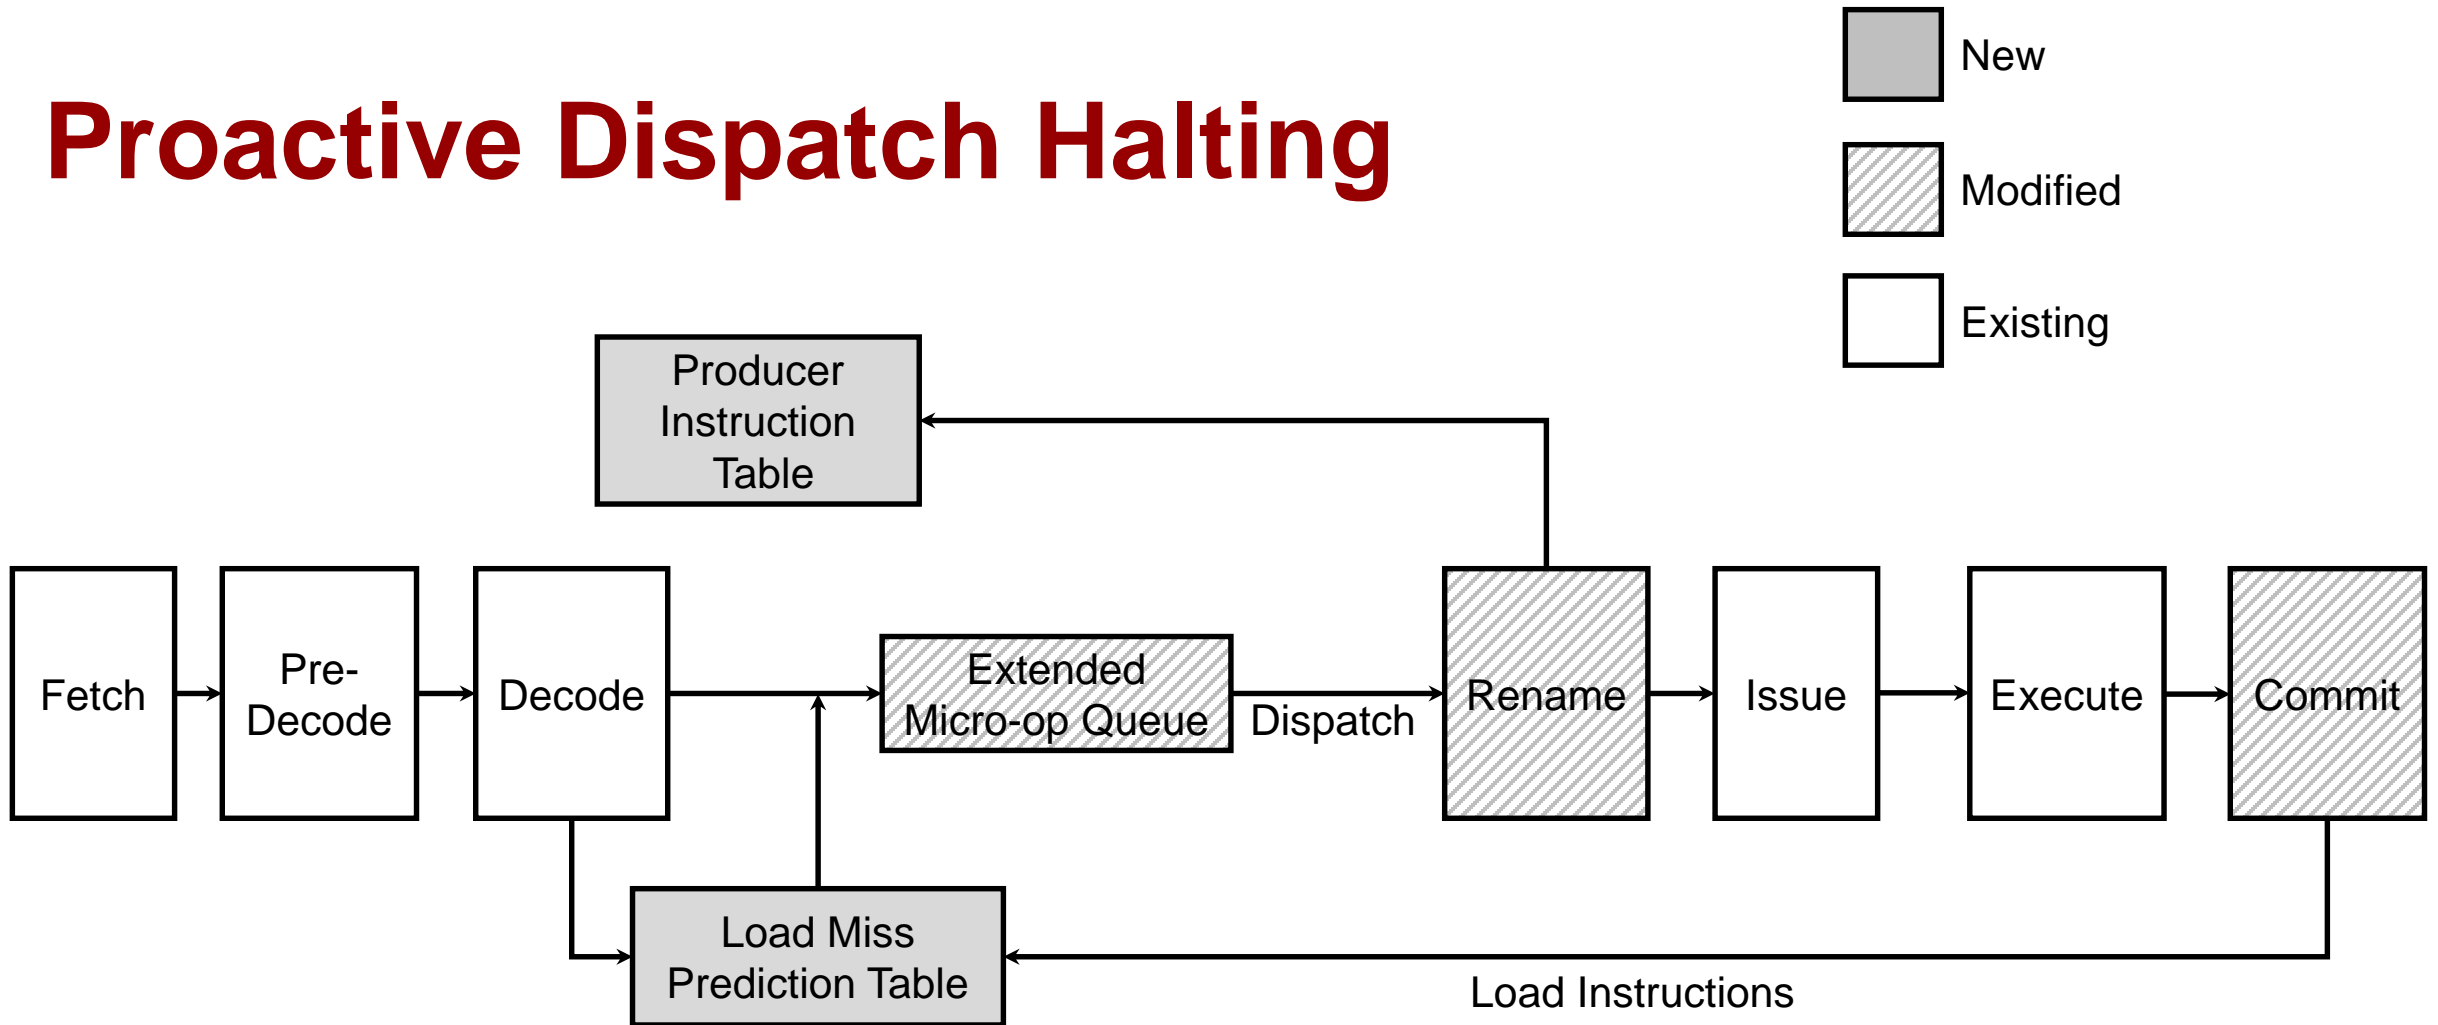
Proactive Dispatch Halting



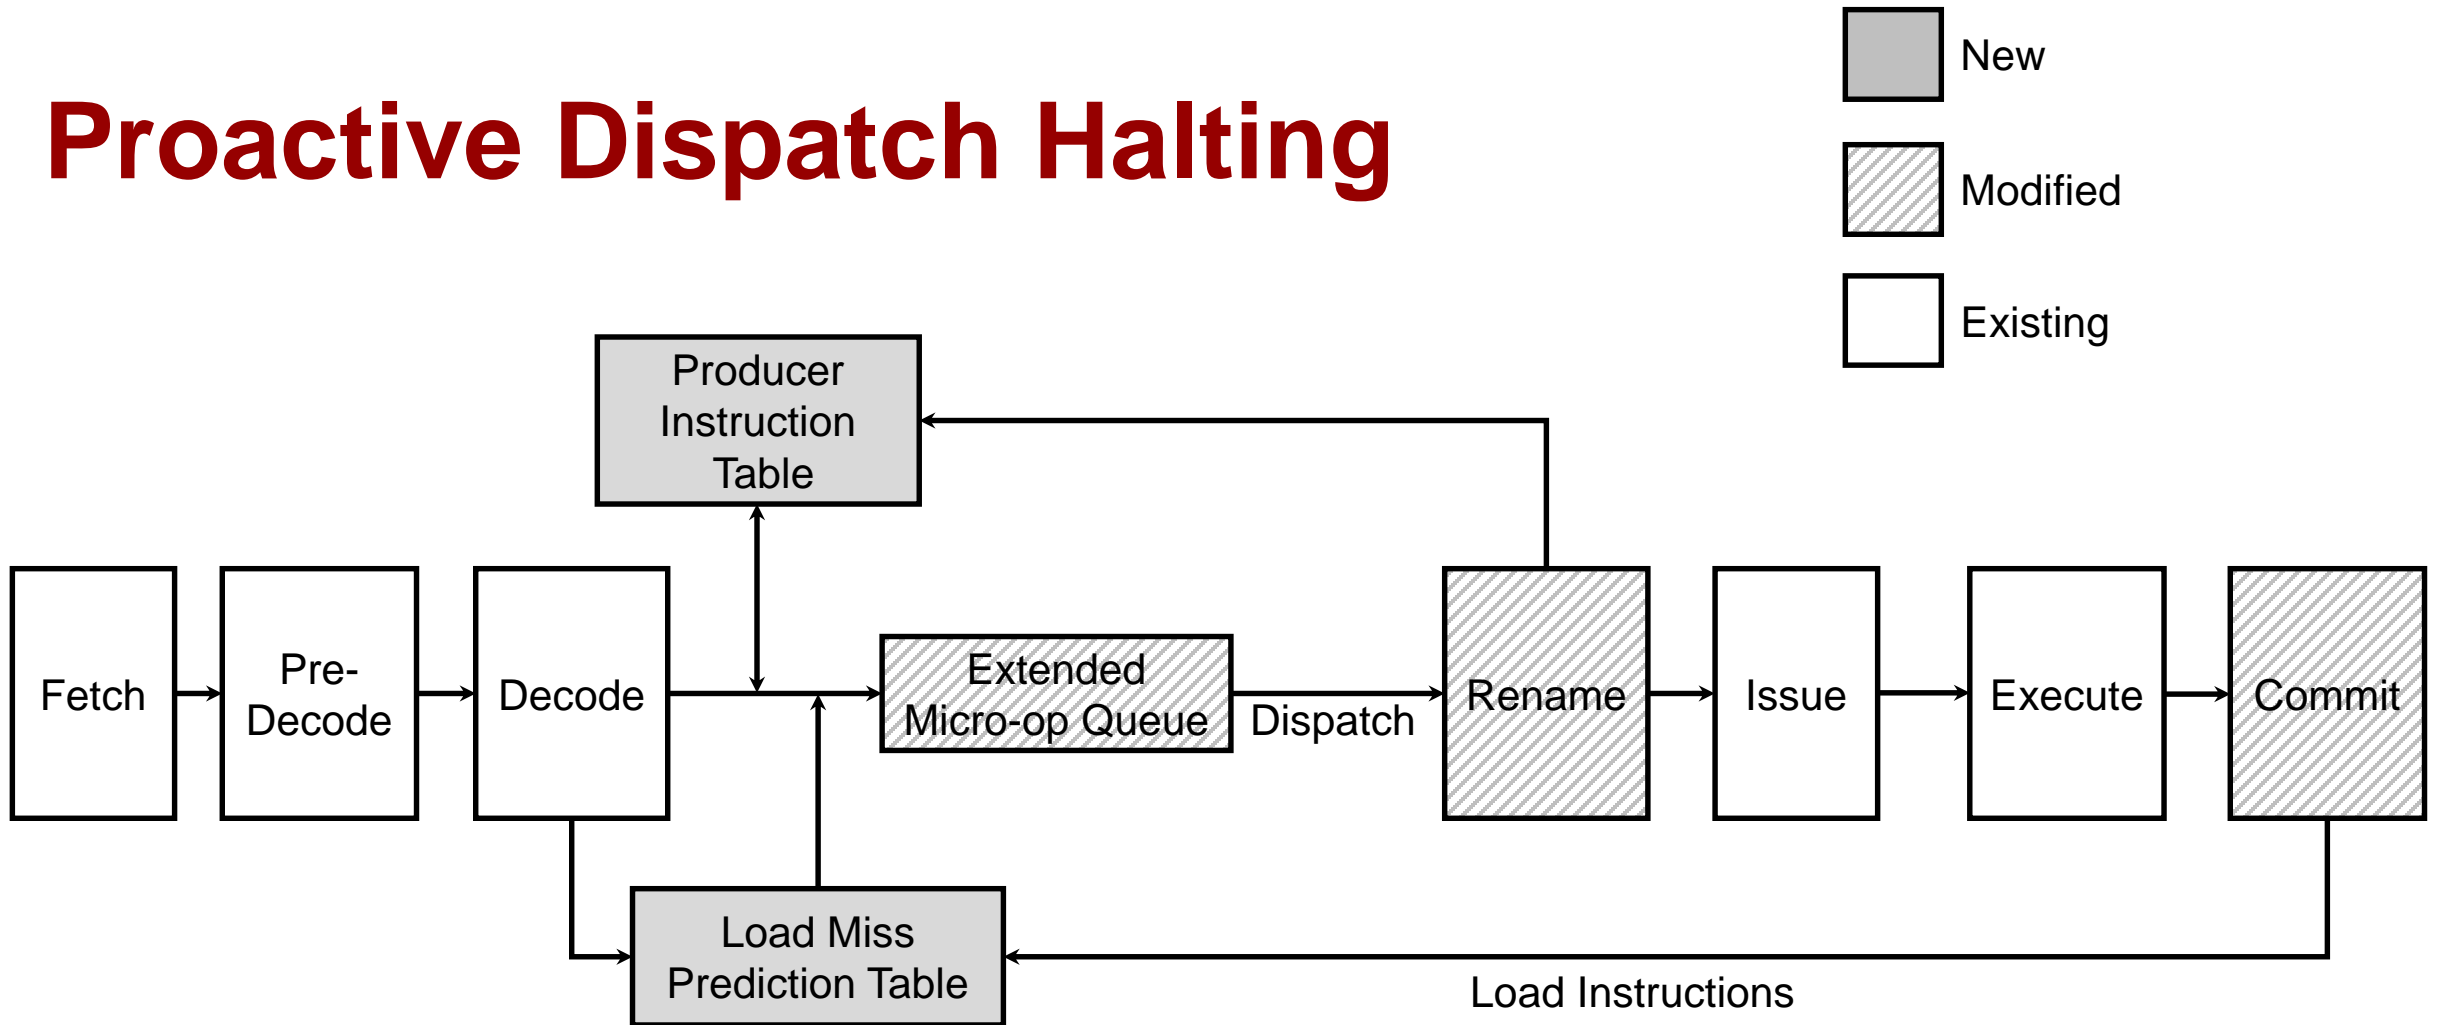
Proactive Dispatch Halting



Proactive Dispatch Halting



Proactive Dispatch Halting



Reactive Dispatch Halting

Reactive Dispatch Halting

- Mark back-end as **speculative** after a load miss

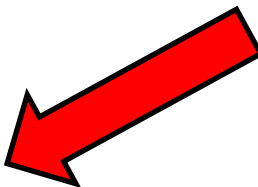
Reactive Dispatch Halting

- Mark back-end as **speculative** after a load miss
- EMQ works as a **replay buffer**

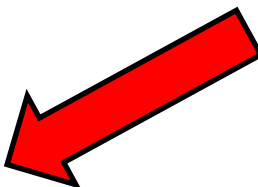
Reactive Dispatch Halting

- Mark back-end as **speculative** after a load miss
- EMQ works as a **replay buffer**
- Upon resume

Reactive Dispatch Halting

- Mark back-end as **speculative** after a load miss
 - EMQ works as a **replay buffer**
 - Upon resume
 - Copy retirement RAT to front-end RAT
- Checkpointing RAT
does not work
- 

Reactive Dispatch Halting

- Mark back-end as **speculative** after a load miss
 - EMQ works as a **replay buffer**
 - Upon resume
 - Copy retirement RAT to front-end RAT
 - **Re-execute all instructions** from the replay buffer or EMQ
- Checkpointing RAT
does not work
- 

Evaluation

OoO: Baseline out-of-order core

Evaluation

OoO: Baseline out-of-order core

NO-SPEC: Proactive dispatch halting without speculative execution

Evaluation

OoO: Baseline out-of-order core

NO-SPEC: Proactive dispatch halting without speculative execution

P-DH: Proactive dispatch halting

Evaluation

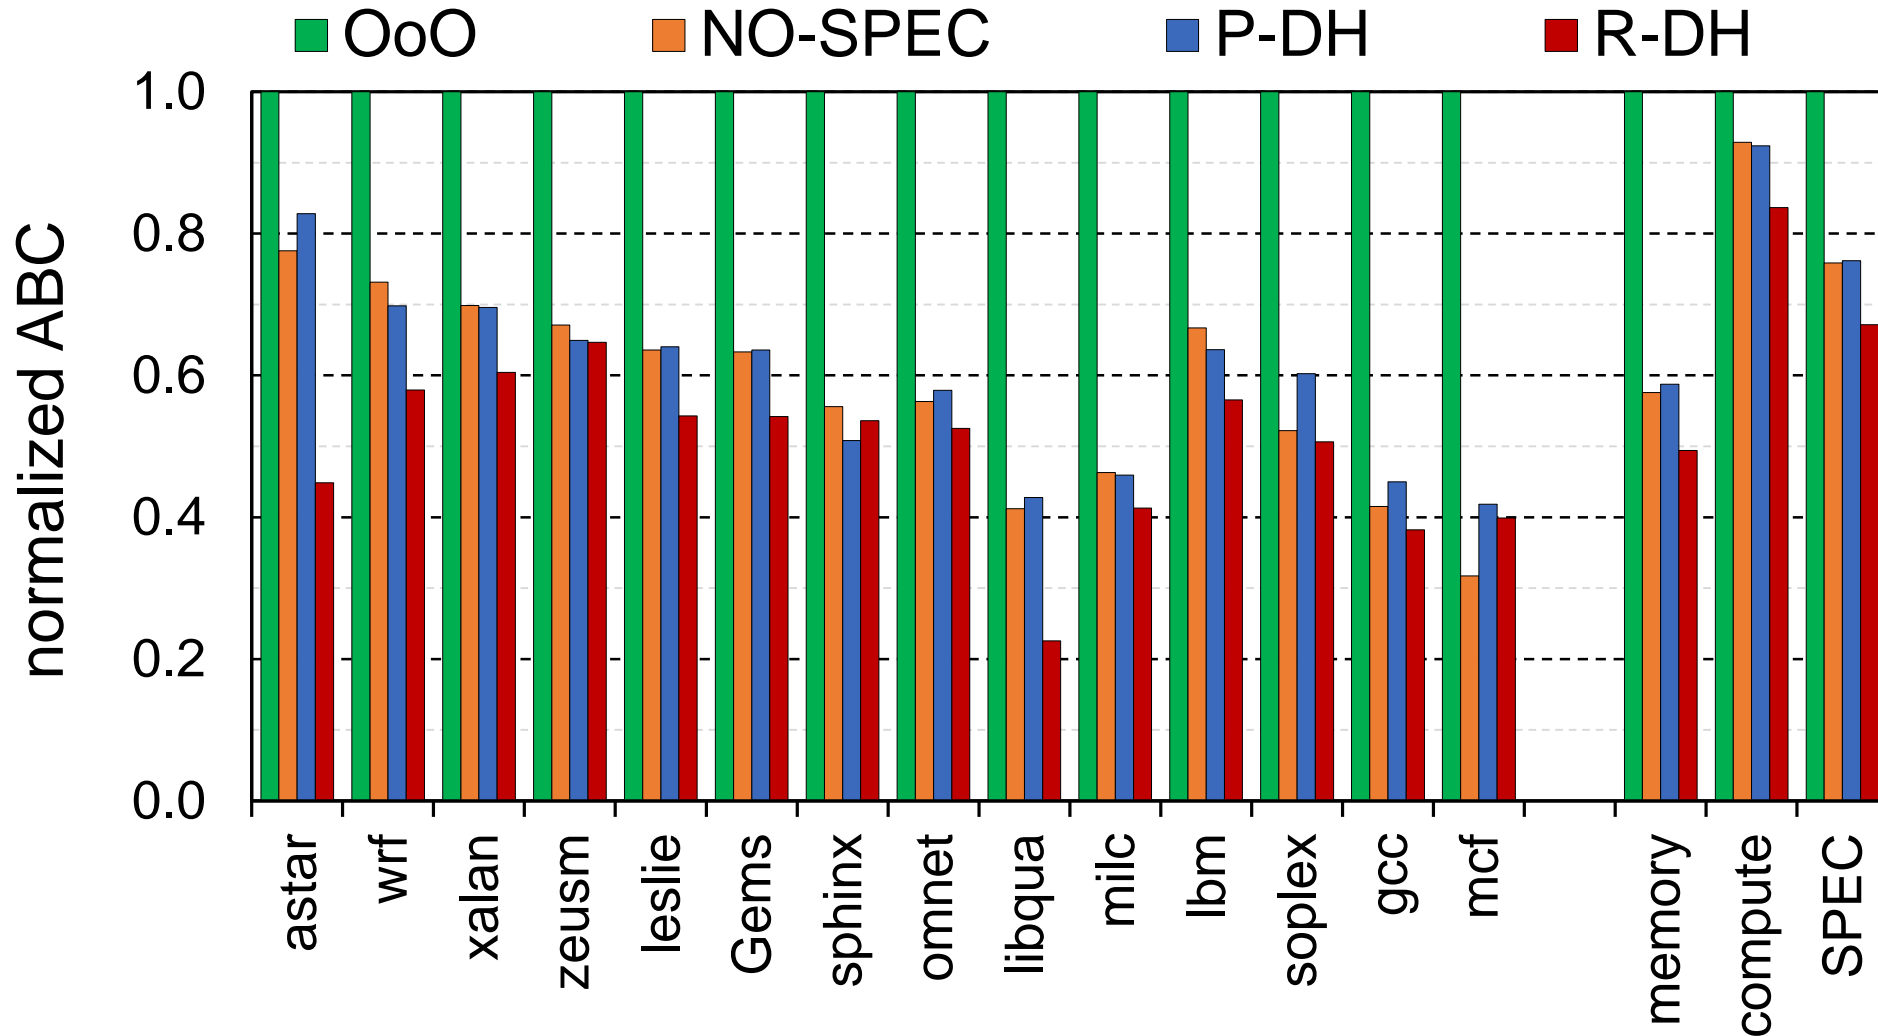
OoO: Baseline out-of-order core

NO-SPEC: Proactive dispatch halting without speculative execution

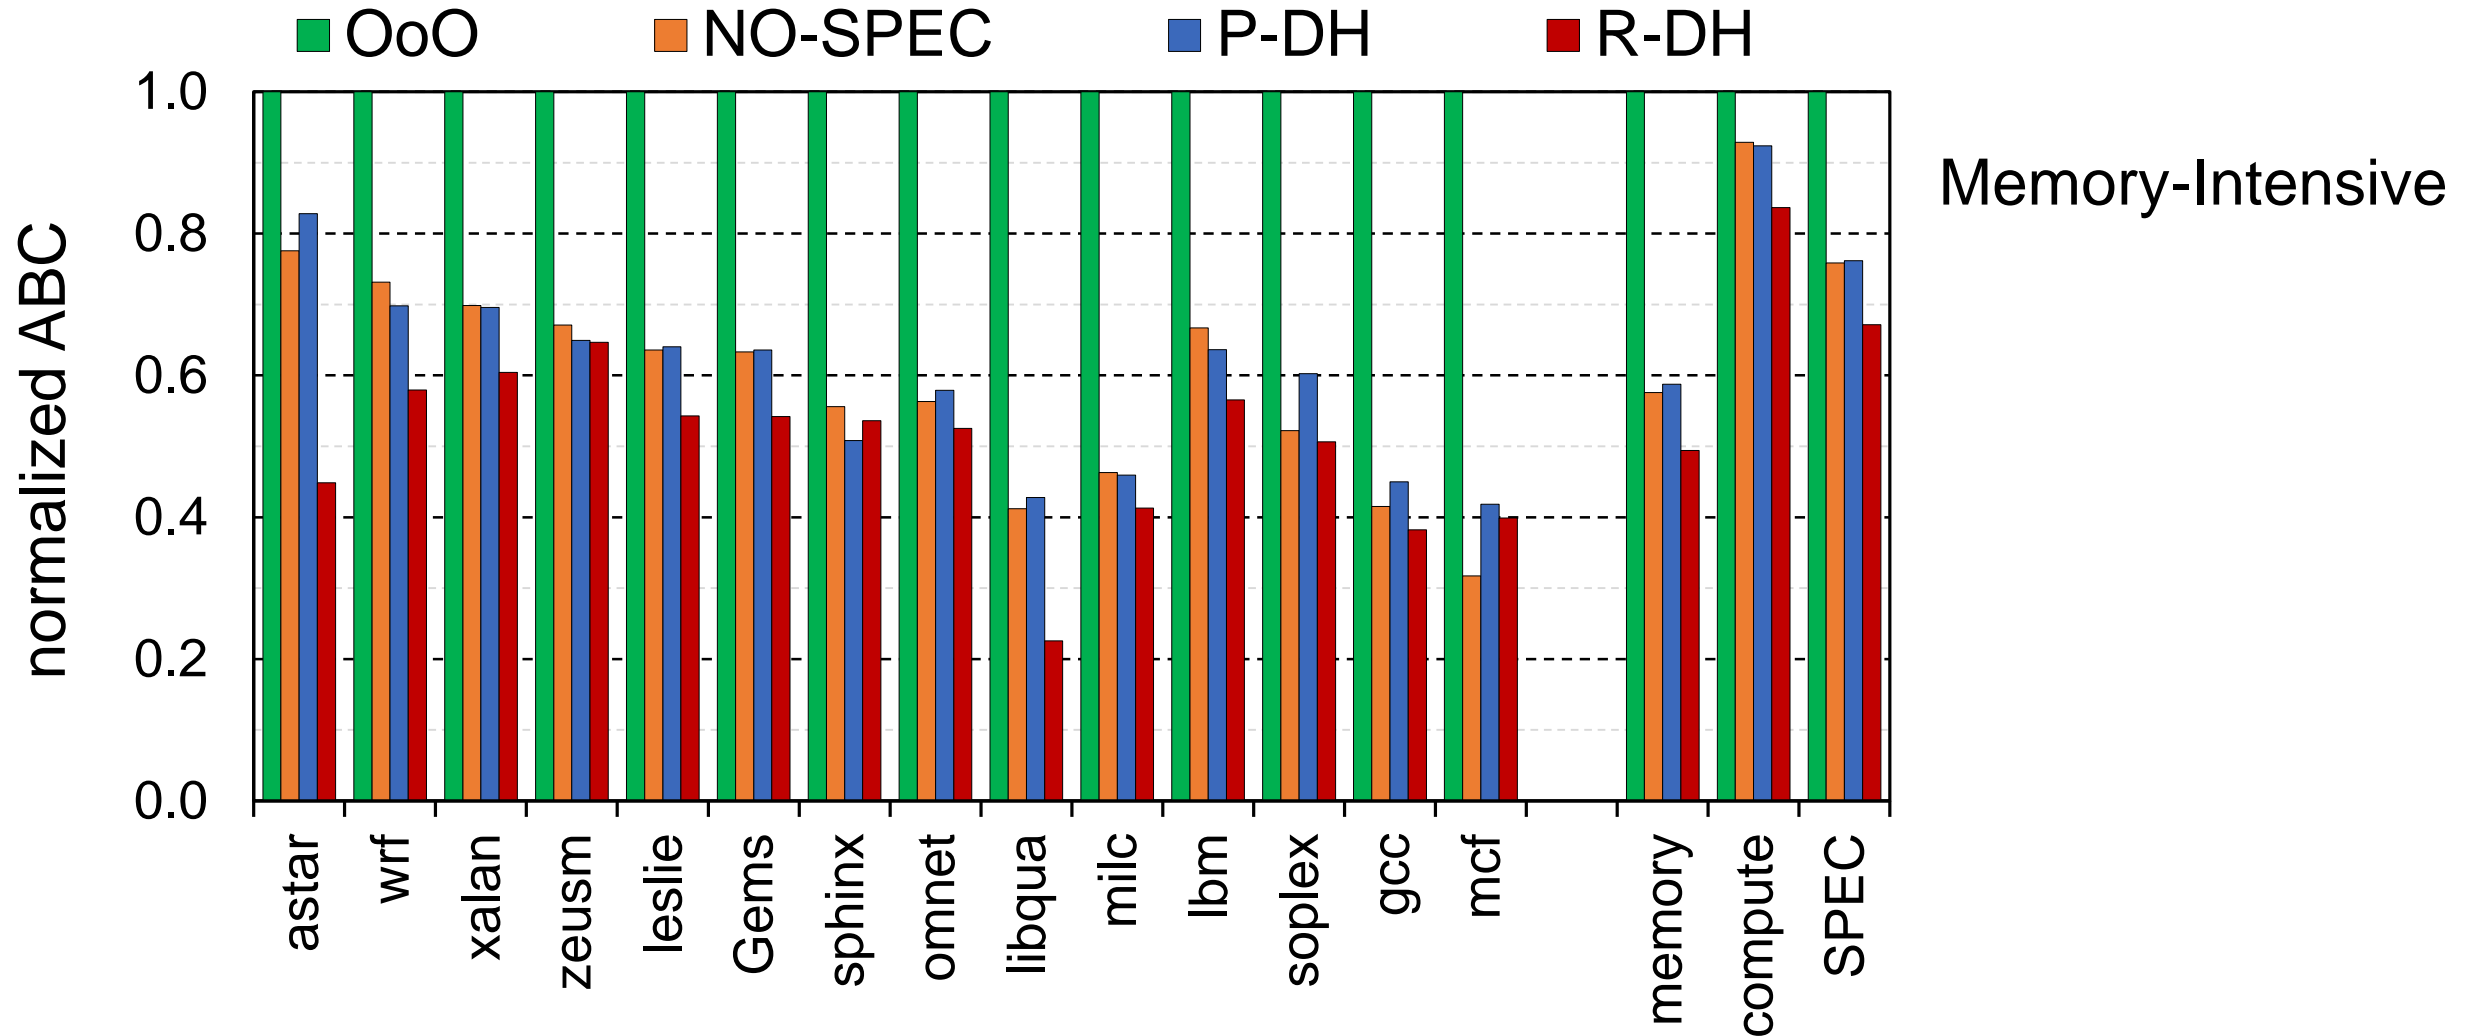
P-DH: Proactive dispatch halting

R-DH: Reactive dispatch halting

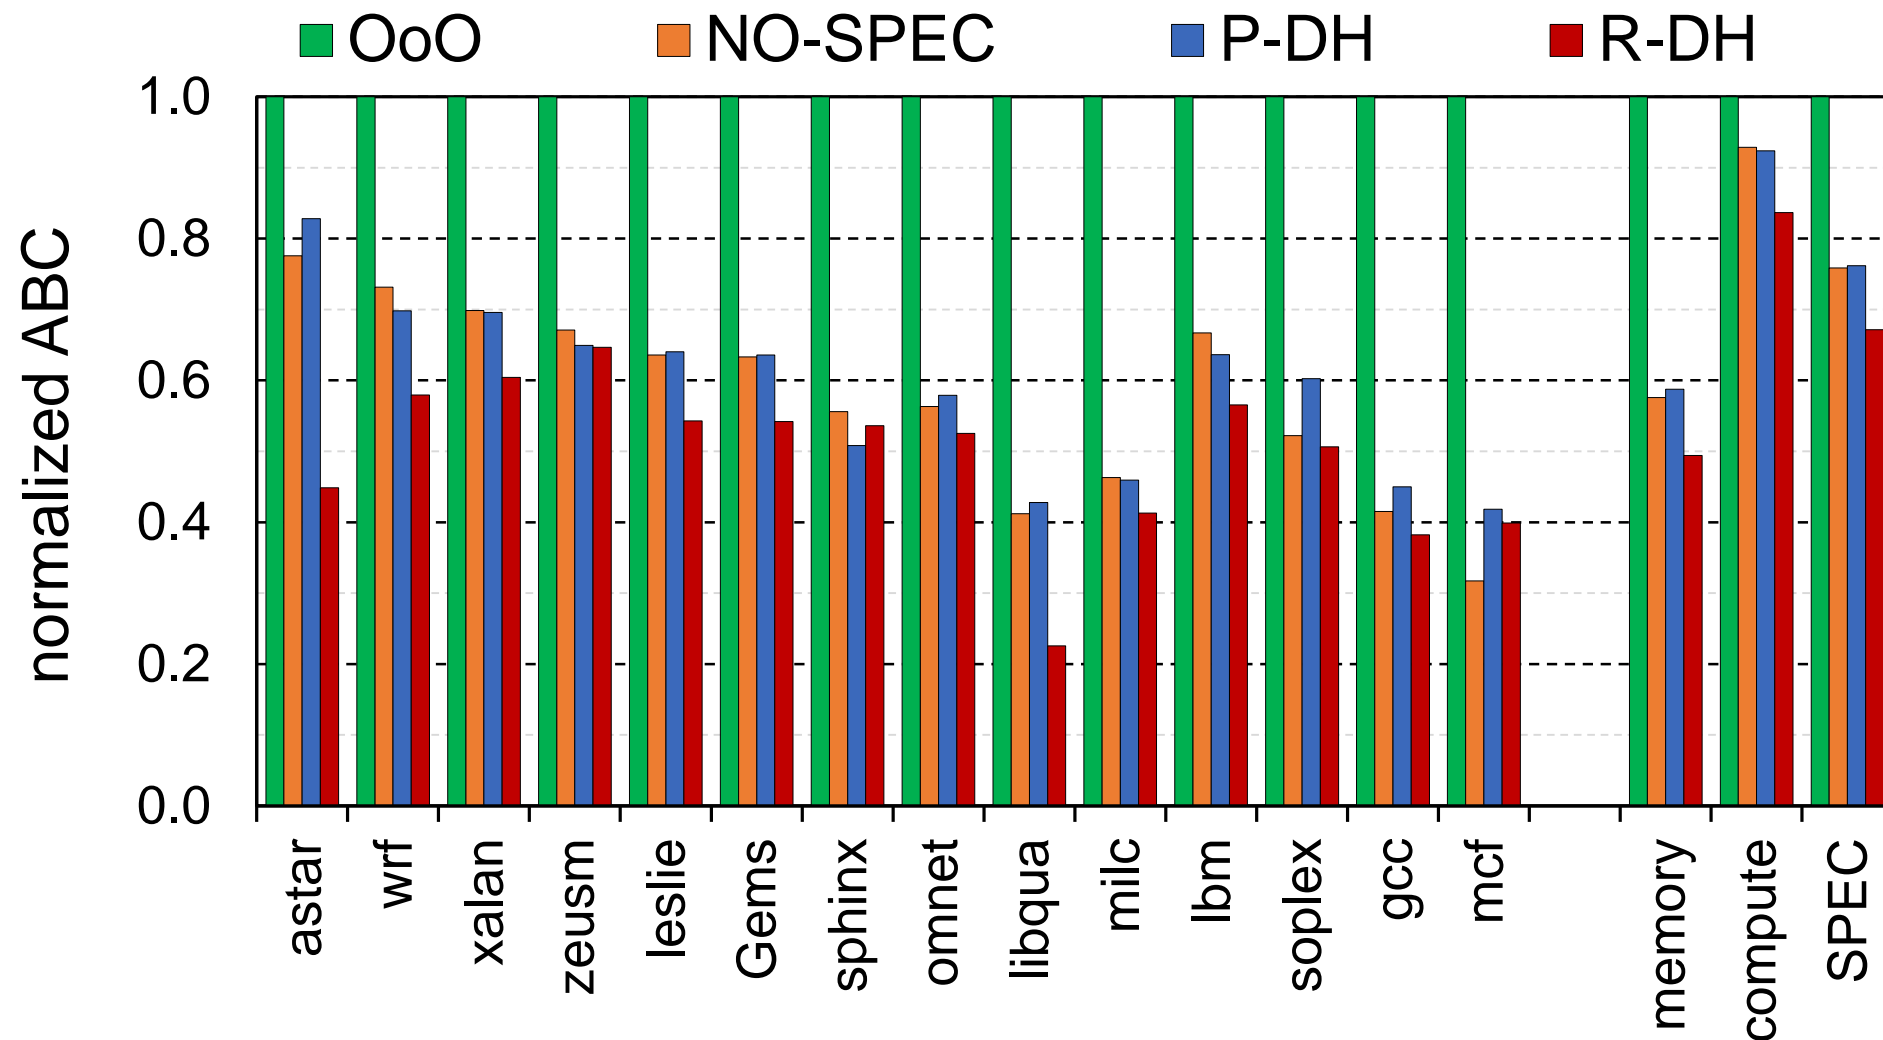
Evaluation -- Reliability



Evaluation -- Reliability



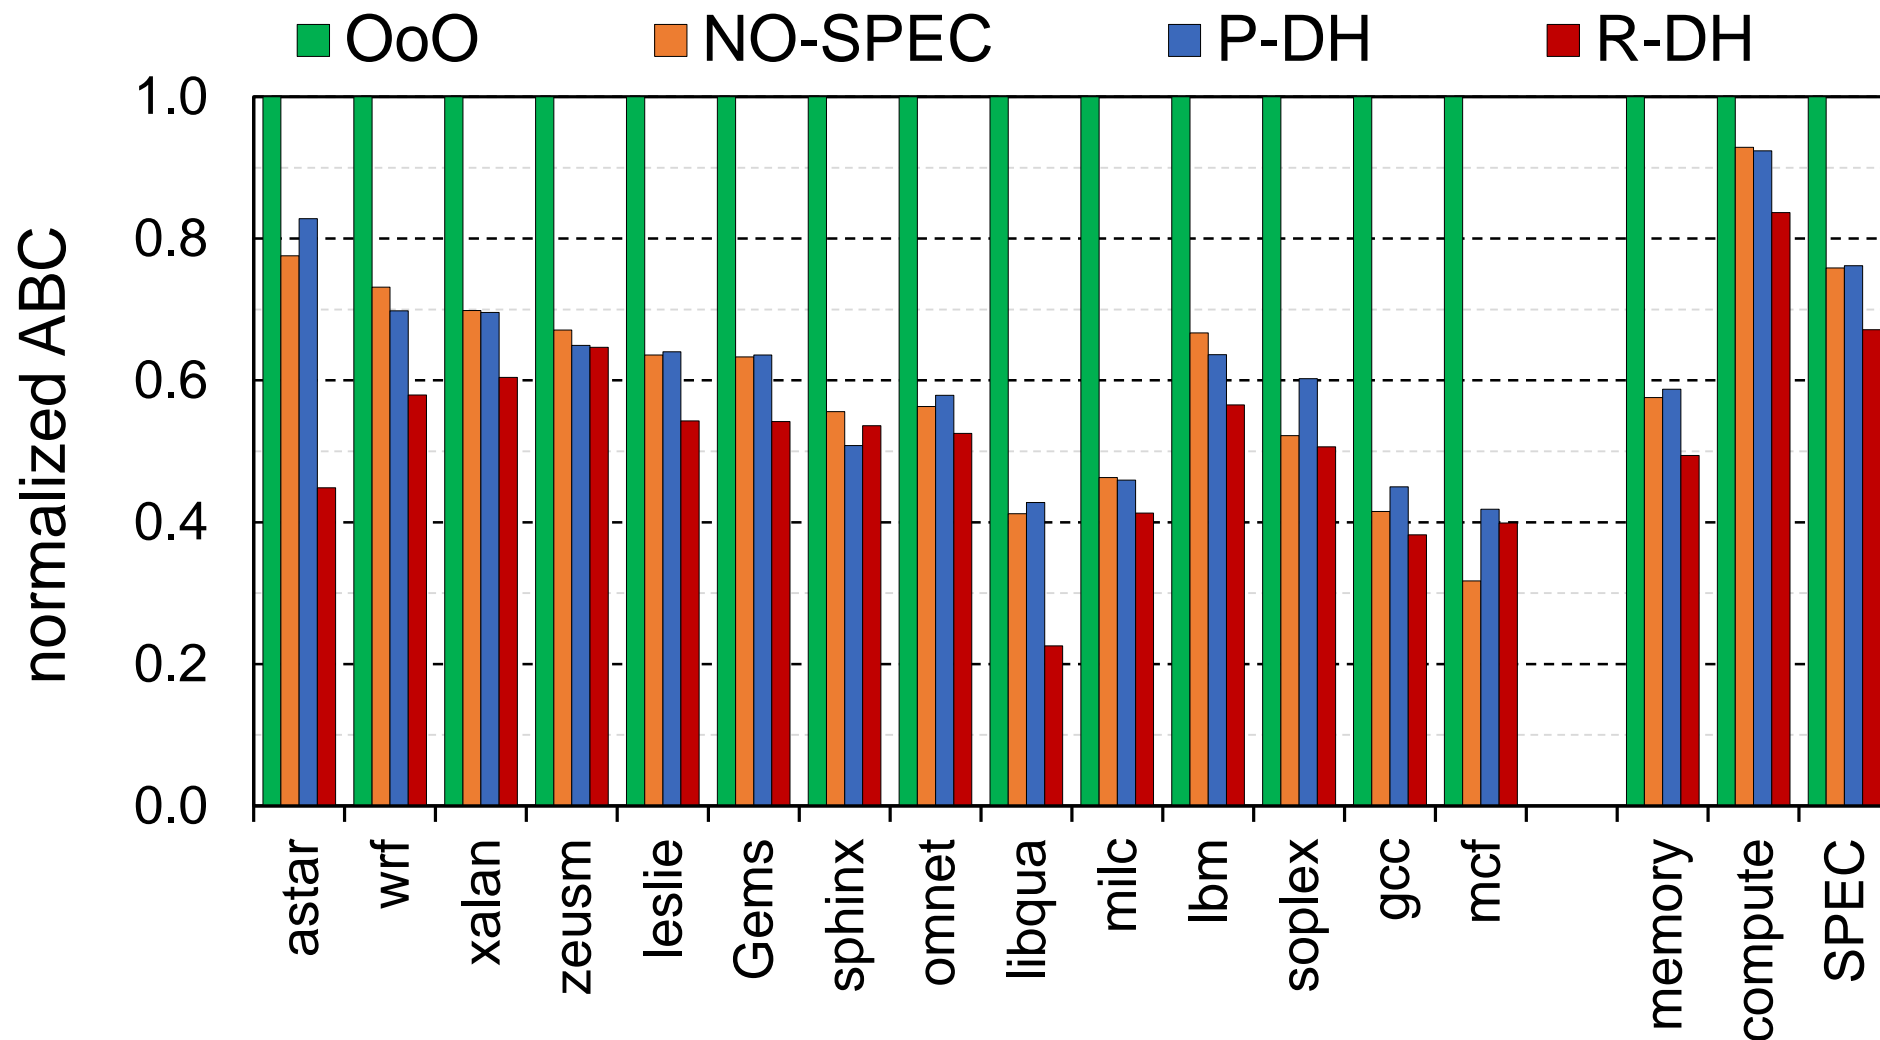
Evaluation -- Reliability



Memory-Intensive

P-DH: 41.3%

Evaluation -- Reliability

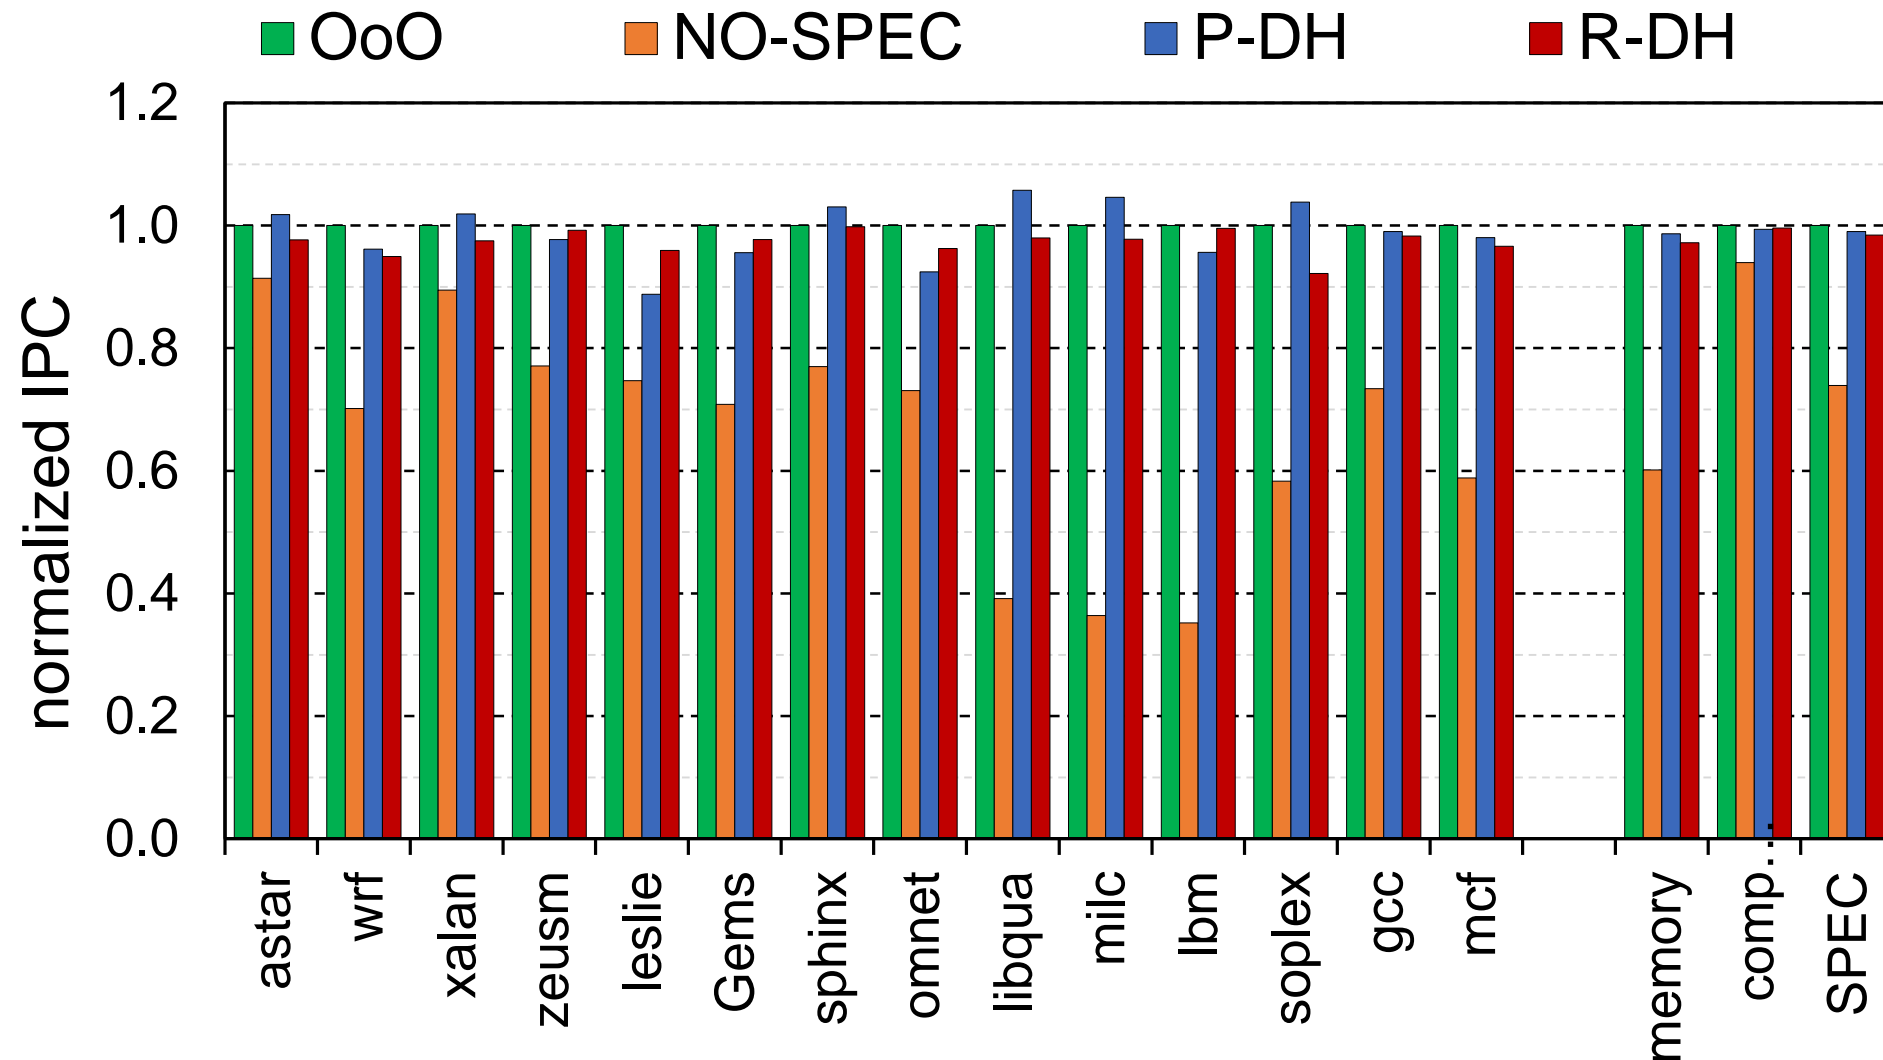


Memory-Intensive

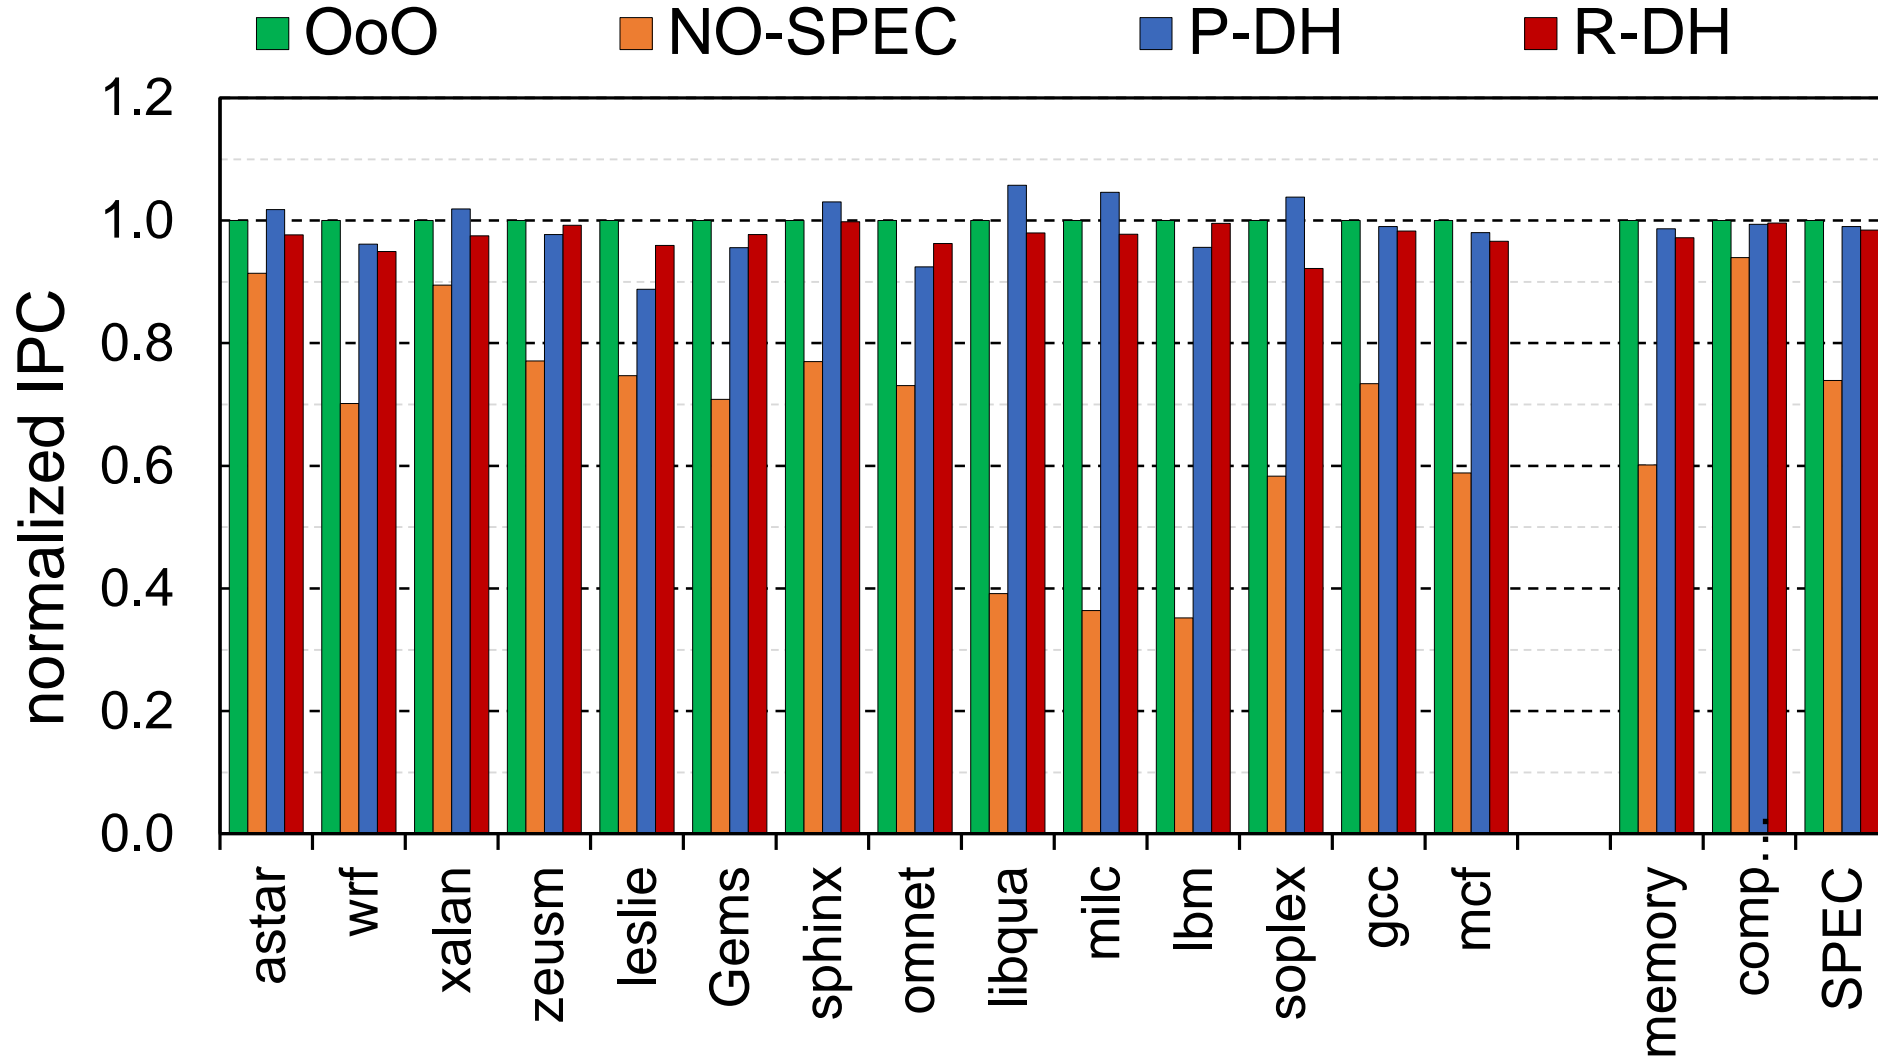
P-DH: 41.3%

R-DH: 50.6%

Evaluation -- Performance



Evaluation -- Performance



Dispatch halting
does not degrade
performance

Contribution #3 [CAL 2019, HPCA 2020]

**Precise Runahead Execution
to Improve Performance
and Evaluate Reliability**

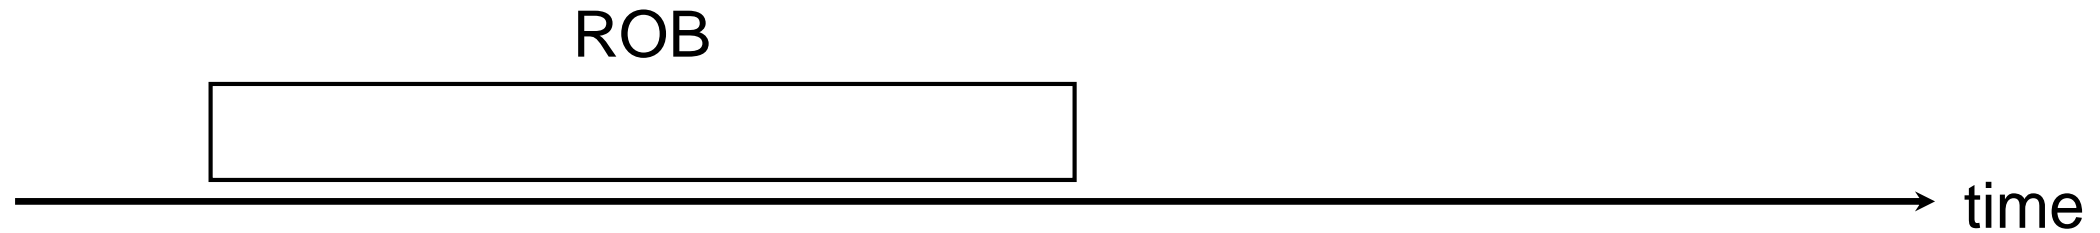
Full-ROB Stalls Degrade Performance

Full-ROB Stalls Degrade Performance

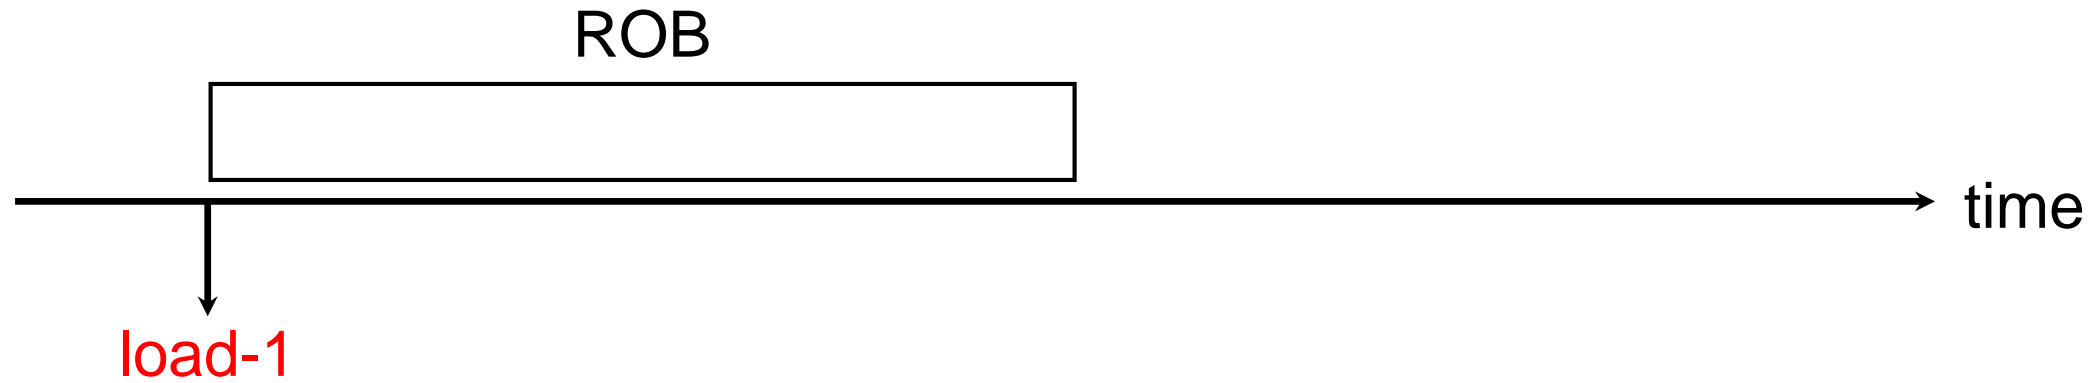


time

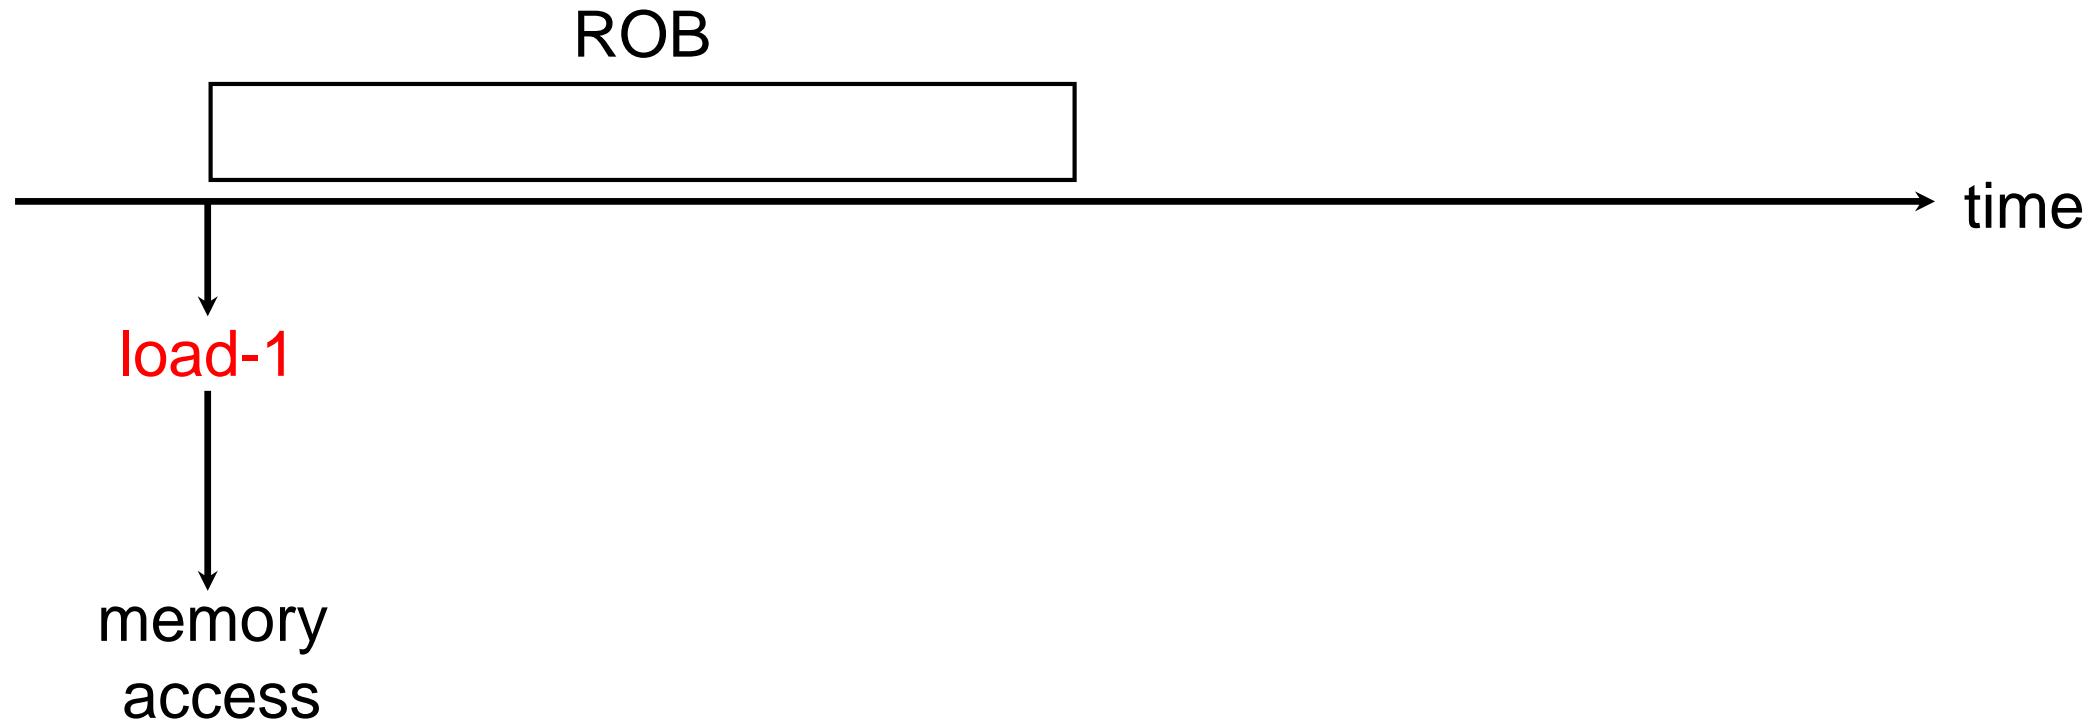
Full-ROB Stalls Degrade Performance



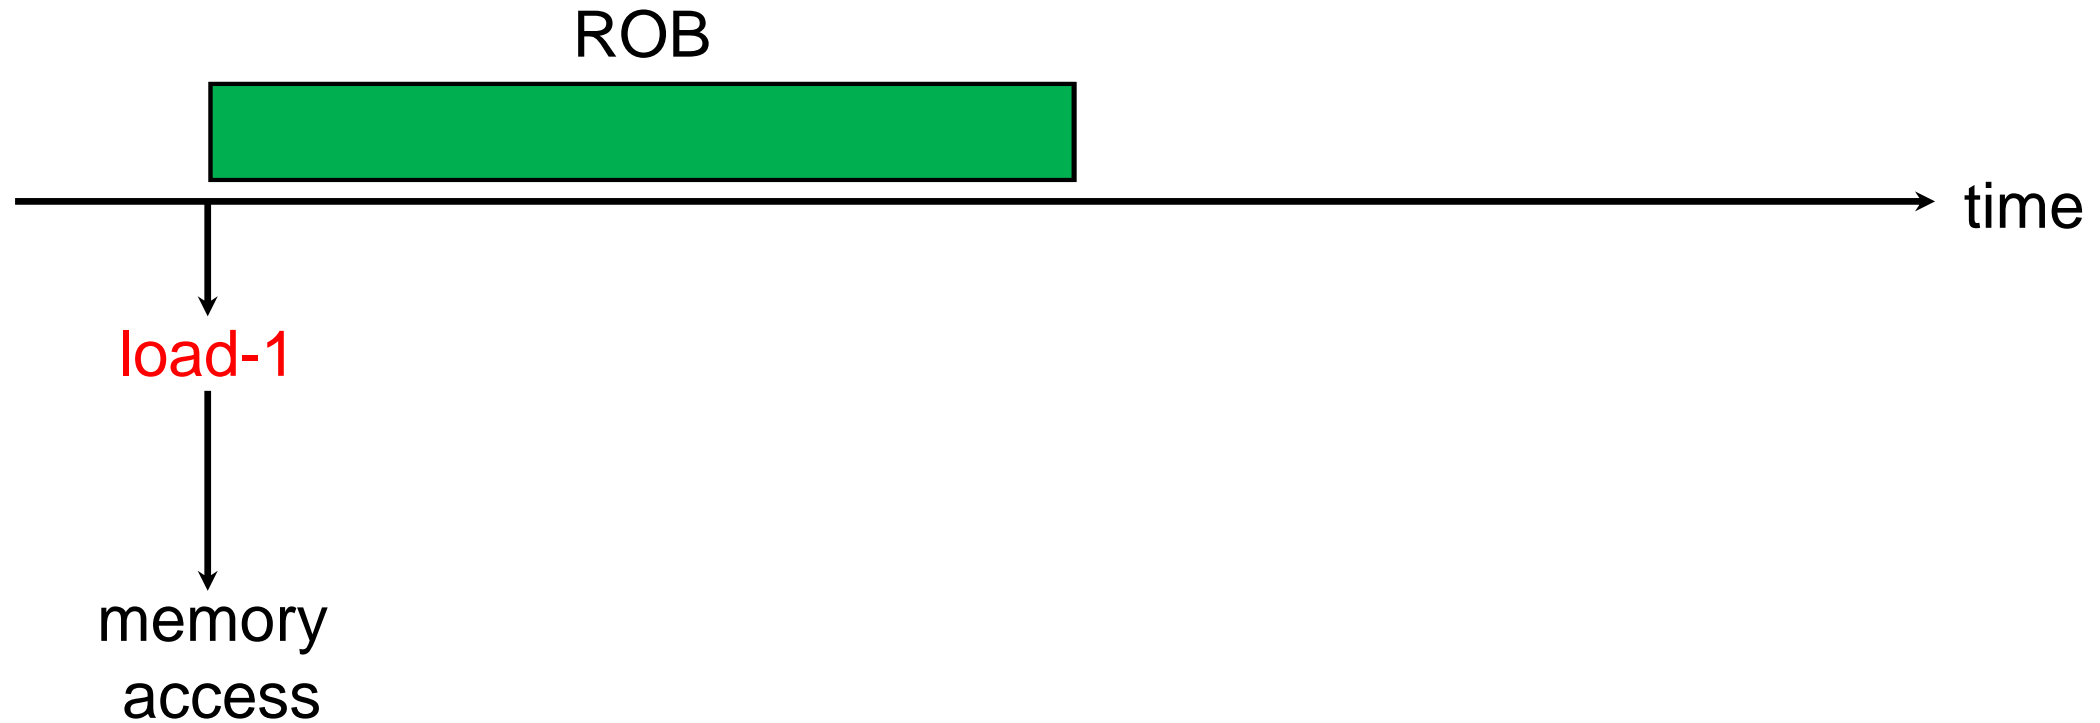
Full-ROB Stalls Degrade Performance



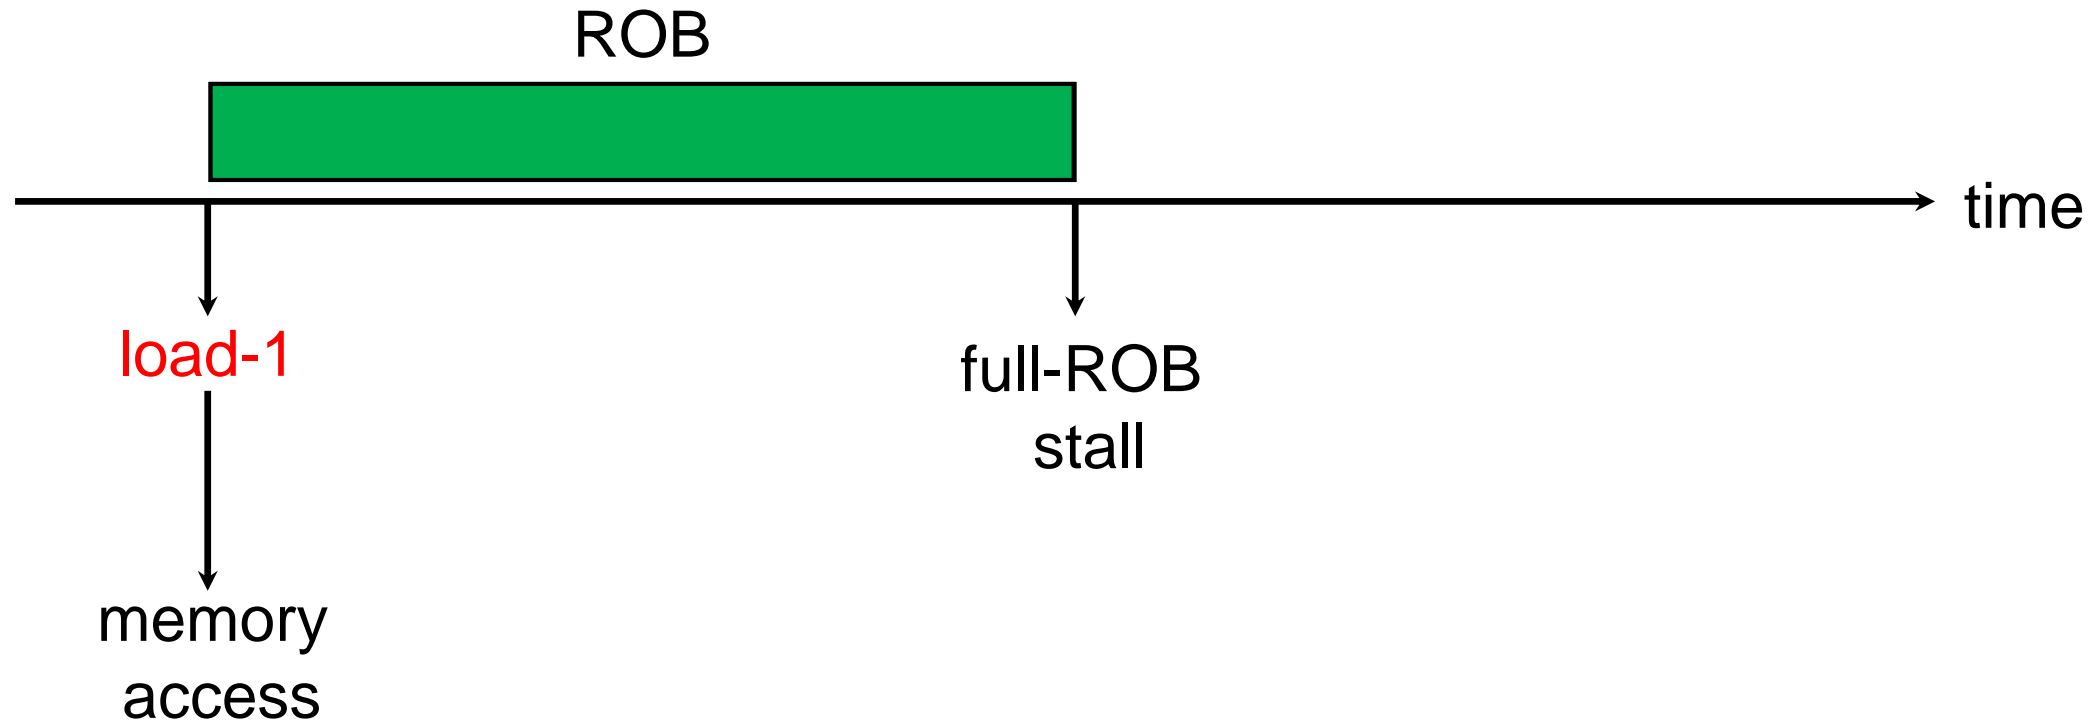
Full-ROB Stalls Degrade Performance



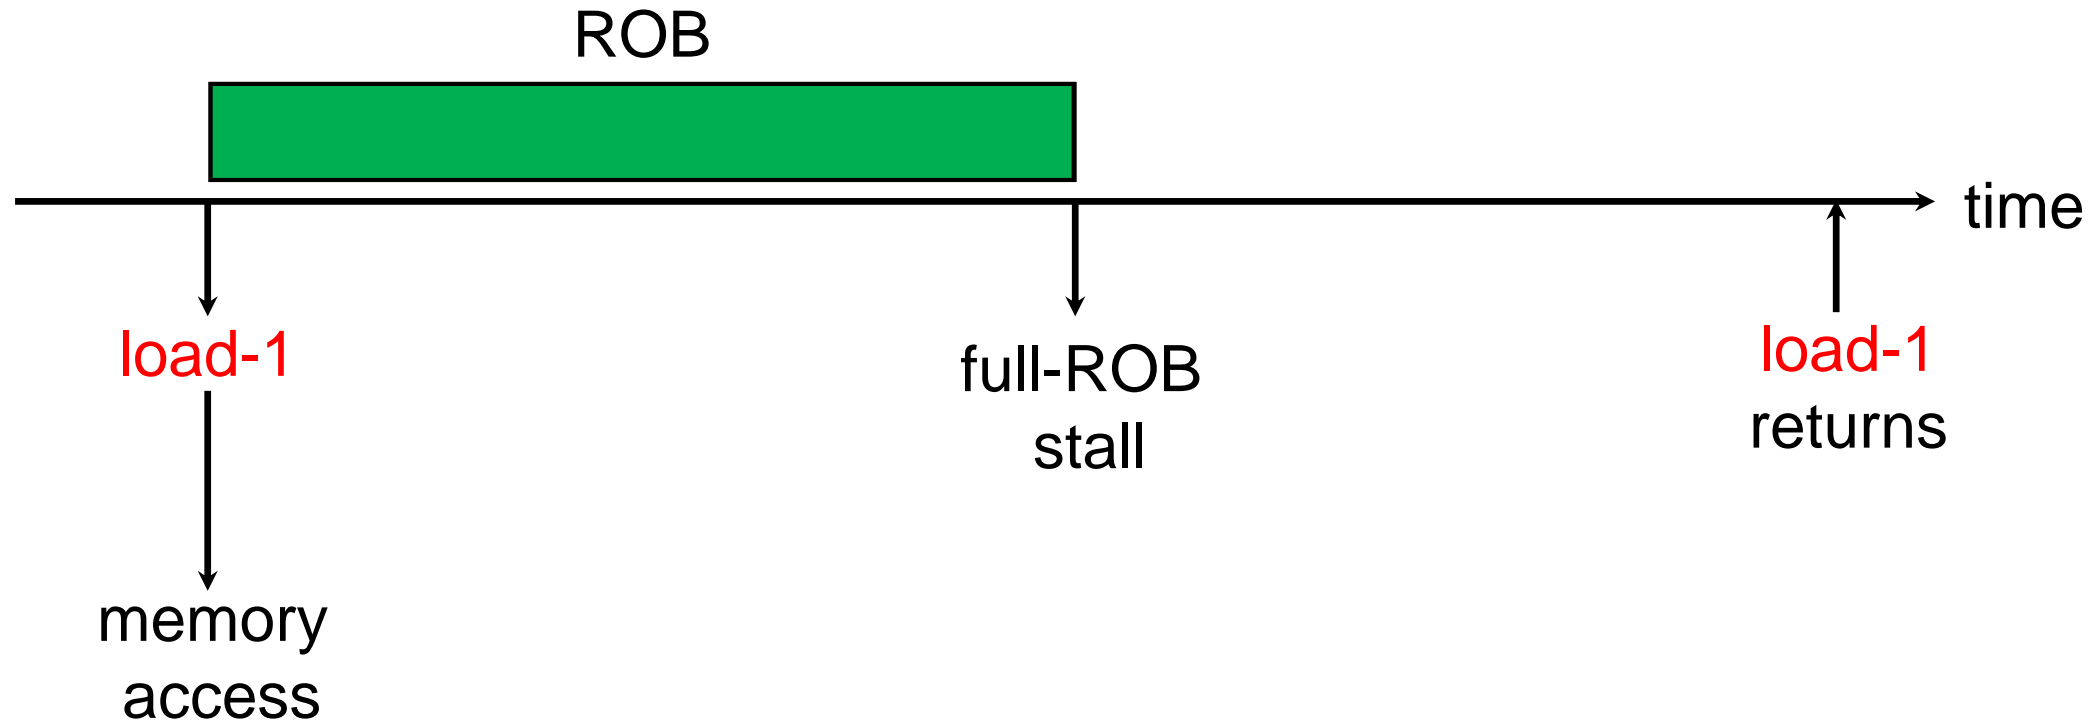
Full-ROB Stalls Degrade Performance



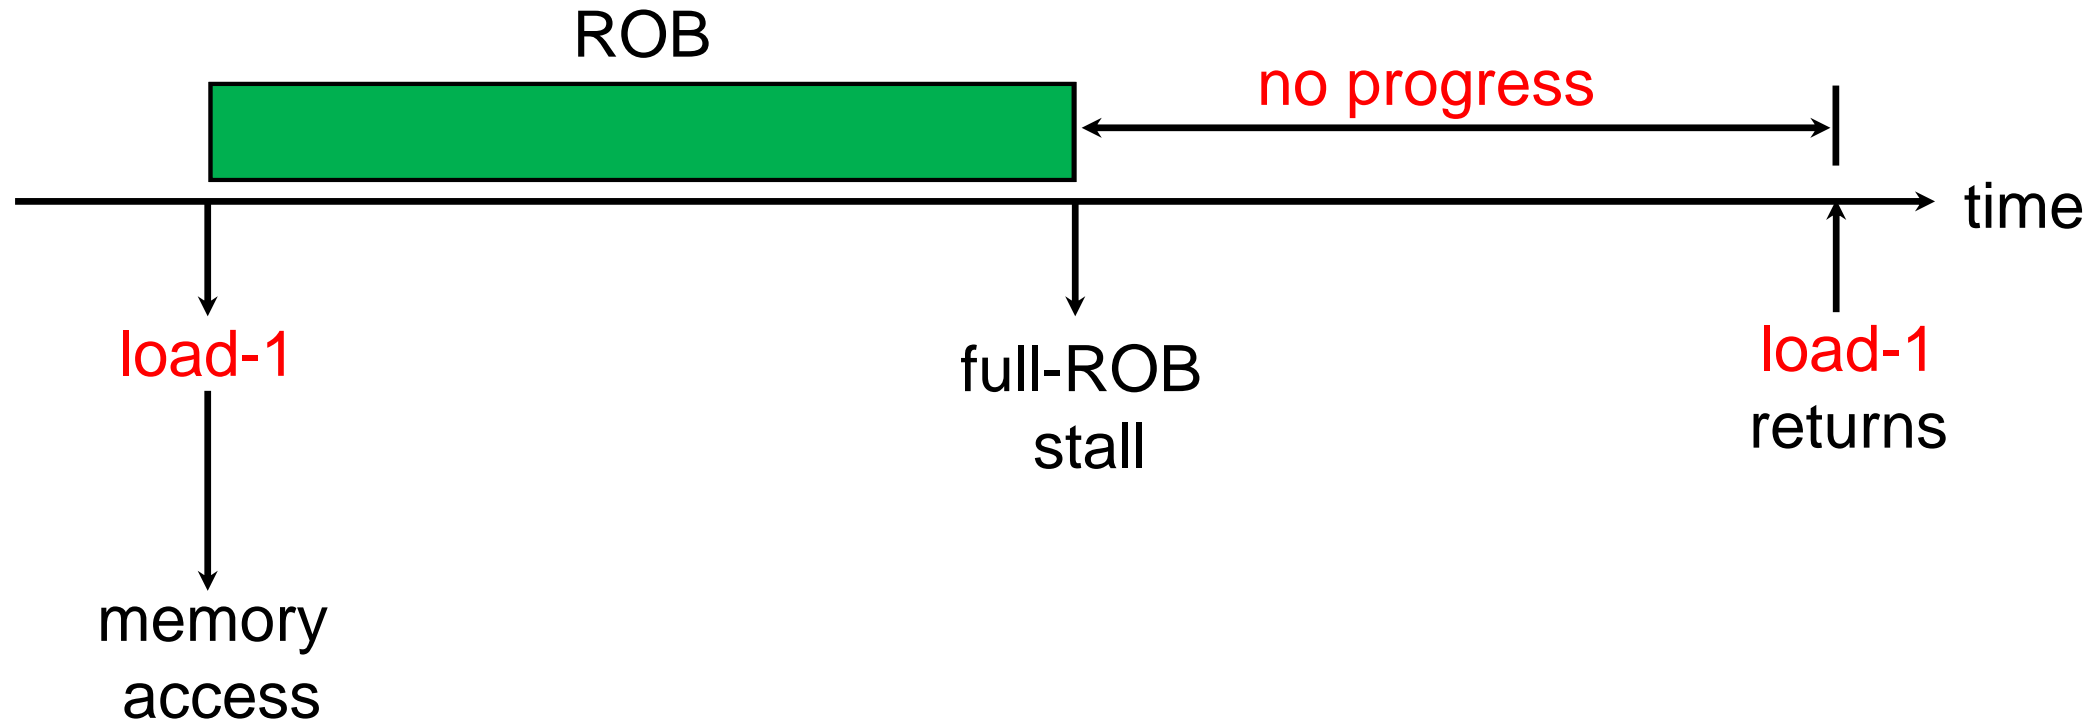
Full-ROB Stalls Degrade Performance



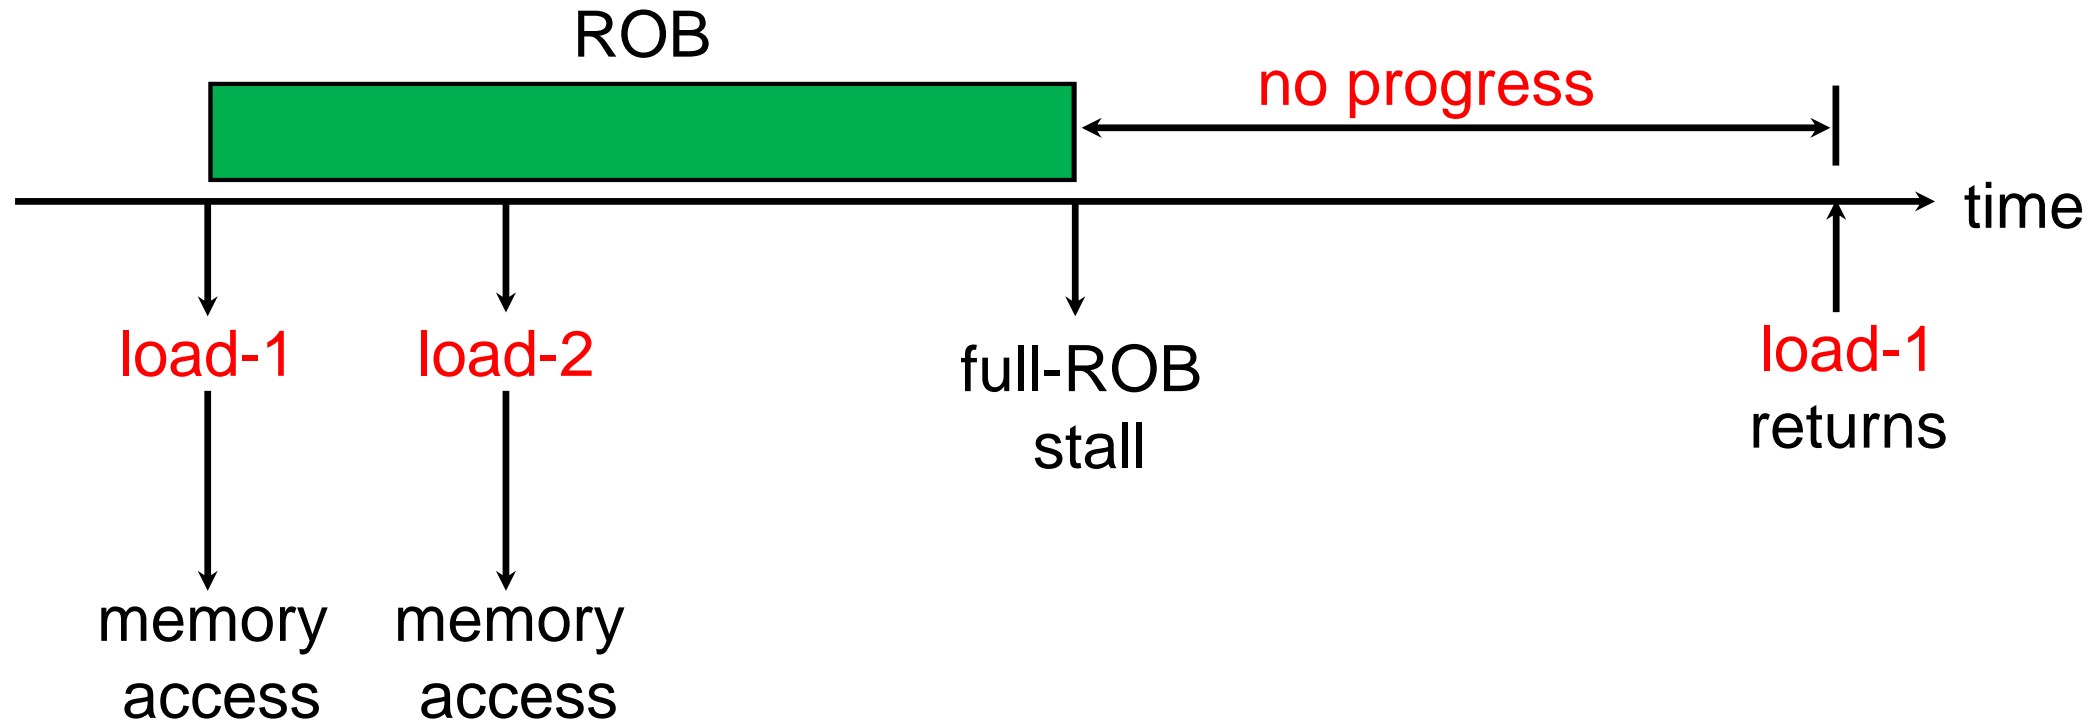
Full-ROB Stalls Degrade Performance



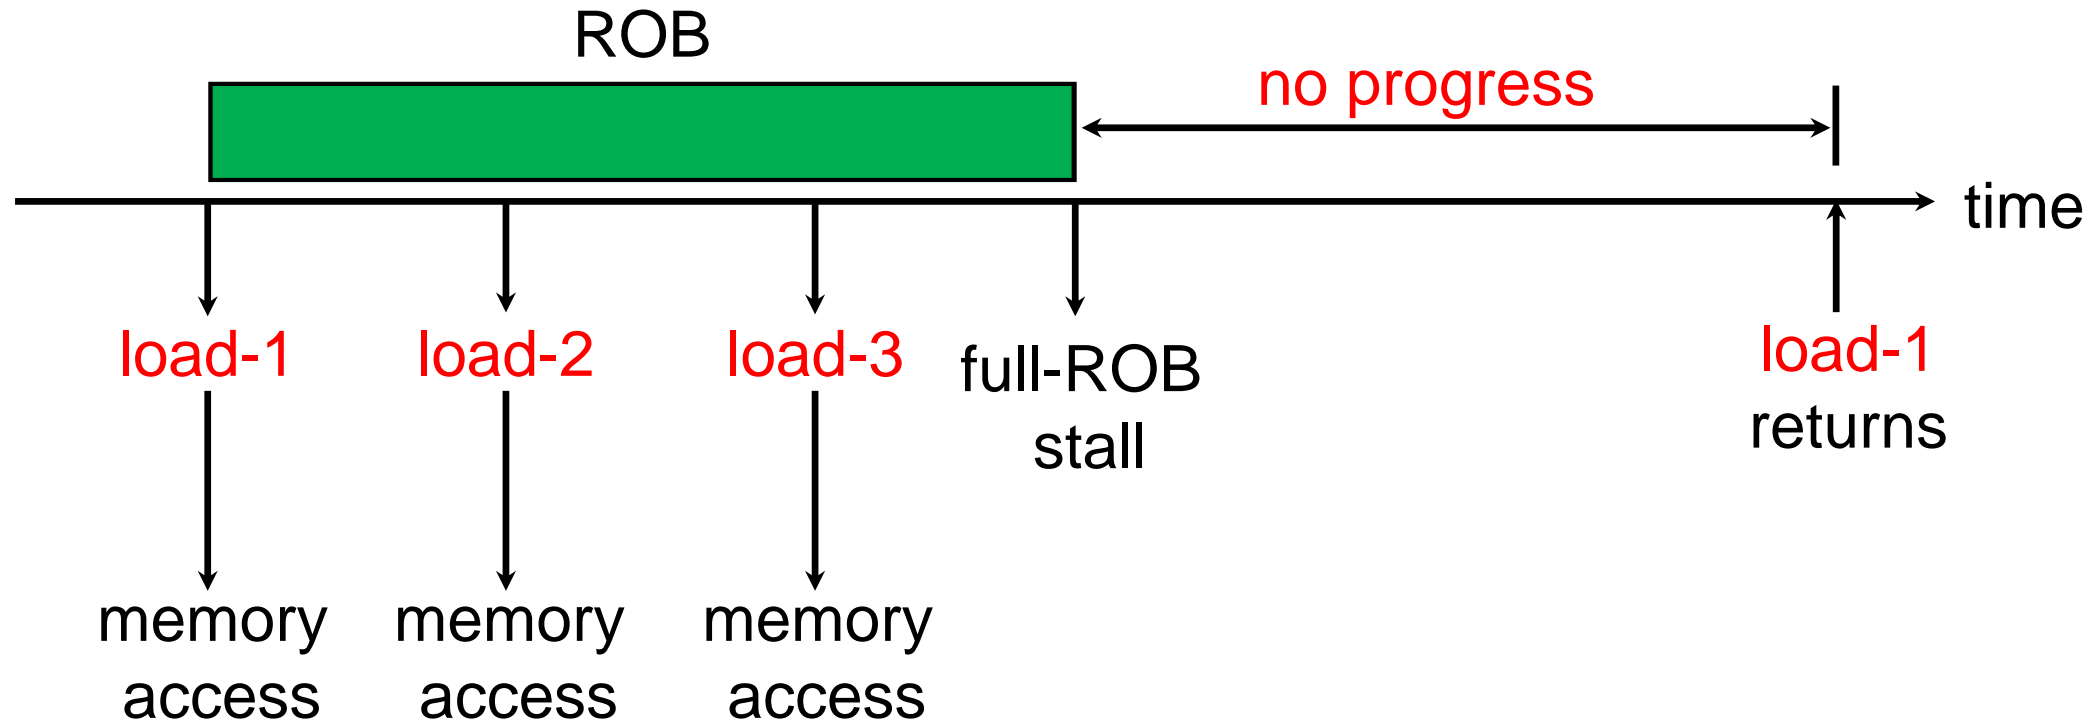
Full-ROB Stalls Degrade Performance



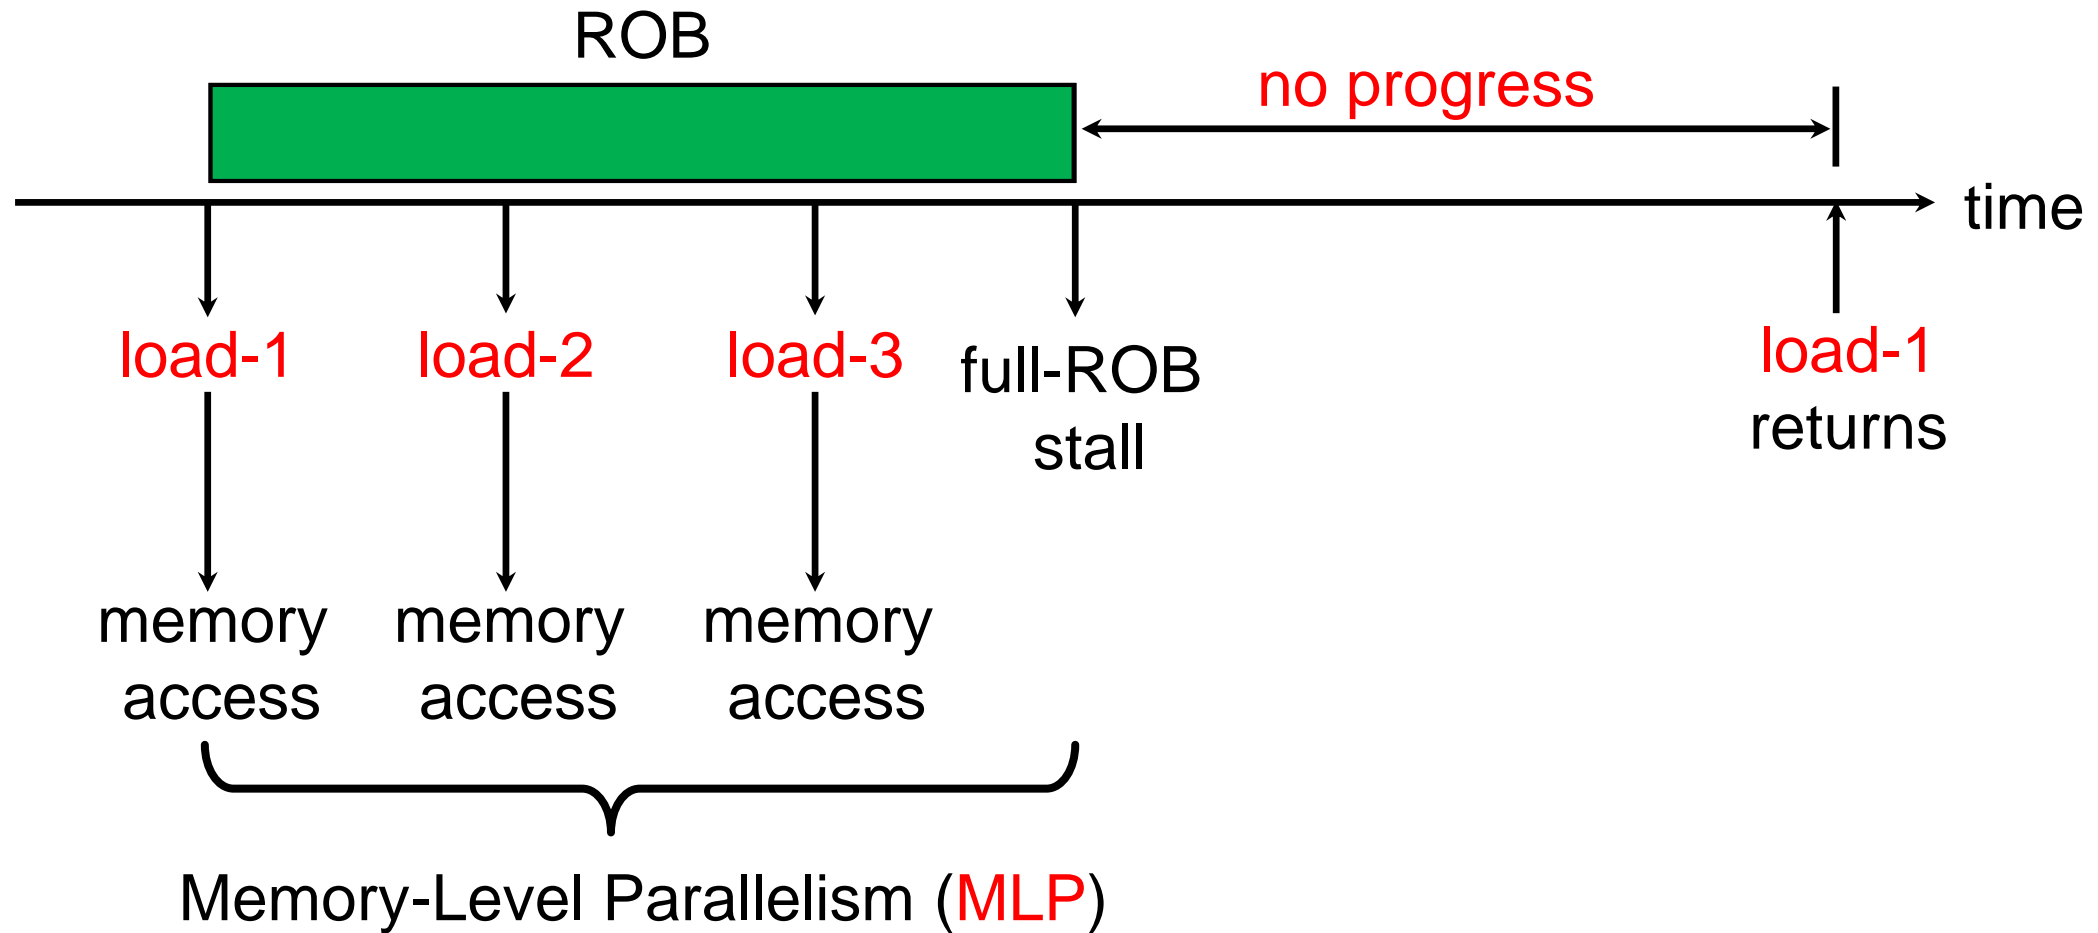
Full-ROB Stalls Degrade Performance



Full-ROB Stalls Degrade Performance



Full-ROB Stalls Degrade Performance

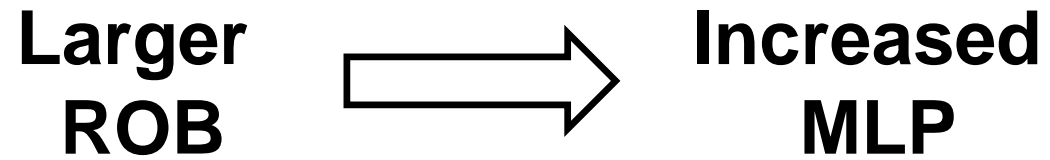


Full-ROB Stalls Degrade Performance

Full-ROB Stalls Degrade Performance

**Larger
ROB**

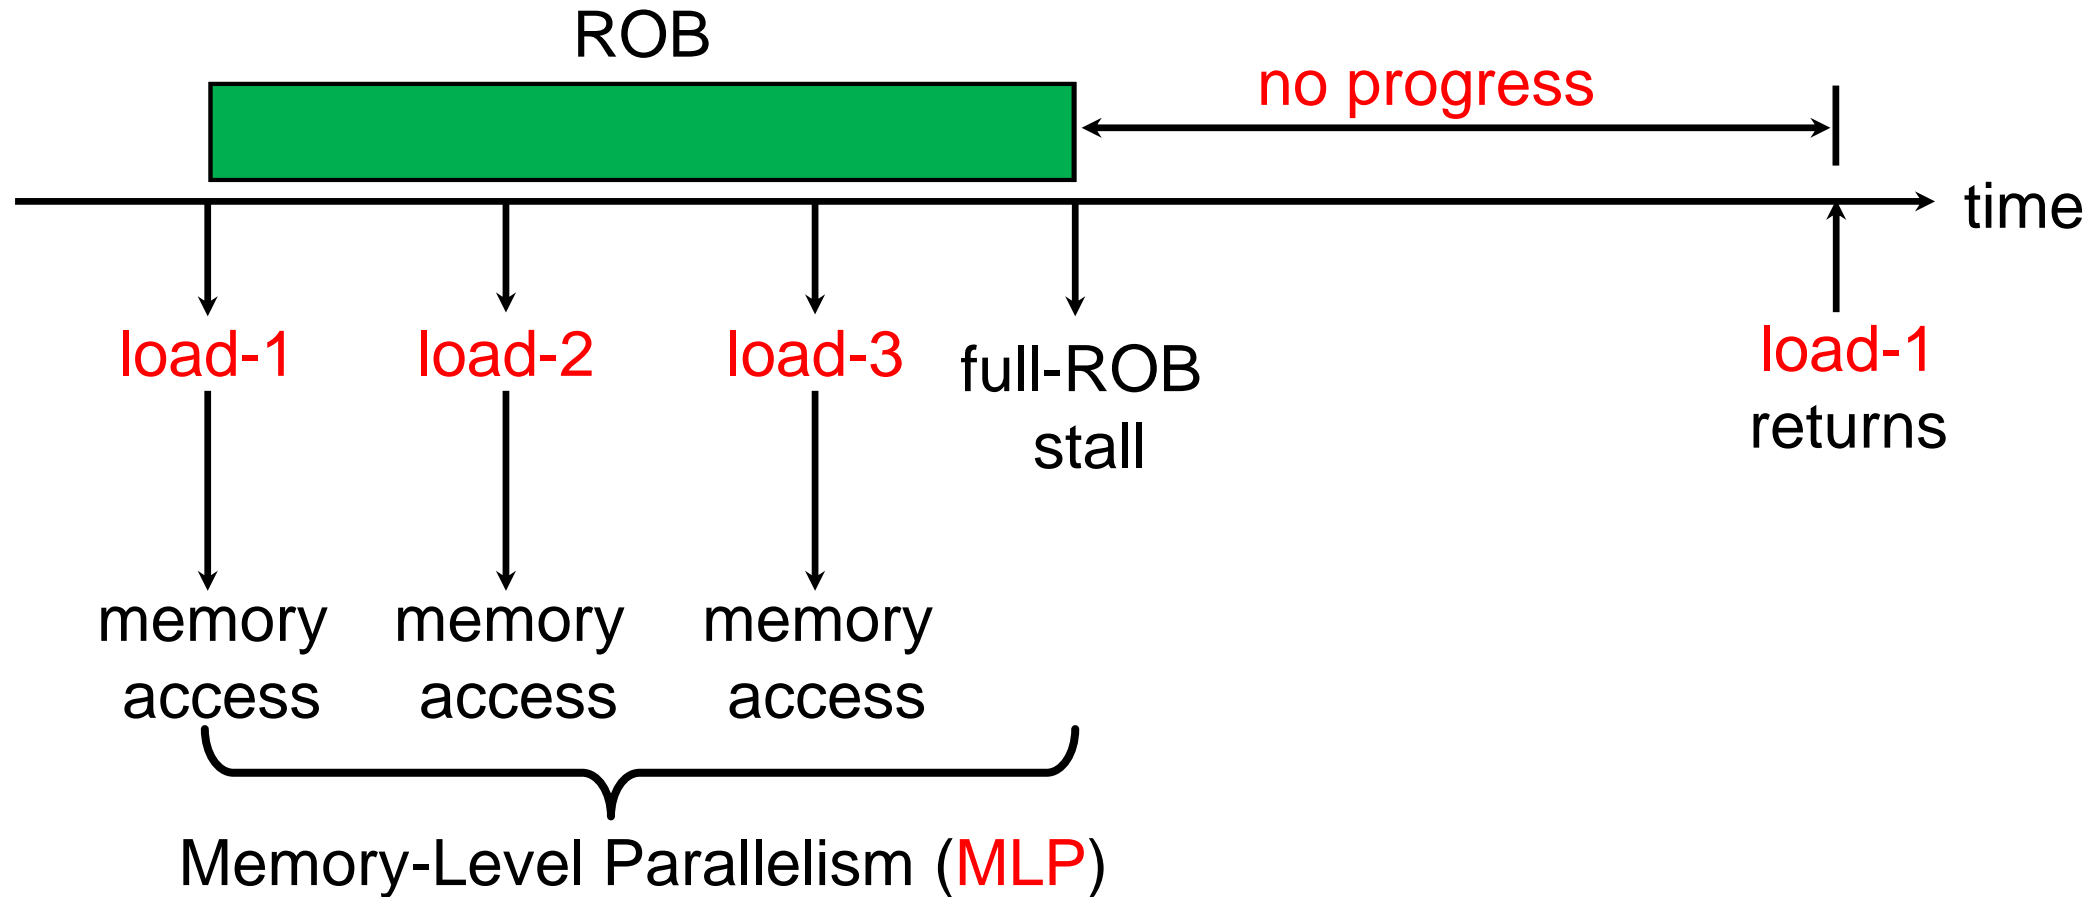
Full-ROB Stalls Degrade Performance



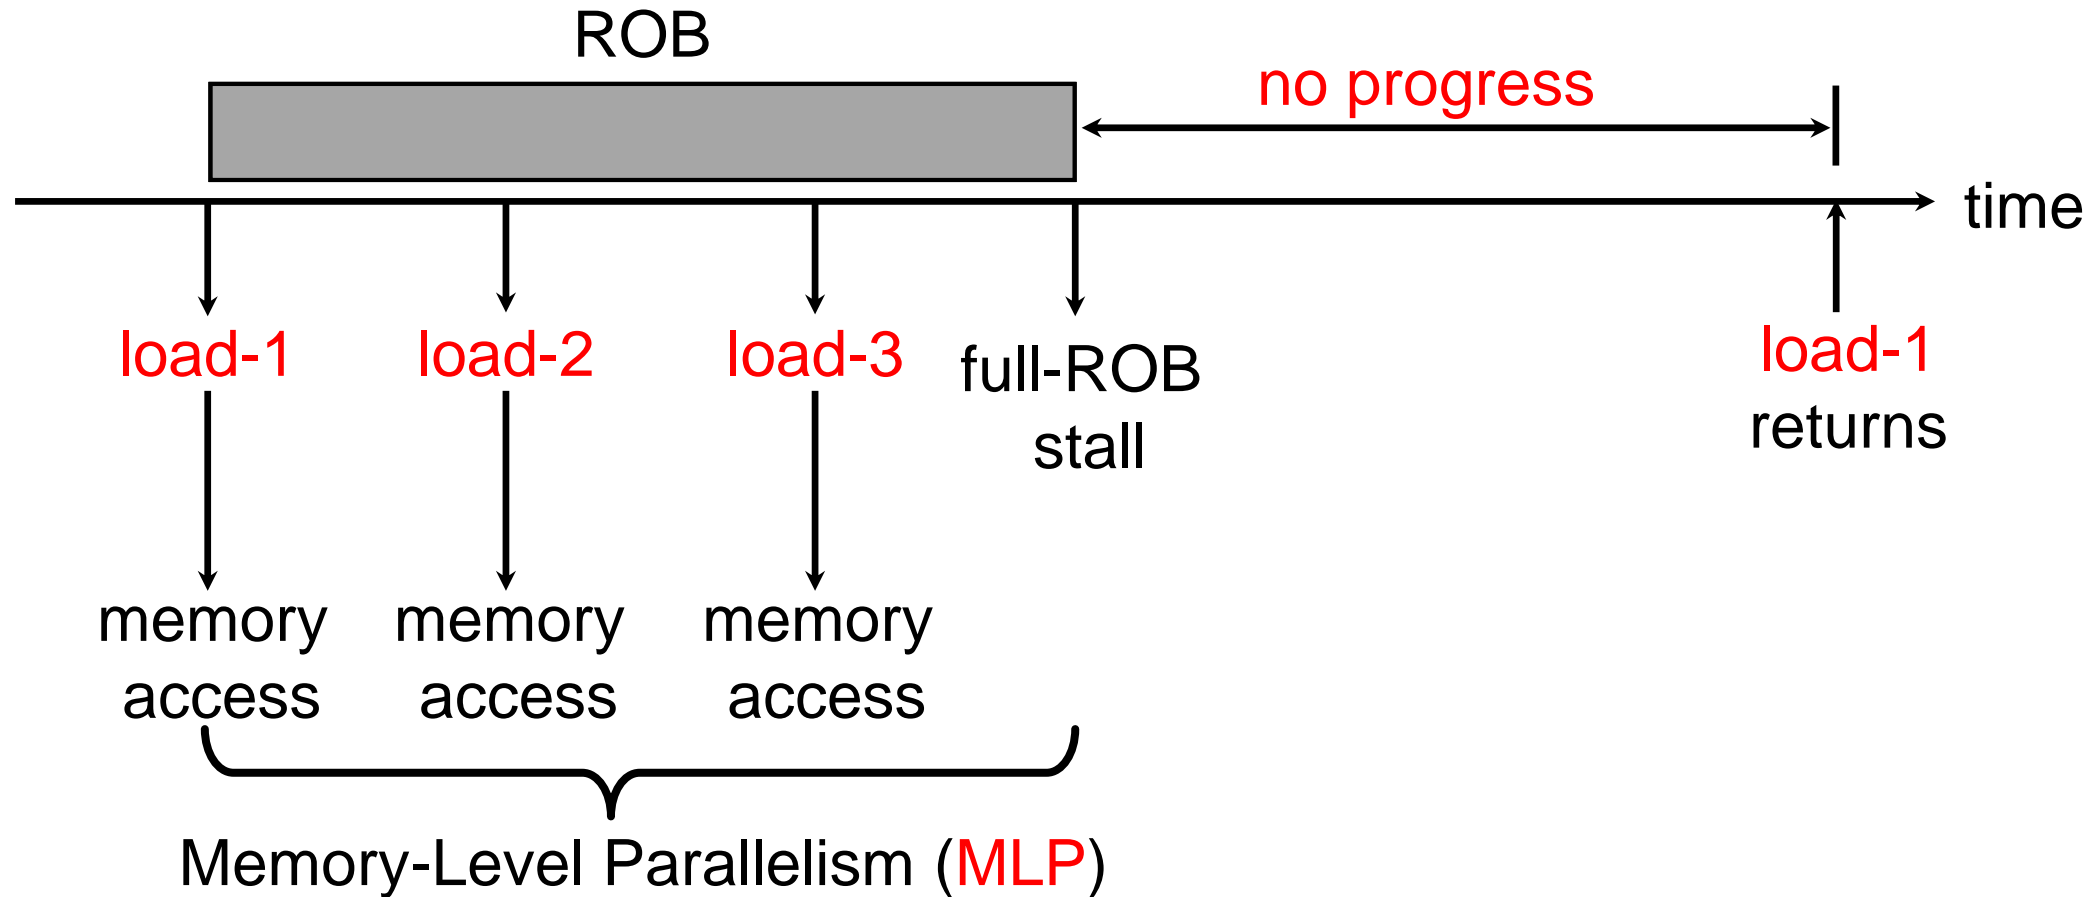
Full-ROB Stalls Degrade Performance



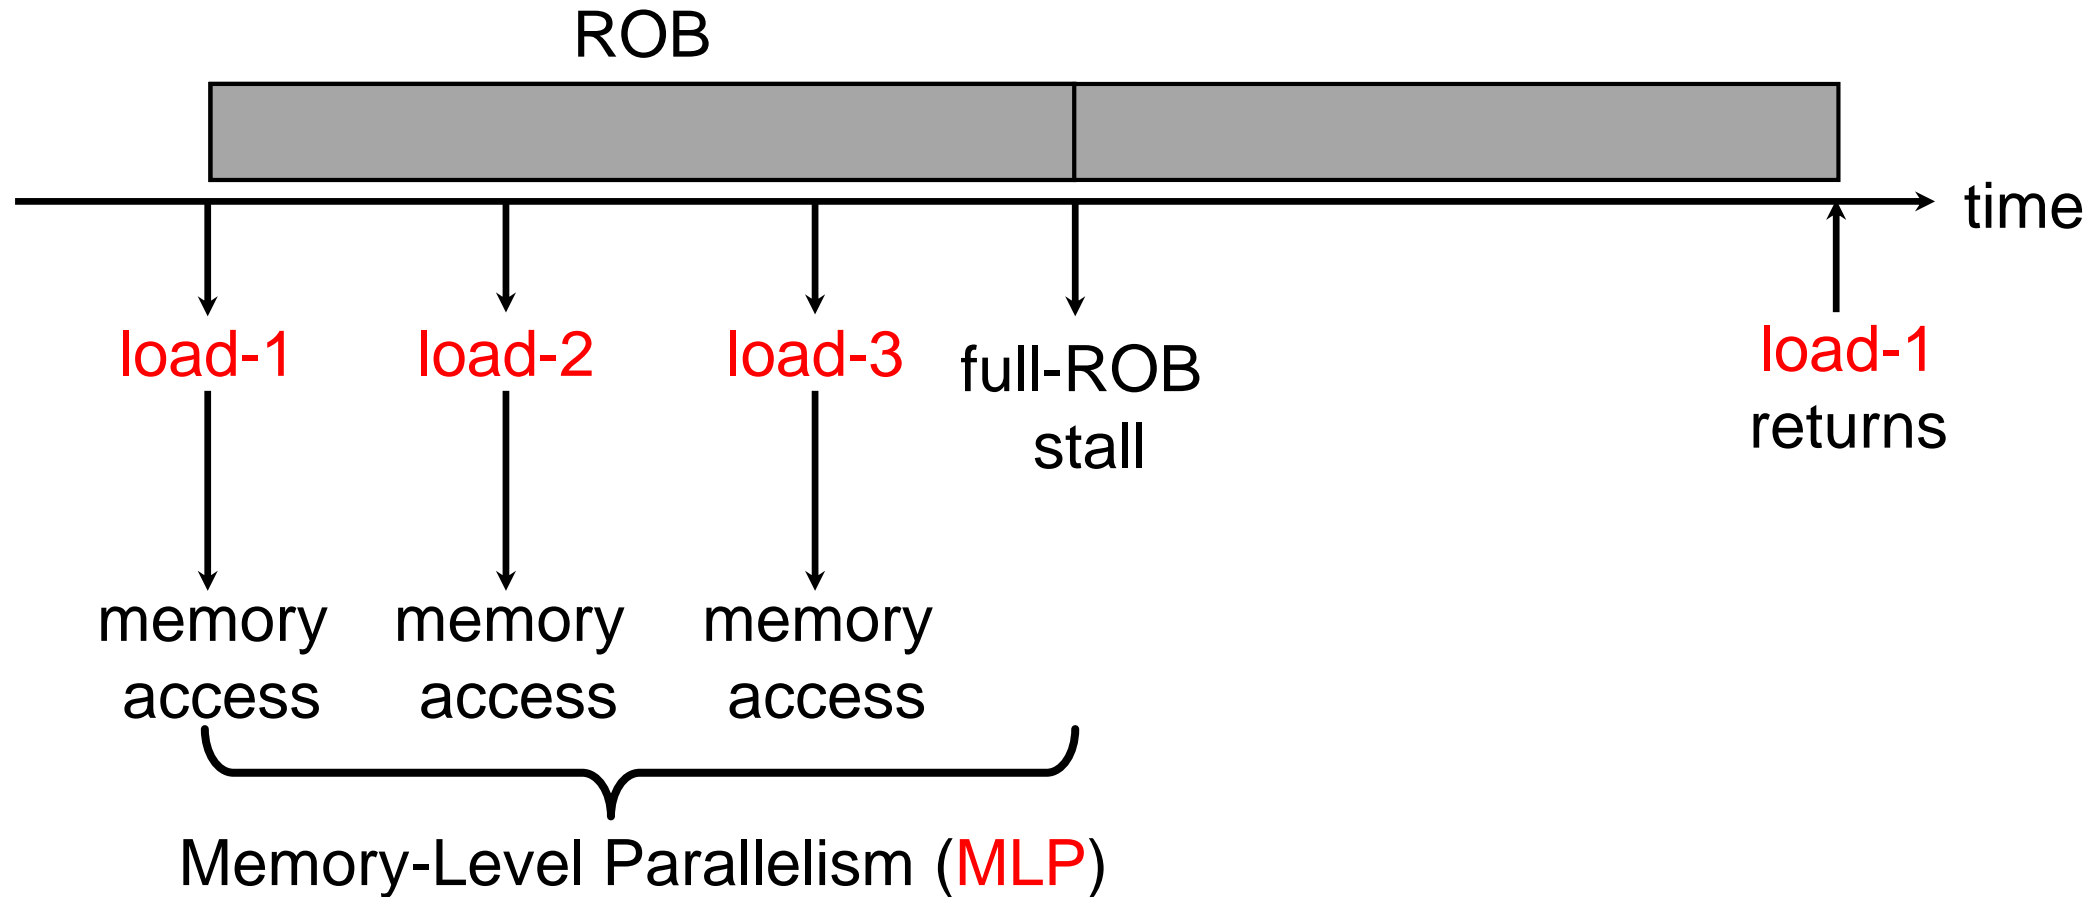
Runahead Execution Prefetches under a Full-ROB Stall



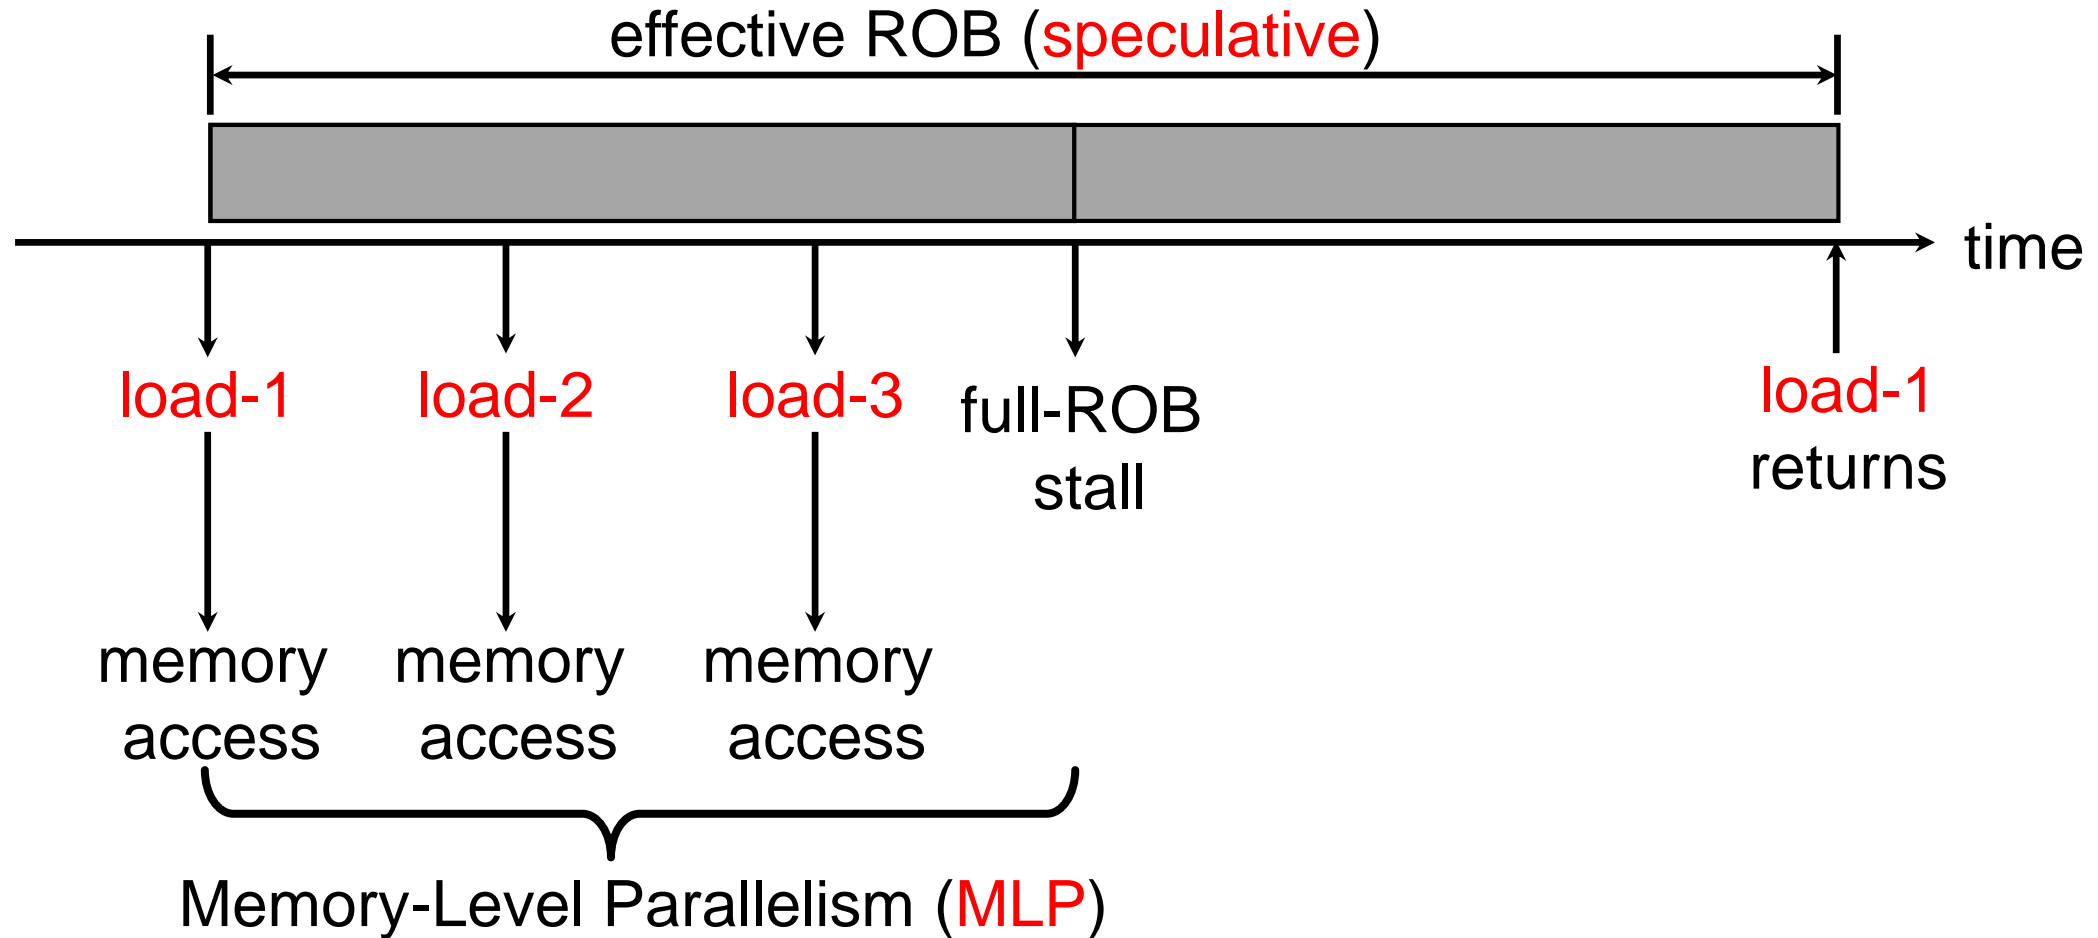
Runahead Execution Prefetches under a Full-ROB Stall



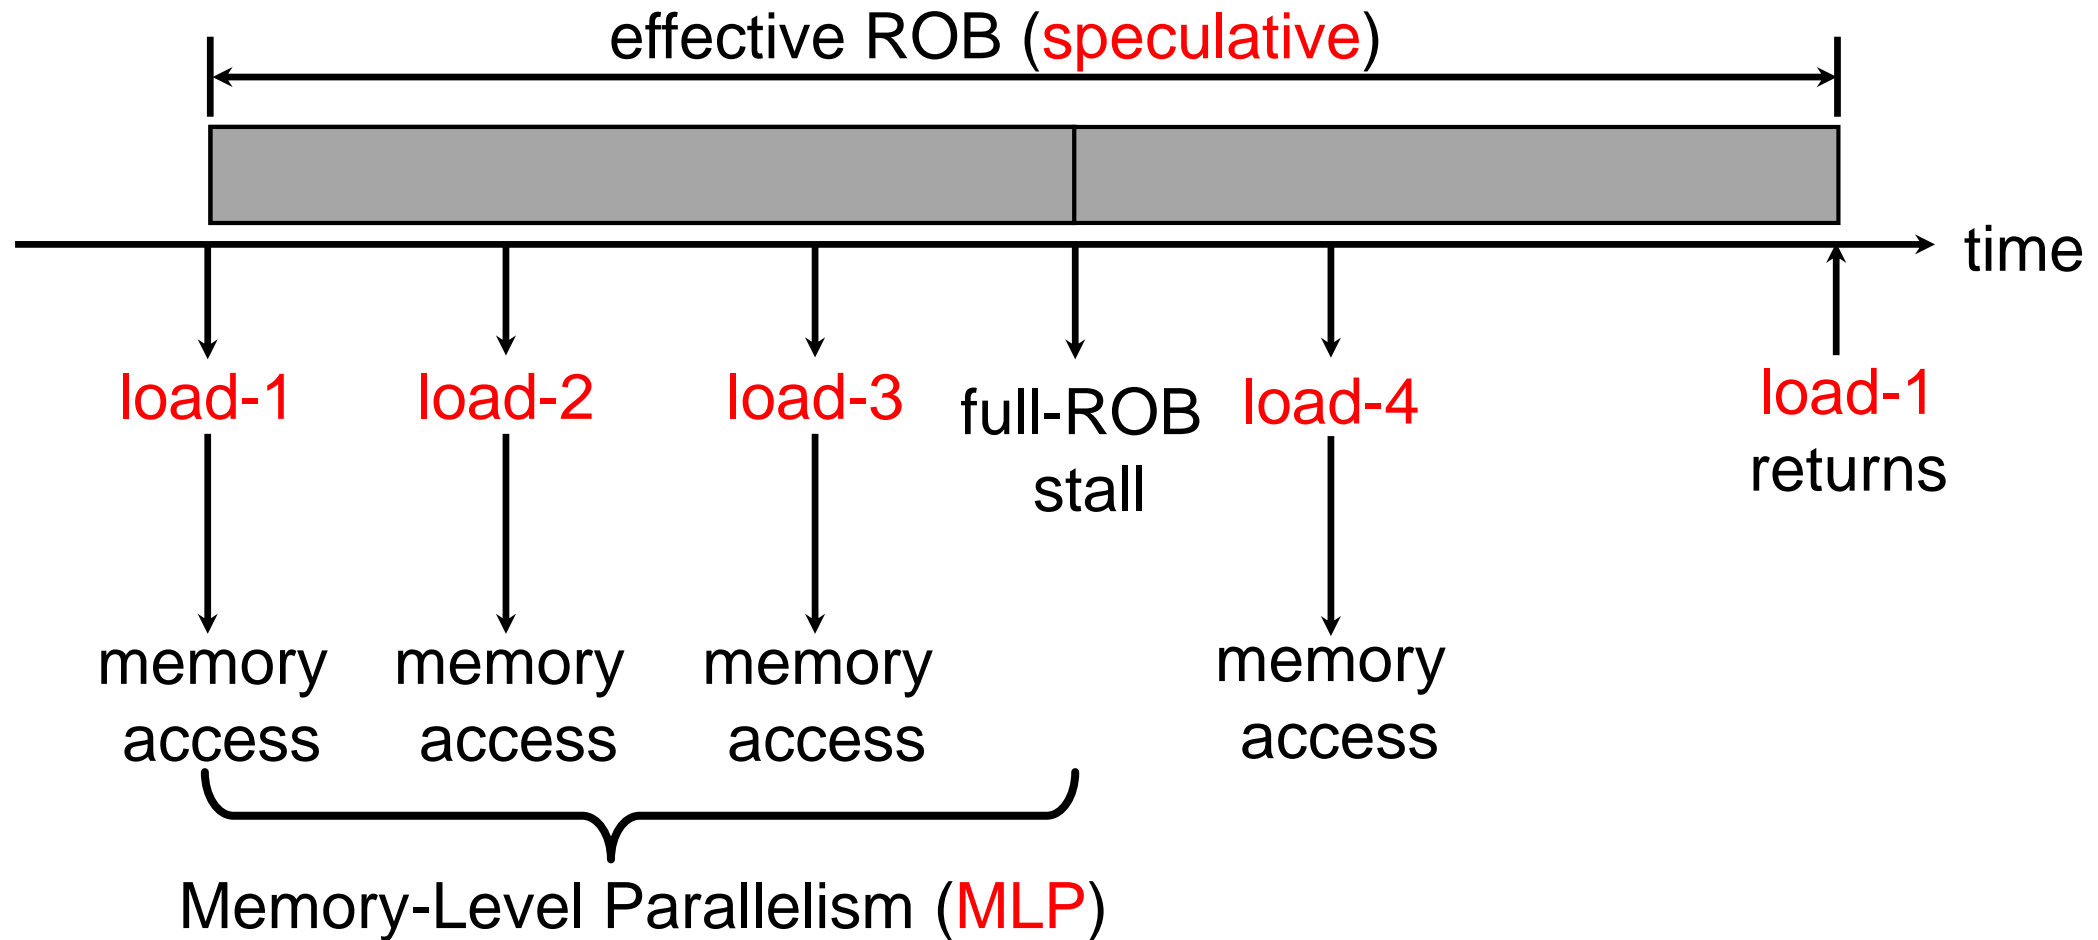
Runahead Execution Prefetches under a Full-ROB Stall



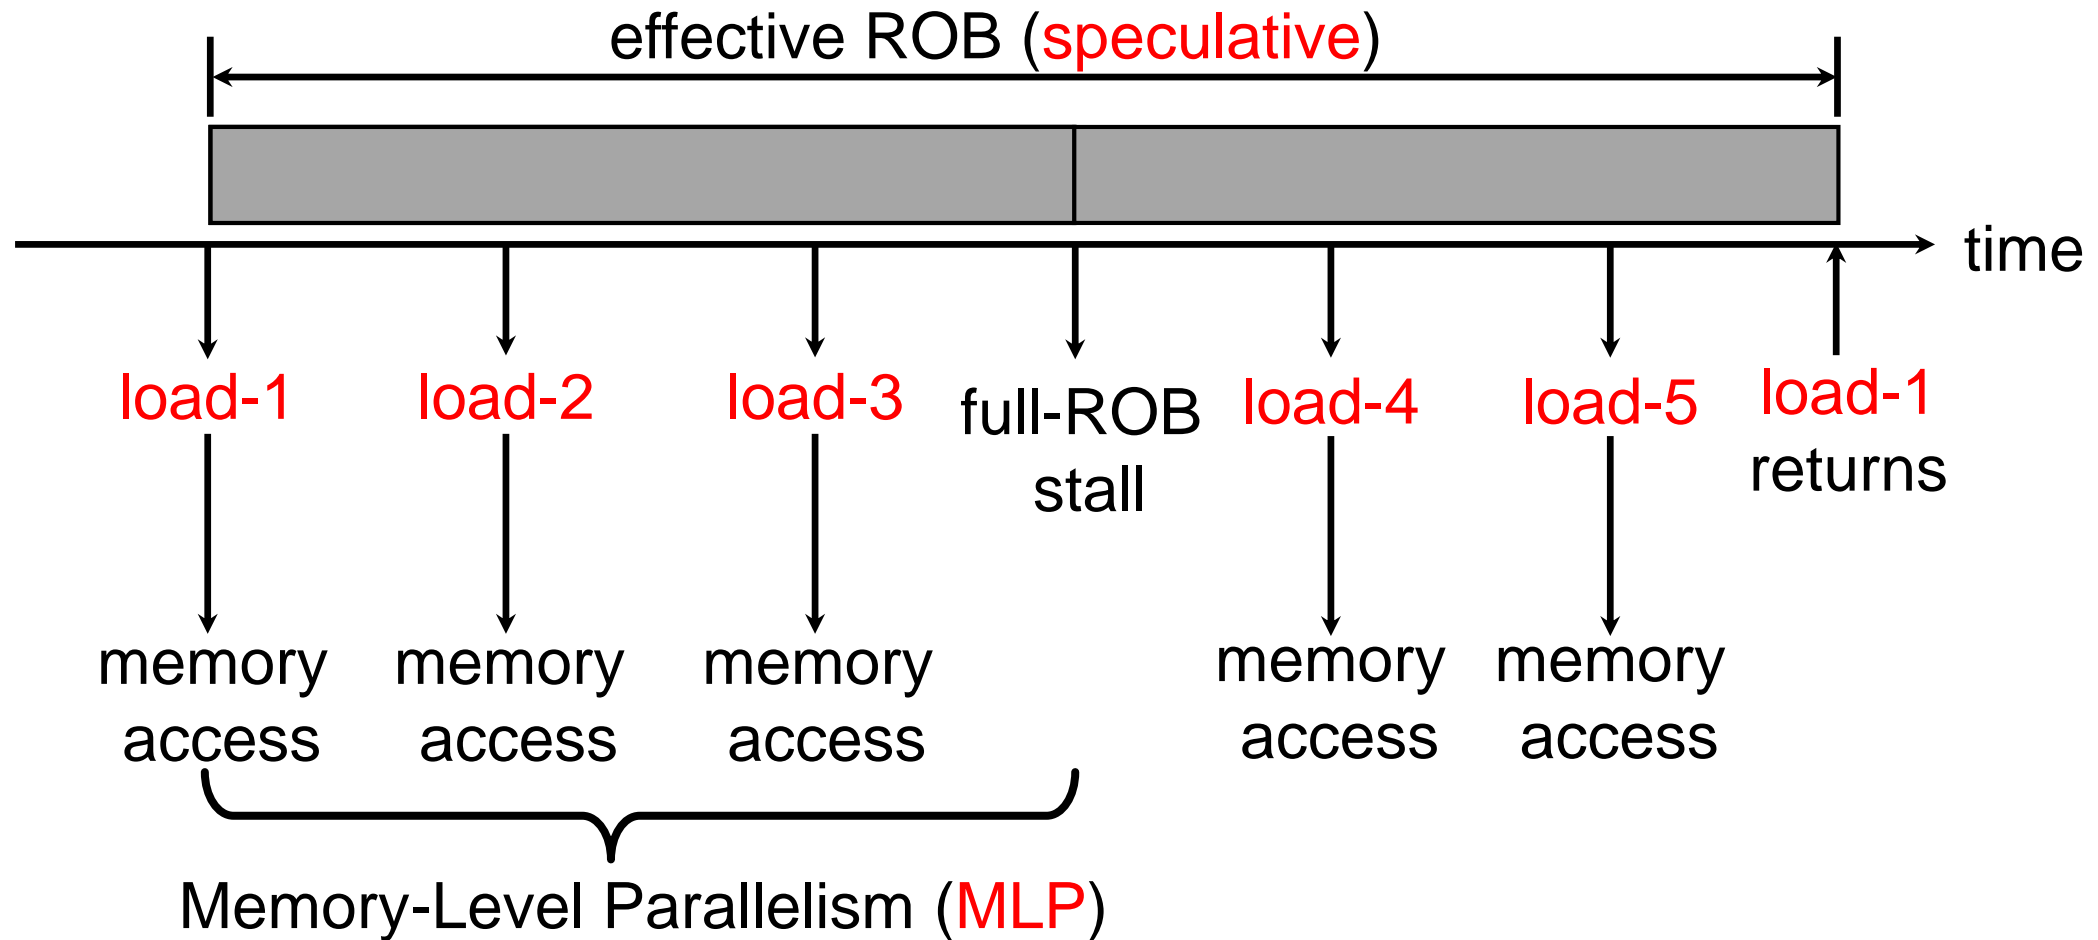
Runahead Execution Prefetches under a Full-ROB Stall



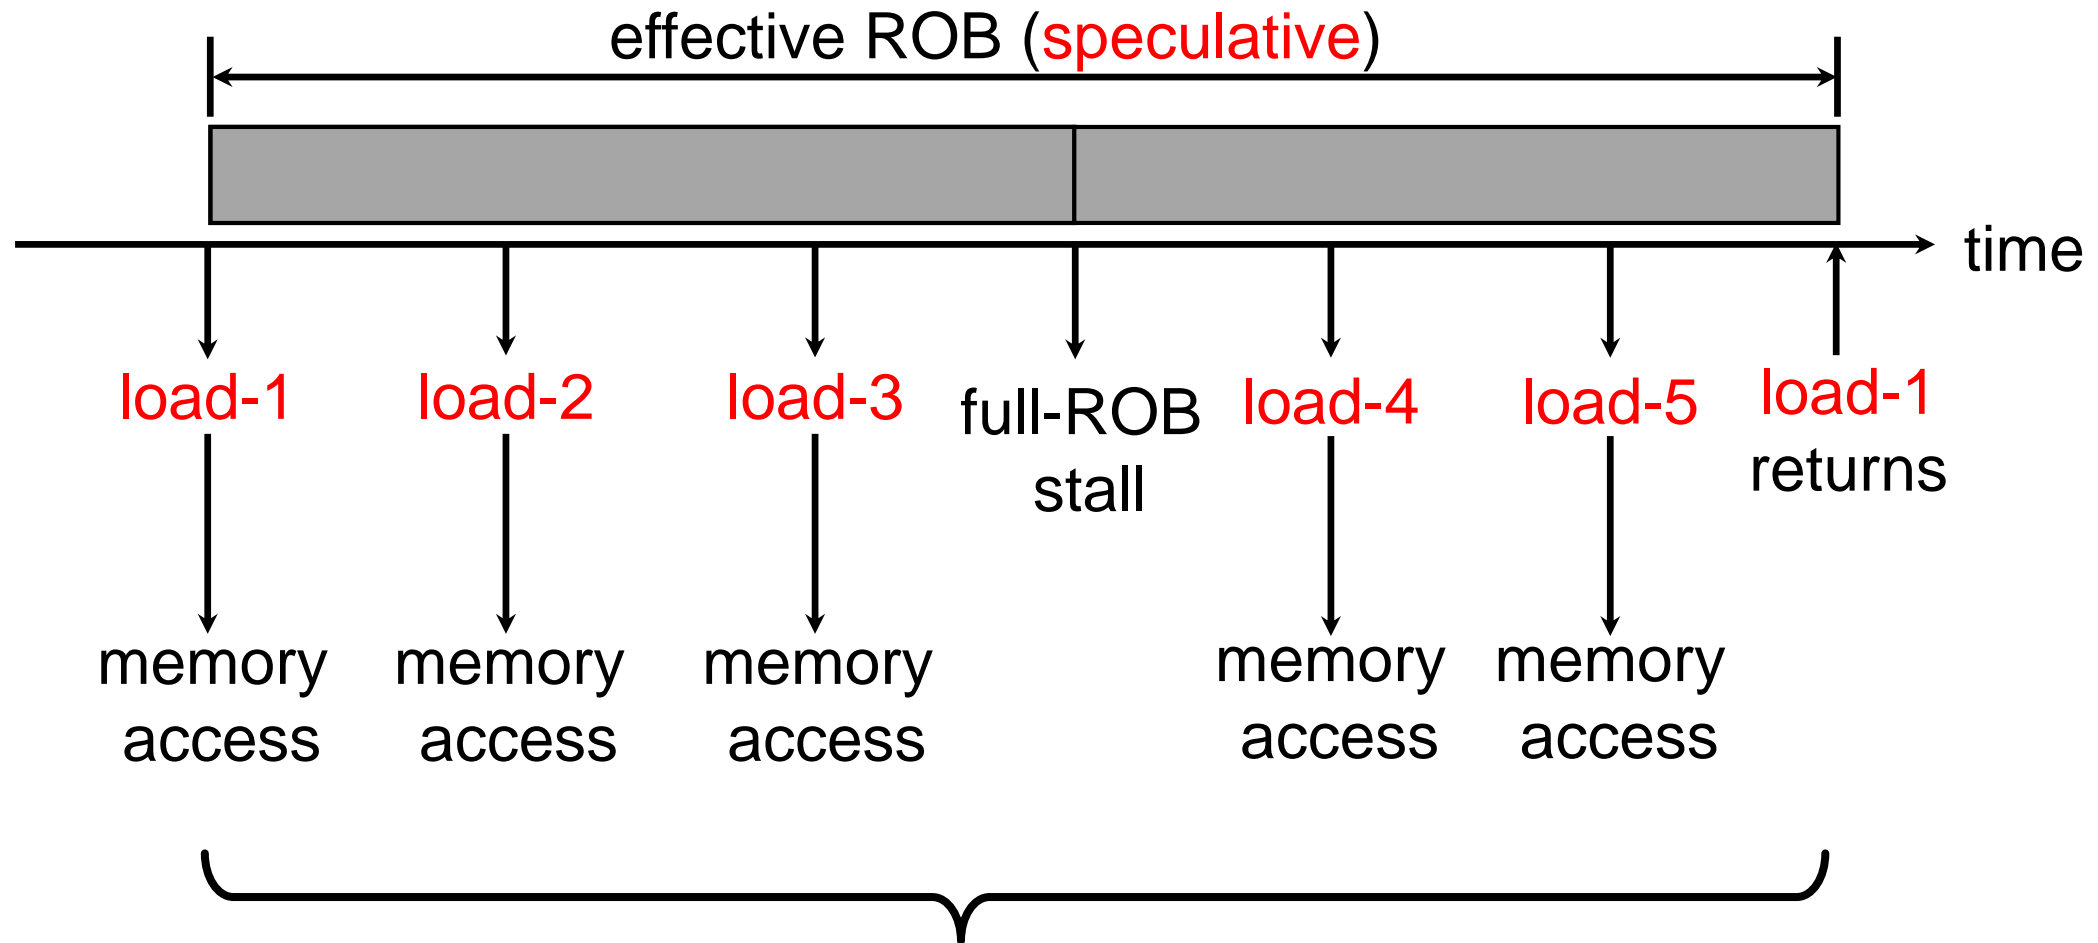
Runahead Execution Prefetches under a Full-ROB Stall



Runahead Execution Prefetches under a Full-ROB Stall



Runahead Execution Prefetches under a Full-ROB Stall



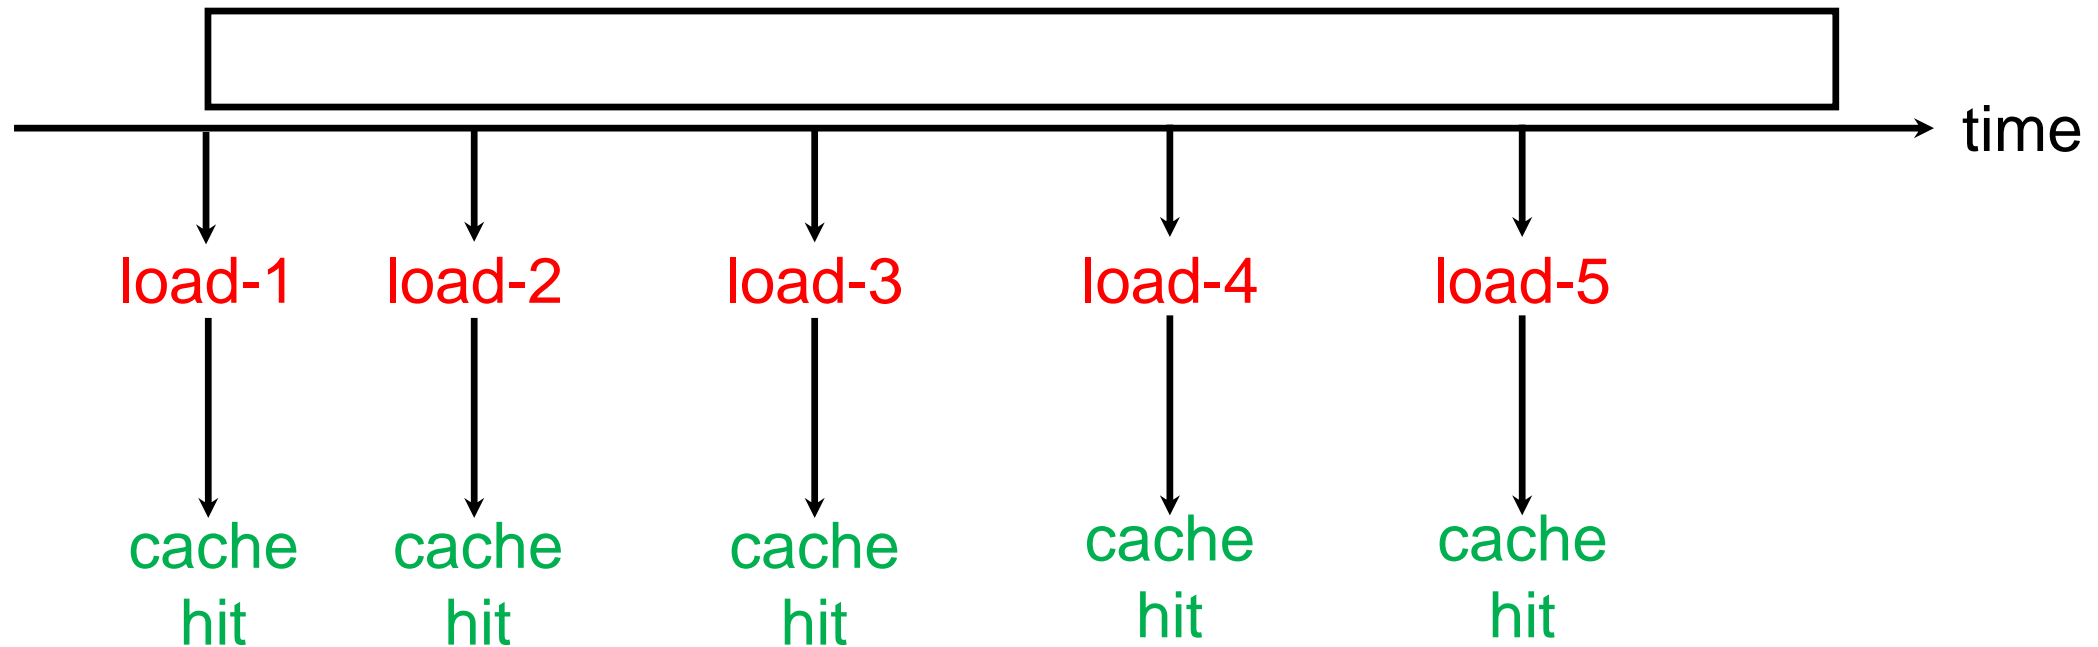
Increased Memory-Level Parallelism (**MLP**)

Runahead Execution Re-Executes Instructions

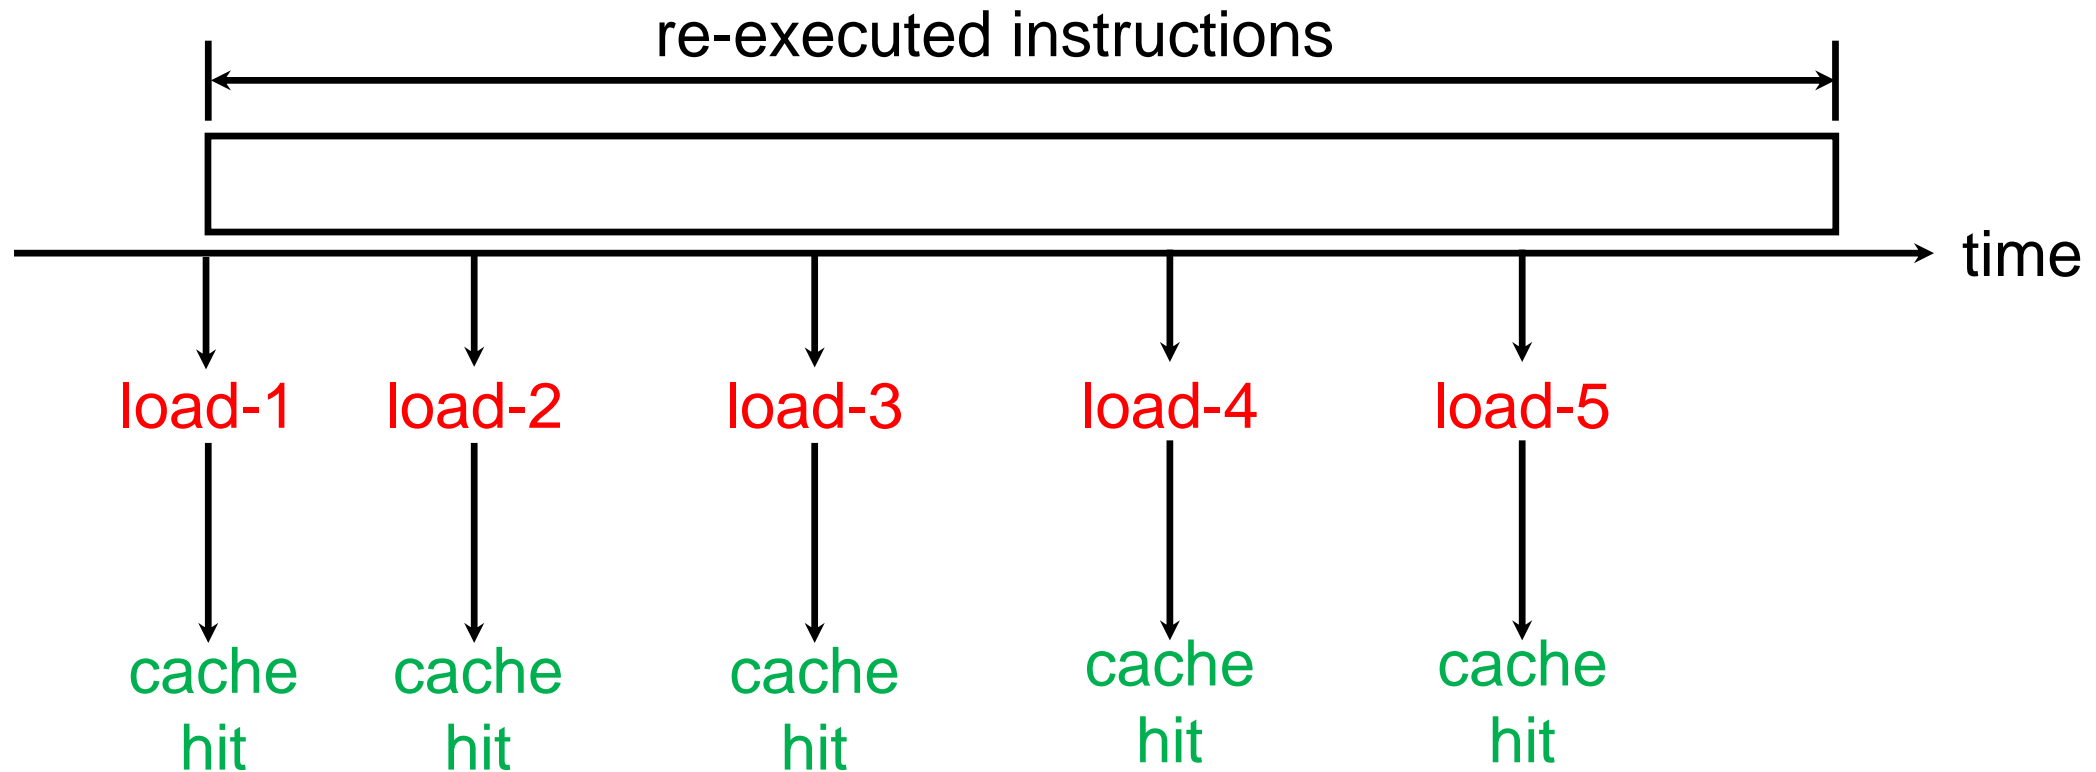
Runahead Execution Re-Executes Instructions

—————→ time

Runahead Execution Re-Executes Instructions



Runahead Execution Re-Executes Instructions



Runahead Buffer Finds the Most Critical Slice of Instructions

Runahead Buffer Finds the Most Critical Slice of Instructions

- Finds the **stalling slice** of instructions **in ROB**

Runahead Buffer Finds the Most Critical Slice of Instructions

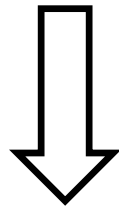
- Finds the **stalling slice** of instructions **in ROB**
- Executes **only** this slice after a full-ROB stall

Runahead Buffer Finds the Most Critical Slice of Instructions

- Finds the **stalling slice** of instructions **in ROB**
- Executes **only** this slice after a full-ROB stall
- Clock-gates the front-end for saving power

Runahead Buffer Finds the Most Critical Slice of Instructions

- Finds the **stalling slice** of instructions **in ROB**
- Executes **only** this slice after a full-ROB stall
- Clock-gates the front-end for saving power



**Better energy-efficiency and
similar performance**

Runahead Techniques Relative to OoO Core

Runahead Techniques Relative to OoO Core

Runahead
execution

Runahead
buffer

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB			

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓		

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	
Execute all future instructions			

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	
Execute all future instructions	✓		

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	
Execute all future instructions	✓	Only one slice	

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	
Execute all future instructions	✓	Only one slice	
Performance			

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	
Execute all future instructions	✓	Only one slice	
Performance	High ↑		

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	
Execute all future instructions	✓	Only one slice	
Performance	High ↑	High ↑	

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	
Execute all future instructions	✓	Only one slice	
Performance	High ↑	High ↑	
Energy Efficiency			

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	
Execute all future instructions	✓	Only one slice	
Performance	High ↑	High ↑	
Energy Efficiency	Low ↓		

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	
Execute all future instructions	✓	Only one slice	
Performance	High ↑	High ↑	
Energy Efficiency	Low ↓	Same	

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	✗
Execute all future instructions	✓	Only one slice	
Performance	High ↑	High ↑	
Energy Efficiency	Low ↓	Same	

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	✗
Execute all future instructions	✓	Only one slice	All slices
Performance	High ↑	High ↑	
Energy Efficiency	Low ↓	Same	

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	✗
Execute all future instructions	✓	Only one slice	All slices
Performance	High ↑	High ↑	Very high ↑↑
Energy Efficiency	Low ↓	Same	

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	
Flush ROB	✓	✓	✗
Execute all future instructions	✓	Only one slice	All slices
Performance	High ↑	High ↑	Very high ↑↑
Energy Efficiency	Low ↓	Same	High ↑

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	Precise runahead
Flush ROB	✓	✓	✗
Execute all future instructions	✓	Only one slice	All slices
Performance	High ↑	High ↑	Very high ↑↑
Energy Efficiency	Low ↓	Same	High ↑

Precise Runahead Exection: Three Key Insights

Precise Runahead Exection: Three Key Insights

1. No need to flush the ROB

Precise Runahead Exection: Three Key Insights

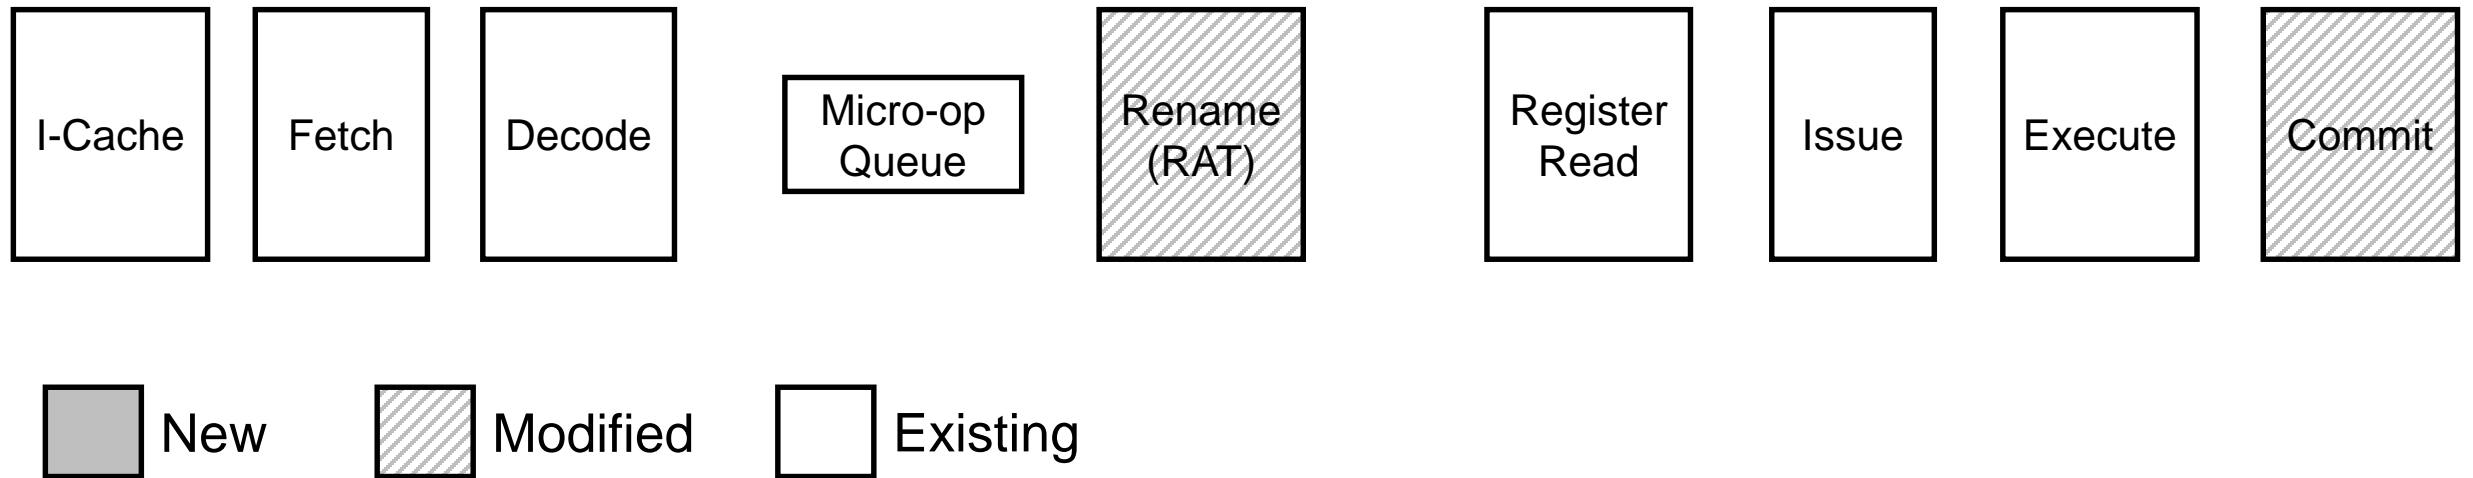
1. No need to flush the ROB
2. Slices execute quickly

Precise Runahead Exection: Three Key Insights

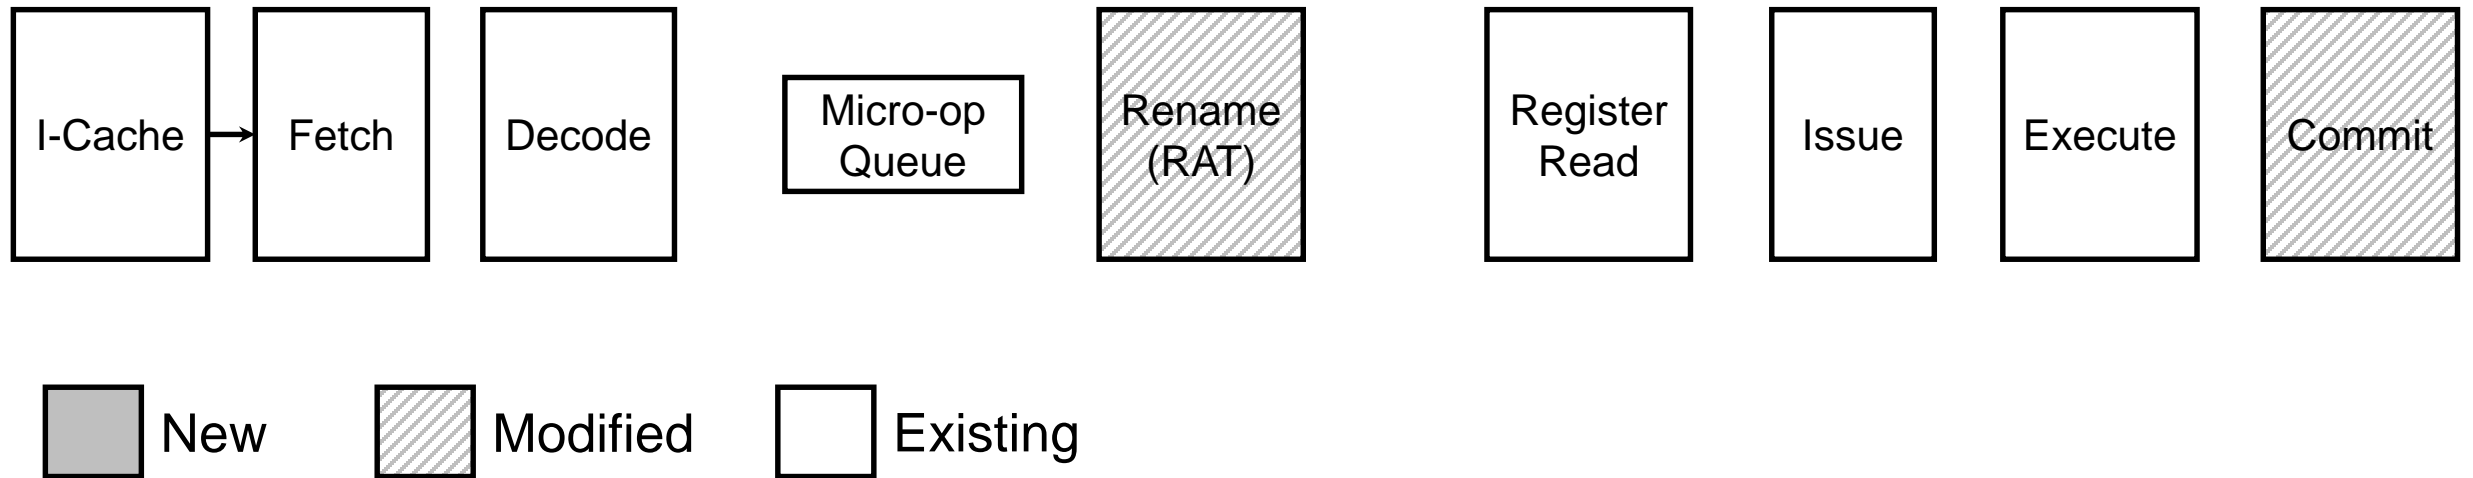
1. No need to flush the ROB
2. Slices execute quickly
3. Only physical registers must be recycled

Precise Runahead Execution

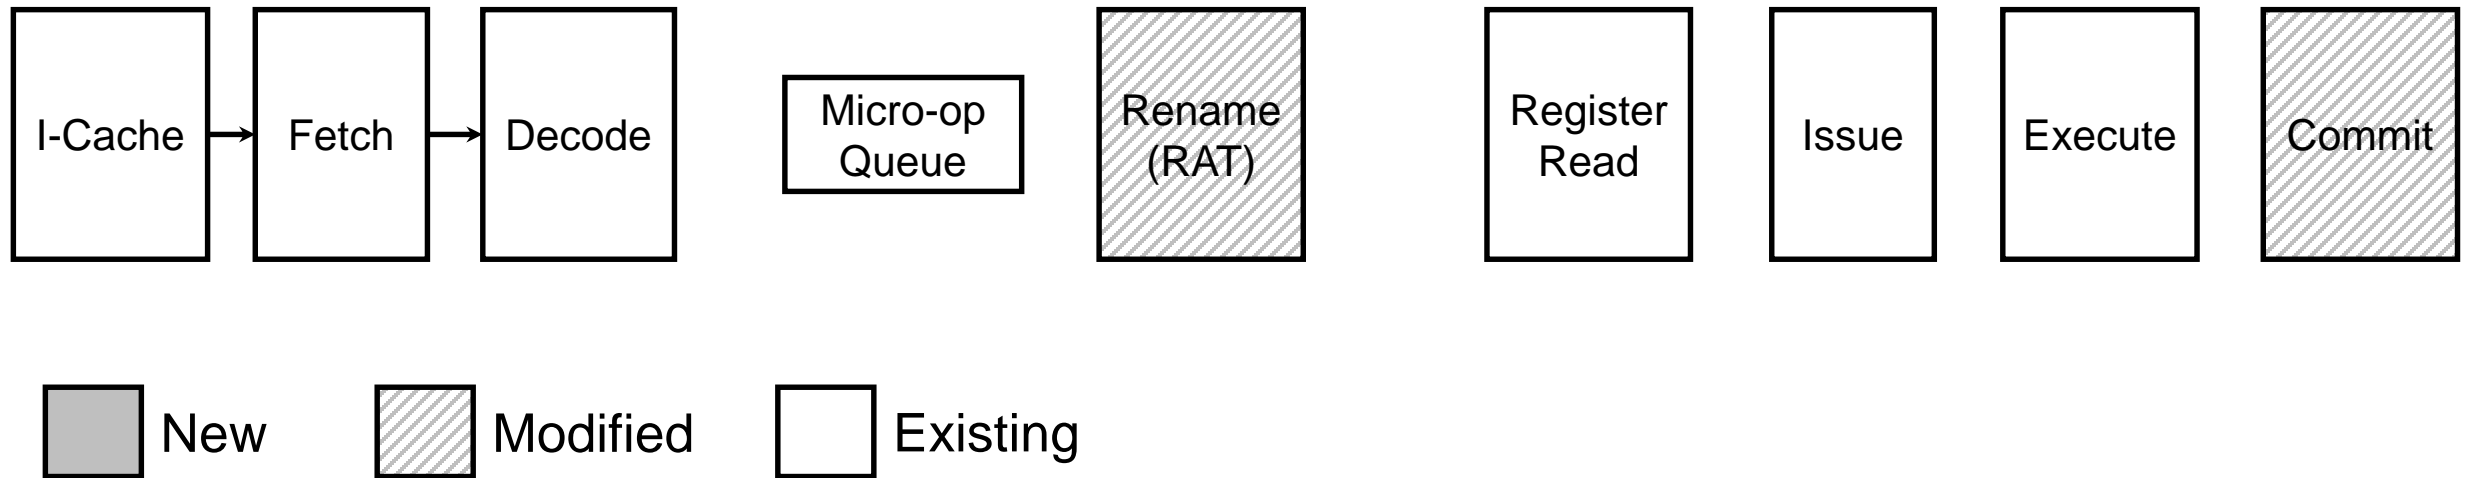
Precise Runahead Execution



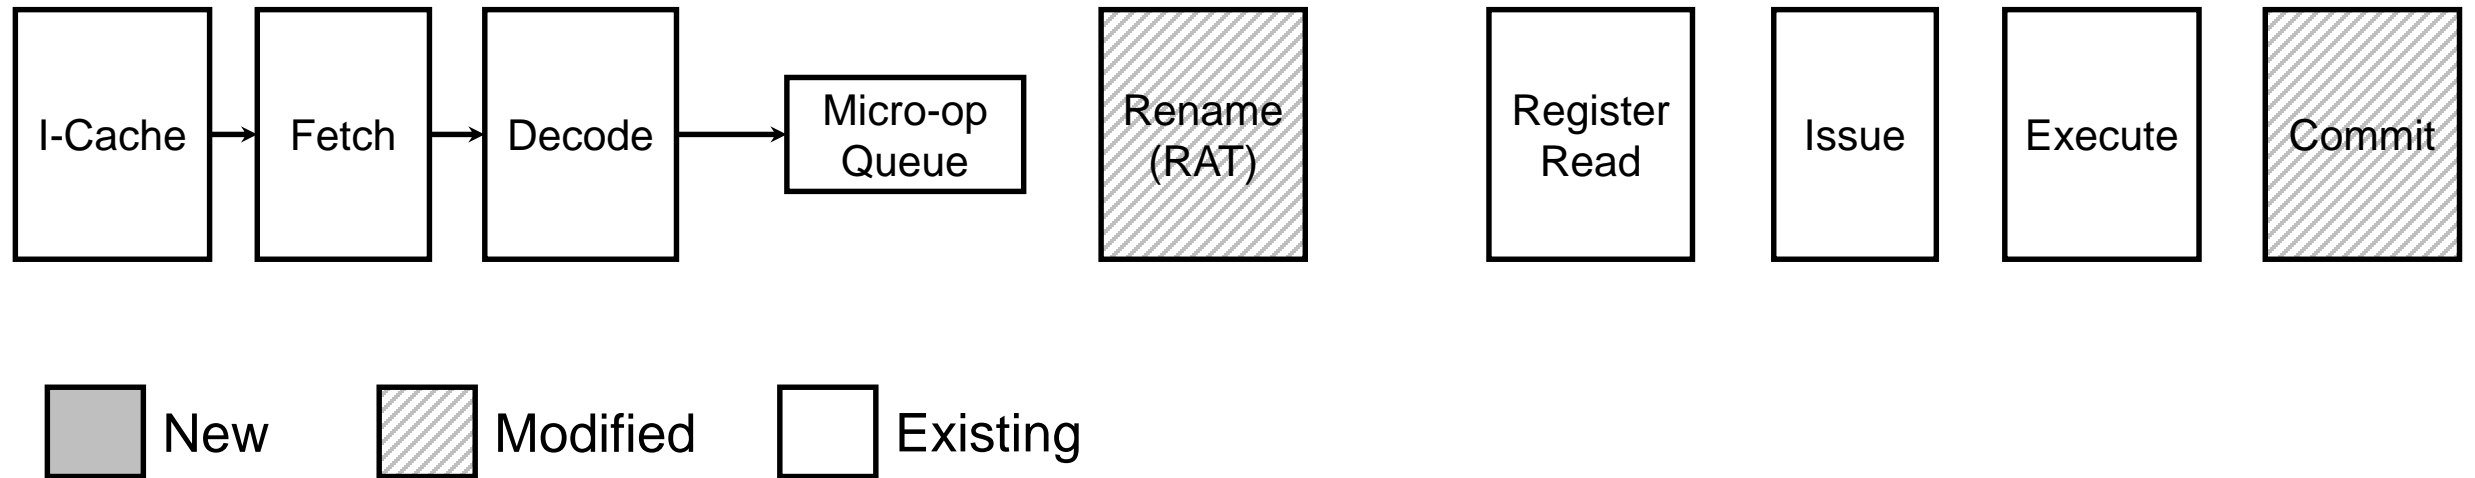
Precise Runahead Execution



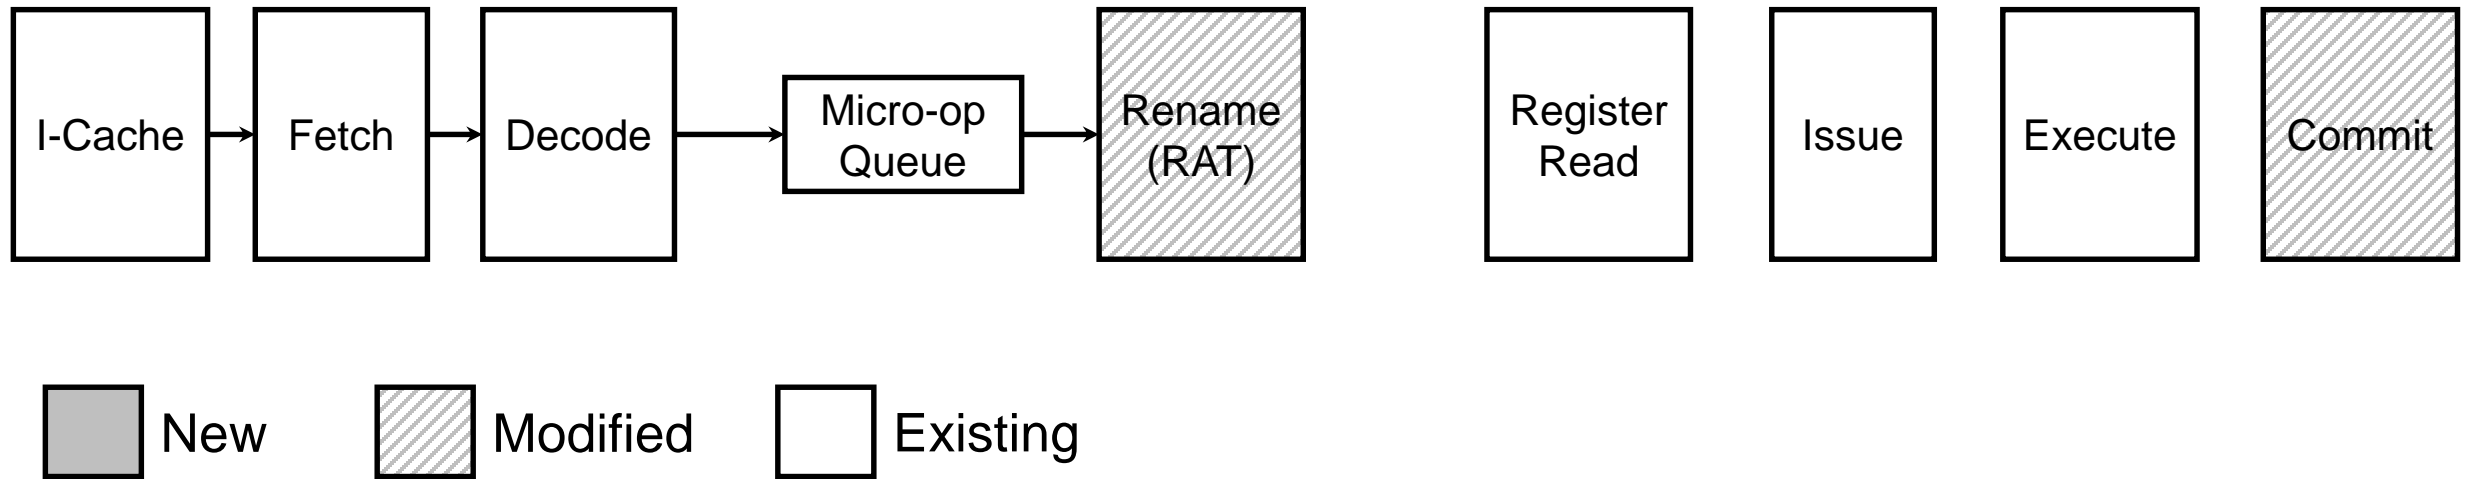
Precise Runahead Execution



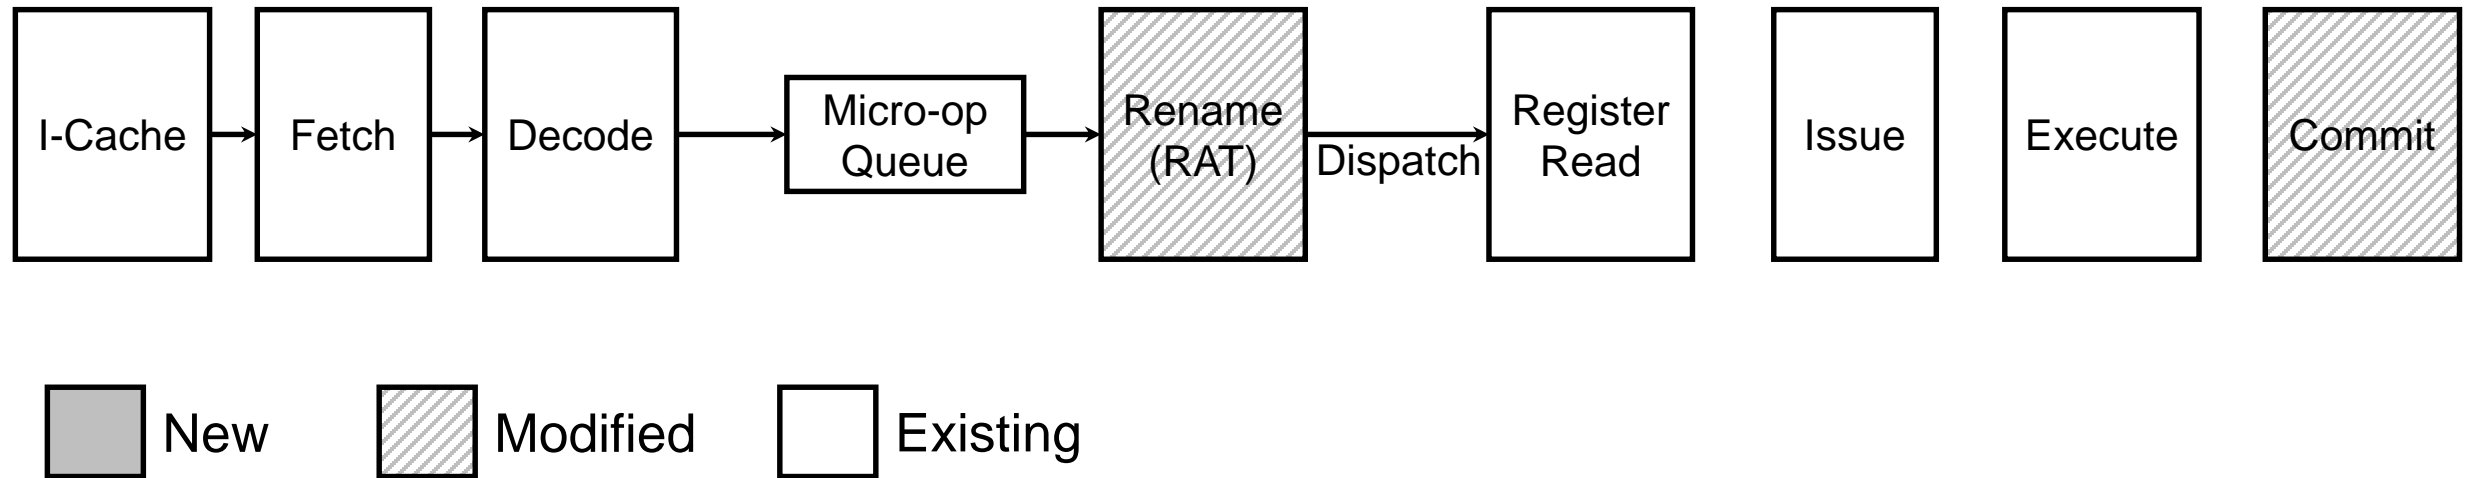
Precise Runahead Execution



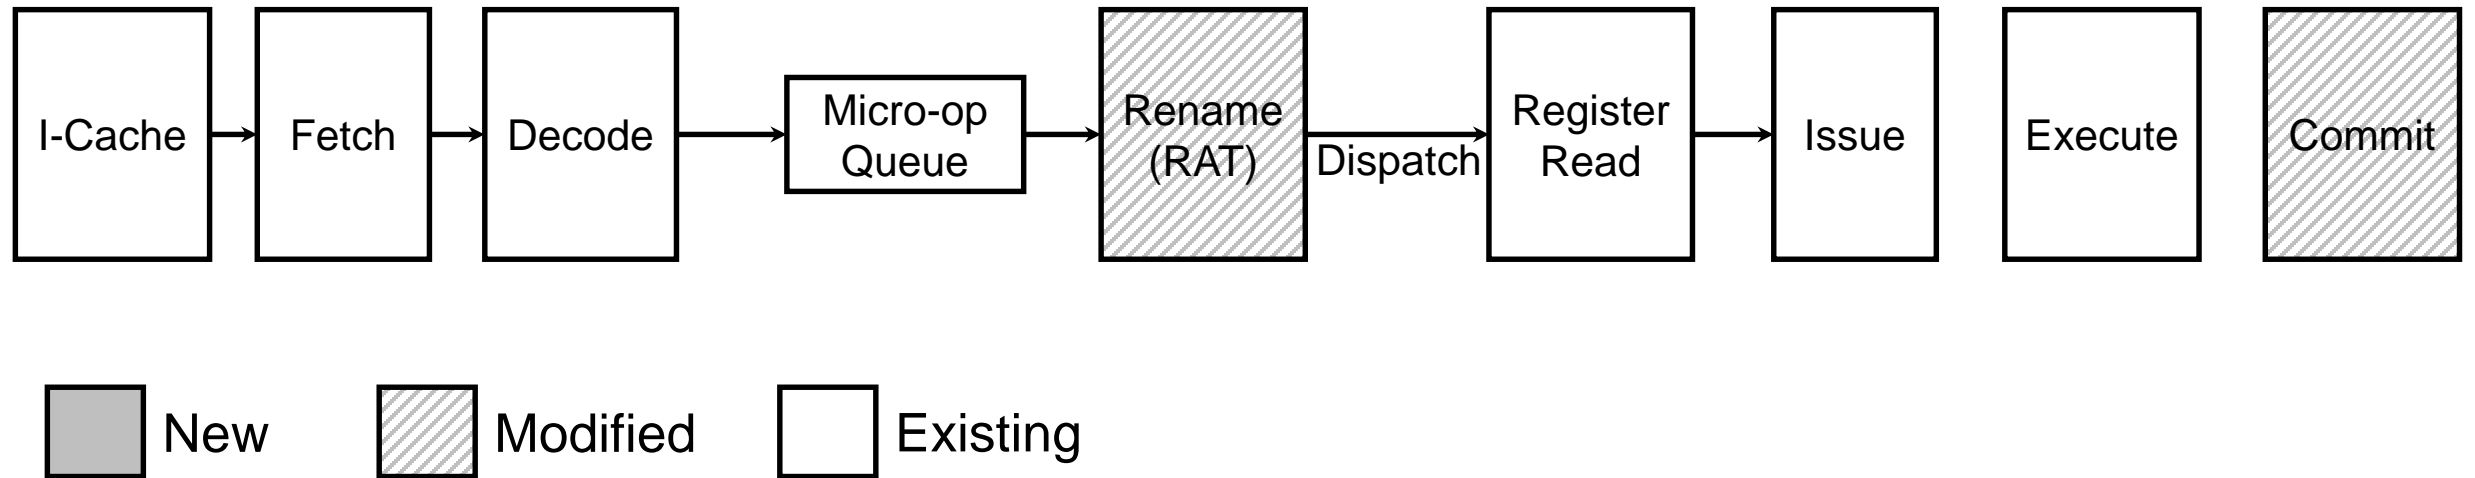
Precise Runahead Execution



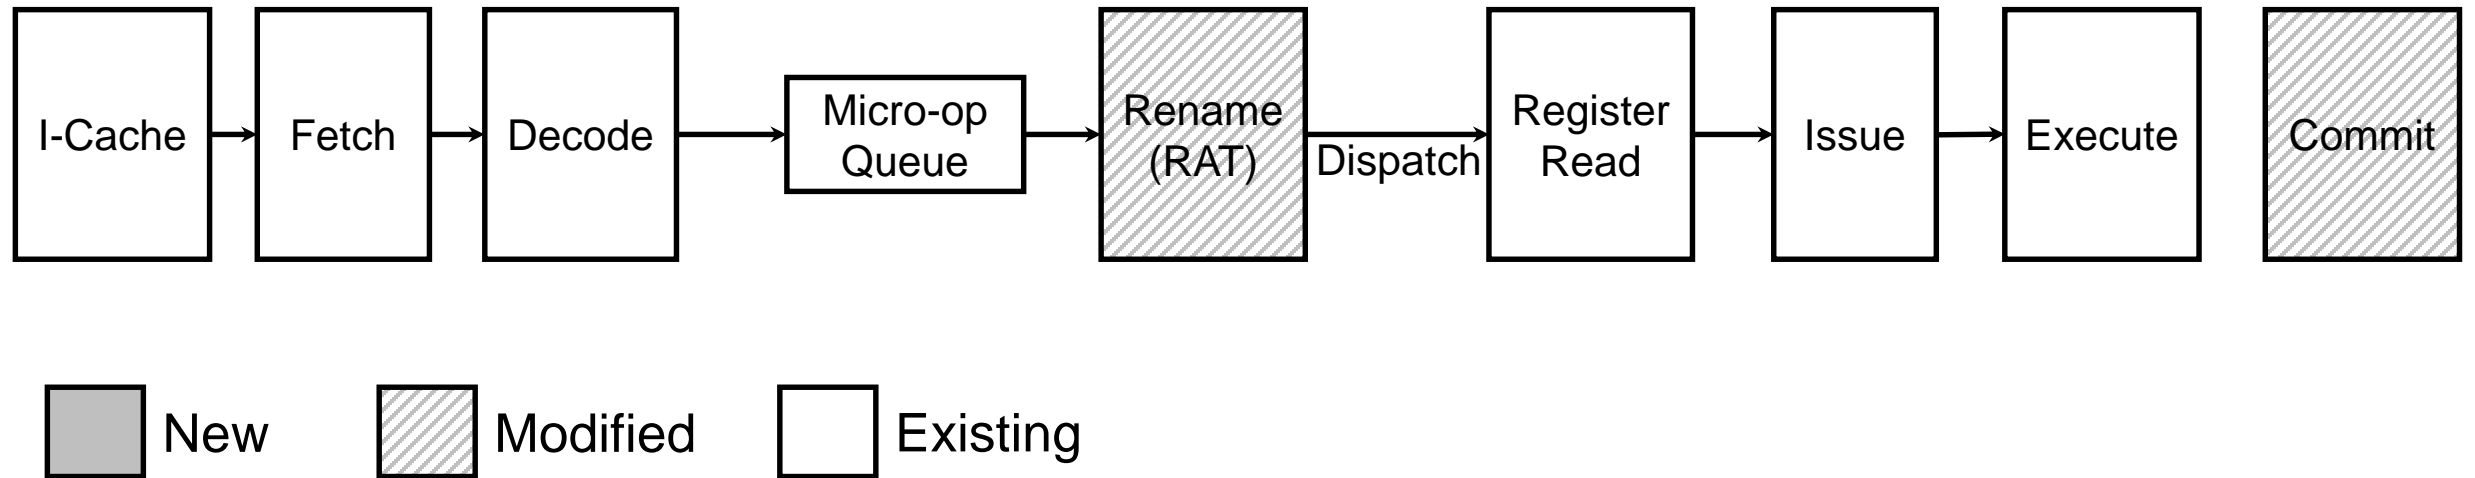
Precise Runahead Execution



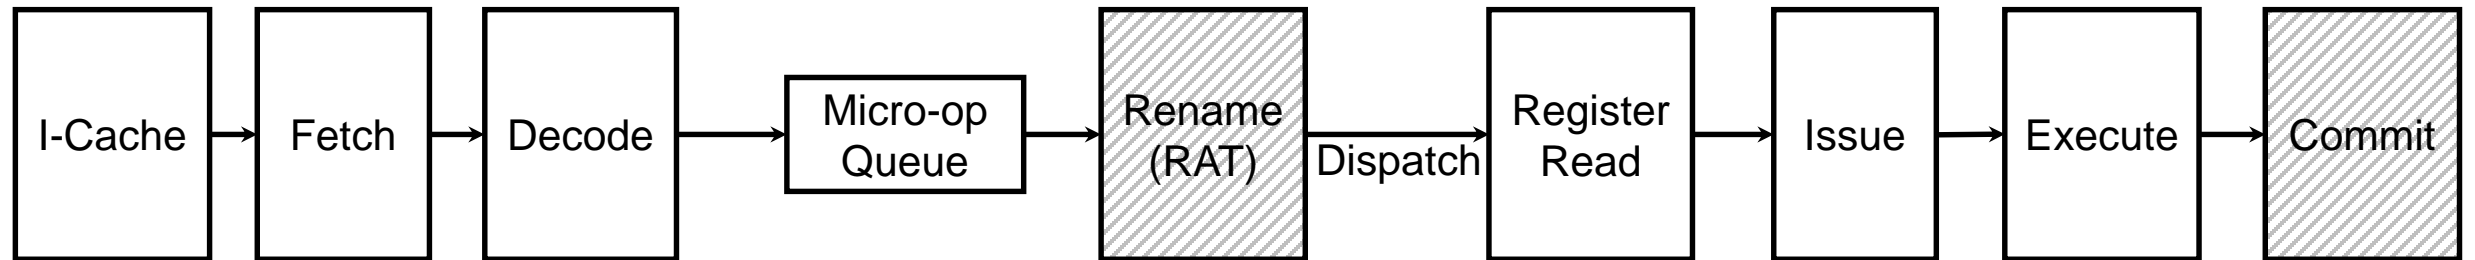
Precise Runahead Execution



Precise Runahead Execution



Precise Runahead Execution



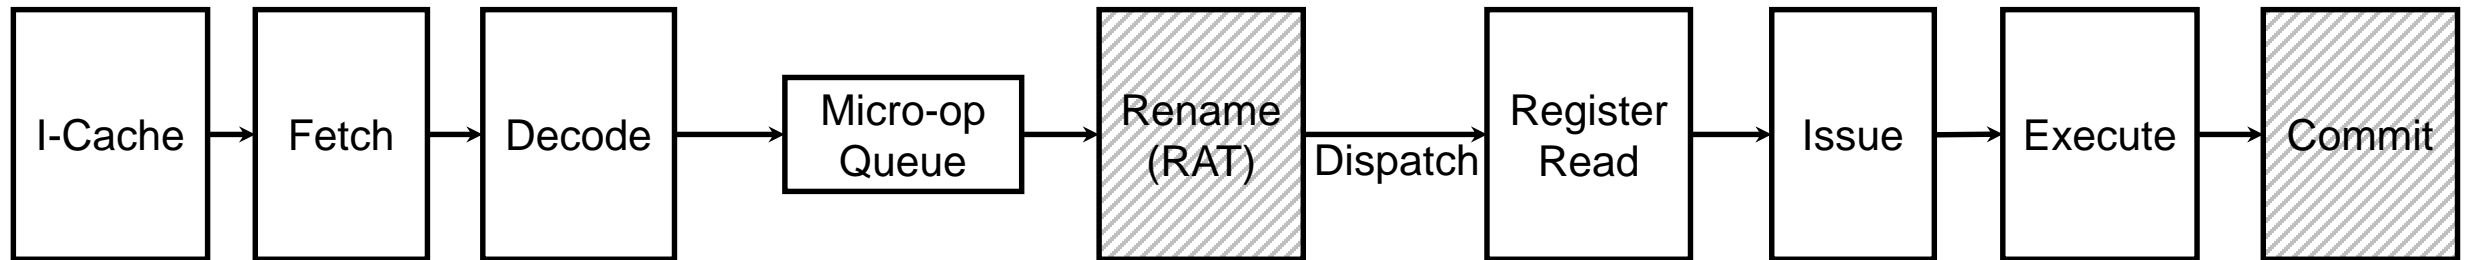
 New  Modified  Existing

→ Normal Mode

Precise Runahead Execution

Stalling
Slice
Table

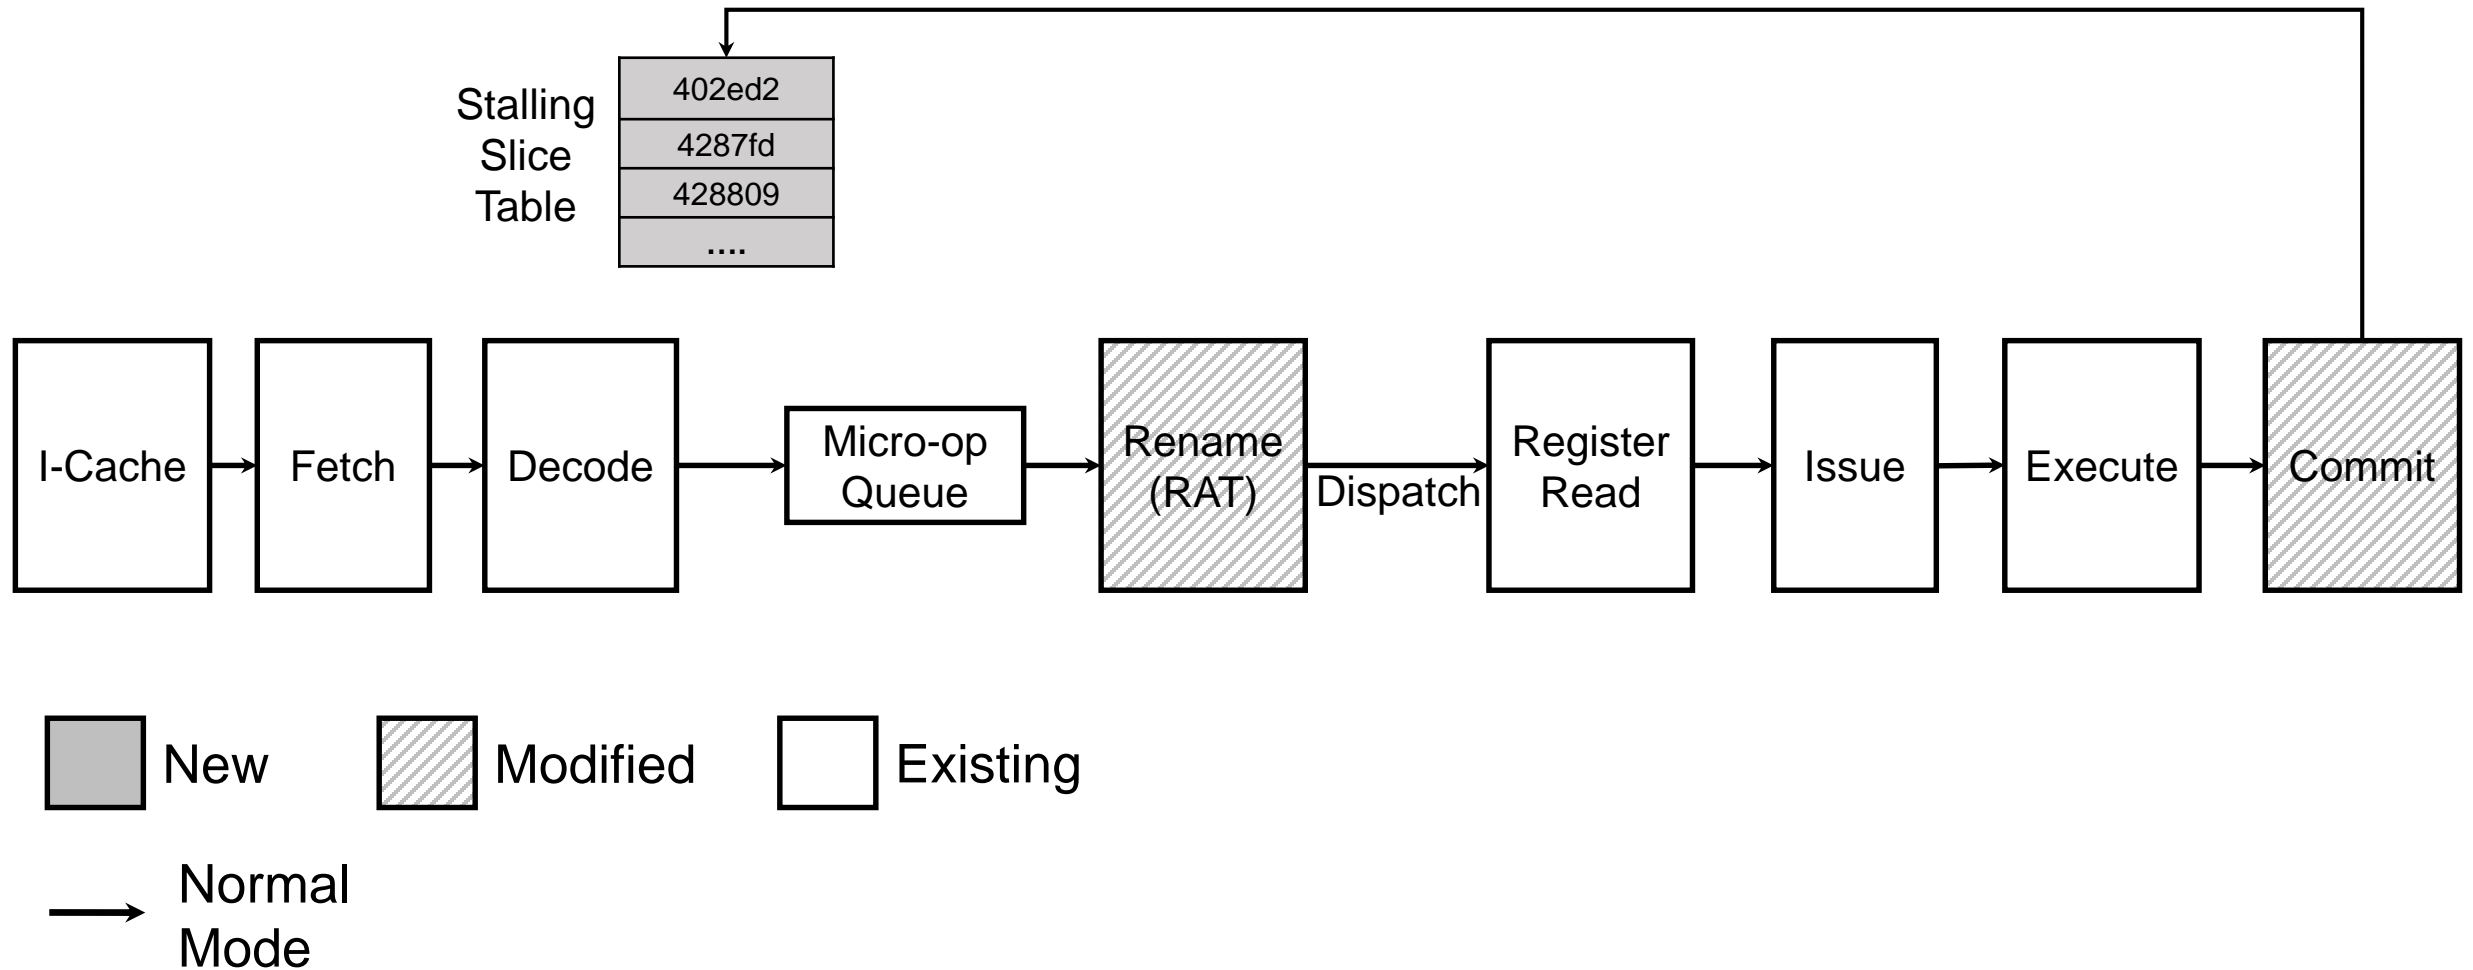
402ed2
4287fd
428809
....



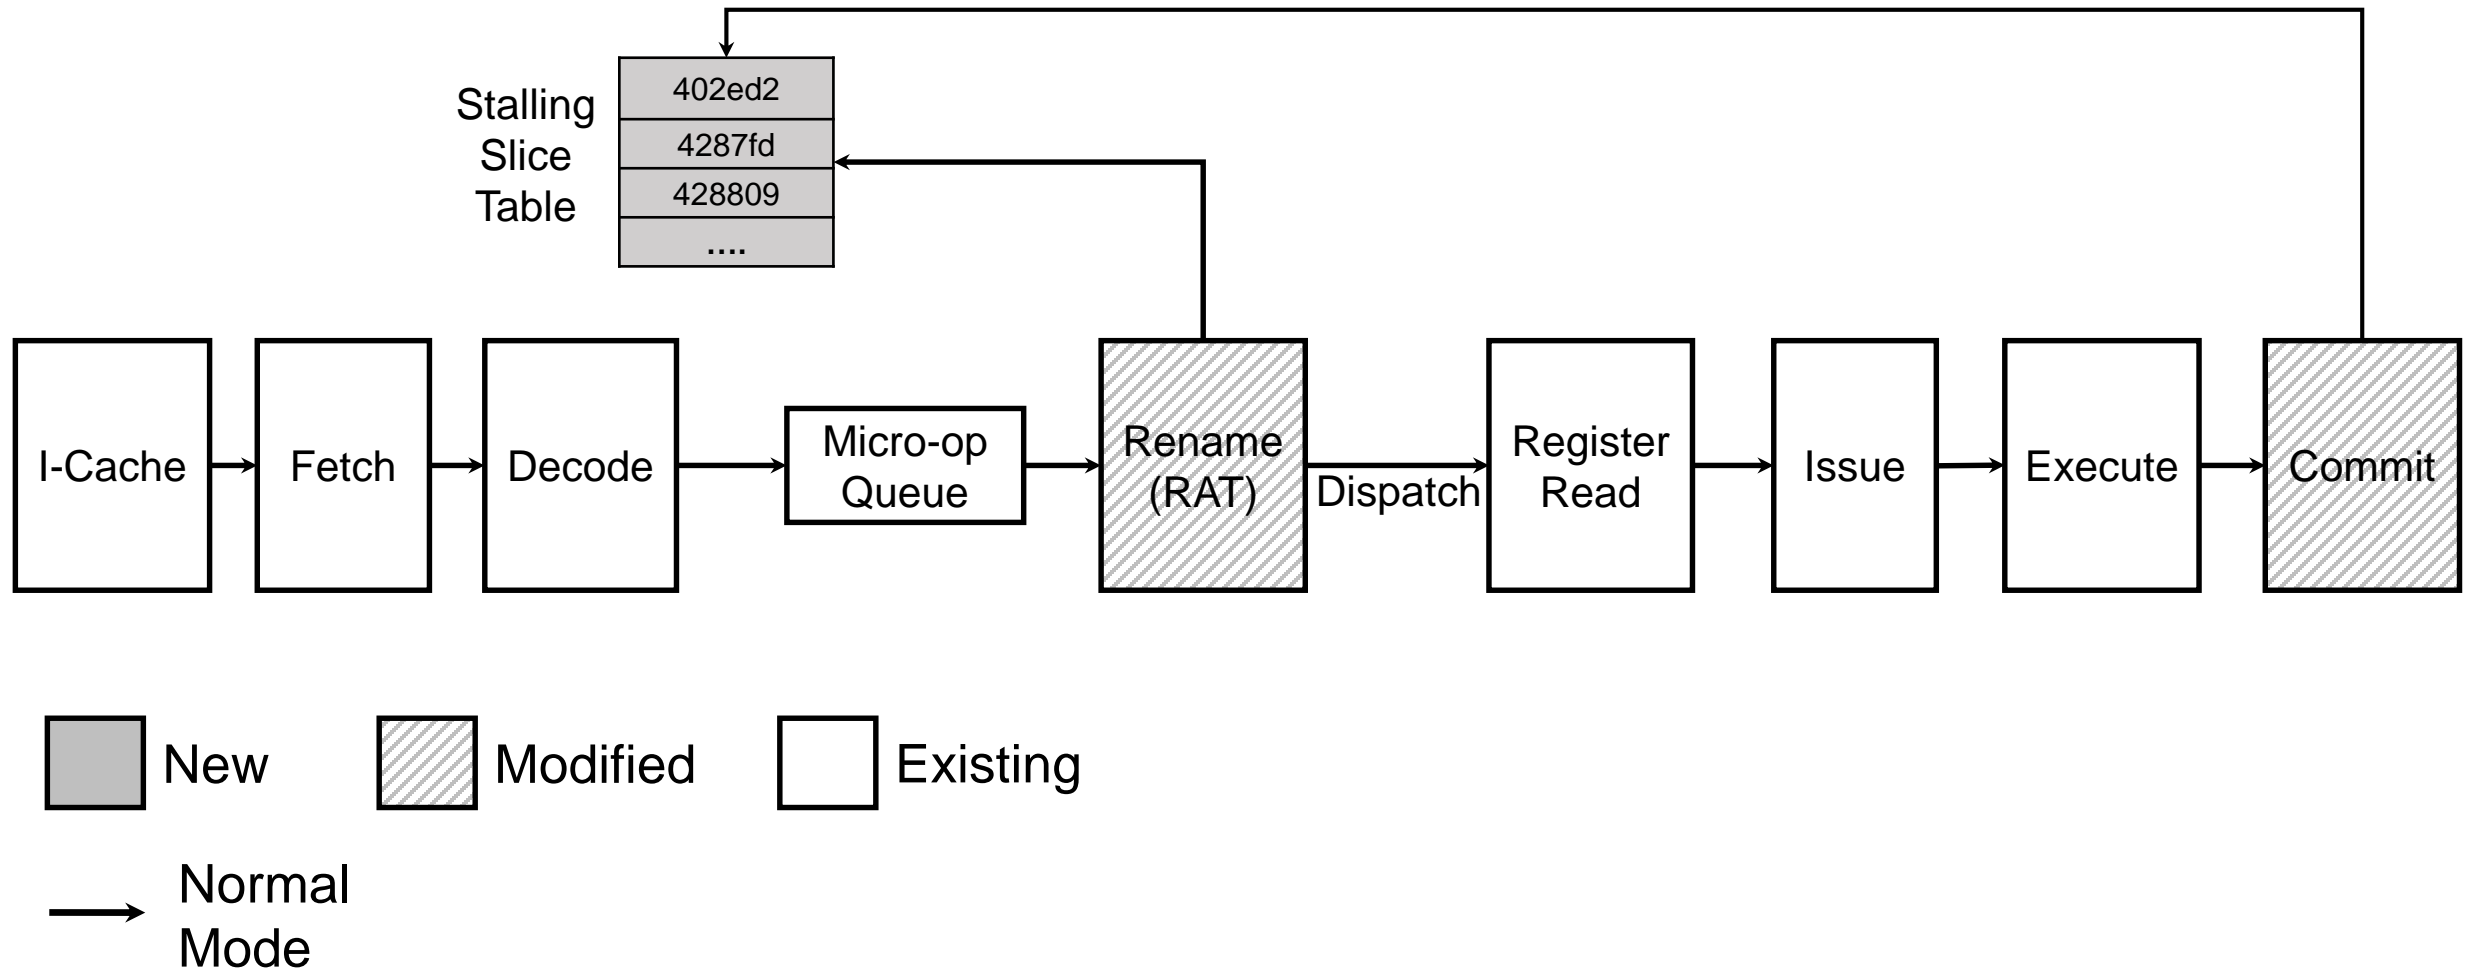
 New  Modified  Existing

→ Normal Mode

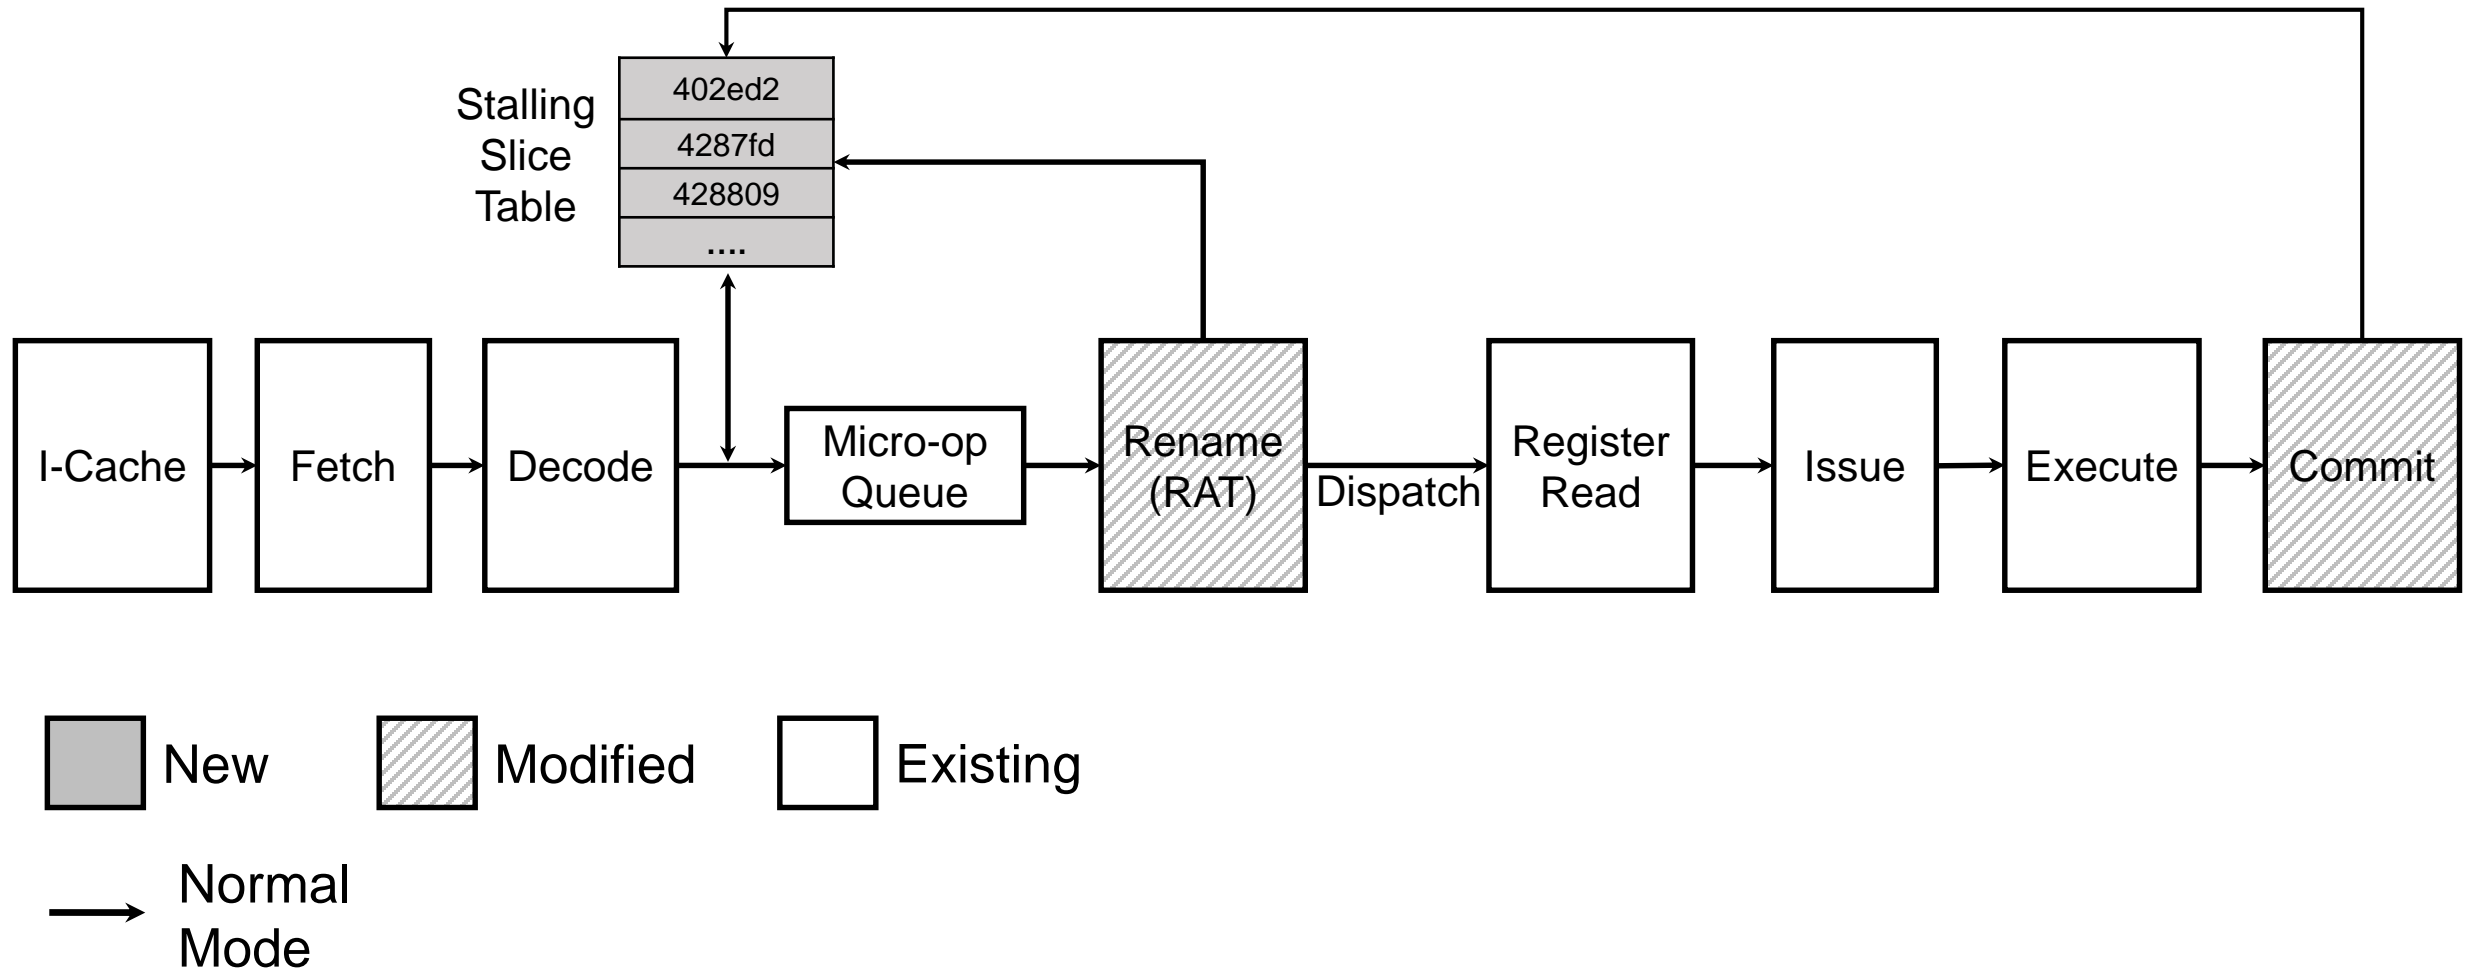
Precise Runahead Execution



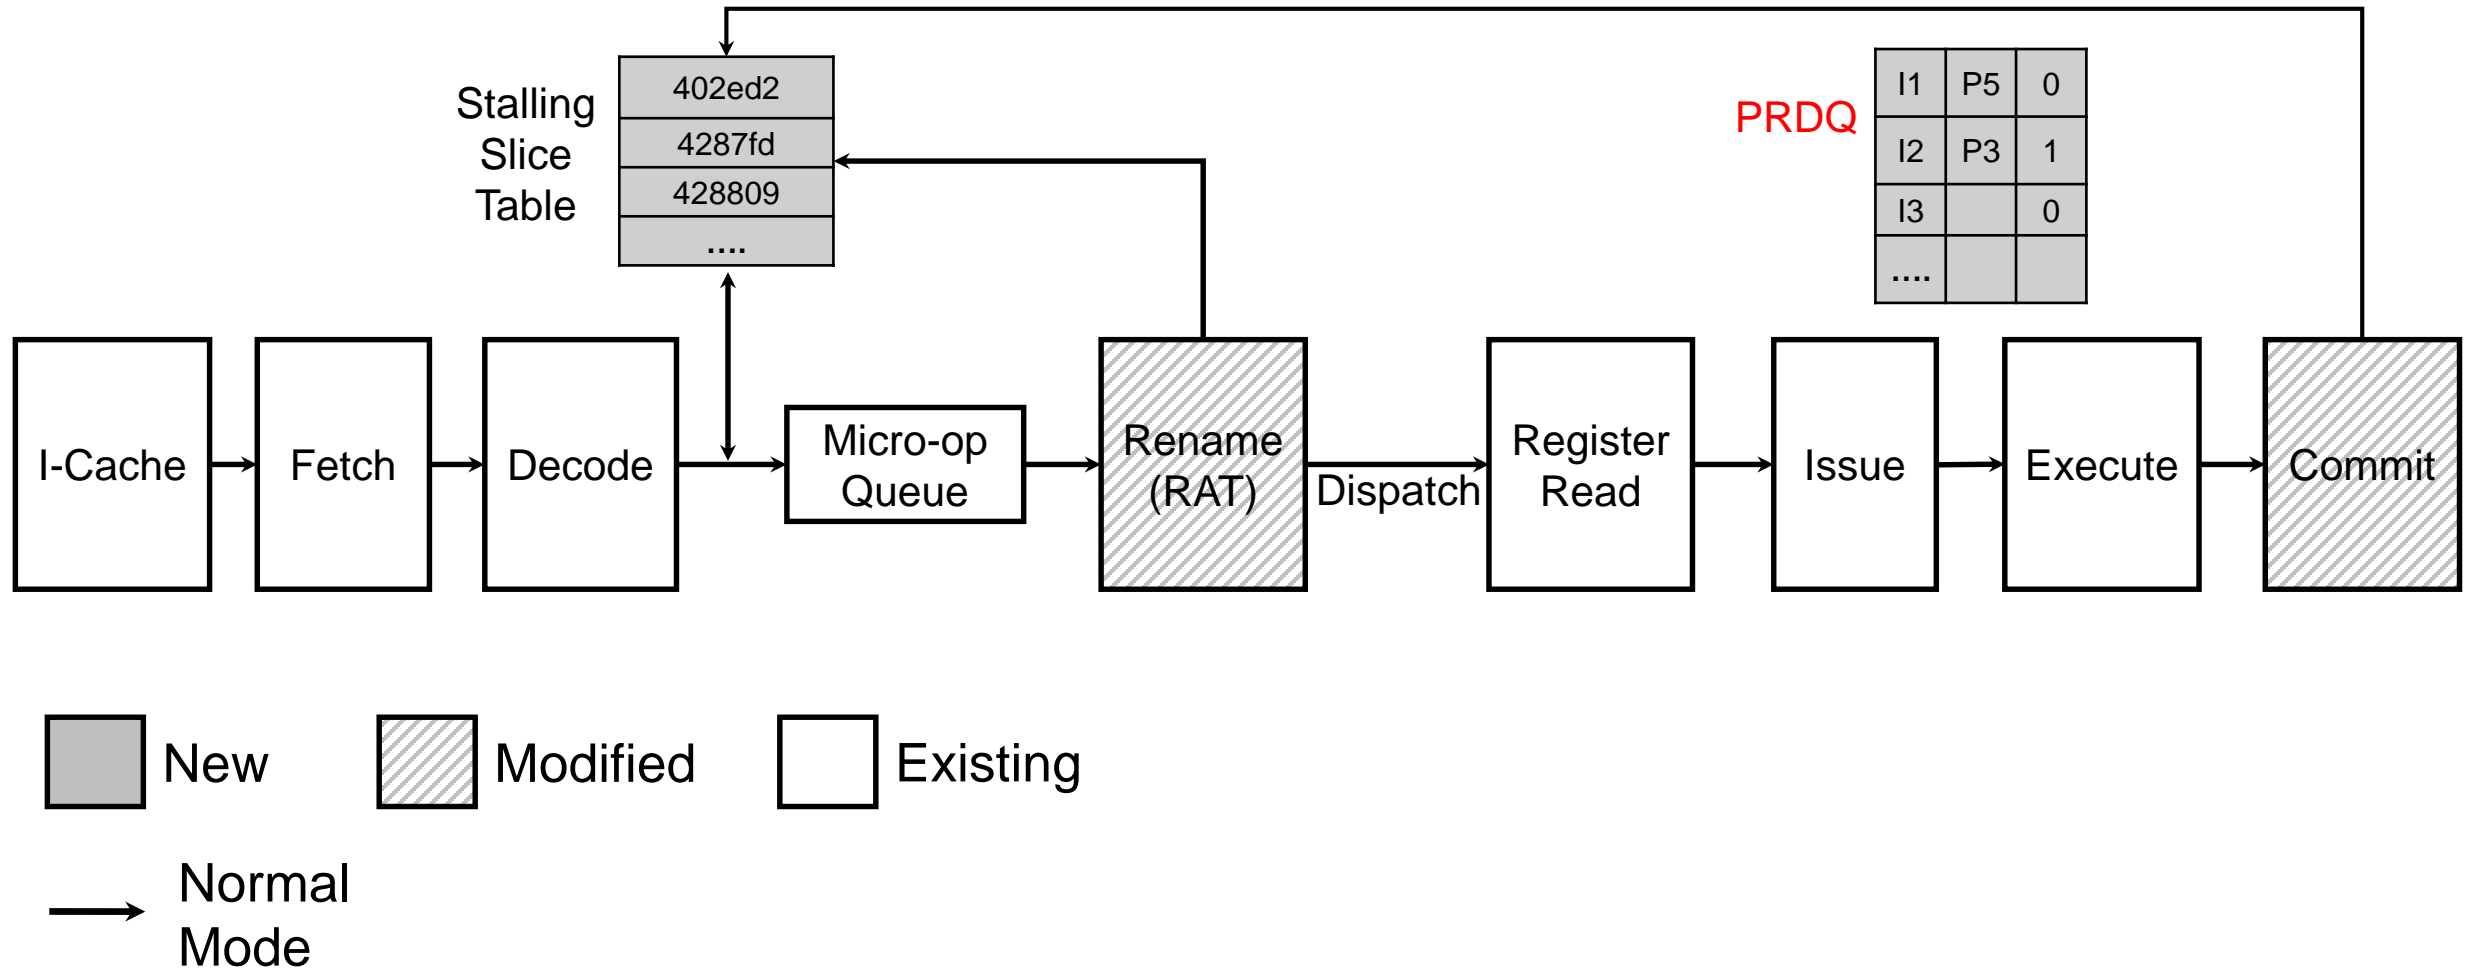
Precise Runahead Execution



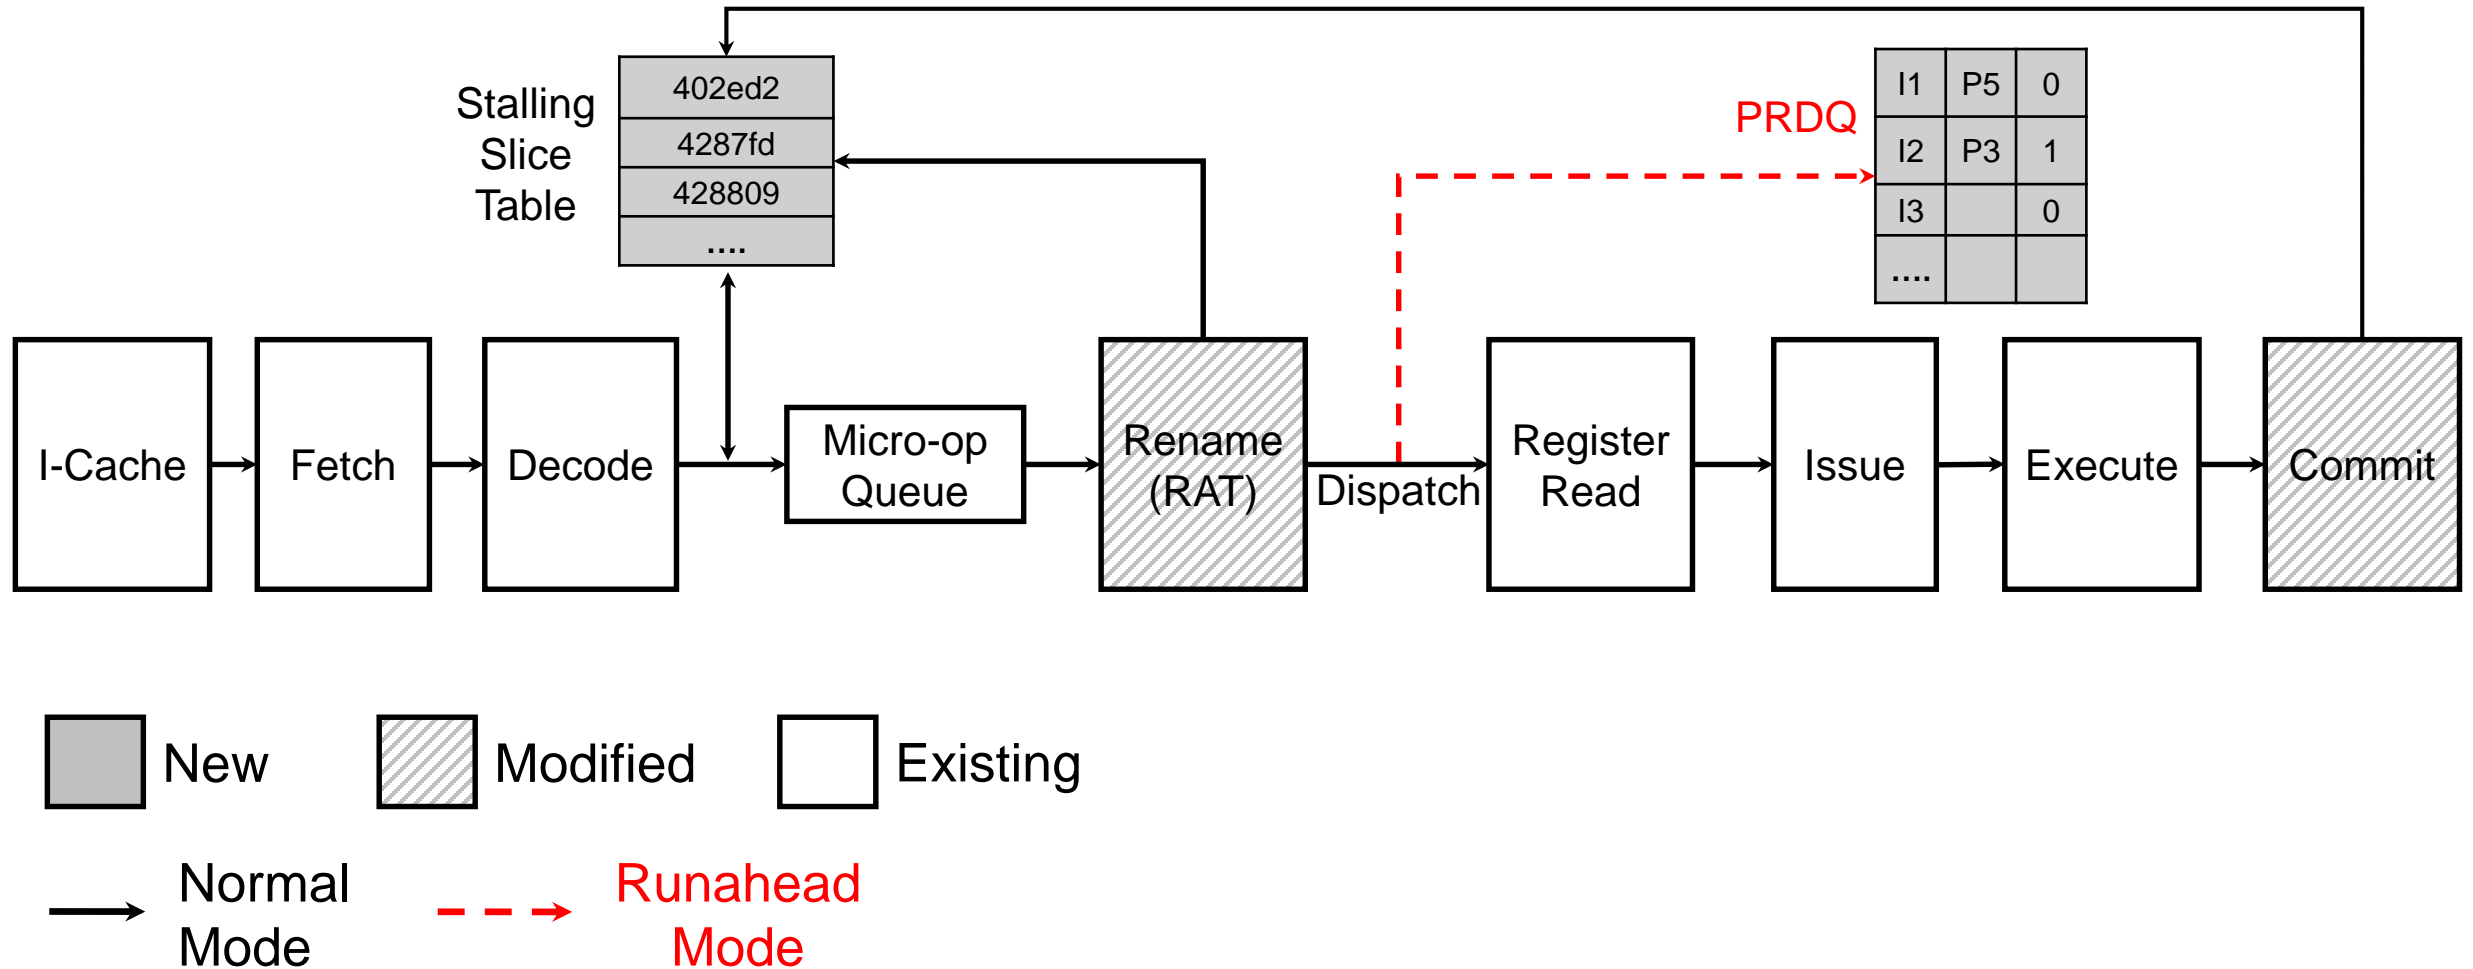
Precise Runahead Execution



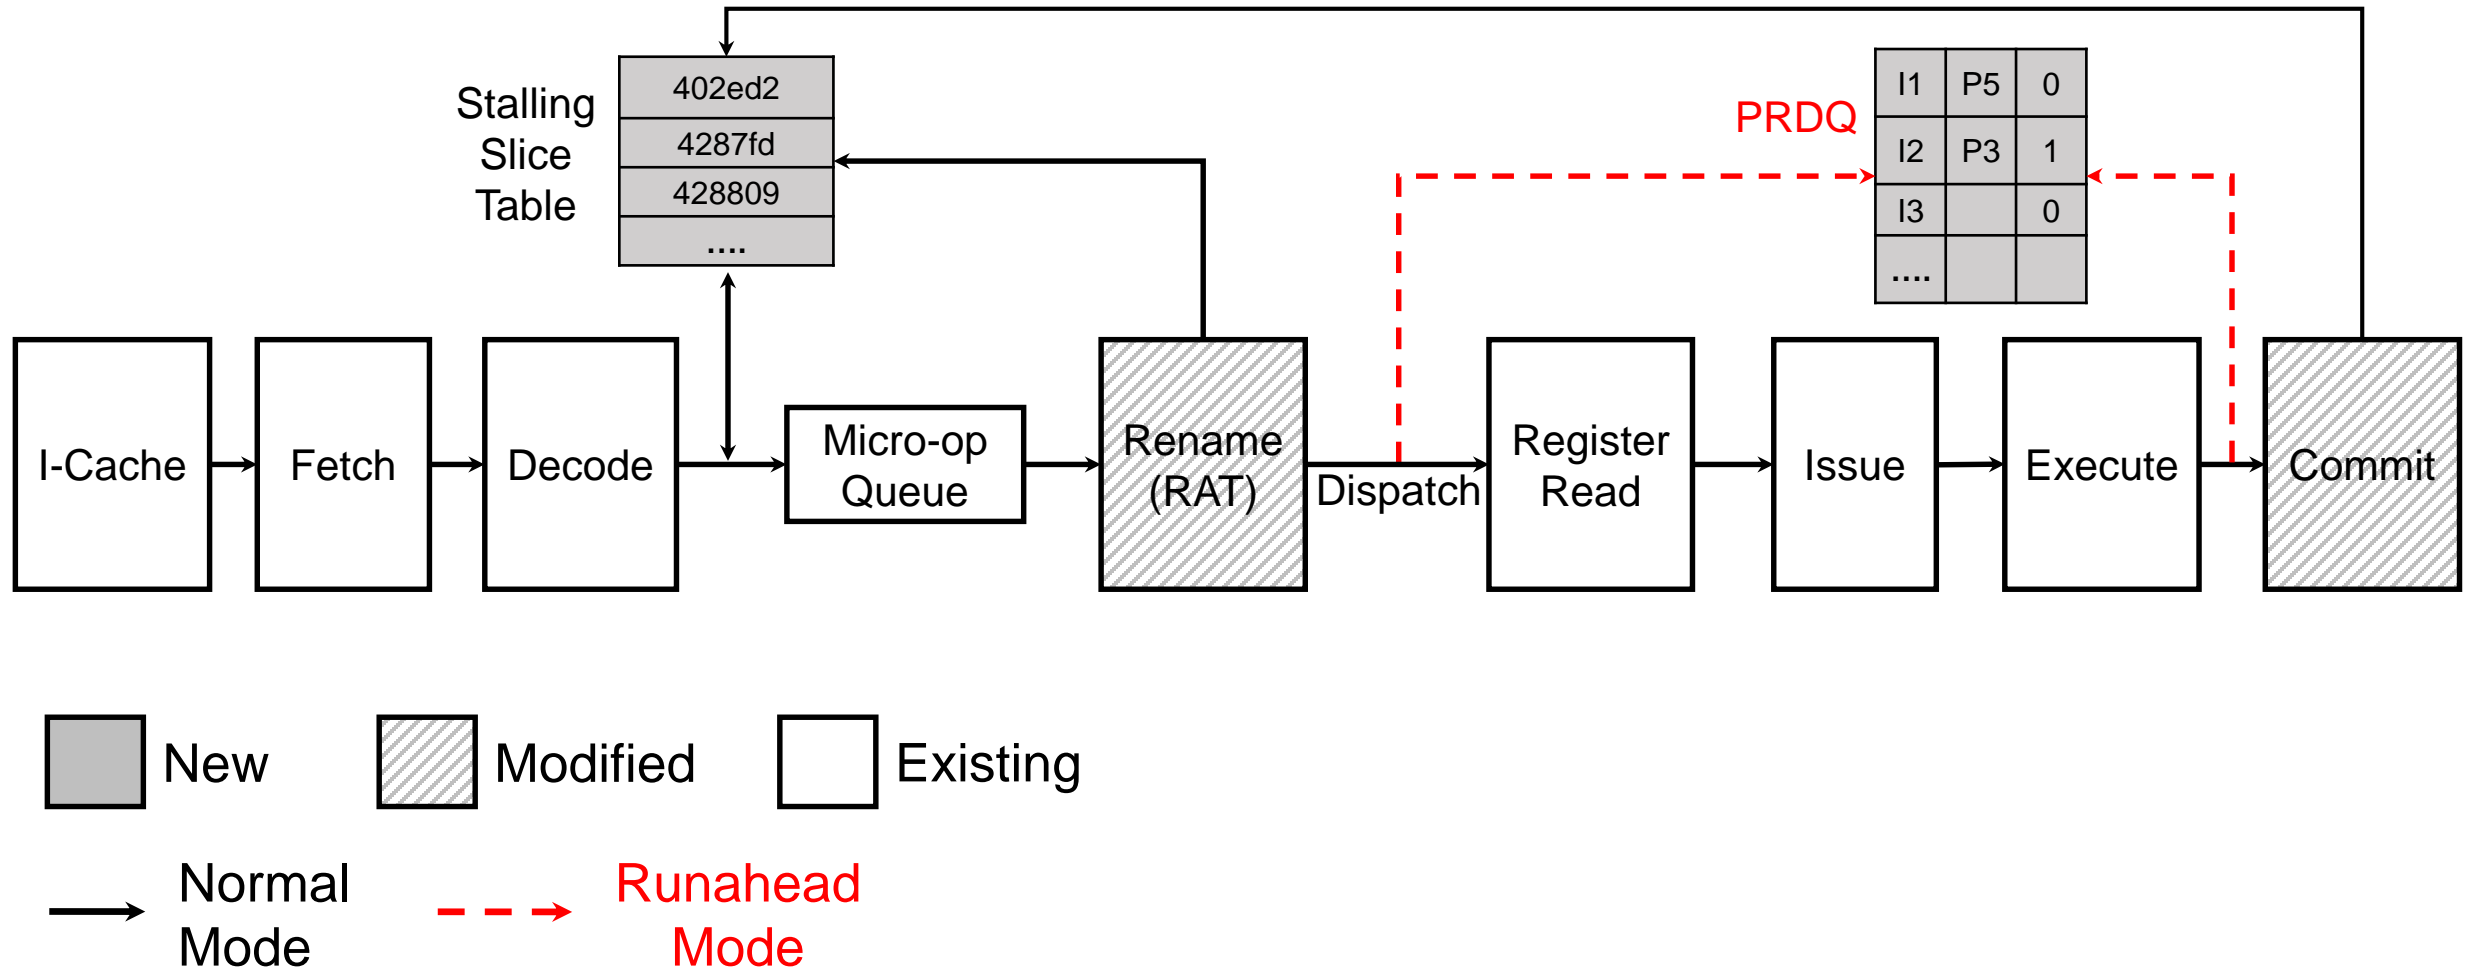
Precise Runahead Execution



Precise Runahead Execution



Precise Runahead Execution



Evaluation

OoO: Baseline out-of-order core

Evaluation

OoO: Baseline out-of-order core

RA: Runahead execution*
-- No short runahead intervals

*[Mutlu et al. ISCA'05]

Evaluation

OoO: Baseline out-of-order core

RA: Runahead execution*

- No short runahead intervals
- No overlapping intervals

*[Mutlu et al. ISCA'05]

Evaluation

OoO: Baseline out-of-order core

RA: Runahead execution*

- No short runahead intervals

- No overlapping intervals

RA-buffer: Runahead buffer**

*[Mutlu et al. ISCA'05]

**[Hashemi et al. MICRO'15]

Evaluation

OoO: Baseline out-of-order core

RA: Runahead execution*

- No short runahead intervals

- No overlapping intervals

RA-buffer: Runahead buffer**

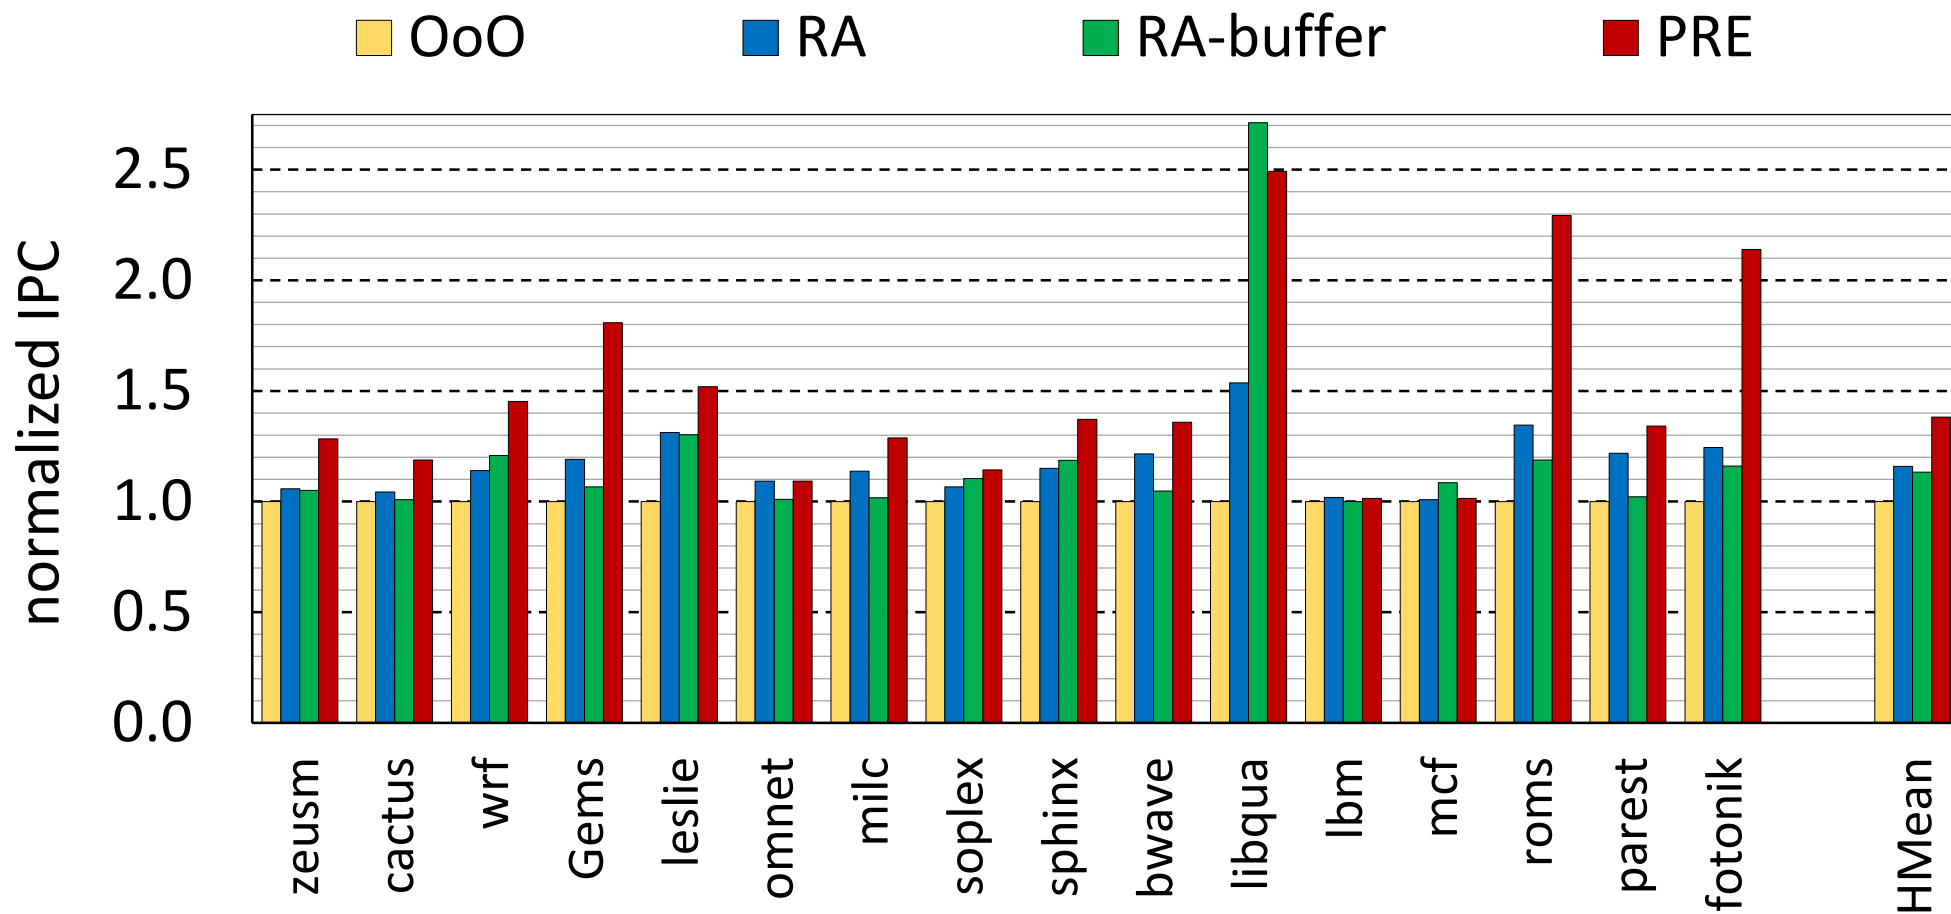
PRE: Precise runahead execution***

*[Mutlu et al. ISCA'05]

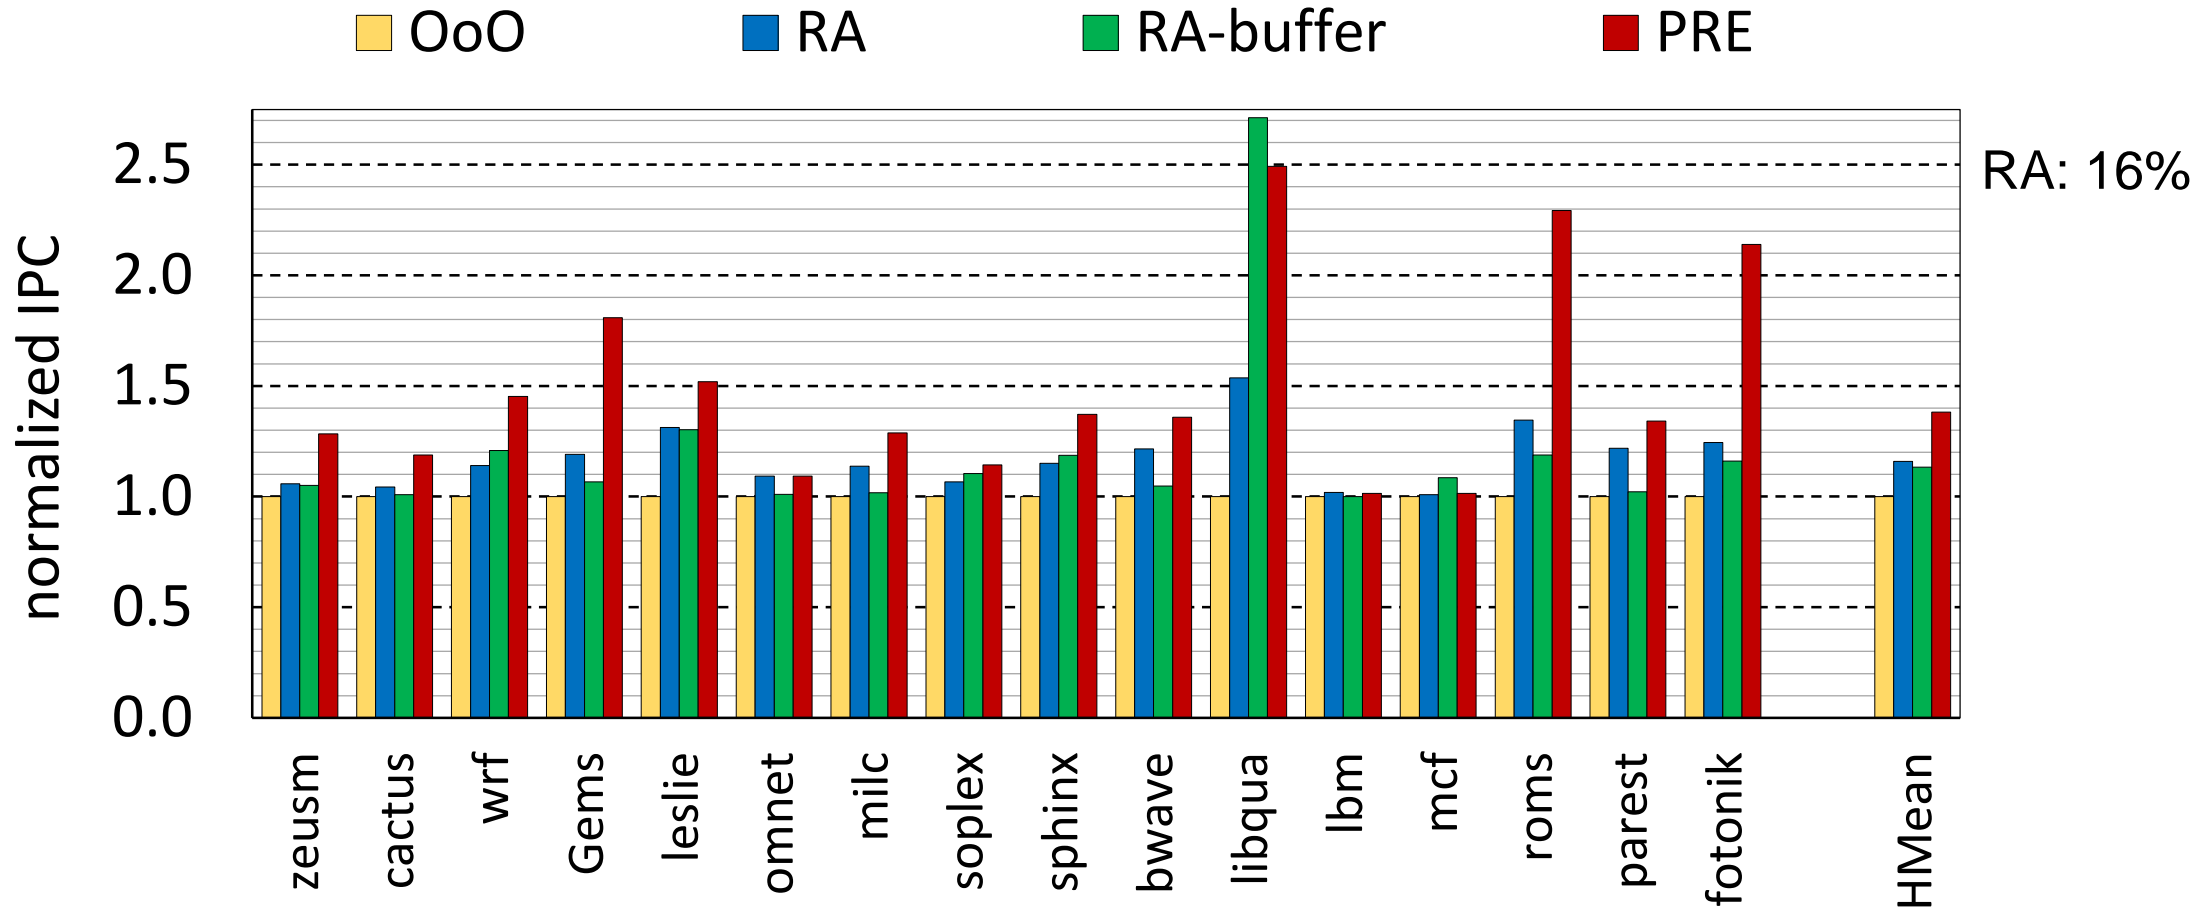
**[Hashemi et al. MICRO'15]

***[Naithani et al. HPCA'20]

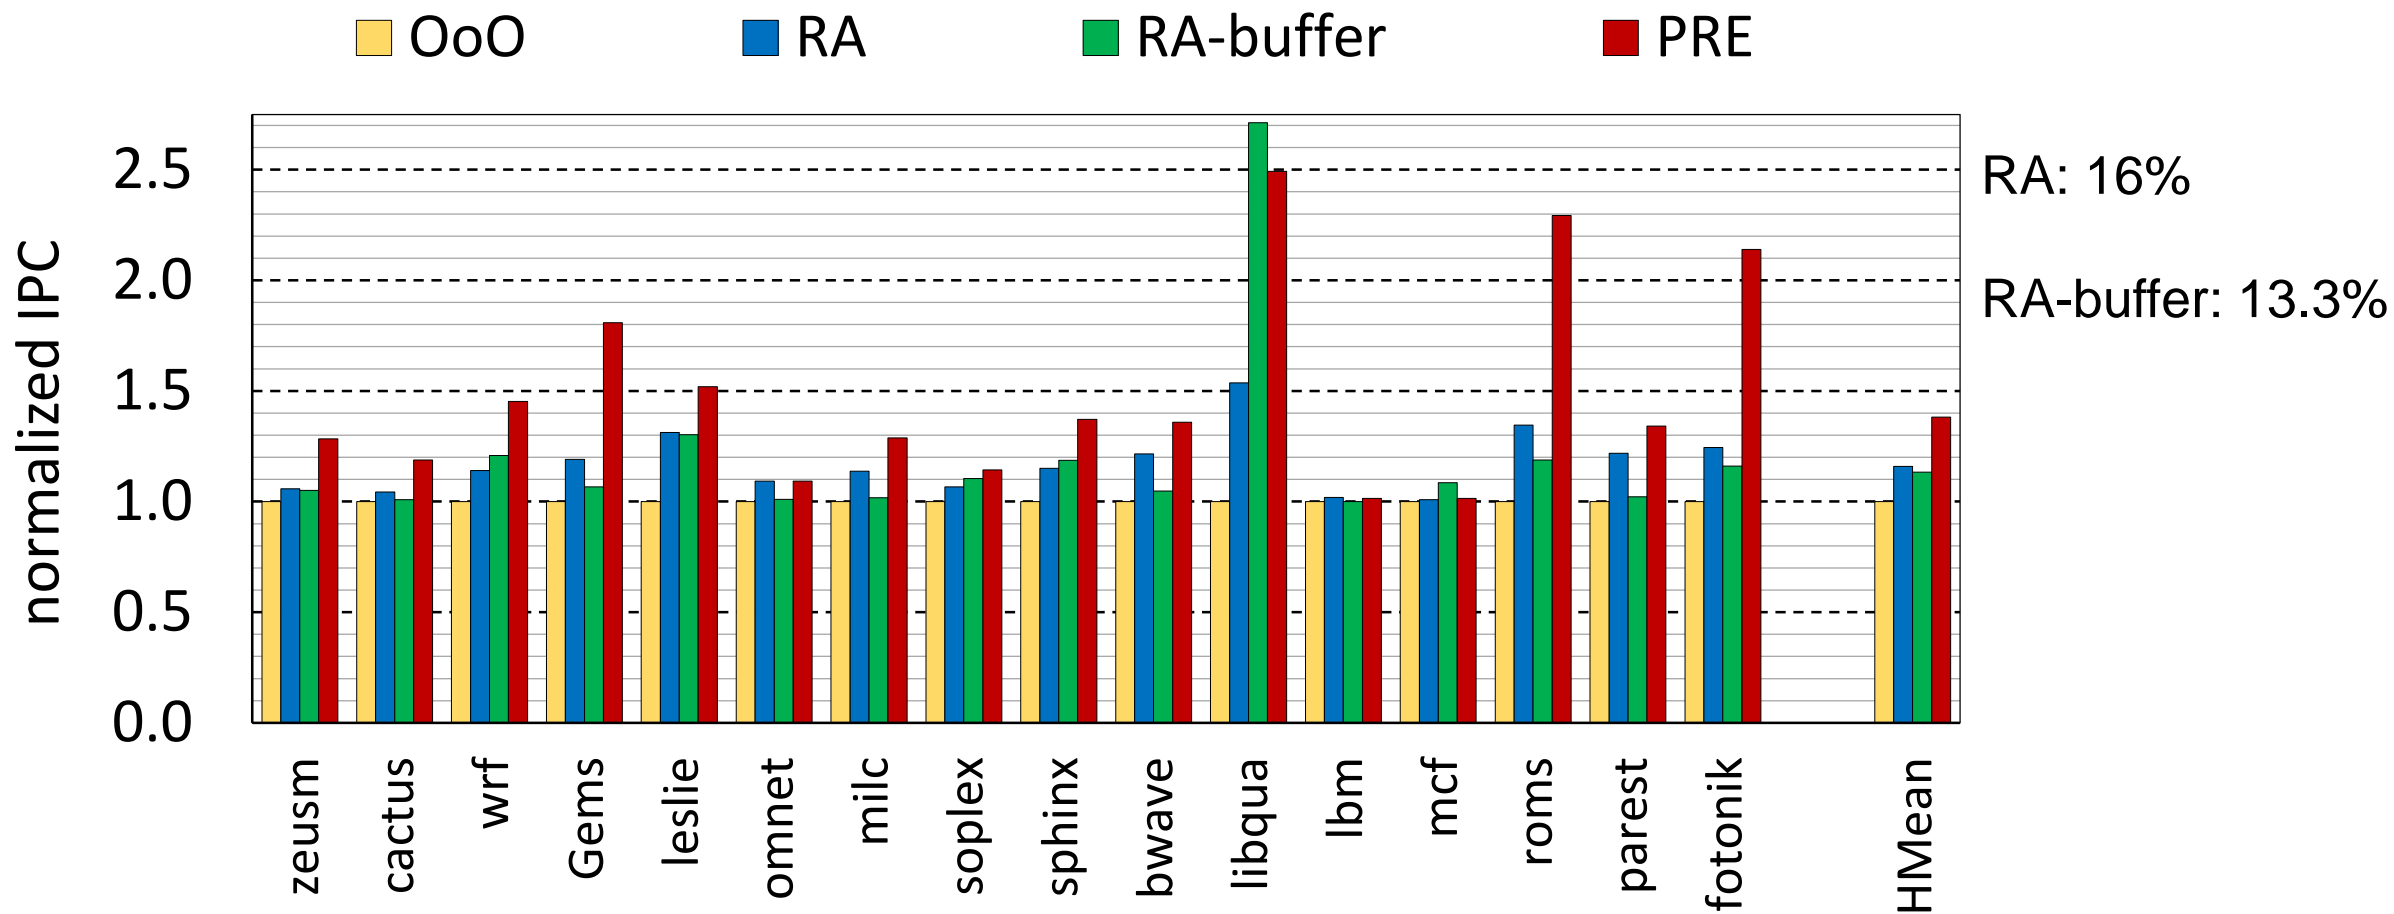
Evaluation -- Performance



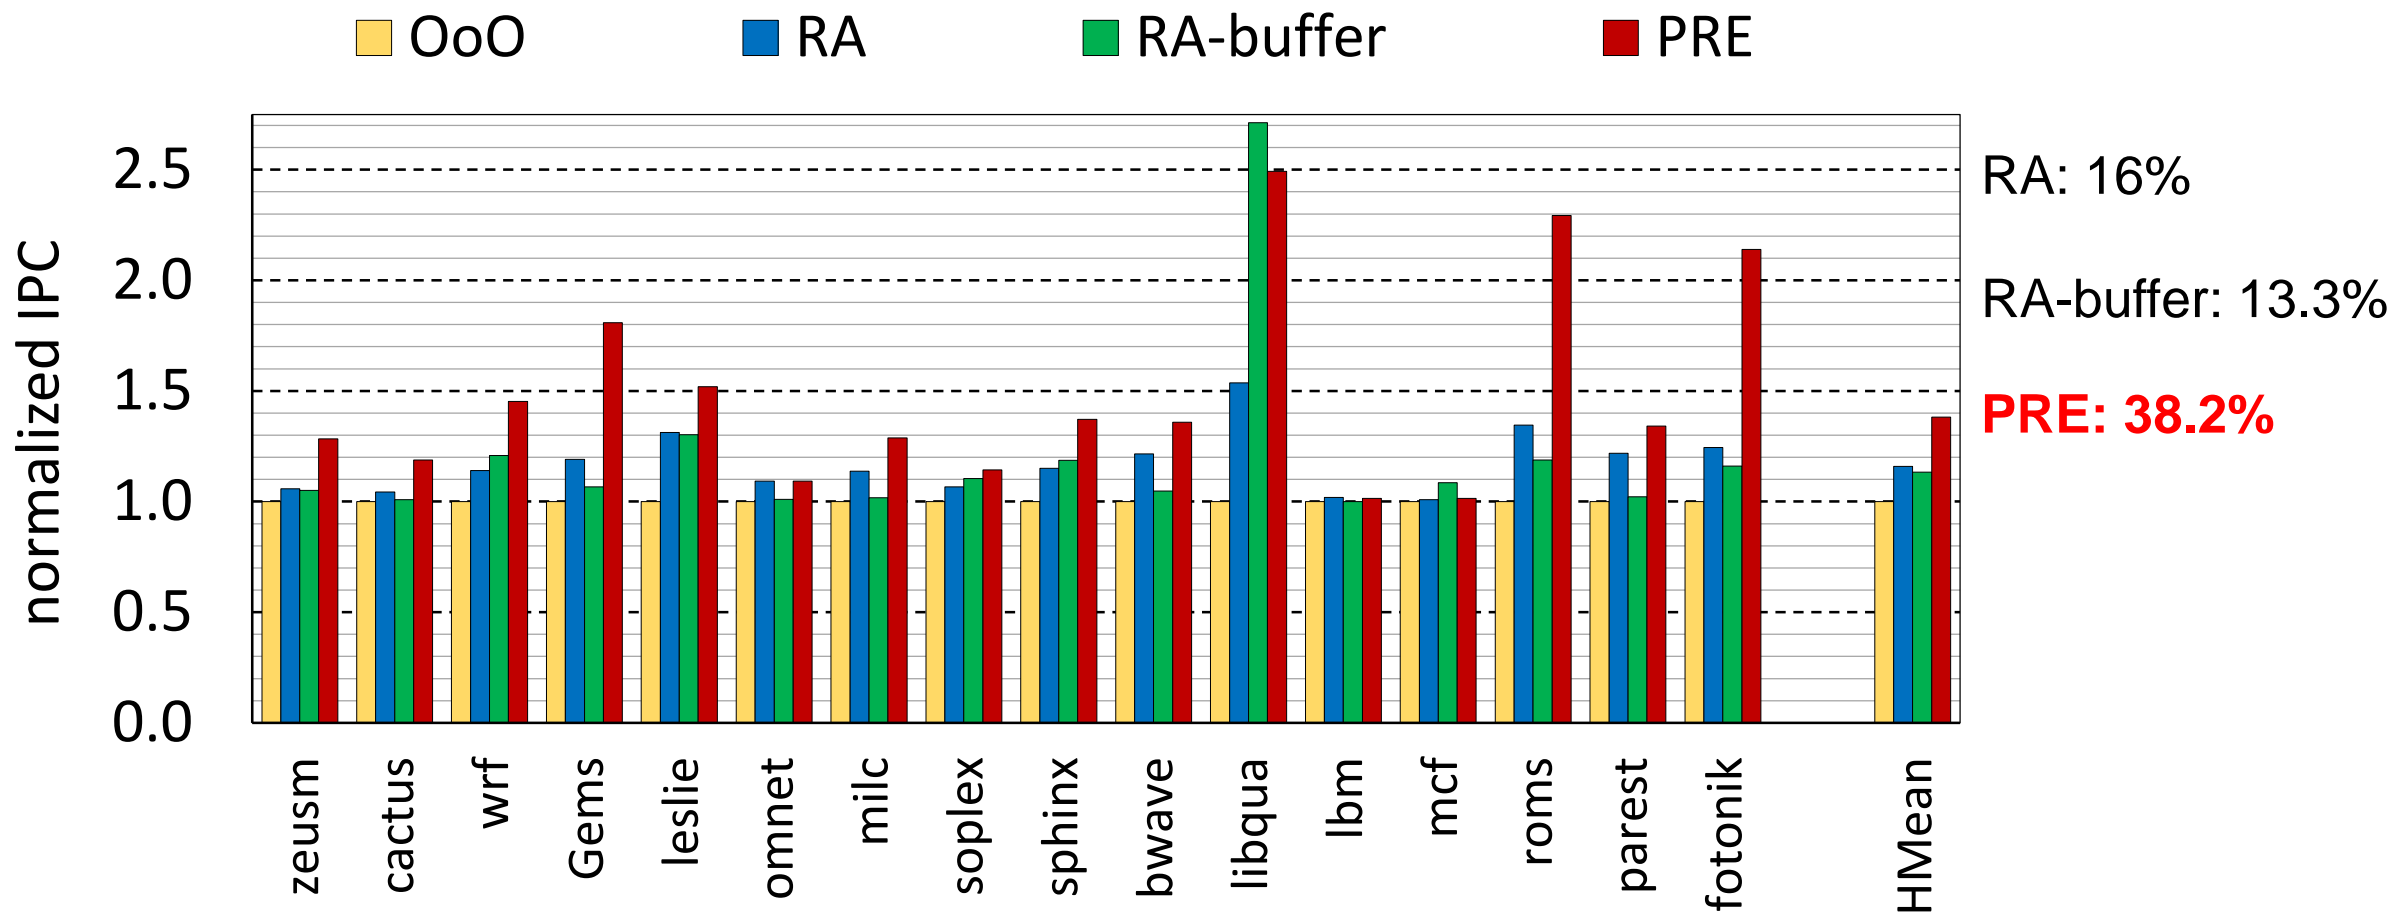
Evaluation -- Performance



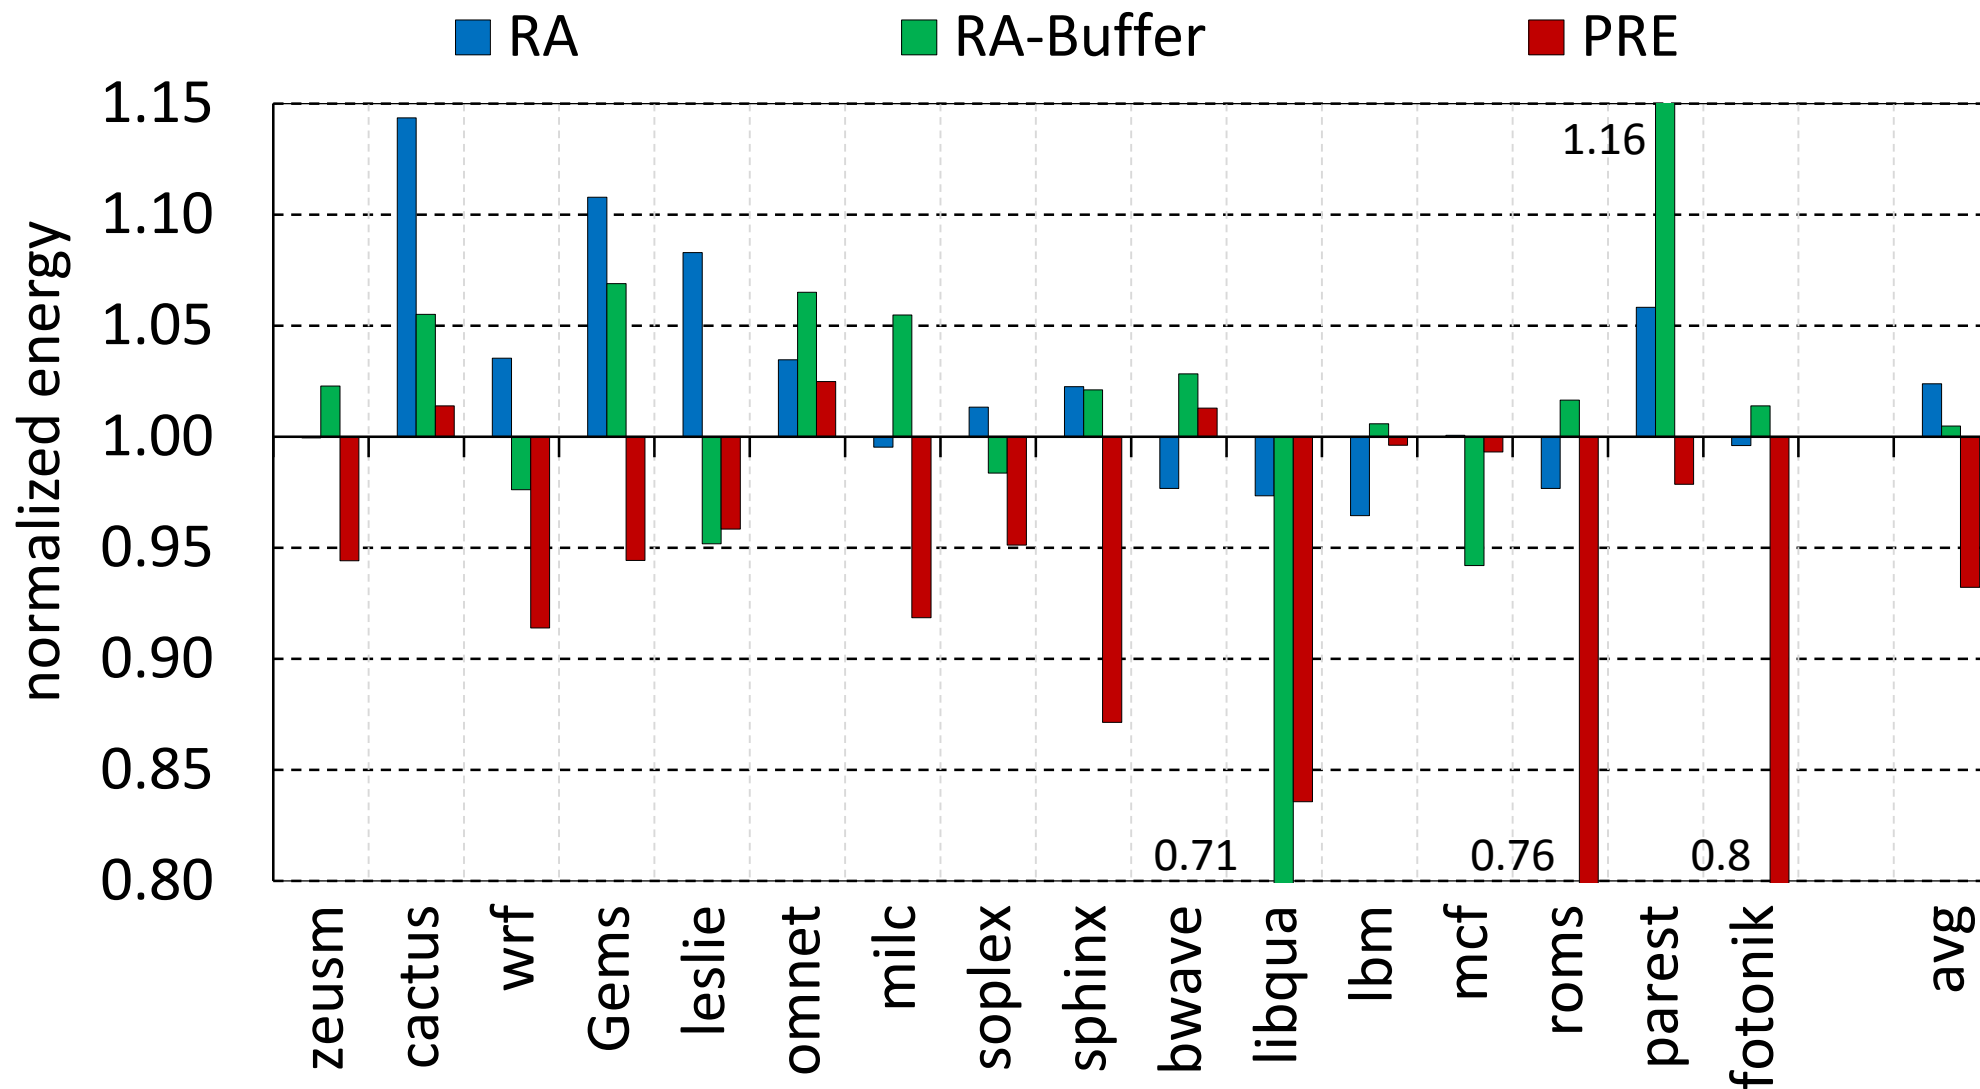
Evaluation -- Performance



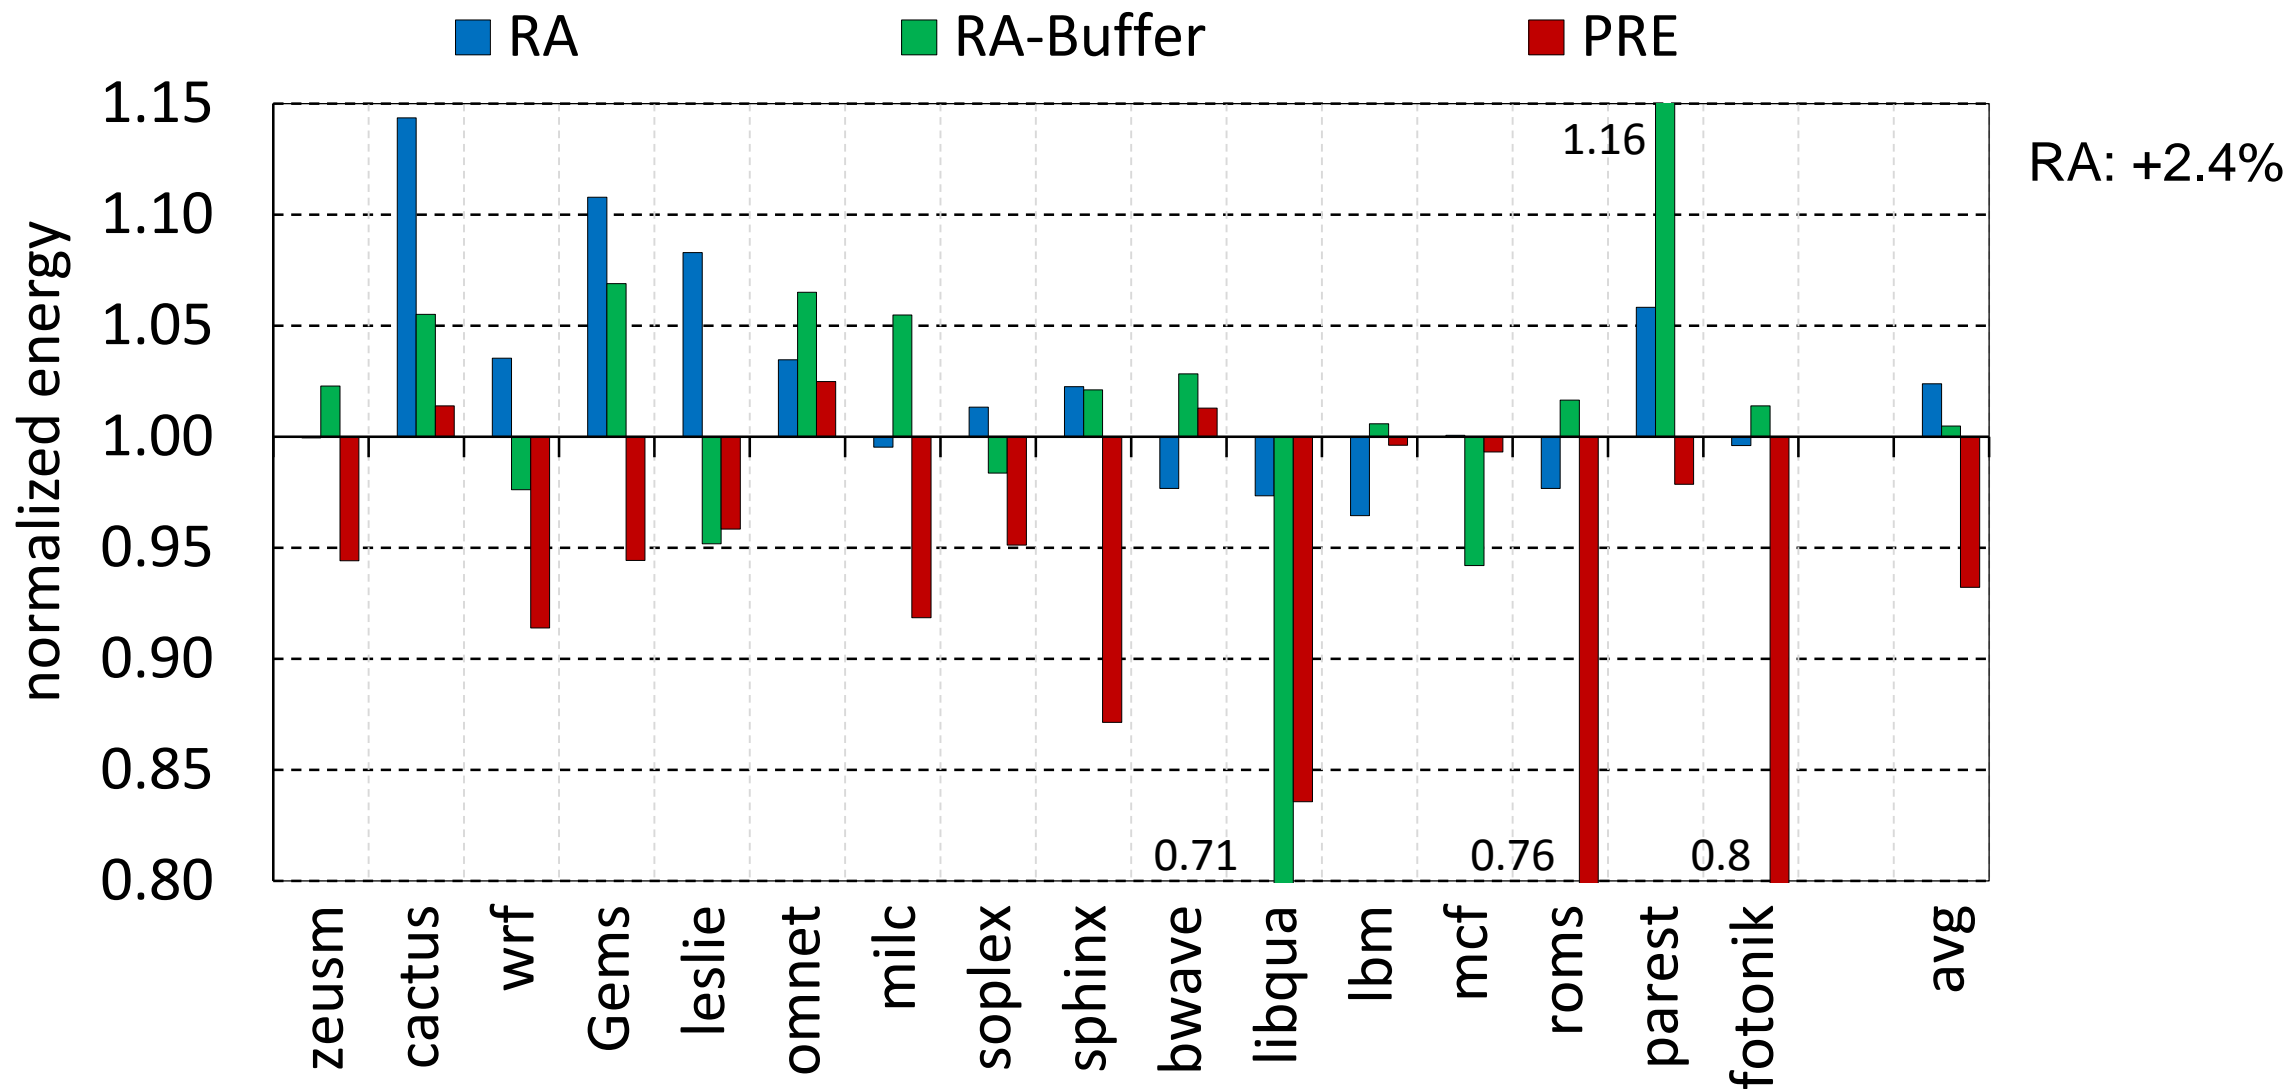
Evaluation -- Performance



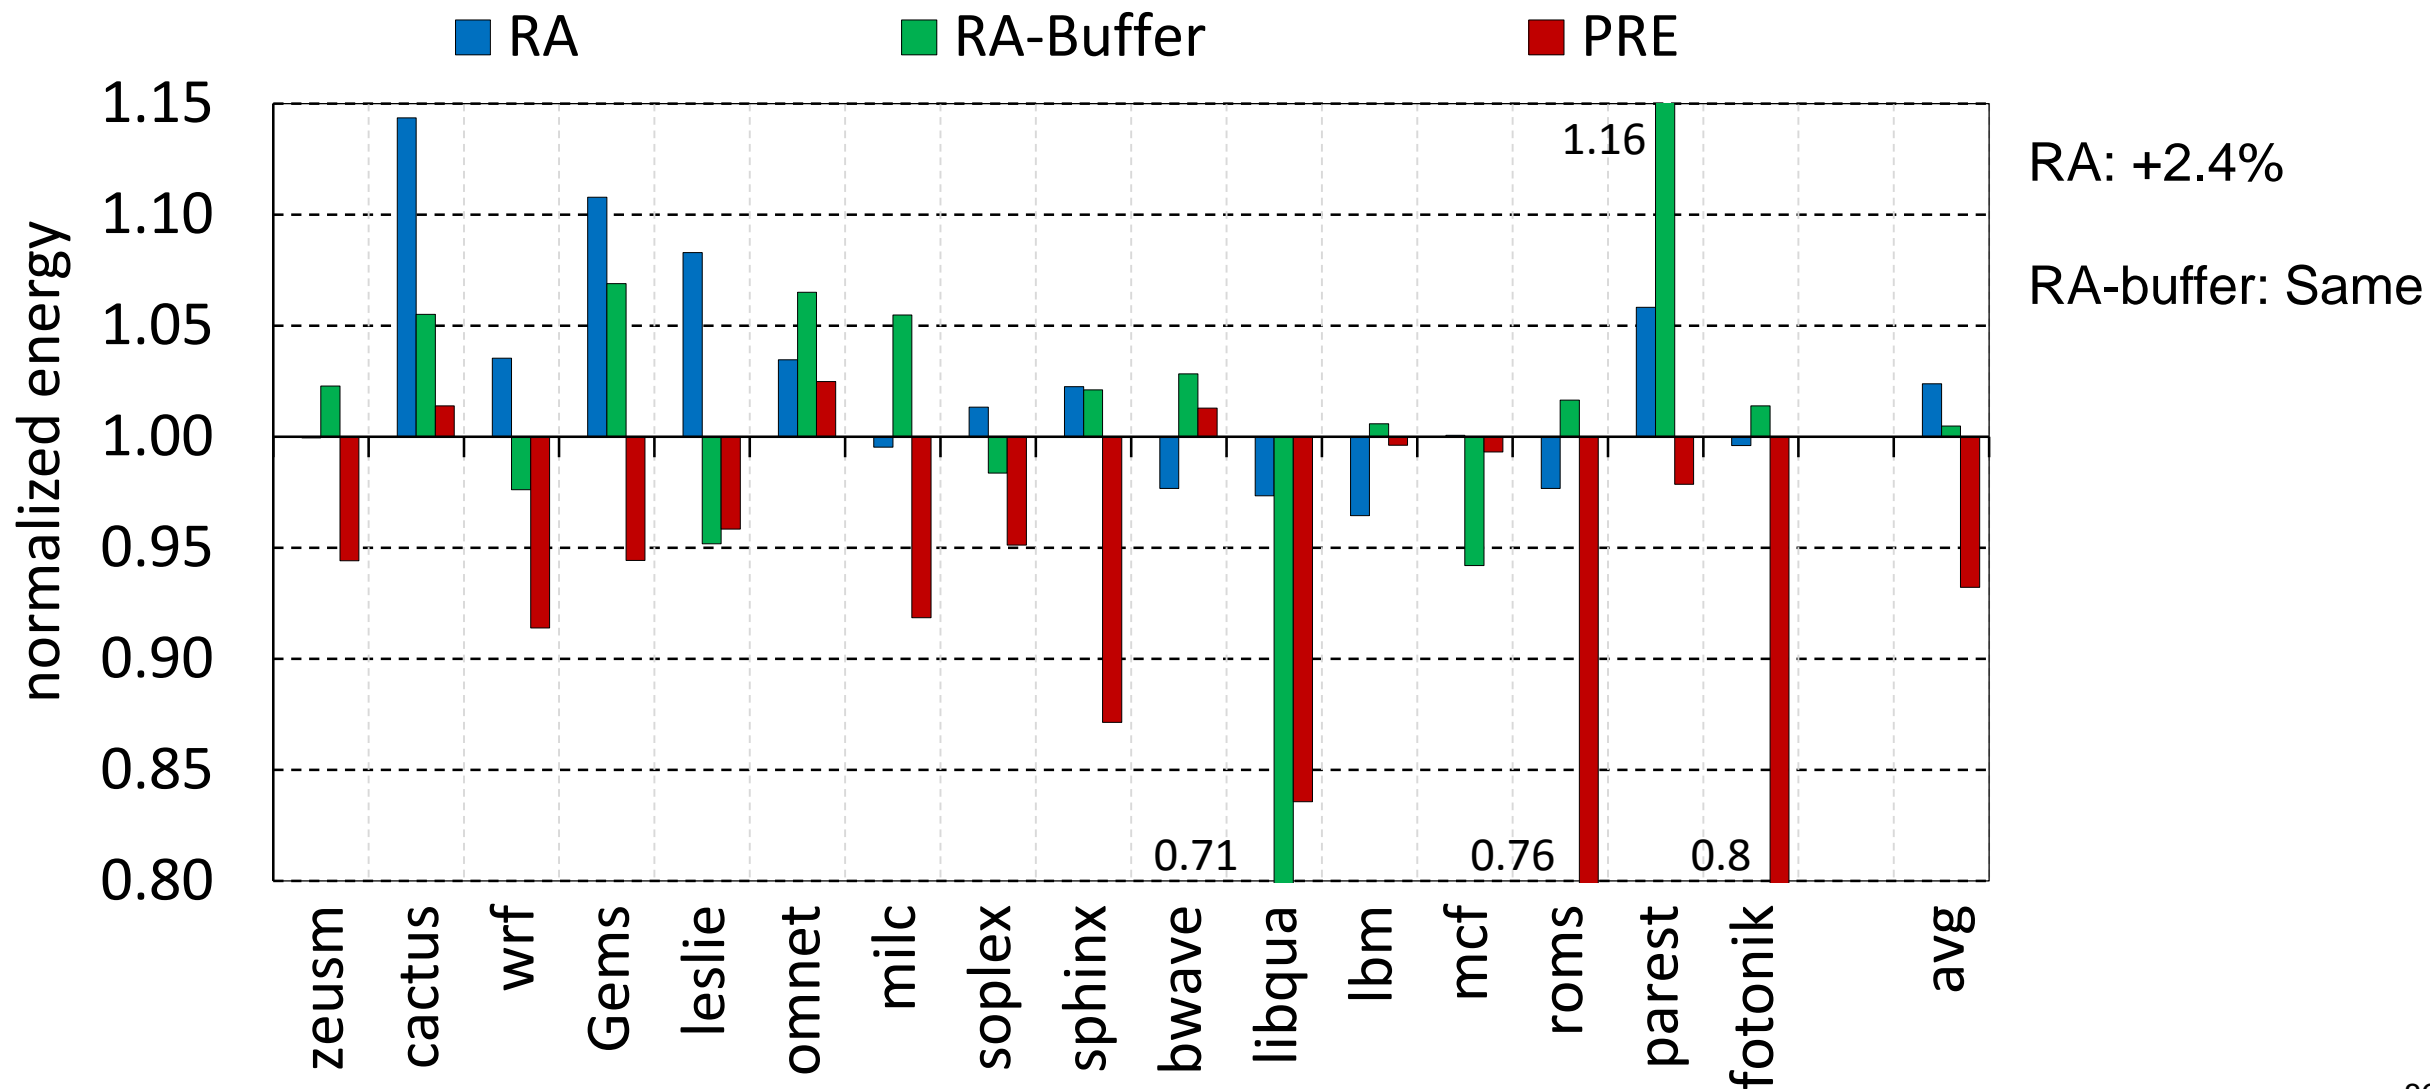
Evaluation -- Energy



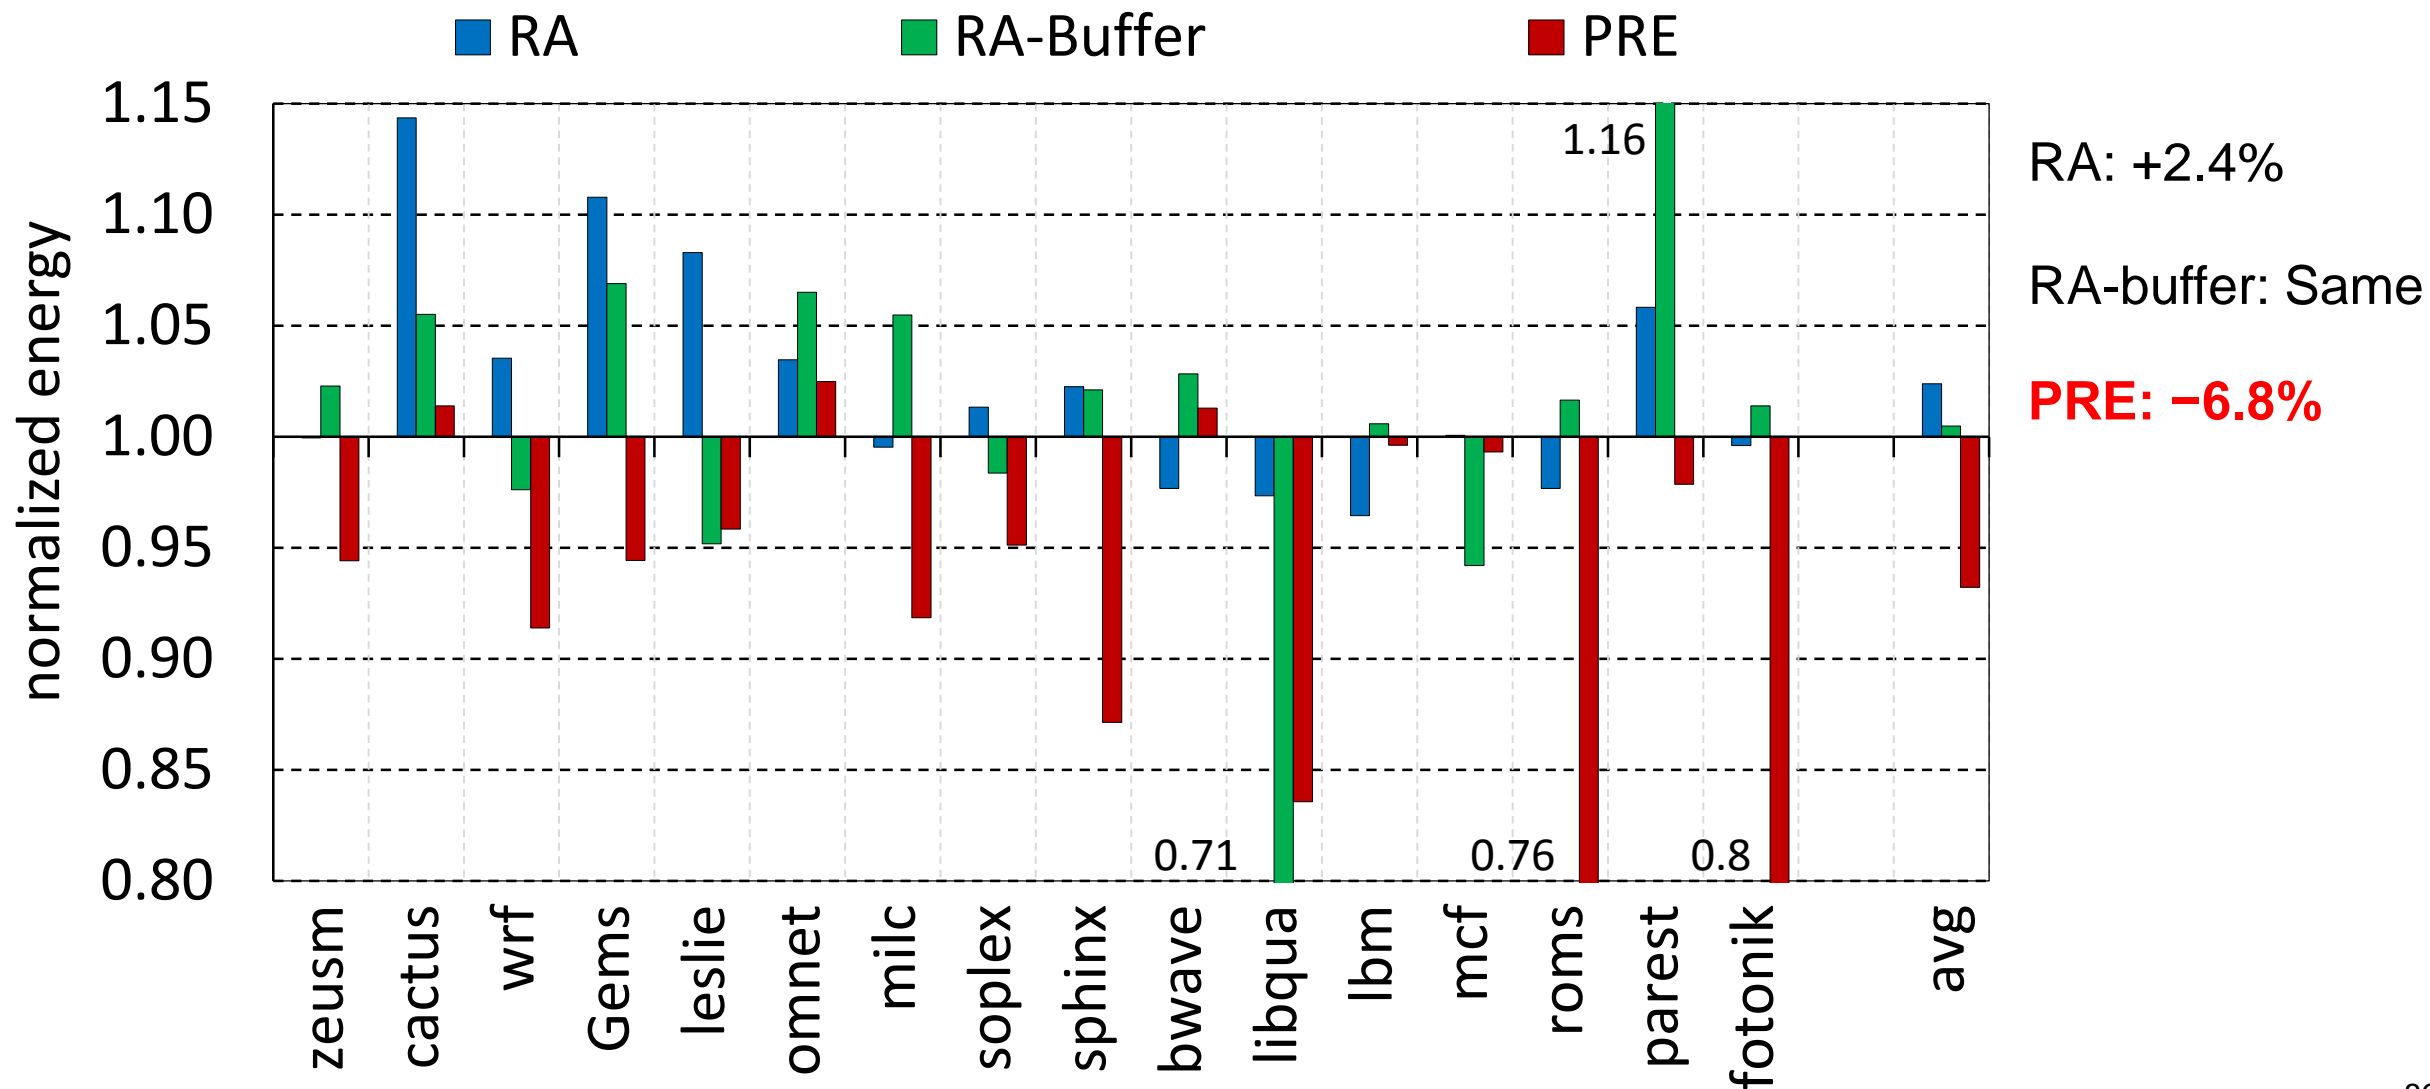
Evaluation -- Energy



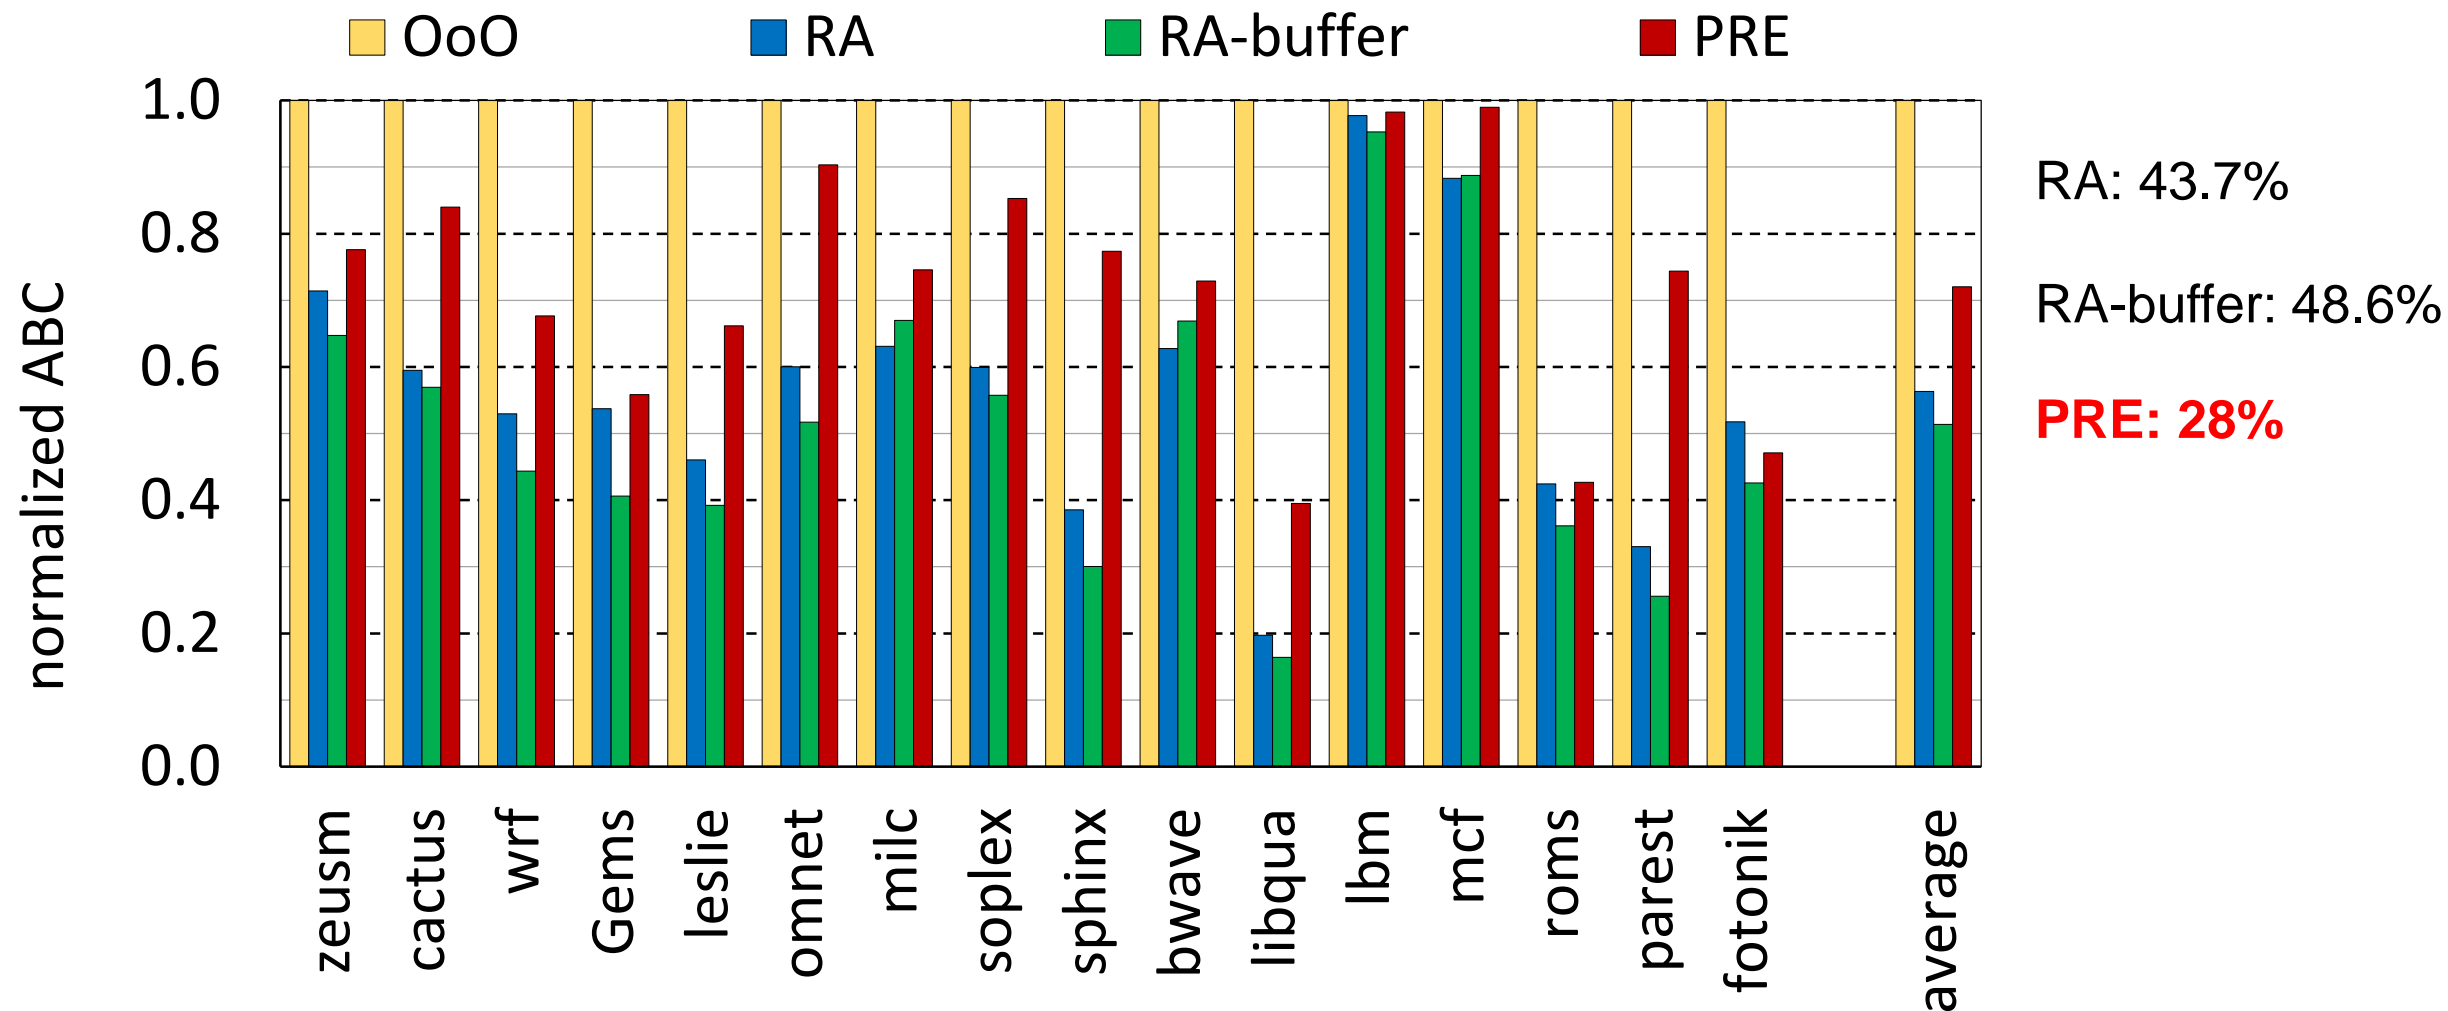
Evaluation -- Energy



Evaluation -- Energy



Evaluation -- Reliability



Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	Precise runahead
Flush ROB	✓	✓	✗
Execute all future instructions	✓	Only one slice	All slices
Performance	High ↑	High ↑	Very high ↑↑
Energy Efficiency	Low ↓	Same	High ↑

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	Precise runahead
Flush ROB	✓	✓	✗
Execute all future instructions	✓	Only one slice	All slices
Performance	High ↑	High ↑	Very high ↑↑
Energy Efficiency	Low ↓	Same	High ↑
Reliability			

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	Precise runahead
Flush ROB	✓	✓	✗
Execute all future instructions	✓	Only one slice	All slices
Performance	High ↑	High ↑	Very high ↑↑
Energy Efficiency	Low ↓	Same	High ↑
Reliability	High ↑		

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	Precise runahead
Flush ROB	✓	✓	✗
Execute all future instructions	✓	Only one slice	All slices
Performance	High ↑	High ↑	Very high ↑↑
Energy Efficiency	Low ↓	Same	High ↑
Reliability	High ↑	High ↑	

Runahead Techniques Relative to OoO Core

	Runahead execution	Runahead buffer	Precise runahead
Flush ROB	✓	✓	✗
Execute all future instructions	✓	Only one slice	All slices
Performance	High ↑	High ↑	Very high ↑↑
Energy Efficiency	Low ↓	Same	High ↑
Reliability	High ↑	High ↑	Medium ↑

Summary of Novel Contributions

Summary of Novel Contributions

1. Reliability-aware scheduling for multiprogram workloads
Exploits the **difference in vulnerability** between core types to improve reliability

Summary of Novel Contributions

1. Reliability-aware scheduling for multiprogram workloads
Exploits the **difference in vulnerability** between core types to improve reliability
2. Dispatch halting for single-threaded workloads
Exploits **invulnerable speculation** to improve reliability under memory accesses

Summary of Novel Contributions

1. Reliability-aware scheduling for multiprogram workloads
Exploits the **difference in vulnerability** between core types to improve reliability
2. Dispatch halting for single-threaded workloads
Exploits **invulnerable speculation** to improve reliability under memory accesses
3. Precise runahead execution for single-threaded workloads
Executes **only useful** future instructions after a full-ROB stall **without flushing the ROB**

Improving Soft Error Reliability in Modern Processors

Ajeya Naithani

Advisor: Prof. Lieven Eeckhout

Doctoral Thesis Defense